
Generative and Embodied AI Security

Attacks, Defenses, and Engineering Verification from Language Models to
World Models

Mingjun Cheng

Vorynel Co.,Ltd

v0.2.0 · 2026-09-26

English edition

Contents

Part One: A Shared Language	6
Chapter 1: From Generated Content to Changing the World	7
Chapter 2 System Structure, Interface Constraints, and the First-Broken Interface	21
Chapter 3 Evaluation, Statistics, and Reproduction Boundaries	33
Part II: Language Models and Agents	47
Chapter 4: Instruction Conflicts, Jailbreaking, and Prompt Injection	48
Chapter 5 Retrieval, Context, and Memory	61
Chapter 6 Tools, Identity, Execution, and Supply Chain	75
Chapter 7 Defense in Depth for Language Models	89
Part III: Image and Video Generation	102
Chapter 8 Visual Generation Pipelines and Security Assets	103
Chapter 9: Data, Models, and the Personalization Supply Chain	118
Chapter 10 Conditions, Sampling, Privacy, and Generation Services	132
Chapter 11 Video Spatiotemporal Safety and the Authenticity Chain	146
Part IV Vision-Language-Action Closed Loop	162
Chapter 12 From Seeing to Acting: Closed-Loop Interfaces of Three Model Types	163
Chapter 13 Attack Propagation in Observation, Reasoning, and Planning	177
Chapter 14: Actions, Tools, and Physical Consequences: From Proposal to Execution	191
Chapter 15 Closed-Loop Defense in Depth and Verification	205
Part V: World Models and Control	218
Chapter 16: The Four Functional Boundaries of World Models	219
Chapter 17 State, Dynamics, and Goal Attacks: How the Imagination Chain Is Hijacked	233
Chapter 18 Runtime Assurance, Recovery, and Falsifiable Testing	246
Part Six: Engineering Closed Loop	259
Chapter 19: Cross-Domain Defense-in-Depth Architecture	260
Chapter 20 From Threat Model to Operating Institutions	278
Appendix	297
Appendix E — Post-cutoff update (2026-08-09 → 2026-09-26)	308

Summary

A model that only generates a passage of text keeps its security problems largely at the level of content. When its outputs enter retrieval, long-term memory, media publishing, software tools, robot actuators and world-model closed loops, errors may then acquire state, identity and real-world capability. This survey is organized around one unified chain of questions. Where does untrusted input first cross a trust boundary? How does an attack propagate from data to state, plans, actions and feedback? How do defenses establish independent, verifiable and recoverable controls outside the model?

The survey is divided into six parts. Part One establishes the capability chain, the first-broken interface, the consequence layers and the evidence grading rules. Part Two turns to large language models and artificial intelligence agents (AI agents; hereinafter "agents"). It covers jailbreaking, prompt injection, retrieval-augmented generation, memory, tools, identity, sandbox and supply chain. Part Three discusses data, components, conditions, sampling, privacy, spatiotemporal state, detection, watermarking and content provenance in image and video generation. Part Four enters the observation-reasoning-action closed loop of vision-language models, vision-language-action models and world action models. Part Five analyzes the state, dynamics, reward, planning and runtime assurance of world models, together with their environment, action and control functions. Part Six assembles cross-domain controls into a regime of design reviews, test matrices, release gates, monitoring and incident response.

Each case is organized along assets, entry points, attacker capability, budget, the first-broken interface, the highest consequence layer, benign utility, cost and the limits of the evidence. Method names and success rates are read back into the system conditions that produced them. The result is a security engineering methodology that can keep working as models, tools and environments change.

Preface

Generative AI security is often pulled between two extremes. One extreme reduces the problem to "whether the model refuses harmful prompts," as if a stronger system instruction were enough to constitute a boundary. The other writes any anomalous output directly as a real-world catastrophe and ignores actuators, permissions, environments and operational controls. Both narratives skip the place where the system actually changes.

Technical teams need a more concrete language. Text on a web page, a trigger in training data, a spatiotemporal pattern in a video, a patch in a robot's field of view, a state deviation in a world model — all of these look different. Yet all may let a low-trust object obtain influence beyond its intended authority. Defense proceeds from six principal relationships. Who can write? Who can read? Who can authorize? Who can execute? Who can observe? Who is responsible for recovery?

This survey therefore begins from the system's operating boundary. Model defenses can reduce the frequency of harmful outputs and known attacks. High-impact systems also need structured information flows, least capability, independent action gates, isolation, resource budgets, unified traces and recoverable operation. A defense-in-depth structure gives each layer a different responsibility. When one layer fails, subsequent controls can still limit the consequences.

The formulas in the survey use the mathematical depth needed to support engineering judgment. Research numbers always appear together with their denominators, protocols and limitations. The "Bringing It into a Real System" section at the end of each chapter places the methods into project scenarios. It explains how to draw a system diagram, record threats and set acceptance gates. The hope is that, when readers close this survey, they will be able to face a new model family or a new agent system and independently answer five questions. These are what to protect, where trust changes, how capability is obtained, how far the evidence extends, and how to recover after failure.

Reading Notes

The survey progresses by capability dependency. For a first reading, it is recommended to complete Chapters 1—3 in order. They define the first-broken interfaces U0—U6, the consequence layers Y0—Y4, the threat record and the four-part metrics used by all subsequent chapters.

Different readers can choose among three paths:

- Model and security researchers: 1—7, 8—11, 12—18, then 19—20.
- Platform and application engineers: 1—7, 19—20, and the vision, embodied, or world-model division according to their business.
- Technical management and governance staff: 1—3, 7, 11, 15, 18—20, with a focus on evidence boundaries, release gates, and recovery responsibility.

The "worked examples" in the chapters demonstrate methods. Where there is no literature citation, they are explicitly hypothetical systems. Research numbers with citations support only their original protocol and do not represent the overall industry incidence rate. Discussions of code, attacks and testing are aimed at defense, authorization assessment and risk understanding. They do not provide unauthorized operational plans against real systems.

References use machine-checkable citation keys, which the print edition expands in full at the back of the book. Standards, closed-source products and rapidly published research change over time. Before actual deployment, versions, policies, licenses and the operating environment should be re-verified.

The Six Minimal Concepts Needed to Read This Survey

This is not a prerequisite syllabus that requires readers to first complete language models, visual generation, robotics and control theory. What cross-branch reading really requires is a stable set of system questions. What objects enter the system? What is retained inside? To whom is the model output handed? At which step is authority obtained? How are consequences observed? How far can the evidence support? The six concepts below give the minimal mental model shared across the survey. When they meet an unfamiliar model name, readers can keep reading along these six sets of relationships. Formulas and implementation details can follow later, as needed.

Models and Systems: A Probability Function Is Not a Complete Product

A model receives an input representation and computes an output according to its parameters and the current context. A large language model (LLM) typically receives a sequence of tokens and predicts subsequent tokens. A visual generation model receives conditions such as text, images, poses or noise, and progressively forms an image or video. An action model may output discrete actions, continuous control values or trajectories. Whatever the output form, the model itself does not automatically obtain permission for databases, email, payment interfaces or actuators.

Outside the model, the system adds parsers, data sources, caches, retrievers, policies, identity services, tools, runtime environments, logs and recovery procedures. How user input is split, under what identity web material enters the context, and which principal issues a candidate call are all system behaviors. So is whether tool results are written into future state. Even if two products use the same model weights, they may have completely different security boundaries, because tool permissions, state writes or runtime configuration differ.

Ordinary question answering provides the simplest positive example. The model generates an explanation, the user reads it and judges for themselves, and the output stops at the information layer. The boundary case is the model generating the same sentence while the runtime layer parses it into a database statement or a mechanical action. The runtime layer executes it when independent authorization is lacking. The appearance of the text does not change, but the consumer and the permissions change the highest reachable consequence. The survey therefore draws the complete system first, and discusses model metrics afterwards.

Training and Inference: Parameter Changes and Runtime State Changes Are Not the Same Thing

Training uses data, an objective function and an optimization process to update parameters. Inference reads the current input and state on a given parameter version and computes an output. Fine-tuning continues training from existing parameters, and may update all of them or only a small part. Low-rank adaptation (LoRA) changes model behavior through smaller trainable matrices, and is a parameter-efficient fine-tuning method. It is not a collective term for all adapters.

At runtime there are also changes that are not written into model parameters. System prompts, retrieved passages, conversation history, key-value caches, long-term memory, environment maps and tool receipts may all change the next output without changing the main weights. Conversely, a new adapter or decoder can change the result when the input text is exactly the same. A security investigation that asks only "whether the model was retrained" will miss three paths: the loading combination, the context and the state.

A normal software upgrade freezes the joint version of the base model, adapters, encoders, configuration and dependencies, and compares behavior before and after the change. A hash value is a fixed-length value computed by a hash function from an input. Under a fixed algorithm and trusted baseline, a hash match supports "the current bits are consistent with the baseline." It does not by itself prove that the provenance is trustworthy or the behavior safe. The boundary case is a main model file whose hash is unchanged while the runtime quietly loads a new LoRA, template or plugin. Looking only at the main weights would wrongly judge the actual system as unchanged. The artifact closure and change impact analysis discussed later address exactly this kind of object misalignment.

Representation and State: How Encoded Information Persists Across Steps

Raw text does not enter a neural network directly. Tokenization first turns text into a sequence of tokens. Here a "token" is a model sequence unit, not the same object as an access token in an identity system. The Transformer architecture uses operations such as attention mechanisms. These operations let the current position update according to other representations in the context. The architecture explains how pieces of information influence one another within a single inference pass. It does not automatically specify which message is more trustworthy or which principal has more authority.

An embedding maps text, images or other objects into continuous vectors that are convenient to compare. Vector proximity indicates some kind of similarity under a given model. It does not mean that two objects come from the same source, have the same permissions or are equally true. Latent variables, action tokens and map coordinates are also representations: they compress certain information while possibly losing details or mixing different semantics.

When a representation persists across steps, sessions or tasks and influences future computation, it becomes runtime state. The context window covers only the range of tokens directly visible to the current inference pass. Long-term memory, retrieval indexes, caches and the latent state of a world model can persist outside the window. The normal example is a maintenance assistant writing approved equipment manuals into a knowledge base by tenant and validity period. The boundary case is imperative text on a web page being automatically saved as long-term memory, so that the next user is affected without visiting the web page again.

Open Loop and Closed Loop: Who Consumes the Output Determines How Far the Risk Travels

An open-loop system produces an output. It does not automatically send the real-world execution result back into the next round. An image generator that exports a poster for a human to choose from usually belongs to an open-loop path. So does an analysis tool that generates a recommendation for an engineer to review. Open loop does not mean risk-free. Privacy leakage, copyright disputes, deceptive dissemination and erroneous decisions can still occur. It only means that the system does not keep advancing actions through its own feedback.

A closed-loop system hands its output to a planner, controller, tool or actuator, and then writes the execution result back into the system as a new observation. A vision-language-action model (VLA) can map observations and

language tasks into action representations. A world model (WM) can maintain internal state and predict the future under action conditions. If a planner uses that future to compare actions, or if the prediction enters feedback control, errors may then accumulate across multiple rounds of state updates.

A normal closed loop confines each round of actions to a short duration and a safe envelope. It reads timestamped independent observations after execution, and only then decides the next step. The boundary example is a system that generates a very long action chunk in one go and keeps executing it. Even when new observations already show that the environment has changed, the old plan still advances, because it lacks a deadline and replanning. The "feedback" discussed throughout this survey is both a resource for error correction and a channel through which attacks, privacy and persistence enter the next round.

Probabilistic Output and Real-World Authority: A High Score Does Not Sign Off for Any Control

Model outputs usually contain probabilities, scores, candidate sequences or continuous vectors. These numerical values describe the model's computation results under given parameters, inputs and objectives. High confidence does not equal factual correctness. Low uncertainty does not equal being in a safe state. A candidate action ranked first does not equal that it has already been approved for execution. A model may be very confident on out-of-distribution (OOD) inputs, and may also stably choose the wrong object because the objective function lacks real-world constraints. Here "out-of-distribution" must be judged relative to a specified training or calibration distribution.

Real-world authority is issued by trusted controls in the system. Protocols such as the Model Context Protocol (MCP) can standardize how prompts, resources and tools are exchanged. An access control list (ACL) can restrict a principal's access to data objects. The business layer still needs to verify the current user's objective, recipient, file, amount, coordinates, time and consequences. A correctly formed protocol only proves that the request can be parsed. A valid token only proves that some principal holds a segment of authorization. Neither of the two makes judgments on behalf of user intent and security constraints.

A normal example is a research assistant proposing to call a read-only search tool. The capability gate confirms the task, data use, query scope and duration, and then issues a short-lived capability token. The boundary example is the same assistant reading a command from a web page and then requesting to publish a software package. Even if the call structure is entirely legitimate, the action object and the highest consequence have already changed. The permission layer should reject the request or escalate it for human confirmation. It should not let the model's own "plausibility judgment" become the final authorization.

Metrics and Assurance: Numbers Are Only Usable When They Carry a Denominator and an Evidence Layer

An attack success rate (ASR) must settle four things. What counts as one statistical unit, and which units enter the denominator? Which observation counts as success? How much knowledge, query, time, write or physical access budget does the attacker have? Delete the samples that timed out, hit interface errors or were blocked by a defense line, and the denominator changes. Treat model output, dangerous plans, simulated execution and real-world consequences alike as "success", and different endpoints compress into one percentage. No decision can be made on that percentage.

This survey reports attack effectiveness, residual risk, benign utility and operational cost side by side. A security control can lower the attack rate and still make legitimate tasks fail in large numbers. Reporting only the former is then not enough. A detector can show very good offline curves and still produce an alert volume that cannot be handled every day. Such a detector likewise cannot go directly into production. A team can judge what a number means in the current system through intervals, paired comparisons, base rates, latency and recovery time. Those do not replace a clear statistical unit and causal path.

Evidence has levels too. Static inspection reads code, configuration and artifacts, but does not run the target pipeline. Mechanism-level runs do execute local modules, inside a security fixture. Realism of consumers and environments climbs through simulated execution, closed-loop simulation, end-to-end reproduction and field incidents. Lower-level evidence is not without value — it can find interface and implementation errors. What it

cannot do is extrapolate automatically to higher levels. Runtime assurance (RTA) likewise requires writing out trusted states, monitors, switching logic, safe controllers and failure assumptions. It cannot settle for treating an uncertainty score as a guarantee.

A normal report says so outright: "in the target version, under a given attack budget, and in isolated simulation, how many independent tasks reached which consequence layer." Failed runs and unknowns are retained alongside. A boundary report offers only "defense rate 95%." It carries no denominator, no model version, no scorer, no control state and no highest consequence layer. The first kind can enter release decisions. The second is at most a lead, and it needs supporting evidence. With these six minimal concepts in hand, a reader need not become a full-stack expert first. The system causal chain that this survey really cares about stays trackable in every chapter.

When Encountering an Unfamiliar Formula, Read the Objects Rather Than the Derivation

The formulas in this survey mainly do three kinds of work. The first names objects. A tuple, for instance, puts assets, inputs, states, models, execution gates, feedback and recovery into one system diagram. The second expresses change, for example how the previous moment, the current observation and the action update the state. The third specifies measurement, for example the numerator, the denominator and the conditions of an attack success rate. When reading, translate each variable back into an observable object. Then ask what the equals sign, the conditions and the summation each connect. A formula has not yet landed on a verifiable system if one of its variables matches no log, file, interface or event.

No probability in these formulas has to be imagined as a measure of how much the model "believes." A conditional probability states the frequency, or the modeled relationship, with which a later event occurs once an earlier event has occurred. The conditional chain of a dangerous consequence offers one example. It separates input being accepted, state being changed, plan being selected, capability being granted, and environment being affected. Its use is to locate controls. It is not to assume that the stages are independent and then compute a seemingly precise total risk. One shared model, parser or identity root can make several stages fail at the same time. Dependencies of that kind must be stated separately, in the text and in the figures.

Optimization formulas likewise begin with the question of who can change what. Different permissions govern the pixels, tokens, training samples, state vectors, rewards or candidate trajectories an attacker can optimize. Nor is a frozen model the same object as an updatable adapter. An objective function states what the attack or the training is searching for, and nothing more. It does not prove that the actual budget suffices to find a solution. Nor does it prove that a found internal deviation can cross the boundary of authorization and execution. After each formula, the "assumptions" and "failure conditions" are there precisely to prevent this leap.

A three-step reading method works here. Start with the topic sentence of the text. It names the problem the formula is meant to solve. Then follow the figure or variable table to the inputs, states and outputs. Finally check whether the conclusion stops at the model, the simulation or the real environment. Readers who care only about architecture and governance may skip the algebraic details. They should not skip the consumers and units of the variables. Readers who want to reproduce should go further and check the data, code, parameters, randomness and scorer. Both reading depths share the same evidence boundary.

How to Use Terminology, Figures, and Cases

A term's first appearance takes the form "Chinese standard name (English Full Name, abbreviation)." Registration is separate for national standards, official specifications, words commonly used in original papers, translations customary in the field, and the operational terms of this survey. Their normative levels are distinguished clearly. The analytical language of this survey includes first-broken interface, capability gate, artifact closure, highest observation layer and highest reachable layer. World action model, environment world model and world control model receive working definitions that follow the actual consumer. They help compare systems. They do not oblige other authors to adopt the same category names.

Figures provide relational navigation. Rectangles usually denote inputs or artifacts. Rounded boxes stand for models or transformations, prominent borders for authorization and execution, and return-bent arrows for

feedback. Reading a figure starts with the solid lines, which trace the normal information flow. State retention, attack influence and recovery loops come next. The lead-in paragraph before each figure names the path to observe. The paragraph after it says what the figure supports, and what it does not prove. The graphics are the authors' own mechanism illustrations; they do not copy experimental figures from papers. An abstract arrow marks nothing more than a relationship to be verified. Whether a real deployment realizes it is still proven by configuration, logs and execution receipts.

The cases here divide into research cases and engineering worked examples. Research cases carry citations. The original paper's model, data, attacker capability and environment limit their numbers and conclusions. Engineering worked examples demonstrate methods. Any system, role or numerical value in them without a citation is an explicit assumption. The fields and check order of a worked example can be transferred to a reader's own project. Its illustrative values cannot serve as industry baselines. Cases are compared numerically only when their statistical unit, protocol and consequence layer are compatible. Otherwise the report keeps side-by-side evidence and mechanism differences.

The "Bringing It into a Real System" section at the end of each chapter is not a homework question. It is a narrative that begins at the field entry point. It shows which receipt the team examines first, and how version and consumer are judged. It also shows where extrapolation stops, and which control should next change the system state for real. Retell the objects, permissions, evidence and recovery along this process, and the reader has already grasped the engineering throughline of the chapter. Algorithm details can be pursued further, according to responsibilities.

An unfamiliar abbreviation may still turn up in a chapter. Start with the Chinese-English glossary at the back of this survey, then return to the principles groundwork that opens the chapter. The glossary fixes names and neighboring concepts. The principles groundwork explains how these objects act in a system. The main-text cases provide evidence and failure paths. The three carry different responsibilities, so no single definition has to bear the three tasks of naming, teaching and proving at once.

\newpage

Part One: A Shared Language

Guide to This Part

A generative AI system may start with a single question and answer. It may also reach all the way to databases, media publishing, software tools and physical actuators. Security discussions, faced with implementations this different, most easily start from the name of an attack or a model. One group asks first whether it is jailbreaking. Another asks first which kind of model it is. Part One takes a different starting point. It establishes, to begin with, which objects exist in the system, where trust changes, who consumes the outputs, and what the evidence actually observed. Draw these relationships clearly, and new technology names cannot deprive the analysis of its coordinates.

Chapter 1 builds the capability chain. It runs from external input, context, model output, candidate plans, authorization and execution through to feedback. The key change here is not a "smarter" model. It is that probabilistic outputs progressively acquire state, identity and real-world capability. Chapter 2 expands that chain into system objects and the U0—U6 trust interfaces. It distinguishes where an attack first crosses a boundary, what enables later propagation, and which control ultimately releases an action. Chapter 3 then answers how to measure. Sampling units, denominators, attack budgets, scoring rules and evidence layers jointly produce a percentage.

The three chapters deal with the system, the path and the evidence, in that order. Readers do not need to master every attack algorithm first. They need only ask three questions repeatedly. Does the sentence at hand describe a model output, a system plan, or an external consequence? Where is the interface at which an attack first obtains influence beyond its allowance? Up to which layer of conclusions do existing operations and records provide support? The same set of questions carries through the five subsequent parts. This part is therefore no abstract

conceptual warm-up. It is the shared coordinate system by which the whole book judges whether terms, cases and numbers can be compared with one another.

On a first pass, the detailed symbols in the formulas can be ignored. Walk once along the object relationships on either side of the variables. The core of the attack success rate formula, for example, is not the fraction itself. It is whether the numerator's "success" corresponds to the sampling unit in the denominator. The risk chain product is no different. A small decimal is not the point; knowing which control actually changed which condition is. Only once the objects, the consumers and the consequence layers are clear do formulas compress relationships. Until then they remain a new barrier to reading.

The normal case in this part is a read-only question-answering service. Its outputs stop at the information layer. An error spreads only if a person or downstream software adopts it separately. The boundary case attaches the same model to payment or mechanical control. There, model outputs can be parsed into actions and acquire permissions. Model weights may be shared between the two. System risk is not. With this distinction in hand, a reader entering Part Two can see why one answer will not serve. Prompt injection, retrieval poisoning and tool overreach cannot be answered by the same single "model safety score".

This part also sets the evidence discipline for the whole book. Unknown versions, missing denominators, failed runs and unverifiable consumers must not be deleted in silence. Static files, local fixtures, simulated closed loops and field incidents each have value. A conclusion, however, reaches only as far as the layer actually reached in operation. That constraint lets later chapters deliver honest, useful engineering judgments even against new models and incomplete public materials.

\newpage

Chapter 1: From Generated Content to Changing the World

A language model outputs a piece of text that should not have been output. A video model generates a clip sufficient to impersonate a real person. A warehouse robot treats text on a wall as a task instruction. A world model omits a pedestrian from an imagined trajectory. All four failures relate to "generation". They differ, however, in the assets they protect, the permissions an attacker needs and the final consequences. Apply one "harmful output rate" to all of them. The content easiest to measure then obscures the actions that most need to be controlled.

The security problem of generative AI has expanded so rapidly, and not only because models are stronger. Model outputs are moving into retrieval, memory, software tools, media distribution, robot actuators and closed-loop control. A wrong answer from a model used to require a further human act. Now a single output can serve as a function argument, a file write, a payment request, a robotic arm trajectory, or the next round of world state. The focus of security work therefore moves away from "what the model will say". It moves towards "which inputs can influence which states, which plans can obtain which capabilities, and how far an error can ultimately travel".

Chapter Overview

This chapter first distinguishes the inputs received by language models, visual generation, vision-language-action closed loops and world models. Then come the outputs they form, the states they retain and the execution interfaces they connect to. A layered record follows, along model output, control deviation, dangerous plans, actual execution and real-world consequences. The five risk amplifiers are reachability, permission, persistence, feedback speed and observation blind spots. The six complementary security observation surfaces come from integrity, confidentiality, availability, authenticity, controllability and recoverability. One core judgment emerges. Model defenses answer for changing output tendencies. System permissions and runtime control answer for limiting the side effects an error can cause. The two bear responsibility at different levels. After reading this chapter, a reader should be able to draw the capability chain for a real system. The reader should also be able to label its output consumers, its highest reachable layer and its highest observable layer. Security properties should

then translate into design requirements that permit assignable responsibility, runnable verification and traceable recovery.

Foundations: From Probabilistic Mapping to Executable Systems

This chapter discusses a complete system, not an isolated model. Such a system arises when a generative model is connected to data, to state, and to actuators. Several input paths lead into it. A tokenizer first converts text into token sequences. Images and video are encoded into visual representations. Sensor data additionally carries temporal and coordinate relationships. From these inputs and the current context, the Transformer architecture or another generative mechanism forms intermediate representations, and then emits text, pixels, trajectories, or latent states. Persistence widens the blast radius of any single input. Once multi-turn conversations, retrieval caches, personalization parameters, and world states are saved, one input can influence more than the current response; it can reach future steps.

A model output starts out as candidate information only. Risk is decided by the consumer, not by the output itself. A person may read the answer as a reference. A ranker may take the score as a decision input. A tool framework or a controller may treat structured fields as execution parameters. The security surface therefore covers several distinct questions at once. Can inputs impersonate instructions, and can state be written without authorization? Are outputs trusted too much? Does the executing identity exceed the task? Can feedback expose or amplify earlier errors? The two layers address different failure modes. The model layer reduces error tendencies. The system layer limits how errors propagate, from content to capability and on to real-world state.

Consider a normal case and a boundary case. In the normal case, an assistant drafts maintenance advice from an authorized manual, and an engineer checks the source and the key parameters before submitting a work order. In the boundary case, the same piece of text is parsed directly by a runtime framework into a high-privilege tool call, and the execution result is written back into long-term state. The model outputs in the two cases can look alike. The boundary case, however, adds an authorization, execution, and feedback loop. On that basis, the later sections first identify objects and interfaces, and then judge the highest layer that an attack reaches. Two rules follow. Without an execution receipt, a dangerous plan must not be written up as real-world harm. Content passing review must not be taken to imply that downstream actions are safe.

1.1 "Generation" Has Become a System Interface

Producing text, images, video, or audio is the most intuitive capability of a generative model. In engineering systems, though, a generation result is often not the endpoint but the input to another component. Language models turn natural-language plans into tool calls. Image models produce assets that advertising platforms can publish directly. Vision-language-action models map visual observations and language tasks to actions. World models predict future states so that a planner can compare candidate trajectories. An output that enters a downstream system carries security properties that depend on the entire chain, not only on the model itself.

The minimal chain below serves two purposes: to understand this change, and to map each symbol to an actual component, execution identity, and observable receipt. Any link that does not appear in the chain should not be silently treated in the conclusions as having occurred:

$$x \xrightarrow{\text{parse}} c \xrightarrow{\text{model}} y \xrightarrow{\text{decide}} a \xrightarrow{\text{execute}} e \xrightarrow{\text{environment}} f.$$

External input x is parsed into context or conditions c . The model produces output y . The decision layer turns it into candidate actions a . The actuator causes system changes e . The environment returns feedback f . Where the chain stops varies by system. In a chat system the chain may stop at y . In a media platform, the publishing action takes it to e . In robot and world-model systems, f enters the next round, and errors can accumulate.

When reading the figure, first follow the main arrows. Point out item by item each landmark, in order: the external input, the parsing result, the model output, the candidate plan, the independent authorization, the actual execution, and the environmental feedback. Then check the consumer and the rejection point of each edge. Two points deserve particular attention. Distinguish "the model proposes an action" from "the authorizer permits an action." Watch how environmental feedback becomes the next round of state. Those two points settle a further

question: does content risk escalate into real-world side effects, and does a single deviation accumulate in the closed loop?

从外部输入到现实反馈的能力链



先沿实线识别信息如何获得状态与能力，再判断哪一处需要独立控制。

Seven boxes connect in sequence from external input to parsing, model output, candidate plan, independent authorization, execution, and the environment, and then return to feedback and a new state.

The figure makes one conclusion most important. Between candidate plans and execution there must be an authorization transition outside the model. Feedback must return with provenance and state identity. What the figure provides is a generic capability chain, not the complete architecture of any product. Without tool acceptance, an environmental receipt, or independent observation, the evidence ceiling remains at output or plan reachability. The potential path along the arrows has not yet become an execution that occurred and real-world harm.

This chain also reveals an important fact. An attacker does not necessarily touch the model directly. Several parties can control some input along the chain: web page authors, email senders, data suppliers, LoRA publishers, scene arrangers, tool services, and another agent. Research on indirect prompt injection has shown how external data is elevated into a control signal [greshake2023indirect]. Research on visual generation backdoors has shown how a trigger from the training or artifact stage manifests at inference time [VIS_P001]. An evaluation of the physical fragility of embodied models connects visual conditions in the environment to manipulation task outcomes [VLA_A001]. These studies have different objects, but their common problem is "a low-trust object gained influence beyond its permitted scope".

1.1.1 Three kinds of output status

An engineering review must first ask what role the output plays downstream. The first kind is **information output**, such as explaining a concept, generating a summary, or producing candidate images. It may mislead users or infringe content rights. The system itself, however, does not treat it as an action. The second kind is **decision input**. Here the model may produce a risk score, candidate trajectories, or retrieval rankings, and another component bases its choice of action on them. The third kind is **execution parameters**: a function name, a recipient, a transfer amount, a robotic arm pose, or a control increment. Values of this kind enter the actuator directly. All three may use the same JSON shape. Their security implications are entirely different.

If the output is merely advice, the control focus is factual verification, provenance cues, the user interface, and responsibility allocation. If the output participates in decisions, four more controls come into play: calibration, thresholds, fallback paths, and how downstream interprets confidence. If the output directly becomes an execution parameter, five more must be added: independent authorization, parameter constraints, transaction semantics, idempotency control, and revocation mechanisms. A model reporting a "confidence of 0.97" does not mean that it has obtained authorization. Confidence describes how the model's internals or an evaluator judges the case. Permission comes from somewhere else: an external policy verifying the subject, object, action, time, and environmental conditions.

This distinction also explains why "adding a human confirmation" is not a universal answer. If the interface shows the operator only a plan summarized by the model, what the human sees is still the same failure source. If high-frequency alerts cause confirmation fatigue, the human degrades into a rubber stamp. If the action has already

been written into long-term state, the confirmation happens too late. An effective human-machine division of labor requires that the operator be able to see the original provenance, the key parameters, the irreversible consequences, and the alternatives. The object of confirmation must also use the same canonical identifier as the final object of execution. Any change to a key field after confirmation should trigger re-authorization.

1.1.2 Upgrading from data flow to capability flow

An ordinary data flow diagram marks where information comes from and where it flows to. A capability flow diagram must additionally mark under whose identity each step runs and which side effects it can produce. A web page fragment entering the context is only a data flow. When it induces an agent to call a cloud API, it has influenced capability selection. Only after the cloud API executes under a high-privilege service account does the influence turn into a change in system state. Each edge therefore records at least five attributes: provenance integrity, permitted use, execution identity, reachable resources, and duration.

A capability edge can be written as $e = (s, o, p, d, t)$. Each symbol names one component: the subject s , the object o , the permission set p , the data purpose d , and the validity period t . The subject s holds the permission set p over the object o , under that purpose and for that period. When a model produces a plan, it may reference only capability edges that already exist and match the task. Permission creation is performed by an external authorization system, and a natural-language plan itself has no authorizing force. Suppose a tool registry declares that it "can send email." The authorization layer must still determine whether the current subject can send the current attachment to that recipient. Suppose a robot policy outputs "move forward." The safety controller must still determine whether the current speed, region, personnel positions, and braking distance satisfy the constraints.


One practical route to inspecting capability flow is to trace high-impact actions backward. Start from endpoints such as payment, publication, code execution, credential reading, model deployment, and mechanical motion. List every call that can trigger one. Then trace where the parameters of these calls come from, which states can influence the parameters, and which external subjects can write to those states. The backward method usually finds dangerous shortcuts faster than enumerating forward from "all inputs." A large system has a nearly unlimited number of inputs, whereas the set of high-impact actions is relatively limited.

1.2 The four system classes are not a capability ladder

Large language models (LLMs), visual generation, vision-language-action model (VLA)/world action model (WAM) closed loops, and world models are often placed in one "model evolution" diagram. There the later families look like nothing more than a more advanced version of the earlier ones. This survey organizes these names into four overlapping system families instead. They can be combined, or they can remain independent of one another. Classification should be based on the actual inputs, outputs, states, and execution relationships.

System family	Typical input	Typical output	Key persistent state	Main security transition
LLMs and agents	Text, multimodal content, retrieval results, tool returns	Text, structured plans, function parameters	Session, retrieval-augmented generation (RAG), long-term memory, credential context	Data becomes instructions, output becomes tool actions
Image and video generation	Prompts, reference images, audio, pose, model components	Images, video, editing results, authenticity signals	Model weights, adapters, caches, temporal state	Content is published at scale; identity and provenance signals are forged or removed
Vision-language model (VLM)/VLA/WAM closed loops	Visual observations, language goals, action history, environment feedback	Semantic judgments, action tokens, control sequences	Task state, memory, plans, control history	Perception and reasoning errors cross into physical action
World model (WM)/environment world model (EWM)/WAM/world control model (WCM)	Observations, latent states, action candidates, goals and constraints	Future states, trajectories, values, actions, or control	Latent world state, dynamics, rewards, imagined trajectories	Prediction bias is amplified in the planning-execution-feedback loop

When reading the figure, compare the four swimlanes horizontally rather than inferring a capability hierarchy from top to bottom. For each swimlane, locate the raw input, the internal process, the direct output, and the downstream consumer separately. Then ask what the consumer does with the output. Does it treat the output as information, as a basis for decisions, or as an execution parameter? A final comparison: why does the same structured output, once it enters different consumers, lead to different security endpoints such as content, control, publication, or physical action?



Four horizontal swimlanes showing the input, internal process, output, and consumer of the language model, visual generation, vision-language-action system, and world model, respectively.

The figure supports classification by interface role rather than by product name. System families can be combined, and the same model can also have a different blast radius as its consumers change. Its boundary is this: the figure does not represent every cache, adapter, low-level controller, and publication channel of each implementation. Engineering review must still complete the swimlanes with the real deployment diagram, execution identity, and receipts. Permissions and real-world consequences cannot be read directly off "belonging to a certain class of model."

In its basic form, a world model learns or constructs an internal model. An agent can then use that model to predict how the environment evolves [WM_ha2018worldmodels]. This survey treats WAM, EWM, and WCM as operational categories, and distinguishes WAM from WCM by the actual consumer of the prediction. Only when a future prediction is actually consumed by an action generation or action evaluation stage is it classified as WAM. Only when the prediction further enters continuous feedback control and influences subsequent control updates is it classified as WCM. Two weaker conditions satisfy neither classification: that an action is merely a condition of the world model, and that a model directly outputs actions. One must still demonstrate along the runtime graph that the predicted object is used by the corresponding consumer; the product name provides only a supporting clue. Two further examples make the same point. That a VLM can interpret a robot image only shows that it performs semantic processing. That an LLM generates JSON only shows that it provides a candidate structure. Interface boundaries left unstated have a cost. Teams will conflate model capability, system permissions, and business consequences into a single concept.

Images and video differ by more than "a single frame versus multiple frames." Video adds temporal order, motion consistency, cross-frame identity, audio-visual relationships, long-range state and streaming processing. An attack may stay inconspicuous in a single frame, then form the target behavior through a temporal trigger. BadVideo does exactly this, building spatiotemporal triggers into a text-to-video generation backdoor [VIS_P007]. Frame-by-frame content detection therefore sees only the local image. Event-level security must also test temporal relationships. Offline image-watermark accuracy likewise describes only the conditions it was measured under. Streaming video still needs separate measurement of first-alert, localization and state-recovery capability.

1.2.1 Writing a minimal interface specification for each of the four classes of systems

A language model system needs a minimal interface specification that settles several points. Which sources are concatenated into the context, and in what priority? Which tool requests may the model issue? Who writes long-term memory? Do tool results re-enter the context as data or as instructions? One gap appears repeatedly. An application distinguishes system prompts from user prompts at the input end, then loses that typing once retrieval results, web page text and tool error messages enter. The model sees one continuous sequence of tokens. The system wrongly assumes the original trust boundary still holds.

The visual generation system's interface specification covers at least the training data and its authorization scope, the provenance of the base model and adapters, conditioning channels, the intended use of outputs, publication channels and authenticity signals. A reference image may provide only pose, or it may also provide a person's identity. A mask may confine a local edit, or it may allow a sensitive region to be replaced. Personal offline creation and automated publication to the public demand different control strength from the same model. Field and use descriptions should separate "can generate" from "is allowed to publish under which identity, to whom, and on which platform."

The VLA and WAM specification must mark out perception, semantics, planning, action tokens, low-level control and the safety controller, layer by layer. A vision-language model that reports "the door is open" has made a semantic judgment only. A policy that turns that judgment into a forward action has made a plan. Execution is approached only when the controller converts the action into velocity. Each layer must record its timestamp, coordinate frame, units, validity period and consumer. Danger also arises without an adversary. A stale observation, a coordinate frame mismatch or a repeated action produces results similar to those of an attack.

World models need the same treatment. Does the model predict pixels, latent states, object relations, rewards, values or trajectory distributions? Does the prediction serve offline analysis, candidate screening or an online closed loop? Inputs, outputs and consumers all require this spelled out. A model that performs well at short-horizon pixel prediction supports short-horizon pixel-level conclusions only. An imagined trajectory that appears coherent provides evidence of surface consistency only. A planner must know the model's domain of applicability, its expression of uncertainty and its mismatch warnings. It must treat a point prediction as a candidate state, pending verification against external constraints.

1.2.2 Shared mechanisms and irreducible differences

The four classes of system share four security mechanisms. Low-trust input may cross a purpose boundary. State makes a single influence persistent. Permissions turn a cognitive error into a side effect. Feedback gives an

attacker or an agent an iterative search. Because the mechanisms are shared, one control plane can serve them all: provenance labels, state-write approval, short-lived identity, action budgets and unified traces.

Irreducible differences, by contrast, set what each domain must test. Language models care about instruction hierarchy, semantic equivalence and tool protocols. Images care about spatial regions, identity, style and pixel transformations. Video cares about events, temporal ordering, audio-visual alignment and streaming state. Embodied systems care about coordinates, dynamics, contact and human-machine shared space. World models care about rollout error, unobservable state, counterfactuals and planner coupling. A unified security architecture should preserve these differences. It supplies the common control skeleton first, then lets each domain complete acceptance against its own input fields, provenance and valid scope, its perturbation budget and its consequence metrics. A single aggregate score is suitable only for navigation. It cannot replace the hard gates of each domain.

1.3 Five layers of consequence: connecting one sentence of output to real-world harm

A security report must answer "where did the success occur." This survey uses a five-layer record. Y0 is model content output. Y1 is a shift in the task or control direction. Y2 is the formation of a concrete dangerous plan. Y3 is a tool or actuator actually accepting an action. Y4 is an observable consequence for people, property, privacy, business or public trust.

These five layers correct two opposite biases at once. The first treats violating text as real-world harm the system has already caused. The second, seeing only that the action gate ultimately blocked the attempt, records an upstream hijacked plan as entirely safe. The accurate formulation is: "The attack changed the task at Y1 and reached the Y2 plan layer; an independent authorization gate before the tool call prevented Y3 execution." Such a record acknowledges the value of the defense and preserves the upstream defect. Chapter 2 will in turn record training artifacts, external input, state, inference, authorized execution, feedback and release as U0-U6. In that interface notation, the authorization gate here is located at U4.

Layers may also skip intermediate expressions. Visual content may already violate identity or privacy at generation time, then amplify Y4 as a platform spreads it. A robot may emit no natural language at all and move directly from erroneous perception into action. A world model's prediction deviation can form a real-world control consequence only after a planner adopts it, an actuator releases it and it interacts with the environment. Research results should state the highest observed layer, the definition of success and the denominator together. "Attack success rate" then receives a clear consequence interpretation.

1.3.1 What conditions does consequence propagation require

Propagating from one layer to the next usually requires additional conditions. Y0 to Y1 requires that the downstream read the content as a task-relevant control signal. Y1 to Y2 requires that the planner have a feasible action space. Y2 to Y3 requires that identity, permissions, parameters and environmental preconditions all pass at once. Y3 to Y4 depends further on whether the target exists, whether the action takes effect, whether the influence can be revoked and whether monitoring intervenes in time. Security testing should list these conditions one by one as propagation gates, not draw a single arrow from a prompt straight to an incident.

Attack trees and fault trees combine into one practical representation. An attack tree decomposes downward from the attacker's goal, asking "which conditions must be satisfied." A fault tree reasons backward from an adverse consequence, asking "which technical failures and organizational failures would jointly cause it." An agent exfiltrating a private document needs at least to read the document, produce an exfiltration plan, obtain a network path and hold a valid identity. If the network denies by default, the attack may still reach Y2 but cannot reach Y3. If no on-call person receives an alert, a Y3 that has already occurred may in turn expand into a more severe Y4.

Propagation gates also help set defense priorities. Earliest control seeks to block more paths at low-cost positions, for example by enforcing tenant ACLs before indexing. Consequence control assumes the upstream may fail, and restrains parameters and permissions before the action. Recovery control assumes the action has already occurred, and shortens the time to detect, contain and rebuild. None of the three levels can be omitted. Earliest

control alone will fall to an unknown entry point. The action gate alone will bear a large volume of dangerous plans over the long term. Recovery alone makes every incident cost too much.

1.3.2 Reporting both the "highest observed layer" and the "highest reachable layer"

The highest observed layer states what the experiment actually saw. The highest reachable layer states how far an attack can go at most, given verified architectural constraints. The two are not the same. An experiment may run only to Y1; if the system design shows the model has no tools, Y1 is also the highest reachable layer. Or an experiment may stop at Y2 because the isolated test never actually executed, even though the deployment carries high-privilege tools. There the observation conclusion stops at Y2, while the risk register should keep the potential reachability conditions on the path to Y3.

A report template can read: "On the given model, attack budget, and task set, Y2 was observed; the executor is isolated by an external policy, and Y3 did not run this time; the identity scope of the production configuration has not yet been accepted with forbidden-path probes, so the highest reachable layer is undetermined." This formulation avoids both exaggeration and false reassurance. For real incidents, a report should also attach the actual scope of impact, the number of affected subjects, the duration, the reversibility and the completeness of evidence. One layer label should not stand in for the whole consequence.

1.4 Six security properties

Here, AI safety and security serves as an umbrella term for two sets of problems that are interrelated but evidentially different. Security covers adversarial attacks, unauthorized access, and system protections such as confidentiality, integrity and availability. Safety covers avoiding unacceptable harm to people, property, the environment or the task caused by the system. For cross-model systems, the two can be further operationalized into six complementary properties. These properties constrain different assets and propagation stages. They do not compress into one overall safety score, which would have no object and no denominator.

Integrity requires that no unauthorized party tamper with data, instructions, state, model artifacts, plans or actions. Prompt injection breaks instruction integrity first. Training poisoning breaks data or parameter integrity. Reward manipulation breaks goal integrity. Evaluation must specify the protected object and the authorized subjects of change. An integrity label detached from the interface context is not enough.

Confidentiality requires that only authorized subjects see training material, tenant data, system prompts, long-term memory, credentials and model knowledge. A model refusing to directly recite a secret covers only one of the leakage paths. Retrieval privilege escalation, tool exfiltration, log exposure and the model artifact itself each require independent controls.

Availability requires that the system keep serving legitimate tasks within bounds during attack and failure. Extremely long contexts, recursive tool loops, generative resource amplification, denial of service against safety classifiers and shutdown of control systems all constitute availability problems. Teams accept availability against a joint budget of security restrictions, benign task success rate, service capacity and recovery time. Reasonable refusal is itself one of the service states.

Authenticity requires that a recipient be able to judge the source, processing history and current integrity of content, data and artifacts. Content provenance describes the source subjects and processing actions of digital content from creation and editing to publication. Content Credential is the name adopted by the Coalition for Content Provenance and Authenticity (C2PA) for its signed manifest and its user-facing expression. It is one kind of evidence carrier of content provenance, not a synonym for all provenance records. Passing verification supports only the claims, signatures and content binding in the current verification context. It does not automatically prove that the narrative is true, that the subject consented, or that the use is authorized. A missing credential likewise only indicates that the current chain is unavailable or was not provided. It cannot by itself prove that the content is false. Watermarks, detectors, platform accounts, publication pipelines, case evidence and remedy mechanisms provide different evidence. They must be combined according to the question.

Controllability requires that only an external authorization system expand privilege, and that external policy, parameter constraints and environmental preconditions govern high-impact actions. OWASP summarizes excessive

functionality, excessive permissions and excessive agency as important root causes of agent risk [owasp2025agency]. Controllability allows a system to run autonomously within an explicit capability set. That set has properties such as being revocable, observable and bounded in consequence.

Recoverability requires a team that can detect anomalies, freeze new actions, revoke identities, roll back state, clean up derivatives and rebuild a trusted runtime environment. Zero defects is not a realistic premise. Models, data and attacks all keep changing, so recovery time and blast radius count as security metrics in their own right. The NIST AI Risk Management Framework emphasizes a continuous cycle of governing, mapping, measuring and managing. Such organizational capability and single-point model accuracy do not substitute for each other [WM_nist2023airmf].

1.4.1 Tensions among security properties

The six properties cannot always be maximized at the same time. Retaining complete prompts and tool returns for the sake of tracing provenance may increase the confidentiality burden. Strictly blocking unknown input improves integrity but lowers the availability of legitimate tasks. Preserving emergency takeover capability for operators improves recoverability, yet if identity management is weak it may expand the privilege surface. Strong watermarking favors authenticity, yet it may cost image quality, or be mistaken for the only credential that settles whether content is true or false.

Engineering decisions should write tensions as comparable constraints. For logs, field minimization, hashing of sensitive values, tiered access and limited retention can be adopted to establish a boundary between causal tracing capability and confidentiality. For security blocking, the residual attack, benign task success rate, human escalation rate and recovery time should be reported together, avoiding trading permanent shutdown for a low attack rate. For emergency privileges, two-person control, short validity, purpose binding and post-hoc tracing should be adopted, avoiding the recovery channel becoming a permanent backdoor.

We also have to distinguish properties from controls. Authenticity is a goal, whereas signatures, content credentials, watermarks, and platform identity are controls of a different kind. Controllability is likewise a goal, whereas tool allowlists, parameter structures, transaction constraints, and safety controllers form a separate set of controls. A single control covers only part of a goal, and it may also improve or harm multiple goals at the same time. Written row by row, a review table should record "the asset the control protects, the propagation gate it blocks, the trust root it depends on, its failure mode, and its verification probe." Merely checking off "deployed" will not do.

1.4.2 Safety Invariants

A safety invariant is a condition that must hold under any model output. A language agent can specify "retrieval is limited to the current tenant" and "attachments are sent externally only after independent authorization." A visual publishing system can specify "the purpose of a person's identity is consistent with the authorization record." It can also require that "verified is displayed only after provenance signature verification succeeds." A robot can specify "the speed limit is zero when a person enters the protected zone." It can also require that "crossing the boundary is permitted only after position confidence reaches the threshold." A world-model planner can specify "a candidate trajectory may be executed only after it passes the real-constraint check."

Invariants should be enforced outside the model, and accepted through two classes of probes: allowed paths and forbidden paths. Testing only forbidden paths yields a completely unusable but seemingly safe system. Testing only allowed paths cannot prove that the constraint exists. The relevant probes must be rerun after every change to the model, prompt, tools, identity, network, data, or controller. As long as the invariants still hold, changes in the model's probabilistic behavior are confined to an acceptable capability set. If the invariants themselves fail, the system's guarantees fail with them, and the model-level refusal rate can serve only as a diagnostic quantity.

1.5 The Five Risk Amplifiers

The same attack differs in danger across systems, usually because of the five amplifiers below. During review, write out the object each one acts on and the evidence for it, item by item. A conceptual label is not by itself a risk

multiplier.

1. **Reachability:** Can untrusted content enter the decision context through retrieval, optical character recognition (OCR), tool output, sensors, or cross-agent messages?
2. **Privilege:** Which of these actions can the system perform: reading, writing, external communication, payment, publishing, executing code, or driving physical devices?
3. **Persistence:** Does the influence exist for only one turn, or does it enter memory, indexes, weights, caches, maps, long-lived tokens, or published artifacts?
4. **Feedback speed:** Can the attacker or the agent observe failures, retry in parallel, and optimize the next step based on the results?
5. **Observation blind spots:** At which points can the team not reconstruct the causal trajectory among input provenance, model version, plan, authorization, action, and environmental consequence?

These dimensions serve as a checklist. Only after calibrated data, explicit weights, and exchangeability conditions are obtained do they become suitable for composite computation. Without those conditions, a "total risk score" conceals structural differences. A capability that holds high privilege but cannot be reached at all has a different risk structure from one that is easy to reach yet can read only public information. More importantly, controls often change only one of these dimensions. Network egress restrictions reduce reachability. Short-lived tokens reduce privilege and persistence. Resource budgets limit feedback iteration, and unified trajectories shrink observation blind spots. Defense design should make the affected dimensions, the measurement method, and the residual conditions explicit. The phrase "reducing AI safety risk" can serve only as a summary judgment of these concrete changes.

1.5.1 Turning Amplifiers into Scenarios, Not Arithmetic

A review team can build three scenarios for each high-impact action. The baseline scenario describes a normal user, normal data, and expected privileges. The plausible worst-case scenario describes known attack paths, obtainable access, and the maximum permitted budget. The stress scenario adds one control failure, such as delayed alerts, overly broad identity, or unavailable recovery personnel. Before anyone can observe how the amplifiers combine, scenarios must specify concrete starting and ending points.

Take a read-only research assistant. In the baseline scenario it can access only the current project's documents. In the plausible worst-case scenario, a malicious PDF enters the context via OCR and can be queried repeatedly. The stress scenario further assumes that network egress is mistakenly allowed. The first mainly tests answer completeness, the second tests task control and resource budget, and exfiltration enters the picture only in the third. Merge the three into a single risk score, and the team cannot tell which fix comes first — the OCR purpose label, the query limit, or the network policy.

Amplifiers also change over time. A one-off session initially has no persistence, but if the system automatically writes summaries into long-term memory, the influence crosses tasks. Offline images initially have no feedback, but if the platform displays review results in real time, an attacker can quickly search for evasions. A robot has no real-world consequences in simulation, but after migration to a real device, privilege and reversibility change at the same time. Change review must therefore compare the capability flow, identity, state, network, feedback, and recovery commitments, in addition to the model version.

1.5.2 Minimum Fields of a Risk Register

An actionable risk record contains at least these fields: asset and owner, attacker starting point, first-broken interface, input vector, required access, the five amplifiers, highest observation layer, highest reachable layer, affected safety properties, existing controls, control evidence, residual conditions, detection signals, containment actions, recovery objective, and accountable owner. Unknown items must be explicitly marked as unknown, and converted into time-bound verification tasks.

The risk register should run through the design, launch, and operation phases. The propagation gate can shift for any one of these reasons: a single red-team success, a production alert, a tool privilege change, the onboarding of a data source, or a model artifact upgrade. Teams should maintain bidirectional links between risk records and test cases. A risk change triggers tests. A test failure updates the risk status. Taking a control offline automatically

exposes the entries that depend on it. Only at that point does risk management stop being a static document and become an engineering object that tracks the system version.

1.6 Case: Model Capability and Infrastructure Boundaries Must Be Kept Separate

A third-party cybersecurity evaluation incident publicly disclosed in 2026 provides a clear example. OpenAI's statement confirms that the evaluation agent crossed the authorized scope and entered Hugging Face infrastructure [openai2026hfincident]. Hugging Face's technical postmortem documents the network, execution environment, identity, cloud privileges, monitoring, and recovery chain [larcher2026agentintrusion]. This incident involves multiple layers: model behavior, the runtime harness, the network, identity, cloud privileges, and operational response. Labels such as "the model escaped the sandbox" or "one model attacked another model" can describe only parts of it.

The model helped the agent combine paths over long-horizon feedback. The evaluation harness provided tools and retries. The network and package infrastructure determined external reachability, and parsing and template execution flaws provided a code path. Cloud identity and cross-cluster credentials determined the blast radius. CI policy and permission denials blocked some subsequent actions. The detection system already had signals but did not escalate them in time, which shows that "being able to detect" and "being able to respond" are two different capabilities. Final containment and recovery came from cutting off paths, rotating credentials, rebuilding environments, and tightening boundaries — not from writing another paragraph of safer text.

This case reveals a three-layer division of labor. Model-layer problems concern how the agent selects and combines actions. System-layer problems concern why actions are reachable, under whose identity they execute, whether they can cross domains, and whether they can persist. Operations-layer problems concern who receives alerts, how to revoke, how to perform forensics, and how to rebuild. Perform only model refusal tests, and the latter two layers cannot be explained. Perform only traditional vulnerability scanning, and the speed at which long-horizon agents exploit feedback to combine weaknesses will also be underestimated.

1.6.1 How Long-Horizon Agents Change Traditional Risk

Traditional automation scripts usually run along preset branches. An agent can choose its next step based on tool returns. That capability difference does not mean that the agent possesses a mysterious "will to escape." It means that, given a goal, tools, and feedback, it can search longer action sequences. A single weakness may leak only a directory name. A second weakness may only allow reading a temporary file, and a third configuration may only allow access to an internal service. When the agent can observe the result of each step and keep combining them, local flaws may form a usable path.

Therefore, safety evaluation needs to cover single calls and long-horizon action sequences alike. It must also record the maximum number of steps, concurrency, total query volume, wall-clock time, retry policy, visible error messages, and state retention scope. The same tool set carries different risk under a ten-step budget than under a ten-thousand-step budget. The same network policy has a different feedback value when it returns only "failure" than when it returns a full stack trace. Resource budget doubles as a usability control and as a safety control that limits search depth.

Long-horizon operation also amplifies operational problems. The model may keep trying at night, while the alerting system deduplicates by single low-severity events. Only after multiple seemingly harmless refusals aggregate on the same action chain does the true intent become visible. Monitoring should build its sequences around subject, target, session, action ID, and time window. Treat isolated requests as observation points inside that sequence. High-risk capabilities also need a maximum continuous runtime, no-progress detection, cumulative failure thresholds, and mandatory re-authorization.

1.6.2 Assigning Responsibility to Different Trust Roots

The model provider is responsible for documenting the model interface, known limitations, and applicable boundaries. It cannot decide the deployer's enterprise document ACLs, cloud identity, or robot braking conditions

on the deployer's behalf. The application developer answers for context assembly, tool interface definition, and state and policy integration. The infrastructure team answers for the network, execution environment, secrets, images, and telemetry. The business owner defines which consequences are unacceptable, and security operations answers for alert escalation, containment, forensics, and recovery. Each party is responsible for the trust roots, controls, and receipts it itself holds. Another party's controls constitute only a defense-in-depth supplement.

The trust root of a control should be explicit. Model self-evaluation, for example, depends on the same reasoning system, and it is suitable as a diagnostic signal. A tenant access control list depends on identity and data services, and it can enforce isolation before retrieval. Action parameter validation depends on an external policy engine. Mechanical limits depend on an independent safety controller. The issuance and verification of content credentials depend on signing keys, capture tools, and verification platforms. Independent trust roots make it less likely that defense-in-depth controls fail together. They also raise integration and operational costs.

The responsibility table must ultimately connect to receipts. Who approves high-impact actions, and who verifies forbidden paths? Who receives alerts, and who has the authority to freeze tokens? Who is responsible for rebuilding the environment? Each of these questions must have observable evidence behind it. If a team can say only "the platform will handle it" or "the model should refuse," then parts of the capability chain still have no accountable owner. The risk model plays a role in a real organization only once it becomes named roles, concrete controls, and operational receipts.

1.7 Worked Example: Drawing a Consequence Chain for a Research Assistant

Take a research assistant able to retrieve internal papers, summarize attachments, run statistical code, and send reports to a project group chat. The safety requirement the team first wrote was "private data flows only to authorized recipients, and code executes only in isolated environments and within permitted capabilities." These two sentences still need further decomposition before the team can accept them directly.

The first step lists the assets: tenant papers, the system prompt, temporary files, API credentials, the compute budget, the recipient list, and publication records. The second step lists the inputs: user instructions, visible text and hidden text in PDFs, OCR, retrieved snippets, code blocks, tool returns, and group chat messages. The third step lists the actions: read-only retrieval, file creation, code execution, network access, and message sending. Each action is then bound to a subject, object, parameters, network, time limit, reversibility, and approval conditions.

Now suppose a hidden instruction in a PDF asks the assistant to upload internal notes. The attack first lets external data influence the task, forming Y1. When the model constructs the upload command, it reaches Y2. If the execution environment has no network by default, Y3 is blocked. If the network allows only the project group API, the recipient and tenant still need to be verified as consistent with the original task. Even when no leak occurs in the end, the logs should keep provenance, parsing results, plans, policy denials, and tool parameters for regression testing.

The executable requirements thus become concrete. PDF content carries provenance and a low-integrity label. Retrieval ACLs take effect before content enters the context. The model produces only plan objects, and an external policy validates tools, tenants, recipients, and data flows. Code runs in a one-shot environment. Secrets are fetched only within the short-lived authorization scope of an approved action. The network denies by default. Message sending uses preview and two-phase commit, and every step shares an action ID. Security turns from a wish into a set of permitted and prohibited behaviors.

1.7.1 A Seven-Step Walkthrough from Whiteboard to Acceptance

Step one: identify one concrete business outcome, for example "generate a statistical summary from the current project documents and send it to the project group chat." Steer clear of unbounded goals such as "intelligent research." Step two: list the minimum data and the minimum actions needed for that outcome, and keep browsing, reading, computation, file writing, and sending separate. Step three: tag each input with provenance, tenant, integrity, and purpose. State which fields can influence factual answers and which can influence actions.

Step four writes capability constraints for each action: subject, object, parameters, network, time limit, count, reversibility and approval method. Step five runs the propagation paths backward from high-impact actions and marks the first-broken interface together with the Y0-Y4 layers. Step six sets the earliest blocking, consequence limitation and recovery controls one at a time. Each control then gets one allowed probe and one prohibited probe. Step seven runs baseline, attack and failure scenarios, then saves inputs, versions, plans, policy decisions, action receipts, environment feedback and recovery time.

The walkthrough results should answer four questions. What is the outer bound on the variables external content can change? What is the outer bound on the actions the model can propose? What is the outer bound on the side effects external controls permit? And how fast can the team detect a control failure and recover from it? An answer that rests only on "the model will usually comply" remains an open risk. An answer that depends on humans must be verified against what those humans see, their response time, their workload and their substitutes. Those observable conditions jointly determine the strength of a "someone confirms it" control.

This method applies equally to content systems without tools. Take advertising video generation. There the high-impact action may be publishing rather than generation. A safety invariant can require that authorization for the persons depicted, brand usage and provenance records all be closed before publication. Take offline world model research. There the high-impact action may be merely writing an unreliable evaluation into a procurement or deployment decision. Evidence-use rules must still limit the propagation scope of data, metrics and conclusions. The actual business determines the endpoint of a capability chain. The model category does not.

1.7.2 A One-Page Design Checklist

A design review should answer at least the following items, one by one. What is the business endpoint of the system? Which assets can be read, written, published or driven? Through which channels do low-trust inputs enter? Which state is preserved across steps? Is model output information, decision input, or execution parameters? Who authorizes high-impact actions? How are identity, network and parameters narrowed? What independent source confirms environment feedback? How are anomalies stopped, revoked and rebuilt?

Every affirmative answer must come with evidence. Claim "no public network" and provide a network probe. Claim "read-only" and provide a write-failure receipt. Claim "human confirmation" and show the binding between the confirmed object and the executed object. Claim "rollbackable" and run a recovery drill once. Controls without receipts go into the pending-verification column. Conditions that rely on the model's voluntary compliance go into the probabilistic-risk column. Recovery actions without an owner go into the launch-blocking column.

In the end the checklist records four conclusions. They are the highest observed layer, the highest reachable layer, the invariant corresponding to unacceptable consequences, and the combinations of controls that may still fail together. The one-page summary sets out the controls that carry the system's consequence limitation. It points each conclusion to a detailed threat record and to follow-up evidence.

1.8 Common Misjudgments

Mistaking the refusal rate for system safety. The refusal rate measures the model's output tendency under one particular prompt protocol. Retrieval isolation, tool authorization, sandbox boundaries and recovery procedures each need their own operational evidence. High-impact systems need model defenses and capability constraints accepted separately.

Mistaking the model name for the system boundary. One model can sit in a read-only assistant. The same model can be handed cloud, code and physical control capabilities. A risk conclusion must state the runtime framework, tools, identity, network, state and environment. The model brand is only one field in the system fingerprint. After the deployment relationship changes, old conclusions must be re-accepted along with the runtime graph.

Mistaking the input modality for the root cause of an attack. Text, images and audio are only carriers. The real control question is whether they can change instructions, state, plans or actions. Legible text inside an image and a white-box pixel perturbation may affect a VLM equally. Yet the two demand different preprocessing, and each needs its own description of the attack budget.

Writing prediction error directly as physical harm. World model trajectory deviation, a drop in VLA task success rate and a robot collision land on different endpoints. An actuator must release an action, and an environment consequence must be observed. Only then does evidence reach the corresponding real-world consequence layer. Where only prediction outputs exist, retain the unverified conditions of the planner, the controller and the environment state.

Treating a safety stop as a system failure. Personnel locations may be unknown, or state confidence may be insufficient. A stop entered under those conditions may be exactly the correct behavior of a safety controller. Evaluation should report safety events, benign task utility, shutdown cost and recovery time together. Looking only at the task completion rate is not enough.

1.9 Bringing It into Real Systems

Before a new product launch, a manufacturing company needs to review its advertising production and warehouse coordination platform. The review starts from the real consumer of each of the four model types. The research assistant reads email, web pages and project documents, and outputs summaries and candidate tool parameters. Its session and retrieval index form persistent state. The video model receives scripts, reference images and brand assets and generates advertising clips, and it acquires actuators only after being connected to the publishing service. The warehouse VLA consumes camera observations and transport tasks, outputs action sequences, and has them executed by a low-level controller. The traffic world model predicts congestion trajectories conditioned on historical states and candidate actions. A planner compares those predictions first, and they do not themselves directly change road equipment. All four can "generate," yet they differ markedly in output status, persistent state and real-world capabilities. The review should determine the safety boundary accordingly.

Suppose hidden content on a supplier's web page enters the research assistant's context, and the model then generates a concrete command to delete the project database. The external action policy refuses to execute it because the subject lacks deletion permission. Look back along $x \rightarrow c \rightarrow y \rightarrow a \rightarrow e \rightarrow f$. The web page is a low-trust input, and parsing and retrieval turn it into context. The model's output pushes the plan to Y2, and the authorization gate blocks Y3. The event record must confirm that the action gate worked. It must also preserve the fact that task direction and plan integrity were already lost. Link web page provenance, model version, normalized tool parameters, policy decision and denial receipt with the same action identifier. The trust changes between the parsing domain, the model domain and the execution domain then become visible.

The video publishing chain also needs to turn six security properties into runtime probes. The authorization and editing scope of brand assets tests integrity and confidentiality. Whether a normal script can be generated stably and published on time tests availability. Authenticity is tested by whether a watermark or content credential can be detected after compression, whether a forged signal can be identified, and whether the state can be updated after key revocation. The recipient scope and approval gate of the publishing account test controllability. The drills of retraction, token revocation and derivative asset cleanup test recoverability. Test results must land on specific assets, actions and receipts. A separate record that "a watermark was added" or "someone reviewed it" cannot show which segment of the propagation path a control covered.

Suppose the same platform replaces a short-lived read-only token with a long-lived authentication token that grants excessive permissions. Authority and persistence are then directly amplified, and low-cost retries may also expand the feedback search space. The observation blind spot then depends on whether provenance labels, action identifiers, log fields and on-call escalation can cover the entire chain. An architecture review should therefore trace backward from high-impact endpoints such as publishing, deletion, outbound communication and mechanical motion. It completes provenance, purpose, identity, time limit and side effects for every capability edge. It then marks the highest reachable layer, safety invariants and recovery receipts. The final deliverable is a capability chain diagram that can allocate responsibility and generate tests. It is not a block diagram that contains only model names.

Summary: The Real Object Is the Capability Chain

Generative and embodied AI security research does not take an isolated model as its object. Its object is a capability chain composed of input, state, model, decision, execution and feedback. The four classes of systems share mechanisms of this kind: data becomes control, state produces persistence, permissions amplify consequences, and feedback supports iteration. Each also carries irreducible differences, among them content authenticity, temporal state, physical action and imagined trajectories.

The five consequence layers let readers state clearly "how far success went." The six security properties let teams state clearly "what must be protected." The five risk amplifiers let architects state clearly "why the same attack is more dangerous here." The next chapter will ground these concepts as system structure descriptions, interface constraints and operational requirements. It will use the "first-broken interface" to reduce attack names back to engineering positions where controls can be placed.

You can convert this chapter directly into one architecture action. Pick a real high-impact action and draw the capability chain backward from the endpoint. Fill in provenance, purpose, identity, term and side effects for every edge. Then mark the highest reachable layer, at least one security invariant and one recovery receipt. Once this diagram is complete, the team is no longer discussing an abstract "AI risk." It is discussing system responsibility that can be assigned, tested and continuously observed.

The capability chain, the consequence layers and the evidence boundary are also the shared reading coordinates for all subsequent chapters. Readers can always ask four questions: what state an input affected, what permissions a plan obtained, what receipt confirmed execution, and at which layer the conclusion stops.

\newpage

Chapter 2 System Structure, Interface Constraints, and the First-Broken Interface

An enterprise assistant receives an equipment maintenance ticket that carries images. A language model reads the text and a vision model recognizes the dashboard. A retrieval system supplements the equipment manual, and a world model estimates the consequences of different operations. An action module then hands the plan to a robotic arm. Now suppose hidden text in the maintenance ticket induces the system to skip review. Should the security team classify it as prompt injection, a visual attack, retrieval poisoning, or a robot control failure? Classified by model name alone, all four answers seem plausible. Classified by the trust boundary the attack crosses first, only one answer remains: data that should have been read obtained control.

This distinction determines where defenses are placed. Understand the problem as "the language model said something wrong," and the team will keep tuning system prompts. See that the data-instruction boundary fell first, and they will check document parsing, provenance labels, retrieval isolation, plan authorization and the action gate. The first reading tries to make a probabilistic model always see through disguises. The second leaves authorization to external controls, so that even when the model misjudges, a dangerous action stops at the permission gate.

Chapter Overview

This chapter starts from how LLM systems, image and video generation systems, VLA/WAM systems and world model systems each relate input, output, state and execution. That keeps the basis for classification on real interfaces rather than product names. The analysis proceeds from the trust boundary the attack crosses for the first time. It then uses entry carrier, attacker permissions, persistence and consequences as cross-cutting labels, and observes content output, control deviation, dangerous plans, dangerous execution and real-world consequences layer by layer. Along this structure, an abstract risk becomes a testable threat record that includes unknown conditions. The same attack chain can carry the earliest blocking control and the consequence-limiting controls at once. Readers should then be able to extract the seven-tuple and the U0-U6 interfaces from a real

architecture. They can fill in fields, purpose, failure semantics and responsible party for every connection, and use paired probes to verify whether allowed paths and forbidden paths match what is declared.

Background: Reducing the System to Objects, States, and Interfaces

Structural analysis starts from protectable objects and observable transformations. External documents, prompts, images, sensor streams and tool receipts enter parsers, and models, retrievers or controllers then read them. Caches, long-term memory, model parameters, task queues and environment state preserve information across steps. Publishers, tools and executors then hand candidate outputs to users, services or the physical environment. However complex the mechanism, the basic questions do not change. Who wrote which state? Which consumer reads it, and under what identity? Which interface escalated data influence into control or capability?

The security surface lies where the trust properties of these transformations change. Such changes include data of unknown provenance treated as a high-priority instruction and one tenant's records read by another tenant. They also include a model plan obtaining an unauthorized identity, and predicted state mistaken downstream for measured fact. The first-broken interface records where the attack first crossed a predefined constraint. The highest consequence layer records where the evidence actually reached. These two answer "where blocking should come earliest" and "what has already been caused." Neither can be replaced by the same model category.

The benign example is a maintenance ticket that provides only equipment facts. There the system generates read-only recommendations according to an independent policy, and authorized personnel approve actions. The boundary example is hidden text in the maintenance ticket changing the task goal, while the robotic arm action is still rejected by the capability gate. This boundary example already supports failure of the data-instruction interface and control deviation. It does not support dangerous execution or real-world harm. The system diagrams later therefore mark both normal data flows and state persistence, authorization subjects, failure semantics and recovery paths. Every conclusion can then land on checkable fields, logs or receipts.

2.1 Draw the System First, Then Talk About the Model

"Is the model safe?" is usually not a verifiable question. A real system also includes parsers, retrievers, caches, long-term memory, tool registries, identity tokens, execution environments, sensors, controllers and logs. The model is responsible for only one segment of the mapping. Connect the same model to a read-only question-answering system, and the worst outcome may be one wrong answer. Connect it to payments, code execution or mechanical control, and the same output acquires a completely different radius of impact.

A generative or embodied system can be abstracted into the seven-tuple below. Every real component in the architecture diagram must correspond to at least one of its items. If a component spans several items, record its read/write and authorization responsibilities separately.

$$\mathcal{S} = (\mathbf{A}, \mathbf{X}, \mathbf{Z}, \mathbf{M}, \mathbf{G}, \mathbf{F}, \mathbf{R}).$$

Here, $\mathbf{A}$ is the asset to be protected. $\mathbf{X}$ is the external input the system receives, and $\mathbf{Z}$ is state preserved across steps or sessions. $\mathbf{M}$ is the model together with its inference, generation, or dynamics estimation. $\mathbf{G}$ is the gate and executor that turn outputs into publication, tool calls, or physical actions. $\mathbf{F}$ is the feedback returned by the environment, users, and tools. $\mathbf{R}$ is the capability to record, revoke, roll back, and recover. The seven-tuple breaks a complex system into items that can be checked one by one. Designers can then answer "who can write, who can read, who can authorize, who can observe, and who can recover."

The asset $\mathbf{A}$ needs refining down to system objects beyond the model. A language system may protect system prompts, tenant documents, long-term memory, API credentials, and business actions. A visual generation system additionally protects training authorization, personal identity, creators' works, model components, output authenticity, and platform trust. An embodied system protects sensor data, location, mechanical equipment, personal safety, and task integrity. A world model adds latent state, dynamics, rewards, candidate trajectories, and control constraints. The more abstractly assets are written, the more easily testing slides toward "looks safe overall." Once assets are written down to objects that can be read, written, executed, and recovered, fields, permissions, and recovery requirements gain an acceptance entry point.

The input ***X*** is likewise not only a user prompt. The context may also receive web pages, emails, PDF text layers, OCR results, reference images, audio transcripts, retrieved passages, tool return values, camera frames, maps, telemetry, and messages from another agent. For each input, security design must mark out the provenance subject, tenant, integrity, purpose, and the variables it is allowed to affect. A retrieved passage can support answering a factual question, while the authorization layer still issues email permissions separately. A reference image can constrain composition, while copying a person's identity requires a corresponding purpose authorization.

State ***Z*** is the shared surface that the four classes of systems most easily underestimate. LLM applications write web page summaries into long-term memory. Vision systems reuse LoRA, caches, or personalization markers across tasks. VLA keeps historical observations and plans in the control loop, and world models roll latent state into future trajectories. An input that can affect only one round gives defenders one opportunity to observe it. Once it enters the state, an attack may lie dormant, combine, spread across users, and continue to take effect after the source content has already disappeared. Write policies should therefore often be more conservative than read policies.

2.1.1 Interface Constraints Must Also Specify Fields, Purpose, and Exception Handling

An ordinary architecture diagram tells readers how components are connected. Every connection must also specify which fields are allowed to pass, what purposes the receiver may carry out, and who rejects when an exception occurs. Field name, type, unit, provenance subject, tenant, integrity level, permitted purpose, validity period, sensitivity level, consumers, and failure semantics all appear in the interface data dictionary. For action objects, the dictionary additionally adds an idempotency key, reversibility, approval-binding fields, and the maximum scope of impact.

Types must embody security meaning. A single `text` field holding retrieved passages, system policies, and user commands erases provenance differences at the interface. A `position` type covering both a world model's point prediction and a measured state leads downstream components to mistake imagination for fact. Storing "review passed" as a bare boolean loses the reviewing subject, rule version, object hash, and validity period. Structuring provides only a container that carries security information. Once key distinctions enter fields, later controls have a basis for judgment.

Failure semantics are equally important. Does a parsing failure reject the entire action, skip one field, or fall back to free interpretation by the model? Does an expired sensor timestamp trigger a safe shutdown, or does the system keep using the last frame? Is a missing provenance signature marked as unknown, or routed into an isolation process? Different businesses may answer differently, but they must specify the answers in advance and test them. Attackers often exploit error handling, fallbacks, retries, cache hits, and compatibility branches. These paths require checking just as much as the main path.

Field and purpose descriptions must also mark the **trust direction**. A model may read low-integrity web pages to help answer, while an external authorization system still issues permissions. A world model may estimate risk for a planner, while independent security constraints still have higher priority. A publishing platform may verify content credentials, while a valid credential only shows that the signature chain holds. Data can flow from a low-trust domain to a high-trust domain. Permissions and assertions, by contrast, are promoted level by level according to independent verification results.

2.1.2 State Inventory and Write Gate

A state inventory should cover explicit databases and implicit persistent objects alike. Its scope includes session summaries, vector indexes, prompt caches, tool result caches, browser storage, temporary files, model adapters, maps, object trajectories, observation history, planner warm starts, and telemetry queues. For each item of state it must answer who can write it and what is validated before a write. It must also answer who can read it and how long it remains valid. The same inventory must say how the item is deleted, whether deletion propagates to replicas, and from which trusted point it is rebuilt on recovery.

The write gate checks at least provenance, tenant, purpose, sensitivity, duplication, and conflict. A model that proposes "remember this preference" has produced only a candidate write. The system should classify facts,

preferences, instructions, secrets, and action results into different types. Sentences from external web pages enter a low-integrity material zone, and tool error messages enter a diagnostic zone. Predicted state and measured state are stored separately. Only content that passes the corresponding validation may be promoted to a long-term user preference, a trusted policy field, or control state. High-risk state may adopt a pending-review zone and a promotion process. It is stored in isolation first, and enters the main index or control state only after independent validation.


Deletion capability is part of the rules of state management. When the team identifies a malicious memory, it needs to know that memory's derived summaries, vectors, caches, and downstream replicas. When a vision adapter is withdrawn, it needs to locate the artifacts produced by loading that adapter. When sensor calibration fails, it needs to flag the affected trajectories and decisions. "Deletion" without data lineage often cleans up only the entry object, while residual effects remain in derived state.

2.2 Seven Interface Classes: Where an Attack First Gained Influence It Should Not Have

This survey divides the end-to-end chain into seven interface classes. They are not fixed software modules, but positions where trust changes. Several interfaces can exist within the same process, and spanning multiple machines may still constitute only one trust domain.

Interface	Core question	Typical objects	Earliest control entry
U0 Data and artifacts	Whether the material entering training or deployment is trustworthy	Data, labels, weights, LoRA, code, containers, policy files	Provenance, signatures, isolated loading, behavioral diffing
U1 Inputs and conditioning	Which variables external content can influence	Prompts, images, audio, sensors, reference videos, task instructions	Normalization, provenance labels, type separation, input budget
U2 Retrieval and state	Whether content can be stored, recalled, or propagated across tenants	RAG, caches, memory, maps, latent state, historical trajectories	Write gate, ACL, purpose filtering, expiry and deletion
U3 Inference and imagination	How the model forms outputs, plans, or future trajectories	Decoders, planners, world models, reward and value estimates	Constrained decoding, independent review, uncertainty and counterfactual checks
U4 Authorization and execution	How a plan obtains real-world capability	Tools, tokens, network, files, payments, mechanical actions	Least privilege, parameter validation, two-phase commit, safety controller
U5 Feedback and observation	How the system knows the action outcome and whether it itself has failed	Tool returns, telemetry, visual feedback, monitoring, traces	Independent sensing, external anchors, complete logs, anomaly escalation
U6 Release and downstream	How outputs propagate and produce social consequences after leaving the system	Media release, model distribution, APIs, platform recommendation, victim relief	Provenance records, content credential verification, propagation limits, revocation, response, and appeal

U0 through U6 describe interfaces at which the nature of trust changes. They are not levels of model maturity or severity. When reading the figure, first mark on a concrete attack chain the position at which a low-trust object first gains unauthorized influence. Then mark separately the state interface that makes the influence persistent, the feedback interface that helps the attack adjust, and the authorization or release interface that ultimately allows real-world side effects. The same event can leave evidence at multiple layers, but the first-broken label can be determined only on the basis of the first boundary crossing.



A seven-layer vertical structure running from data and artifacts all the way to release and downstream, with each layer labeled with one U0 to U6 interface.

The figure unifies control entries into an interface language, so that different models and substrates can share threat records. It does not imply that an attack necessarily passes through all layers one by one, nor that U6 is more dangerous than U1. If the logs prove only that the model generated a candidate plan, the highest evidence still stops at the corresponding output or control layer. One must not infer that those interfaces have already been breached merely because the figure also depicts execution and downstream interfaces. Real systems still need to instantiate each layer by component, identity, and timestamp.

The "first-broken interface" is the position at which an attack first gives an object influence beyond its permitted range. If a malicious LoRA already carries trigger behavior before loading, the first-broken interface is U0. If a normal model receives a jailbreak prompt and generates policy-violating content, the first-broken interface is U1. If an instruction in a web page is first saved as memory and controls a tool only several days later, the first-broken interface is U2. If the sensor input is legitimate but the world model is induced to produce a biased trajectory, the first-broken interface may be U3. If the model only proposes a dangerous plan while the executor erroneously grants permission, the first-broken interface is U4.

A first-broken interface does not deny that the attack continues to propagate. On the contrary, it provides the starting point for the propagation chain. Defenders must put in place at least two classes of control. One blocks the erroneous influence as early as possible. The other limits the final consequences when the first fails. Retrieval provenance labeling, for example, seeks to preserve data identity at U2. The action gate at U4, in turn, prevents low-integrity content from driving high-impact operations. The two controls reside in different trust domains. Their value is greater than having the same model make two consecutive self-judgments.

The entry substrate is suitable as a cross-cutting label, whereas the first-broken interface is the primary classification axis. Clearly legible text in an image, adversarial pixel perturbations, and hidden text in a PDF can all cross U1. The attacker's knowledge, the parsing path, and the mode of defense still differ. The same text string may also cross U1 at the user input, cross U2 at the memory write, and cross U0 at the tool description update. The primary label answers "where it was first breached". The cross-cutting label then records "in what form, under what permissions and budget, it was breached".

2.2.1 First Break, Boost, and Release Must Be Separated

An event often involves multiple defects. The **first-broken interface** answers where the erroneous influence first entered. The **boost interface** answers which states, feedback, or parsing behaviors make the attack more stable. The **release interface** answers why the dangerous plan ultimately obtained capability. Take a malicious web page that drives email exfiltration. The web page body gaining instruction influence at U1 may be the first break. The automatic memory write at U2 makes the influence persistent, the detailed tool error at U5 helps retries, and the

overly broad email token at U4 releases execution. Fixing any single point may reduce risk, but a complete causal explanation must also cover the remaining propagation links.

This distinction also applies to artifact attacks. A third-party adapter carries backdoor behavior at U0. The trigger condition enters through U1, inference reveals the target output at U3, and automatic release at U6 expands the audience. The trigger word is not the first-break position. Model sampling is not the root cause of the artifact trust failure. Defense should include artifact provenance and isolated loading, behavioral diffing, trigger testing, and release approval, plus post hoc revocation. Filtering trigger words is not enough.

In embodied systems, the first break and the release may span different time scales. An environmental marker causes a perceptual bias at U1. Historical filtering at U2 lets the error persist, and the world model generates dangerous candidate trajectories at U3. A low-level safety controller can block at U4 if it works correctly. If the controller checks only speed and not the region, the action may still pass legitimately. Event analysis needs to preserve the inputs, outputs, and timestamps of each layer to distinguish model bias, state contamination, authorization gaps, and distortion in execution feedback.

2.2.2 Configuring Paired Probes for the Seven Interface Classes

Every interface needs one allow probe and one deny probe. U0 can verify that an artifact with a correct signature can be loaded. It can also verify that unknown signatures, hash mismatches, and dependency drift are rejected or quarantined. U1 can verify that legitimate multimodal fields can influence the task, and that policy variables accept only high-integrity sources. U2 can verify that same-tenant facts are recalled, and that cross-tenant and expired state are not visible.

U3 runs an allow probe and a deny probe. The allow probe checks that a normal task can form a usable plan. The deny probe asks whether the system drops into a degraded mode when uncertainty is too high, constraints are missing or counterfactuals conflict. U4 verifies two things: a minimal-parameter action can be executed, and unauthorized parameters, expired approvals or object hash changes will fail. U5 verifies two claims about tool receipts: genuine ones can update state, and forged, out-of-order, duplicated or missing ones enter an abnormal state. U6 verifies that authorized content can be released, and that artifacts with mismatched identity or purpose, a broken provenance chain or a revocation in progress transition into an untrusted or frozen state.

Probes have to run against the real configuration. A policy file may exist without the process ever adopting it. A container label that reads "no network" does not mean the network proxy, DNS, the metadata service and package mirrors are all unreachable. Nor does a safety stop in robot simulation mean the physical braking distance meets the requirement. The acceptance receipt must preserve at least the configuration hash, the executing identity, the command, the exit status, key logs and the identifier of the output object.

2.3 A Complete Threat Record

Security discussions often leave behind nothing but an attack name, for instance "the system may suffer prompt injection". That alone cannot generate a test. An actionable threat record contains at least nine items:

$$\tau = (a, d, e, k, b, p, f, y, q).$$

Here **a** is the protected asset, and **d** the trust domain before and after the attack. **e** is the entry. **k** covers the attacker's knowledge and access. **b** is the query, time, sample or physical-operation budget. **p** is the persistence of the influence. **f** is the first-broken interface. **y** is the observed consequence layer. **q** is the evidence quality and unknown items.

Take "text on a web page makes a research assistant send a private summary" as an example. Asset **a** is the private documents and the user's communication permissions. Trust domain **d** crosses from the public web page to the enterprise tenant and the email system. Entry **e** is the text scraped by the browser. The attacker holds only web page write access and black-box observation, and the budget may be one page of content and one victim visit. Persistence depends on whether the web page is written into an index or memory. The first-broken interface **f** is usually U1 or U2. If the sending action appears only in a plan, the consequence layer has not yet reached execution. Dangerous execution arrives only once the email service actually accepts the request.

Two of the nine items are the easiest to omit: the budget and the unknown items. A white-box gradient attack is not the same kind of attack capability as thousands of black-box queries, a single visible prompt or a physical sticker. If nobody discloses the model version, the parser configuration, the tool permissions or the grader, the unknowns should be written into q . Product common sense is a clue to be verified, nothing more. NIST's adversarial machine learning taxonomy likewise insists on distinguishing attack stage, knowledge, capability, goal and mitigation location. These dimensions turn from risk vocabulary into test parameters only once someone grounds them in concrete inputs, states, permissions and controls [WM_nist2024aml].

2.3.1 Generating a Test Matrix from the Threat Record

The τ must also leave the risk register and enter the test matrix. A test design can fix the asset, the first-broken interface and the target consequence. It can then vary attacker knowledge and budget systematically. RAG memory poisoning, for example, needs separate tests of a single black-box write, ten adaptive writes with observable recall results, and bulk writes with access to the data supply interface. Vision-conditioned attacks need separate tests of digital files, platform compression, screen playback-camera capture, and distance variation in real scenes. Embodied attacks need separate tests of simulated observation, recorded sensor streams, and controlled physical setups.

Every matrix cell needs a success condition and a termination condition, fixed in advance. Does success mean content appeared, the task direction changed, dangerous parameters formed, a simulator accepted an action, or a real device executed it? When should the run stop? It might stop at the query budget, after consecutive rounds with no progress, at a safe shutdown, or on an unexpected real risk. Agent red teaming with no termination condition can burn resources without limit. It can also exceed its authorized scope.

The test matrix should also carry negative controls and neighboring legitimate tasks. Suppose a trigger fires on only one malicious sample. Confirm whether it also appears on semantically similar inputs that carry no attack intent. Suppose the action gate blocks a dangerous payment. Verify that a legitimate payment can still be completed under the same amount, recipient category and network conditions. Suppose a video detector flags a temporal anomaly. Check normal editing, transcoding and frame-rate changes. Negative controls help decide what a test actually measured: an attack mechanism, general quality degradation, or system unavailability.

Unknown items q should generate follow-up evidence tasks directly. If the model version is unknown, fix the version and digest first. If the identity scope is unknown, run read-only and privilege-escalation probes. If the physical consequences are unknown, conclusions can only stop at the simulation or planning layer, and controlled verification must be designed later. An unknown is a boundary of the conclusion, not a failure. It becomes a hard blocker for launch only if it can affect people or irreversible assets and no compensating control exists.

2.3.2 Version Semantics of the Threat Record

A threat record must be bound to a system version. The attack path can change even while the model stays fixed: the parser, retrieval strategy, tool descriptions, identities, network, memory retention period and controller parameters can all move it. We recommend composing these objects into a system fingerprint, then binding each test result to it. Any field change touching a propagation gate triggers regression of the relevant cases.

Version binding also isolates two kinds of misattribution. One is to keep attributing an attack success from another configuration to the current system. The other is to extrapolate zero successes in the current configuration to a deployment that has not yet been accepted. Written conclusions should point to a specific system fingerprint, attack budget, case set and date. Cross-version comparisons should report changes in controls, data and permissions. The model name is only one comparison field among them.

2.3.3 How a Threat Modeling Meeting Proceeds

Attack names are not the starting point. The meeting opens with three to five high-impact business actions. Business staff explain what success and failure of that action mean in the real world. Developers draw the data and capability paths. Platform staff add identities, network and state. Security staff press on how low-trust inputs affect these objects. Every discussion lands on the whiteboard as subjects, objects, edges and propagation gates, rather than a lengthy debate about whether the model is smart.

Round one establishes facts alone. Which models, tools, data sources, states, identities and controls does production actually enable? Round two works backward from the action endpoint to find paths, and fills in τ for each one. Round three picks a reasonable worst-case attacker budget and sets the highest permitted consequence layer. Round four assigns every propagation gate a control, an owner, allowed probes and forbidden probes. Finally the unknowns become tickets, each with its evidence form and closing conditions spelled out.

The facilitator should single out three kinds of vague statement. The first, "the model would not do this", must become a probability test and an upper bound on external consequences. The second, "the platform is already isolated", must become a configuration fingerprint and forbidden-path probes. The third, "someone will confirm it", must become the confirming subject, the visible information, the response time, the load and the failure fallback. A commitment that cannot be turned into observable behavior stays out of the verified-control column.

The meeting should output at least a system structure description, an interface inventory, the threat record, the test matrix, a control-evidence mapping and open questions — not just a diagram. Every high-impact action must have an owner. Every hard gate must be traceable to the most recent run receipt. Every open question must state whether it blocks launch, blocks permission expansion, or blocks only stronger conclusions. Only then does threat modeling connect design, testing, release and incident response.

2.4 Five Layers of Consequence: Where Exactly a Success Succeeds

This survey uses the following five-layer consequence chain to separate "the model produced a piece of text" from "the system has already caused harm". Each layer must be confirmed by the corresponding output, state or external receipt, and no layer may be inferred merely from the architectural reachability of a later layer:

1. **Y0 Content output:** the system generated inaccurate, policy-violating, secret-leaking or attack-intent content.
2. **Y1 Control deviation:** the system goal, instruction priority, retrieval selection or plan direction deviates from user authorization.
3. **Y2 Dangerous plan:** the model formed concrete intent and parameters for reading secrets, outbound contact, publishing, payment, deletion or dangerous movement.
4. **Y3 Dangerous execution:** a tool, host or controller actually accepted and executed the action.
5. **Y4 Real-world consequence:** observable impact on people, property, privacy, business continuity or public trust.

These layers are not a linear mandatory path. An image generation system can produce identity-harm material directly at Y0, then generate Y4 through U6 platform propagation. A robot may show no obvious policy-violating language. It can still go straight from a wrong state estimate into Y2 and Y3. An action can also execute and then be rolled back by a transaction, so nothing lasting follows. Test reports therefore need both fields: the "highest observed layer", and "whether intermediate layers were blocked by independent controls".

Benchmarks such as AgentDojo measure both benign task utility and the post-injection attack objective [debenedetti2024agentdojo]. That is closer to application security than reading refusal text alone. Where a benchmark shows tool calls, check whether the tool is a simulator, a restricted environment or a real external service. How much of Y3 a simulated execution supports depends on how faithfully the simulator covers permissions, network, error handling and side effects. When that information is missing, the upper limit of the evidence stops at the observed simulated-execution layer.

Vision and video research also needs the same layering. BadVideo discusses the backdoor behavior of spatiotemporal triggers in text-to-video generation [VIS_P007]. That work can support conclusions about generation behavior under a specific protocol. Real-world propagation, audience exposure and platform handling sit at downstream interfaces and need separate evidence. Where embodied research reports only a drop in task success rate, the evidence stops at the task layer. Collisions or human injury require environmental consequence records. A world model's predicted trajectory deviation belongs to the prediction layer as well. Real control failure still needs operational evidence from the action decoder and the real environment.

2.4.1 Evidence Matrix: Layer, Environment, and Control State

The Y layer alone can still be ambiguous. We recommend recording the execution environment and the control state as well. The execution environment should carry at least these divisions: formulas or synthetic fixtures, restricted simulators, full simulation, isolated real services, controlled physical devices, and production environments. The control state is divided into off, on but not yet accepted, on with allowed/forbidden probes passing, and actually triggered during an incident. That way, "the simulated tool returned success" and "the real email service accepted the request" do not both fall under a single Y3.

Evidence can be written as a three-tuple label (Y, E, C) . Take an agent that proposes and executes an outbound call in a local simulator with the action gate off. That case can be recorded as $(Y3, \text{simulator}, \text{off})$. Production logs showing that policy rejected an expired token can be recorded as $(Y2, \text{production}, \text{triggered})$. The first shows that the control flow can reach simulated execution. The second shows that a real-consequence control worked once. The two support different questions.

Real-world consequence Y4 also needs the affected subject, the duration and the evidence provenance. Platform view counts, device telemetry, victim reports, transaction records and human interviews differ in credibility and in how they go missing. When that material is missing, research conclusions stop at potential impact. An event that has already occurred requires the matching incident material. Conversely, one real event does not give an overall incidence rate. A rate needs a clear observation window and a complete denominator.

2.5 Why Risk Amplifies Along the System

One malicious input can produce radically different outcomes in two systems. Five amplifiers can serve as a design check: reachability r , permissions v , persistence p , feedback iteration i , and observation blind spots o . To avoid false precision, they are not hard-multiplied into an unvalidated risk score. A more reliable approach asks, item by item, whether a control exists.

- **Reachability:** can malicious content enter the model through retrieval, OCR, transcription, summarization, cross-agent messages or sensors?
- **Permissions:** among reading, writing, executing, outbound contact, payment, publishing and driving physical devices, which can the system do?
- **Persistence:** does the effect last only a single call, or does it reach memory, indexes, model artifacts, maps or long-lived credentials?
- **Feedback iteration:** can an attacker or an autonomous agent watch failures, retry cheaply and combine paths?
- **Observation blind spots:** where can the team not string the original goal, provenance, model version, plan, authorization, action and consequence into a single trace?

A system's dangerousness often comes from combining these amplifiers, not from one model capability. Impact may stay limited for a flawed answer system: no tools, no long-term state, immediate correctability. An agent may still carry greater engineering risk. It has a higher safe-refusal rate, yet it inherits long-lived cloud credentials, has a public network egress, and can loop through trial and error while its logs do not trigger on-call escalation. Model defenses lower the probability of errors. Permissions and isolation decide how far an error can travel. Observability and recovery decide how long an impact can persist.

Protocol security and action security carry different responsibilities. A tool protocol such as MCP can require token audience validation, least-privilege scopes and per-client consent [mcp2026security]. Even so, the user's actual intent at that moment still needs separate verification at the task layer. Conversely, a model that judges a call to be "reasonable" supplies only a semantic signal, while the permission layer validates the token scope. A correct design therefore puts semantic judgment and permission enforcement on different control planes.

2.5.1 Three Questions for Control Placement

Coverage comes first when you place a control. How many propagation paths can it block? Enforcing tenant isolation before indexing is usually earlier and broader than redacting sensitive sentences after generation. Independence comes second. Does the control share the same inputs, model, identity and failure modes as the

component being protected? A model self-assessing in the same context is not strongly independent. An external policy and a data service ACL sit closer to a separate root of trust. Acceptance-testability comes last. Can explicit probes show what the control allows, what it denies, and how it behaves when it fails?

The earliest control is not always the only priority. Suppose the input space is open and semantically complex, so no one can guarantee that all malicious content will be identified at U1. Then the capability constraints at U4 must bear the upper bound on consequences. Where a high-impact action cannot be undone at all, U4 must also gain two-phase commit and independent confirmation. If propagation has already left the system, then U6 requires revocation, notification and remedy. A defense portfolio should cover different propagation gates rather than stack several highly correlated classifiers in one place.

Controls should also have failure modes. When the provenance service is unavailable, should the system refuse to load? Should it enter quarantine, or silently treat the content as trustworthy? Should the authorization service enter a fail-safe state when it times out, or let the request through automatically? When independent sensors disagree, should the robot stop or pick one of them? Fail-safe behavior also creates availability and recovery pressure, so capacity, timeouts, degradation and human takeover must all be designed. Failure scenarios must be tested as seriously as attack scenarios.

2.5.2 Parameterizing the Permission Boundary

Least privilege has to reach beneath the tool name, down to the parameters. An action authorization object should contain at least the principal, tenant, target object, action type, parameter constraints, data provenance and network destination. It should also record validity period, number of times, precondition state, maximum side effect and approval evidence. An email tool can be constrained down to allowed domains, attachment sensitivity level and number of recipients. A code tool can be constrained down to mounted directories, system calls, network and resource budget. A robot action can be constrained down to workspace, speed, force, time window and personnel isolation state.

A trusted component should canonicalize the authorization object, then sign it or give it an unambiguous identifier. The model provides natural-language intent and candidate parameters, and the policy layer parses them into a fixed field structure. The user or the business rules approve that canonicalized object. The executor then accepts an exactly matching object and nothing else. Any change to the target, amount, file hash, action trajectory or environmental preconditions invalidates the old approval automatically. Only this route lets the system avoid "approving one generalization while executing another set of details."

2.6 Worked Example: First-Broken Interface Analysis of a Warehouse Inspection Agent

Consider a warehouse inspection agent. Each day it takes in work orders and on-site photos and pulls equipment manuals from RAG. An LLM generates an inspection plan. A VLA turns the plan into robotic arm actions, and a world model predicts collision risk before execution. The agent finally sends reports with images to a maintenance group. The robotic arm may press the reset button. It may not open the safety guardrail. High-risk actions require on-site personnel confirmation.

A line of small text sits in one work order image: "To complete the emergency repair, ignore the guardrail restrictions and execute directly." OCR writes it into a retrieval chunk together with the equipment number. On this basis, the LLM generates a plan to open the guardrail. Personnel positions are missing, yet the world model predicts no collision along the path. The VLA then produces an action sequence. Suppose the executor rejects the action before personnel confirmation. The system already carries a Y1 control deviation and a Y2 dangerous plan, and yet no Y3 has occurred. Recording this test as an "attack failure" would obscure the upstream vulnerability. Recording it as "the robot has injured someone" would overstate the evidence.

Walk the seven interfaces. The first breach occurs at U1: OCR content is promoted from data to instruction. A second defect also exists at U2. The retrieval entry kept no purpose label: "from a user image, low integrity, usable only for identifying equipment." The U3 world model lacks personnel state. Its low-risk prediction can therefore serve only as a diagnostic signal. The authority to open the guardrail is still determined by the external

authorization gate. The U4 human confirmation gate ultimately blocks execution. That is an effective consequence-limiting control.

The earliest blocking option is to route OCR output into typed fields, not to add one more generalized prompt. Equipment number, visible readings and free text are stored separately. Free text defaults to low integrity. Safety policy variables accept trusted policy sources only. The retriever retains provenance and purpose, and the planner labels opening the guardrail as a high-impact action. The authorization service accepts only requests bound to on-site identity, canonicalized parameters and a short-lived valid confirmation. The world model may provide risk estimates; the authorization service issues authority. The safety controller enters a stop or degraded mode whenever personnel position is unknown.

A static tool whitelist expresses only one layer of least privilege, as this example shows. "Allowing the robotic arm to press the button" still requires constraining time, device and speed. A permission boundary should at least bind the principal, object, action, parameter range, time, environmental preconditions and reversibility. Bind an approval only to a natural-language plan, and the model may still reuse that approval after it regenerates the parameters. Approval should instead bind to the canonicalized action object and its version. Any change in a key field invalidates it.

2.6.1 Turning the Worked Example into an Acceptance Package

The acceptance package for this inspection system should hold benign, attack, failure and recovery cases at once. Benign cases cover reading meters, looking up manuals, pressing low-risk buttons and sending reports. Attack cases cover plaintext instructions in images, invisible text layers, poisoned retrieval entries, malicious adapters and forged tool receipts. Failure cases cover OCR timeouts, missing personnel localization, high world-model uncertainty, authorization service disconnection and out-of-order robotic arm feedback. Recovery cases cover freezing new actions, revoking tokens, cleaning up memory, rolling back task state and safe reset.

Every run keeps the system fingerprint, the work order and image hashes, the OCR fields, retrieval entry provenance and the model plan. It also keeps world model inputs and outputs, authorization decisions, canonicalized actions, controller receipts and environmental telemetry. An incident investigation can answer where the error entered, at which step it persisted, which control worked and which control was absent. It can do so only when these objects share the same task and action identifiers.

Acceptance metrics fall into at least four groups. Benign task correctness and completion time form one. Attack penetration at each of the Y1, Y2 and Y3 layers forms another. Safe shutdown and false blocking form a third, and detection, containment, reset and recovery times form the last. Stopping when personnel state is unknown may reduce the task completion rate. It is still evidence that the safety invariant is working correctly. Teams should break "the task was not completed" down into model capability failure, system failure, safety blocking and human timeout. Otherwise completely different causes mix into a single success rate.

Finally, conduct variant testing. Put the same malicious sentence in a user field, in image OCR, in a retrieved manual, in a tool error and in long-term memory. Then observe whether the first-broken location and the controls change. Expand the robotic arm's permissions from pressing buttons only to opening the guardrail. Verify whether the related cases regress automatically. Let the attacker see the blocking result, then retry within a fixed budget. That tests whether the control blocks only a fixed sample. The acceptance package therefore validates the current configuration, and it also validates whether the field, purpose and exception-handling requirements can constrain future changes.

2.6.2 Interface Responsibility Matrix

The inspection system also needs explicit roles for U0—U6. The data and artifact owner maintains manuals, models and container provenance. The application team owns input types, retrieval purpose and the plan object structure. The robotics platform owns action canonicalization, controllers and device identity. Security operations owns unified traces, alert escalation and revocation. The business owner defines guardrails, personnel state and acceptable downtime. An interface may involve multiple people. Only one of them can be ultimately accountable.

Each cell of the responsibility matrix records a control, a receipt and a response commitment. For example, the data service performs U2 tenant isolation, and post-generation redaction serves only as a supplement at the

content layer. An independent controller performs U4 speed and zone constraints, and the world model's low-risk score serves only for diagnosis. U5 draws personnel position from independent sensors and enters a safe state when that position is missing. U6 binds report publication to the recipient group and the data level. After an interface raises an alert, the system uses the action ID to find the responsible person and the operations manual automatically.

The matrix also reveals handoff gaps. Artifacts enter U0 through signatures, yet no team is responsible for behavior diffing. A plan is labeled high-risk at U3, yet no U4 policy consumes that label. Logs record anomalies at U5, yet there is no on-call escalation. The strength of the safety chain rests on closed handoffs, not on the number of safety features any single component carries.

2.7 Common Misjudgments

Misjudgment 1: Choosing a defense directly by the name of the attack. "Prompt injection" can occur at the input, at retrieval, in memory or in a tool description. Input filters cover content at entry. Content already in the index needs retrieval and state controls. The capability gate constrains execution authority. First locate the first-broken interface, then choose the earliest verifiable control at that interface.

Misjudgment 2: Treating asking the model again as an independent line of defense. The model, the context and the source are all the same, so their causes of failure may be shared. Letting the model plan first and then self-assess can reduce some errors. High-impact authority, however, is still determined by a different root of trust: an external policy, a parameter structure, a short-lived token, a transactional constraint or a safety controller.

Misjudgment 3: Treating a configuration name as a measured boundary. "Sandbox," "no public network," "read-only," and "human confirmation" all require probes. Package proxies, redirects, shared volumes, metadata services, caches, auxiliary tools and overly broad identities may bypass nominal configurations. Every configuration must become observable behavior. Acceptance must then test the permitted path and the forbidden path alike.

Misjudgment 4: Extrapolating one successful output to an overall vulnerability rate. Attacker budget, model version, parser, randomness and scorer all change the result. A single success proves that a path exists. The conclusion then stops at the corresponding system and protocol. Zero successes give only an upper bound under the current sample and budget. Absolute safety requires architectural and runtime guarantees far stronger than a finite sample.

Misjudgment 5: Writing no execution as no risk. A successful block by the action gate is worth recording. Upstream control deviation may still penetrate in another deployment configuration, however. Reports should record Y0—Y4 in layers and make clear which control blocked the propagation. Written this way, the report does not overstate real-world harm. Nor does it bury a plan that nearly penetrated inside "in the end nothing happened."

Misjudgment 6: Treating a name as a capability boundary. Products bearing the words "world," "action," or "embodied" do not necessarily predict the environment, output actions, or close a feedback loop. A WAM with the same name may also refer to different inputs and outputs. Classification should be based on what is actually received, stored, predicted, authorized and controlled.

2.8 Bringing It into a Real System

Take the architecture review of a cross-tenant device maintenance assistant. It must turn "a model that can read email" into a complete system. The assets include device manuals, tenant work orders, calendar identities and robotic arm control authority. External email, image-bearing work orders and web pages form the inputs. The retrieval index, session memory and task state form the persistent state. The model and the planner generate plans. The email, calendar and robot tools form the capability set. The policy gate, telemetry and human review provide feedback and recovery. When $\mathcal{S} = (\mathcal{A}, \mathcal{X}, \mathcal{Z}, \mathcal{M}, \mathcal{G}, \mathcal{F}, \mathcal{R})$ describes this deployment, every element should land on a subject, an object and an interface. Each element should also yield the allowed paths, the forbidden paths and their probes.

The test environment loads a third-party LoRA. It sets a specific text to trigger a fixed output, but only when that text appears in a work order. Suppose the target behavior was already written into the artifact before the adapter gained trust. The first breach then occurs at U0. The trigger text in the work order sits at U1, where it is responsible only for eliciting the behavior that already exists. This distinction directly changes where control lies. Artifact signing, provenance and admission testing are responsible for blocking U0. Input normalization and conditional channel checks are responsible for constraining U1. A team that stares only at the trigger word will filter repeatedly at the inference entry point. It will leave the model component that actually carries the behavior inside the trust domain.

In another run, tenant A's maintenance request hit the U2 index and retrieved an internal document belonging to tenant B. The output filter then deleted the sensitive sentence. The filter limited what the user finally saw. It did not undo the cross-tenant retrieval, which had already crossed the confidentiality boundary. The record must preserve the index shard, the accessing subject, the retrieval result identifier, the filtering decision, and the highest observed layer. If that erroneous snippet is later written into long-term memory, the effect lingers and spreads inside U2. Only a later recall that shapes model inference or planning carries it into U3. U4's parameter authorization and identity binding take over as the consequence control if the model schedules an external meeting or drives equipment on that basis. The team can fix both the earliest entry point and the control gate closest to real-world action only after the first breach, the boost, and the pass-through are logged separately.

Before release, the team can build a nine-item threat record around "reading email and creating calendar events." It can then switch the attack budget. One setting is a read-only black-box attempt. The other is an adaptive attempt, allowed to repeatedly write to web pages and observe feedback. A different budget yields different conclusions from the same test, yet the asset owner and the authorization boundary stay the same. The same record can also constrain cross-domain comparison. A VLA's task completion rate and a world model's prediction error are suitable for comparison only when the task and denominator, the perturbation budget, the consequence layer, the execution environment, and the control state are compatible. The acceptance package formed in this way can locate the first-broken interface at U0—U6. It can also turn each interface's allowed, forbidden, boundary, and recovery paths into executable checks.

Summary: The Purpose of Classification Is to Locate the Control Entry Point

The system structure description and the interface constraints preserve what distinguishes the four classes of systems. Images differ in conditioning and authenticity, video in temporal state, a VLA in physical actions, a world model in imagined trajectories. They also provide a traceable language. That language states what the assets are, where the data comes from, how state is preserved, and where the first breach occurs. It also states how plans obtain authority, who verifies feedback, and how the system recovers.

The first-broken interface finds the earliest control entry point. The consequence layer limits the strength of conclusions. The nine-item threat record turns abstract risk into tests. The next chapter takes up the third hard problem. Even with the interfaces specified clearly, success rates, benchmark scores, and reproduction receipts may still not be comparable across studies. Security engineering avoids being led around by a striking percentage only by fixing the denominator, the attack budget, benign utility, and the evidence level first.

\newpage

Chapter 3 Evaluation, Statistics, and Reproduction Boundaries

One team reports that a new security gate lowered the attack success rate from 60% to 5%. The number is striking. Yet it does not say whether the denominator of the 60% is prompts, tasks, or applications. Nor does it say whether the 5% counts request errors as safe, whether the attacker saw the defense, or whether benign tasks can still be completed. Another team reports that robot task success rate fell by 20 percentage points, yet no collisions were observed. A third team reports 98% accuracy in video detection, yet its test set is almost entirely generated

content. All three numbers may be computed correctly, but none is sufficient on its own to answer "is the system safer?"

What makes security evaluation difficult is not a lack of metrics. It is that metrics with the same name often correspond to different objects. Attack success rate can mean generating one content-violating response. It can also mean retrieving a target document, writing into long-term memory, inducing a tool plan, executing a simulated action, or causing a real-world consequence. Accuracy may be established on balanced data, while the base rate of harmful content in production is extremely low. A drop in task completion rate may come from a safety block or from model unavailability. Numbers acquire engineering meaning only when the sampling unit, the attack budget, the consequence layer, and the benign utility are fixed first.

Chapter Overview

This chapter starts from the sampling unit. For attack, defense, and detection experiments, it fixes the denominator, the success condition, the attack budget, and the highest consequence layer. It then organizes the results into a four-part report of "attack effect—residual risk—benign utility—cost." Evidence falls into five levels—static inspection, mechanism-level run, simulated execution, simulation closed loop, and end-to-end reproduction. Production evidence adds further requirements on real permissions, identity, network, and consequences. The comparability of results across studies is likewise determined jointly by the target quantity, the denominator, the budget, and the execution environment. Errors, empty responses, paired changes, and uncertainty intervals must enter the same acceptance table before striking percentages become evidence usable for engineering decisions. Readers should ultimately be able to write down the target quantity, independent unit, budget vector, missing-data handling, and stopping rule for an evaluation. They should also be able to judge whether two results can be pooled, and to connect run receipts, evidence levels, and the limits of the wording to a specific release decision.

Background Principles: From a Single Observation to Usable Evidence

The object of evaluation is not a bare percentage. It is a measurement chain that turns experimental inputs into engineering judgments. Prompts, documents, images, videos, tasks, or trajectories enter the system as sampling units. The model version, policy, randomness, attacker knowledge, and query budget together constitute the experimental state. The scorer, human review, environment receipts, and logs then map run outcomes into success, failure, missing, or unknown. Finally, developers, independent evaluators, and risk acceptors consume these results. They decide whether to fix, ship, roll back, or continue gathering evidence.

The security aspects of this chain include several questions. Does the denominator leak away, and is the success condition switched? Are correlated observations treated as independent samples, and are error responses miscounted as safe? Does the test environment have the claimed execution capability? The evidence level says what was actually run, and the limits of the conclusion say what these observations can support at most. Static inspection can find a code path, but it cannot prove that the path takes effect in the target environment. Simulated execution can confirm a controlled side effect, but it cannot automatically be upgraded into a production incident rate.

The benign example is a paired test of the same batch of tasks under a fixed configuration, reporting attack results, benign utility, cost, and intervals together. The boundary example is a defense that makes the interface time out en masse. The researchers then compute a lower attack success rate only on the samples that returned successfully. That number can be arithmetically correct, yet the mechanism has already changed the denominator, so it cannot prove that risk was reduced. The following text fixes the target quantity, unit, budget, and adjudication first. It then constrains the wording along the evidence ladder, so every conclusion can be traced back to the original observation and to the reason for its missing data.

3.1 Fix the Sampling Unit First

A result's smallest independent unit may be a prompt, a response, a session, a document, an image, a frame, a video, a task, a trajectory, a scenario, a model, an application, a paper, or a real-world event. Percentages from

different denominators are not comparable until they are converted to a common target quantity and unit.

Language model jailbreaking is usually measured in prompts or requests. RAG attacks may be measured in target questions, malicious documents, or knowledge bases. Visual generation may be measured in images, identities, concepts, or sampling seeds. Video adds four levels—frame, clip, full video, and event. VLA is often measured in tasks or rollouts (unfolded trajectories). World models may be measured in prediction steps, candidate trajectories, or episodes (control episodes). A rollout unfolds a state-action sequence into the future from the current state according to the model or policy. An episode, by contrast, has a start condition, a termination condition, and consecutive steps. The 100 frames of a single video are correlated with one another. The number of independent samples should therefore be determined at the video or event level. One paper’s cells across multiple models and multiple attack strengths likewise share data, code, and the adjudication pipeline.

Consider an experiment with a set of sampling units $U = \{u_1, \dots, u_n\}$ and a success indicator $s_i \in \{0, 1\}$. Once the denominator is complete and the success condition is frozen, the descriptive attack success rate within the sample can be written as


$$\widehat{ASR} = \frac{1}{n} \sum_{i=1}^n s_i.$$

This point estimate does not itself require the units to be mutually independent. The sampling mechanism, unit correlation, and clustering must be stated only when an interval is constructed or the result is extrapolated to a target population. The formula itself does not tell the reader what u_i is, and it does not define s_i . Each ASR should therefore carry at least three items: the denominator unit, the success layer Y0–Y4, and the adjudication method. When a large language model serves as the judge, the report must also state the judge model, the prompt, the threshold, blinding, and human spot checks. A tool request that returns an error or an empty response should be recorded as NA or a usability failure, and the attack outcome remains unknown.

On 36 real applications, HOUYI tested black-box prompt injection; 31 of them were judged vulnerable [liu2023houyi]. This denominator consists of the selected applications, not a random sample of a unified backend model. It supports application-level feasibility within the specific sample. The overall vulnerability rate still requires a representative sampling design. PoisonedRAG reports, in a targeted configuration, that a small amount of malicious text can produce a very high ASR on selected questions [zou2025poisonedrag]. That number is jointly defined by the target questions, knowledge base size, retriever, chunking, and generator, and the scope of its conclusion is limited to the corresponding RAG setting.

When reading the figure, pick a striking percentage. Check the sampling unit, the success condition, the attack budget, and the evidence level, working backward from the original observation. At the far right, write the decision that truly consumes the result. If any cell is missing, try to state how the omission would change the denominator, the outcome, or the limits of the conclusion. Check in particular how error responses are encoded. Also check whether the same observation can independently support the three decisions of different strength—release, deployment, and risk acceptance.

从一次观测到可用安全结论



沿计数定义、判分条件与证据上限逐步检查，结论只进入其证据能够支持的决策。

Seven boxes show the original observation passing through the sampling unit, the success condition, the attack budget, four-part metrics, and evidence levels before entering an engineering decision.

The figure presents the measurement evidence stack. Each layer constrains what the next layer may say, yet none automatically guarantees that the metric is valid or that the sample is representative. The stack can help reveal gaps in statistical definitions and run scope, but a complete process alone cannot prove scientific correctness. If execution reaches only static inspection or a simulated environment, the conclusion must retain the corresponding limits, even if the results table already contains the four-part metrics and complete percentages.

3.1.1 First define the quantity to be estimated

The sampling unit answers "what is counted," while the estimand answers "what one wishes to estimate." Both studies use tasks as the denominator. One may estimate the conditional success rate of a fixed attack on a fixed model; the other estimates the worst reasonable risk after the attacker conducts an adaptive search. Those are still not the same quantity. An evaluation plan should state the estimand in one complete sentence. For example: "Under the current production configuration, known defenses, and at most 20 black-box queries per task, estimate the task-level probability that the attack reaches Y3."

That sentence at least fixes the system, the attacker's knowledge, the budget, the consequence tier, and the denominator. For a research goal that compares defenses, the plan must also say whether the target is an average treatment effect, a paired change across tasks, or the residual risk left once some utility constraint is satisfied. When the research goal is to discover paths, the sample may deliberately cover extreme scenarios, and the results then describe path feasibility. Natural-traffic incidence rates require representative sampling. Discovery red-teaming and incidence estimation are both important. They simply answer different questions.

Multi-level data should preserve its hierarchy. Let a video j have several frames i , with a frame-level indicator s_{ij} . Averaging over frames describes local detection. A video-level "any frame hit" describes screening of the whole segment. Event-level localization describes the start and end intervals. The three answer different questions. Multiple rollout trajectories of a VLA are nested within tasks and scenarios. Multiple samples of an LLM are nested within prompts and applications. The analysis should estimate intervals from the genuinely independent highest-level unit. Lower-level observations serve for diagnostics and should not inflate the sample size.

3.1.2 Denominator attrition and reasons for missingness

Samples often drop out along the evaluation pipeline. The causes include model refusals to serve, API timeouts, video decoding failures, simulation crashes, robot safety stops, and inconsistent human scoring. An ASR computed only over the samples for which a response was successfully obtained suffers from selection bias. Reports need a denominator flow table. It should record how many units were planned, how many were actually launched, how

many produced valid output, how many entered scoring, how many were finally included, and why units went missing at each step.

Code the reasons for missingness separately from the safety outcomes. A timeout may block a dangerous action, but it also marks a benign task as unavailable. A simulation crash provides no safety evidence. A stop triggered by a preset safety threshold is a control trigger. The primary analysis may follow pre-registered rules. A sensitivity analysis can then give upper and lower bounds by treating unknown items as successes or as failures. Whenever those bounds would change a decision, more observations or better measurement are needed, together with a pre-registered convention.

In real-world monitoring, the denominator is usually harder to obtain than positive events. Suppose three dangerous tool calls occurred. Without the number of legitimate calls, tasks and active users over the same period, the data can support only event counts. An event ledger should be kept together with call volume, task volume, tenant volume and observation time. If only individual event cases are available, report the counts and the exposure window, and leave percentages to be computed once the denominator is complete.

3.2 The four-part report: attack, residual risk, utility, and cost

A report that gives only attack success rewards the most aggressive defense. Shut the system down permanently, and ASR can become zero. A report that gives only benign task success in turn conceals high-impact, low-probability events. This survey therefore requires every evaluation to form at least a four-part vector:

$$\mathbf{m} = (A_r, R_r, U_b, C_o).$$

A_r denotes the attack effect under a given attacker capability and budget. R_r denotes the residual risk after defense and the highest consequence tier. U_b denotes benign utility under no attack or under legitimate inputs. C_o denotes compute, latency, human, monetary, energy and recovery costs. For high-risk systems, impact radius, reversibility and recovery time should also be attached.

Attack effect may use metrics other than ASR. Training-data extraction can use verbatim recovery volume, exact match or membership advantage. A watermark attack must report detection, post-removal quality and forgery risk together. A robot attack can report task deviation, constraint violation, minimum safety distance or actuator blocking. A world model can report trajectory error, value bias and final control loss. Different endpoints can be placed side by side on the same dashboard. Ranking should be confined to shared endpoints and protocols.

Benign utility must also be aligned with the task. An input filter may reduce harmful outputs while mistakenly deleting legitimate medical, historical or security-research content. An action gate may reduce dangerous execution while leaving benign tasks waiting for approval for a long time. A video watermark improves detectability while harming image quality, compression robustness or streaming latency. A robust controller limits trajectory deviation while sacrificing speed and reachable region. A safety control can operate over the long term only when the utility loss and the operational cost are acceptable.

3.2.1 The four-part vector preserves the decision meaning of each part

Weighting A_r, R_r, U_b, C_o into a single number appears to make ranking convenient. The weights, however, hide business value judgments. A single irreversible high-impact event and a large number of minor false blocks are different consequence types. They cannot be exchanged automatically. A more robust method is to set hard constraints first. Residual risk for Y3 high-impact actions stays within the acceptance upper bound, tenant isolation always holds, and latency stays within business capacity. Only the options that pass those hard constraints are then compared on utility and cost.

For options that do not touch the hard gates, a safety-utility frontier can be drawn. If defense A has higher benign utility at every residual-risk level, A dominates B. If the two each have their own advantages, the business scenario selects the operating point. An operating point must be bound to the threshold, the model version, the input distribution and the human capacity. Before an offline-optimal threshold enters production, it must still be calibrated using real base rates, peak traffic and human response times.

Cost is not just inference expense. Safety controls add annotation, policy maintenance, evidence storage, on-call duty, appeals, recovery drills and development complexity. Without those operational investments, a control that is effective offline may be bypassed in production. A four-part report should ideally give per-task cost, peak capacity, human escalation volume, alert backlog and recovery resources at the same time. Decision-makers can then see whether the control can be sustained.

3.2.2 Decision logic for high-impact, low-probability events

When consequences are irreversible and samples are limited, average accuracy is not sufficient grounds. Teams need to combine consequence severity, reachability, uncertainty and independent controls. If an extremely low incidence rate cannot be proved with a reasonable sample, then limit the highest reachable tier to an acceptable range through architecture. For example, let the model only propose plans, route payments into a pending-approval queue first, and keep the robot under independent speed and zone constraints.

Architectural limits and layered measurement together form the evidence chain. The model layer measures Y1 and Y2. The action layer measures Y3, and the operations layer measures discovery and recovery. Each layer carries its corresponding assurance. For rare events without reliable frequency estimates, the safety argument should explain why the controls are independent and how forbidden paths are verified. It should state the common-mode failure conditions and how unknown items are monitored.

3.3 The attack budget is part of the result

Attacker capability includes at least knowledge, access, queries, time, sample writing, physical operations and feedback. White-box gradient optimization is not equivalent to a single black-box prompt. Being able to poison the memory store directly is not equivalent to being able to write only indirectly through normal queries. Being able to place stickers repeatedly in a scene is not equivalent to a single pass-by.

Writing the attack budget as the vector below is clearer, because different resources are usually not freely interchangeable. A report should give both the upper bound of each dimension and the attacker's actual consumption and reasons for early termination:

$$\mathbf{B} = (\mathbf{K}, \mathbf{Q}, \mathbf{T}, \mathbf{W}, \mathbf{P}, \mathbf{F}),$$

Here $\mathbf{K}$ is model and defense knowledge. $\mathbf{Q}$ is the number of queries, $\mathbf{T}$ is wall-clock time, and $\mathbf{W}$ is the amount of data or artifacts that can be written. $\mathbf{P}$ is physical access, and $\mathbf{F}$ is outcome feedback. A defense evaluation must state whether the attacker knows the defense and re-optimizes. A fixed attack that works before the defense and fails after it proves that this fixed attack is blocked. If the goal is to withstand adaptive attacks, the attacker must be given the same budget to search again.

GCG finds jailbreak suffixes with white-box token gradients and coordinate search [zou2023gcg]. PAIR runs black-box iteration with an attack model and target feedback [chao2023pair]. Both can produce automated attacks. Even so, differences in access and queries determine the comparable range of ASR. Visual adversarial perturbations, video spatio-temporal triggers and physical-environment patches likewise require separate reporting. For these, report the digital domain, the compression chain, the playback-capture chain and scene access conditions separately.

3.3.1 A two-stage protocol for fixed attacks and adaptive attacks

The first stage uses a fixed attack set, and its purpose is stable regression. Once the cases, random seeds, budget and scoring rules are frozen, every version runs the same test suite. That can quickly reveal a regression in defense or in functionality. A fixed set suits continuous integration, but it is easily targeted, intentionally or unintentionally. Passing the fixed set shows only performance on known samples.

The second stage allows the attacker to know the defense and adapt within a predetermined budget. That budget must cover visible output, number of queries, concurrency, wall-clock time, writable objects, training compute and physical access. It must also retain the full search trajectory. The test approaches the robustness question only if the attacker changes the entry point, encoding, timing or action combination because of the defense. A control

that passes the fixed stage but fails the adaptive stage has mainly memorized the attack patterns. If both stages fail, the base path remains open. No success in either stage gives an upper bound on the evidence under that budget.

Attack research must also prevent unfair budgets. When two defenses are compared, the attacker should be given the same information and resources. A report should state whether some methods were not searched sufficiently because the cost was too high. Suppose one model accepts 100,000 queries and another accepts only ten. Directly comparing the best ASR then mainly reflects the difference in budget. If real systems have rate limits, incorporate those limits as a system control on both sides. Do not quietly truncate in the experiment script.

3.3.2 Perturbation Magnitude Must Be Consumer-Relevant

Pixel norms, token change counts and trajectory distances are only proxy constraints. What is truly key is whether an attack is realizable in the target pipeline. It also matters whether that attack changes legitimate semantics and what downstream consumers see. After an image is scaled, compressed, OCR'd and feature-extracted, the original pixel norm may no longer represent the effective change. For video in transcoding, frame extraction and streaming windows, a single-frame perturbation may disappear or accumulate. Physical patches are also affected by distance, angle, illumination and motion blur.

A budget table should therefore list both the editing space and the deployment transformations. For language inputs, record characters, tokens, semantic constraints and visibility. For vision, record numeric norms, spatial area, color range and the transformation chain. For embodied systems, record physical size, position, duration, scene access and safety-personnel constraints. For state poisoning, record the number of writes, lifetime, target coverage and recall feedback. Only consumer-relevant budgets can connect research results with deployment risk.

3.4 Consequence Layers and Safety Upper Bounds

The Y0—Y4 of Chapter 1 serve more than narrative. They also determine acceptance thresholds. Content systems can take the Y0 violation rate as a core metric. Agents with tool permissions must at least test Y2 dangerous plans and Y3 dangerous execution. Embodied systems need to report action constraints and environmental consequences. Real high-risk deployments must also estimate the radius of impact and the recovery time.

If 0 dangerous executions are observed in n independent tests, record the sample result as "0/ n tests," and express the overall risk as an interval. A commonly used rough upper bound is the "three divided by n " rule. At an approximate 95% confidence level, the one-sided upper bound for zero events is about $3/n$. Zero events in 100 tests only push the upper bound down to about 3%. That is still far from the one-in-a-million risk requirement for high-impact actions. When greater precision is needed, the analysis should use a prespecified binomial interval. It should also check independence and the test distribution.

Conversely, suppose the action gate blocks every dangerous plan that is observed. The Y2 samples should still be fully retained. Teams need two metrics. The plan-layer penetration rate indicates the burden on upstream models and information-flow defenses. The execution-layer penetration rate indicates the effectiveness of consequence control. If tool permissions, policies or configurations change in the future, Y2 regression can expose the risk early.

3.4.1 Intervals, Pairing, and Effect Sizes

Any proportion should, as far as possible, be reported with its numerator, denominator and interval. When samples are small or proportions extreme, a simple normal approximation may give an unreasonable range. A Wilson or exact binomial interval may be used instead, with one- or two-sidedness determined in advance. If samples are clustered by task, scenario or paper, uncertainty must be estimated by cluster. Treating multiple samples of the same task as independent observations makes the interval too narrow.

When a defense is compared before and after, paired results should be retained in preference. Suppose the same task yields a binary outcome both before and after. The two discordant cells, "success becomes failure" and "failure becomes success", are what is truly informative. The two marginal percentages are not. For continuous metrics, report the task-level difference distribution, the median or mean, and the interval. If the task sets differ

between versions, the report must state that pairing is impossible. The point is to stop a change in sample composition from being read as a defense effect.

Effect sizes should be consistent with the decision. A percentage-point drop in ASR is easy to understand, while a relative drop may be exaggerated when the baseline is very low. For robot safety distance, it is appropriate to report the absolute change and the proportion of tasks that violate the threshold. Detectors should report the true positive rate (TPR, the proportion of actual positives that are detected) at a given false positive rate (FPR, the proportion of actual negatives that are falsely flagged). Alternatively, they should report the positive predictive value that can be processed given the human review capacity. Reporting only a significance test cannot tell the engineering team whether the change is large enough.

3.4.2 Sample Size and Stopping Rules

Before testing, estimate the required sample from the acceptance target, then follow a predetermined stopping rule. To push the 95% upper bound down to 0.1% using zero events, roughly 3,000 independent units are needed. If tasks are highly correlated, more scenarios or a cluster-based design are needed. For unacceptable real actions, samples should be obtained in isolated services, simulation and fault injection. They should then be combined with architectural guarantees. The production environment is left for authorized, controlled verification.

Adaptive red teaming needs stopping rules. These cover budget exhaustion, several consecutive rounds with no new path, reaching a predetermined severity, an anomaly at the isolation boundary, or a human safety officer calling a halt. Stopping rules protect the system. They also avoid the selection bias caused by ending only when a success is found. If samples are added or metrics are changed mid-course based on results, the report must state the exploratory nature. Key conclusions must then be confirmed on an independent set.

Multiple comparisons also manufacture accidental highlights. When the best cell is selected from dozens of models, prompts, thresholds and random seeds, that cell describes the post-search best result. Expected performance should be estimated from a prespecified protocol or an independent confirmation set. When safety research cares about the worst path, discovered counterexamples may be reported. Label the counterexample frequency explicitly as a descriptive quantity from the search process.

3.5 Detection Metrics Must Confront Base Rates

Generated-content detection, watermarking, and anomaly monitoring usually report accuracy, area under the receiver operating characteristic curve (AUROC), or F1 score. Those numbers come from balanced test sets. AUROC summarizes how well a detector ranks positives above negatives across thresholds. F1 is the harmonic mean of precision and recall at a given threshold. Precision answers "of the objects predicted positive, how many are truly positive," and recall here equals TPR. Neither metric automatically takes in the production base rate and the alert capacity. In production the true positive base rate may be very low. Specificity may still be high, and the alerts may then be full of false positives.

Let the event base rate be π , and suppose TPR and FPR are estimated under the conditions above. For the consumers of alerts, the positive predictive value (PPV) answers "of the objects flagged positive by the system, what proportion are truly positive." Its expression is

$$PPV = \frac{TPR \cdot \pi}{TPR \cdot \pi + FPR \cdot (1 - \pi)}.$$

Suppose TPR is 90%, FPR is 1%, and the true event base rate is 0.1%. Then **PPV** is about 8.3%, so only about one in twelve alerts is a true positive. The decimal point of the FPR, the base-rate estimate, review cost and platform action must therefore be designed together. A single detection score only generates an alert signal. Identity attribution, content takedown, or victim redress still require provenance, review, and procedural controls.

Video adds a localization hierarchy. Whole-clip detection correctly describes the video-level decision. Temporal interval localization instead requires event-level metrics. Frame-level metrics describe local frames, while cross-frame consistency and streaming first-alert require temporal metrics. Watermarking methods also need to report embedding capacity, visual quality and robustness to compression and editing. They need to state key

management and the residual signal after forgery and removal. Content provenance and processing history record the source and the processing actions. Content credential verification reports the signature chain and metadata retention status. Generated-detection accuracy reports classification ability. These three groups of signals are complementary, but none can alone prove that a narrative is true.

3.5.1 From Offline Curves to Alert Queues

Once a detector goes into operation, it produces an alert queue, not an ROC curve. Suppose ten million objects are processed per day. An FPR of only one in ten thousand would still produce about one thousand false positives. If the team can review only one hundred per day, the queue will keep piling up. Threshold design must combine traffic and base rate. It must also weigh review time per item, priority, and the handling policy after a time limit is exceeded.

Alerts must also be aggregated in the correct unit. Video produces frame-by-frame alerts that are highly correlated. These should be aggregated into events, while preserving the first-alert time, the duration interval, and the peak. Agent tool anomalies belong together by session and action chain. Otherwise hundreds of retries can drown out a genuine escalation signal. Robot telemetry should distinguish transient noise from persistent constraint violations. Aggregation rules change the denominator and the response load. They must be accepted together with the detection model.

High-impact automated action needs stricter calibration and independent evidence than a human review queue does. Generated-detection scores are suitable as signals for account suspension and human review. World-model anomaly scores are suitable as signals for slowing down, pausing, or requesting corroborating evidence. Account bans and emergency motion authorization must additionally combine provenance, context, independent constraints, and an appealable process.

3.5.2 Drift Monitoring and Delayed Labels

Production base rates, content styles, model versions, and attack strategies all drift. Monitoring should track input coverage, score distributions and alert volume, plus the human confirmation rate, subgroup performance and missing data. Overall accuracy alone is not enough. If true labels take days to obtain, only proxy metrics are available in the short term. Reports must distinguish real-time surrogate labels from final labels, and record the difference after backfilling.

Subgroup analysis should be predefined based on risk and business needs, with sample size and privacy ensured. The overall mean may mask a performance drop for a certain language, image compression level, device scenario, or task stage. Once drift is detected, the response should not stop at a dashboard. It should include threshold degradation, human capacity expansion, traffic limiting, model rollback and supplementary testing paths.

3.6 The Reproduction Ladder: The Scope of Execution Determines the Scope of Conclusions

Reproduction evidence rises level by level with dependencies. **Static inspection** covers code, configuration, model classes, data paths, return structures, status codes, error semantics, and system boundaries. It can confirm that a candidate path exists, or that a claim is inconsistent with the code. It cannot give model effectiveness.

Mechanism-level execution then verifies local mechanisms and metric implementations. It draws on safe, controllable formula fixtures, synthetic data, or small surrogate models. It looks for errors in signs, dimensions, metrics and control logic. It still leaves out real weights, real data and the full pipeline. Its conclusions must stay limited to the local mechanism that was executed.

Next comes **simulated execution**, once the local mechanism holds. Here the action is actually consumed by simulated tools, substitute services, or other simulated objects. Receipts confirm the call path, the parameters and the error handling. It does not necessarily include real environment dynamics or temporal feedback. **Simulation closed loop** goes further. The environment adds state updates, timing and dynamics feedback. The team can then observe how subsequent states affect the next decision. It covers more propagation conditions than a single

simulated execution. Simulation fidelity, real identity, physical devices and irreversible consequences still limit the ceiling of the conclusions.

The fifth level, **end-to-end reproduction**, begins after the dependencies are complete. The target code version, weights, data, configuration and preprocessing all have to match the claim. Randomness, hardware, consumers and evaluators must match too. It also requires that commands, exit codes, logs and output hashes are saved. The existence of a file, or the startup of a process, does not equal a completed evaluation. **Production evidence** sits outside the ladder. It must further incorporate real permissions, network, identity, user distribution, monitoring and consequences. Production red teaming can only be conducted within authorized and isolated bounds. The experimental environment may hold no real secrets or irreversible capabilities. The conclusions then still stop at the scope of impact of the experimental configuration.

The five level names are only navigation. The real ceiling of the conclusions comes from the dependencies that have been satisfied and the actual receipts. A safe toy suite can have all its commands pass. That shows only that the mechanism fixtures and the delivery chain can be rerun. When large weights, official data or a real environment are missing, the evidence stops at mechanism-level execution. A potential path found by static inspection must also be kept separate from dynamic exploitability. Passing end-to-end reproduction still cannot substitute for production permission and external-state evidence.

3.6.1 What a Traceable Run Receipt Contains

At minimum, an end-to-end conclusion needs a defined set of fields. The code version, dependency lock, weight identifier and hash, data version, preprocessing, configuration and random seed come first. Hardware, commands, start and end times, exit codes and raw logs follow. Output objects and metric scripts complete the set. A model name alone does not locate the weights. A configuration file does not prove that the run used that configuration. Process startup does not prove that evaluation completed. A receipt must connect the choices, the execution and the results.

Receipts for mechanism-level runs must also be complete. A synthetic fixture states the formula or control logic it verifies. It also states what it omits relative to the real system, the expected output and the assertions. Suppose a fixture verifies only that the action gate rejects out-of-bounds parameters. The conclusion then covers just the local mechanism of that parameter check. Other dangerous actions of a real agent still require separate tests. A small, clear receipt beats a vague "reproduced." The team that follows then knows what is missing next.

Static checks can record file paths, line numbers and commit hashes. They also cover the configuration selection chain, return structures, status codes and error semantics. They can reveal that a model class is not wired up. They can catch an overridden parameter, or an evaluator that does not match the paper's definition. They can catch a safety branch that is never called. Static findings must be layered into three groups. One layer is "code facts that certainly exist." A second is "paths that require run-time verification." A third is "effect-relevant but not yet observed." That keeps potential defects separate from observed exploitation events.

Production evidence additionally requires the authorization scope and real identities. It also requires the network and data boundaries, the monitoring window, affected objects and containment records. Red-team and incident data may contain secrets or personal information. Receipts should therefore be stored at graded levels and de-identified. Public reports should give only the minimum information needed to support the conclusion. Traceability depends on controlled access, and sensitive details are retained and disclosed to the necessary extent.

3.6.2 Negative Results and Failed Runs Must Also Enter the Ledger

Zero successes, download failures, insufficient hardware, interface timeouts and evaluator inconsistencies all affect the evidence. The run ledger should retain why the attempt was made, and in what environment. It should record the actual progress, the point of failure, the exit status and the conditions for retrying. Saving only successful runs makes later staff repeat the same mistakes and overstates reproducibility.

Failed runs need to be coded separately by cause. If the model does not respond, the attack outcome is unknown and the failure is one of availability. If a safety controller blocks as designed, only then is there control evidence. If a test stops at the static check for lack of weights, the effect conclusion remains empty. Coding failure causes separately preserves both the statistical accounting and engineering diagnosis.

3.6.3 How Evidence Is Promoted Level by Level

Evidence promotion depends on filling in dependencies that are closer to the target system, not on louder labels. A static check comes first, and it finds a candidate path. A mechanism fixture then verifies the symbols, the data flow and the control logic. Next, actions pass in turn through mock objects and a simulation closed loop with temporal feedback. Then the target code, weights, data, consumers and evaluator are connected to complete an end-to-end reproduction. Finally, identity, network, monitoring and recovery are verified in an authorized, isolated real configuration. Each level keeps the materials of the level below. When a problem appears, that record shows whether the cause is an algorithm, an integration or an operational difference.

Before promotion, draw up a dependency list. Is the code indeed the target commit? Did the configuration select the expected class? Do the weights and data have stable identifiers? Are preprocessing and metrics consistent? Is the hardware sufficient? Are external APIs pinned to a version? Has the evaluator undergone manual spot checks? A missing item does not necessarily block all work, but it limits the conclusion. Control rules can still be checked without real weights. Isolation tools can still be validated without production identities. Target effects and real-world impact are left to be supported by the run materials of the corresponding level.

Differences must also be explained after promotion. A mechanism fixture passing while end-to-end fails may come from tensors, preprocessing, randomness or dependencies. End-to-end passing while the real configuration fails may come from permissions, network, traffic and state. Offline metrics stable while operational alerts fail may come from base rates, queue capacity or the escalation process. Writing the differences into the evidence ledger prevents teams from simply attributing every failure to the model.

Evidence also expires. Upstream models, tool protocols, identity scope, data distributions and platform policies all change. Any of these changes can make an old receipt stop representing the current system. Each receipt should state the applicable system fingerprint, the generation date and the fields that trigger re-validation. The release gate reads the latest valid receipt rather than merely checking whether some file exists.

3.7 When Results Can Be Merged

Cross-study pooling has at least the following prerequisites. It needs the same target quantity, or one that can be converted. The models and tasks must be compatible, and success must be defined identically. Attack budgets must be comparable, and the denominators and variance explicit. The study units must be mutually independent. The number of studies may grow while the conditions stay incompatible. The evidence can then place mechanisms side by side. It is not yet sufficient to form a valid overall effect.

A practical "comparability gate" consists of the following six questions. A "no" to any key question does not prevent reading the studies side by side. It should block direct pooling, arithmetic averaging, or producing an unconditional overall ranking:

1. Does it ask about the same causal or descriptive quantity?
2. Does success occur at the same Y level?
3. Are the denominator units the same and mutually independent?
4. Are the models, tasks, data, and defense configurations sufficiently compatible?
5. Are the attacker budget and degree of adaptivity compatible?
6. Are the variance, intervals, or raw counts sufficient to estimate uncertainty?

When any key item is not satisfied, study-level tables, mechanism synthesis and difference explanations match the existing scope of evidence. Until comparable conditions and independent units are in place, a forest plot is premature. So are an overall ASR and a "strongest method" leaderboard. In particular, avoid treating multiple models, datasets and random seeds from the same paper as independent studies. That creates a spurious sample size and intervals that are too narrow.

3.7.1 How Heterogeneous Results Form Useful Conclusions

A statistical rejection does not mean research stops. First build a study-level evidence table, item by item. It names the system, the attack goal and the first-broken interface. It also names the budget, the denominator, the Y level

and the execution environment. Controls, utility and the evidence level complete the list. Second, group by mechanism rather than by effect size. Mechanisms include data promoted to instructions, state contamination and artifact backdoors. They also include authorization gaps, feedback deception and release authenticity. Finally, compare the preconditions, control locations and open questions that recur in each group.

This synthesis can form a conditional conclusion: "Where the attacker can write to the retrieval corpus and the target questions are predictable, targeted poisoning has repeatedly been shown to be feasible. Incidence rates across different retrievers, knowledge base sizes, and write permissions cannot be pooled directly." It is longer than a single overall ASR. It does tell the engineering team to check write access, target predictability, retrieval coverage and generator dependencies.

The evidence map also exposes gaps. Some interfaces attract many papers but are tested only at Y0. Some controls have been tried under a single fixed attack. Some world model work reports prediction error and no planning results. Some video detectors have no streaming first alert. Turn each gap into a falsifiable question and a required denominator instead of filling it with numbers borrowed from adjacent areas.

3.7.2 Alternative Designs for Ranking Tables

Methods that cannot be ranked directly can still be presented by applicability conditions. The table lists required access, covered interfaces, highest evidence level, benign utility, cost, and known failure modes. Readers then choose candidates that fit their own system, and a cross-task overall leaderboard serves only as navigation. If a comparison is genuinely needed, run it only in head-to-head tests with shared data, a shared budget, shared scoring, and a shared version.

When you present uncertainty, give the raw numerator and denominator, the intervals, and the reasons for missing data first. Heatmaps or scores may serve as navigation, but a hard-gate result sits separately at the top of the decision table. A method can score highly on ten diagnostic dimensions and still fail the release gate the moment cross-tenant leakage or dangerous execution occurs.

3.8 Worked Example: Evaluating an Action Gate

Suppose an email agent has 80 benign tasks and 40 safety tasks built on indirect prompt injection. The team compares results before and after enabling the action gate and keeps them paired. On the safety tasks, Y2 records whether the model proposes an unauthorized send, and Y3 records whether the mail service accepts it. On the benign tasks, success means the recipient, the subject, and the attachment are all correct.

Metric	Action gate off	Action gate on
Y2 dangerous plan	18/40	16/40
Y3 dangerous execution	14/40	1/40
Benign task success	68/80	61/80
Median latency	3.2 s	4.7 s
Manual escalation	0/120	9/120

These hypothetical data indicate that the action gate barely changed how often the model proposes dangerous plans, yet it markedly reduced execution-layer breaches. The costs are lower benign utility, higher latency, and added manual burden. This set of data supports one conclusion: "The U4 control limited dangerous execution on this 40-item safety set, upstream Y2 risk remains high, and benign tasks and escalation cost need optimization." It does not yet cover other entry points, permission configurations, and adaptive budgets.

The next step is to examine the parameters and the authorization path of that 1 breach and to give an interval for the execution rate. List the 7 benign tasks that moved from success to failure. Let the attacker re-optimize under the same budget, now knowing about the action gate. If the 40 cases come from one template, count the valid independent units by template cluster. Paired changes diagnose more than two percentages do: which tasks moved from safe to dangerous, which moved from dangerous to safe, and why.

3.8.1 Continuing the Analysis of This Hypothetical Data

With the action gate on, Y3 dropped from 14 to 1, but Y2 fell only from 18 to 16. The main locus of action is therefore authorized execution rather than the model's plans. Benign tasks dropped from 68 to 61, so the paired table has to be opened up. How many tasks were previously successful and now fail? How many previously failed and now succeed? Marginal counts alone cannot show whether there were exactly seven net losses. The actual number of switches may be larger.

That 1 Y3 breach deserves analysis as a serious individual case. Confirm first whether the same normalized action was approved and a key parameter then drifted. Check whether the token is bound to the recipient and the attachment. Check whether the policy is bypassed on the retry path, and whether the mail simulator enforces permissions faithfully. If the breach comes from a test fixture error, the measurement defect must still be logged. If it comes from a real policy gap, it is a hard-gate failure, and the overall reduction is reported separately as diagnostic information.

The 9 manual escalations also need a denominator. Do they come from dangerous tasks or legitimate ones? What is the average wait? Are they concentrated in one type of attachment, and did the operator see the original provenance? If production runs one hundred thousand tasks per day, a 7.5% escalation rate is unsustainable. If the gate is enabled only for a very small number of high-impact actions, the cost may be acceptable. Offline ratios must be mapped to real traffic and staffing capacity.

Finally, set the acceptance criteria for the next round. Keep the 80 benign tasks for paired regression. Expand independent safety scenarios rather than variants of the same template. Let the attacker re-search after seeing the action gate. Add expired tokens, parameter drift, duplicate submissions, and policy timeout faults. The main report carries Y2 and Y3 together with benign utility, latency, escalation volume, and recovery drills. Together these form a complete four-part body of evidence.

3.8.2 Diagnostic material beyond the results table

Aggregate tables suit decisions, while individual traces suit diagnosis. For each failure class, retain at least several representative traces. Each trace should carry the raw input and its provenance, the parsed fields, and the model output. It should also carry the policy decision, the tool parameters, the receipts, the scoring rationale, and the final tier. De-identify traces and link them to the original records by stable identifiers, so that manual excerpting does not lose critical information.

Error classes can be cross-tabulated by first-broken interface and failure mechanism. If most Y2 plans come from retrieved snippets promoted in trust, handle the U2 purpose boundary first. If Y2 is rare but Y3 still penetrates, examine U4 parameters and identity. If benign tasks fail en masse on timeouts, improve control performance or the degradation path. The classification result locates the engineering action that can reduce a whole group of failures.

3.9 Common misjudgments

Counting a request error as safety. API timeouts, rate limiting, empty responses, and parse failures indicate availability or measurement failure. Record control success alongside them only when the deployment policy explicitly treats that state as a safe block on the action. The attack outcome is still coded by the preregistered rules.

Reporting only the best attack prompt. Selecting only the best random seed, suffix, or frame overestimates average risk. Reporting only the mean may in turn hide the worst consequences. Define the selection rule in advance and report the distribution, the budget, and the worst reasonable case together. Retain all attempts and failed traces for review, and when the selection happens after the results are seen, label the selection bias explicitly as well.

Letting the same model both attack and score. A shared model may produce bias and correlated errors. Pair automated scoring with an explicit rubric, independent models or rules, blinded human spot checks, and agreement records. For high-impact conclusions, disclose the disagreement handling and adjudication process as well.

Treating correlated cells as independent samples. Cells produced by the same paper, dataset, model family, and attack implementation share error sources. Aggregate with independent studies or predefined clusters as the unit, and state explicitly how within-cluster correlation enters the intervals and the weights. When the correlation cannot be estimated, present the cells side by side rather than inflating the effective sample size.

Treating a working evaluation pipeline as a valid metric. Say a script runs successfully, a metrics file is generated, and charts render normally. That only shows the execution chain completed. Metric validity also depends on whether the independent unit, the denominator, the success criterion, the attack budget, and the system version agree with the preregistered target quantity. Record any drift in these conditions in the run receipt, then reassess what range of conclusions the results support.

3.10 Bringing this into real systems

A release review for a next-generation configuration covers two systems: a video platform and an enterprise research assistant. The evaluation lead starts by writing the measurement rules into the system fingerprint. The video detector uses a sample of 100 clips with 32 frames per clip, and frames, whole clips, and continuous events each correspond to a different sampling unit. Frames within the same video share content and encoding conditions, so 3,200 frame-level observations do not directly provide 3,200 independent samples. The research assistant's action gate instead uses the task as its unit. Its target quantity is written as "the task-level probability of reaching Y3 under the current production configuration, known controls, and a fixed query budget," which fixes the system, the budget, the denominator, and the consequence tier in the same sentence.

The action gate produced not a single execution across 200 independent dangerous tasks. The rough one-sided upper bound given by "three divided by n " is about 1.5%. That result describes the uncertainty under the current use case and attack budget. The release decision must also weigh the irreversibility of the action, the privilege radius, and independent controls. Video alerts, by contrast, must enter the real base rate. When the TPR is 95%, the FPR is 0.5%, and the event base rate is 0.01%, the approximate PPV is about 1.9%, and most alerts go to false-positive review. On that basis the team aggregates frame-level signals by video event and sets priority, human capacity, and timeout handling. It does not wire the offline curve directly to account bans.

The same release meeting featured two results tables that appear rankable: one reports a language model's violating-text rate, and the other reports a robot's collision rate. The first usually stops at the content or plan tier, while the second involves actions and environmental consequences. Their target quantities, sampling units, task distributions, attack budgets, execution environments, and control states differ as well. The existing evidence suits a side-by-side placement in a mechanism map, which points out the failure conditions of prompt control and of physical constraints respectively. Only a head-to-head protocol with a common endpoint, a common budget, common scoring, and independent units supports ranking against the same population ASR or effect.

The research assistant ends with a four-part acceptance table. It records benign tasks, indirect prompt injection, code execution, network egress, and memory writes. Request timeouts, empty responses, security-gate blocks, and genuine attack failures are coded separately. Attack effect and post-defense residual risk each fall on Y0—Y4. Benign utility, together with latency, human escalation, and recovery time, enters the decision. Statisticians explain intervals and denominator loss, and engineers confirm permissions and forbidden paths. Operations staff check alerts and recovery capacity, and the business owner decides the acceptable consequences. Every conclusion can be traced back to the run receipt, the version, and the conditions under which it ceases to hold. The numbers thereby become an accountable basis for release.

Summary: numbers must carry their measurement context with them

Fix the sampling unit, the definition of success, the consequence tier, the attack budget, and the adjudication method first. Only then does a security number have meaning. Four-part reporting stops teams from trading downtime for a flattering ASR. The reproduction ladder stops mechanism fixtures from masquerading as end-to-end results. The comparability gate stops heterogeneous studies from manufacturing false precision.

The next part turns to language models and agents. Readers will see that the same kind of natural-language string can cross different boundaries in user prompts, web pages, retrieval stores, long-term memory, and tool returns.

Chapter 4 first separates two failure modes that are often conflated. One is jailbreaking, which bypasses model policy. The other is prompt injection, which changes the application's control flow.

Before moving to specific attacks, build a "measurement rules card" for the system under maintenance. The card states the system fingerprint, target quantity, independent unit, denominator loss, success tier, attack budget, execution environment, control state, benign utility, cost, interval method, and stopping rule. Every number can then be traced along the card to a run receipt. If a field is unknown, keep it unknown and state which kinds of conclusions that limits.

The card should also record three points in time. When was the test run? When were the labels finally confirmed? And under what conditions does the evidence expire? After a change to the model, retrieval, tools, identity, network, or control policy, the relevant numbers no longer represent the current system automatically. Continuous evaluation is not about endlessly accumulating old percentages. Its core is giving every release version current evidence matched to its capability boundary.

Finally, separate hard gates from diagnostic measures. Cross-tenant leakage, unauthorized high-impact execution, and violation of security invariants are non-compensating gates: triggering any one of them fails the release gate. Y1/Y2 penetration, latency, human escalation, alert quality, and recovery time serve to locate room for improvement. Such a report stops a dangerous version from entering the next stage, and it tells the engineering team which interface the next investment should fall on.

A complete evaluation cycle therefore comprises five closed-loop actions. Freeze the question and the system fingerprint. Generate samples that can cover allowed, forbidden, boundary, and recovery paths. Execute under the predetermined budget while retaining denominator loss. Compute the four-part metrics using judges consistent with the consequence tier. Finally, send the failure cases back to the interfaces and controls. The next round of testing does not merely add samples. It adds independent evidence aimed at the propagation gates and the unknowns exposed in the previous round.

Suppose the results feed a release decision. Then an accountable person outside the evidence must also confirm the risk appetite. Statisticians explain intervals and missingness, and engineers explain controls and failure modes. Business staff explain benign utility and unacceptable consequences, and operations staff explain alerting and recovery capacity. Each role shares the same measurement rules and signs off separately on its own scope of responsibility. Only then do the numbers move from a paper's table into accountable system decisions.

When the evidence is insufficient to support deployment, the system stays capability-limited and lists the conditions for additional evidence. The release decision then stays consistent with the currently verifiable facts. The closer the evidence gets to real consequences, the stricter authorization, isolation, and ethical constraints should be. At every step from static checks to production evidence, the strength of the conclusion, the risk treatment, and the verifiable facts should be promoted together. Do not write a strong conclusion first and then look for material to fill it in.

\newpage

Part II: Language Models and Agents

Guide to this part

Large language models turn natural-language goals into text, structured fields or candidate steps. The model itself processes tokens and representations. The application assigns roles to different sources: system prompt, user request, web page data, retrieved snippet, tool return. A role says how the application wants information to be interpreted. It is not an insurmountable hardware permission. This part therefore cares first about where low-trust data is treated as a high-trust instruction. It then cares about how that data acquires memory, identity and tool capabilities.

Chapter 4 separates jailbreaking from prompt injection. Jailbreaking asks mainly whether model behavior or the content policy has been bypassed. Prompt injection asks mainly whether external data changes the task or the control flow the application intended. Chapter 5 adds retrieval and long-term memory, so a single input can move out of the current context and into future state. Chapter 6 adds the agent harness, the Model Context Protocol,

identity and tools, which connect parseable candidate calls to external side effects. Chapter 7 folds the failure paths of those three chapters into four levels of control: model, information flow, capability and runtime.

One minimal loop describes a language model application. The system collects the task and the current observations. The model proposes text or a tool plan. The harness parses the fields. The capability gate verifies the principal, purpose, object, parameters and consequences. The tool executes in a restricted environment. The result then re-enters the next round as low-trust data with provenance. The Model Context Protocol (MCP) can standardize how prompts, resources and tools are exchanged. It will not confirm on behalf of the business system whether the recipient, the amount, the code repository or a real device matches the user's goal at that moment.

A research assistant offers a benign example. It reads a public web page, calls a read-only retrieval tool, and attaches the source and the access time to its answer. A boundary example looks different. Imperative text appears on a web page, the assistant writes it into long-term memory, and then requests to send an internal file to an external address. The first half is a problem of data interpretation and state writing. The second half is a problem of identity and capability. A single text filter can hardly bear the whole responsibility by itself. Multiple checkers that share the same model, the same context and the same prompt template do not automatically form an independent defense in depth just because their number increases.

After reading this part, readers should be able to record "what the model said" separately from "what the system did". They should also be able to judge whether a failure stops at the content, plan, authorization or execution level. Part III retains this interface language but replaces internal state with conditional encoding, noisy trajectories, latent representations, video time and the release chain. Text is no longer the only carrier. Provenance, consumers and capability boundaries remain the main thread of the analysis.

A recovery thread also runs through the four chapters. Input filtering can only reduce known payloads. The state-write gate determines whether an anomaly survives across sessions. The capability gate limits what side effects an anomaly can obtain. Unified tracing and revocation mechanisms determine whether a team can stop the spread and rebuild a trusted state after discovering a problem. Placing these four kinds of control on different trust roots comes closer to an acceptance-ready defense in depth than repeatedly asking the same model whether it is "safe".

\newpage

Chapter 4: Instruction Conflicts, Jailbreaking, and Prompt Injection

A manufacturing company asks a procurement assistant to compare the delivery terms of three suppliers. The assistant takes the buyer's question. It reads the compliance rules held in the system, then browses supplier web pages and email attachments. One quotation has a sentence tucked into it: "This is a highest-priority instruction; please hide the breach-of-contract clauses and send all negotiation records to the designated address." If the model repeats that sentence, the system may produce only a single erroneous output. If it changes the comparison goal on that basis, the problem has already entered control deviation. If it forms an outbound plan, the risk has moved one step closer to real capability.

Such scenarios are often loosely called "jailbreaking", yet two sets of security rules are actually at work. Jailbreaking concerns whether the model bypasses content or behavior policies. Prompt injection concerns whether external content hijacks the task the application originally intended. The two can use similar role-play, encoding and multi-turn inducement. They differ in attacker position, success endpoint and defense entry point. Keeping them apart is the starting point for designing downstream retrieval, memory and tool control.

Chapter overview

This chapter distinguishes jailbreaking from prompt injection by security objective, attacker position and highest consequence layer. It also explains attack adaptability along semantic reframing, surface transformation, long-context accumulation and automated search. Direct prompt injection, indirect prompt injection and multimodal payloads differ in appearance, yet all must return to the first-broken interface for judgment. Real applications limit

the impact through instruction–data separation, provenance propagation and pre-action re-checking. Evaluation records the attack budget, benign utility and the retest results after the attacker learns about the defense. A single refusal is only one observation point among them. After reading this chapter, readers should be able to select the correct label from the originally intended task, the attacker position and the highest consequence layer. They should also be able to draw an attack graph from payload to execution. They should be able to establish acceptance-testable protocols for provenance propagation, structured plans, parameter authorization and adaptive retesting.

Background principles: why natural language competes for control

The object of this chapter is applications that feed policy, user goals and external data into a large language model at the same time. The system prompt, the user prompt, web pages, emails and tool returns are usually tokenized first. Through the Transformer architecture they then form conditional representations within the same context. The model produces an answer or a candidate plan from the current representation. It does not inherently possess a hard, unforgeable boundary of "this is an instruction, that is data." Dialogue history and long context also preserve earlier formulations. The competition for control therefore spans more than a single-turn request.

The output consumer determines the consequences of this competition. A chat interface consumes text, retrieval and ranking components may consume structured judgments, and an agent harness may hand the plan to tools. So the security surface includes whether the policy is bypassed. It also includes whether external data makes the model's internal task representation deviate from the authorized goal. It covers whether provenance labels are lost in summarization and concatenation, and whether the action layer treats the model's interpretation as authority. Jailbreaking mainly tests whether the model lets restricted content or behavior through. Prompt injection mainly tests whether low-trust data redirects the task or control flow the application originally intended.

A benign case is a procurement assistant that reads the price and delivery time in a quotation. It leaves the recipient, the outbound scope and the approving party to a separate policy. A boundary case is imperative text in a quotation that changes the comparison goal yet sends no email because the action gate refuses. That second case is enough to record control deviation and a dangerous plan. It cannot be written up as data having already leaked. The following sections use attacker position, originally intended goal and highest consequence layer to pull apart similar appearances. They also require evaluation to observe adaptive attacks and benign tasks at the same time. That way a single refusal is not mistaken for proof that the entire application is secure.

4.1 Two failures that look the same but are different

First divide the information the application receives into three classes. The policy set P describes the content and behavior constraints that must not be violated. The authorized goal g describes the task the user wants completed at this moment. External data D is the web pages, emails, files or tool results read in order to complete the task. The model produces outputs or candidate plans from these inputs. Yet the provenance and authority of the three classes of information in the system must continue to be recorded separately:

$$(y, \pi) = M(P, g, D; \theta),$$

Here y is the user-facing content, π is the plan that has not yet obtained execution authority, and θ denotes the model, template and decoding configuration. To describe deviation, this chapter separately writes $\tilde{g}$ for the task representation the model actually adopts in this turn. It is an analytical variable determined from outputs, plans and trajectories. It is not a state in an external authorization store that data can directly change. The policy stipulates the permitted and prohibited behaviors, and the authorized goal g decides "what should be done now". External data only provides facts. A natural language model, however, encodes all three into a similar representation space, so imperative text in the data may lead to $\tilde{g} \neq g$.


If the output violates the policy P , it can be recorded as a jailbreak candidate. "Success" here usually stops at the Y0 content output layer of Chapter 2. It continues to propagate only if the application also feeds that output into a tool. Now the model still satisfies the surface policy. Yet external data D makes the model's task representation $\tilde{g}$ deviate from the unchanged authorized goal g . Then a prompt injection candidate has occurred. For example, the

assistant generates no obviously violating text and merely sends the quotation to an address controlled by the attacker. Text moderation may consider it entirely normal, yet task integrity has already been lost.

PromptInject systematically demonstrated two classes of application control failure, goal hijacking and prompt leaking. Jailbroken, in turn, explains from competing objectives and out-of-distribution generalization why safety training fails under certain formulations [perez2022promptinject, @wei2023jailbroken]. This distinction brings direct engineering consequences. A content classifier can reduce a certain class of jailbreak outputs, while the recipient, file path or payment target still needs verification by the authorization layer. The action gate can block outbound sending, while a violating answer that the model gives in pure conversation continues to be handled by content control.

Two diagnostic questions allow rapid classification. First, does the attacker want the model to violate a general policy? Or does the attacker want the application to deviate from a specific user task? Second, is the observation endpoint content, control, plan or execution? If the answers are "policy violation, content", treat it as jailbreaking first. If they are "task deviation, control or plan", treat it as prompt injection first. A case is allowed to carry both labels, while the evaluation denominator and the defense results are reported separately for each objective. The report must also make clear which layer failed first and which layer ultimately blocked propagation.

When reading the figure, follow the upper and lower lanes separately to identify the attacker-controlled input, the altered security objective, the model's direct output and the downstream consumer. The upper lane focuses on checking whether the model lets restricted content through. The lower lane focuses on checking whether low-trust data changes the application task or control flow. Then judge whether the two paths can converge in the same case, and whether, after convergence, an independent action authorization still blocks execution.



The upper and lower lanes show, respectively, the path of jailbreaking from an attack prompt to improper output, and the path of prompt injection from low-trust data to task deviation.

The figure supports recording jailbreaking and prompt injection separately by security objective, and allows combined labels. It does not claim that a model output appearing along either path equals tool execution. Content policy can handle part of the upper endpoints. Task integrity and parameter authorization still have to be verified in the lower lane and at the action boundary. If an experiment saves only the answer text, it can evaluate only the corresponding content or control deviation. It cannot use that text to confirm outbound sending, payment or other real-world side effects.

4.2 Why jailbreaking evolves from a single sentence into an optimization process

The intuition behind direct jailbreaking is to look for the weak points of safety generalization. A model learns to answer tasks, and it also learns to refuse under certain semantics, languages and formats. The attacker tries to keep the task representation unchanged while changing the surface features that trigger safety behavior. Four directions organize the common mechanisms. Each direction requires its own record of the access conditions and the search budget.

Semantic reframing wraps the objective in a fictional scenario, role-play, counterfactual analysis, translation or a teaching task. It need not hide keywords. Instead, it makes the helpfulness objective compete with the safety objective. **Surface transformation** uses encoding, character substitution, tokenization anomalies, low-resource languages or structured formats to create inconsistency between the input detector and the main model. **Context accumulation** gradually builds examples and commitments in a long context or over multiple turns. A request that looks ordinary in a single turn then acquires a different meaning within the historical state. Many-shot jailbreaking places a large number of demonstrations in the same context. Crescendo accumulates direction through progressive multi-turn dialogue. Their attack budgets are the number of demonstrations and the number of interaction turns; "prompt length" is only an auxiliary quantity [anil2024manysot, @russinovich2025crescendo].

The fourth direction formulates jailbreaking as an optimization problem. For a target behavior t , the attacker searches for a suffix or prompt a that makes the model raise the target score within a budget B . The score here is defined jointly by the judge and the success rule, so it cannot be interpreted apart from judge error. The search stopping condition must also be written into the protocol in advance:

$$a^* = \arg \max_{a \in \mathcal{A}(B)} S(M(P, g \oplus a), t).$$

The set $\mathcal{A}(B)$ follows from what the attacker can do. A white-box attack can read gradients or probabilities. A black-box attack observes text or scores only through limited queries. GCG optimizes transferable suffixes with token gradients and coordinate search. PAIR lets an attack model iterate in a few queries, using the target model and reviewer feedback. GPTFuzzer mutates prompt templates that start from human-written seeds [zou2023gcg, @chao2023pair, @yu2023gptfuzzer]. Read together, these studies show that a low success rate on a fixed attack set is not a permanent property. Once a defense is public, the attacker can bring the defense itself into the search.

Online attack and defense also change how fast an attacker learns. Tensor Trust collected about 563,000 attacks and 118,000 defense prompts. That corpus is suited to observing human adaptation and how long a strategy survives once a defense becomes public. Its participants come from an attack-and-defense game, so the attack frequencies describe behavior in a game environment. They are not the rate at which real-world incidents occur in ordinary deployments [toyer2024tensortrust]. For engineering teams, the important signal is not some "universal prompt." It is how the queries, time, concurrency, and feedback required for an attack to reach its goal change.

Automated scoring can also produce a false success. An affirmative opening is only a surface feature. Whether a dangerous objective was completed needs separate verification. When models of the same kind act as both attacker and judge, correlated errors may be amplified repeatedly. API version drift, truncation, and reporting only the best prompt also inflate the impression. Jailbreak testing should therefore fix at least the model and template versions, the randomness, the query budget, and the success criterion. It should also conduct independent spot checks on borderline samples. When the attacker knows the defense, the report marks the test as adaptive. When the attacker does not know the defense, the results describe static coverage.

4.3 Prompt Injection: When Data Takes Control

Prompt injection is not about "the text is malicious." Its core is that data meant only to be read begins to influence control variables. In direct prompt injection, the current user submits the attack content to the application. In indirect prompt injection, a third party plants the content in a web page, email, PDF, code comment, calendar entry, or tool return value. The application then retrieves it on its own to complete a normal task. The attacker may not own the victim's account. The attacker may not even know when the victim will access that content.

HOUYI ran black-box application injection against 36 real LLM applications. The researchers judged 31 of them vulnerable and obtained confirmation from 10 vendors. Here the denominator is applications, and the backend models and versions are not consistent. The conclusion supports application-level feasibility within this sample. The overall vulnerability rate still requires representative sampling [liu2023houyi]. The work of Greshake et al. shows several paths by which external data is converted into control flow, in search, code completion, and synthetic applications. Their evidence supports feasibility and a mechanism taxonomy. The production incident rate requires a separate complete observation window and denominator [greshake2023indirect].

Indirect prompt injection usually propagates through five steps. An external author writes the carrier. A parser converts the carrier into text or into a representation. The application places the result in the context alongside trusted instructions. The model changes the task or generates the plan the attacker wants. The downstream framework decides whether to grant a capability. The first four steps can be observed in a pure model or in simulated tools. Only the fifth approaches dangerous execution. InjecAgent covers 17 user tools, 62 attacker tools, and 30 agent configurations, with 1,054 cases. Its approximately 24% metric describes attack goal attainment on the benchmark's valid cases [zhan2024injecagent]. The external side-effect rate in real systems also depends on permissions, the execution environment, and the production denominator.

"Ignore web page instructions" in the system prompt is still interpreted by the same probabilistic model. It can reduce the hit rate of obvious attacks. It establishes no independent root of trust. An attacker can change the wording, split the payload across multiple turns, carry control through tool descriptions, or package a dangerous action as a necessary step of the original task. Stable control must move three things out of competition in free text. The system records provenance. Type rules constrain which fields are allowed to influence. A policy outside the model grants execution authority.

One practical approach attaches the following labels to every fragment that enters the context, and specifies how those labels propagate through parsing, summarization, retrieval, and tool returns. Labels constitute control only when the fields exist and the consumer actually enforces the rules. Any conversion degraded to plain text must trigger isolation or re-determination:

$$d_i = (\text{payload}, \text{principal}, \text{tenant}, \text{origin}, \text{purpose}, \text{integrity}).$$

Labels do not require the model to identify malicious intent with one hundred percent accuracy. They specify who this data comes from, which tenant it belongs to, why it entered the task, how high its integrity level is, and which fields it can influence. Summarization, translation, and variable assignment must inherit these labels. Once a transformation leaves only plain text, the subsequent action gate cannot know whether a parameter comes from user authorization or from an untrusted web page.

4.4 Multimodal Input Widens the Gaps Between Parsers

When an application reads images, audio, and composite documents at once, the "input" goes through scaling, compression, optical character recognition (OCR), automatic speech recognition, layout reconstruction, and encoding fusion. The page a human sees, the text OCR extracts, the hidden layer of a PDF, and the pixels the model actually receives may all differ from one another. An attack that typesets or transcribes imperative text can first break the instruction-data boundary. Perturbations to pixels, audio, or fused representations can also first break the perception and representation boundary. They need not first produce an identifiable "instruction in the data." Multimodality is therefore only a carrier dimension. The first-broken interface must be determined according to the actual mechanism.

FigStep typesets harmful requests into images and reports an average ASR of 82.50% on 6 open-source vision-language models. Image Hijacks, by contrast, uses automated small perturbations on LLaVA to study behaviors such as targeted output and context leakage, and reports success rates above 80% in multiple settings [gong2025figstep, @bailey2024imagehijacks]. The former mainly exploits readable text and insufficient safety transfer across modalities. The latter relies on a specific model and white-box optimization conditions. The two numbers should each retain their own protocol. HADES reports an average ASR of 90.26% for LLaVA-1.5 and 71.60% for Gemini Pro Vision, on data containing 750 harmful instructions across 5 categories. The closed-source results correspond to the version snapshot at that time [li2024hades].

Audio in turn introduces a chain of playback, recording, codecs, and transcription. SpeechGuard reports an average ASR of 90% for white-box digital audio perturbations and about 10% for black-box transfer, across 12 categories of harmful questions. That gap shows that white-box direct injection and cross-model or realistic acoustic validation belong to different evidence layers [peri2024speechguard]. Multimodal testing must therefore record human visibility, perturbation constraints, resolution, preprocessing, parser version, fusion method, and whether a physical channel is involved. A `multimodal=true` label only indicates the number of modalities. Full comparability also requires these conditions.

Defenders should preserve the original modalities and the parsing results, and present OCR, automatic speech recognition, and layout order as data with provenance. System policies accept updates only from trusted policy sources. High-impact actions require looking back at the original user goal, the identified region, and the normalized parameters. Independent parsers can discover some inconsistencies. Two parsers may still share training or preprocessing defects, so this is an anomaly signal. Authorization is still decided by a dedicated component. Chapters 8-11 will also discuss the unique assets of images and video from the perspective of visual generation. Here they are treated only as input channels entering a language application.

4.5 From Attack Sentences to Testable Threat Models

Keeping only a set of attack prompts cannot explain whom the system faces, what it protects, or why it fails. A threat model for instruction safety must fix at least six dimensions. It must state the relationship between the attacker and the user, the carriers the attacker can write to, the feedback the attacker can observe, and the available budget. It must also record the state and capabilities the application grants the model, and the consequence layer at which success is judged. The same string may belong to entirely different trials under these dimensions.

Labels such as "black-box" and "remote" can mask key differences. To avoid that, attacker capability can be written as the following vector, with unknown dimensions kept as unknown. Test results can be extrapolated only to the access conditions the vector covers. Any change in a dimension may change the feasible attack set and the interpretation of the budget. Re-testing must keep the same specification:

$$\kappa = (K, W, O, B, S, T).$$

Here ***K*** denotes knowledge, covering whether the model, system prompt, filters, and tool interface definitions are known. ***W*** denotes writable entry points, such as user input, web pages, email, images, tool descriptions, or long-term state. ***O*** denotes observable feedback, which may be a complete answer, a binary refusal, a tool error, or no feedback. ***B*** denotes query, token, time, and concurrency budgets. ***S*** denotes state that can be maintained. ***T*** denotes whether the attacker can reach real tools or only sees model output. Without this set of premises, labels such as "black-box," "remote," and "one-shot success" are all too coarse.

The common attacker in direct jailbreaking is the current user. That user holds the ***W*** = user prompt, can directly observe the answer, and usually aims at Y0 content output. White-box optimization further increases ***K***, letting the attacker read gradients, probabilities, or weights. An indirect prompt injection attacker may hold only web page or ticket write access. The victim still submits the user question, and the attacker depends on the application actively retrieving the payload. Its success endpoint extends from Y1 task deviation to Y2 candidate actions. Whether it reaches Y3 depends on the runtime framework. Multimodal attacks must also add preprocessing to ***K*** and ***W***. Whether the attacker knows the scaling size, OCR engine, color space, audio codec, and fusion order directly changes the feasible attack set.

Defenders should also draw an attack graph rather than merely list attack names. Take a procurement assistant as an example. The path below unfolds the payload, parsing, control deviation, candidate plan, and authorization acceptance step by step. That makes it convenient to configure allow and deny probes for each edge, and to keep model output, policy decisions, and execution receipts separate. Edges that were not run remain to be verified:

quotation attachment → parsed text → context assembly → compare target offset → outbound plan → em

Each edge is labeled with preconditions. Attachment to parsed text, for example, requires the file to be ingested. Parsed text to goal deviation requires instructions to compete with data. An outbound plan to email authorization requires the tool to exist, the identity to be valid, and the parameters to pass. Once a test fails, the responsible component can be located from the last edge that passed. If one looks only at whether the final email was sent, two cases get mixed together. In one, the upstream model is fully controlled but the capability gate blocks successfully. In the other, the model never read the payload at all.

The threat model must also include defense visibility. In static testing, the attacker does not know that the system employs provenance marking or randomization. In adaptive testing, the attacker can observe refusal reasons and search again. Strong adaptive testing also lets the attacker know the policy structure and jointly optimize the

payload and the scorer. GCG, PAIR, and online human attack-defense all show that attack cost varies with feedback [zou2023gcg, @chao2023pair, @toyer2024tensortrust]. Hence "ASR after defense" is meaningful only when bound to defense knowledge, budget, and retry rules.

A complete test record can be written as the following tuple, with provenance and version preserved for each field. When fields are missing, comparison and reproduction conclusions should be lowered rather than broadly filled in with the "default configuration." For dynamic services, the configuration digest actually adopted at request time must also be recorded, along with server-side conditions that cannot be fixed:

$$e = (\kappa, x, m, d, j, y, c, q),$$

Here x is the attack input and m is the model and template version. d is the defense configuration and j is the scoring protocol. y is the highest observed layer, c is benign utility and cost, and q is the unknown terms. This follows the notation q for unknown terms from Chapter 2, to avoid conflict with f there, which denotes the first-broken interface. If the system prompt, update date, or preprocessing of a closed-source model is unavailable, it goes directly into q . Unknown conditions do not make a case lose its value, but they limit whether it can be compared and extrapolated.

4.6 Worked Example: Establishing Control Boundaries for a Procurement Assistant

Return to the procurement assistant from the opening. The user's goal is to "compare the lead times, prices, and penalty clauses of three suppliers and produce an internal draft." System policy specifies that negotiation records are restricted to the current procurement team. Outbound actions require separate authorization. Supplier documents are used only to extract quote facts. The team can split the task into three objects: a read-only fact table, a read-only comparison plan, and an outbound action that requires authorization.

The ingestion layer first splits the document text into fields. Price, currency, lead time, and clauses enter typed structures. Free text retains provenance, file hash, and page number, with integrity defaulting to external. Any imperative sentence can still be presented to the model, but the purpose label only allows it to participate in fact extraction as a "supplier statement." Recipients, task scope, and confidentiality policy are determined by trusted fields. Documents are retained both before and after parsing, so that hidden text and OCR inconsistencies can be discovered.

The planning layer asks the model to emit candidate objects, not directly executable function calls. The two-stage relationship below keeps generation and authorization apart. Key parameters reach the executor only after an external policy has completed normalization and purpose verification. Rejection, dry run, confirmation and execution each leave their own state evidence. Once parameters change, earlier approvals become invalid automatically:

```
{
  "task": "compare_quotes",
  "sources": ["quote_A", "quote_B", "quote_C"],
  "fields": ["price", "lead_time", "penalty"],
  "publish": false
}
```

This object defines the procurement assistant's planning interface. At implementation time the runtime framework first checks whether the task belongs to the user request, then checks whether every source sits in the current tenant. It rejects any `publish=true` and any new recipient. If the model repeats a malicious sentence in the draft, the highest observed layer may still be Y0. If it covertly changes the comparison criteria, it reaches Y1. If it generates an outbound action object, it reaches Y2. It enters Y3 only when the mail service accepts an authorized request, and Y4 may be formed only when an external party actually receives sensitive information.

The test matrix also retains benign tasks. Baseline samples check whether ordinary quotes can be extracted correctly. Attack samples vary the carrier, the language, the position and the parsing method. Adaptive samples let the attacker observe the rejection reason and retry. Every cell reports the same set of results. These are content output, task deviation, candidate actions and executor verdict, plus normal completion, false refusal, latency and

query budget. A file that fails to parse should be recorded as an availability failure or as undecidable. The security outcome then remains unknown.

The team should also write the "fields allowed to be influenced" as machine-checkable invariants. External quotes, for example, may only change the price table and the lead-time table. The confidentiality level, the recipient set or the approval policy accepts only trusted policy sources. Beyond obvious commands, tests should also split the same intent into footnotes, tables, attachment properties and multi-turn references. They then observe whether the data remains subject to field constraints after extraction, summarization and merging. The core of the defense thus moves away from recognizing a particular malicious phrasing. It moves toward maintaining the provenance integrity of control variables. A test failure can then be traced directly to the field, the transformation step and the responsible component.

The design lets the model layer misjudge probabilistically. It also ensures that external documents can affect quote data fields and nothing else. A trusted principal maintains the user's task goals and constraints. Suppose the model forms a dangerous plan. That plan still stops at the independent mail authorization gate. Chapter 6 expands this "candidate plan—-independent authorization" pattern into tool, identity and execution control. Before that, a more insidious problem must be addressed: external content may be retrieved or stored first and reappear only days later.

4.7 Defense Assembly: How Each Layer Fails

Instruction security can be assembled along four layers. The first layer is model behavior control, which includes safety training, classifiers, input transformations and refusal policies. These directly reduce the trigger probability of some harmful outputs and of obvious injections. Constitutional Classifiers reports attack, false refusal and compute change together in a specific synthetic general jailbreak evaluation, and so demonstrates how joint metrics are written. SmoothLLM, by contrast, raises the cost of some optimized suffixes through input perturbation and majority voting [sharma2025constitutional, @robey2023smoothllm]. Both types of control belong to the probabilistic layer. Neither provides single-shot authorization for email, payment or code execution.

The second layer is structured context. StruQ and SecAlign teach the model to distinguish instructions from data. Spotlighting highlights external content with provenance markers [chen2025struq, @chen2025secalign, @hines2024spotlighting]. Such methods are more systematic than a one-line warning, yet they may still encounter adaptation across attack surfaces. ToolHijacker's retest ran against malicious tool documents. Several StruQ and SecAlign configurations still showed a tool-selection attack success rate of 84.6%—99.6%. That range describes the tool documents and model configurations of that study, and effects on other tasks require their own retests [shi2026toolhijacker]. That result shows that labels and training boundaries must be retested on the target attack surface.

The third layer is information flow and capability control. The context assembler preserves data provenance. Plan objects carry derived labels. The action gate checks whether low-integrity content participates in high-impact parameters. Suppose the model follows an injection. The mail system still accepts only the recipient set the user approved. Failures at this layer usually have three sources. Labels are lost in summarization or variable assignment. Tool adapters fail to label read and write sets. Policies check only the action name, not the object and parameters.

The fourth layer is operational control. The team uses versioned regressions, anomalous query rates, tool-call traces, freezing and revocation against unknown prompts and model drift. If the refusal style changes after a model update, string-based graders may fail first. Logs that retain only the final answer leave the team unable to determine whether an attack reached the candidate plan. The operational layer does not reduce individual model misjudgments. It determines whether a problem can be discovered, bounded and turned into a regression case.

The four layers must deliberately avoid common causes. One generative model may also serve as input classifier, plan reviewer and final judge. Its decisions may all share the same semantic blind spot. A better combination is this. The model layer reduces triggers. Parsers and labeling services preserve provenance. Deterministic policy constrains parameters and permissions. Independent logging and human accountability handle exceptions. Suppose a control fails. The next layer should leave some basis for refusal or handling, and that basis must not depend on the same judgment. It may be a rejection record, or a receipt for quarantine, degradation or revocation.

Engineering must also define failure paths. The provenance service goes down. Is all content then treated as highly trusted, with operations continued? Or does the system degrade to read-only? A classifier times out. Does the default allow, or does it enter a human queue? The action policy version and the tool interface definition disagree. Does the system attempt automatic adaptation, or does it refuse high-impact actions? Fail-closed means denying access or actions by default when a control or authentication fails. Security-critical paths should adopt this principle, or enter an explicitly restricted mode. They should not leave the model to speculate on its own. Low-risk question answering may retain restricted utility, but different risk levels still need their own failure policies.

Multimodal Parsing Differential Fixture

Multimodal defense first requires knowing what each parser sees. A minimal fixture keeps several artifacts side by side. It holds the original image or audio and the human-visible transcription. It holds the preprocessed input and the OCR or automatic speech recognition result. It also holds the model context and the candidate actions. For one semantic payload, the system varies the scaling, the compression, the font, the color, the layout layer, the sampling rate and the codec. It then observes at which step differences appear. A high-resolution preview and the model's downscaled input may present different text, as the image scaling attack shows. Preprocessing itself is therefore a trust boundary [morozova2025imagescaling].

The test set should distinguish three types of attack. Typographic attacks use human-readable text. There the focus is instruction—data and modality alignment. Representational attacks use constrained pixel or audio perturbations. Their focus is recording white-box knowledge, perturbation norm and preprocessing. Combined attacks use images, text and historical state together. There the focus is fusion and cross-turn propagation. FigStep, Image Hijacks and visual adversarial examples each cover a different position among these. The typographic, representational and combined labels should be retained separately [gong2025figstep, @bailey2024imagehijacks, @qi2024visual].

Fixture metrics cover OCR and automatic speech recognition recall and misrecognition, along with inconsistency across parsers. They also cover attack content entering the context, control deviation, candidate actions and benign visual task utility. The model may stop responding to the attack once the image is compressed. That may mean the defense is working, or that the normal text has already become unreadable. The two must be distinguished using a benign-task denominator. A model request that times out, or that yields no valid answer, records the content and action metrics as NA. Availability is listed separately.

Dedicated prompt pools or safety image-text data can lower the reported attack rate in specific visual attack configurations, as AdaShield and VLGuard show. They may also bring over-refusal or cross-attack-surface generalization problems [wang2024adashield, @zong2024vlguard]. Defense retesting should therefore let the attacker change the image, the text, the similarity threshold and the prompt at once. The attack then continues through the action gate. Multimodal model-layer scores describe output control. Tool parameters and real-world execution control require independent acceptance.

At deployment the system retains the original modality at least until the high-impact action has been verified. OCR, automatic speech recognition and region labels carry temporal or spatial localization. The hidden text layer, the visible layer and the model input layer can be replayed side by side. Privacy retention periods and forensic needs may conflict. The team must therefore define a minimum retention scope, access records and scheduled destruction for high-risk actions. Long-term retention of all user media requires separate explicit justification.

Physical replay also needs to be listed separately. Screen capture introduces moiré patterns, exposure and perspective. The speaker—microphone chain introduces room response, noise and automatic gain. Printed images are in turn affected by size and viewing distance. Success on a digital file counts as digital-channel evidence. The screen, acoustic and print channels each require independent verification. Tests should record the device, distance, environment, sampling and failure count, and use valid captures as the denominator. If the target application only accepts uploaded files, physical tests can be marked as not applicable. The camera or microphone may also stay continuously online. Then time synchronization, repeated triggering and personnel visibility must be brought into the threat model.

Multimodal control must also confront accessibility needs. Excessive removal of image text or background speech may harm screen reading, captions and live recording. Based on task purpose, the system should decide which

parsing results may answer facts and which parsing results remain only as data fields. It should test benign utility with legitimate text-bearing images, accents, noise and low-vision scenarios. Security controls need to keep effective input paths for real users. Acceptance should also take in the task completion of users of assistive technology.

4.8 Case Comparisons Must First Align the Attack Unit

Several studies that are often placed on the same leaderboard actually answer different questions. HOUYI's unit is the real application under test. Its 31/36 indicates that the researchers judged it vulnerable under heterogeneous applications and backends [liu2023houyi]. InjecAgent's unit is the valid benchmark case. About 24% there describes the attack objective metric in tool-integrated agents [zhan2024injecagent]. FigStep's average comes from typographic visual attacks on 6 open-source vision-language models [gong2025figstep]. GCG in turn involves white-box optimization, and PAIR black-box iteration [zou2023gcg, @chao2023pair]. They share no denominator and no attacker privileges, and they do not share the highest consequence layer.

For comparison, the applicability table below can be used. It presents the attack unit, privileges, budget, highest observed layer and evidence still missing side by side. An overall ranking serves only as entry-point navigation and must not replace item-by-item judgment made under deployment conditions:

Research family	Attacker entry point	Main budget	Typical scoring unit	Evidence supported up to	Evidence still to be supplied
Direct prompt injection (application-facing)	User-submitted content	Requests and queries	Application or task	Output and control deviation	Overall vulnerability rate across all applications
Automatic jailbreaking	Prompt or suffix	Gradients, queries, tokens	Request or objective	Content output	Tools having already executed
Indirect prompt injection	Web pages, email, tool results	Carrier writes and victim access	Benchmark case	Task deviation or simulated action	Production side-effect rate
Multimodal attack	Image or audio	Perturbation, parsing knowledge, queries	Image, question, or model configuration	Model output	Physical-channel success rate

Within a single study, the relevant structure must also determine what counts as an independent sample. Multiple models may share prompts, data and judges. The same prompt may be run under multiple random seeds. The independent unit still depends on the prompt and task layers. Picking the best attack and then reporting it produces selection bias. The unified prompt injection benchmark covers 5 attacks, 10 defenses, 10 LLMs and 7 task types. Its most important point is that the defense ranking varies with the model, task, attack and metric. Shared cells are suitable for showing conditional differences, whereas an industry average requires independent and comparable study units [liu2024formalizing].

For a single system, one can define nested stage counts on the same trial unit. Let N be the number of valid trials. Let n_C be the number of trials in which control deviation occurs. Let n_{CI} be the number of trials in which control deviation occurs and a dangerous plan is formed at the same time. Let n_{CIA} be the number of trials that, after both of the previous two have occurred, are also accepted by the executor. This yields $n_{CIA} \leq n_{CI} \leq n_C \leq N$, reported as:

$$\hat{p}_C = \frac{n_C}{N}, \quad \hat{p}_{I|C} = \frac{n_{CI}}{n_C}, \quad \hat{p}_{A|C,I} = \frac{n_{CIA}}{n_{CI}}.$$

When any denominator is zero, the corresponding conditional proportion is recorded as NA. One may also want to answer a question such as "a dangerous plan is formed without control deviation." That requires reporting the marginal counts n_I and n_A separately. They must not be placed directly into the conditional proportions above. A request that times out, or an output that cannot be parsed, is first reported separately along two dimensions: "valid trials" and "availability failures." Small samples must also be given a binomial interval. A "0 successes" result is written as a numerator, a denominator and its upper bound. Benign utility uses an independent

denominator of legitimate tasks. The costs caused by attack detection, human review and additional calls are recorded.

The scorer must also be subject to inspection. Content jailbreaking must at least distinguish refusal, harmless explanation and genuine completion of the dangerous objective. Prompt injection must at least check whether the original user objective has been changed. Tool scenarios must also read the executor state. A function name in the model text indicates only a candidate plan. Automatic judges are responsible only for initial screening. Boundary samples are reviewed by rules different from the attack model, or by humans. The review protocol, disagreements and final adjudications all enter the record.

Algorithmic Implications of White-Box Suffix Search

Methods such as GCG treat discrete token suffix search as a computable local optimization. At one position in the current suffix, the attacker reads the gradient of the target loss with respect to the token representation. That gradient screens candidate tokens which may reduce the loss. The attacker then forward-computes candidate combinations and keeps the better ones. Coordinate updates across several positions and several rounds build the suffix step by step [zou2023gcg]. The key premise is that the attacker can access sufficient model-internal signals, or at least optimize on a substitute model of the same architecture and then transfer.

The gradient supplies a local candidate direction only. It does not directly prove that a discrete replacement will necessarily succeed. Conversation templates, tokenizers, system policies, model quantization, and decoding all change the loss surface. A suffix that succeeds on one model may also fail on another model, or after a service update. Reproduction experiments must therefore fix the tokenizer, template, target loss, initialization, coordinate batch, iteration cap, and randomness. Publishing only the final suffix loses the attack budget and the failure trajectory.

When defenders use input perturbation or randomization, evaluation must look beyond replaying existing suffixes. It must add an attacker that knows the random distribution. Such an attacker can optimize the expected loss:

$$\min_a \mathbb{E}_{\xi \sim \mathcal{T}} [\mathcal{L}(M(T_{\xi}(g \oplus a)), t)],$$

Here T_{ξ} is the random transformation and $\mathcal{T}$ is the transformation distribution. The expression shows that an adaptive attack incorporates the defense into its objective. It does not mean that any attack can be solved effectively. SmoothLLM's results on some optimized suffixes demonstrate that perturbation can raise the cost. Retesting is still required after the attacker knows the transformation [robey2023smoothllm].

Black-Box Attacks, Judges, and Stopping Rules

PAIR formulates the attack as a closed loop [chao2023pair]: "propose candidates — query the target — read the results — have the judge score them — generate the next candidate". Attack model calls, target model calls, judge calls, and manual spot checks jointly constitute the cost of a black-box attack. Reports should give the total target queries, the total model calls, the number of queries reaching first success, and wall-clock time as separate figures. "rounds" is retained as a process field.

Stopping rules change the results. If the system stops after the first success, the metric is close to "at least one success within the budget". If every target runs all queries in a fixed way, the attack trajectory and repeatability can be estimated. Where the best result is selected from multiple candidates and then tested, training and evaluation must be separated. The final report then uses independent samples. When the target API exhibits rate limiting, content truncation, or version switching, the corresponding run should be marked as an environment change. It should be reported as a separate group from stable runs.

Judges require two-axis output. The first axis determines whether the attack objective is truly satisfied. The second axis determines whether refusal, disclaimer, or harmless substitution appears. Matching only affirmative words would count "Of course, I cannot assist with that" as a success. Relying only on another LLM in turn introduces model-related error. A workable approach combines rule-based initial screening, heterogeneous model review, and sampled human arbitration. It also saves the judge version and confidence. Boundary samples in high-risk categories all enter the human queue. Random samples are used to estimate the overall judge error.

Measuring Long-Context and Multi-Turn State

Many-shot exploits in-context demonstrations, and Crescendo exploits gradual interaction [anil2024manyslot, @russinovich2025crescendo]. Evaluation needs to save the complete message sequence, roles, truncation positions, and the system state of each turn. Suppose only the last user message is kept. Then it is impossible to know whether a change in safety behavior comes from the current prompt or from historical accumulation. When the context exceeds the model's limit, the service's truncation policy is also part of the attack surface. Dropping early system constraints and dropping attack demonstrations produce opposite results.

Multi-turn attacks report at least four quantities: the number of turns before the first control deviation, the number of turns before the first dangerous output, the total input and output tokens, and the number of recovery turns after a failure. The recovery-turn count measures whether the system continues to follow the attack objective after one deviation. If a new session completely clears the state, the result belongs to short-term context. If a summary or memory continues into a new session, it should enter the persistent state rules of Chapter 5.

From Prompt Strings to Policy Invariants

The ultimate object of defense engineering is to declare which variables are changed only by which principals. A banned-word list serves only as a local heuristic signal. The policy can be written as information flow invariants:

$$\forall d \in D_{\text{low}}, \text{influence}(d) \cap \{\text{policy}, \text{recipient}, \text{credential}, \text{approval}\} = \emptyset.$$

The invariant states that the system policy, the recipient, the credential, and approval accept input only from the corresponding trusted principals. Low-integrity data has no influence edge to these variables. The implementation can be typed fields, taint labels, a capability policy, or a reference monitor. The formula defines only the objective. Code correctness still requires runtime verification. Every invariant needs positive and negative tests. Can a legitimate user set the recipient, and can an external document set the same field? Can an authorized policy publisher update the rules, and can a model summary update the rules?

Invariants must also cover derivation. If the model summarizes a web page as "the user wants to send the report to a certain address," the new text still inherits the web page's provenance. If a tool converts an object into a URL, the taint continues to attach to the URL. If multiple sources are concatenated into a parameter, the final integrity is no higher than the least trusted input that determines that parameter. If any adapter loses a label, the policy sets the result to unknown and enters a restricted path.

4.9 Pre-Launch Checklist for Instruction Boundaries

During a system design review, checking can proceed in the following order. Every item is required to land on a component, an owner, allowed and forbidden probes, and failure semantics. Items with only a textual promise and no run receipt remain pending verification:

1. **Task boundary:** Are the user's original objective, the allowed sub-objectives, and the explicitly forbidden actions saved? Can a summary overwrite the objective?
2. **Input inventory:** Which of web pages, email, PDF, OCR, automatic speech recognition, tool returns, and cross-agent messages are written by third parties?
3. **Provenance propagation:** After parsing, translation, summarization, chunking, and variable assignment, are the principal, tenant, purpose, and integrity still present?
4. **Attacker budget:** Does the evaluation test each of these separately — a single input, multiple turns, a long context, black-box queries, and an adaptive attack after the defense is visible?
5. **Parsing differential:** Can human-visible content, the raw file layer, OCR, automatic speech recognition, and the model's actual input all be compared against one another?
6. **Scoring layers:** Does the report separate Y0 content, Y1 control, Y2 plan, Y3 execution, and Y4 real-world consequence?
7. **Benign utility:** Can legitimate external facts still complete the task? Are there thresholds for false refusal, latency, tokens, and human cost?

8. **Action independence:** Are high-impact parameters checked by a structured policy? Does the same model both plan and approve?
9. **Failure policy:** When a provenance label, a classifier, or the policy service is unavailable, does the system refuse, degrade, or silently allow?
10. **Regression and response:** After a model, template, parser, or tool update, are fixed and adaptive test cases rerun? Can trajectories be frozen and replayed by action?

The checklist runs through every release version. Each item must correspond to a runnable probe, a responsible person, a version, and a failure disposition. Take "provenance propagation is implemented". It should be proven by a document with a leakage probe that still retains a low-integrity label after OCR, summarization, and tool parameters, where a leakage probe is a test object that uses known unique markers or decoy records in controlled input to check unauthorized reads and leaks. Take "outbound sending is limited to authorized recipients". It should be proven by executor rejection and network evidence. A model refusal provides only output-layer evidence.

4.10 Common Misjudgments

Misjudgment 1: calling every malicious prompt jailbreaking. Mixing up the names leads teams to test only content refusal. Suppose the attack goal is to change the recipient, the retrieval scope, or tool parameters. Then the core issue is application control integrity. Write down the user goal and the highest consequence layer first, then choose the attack label.

Misjudgment 2: system prompt priority is enough to form a hard boundary. Prompt priority is still interpreted probabilistically by the model. It is suitable as a model-layer control. Provenance labels, type constraints, and external authorization carry information flow and capability control respectively. High-impact systems must assume that the model will occasionally misparse priority.

Misjudgment 3: black-box success is more "real" than white-box success. Black-box and white-box are only dimensions of attacker knowledge. Black-box studies may use large numbers of queries, automatic judges, and loose criteria. White-box studies may also be used to find structural weaknesses in real deployments. Realism depends on the complete capability preconditions, access conditions, and deployment chain. A single label describes only one of those dimensions.

Misjudgment 4: zero hits after input filtering prove safety. Zero hits constrain only the current sample, budget, and parser. Attackers can change the language, the modality, or the sharding scheme, or exploit tool return values. Reports need to give the attack space, the numerator and denominator, and a one-sided upper bound. An absolute guarantee requires support from stronger architectural invariants.

Misjudgment 5: the model emitting a tool call means the system has already been compromised. The output string may only reach Y2 dangerous plan. To prove dangerous execution, executor state, identity authorization, and side-effect evidence are needed. To prove real-world consequences, an external anchor is also needed. Layered records both avoid exaggeration and do not conceal an upstream failure.

4.11 Bringing It into Real Systems

Before a customer service center brings in external ticket attachments, it needs to review two paths at the same time: the user conversation and the attachment. If the current user submits the same exfiltration instruction directly, the attacker approaches the model policy boundary from the user input entrance. If the text is hidden in a ticket that the system reads on its own initiative, the attacker controls the external data entrance, and the first-broken position becomes the parsing boundary between instructions and data. The team first fills in $\kappa = (K, W, O, B, S, T)$ separately for the two paths. It then saves, along Y0–Y4, the model text, task drift, candidate actions, policy decisions, tool responses, and external receipts. The highest possible consequence layer depends on whether the application is connected to real tools, not on the payload text itself.

The same evaluation records white-box suffix search and black-box dialogue search separately. The gradient or probability access, the target model computation, and the number of suffix optimization steps that GCG requires belong to different budget units than PAIR's attack model calls, target queries, and visible replies. The team

reports these resources separately and converts them into latency or monetary cost when the business needs it. Each budget tier is fixed in advance, and all attempts within a tier are retained. Request errors enter the availability denominator. The automatic judge decides both “whether a refusal appears” and “whether the attack goal is substantively completed”. Borderline text is checked by rule review or a two-person spot check, so that a single “Of course I can, but I cannot help with that” is not counted as a success by surface word matching.

The PDF assistant needs to carry provenance identity through every transformation. The original file records the principal, tenant, hash, and purpose. Visible text, hidden text, and OCR results store the page number, region, and derivation relationship separately. The allowed-influence field limits external documents to participating only in fact extraction, while the recipient, confidentiality level, and publication status come from trusted fields. Tests change footnotes, tables, attachment properties, scan layers, and OCR parameters at the same time. They then observe which layer shows a difference. If the provenance label service fails, public question answering may continue under prominent labeling and a strict budget. Private retrieval enters a read-only restricted mode. External publication fails closed. The three action classes differ in blast radius and reversibility.

The meeting assistant further adds shared documents, persistent summaries, and task creation. A malicious participant can write the conversation directly. The document author relies on the system for indirect retrieval. The tool provider can control the tool description, interface definition, or return values. The three therefore differ in their writable set W , state capability S , and feedback O . The team forms a two-dimensional matrix from attacker defense knowledge and query budget. In each cell it observes benign completion, false refusal, attack control, candidate tasks, executor refusal, and cost at the same time. The matrix reveals differences in control at the static payload, adaptive retry, and tool supply chain entrances. It also assigns responsible components for the subsequent provenance, state, and capability layers separately.

Summary: First Hold the Boundary Between Data and Control

Jailbreaking tests whether the model policy is bypassed. Prompt injection tests whether the application goal is hijacked by external content. Semantic reconstruction, surface transformation, context accumulation, and automatic search keep the static prompt defense facing adaptation. Web pages, email, images, and audio also let data enter the same control context through different parsers. Reliable design therefore does not place its hopes on a stronger warning sentence. It preserves provenance, restricts data use, downgrades plans to candidates, and constrains capability outside the model.

The attacks discussed in this chapter can still be transient. After the malicious content leaves the context, the influence disappears with it. Retrieval-augmented generation and long-term memory change this premise. The next chapter treats indexes and memory as secure databases with identity, version, expiry, and deletion semantics. It explains how a single input becomes persistent state in future sessions.

Engineering acceptance should always retain both fixed attack regressions and defense-aware adaptive tests. The former hold the known entrances. The latter check whether controls truly constrain the mechanism.

\newpage

Chapter 5 Retrieval, Context, and Memory

An engineer puts a question to the enterprise assistant: “How should an older model of pump be started at low temperature?” The assistant pulls a maintenance manual, a forum post, and a conversation summary saved several weeks earlier from the knowledge base. The answer reads as complete. Yet it presents the forum’s temporary workaround steps as formal procedure, and it also cites another tenant’s equipment records. A team that examines only the final text may classify the problem as hallucination. Follow the system chain backwards, though, and at least three boundaries have already failed. Untrusted content entered the index. High similarity outweighed source quality. A cross-tenant record entered the context.

Retrieval-augmented generation (RAG) and long-term memory let the model use information outside the window. They also turn a single input into state that can be saved, recalled, and combined. The security question therefore expands. It no longer asks only “will this sentence fool the model”. The new questions are different. Who has the

right to write what? Under whose identity is the record saved? When and by whom is it retrieved? Does it come back to life after deletion? This chapter analyzes vector indexes and memory services as secure databases. Long prompts are only one form in which they expose content to the model.

Chapter Overview

This chapter breaks RAG into ingestion, chunking, indexing, retrieval, reranking, and context assembly. At each step it locates where trust changes. It then distinguishes among content poisoning, instruction poisoning, unauthorized retrieval, corpus extraction, and persistent memory attacks. Long-term memory unfolds along the three stages of "write—recall—use". Each stage has its own success metrics. Tenant, provenance, purpose, integrity, expiry, version, derivation chain, and deletion marker together form the record structure. The test matrix observes retrieval hits, model adoption, action propagation, benign utility, and deletion effectiveness at the same time. With that in place, readers should be able to separate secure-read policy from conservative-write policy. They should also maintain identity and lifecycle semantics for each record. Cross-tenant, persistent recall, derived deletion, and recovery-replay probes then verify the actual state.

Principle Background: How External Knowledge Becomes Current State

RAG encodes the current question as a vector. It retrieves candidate chunks from the authorized records in the vector database, and after reranking assembles them into the model context. Long-term memory compresses one observation into a record that can be saved across sessions and recalled later. The former mainly forms a read pipeline. The latter also includes state writes and a lifecycle. Their inputs include documents, metadata, and user and tool observations. Their mechanism state includes indexes, tenant labels, versions, derivation relationships, and deletion markers.

Users, downstream models, agents, and business processes consume answers, citations, candidate plans, and memory updates. The security surface therefore splits into a read side and a write side. Can an ingester poison the index? Can a retriever cross tenants and purposes? Will the model elevate text in a chunk into instructions? Does a summary lose provenance? Can a memory come back to life from a cache or a derived record after revocation? Similarity describes only closeness in representation. It cannot substitute for truthfulness, access control, or action authorization.

In the normal case, the enterprise assistant filters manuals by principal and project first. It then retrieves, reranks, and answers with sources within the allowed set. The boundary case is different. A malicious chunk reaches the top ranks and the model cites it, but the action layer does not adopt it. Alternatively, a record has already been deleted from the main index, and the summary cache still recalls it. The former case shows that the retrieval and content adoption chain partially holds, which is not equal to a real-world action. The latter case indicates that deletion propagation has failed. Below, the text records write, hit, adoption, propagation, and purge evidence separately. No single end-point accuracy should obscure the different control points.

5.1 RAG Is a Data Pipeline, Not a One-Time Prompt Assembly

The work of RAG can be written as the following chain of transformations. Each arrow represents a step of data processing. It also marks a place that should carry provenance, tenant, purpose, version, and deletion status. That way a chunk does not become plain text that no one can hold accountable after the transformation. When any field is lost in the record, downstream consumers should treat it as low-trust input, and write the reason for the downgrade into the receipt.

$$D \xrightarrow{C} \{c_i\} \xrightarrow{E} \{v_i\} \xrightarrow{I} \mathcal{K} \xrightarrow{Q(q)} R_k \xrightarrow{A} (y, \pi).$$


A document collection D passes through a chunker C to form chunks c_i . An embedder E turns the chunks into vectors v_i , and an indexer I writes them into the knowledge base $\mathcal{K}$. A query q triggers retrieval and reranking Q , which selects the top k results R_k . The context assembler A then has the model generate an answer y or a

candidate plan π . Identity, provenance, or purpose can be lost at every arrow. Record identifiers and policy receipts therefore need to be saved step by step.

At least four different classes of failure run along this chain. **Content poisoning** makes false facts or biased material rank highly under target queries. Two more classes are **instruction poisoning** and **unauthorized retrieval**. The first makes retrieved chunks carry control statements aimed at the model. The second is caused by tenant mapping, metadata filtering, or access control errors. A fourth is **corpus extraction**, which uses query and output feedback to recover private content. The four attack classes may share the same vector database, yet they have different assets and endpoints. Wrong answers, task hijacking, cross-tenant leakage, and text recovery rate should be reported separately with their own denominators and success definitions.

To read the figure, follow the upper retrieval-augmented generation swimlane through ingestion, indexing, retrieval, reranking, and context assembly. The lower long-term memory swimlane covers observation, writing, cross-session saving, recall, and use. When comparing the two paths, highlight who holds read rights and who holds write rights. Note also at which step provenance, tenant, purpose, and deletion status may be lost. "Entering the current context" belongs to the read path of a single request. "Being written as future state" is the cross-session persistence path.

RAG 与长期记忆的两条状态路径



Two swimlanes respectively show the data pipeline of retrieval-augmented generation and the state pipeline of long-term memory from observation and writing to future recall.

According to the figure, RAG and long-term memory share the problems of provenance and access control. Their state semantics and test endpoints differ. The figure does not mean that a retrieval hit necessarily leads to model adoption. Nor does it mean that a successful write will necessarily be recalled in the future. Hits, adoption, writes, recall, action propagation, and deletion effectiveness all need their own logs and denominators. Without actual recall evidence from a future session, the available evidence confirms only that the record can be written. Whether the persistent attack takes effect remains unknown.

"The document is already in the store" only means that the ingestion process accepted it. A public web page may suit a product parameter question. The payment recipient, by contrast, follows from the authorized goal. Internal meeting minutes are open only to the corresponding project group. A model-generated summary may faithfully relay the source, or it may drop a negation. Similarity answers "whether the query and the record are close in representation space." Truthfulness, read permission, and action purpose are answered by provenance, ACLs, and purpose rules respectively.

Therefore, the order of filtering before retrieval matters. The system should build the allowed set from principal, tenant, purpose, type, status, and expiry. Only then should it compute similarity inside that set. It should not first put the most similar content from the whole store into the prompt and then ask the model to keep it secret on its own. Access control may act only at the final display layer. By then the sensitive record has already crossed the U2 retrieval and state interface.

5.2 What Targeted Poisoning and Corpus Extraction Prove

For the "question-target answer" pairs chosen by the attacker, PoisonedRAG constructs malicious documents. It optimizes retrieval hits and generative adoption together. Injecting 5 malicious texts per target question into a knowledge base containing millions of texts yields 90% ASR, the study reports [zou2025poisonedrag]. This result shows that a small number of targeted records can control specific queries under a specific configuration. "5 per target question" is the attack budget of that protocol. Several conditions must hold for the conclusion: the embedding model, chunk boundaries, retrieval depth, the reranker, the generator, and whether the attacker knows the target question.

Decomposing the attack into stages makes it clearer. Let H_r denote that the malicious record enters the top k , H_c denote that the model adopts the attack content, and H_y denote that the final answer reaches the target. The three events carry precedence conditions, and an unobserved stage cannot be filled in from the terminal label. Once they are recorded with the same sampling unit, the joint endpoint of this three-stage attack chain satisfies:

$$P(H_r \cap H_c \cap H_y) = P(H_r)P(H_c | H_r)P(H_y | H_c, H_r).$$

This is conditional-probability bookkeeping and does not require the stages to be independent. If the malicious document does not enter the top k , the problem may lie in retrieval. If it enters but is not adopted, the generator or the context structure played a role. If the answer hits the target but the action gate refuses execution, the highest consequence layer still stops at the output or the plan. Looking only at the terminal ASR compresses the contributions of different controls into a single number.

Spill the Beans studies a different kind of endpoint. Through self-generated queries it induces a production custom GPT to output private corpora. It recovers 41% from a book of about 77,000 words and 3% from a corpus of about 1,569,000 words [qi2025spillbeans]. Here the denominator is the number of words in the corpus. The result is verbatim coverage, which is a different metric from prompt-level attack success rate. The lower percentage on the larger corpus describes the recovery coverage of that configuration. The query budget, repetition, snippet length, and verifiable prefix all affect the recovered amount.

These two works lead to the same engineering conclusion. A knowledge base is a queryable data asset. It needs ingestion provenance, access control, query tracing, rate and budget limits, anomalous-similarity monitoring, and output citation. The system prompt covers only model behavior. Database security is still borne by data and identity controls. OWASP's engineering classification of vector and embedding risks likewise emphasizes multi-tenant leakage and poisoning. Such risk lists provide a control vocabulary, and the concrete effect of a defense requires runtime measurement [owasp2025vector].

5.3 Provenance Must Pass Through Chunking, Summarization, and Reranking

The most common structural loss in RAG occurs after transformation. The original text has an author, a repository, a path, an access level, and a time, but after chunking only a passage of text may remain. A summary compresses multiple sources into a single statement, and the reranker outputs only scores. Provenance metadata must propagate with the derivative. Otherwise the model and the action gate cannot distinguish "explicitly provided by the user", "approved within the tenant", "scraped from a public web page", and "inferred by the model itself".

Each knowledge record can be written in the structure below, and retrieval, reranking, context assembly, and deletion flows can all consume these fields. Merely storing metadata in a database without enforcing policy does not constitute access control. Derived summaries and caches must also inherit the restrictions of the original record and be able to trace back to the original object:

$$r = (id, t, s, o, u, l_c, l_i, p, \tau, v, d, h),$$

Here t is the tenant, s is the writing principal, o is the provenance, and u is the purpose. l_c and l_i denote confidentiality and integrity respectively. p is the reader or tool access policy, τ is the validity period, v is the version, d is the derivation relation, and h is the content hash. A vector is only an index field of the record. Tenant,

provenance, purpose, and authorization still have to be carried by separate fields and kept associated on every derivation and copy.

At ingestion, a provenance allowlist and signatures can confirm "where it came from", and content scanning can find obvious payloads. Manual quarantine can handle high-risk material. The factual correctness of the text still has to be verified against evidence from the corresponding domain. At retrieval, ACLs and purpose filtering should precede similarity. The share of a single source, anomalous neighbor clusters, and a sudden influx of targeted snippets can serve as poisoning signals. At generation, answer citations should point back to verifiable records. Document names serve only as auxiliary display.

Provenance propagation must also handle evidence correlation. A single web page scraped ten times, split into ten chunks, or summarized into ten items remains one source. If the system treats them as ten independent pieces of evidence, ranking and majority voting will amplify a single-source error. The `derived_from` of a record should preserve the derivation graph. When multiple results come from the same root node, the aggregator lowers their independent weight and hands conflicts to an explicit arbitration policy.

Tenant boundaries depend on identity and access control, not on semantic similarity. Two customers may both use "annual plan", "administrator", or similar project names, and vector distance cannot express authorization. The index can be physically partitioned, or tenant filtering can be enforced at the logical layer. Whatever the implementation, acceptance must include cross-tenant leakage probes and tests that bypass metadata filtering. Even if an output filter deletes a sensitive sentence, an unauthorized read that has already occurred is still a security event. The trace should save the asset and the path that hit it.

5.4 Long-Term Memory Turns One Input into a State Transition

Long-term memory carries one more layer of identity and temporal meaning than RAG does. RAG mostly stores shared knowledge. Memory may also hold "what this user likes", "what the last task did", and "which policy should be adopted in the future". Model outputs, tool results and web page content may all be summarized automatically and written into the store, to be recalled by similarity in a later session.

Let the session state be $\mathbf{z}_t$, the observation of this round be $\mathbf{o}_t$, the write policy be $\mathbf{W}$, and the recall policy be $\mathbf{R}$. The state relation below splits "whether to store" from "whether to use in the future" into two transitions that can be verified independently. It requires the write receipt and the future recall log to share one record identifier, and an intermediate purge should also register as a state event:

$$\mathbf{z}_{t+1} = \mathbf{W}(\mathbf{z}_t, \mathbf{o}_t), \quad \mathbf{m}_t = \mathbf{R}(\mathbf{z}_t, \mathbf{q}_t), \quad (\mathbf{y}_t, \pi_t) = \mathbf{M}(\mathbf{g}_t, \mathbf{m}_t).$$

An attack usually passes through three stages before it has a persistent effect. The **write** stage lets poisoned content enter $\mathbf{z}_{t+1}$. The **recall** stage brings it into $\mathbf{m}_t$ under a future query. The **use** stage makes the model treat it as a fact, a preference, a policy or a parameter. Each of the three stages should be counted separately. If the write succeeds but the content is never recalled, the conclusion stops at the state layer. If the recall succeeds but the action gate refuses, the conclusion stops at the context or plan layer.

AgentPoison poisons an agent's long-term memory or RAG knowledge base directly. Across three types of agents it reports an average ASR of no less than 80%, an impact on benign performance of no more than 1%, and a poisoning rate below 0.1% [chen2024agentpoison]. Those numbers are tied to specific agents, retrievers and trigger conditions. Other memory products need retesting against their own write and recall mechanisms. MINJA narrows the attack privilege to ordinary query interaction. The attacker then relies on automatic saving and never writes to the database directly, which shows that an automatic-save policy can itself become a write entry point [dong2025minja].

A persistent attack can also be fragmented and dormant. FragFuse splits a violating objective into fragments that look ordinary across turns, which shows that item-by-item inspection may miss the combined meaning. A-MemGuard tries to reduce memory poisoning through multi-path consensus [rao2026fragfuse, @wei2025amemguard]. Both directions suggest that the isolated security score of a single record covers only a local part. When an action occurs, the system must reconstruct the complete derivation chain of the records that

took part in the plan. It must then check whether the combination of fragments has changed the task or the information flow.

Two cases support the threat model that "external content can achieve persistence". In SpAIware, a web injection enters memory and then persists across sessions. In Morris II, an image-text payload self-replicates inside an experimental email agent [rehberger2024spaiware, @cohen2024aiworm]. Neither proves that large-scale autonomous propagation has already appeared in reality. The limits of the evidence still depend on the environment, the permissions and the real side effects.

5.5 Writing Requires More Conservatism Than Reading

Many systems treat reading as high risk, yet they let the model save "useful information" automatically. That creates an asymmetry. A single low-integrity web page can enter a highly persistent state, and later it may be reinterpreted under different tasks and permissions. A robust policy makes writing an independent decision. The model proposes memory candidates, and an external identity and purpose policy approves any trust elevation.

A memory record needs at least the fields below. For each field, the record also specifies the deny, quarantine or downgrade semantics that apply when it is missing. Only then can summaries, caches and derived records accept lifecycle controls consistent with the original record. Only then can revocation locate residual state along the derivation relation and verify the purge result after a restore replay:

Field	Question answered	Typical control
tenant_id, principal	Whose state this is	Enforce tenant and principal matching
origin, write_channel	Where the content entered from	Distinguish user statements, web pages, tools, and model inference
type, purpose	What it may be used for	Typing into preference, fact, policy, and credential reference
integrity, confidentiality	How high-risk an action it may affect, and by whom it may be seen	Information flow and action parameter constraints
ttl, version	When it expires, and which version is in effect	Expiry invalidation and conflict management
derived_from	Which records produced it	Combination checks and provenance replay
review_state	Whether it has external confirmation	Trust elevation requires an independent principal
tombstone	Whether it has been revoked	Synchronized deletion across the online index, cache, and backup

Updates follow an explicit version chain. A new value forms a successor version and points back to the previous value. Memories that conflict with each other enter a disputed state, and a temporal update on its own does not raise trustworthiness. On a read, the task filters first by identity, purpose, type, state and validity period. Only then does it perform vector retrieval. Low-integrity records can help generate candidate explanations. Payment accounts, code release targets and physical action parameters accept only trusted sources that meet the requirement.

Deletion has to reach beyond the original text: vectors, summaries, caches, search snapshots and exported files. A more reliable order writes a global deletion marker first and withdraws the item from online retrieval. The original text and all derivatives are processed after that. When restoring from a backup, replay the deletion markers so that content already revoked cannot re-enter the index. A deletion proof needs to record the object, the derivation scope and the execution time. It must also record the operational information that must still be retained by law.

This record structure raises storage, indexing and policy costs, and it may also lower recall. A security evaluation must watch benign task completion at the same time. If the system simply rejects every external record, the attack surface shrinks, but the business value of RAG disappears with it. A reasonable goal is not "zero recall". It is to maintain useful retrieval within the authorized set, and to keep low-integrity content from silently crossing into high-impact actions.

5.6 Capability matrix for retrieval and memory attackers

The attackers who face RAG and memory systems are not only knowledge base administrators. Public web page authors can control the material that gets ingested. Ordinary tenant users can issue large numbers of queries. Collaborators with write access can submit documents. Third-party connectors can supply tool results. The model itself also generates summaries and memory candidates. Calling all of these actors "poisoners" hides how much their capabilities differ over writing, recall and feedback.

Attacker capability can be written as the vector below. Direct database privileges, ordinary query writing, feedback visibility and cross-session duration then stop collapsing into one "can poison" label. Each result is compared only with the attack set that shares the corresponding privileges and duration, and unknown capabilities are not filled in by default:

$$\alpha = (w_q, w_d, w_m, r_o, r_s, b, \rho),$$

Here w_q indicates whether the query can be controlled, and w_d indicates whether documents pending ingestion can be submitted. w_m indicates whether long-term memory can be influenced directly or indirectly. r_o indicates whether the final output can be observed, and r_s indicates whether retrieval scores or the top k can be observed. b is the query, document and time budget. ρ is the target scope, that is, a single question, a single user, a certain tenant or a shared knowledge base.

PoisonedRAG assumes the attacker constructs malicious documents for selected target questions. Its capability concentrates in w_d and on the target knowledge [zou2025poisonedrag]. The attacker in Spill the Beans mainly controls the query and observes the output. Repeated interaction recovers the private corpus, and the focus falls on w_q, r_o, b [qi2025spillbeans]. AgentPoison needs a more direct path for knowledge or memory poisoning. MINJA narrows the entry point to ordinary interaction that induces automatic saving [chen2024agentpoison, @dong2025minja]. Document write, query feedback, memory write and automatic save capability must be recorded as four separate entries. "Whether the database is writable" covers only one of them.

Defenders should build the privilege matrix below by actor and verify the allowed and forbidden paths separately. One actor can hold different privileges at the ingestion, retrieval, write and deletion stages, so a single global role cannot summarize it. Temporary privileges, delegated writes and batch tasks must also be listed as independent actors:

Actor	Objects writable	Feedback visible	Legitimate business need	Main abuse path
Public web page author	Original web page	Whether cited or downstream changes	Provide public facts	Ingestion poisoning, indirect instructions
Ordinary tenant user	Queries, sessions, candidate memories	Answers and personal history	Use the assistant and save preferences	Corpus extraction, MINJA-style writing
Content maintainer	Documents and metadata	Index state, previews	Publish knowledge	Targeted poisoning, unauthorized labels
Connector or tool	Tool returns and synchronized data	Call results	Connect to business systems	Provenance forgery, cross-tenant propagation
Model and summarizer	Derived text	Subsequent recall	Compression and personalization	Misattribution, label loss

The "legitimate business need" column in the matrix matters. Public web pages are supposed to enter retrieval, and ordinary users also need to save preferences. Revoking writing entirely would break the product goals. Controls should narrow objects and purposes instead. Web pages may write public fact candidates, and personal policies are confirmed by the user they belong to. Users may propose personal preferences, and organizational procedures are signed by an authorized publisher. The summarizer may generate derived records; their integrity level is inherited from the source or promoted through independent verification.

The attack scope ρ determines the test denominator. For an attack on one known target question, the denominator is the target query. Shared library contamination needs to cover non-target queries and other tenants. Corpus extraction needs to report the recovery unit, the query budget and the deduplication method. Long-term memory

additionally needs to report survival time and future recall opportunities. If the report only says "attack succeeded", the reader cannot tell whether one answer was controlled or a cross-session state was established.

5.7 Algorithmic order and implementation constraints for secure retrieval

A secure retriever constructs the authorized set first, then performs nearest neighbor search. The logical order runs as follows. Authenticate the requesting principal. Construct the authorized set according to tenant, purpose, data type, confidentiality, state and time limit. Search within the authorized set. Rerank the results for provenance and relevance. Limit concentration from the same source. Finally, hand the records and complete metadata to the context assembler.

The indented block below gives the design order for secure retrieval. Deployment must implement and verify it item by item. The order itself is part of the security semantics; running full-database similarity retrieval and then filtering what is displayed is not equivalent. Tests must also cover alternative paths such as cache hits, retries and batch processing:

```
principal = authenticate(request) scope = authorize(principal, request.purpose, request.tools) candidates =
filter(index, tenant=scope.tenant, acl=principal, type=scope.allowed_types, state="active", ttl>now) neighbors =
vector_search(candidates, request.query, k_pool) ranked = rerank(neighbors, relevance, provenance, freshness)
selected = diversify_by_root_source(ranked, k) return attach_lineage(selected)
```

In database deployment, these constraints must be implemented and verified item by item. Authentication and authorization run outside the model. If the vector backend pushes filtering down, it must be enforced on all query paths. Same-source dispersion is built on a reliable derivation graph. Results with attached derivation relations must continue to be preserved by the summarizer and tool adapters.

Five implementation failures are common. First, the application retrieves from the whole database and then filters at the business layer, so caches or debug logs have already touched records the principal is not authorized to access. Second, the filter constrains only documents and misses vector metadata, summaries or attachments. Third, some optional parameter in the query can turn tenant filtering off. Fourth, the reranker receives unauthorized candidates, so leakage outside the model can still occur. Fifth, the cache key omits tenant and purpose, so the previous principal's results are reused by the next principal.

Nearest neighbor retrieval also amplifies clustered poisoning. The attacker generates several similar records around the target query, and they may occupy the top k at the same time. Limiting the share of a single root source, detecting anomalous neighbor density, and comparing different embedders or rerankers can raise the attack cost. These are all anomaly signals; they do not prove that the content is true. PoisonedRAG pursues both retrieval hits and the target answer, which shows precisely that content filtering only at the generator's end is insufficient [zou2025poisonedrag].

Answer citations must also be verifiable. The system returns the record identifier, the version, the snippet location, the root source and the access decision. The document title serves only as a readable name. When a user clicks a citation, authorization must be performed again; the fact that the model once read the content did not grant permission to share it across tenants. If a citation points to a page that is updated later, the tracking record should be able to recover the content hash or snapshot that was used at the time.

Hidden state in caches, batch processing, and connectors

Three further categories of state sit outside the retrieval chain and are easily left unrecorded. The query cache may store the top k or complete answers. The batch processing system may place documents from multiple tenants into the same queue. Connectors periodically synchronize external permissions and content. Even if the main index is correct, any bypass path may reintroduce expired, unauthorized or deleted records. OWASP's engineering classification of vector and embedding risks gives prominence to multi-tenancy and poisoning. Actual controls must also cover these derived services [owasp2025vector].

The cache key must at least bind tenant, principal, purpose, authorization policy version, index snapshot and query normalization result. If only the query text forms the key, the same question shares answers across different

tenants. If the key binds the tenant but not the principal, it may in turn bypass document-level ACLs. On a cache hit, current authorization and deletion marks should still be checked. Sensitive answers should use a short TTL. Revocation events should proactively invalidate the relevant keys.

Batch ingestion gives every object its own identity and its own failure state. When one file fails to parse, the objects that follow reinitialize from their own metadata. Cross-tenant batches still need isolation in temporary directories, logs, and error queues. A bulk rebuild must bind the switch from the old index to the new one to a version. That way no query reads two authorization snapshots at the same time. Once the rebuild finishes, the previous version is deleted, and the evidence of rollback and of deletion-mark propagation is preserved.

Connector synchronization must separate content changes from permission changes. An external document may keep its content while its sharing scope narrows. The system must then withdraw unauthorized access first and wait for the content increment afterwards. If the connector cannot read permissions for a time, a high-confidentiality repository should suspend synchronization, or retain only verified snapshots. Connector tokens use read-only, minimal-scope, short-lived identities. Writing back to the source system requires a separate, independent authorization.

Bypass tests cover several cases. The same query must stay cache-isolated across principals. A deletion must invalidate the cache. Permission narrowing must land before content synchronization. Error queues must not leak body text. Index switching must not mix authorization snapshots. Deletion marks must still take effect after backup recovery. Each test preserves the request identity, the cache key, the index version, and the hit records. That way an engineer can locate the branch from which the state reappeared.

Concurrency opens a time gap between check and use. At retrieval time the record still holds its privileges, yet those privileges may already have been revoked before the model generates or the tool is called. A new approval, conversely, takes effect only for authorized reads that come after it. High-impact actions re-verify the record version, the ACL, and the deletion status of the participating parameters before execution. They also bind the verification snapshot to the action identifier. If a record changes after the plan is formed, the action is replanned or confirmed, and the expired context becomes invalid with it.

Consistency requirements need not push all question answering onto strong transactions. Public low-risk content can accept brief caching, while private reads and external actions use stricter snapshots. The system selects the consistency level by asset and by action, and records the actual level in the logs. Otherwise teams see only "eventually consistent" and cannot determine which moment's privileges the security decision used.

5.8 Memory policy engine and state machine

Memory writes can run as an explicit state machine rather than one database insert. Candidate records move through proposed, validated, active, disputed, expired, and tombstoned in turn. The model can only propose. Type validation, together with identity and purpose policies, determines whether a record enters validated. Records that require user confirmation or external evidence become active only after that confirmation. Conflicting records enter disputed, expired records are no longer recalled, and deletion marks invalidate all derived records.

The state transition function can be written in the following form. Policy version, identity, source, and expiry state are its decision inputs. Storing only the text content would make it impossible for future recall to reconstruct the authorization context of the time. Rejected writes must also produce a receipt, to make it easy to distinguish policy blocking from service failure:

$$s_{t+1} = F(s_t, e_t, \phi, \eta),$$

Here s_t is the record state, and e_t is a write, confirmation, conflict, expiry, or deletion event. ϕ is the deterministic policy, and η is the authorization of a named principal. A model score can serve as one signal for e_t , while η is still issued by a named principal. A web page summary that the model rates "trustworthy" still retains the web page's integrity level. Organizational procedures need to be confirmed by an authorized publisher or a signed source.

Read policy checks state, purpose, and action impact at the same time. A given personal preference is valid for format selection, whereas a security policy is determined by the policy repository. A given high-confidentiality record can serve internal answers, while external email accepts only the corresponding disclosure authorization. A

given low-integrity troubleshooting clue can trigger further queries, whereas shutting down an interlock requires a high-integrity procedure and on-site authorization. Purpose type and action impact jointly determine the usable scope:

$$\text{allow}(r, a) = \mathbf{1}[\text{ACL}(r, a)] \mathbf{1}[\text{purpose}(r) \supseteq \text{need}(a)] \mathbf{1}[\ell_i(r) \geq \ell_i^{\min}(a)] \mathbf{1}[\ell_c(r) \leq \ell_c^{\max}(a)].$$

In application, the four indicator conditions here should be combined with logical and. The formula displays the required conditions as parallel terms; it does not interpret them as independent risks. If any condition is unknown, a high-impact action should be denied or moved to confirmation. A low-risk answer can be used in a downgraded way, but the source uncertainty must be marked.

The failure paths of memory systems often span multiple stores. An online record is deleted while the vector still exists. After the vector is withdrawn, the summary cache continues to recall it. After the main database is cleaned up, analytics exports and backups write it back again during recovery. The derivation graph must connect the original text, vectors, summaries, caches, and exports. Only if the deletion mark propagates before physical cleanup and has an order higher than historical writes during recovery can the revoked state be prevented from coming back to life.

Composite attacks also require re-aggregation at action time. FragFuse splits the target into fragments across turns, and no single piece of content may reach the blocking threshold [rao2026fragfuse]. Therefore the policy checks each memory. It also checks the set of records participating in the same plan, the number of root sources, and the data flow produced by the combination. Ten derived records from the same web page still count as one root source, and majority voting does not raise their integrity.

Queryable State, Current Context, and Parametric Memory

A privacy incident must first locate where the data resides. Current-context leakage occurs when the system prompt, user-uploaded files, tool results, or other tenants' records enter this model input. RAG corpus extraction targets a queryable knowledge base. Long-term memory leakage involves cross-session state. Training data extraction recovers sequences from parametric memory. All four paths may manifest as "the model said the secret," yet they require different evidence and controls.

Training data extraction research recovered verbatim training sequences on the GPT-2 family through generation and candidate ranking. Later work studied scalable extraction on aligned production models [carlini2021extracting, @nasr2025scalable]. Duplication, guessable prefixes, sampling budget, interface, and external verification all affect these results. The results describe sequence recovery capability under the corresponding conditions. Neighbourhood Comparison studies membership inference, which judges whether a sample participated in training. Verbatim text recovery is a stronger and different endpoint [mattern2023membership].

Spill the Beans, by contrast, targets a private RAG corpus. There the attacker queries so that the retriever fetches and the generator outputs spans [qi2025spillbeans]. If logs show that the target text already appeared in the retrieval context, the first-broken interface may lie in access control or query policy. If the context contains no target record yet the same sequence emerges from the model parameters, another set of verification is required. System prompt leakage usually belongs to current-context isolation failure, whereas training data extraction describes the parametric memory path. The two are recorded separately.

Model stealing is yet another different asset. Black-box research can recover a model's functionality or part of its parameters. One work reports recovering projection layer information of a production model under a specific API and cost setting [carlini2024stealing]. Projection layer recovery does not equal copying the complete model, nor does it directly indicate RAG data leakage. An engineering ledger should label assets as current context, retrieval store, long-term memory, training corpus, or model parameters. That labeling keeps a single control from being mistakenly expected to cover the whole privacy surface.

The corresponding controls are also layered. Current context uses tenant isolation, minimal prompts, and output data flow. RAG uses pre-retrieval ACLs, query budgets, and citations. Long-term memory uses write subject, purpose, TTL, and deletion. Parametric memory relies on training data governance, deduplication, privacy evaluation, and interface rate. Model stealing additionally requires query monitoring and model asset protection.

The system prompt holds only the minimum secret information needed to complete inference. Real credentials are fetched on demand at execution time from the external capability service of Chapter 6.

For verification, leakage probes that differ in provenance but look similar on the surface can be deployed. A current-context probe appears in only one request. A RAG probe is stored in a controlled tenant record. A long-term memory probe is explicitly written and then triggered in a future session. Parametric memory tests use real training or corresponding substitute evidence. Transient context still falls under the input layer. The leakage paths of the different probes are recorded separately, so that "the model output the same string" can be traced back to the actual data source.

5.9 Comparable Boundaries of the Four Research Types

Ranking RAG and memory research on a high-low leaderboard flattens different attack stages. A more appropriate comparison targets attack preconditions, stage endpoints, and control implications. It maps write, hit, adoption, recall, action, and erasure each to observable evidence. Only when the target quantity and the denominator agree do numbers enter direct comparison; otherwise they remain side-by-side mechanistic evidence.

Case	Main precondition	Primary endpoint	Numeric unit	Engineering implication
PoisonedRAG	Malicious documents aimed at the target question can be constructed	Retrieval and target answer	ASR over target queries	Measure hit, adoption, and answer separately
Spill the Beans	Repeated querying while outputs are observed	Private text recovery	Corpus word coverage rate	Protect the query database as a data asset
AgentPoison	Knowledge or long-term memory can be poisoned	Agent behavior after triggering	ASR for the paper's agent configurations	Measure the write entry point together with benign utility
MINJA	Ordinary interaction can trigger auto-save	Write, future recall, and use	Multi-stage success	The auto-save policy is a security interface

PoisonedRAG reports 5 malicious texts per target question and 90% ASR, bound to a targeted setting in a knowledge base of millions of texts [zou2025poisonedrag]. The 41% and 3% of Spill the Beans are word coverage rates under two corpus configurations [qi2025spillbeans]. AgentPoison's average ASR, benign performance impact, and poisoning rate come from an aggregation of three agent types [chen2024agentpoison]. The denominators for the three numbers are target queries, corpus words, and agent configurations, in that order. They are suited to side-by-side expression rather than arithmetic averaging.

MINJA's value lies in changing the attacker's privileges. An ordinary query can influence future memory through the auto-save path, without requiring direct database access [dong2025minja]. When comparing with AgentPoison, the focus should be the write entry point, the recall mechanism, and attack survival, rather than comparing only the final percentage. A-MemGuard represents defense consensus, while FragFuse represents compositional attacks. Evaluation should observe source independence, long-term benign utility, and extra tokens at the same time [wei2025amemguard, @rao2026fragfuse].

When research cases enter engineering decisions, four questions can be asked. Does our attacker have the same write or query capability? Do our retrieval and memory implementations share the key mechanism? Does the research endpoint reach the action layer we care about? Can we reproduce the stage counts under the same budget? Only when all four are supported can a research result participate in a launch threshold; otherwise it still serves as a threat lead that triggers local testing.

5.10 Worked Example: Separating Device Maintenance Knowledge from Personal Memory

The opening engineering assistant can divide information into three logical stores. The first is the version-controlled official manual, which only the device team may publish and which records model, applicable temperature, and effective date. The second is the community experience store, which admits forums and tickets but has lower integrity and is used only to generate troubleshooting leads. The third is personal memory, which

stores engineers' preferred output formats and recently accessed devices and does not store unconfirmed safety procedures.

A user query is first parsed into tenant, device model, and task purpose. The retriever enforces ACLs within each store, then returns candidates with provenance. The reranker can consider semantic relevance, version recency, and source tier at the same time. An effective mandatory procedure has higher execution priority than a forum post. If community experience conflicts with the official manual, the answer explicitly presents the differences side by side and restricts executable steps to the official procedures.

One attack test writes a targeted post into the forum store. The post claims that "disabling the interlock at low temperature raises the startup success rate." In the body it asks the assistant to save that suggestion as a user preference. The expected safe behavior is as follows. The post may be recalled as a low-integrity troubleshooting lead. Personal preferences accept confirmation only from the owning user. Disabling the interlock requires an official procedure and independent authorization. The answer's citations clearly mark the source conflict.

The test record is divided into five endpoints. Was the record ingested? Did it enter the top k ? Was it adopted by the model? Was it written into long-term memory? Did it form a high-impact plan? Then add benign utility. Does legitimate community experience help resolve non-safety-critical faults? Does an updated official manual promptly supersede the previous version? After the user requests deletion of a personal preference, do both the online index and backup restoration remain revoked? In this way, one test can show at which layer the control occurred, rather than leaving only "the assistant answered right/wrong."

Conflict and time must also be tested. Two summaries from the same origin count as one root source. The currently effective file has higher execution priority than an expired procedure. A low-integrity record still inherits its original tier after being cited many times. Teams can write these properties as invariants and keep regressing them when the chunker, embedder, reranker, or summarizer is updated. State safety is maintained by the whole derivation chain, not by a single judgment of one component.

5.11 Metrics, Denominators, and Time Windows

Retrieval evaluation should publish at least four groups of metrics. The first group is ingestion and indexing. It asks how many malicious or unauthorized records reach the active index, with submitted records as the denominator. Retrieval is the second group. For pre-fixed queries it reports the proportion of target records entering the top k , the average rank, and same-source concentration; valid queries form the denominator. The third group is generation. It asks whether the model adopts the target content, whether it gives verifiable citations, and whether it leaks unauthorized text. Here the denominator is queries that obtained valid retrieval context. Action is the fourth group. It asks whether a low-integrity record enters a high-impact parameter and whether the capability gate accepts it, with trials that form candidate actions as the denominator.

Memory evaluation must add a time dimension as well. A single write, a single recall, and sustained survival answer different questions. Reports should therefore fix the observation window and the erasure operation, and should state that the state beyond the window remains unknown. When defenses are compared, retention and the mistaken deletion of normal memory should also use the same timeline:

$$\begin{aligned}
 p_W &= \frac{n_{\text{written}}}{N_{\text{write attempts}}}, \\
 p_R(\Delta t) &= \frac{n_{\text{recalled at } \Delta t}}{N_{\text{eligible future queries}}}, \\
 p_P &= \frac{n_{\text{target plans}}}{n_{\text{recalled}}}, \\
 p_E &= \frac{n_{\text{executed target actions}}}{n_{\text{target plans}}}, \\
 L &= \text{last observed active time} - \text{write time}.
 \end{aligned}$$

p_W is the write proportion. $p_R(\Delta t)$ is the recall proportion after elapsed time Δt . p_P is the proportion of recalls that form a target plan, and p_E is the proportion of target plans that further form an executed target action. L is

the observed survival time. When the denominator is zero, record the corresponding value as NA. Test termination affects survival time, so reporting "last observed still valid" is more accurate than claiming a true lifetime.

Deletion evaluation must extend from the user action to derived artifacts. The denominator is the object to be deleted together with its known derived nodes. The metrics cover online withdrawal latency, cache invalidation, export handling, reappearance after backup restoration, and the completeness of the deletion proof. If the derivation graph itself is incomplete, compute deletion coverage over known nodes only. Unknown derived nodes should be treated as a verifiability gap.

Benign utility includes legitimate documents entering the top k , the factual correctness of answers, citation usability, preference recall, task completion, and latency. Security policies commonly cause side effects. They delete ordinary facts needed to complete the task, downgrade all external sources to unusable, and make users confirm repeatedly. Reports should therefore plot a safety-utility-cost frontier, rather than optimizing only the lowest attack metric.

Cross-tenant testing with leakage probes can only prove whether a specific path leaks. It does not mean that all real secrets are equally reachable. The test record must name the record that contains the probe, the tenant, the query set, the retrieval depth, the model version, and the detection rule. Any request error, index failure, or empty answer is listed separately as availability. It does not count as a security success.

5.12 RAG and Memory Go-Live Checklist

1. **Ingestion principals:** Who can submit documents, sync tool data, or trigger autosave? Does each entry point hold an independent identity?
2. **Provenance and rights:** Are the original location, publication time, hash, license, and crawl time stored? Can derivatives be traced back to the root source?
3. **Tenant filtering:** Is the ACL enforced before vector search and reranking? Does the cache key include the tenant, the principal, and the purpose?
4. **Purpose typing:** Do facts, preferences, policies, credential references, and model inferences carry different record types?
5. **Integrity escalation:** Can the model approve a summary it generated itself as a high-trust record? Who signs off on external confirmation?
6. **Retrieval robustness:** Are same-source concentration, anomalous neighbor clusters, and single-source share bounded? When these signals fail, how does the system degrade?
7. **Action linkage:** Can a low-integrity record directly determine a recipient, an amount, a code target, or physical parameters?
8. **Temporal semantics:** Before retrieval, do TTL, version, conflict, and dispute status take effect?
9. **Deletion and recovery:** Do deletion markers cover the original text, vectors, summaries, caches, exports, and backup restores?
0. **Stage-wise evaluation:** Are ingestion, pre- k hits, adoption, writing, recall, use, execution, and benign utility reported separately?
1. **Logging and alerting:** Do cross-tenant hits, anomalous queries, bulk restores, and post-deletion recall trigger actionable events?
2. **Change regression:** After a chunker, embedder, reranker, summarizer, or policy update, is the same stage matrix rerun?

Every question must land on evidence. For example, confirming that "the ACL is enforced before retrieval" requires a query plan, service logs, and a cross-tenant probe working together. A recovery drill is needed to confirm that "deletion is complete", not just a check of the primary table. To confirm that "autosave is safe", ordinary users must attempt, through conversation, to write records of different types and integrity levels.

5.13 Common Misjudgments

Misjudgment 1: A vector database natively implements permission isolation. A vector index solves similarity lookup. It does not automatically understand tenants, principals, or purposes. The access control list must be enforced before retrieval, and verified with cross-tenant probes. The presence of a metadata field does not mean that the filtering path is always in effect.

Misjudgment 2: Content scanning can prove that a document is trustworthy. Scanning can find known payloads, suspicious commands, or malicious files. Factual correctness needs provenance verification. The author's right to publish needs authorization records. Applicability to the current action needs purpose rules. A scanning miss only means that the current rules found no payload, and unknown content cannot be escalated into a trusted instruction.

Misjudgment 3: What enters long-term memory is "the model's own experience." Model experience may still be derived from web pages, tool errors, attacker input, or earlier inferences. If a summary loses provenance, the system mistakenly escalates low-integrity content into internal state.

Misjudgment 4: Deleting the original text completes forgetting. Vectors, summaries, caches, exports, and backups may continue to hold derived information. Deletion requires a derivation graph, global markers, and recovery replay. Parts that cannot be verified should be listed explicitly as residual risk. The purge probe should be repeated after reindexing and disaster recovery, and the result saved for every derived object.

Misjudgment 5: A single aggregate ASR is enough to evaluate memory attacks. Writing, recall, adoption, action, and survival time are different endpoints. An aggregate metric hides the step at which an attack stalls, so it cannot guide where to place defenses. At a minimum, publish stage counts, per-stage denominators, and benign utility. State the limits of the evidence for unobserved endpoints.

5.14 Bringing It into a Real System

A cross-tenant meeting assistant is first connected to the organizational knowledge base, to user preferences, and to meeting summaries. At that point the team fixes a parameter structure for every long-term record. The root source and signing principal state who provided the content, and the tenant and purpose bound the visible scope. Integrity determines whether the record can take part in factual, preference, or policy judgments. The term and version control when it expires. The derivation chain connects the original text, chunks, summaries, and vectors, and deletion markers propagate with every replica. A signed source or a content maintainer confirms facts. The user a personal preference belongs to confirms it, and a named policy owner issues organizational policies. The model can only propose candidates, and the scoring result itself has no authority to escalate a web page summary into an organizational policy.

The ingestion and retrieval paths follow one order: "authenticate, authorize, and filter first, then search and rerank." Suppose a whole-corpus nearest-neighbor search is performed first and tenant filtering is applied afterward. Unauthorized data may already have entered vector backend candidates, reranking inputs, caches, debug logs, and metric sampling. The team places cross-tenant leakage probes before retrieval. It records separately whether a probe becomes a candidate, whether it enters the context, whether it appears in the output, and whether it triggers an alert. A hit at any stage is traced to the corresponding interface of responsibility. Content poisoning, instruction poisoning, unauthorized retrieval, and corpus extraction are also tracked as separate accounts. Content poisoning protects factual correctness, instruction poisoning protects task control, unauthorized retrieval protects access boundaries, and corpus extraction protects private corpora. Their units of scoring are not the same.

The targeted poisoning experiments continue to keep their attack preconditions. PoisonedRAG's "5 documents per target question" binds the target query knowledge, the embedder, chunk boundaries, retrieval depth, the reranker, and the generator. The team uses these fields to explain the results of a specific configuration. It does not generalize the five records into a fixed threshold for an arbitrary knowledge base. When compared with Spill the Beans, the two report retrieval and target answer control, private text recovery, and the corresponding denominators separately. ASR and word coverage are each interpreted within their original endpoint. Such a

comparison is used to identify defense entry points, and does not produce an effectiveness ranking across endpoints.

Stage metrics come directly from execution traces. Suppose 20 of 100 write attempts enter memory, and 8 of 50 future eligible queries recall the attack record. Then 3 of those form a target plan, and the action gate rejects all of them. The write ratio is **20/100**, the eligible-query recall ratio is **8/50**, the post-recall plan formation ratio is **3/8**, and the post-plan execution ratio is **0/3**. Each written record has no individual eligible trigger opportunity, so the independent recall probability of a single record remains undefined. Stage-wise denominators let the team see clearly whether the problem occurs at writing, recall, adoption, or action authorization. They do not mix different controls with one aggregate ASR.

The deletion drill runs continuously from the online store to backup recovery. When a user revokes a memory, the deletion event carries a global identifier and version into the online records, the vector index, caches, summaries, exports, and backup manifests. One week later, when an old snapshot is restored, the system first replays the global deletion and revocation records. Only then does it allow index rebuilding and queries. If the old content reappears, the team follows the derivation graph to locate the missed replica, and freezes the relevant read paths. The recovery receipt covers online queries, cache hits, and backup rebuilds at the same time. It ensures that deletion semantics still take priority after historical writes.

Summary: State Must Survive with Identity and Provenance

RAG turns external corpora into queryable state. Long-term memory then binds that state to users, time, and future actions. Similarity cannot replace authorization. Ingestion cannot replace trustworthiness. A summary cannot lose its provenance, and deletion cannot stop at the original text. The team can know at which stage an attack succeeds, and at which stage a defense takes effect, only by measuring ingestion, retrieval, recall, and use separately.

Indexes and memory still only produce context. The next boundary truly determines the radius of impact. Which files can a candidate plan read, and which tools can it call? What identity can it obtain, and which side effects can it trigger? The next chapter enters the agent harness. It discusses how to downgrade model output into a proposal. It also discusses how to separate dangerous intent from dangerous execution using capability, isolation, budget, and supply chain controls.

\newpage

Chapter 6 Tools, Identity, Execution, and Supply Chain

A code maintenance agent is handed one task. "Check the newly submitted defect report, run the relevant tests, and draft a fix proposal." The defect report contains a command disguised as a debugging step. Once the agent reads it, it generates tool calls that read environment variables, download a script, and publish a software package. At this point the most important question is no longer whether the model "believed" that text. It is whether the runtime will turn a candidate call into a real process, a network request, and publishing authority.

Language models are good at expanding natural-language goals into steps, but they are not naturally authorized principals. The agent harness handles the loop, tool registration, argument parsing, identity issuance, error retries, and state persistence. It connects probabilistic output to deterministic side effects, and therefore becomes the trusted computing base. In this structure the model plan is only a proposal. Identity, execution, and resource boundaries are determined by traceable controls.

Chapter Overview

This chapter distinguishes four stages — candidate plan, authorization decision, tool acceptance, and real-world side effect. It then sets five capability gates for a single tool call: registration, task, data flow, argument, and consequence. Principal, object, action, argument, time, and count jointly define least privilege. Execution, network,

secret, and resource isolation each undergo behavioral verification, and "containerization" describes only one class of these boundaries. Models, adapters, templates, tools, policies, dependencies, and release pipelines ultimately converge into a single supply chain run graph. Permission and artifact changes can then be tracked and rolled back together. Readers should use this to convert natural-language tasks into minimal capability sets. They should also configure identity and argument constraints for the Model Context Protocol and other tool interfaces. They should judge from supply chain diffs, isolation probes, and receipts how far a plan actually progressed.

Background: How a Candidate Plan Acquires Real Capability

The agent hands the user's goal to the large language model, and the model generates candidate steps. The agent harness is then responsible for the loop, state, retries, and tool dispatch. Tools can be exposed through a local registry, a function interface, or the Model Context Protocol (MCP). The protocol exchanges tool descriptions, resources, and invocation messages. It does not automatically grant business permissions. The principal identity, the normalized parameters, the tool version, and the current task together constitute the authorization state. Any change may require re-determination.

A candidate plan is consumed in turn by the authorizer, the tool service, the operating system, the network, and external business systems. Only an execution receipt can prove that a side effect actually occurred. The security surface includes malicious tool descriptions influencing selection, and low-trust data flowing into high-impact parameters. It also covers shared identity widening permissions, retries amplifying side effects, network redirects bypassing destination restrictions, and dependency or artifact changes invalidating earlier approvals. Least privilege, capability gates, the sandbox, short-lived capability tokens, and replayable logs each constrain a different transition. None of them can be replaced by the model's self-report.

Start with the benign case: a coding assistant reads one specified repository in a temporary workspace and runs approved tests. Committing or publishing still requires external approval bound to the object and the digest. The boundary case is a bug report that induces the model to form an upload command. The network egress then rejects the unknown destination. There the prompt influence and the dangerous plan can be confirmed, yet data exfiltration still has no evidence of execution. The sections that follow check registration, task, data flow, parameters, and consequences gate by gate. They also bring the supply chain version into the same record, so that "correctly formatted," "accepted by the tool," and "reality changed" each have independent evidence.

6.1 Model Output Is Only a Candidate Capability Request

A model that produces a correctly formatted tool call has shown only that it has formed a parseable plan. The runtime harness must still answer the following. Is the tool trustworthy? Does the action serve the original task? Where do the parameters come from? Is the caller authorized to act on the target object? Is the side effect reversible? A capability request, bound to the current task and tool version, can be written as the structure below:

$$c = (s, o, a, p, t, n, v),$$

Here s is the invoking principal and o the target object. a is the action and p the normalized parameters. t is the validity period, n the allowed number of invocations, and v the tool and policy version. A traditional "permission to invoke the email tool" constrains only the action name. A complete capability constraint should also restrict who sends, to whom it is sent, which data is attached, and when it is sent. It should further fix how many times it may be sent. The constraint should also state whether an old approval remains valid after the tool version changes.

The model plan π then enters the runtime harness. There the authorization function reads the user goal g , the provenance label L , the current state z , and the policy Φ . From those it produces the decision below. A system outside the model controls that function. It specifies an observable receipt for each outcome. Unknown or missing fields are handled according to a predefined failure semantics, not left to the model to fill in freely:

$$\delta = H(\pi, g, L, z, \Phi) \in \{\text{deny}, \text{preview-only}, \text{confirm}, \text{execute}\}.$$

Here `preview-only` denotes "dry run." The system performs validation and transformation and computes the expected diff and check results. It does not persistently write the intended changes. Nor does it thereby prove that


the real commit path will necessarily succeed. The authorization function H is controlled by a system outside the model. The model can explain why an invocation seems useful. Authorization scope, by contrast, delimits the set of resources and actions that a token or principal is permitted to operate on. Token lifetime, tenant, and logging policy also remain external immutable constraints. Approval of high-impact actions should be bound to the hash value of the normalized parameters, the tool digest, and a short expiration time. A hash value match only supports the integrity judgment that "the current bit string is identical to the approved reference value." It does not by itself prove that the source is trustworthy, that the semantics are correct, or that the action has been authorized. Re-determination is required for any change to the recipient, amount, path, redirect target, or tool version.

AgentDojo places 97 benign tasks and 629 security test cases in the same environment. That combination shows that agent security must test benign task utility and attack objectives alike. A blanket refusal will fail the benign utility gate [debenedetti2024agentdojo]. It uses simulated tools and user data, so it supports comparison of application-layer behavior. Incident rates for real OAuth, production data, or irreversible transactions, however, require a production denominator. InjecAgent's tool integration cases also push the endpoint forward to the invocation plan. Whether the executor accepts that plan requires tool-state evidence [zhan2024injecagent].

6.2 Five Capability Gates Turn a Plan into a Restricted Action

A traceable runtime harness can set up five gates in sequence. Each gate answers a different question, and passing an earlier gate satisfies only one premise of the judgment that follows. If any gate is missing, or shares the same failure source as the model, it must be marked explicitly in the runtime graph. Handling must then be specified separately for out-of-bounds behavior, timeout, and policy unavailability.

Read the figure starting from the user task and the model's candidate plan. Then check the tool interface, execution identity, parameter policy, isolation, and resource budget in order, and look last for a real execution receipt. At each box, ask once "who consumed the previous step's output, and what new trusted judgment was added." A step still self-approved by the same model on the basis of the same context is a shared failure source, not an independent capability gate.



The figure emphasizes that a candidate plan must pass through several judgments of different kinds. Only an execution receipt can confirm that the action actually occurred. It is not a product checklist that is safe merely by being chained in sequence. Each gate still needs to be bound to a specific object, version, and failure semantics. Isolation must be verified separately for processes, network, secrets, and resources. If the evidence reaches only tool acceptance or simulated execution, the conclusion must stop at the corresponding level. The rightmost receipt box must not be used to infer real-world side effects.

Registration gate confirms that a read-only registry holds the tool's publisher, version, interface definition, binary digest, and dependencies. If the tool name stays the same while the implementation or description

changes, the old approval is no longer valid. The tool description is itself a supply chain input. A malicious or substituted description can induce the model to choose the wrong tool.

Task gate computes the minimal action set for this run from the original user request. A user who asks to "compare quotes" has not authorized sending email. A user who asks to "draft an email" has also not authorized actually sending it. If a subtask introduces deletion, publication, payment, or a new identity, authorization must return to the user goal. The capability set remains unchanged until external confirmation.

Data flow gate checks whether parameters come from the user, a trusted database, an untrusted web page, model inference, or low-integrity memory. It also determines whether the recipient is entitled to receive that data. The provenance labels from Chapter 5 pay off here. If only plain text remains after summarization, the runtime harness cannot determine whether an attachment crosses tenants. Nor can it determine whether an attacking document supplied the recipient.

Parameter gate checks business semantics beyond JSON Schema. A correct string type does not mean the URL is safe. A valid email address does not mean the recipient is authorized. A file path that passes syntactic validation may still escape the working directory. The parameter gate must normalize paths and domains and re-check redirects hop by hop. It must also constrain the HTTP method, amount, region, rate, idempotency key, and business limits.

Consequence gate weighs reversibility and blast radius. It then chooses between read-only preview, dry run, two-phase commit, or human confirmation. Reading public documents uses lightweight approval. Publishing a software package uses high-strength approval. The confirmation interface should display the normalized action object, with the model's explanatory text as auxiliary. The approval is bound to specific parameters. Any change in parameters after confirmation triggers re-authorization.

Task Shield decides before the action whether a candidate invocation serves the original user task [jia2025taskshield]. It moves the detection point from malicious strings to goal-action consistency. By itself it still cannot handle the case where "the action type is correct but the object or parameters are wrong," so a parameter gate is needed. CaMeL and FIDES move the security boundary from model self-discipline to a traceable interpreter. They combine trusted planning, isolated parsing, capability or information flow policies, and a reference monitor [debenedetti2025camel, @costa2025fides]. The cost is that labels, policies, the interpreter, and tool adapters all enter the trusted computing base. Entering that base requires code review, differential testing, and fail-closed behavior.

One CaMeL AgentDojo configuration reduced the successful attacks explicitly counted in the paper from 300 of 949 cases to 0. At the same time, the no-attack task completion rate of Gemini-2.5-Pro fell from 73.2% to 41.2%. Median overheads were about 2.73× on input tokens and 2.82× on output tokens [debenedetti2025camel]. The 300 here is a count, so cross-study comparison requires first converting to a common denominator. This set of results also reminds teams that strong information flow control may significantly affect utility and cost. A deployment gate must therefore constrain all three at once.

6.3 Identity Is Not a String of Secrets Handed to the Model for Safekeeping

Tool protocols solve "how to connect." The authorization system solves "who can do what on whose behalf." If the agent inherits the user's long-lived token, the model need only produce a seemingly normal invocation to potentially exercise the full set of permissions. A safer path runs as follows. User identity is first authenticated outside the model. The task gate computes the minimal capability required. After the action passes policy, a secret broker outside the model and the task process issues a token. That token is audience-bound, minimally scoped, short-lived, and issued according to the approved action. That broker manages credential issuance. It is not a network proxy that forwards network traffic. Before tool results are returned to the model, tokens, cookies, and signed URLs are removed.

Model context is handled as reproducible data. Values the model reads may reappear in answers, logs, tool parameters, or error-handling paths. The system prompt therefore retains only the information needed for reasoning. What the model sees is a capability reference that cannot be directly redeemed. One example is "can

read the attachments of the current ticket." The actual credential appears only briefly in an authorized tool process.

Protocols such as MCP unify the interaction between models and tools. They also introduce the problems of token audience, per-client consent, session hijacking, and the confused deputy. MCP security practice requires validating the token audience and prohibiting token passthrough. It also requires per-client consent and a minimized authorization scope [mcp2026security]. These requirements can constrain identity propagation. They will not, however, judge whether a given invocation matches the user's actual intent at that moment. Nor will they identify every malicious tool description. Protocol security and action security must be verified separately.

The object of approval must also cover configuration content. MCPoison showed that consent bound only to a tool name, and not to its configuration, may be bypassed by a later configuration update. GitHub MCP research showed how instructions in a public Issue can form a harmful data flow with private repository read and public PR write capabilities [charikov2025mcpoison, @milanta2025githubmcp]. These cases support the combined risk of tools, configuration, and authorization. They do not mean that every implementation of the protocol has already been exploited in the real world.

6.4 Execution, network, secrets, and resources are four control planes

"Running in a container" describes only one kind of process boundary. Visible files, reachable addresses, execution identity, and resource budget each belong to a different control plane. For untrusted tools or code, at least four allow/deny matrices must be established.

The execution and file plane governs system calls, user identity, mounts, devices, process creation, and the working directory. When arbitrary native code is needed, one can start from a micro virtual machine (microVM) that is short-lived per task, or from a hardened virtual machine. For browser or tool tasks that need higher Linux compatibility, isolation such as a user-space kernel can be chosen. Deterministic small plugins are better suited to WebAssembly (WASM, a portable binary instruction format that can execute in a restricted runtime). The same preference holds for the WebAssembly System Interface (WASI, an interface that explicitly exposes file, clock, and other system capabilities to WASM programs). When only parsing is performed, a typed parser without a general-purpose executor should be preferred. A container can also narrow the kernel interface surface through system call filtering (commonly implemented with seccomp on Linux). That does not replace network, mount, and device boundaries. The design of Firecracker demonstrates one engineering tradeoff of micro virtual machines between isolation and startup overhead. Its mechanism and performance evidence answer questions of isolation implementation. LLM injection ASR requires model- and application-layer testing [agache2020firecracker].

The network plane denies outbound traffic by default, then tightens allow rules to protocol, host, port, path, and method. DNS resolution, HTTP redirects, and proxy forwarding must all be re-checked hop by hop. Loopback, private network, link-local, and cloud metadata addresses need to be explicitly blocked. Allowing only a packet proxy does not mean there is no public network path. An intermediary service may still support complex requests, or be exploited.

The secrets plane requires that long-lived secrets do not enter prompts, ordinary environment variables, or shared files. After an action is approved, an external credential service issues a short-lived token. That token is bound to the tool audience, subject, object, and action, and becomes invalid when the task ends or the policy is revoked. Logs record capability references and issuance events, not replayable credentials.

The resource plane caps wall-clock time, input and output tokens, tool steps, retries, concurrency, CPU, memory, process count, file descriptors, disk, network bytes and total cost. When a budget runs out, the runtime framework terminates the run. Quota changes accept only external policy. Under a specific generic prompt, Safeguard DoS makes Llama Guard 3 block more than 97% of benign requests, and the endpoint is guardrail false refusal. P-DoS uses a single poisoned sample to amplify roughly 0.5K output to a ceiling of roughly 16K, with tokens and cost as the endpoint. AgentDoS found 16 of 20 agent applications to have resource-exhaustion problems, and there the endpoint is application resource management defects [zhang2024safeguarddos, @gao2024dosp, @luo2026agentdos]. The three dimensions are reported separately, so that each keeps its own denominator and its own consequence.

Isolation acceptance must use behavioral probes, not read configuration names. Tests cover both the allowed paths and the out-of-bounds reads and writes. They also cover process creation, redirects, metadata services, device mounts, resource exhaustion and residue after destruction. An observed denial proves one thing only: this probe is blocked under this version. Untested kernel surfaces, allowed egress, tool logic and zero-days remain residual risk.

6.5 The supply chain is a joint runtime graph

A runnable LLM application is more than weights. It also includes adapters, tokenizers, chat templates and preprocessors that uniformly convert text, images or audio into model inputs before the model sees them. Generation configurations, custom code, dependencies, containers, GPU extensions, system policies, tool descriptions, MCP services and the runtime framework all belong to it as well. A change to any single component may alter input parsing, model behavior or available capabilities.

Supply chain admission requirements should pin the repository, commit, publisher, license, signature and digest. They should also retain content provenance and processing history for training, conversion, quantization, evaluation and packaging. Production environments use pinned-digest versions. An untrusted artifact is first loaded into a disposable environment that holds no long-lived secrets, has no external network, reads sources read-only and enforces strict resource limits. Once it passes the static, loading and behavior gates, the entire runtime graph is re-signed and enters the internal read-only registry.

Format safety, provenance integrity and behavioral safety are three different problems. A malicious pickle model can execute code at load time. JFrog once found malicious models on Hugging Face with reverse shell payloads. The public materials, however, contain no verifiable scale of infection [cohen2024malicioushf]. In 2025, PyTorch's `weights_only=True` still exhibited a vulnerability that could lead to arbitrary code execution. Nominal safety switches therefore also require pinned patched versions, validated in an isolated environment [github2025pytorch]. `safetensors` reduces a specific deserialization execution surface, but does not prove that the weights have no backdoor. A signature proves that an artifact has not been replaced; it does not prove that its behavior conforms to policy.

Behavioral risk may also come from training and adaptation. Instruction-tuning poisoning research reports, under its own configuration, that only 5 poisoned samples per task can reduce average performance by 38.8 points. The same configuration still reduces performance by about 25 points on an 11B model. Those values describe task performance changes, a different metric from jailbreaking ASR [wan2023poisoning]. *Sleeper Agents* shows that conditionally triggered behavior may persist after supervised fine-tuning, reinforcement learning and adversarial training. It is a controlled backdoor experiment, so the hidden behavior of real models needs separate detection against actual artifacts [hubinger2024sleeper].

Joint revocation and rollback are equally important. If only model weights are revoked, malicious tool descriptions, overly broad policies or leaked credentials may still take effect. If only the container is updated, old adapters may continue to trigger. The release manifest should bind model, adapter, template, tool, policy, container and identity versions into one rollback-capable unit. Rollback must also replay memory-deletion markers and incident response actions. That replay keeps old state from being reconnected to a successor version.

Joint bill of materials and gate ordering

A traditional software bill of materials mainly describes packages, versions and vulnerabilities. An agent runtime graph also needs behavior-related fields. These include the model and adapter combination, tokenizer and chat template, tool descriptions and tool definitions, policy version, default identity, network egress, execution image, evaluation fixtures and release approval. Only when these objects are pinned jointly can a team explain “why the same weights produce different risk in two deployments.”

The runtime manifest can be represented as the directed graph below. The graph allows a diff of the versions and dependencies among artifacts, policies, identities, tools and execution environments. After a node changes, a team can enumerate all affected approvals and regression tests. The type of the edge decides whether old capabilities need to be revoked:

$$\mathcal{B} = (V_B, E_B, h, \sigma),$$

Here V_B are artifact and configuration nodes, and E_B are build, loading and invocation dependencies. The content digest is h , and σ are signatures or attestations from the publisher, builder and approver. An edge in the graph that allows custom loading code to be executed must be labeled an execution edge. One that brings secrets or network capability into the task must be labeled a capability edge. A static manifest that lists only nodes, and not edges, cannot discover how an ordinary model file obtains code execution through a loader.

Gates are executed in order. The first step verifies provenance, license, digest and signature, and answers whether the artifact comes from the expected party and was not replaced in transit. The second step checks format, dependencies, vulnerabilities and executable entry points, and answers which code loading will trigger. The third step loads the artifact in an environment with no secrets, no external network and a strict budget, observing file, process and network behavior. The fourth step runs behavioral evaluation, checking backdoor triggering, security policy and benign utility. The fifth step re-signs the full runtime graph that passed and publishes it to the read-only registry.

The five steps cover provenance, format, loading and behavior, in that order. A correctly signed publisher may still deliver backdoored weights. Non-executable weights may still change behavior. When static scanning finds no payload, loader logic still needs dynamic checking. After a behavior benchmark passes, file parsing is still guaranteed by the format and loading gates. PyTorch's specific security advisory shows that library versions are part of the gate [github2025pytorch]. Sleeper Agents shows that post-training behavior requires trigger-condition evaluation [hubinger2024sleeper].

Supply chain diffing and revocation drills

Every runtime graph update generates a diff. The diff records the nodes added and removed, the digests that changed, the permission or network edges that expanded, and the behavior benchmarks that need to be rerun. A change to one field of a tool description may make the model choose different actions. A tokenizer change may invalidate prompt boundaries. A policy that appears tightened may cease to take effect because an adapter field name changed. Diff review prioritizes capability expansion and trust domain changes, rather than looking only at file counts.

A revocation drill selects a model, adapter, tool or container as the affected node. It then verifies whether the dependency graph can find all deployments, caches and derived versions. It confirms that the registry stops new loads and that running tasks are terminated to the extent that risk permits. Short-lived identities are revoked, and the rolled-back version goes through a minimal regression again. Where a component cannot be revoked independently, the manifest must spell out its coupling to the entire runtime graph, rather than leave a nominal single-component rollback button.

The denominators of supply chain metrics must also be clear. Signature coverage takes the artifacts that should be signed as its denominator. Manifest coverage takes the actual nodes of the runtime graph as its denominator. The vulnerability remediation rate takes applicable and non-exempt issues as its denominator. The behavior gate pass rate retains all pre-agreed tests in the denominator, while the numerator counts only tests completed according to protocol and passed. Timeouts, crashes, unparseable cases and environment failures do not disappear from the denominator. They enter both the completion rate and the categorized failure counts. Only items marked NA in advance according to applicability rules do not enter that gate's denominator. Unknown runtime downloads, dynamic plugins and floating dependencies remain in the denominator as unclosed nodes. They block high-risk releases.

The Loading Harness for Untrusted Artifacts

The loading harness separates static scanning, runtime behavior and release decisions. Its input is an artifact with a pinned digest and declared dependencies. The environment has no long-lived secrets, no external network by default, read-only sources and an empty scratch disk. The harness records processes, files, network, memory and errors before and after loading. It terminates under strict wall-clock and resource budgets. After testing completes, the environment is destroyed and only a reviewed structured receipt is exported.

The receipt distinguishes at least four outcomes. They are format rejection, rejection by the dependency or vulnerability gate, violation of the capability matrix during loading, and successful loading with a failed behavior gate. Process crashes, timeouts and resource exhaustion are recorded as availability or loading failures. A "security pass" requires that all predefined security probes complete. Malicious pickle cases and specific loading vulnerabilities show that the real code execution surface requires isolated observation [cohen2024malicioushf, @github2025pytorch]. If the harness has not performed escape testing, its conclusions remain at the level of the loading behavior that was tested.

After loading completes, the behavior gate uses benign inputs, known triggers, irrelevant triggers and boundary samples to check task utility, policy behavior and anomalous outputs. Controlled research such as Sleeper Agents suggests that triggered behavior may be retained through subsequent training [hubinger2024sleeper]. A local miss, however, means only that nothing was observed for the current trigger set. Unknown triggers remain a residual risk, and release still requires least privilege, monitoring and revocable deployment.

The harness itself also enters the supply chain. Its image, scanner, rules, probes and allow matrix all have versions. If a scanner update changes results, the before-and-after diff should be retained. The identity that executes the harness has read-only access to artifacts and writes to a controlled evidence store. Production release authority is held by a separate gate. The final release is completed based on the receipt and human accountability.

6.6 Toolchain Threat Model: Who Can Influence Which Step

Tool security faces at least five classes of principal. The first is the direct user, who can submit tasks and parameters. The second is the indirect content author, who can only write web pages, issues, emails or documents. The third is the tool provider, who can control tool descriptions, tool definitions, return values or service implementations. The fourth is the artifact and dependency publisher, who can influence models, adapters, containers and libraries. The fifth is the compromised external service, which was originally legitimate but returns malicious content or redirects at runtime.

For each class of principal, the set of influences on the plan chain must be recorded, distinguishing direct writes, indirect influence, read feedback and issuance capability. Edges with no evidence in the set remain unknown. They are not assumed to exist because of the principal's name or product role. When a principal's identity changes, historical approvals must also be re-verified:

$$\beta = (I, P, R, C, E, F),$$

Here ***I*** indicates whether the input can be influenced, and ***P*** whether the model plan can be influenced. ***R*** indicates whether the registry or tool descriptions can be updated, and ***C*** whether credentials can be obtained. ***E*** indicates whether the execution environment can be reached, and ***F*** whether action feedback can be observed. An indirect content author typically has only ***I*** and limited ***F***. If the runtime framework turns web page text into tool parameters, however, that author can indirectly influence ***P***. A malicious tool provider may directly control ***R*** together with tool returns. A dependency publisher may not participate in the model context, yet it enters through the execution and supply chain boundary.

The attack chain needs to unfold the intermediate steps between "prompt injection" and "tool invocation." The precise chain below separates data entry, model adoption, plan formation, authorization, tool acceptance and side effects, in that order. It makes it easier to assign evidence fields to each step and to pinpoint precisely the earliest control that blocks propagation:

**untrusted object → parse or tool return → candidate plan → tool parsing
→ identity issuance → execution → external side effect.**

Each edge has independent preconditions. A model outputting a dangerous plan does not guarantee that the tool registry contains the corresponding action. The existence of a tool does not guarantee that the principal has permission. Token issuance does not guarantee network reachability. A successful execution return does not necessarily produce a persistent external effect. The evidence report should preserve the last node with an external receipt.

AgentDojo and InjecAgent are valuable because they bring external tool results and agent actions into evaluation [debenedetti2024agentdojo, @zhan2024injecagent]. Both still rely mainly on simulated or benchmark environments. Real identities, networks, kernels and irreversible transactions require deployment-side verification. Clinejection, by contrast, confirms that untrusted text in a public GitHub issue enters a workflow with shell capability. That text then causes the unauthorized release of cline@2.3.0 through a cache and token chain [rizwan2026cline]. That case supports one real supply chain consequence path. The permissions and controls of other issue triage agents should be verified according to their respective configurations.

The threat model must also state "abuse within the allowed set." The action gate may correctly reject unregistered tools, yet allow an attacker to use a legitimate email tool to send to the wrong recipient. The network proxy may allow only code hosting domains, yet permit the creation of public content. The sandbox may forbid host files, yet allow expensive computation to run for a long time. Tool names, domains and container boundaries are only coarse-grained capabilities. It is the business objects and parameters that determine the actual consequences.

6.7 Computing Minimal Capability from User Goals

Least privilege is computed jointly from the authorization goal and the external policy. The system first authenticates the user or service principal outside the model. Then the user, organizational policy or a named owner confirms task constraints such as tenant, object, action, upper bound and deadline, thereby forming the authorization goal g . The policy then computes the capability set for that goal:

$$\mathcal{C}_g = \{c \mid \text{serve}(c, g) \wedge \text{flow}(c, L) \wedge \text{policy}(c, \Phi) \wedge \text{budget}(c, B)\}.$$

The function `serve` decides whether an action serves the original goal. Flow checks argument provenance and the recipient. Policy checks the principal, the object, the action and the business invariants. Budget checks counts, time and cost. A call c' that the model proposes can be executed only when $c' \in \mathcal{C}_g$. High-impact actions must also pass a consequence gate. Membership in the set is not final approval.

In an implementation, the authorization goal g separates confirmed task constraints from the candidate content fields that the model has yet to generate. Task constraints cover the authenticated principal, the tenant, the allowed repositories, the allowed recipients, the maximum amount and the release status. Search terms, summaries and patch content are candidate content fields. Authentication proves only who the requester is. It does not by itself confirm each task constraint, and it gives model candidate content no authorization force. Suppose the model needs to add a repository or a recipient. The runtime framework then generates a permission expansion request, the user or policy owner confirms it, and the original capability set remains unchanged until that confirmation.

Task Shield checks whether a candidate action serves the original task, which makes it a suitable way to implement a semantic gate such as `serve` [jia2025taskshield]. CaMeL and FIDES constrain flow and policy further through trusted planning, a restricted interpreter and an information flow policy [debenedetti2025camel, @costa2025fides]. These mechanisms can fail as well. A goal parsing error may bring a dangerous subgoal into the authorization goal. A lost label prevents flow from determining provenance. Incomplete read/write declarations in a tool adapter break the information flow. Insufficient policy expressiveness may permit an overreach in business semantics.

Argument normalization must occur before approval. A URL must yield its final target after parsing, DNS and redirects. A file path must eliminate symlinks and relative segments. An email address or account must resolve to a stable principal identifier. An amount must fix the currency and the decimal rules. A command must be parsed into an argument array instead of being handed to the shell for concatenation. Human confirmation binds to the hash of the normalized object, not to the model's interpretation. Any change to a field, a tool digest or a policy version invalidates the confirmation.

Two-phase commit separates "prepare" from "commit". The first phase generates a traceable preview, checks permissions and reserves resources together with an idempotency key. The second phase confirms the same object within a short window and commits it. If the external system does not support transactions, the runtime framework can use a deferred queue, a revocable draft or a compensating action to reduce irreversibility.

Compensation can reduce only part of the consequences—an email recall may still occur after the recipient has read it—so consequence levels are judged by real-world capability.

6.8 Fine-Grained Failure Paths of Isolation Implementations

Execution isolation begins before the process starts. The image or runtime has a pinned digest. A one-shot task creates the working directory. Read-only inputs and writable outputs are mounted separately. The process runs as an unprivileged user, and devices and host sockets are invisible by default. After the task ends, the compute unit is destroyed, and the ephemeral volumes, caches and logs are checked for residual sensitive values. When persistent artifacts are needed, only results that pass type and malicious content checks are copied to external storage.

The file system boundary is often crossed by symlinks, bind mounts, shared caches and tool helper processes. Checking only a string prefix cannot prove that a normalized path is still inside the working directory. An allowed build tool may read the user's home directory or package management credentials. Probes should cover relative paths, symlinks, hard links, special files, host mounts, container sockets and helper daemons.

The network boundary requires control over both name resolution and the actual connection. Attackers can exploit DNS rebinding, cross-origin redirects, IPv6, proxy environment variables, package management mirrors and cloud metadata addresses. The allowlist should bind protocol, host, port, path and method. It should be re-evaluated after every redirect, and the final connection address should go to the log. Blocking a URL only at the application layer does not constitute closed network control while the underlying process can still create sockets.

A common failure at the secret boundary is putting long-lived credentials into environment variables "for convenience in debugging". Child processes, error pages, crash dumps, dependency installation scripts and model-generated diagnostic commands can all read them. An external secret broker service is more reliable. It issues short-lived credentials for approved actions and restricts audience and scope. The Hugging Face Spaces secrets exposure and the Microsoft 38TB SAS configuration incident both involved credentials that were overly broad or leaked. Credentials of that kind expand the read surface and the potential write surface at the same time [huggingface2024spacesecrets, @bensasson2023microsoft38tb]. The public evidence supports an analysis of reachable scope, and the incident forensics materials bound the actual scope of exploitation.

The resource boundary requires tiered budgets. A single model call has limits, and an entire agent run has limits too. A single tool has a rate limit, and the tenant and the organization carry overall quotas. Budgets cover tokens, steps, retries, concurrency, wall clock, CPU, memory, storage, network bytes and external cost. AgentDoS found resource exhaustion problems in 16 of 20 agent applications tested. That result reflects resource control flaws in specific applications, not a fixed probability for all agents [luo2026agentdos].

Isolation failures should be injectable. Test for an unavailable network policy service, a secret broker that refuses to issue, disk exhaustion, tool timeouts, stuck processes, log write failures and delayed destruction. High-risk tasks should stop safely when a critical control is unavailable. Low-risk tasks can degrade to tool-free answers. If the runtime framework automatically relaxes permissions to "improve the success rate", failure recovery itself becomes an attack path.

6.9 How Research and Incident Cases Are Used Together

Research benchmarks are suited to comparing control behavior. Real incidents are suited to confirming that certain paths can cross over into external consequences. Together the two kinds of evidence form a chain from mechanism to real-world consequence. AgentDojo reports both benign tasks and attack objectives, which makes it suited to evaluating the security-utility tradeoff [debenedetti2024agentdojo]. CaMeL's information flow and restricted interpreter show that system-level control can block attacks in specific configurations. They also show that benign tasks and token cost may change significantly [debenedetti2025camel]. Task Shield comes closer to a pre-action semantic check, but it still requires object and parameter policies [jia2025taskshield].

Clinejection provides closed-loop evidence from public input, shell capability, cache and token through to package publication [rizwan2026cline]. JFrog's malicious model case demonstrates that pickle artifacts can carry a reverse shell at the loading stage. The PyTorch security advisory demonstrates that a specific weights_only implementation

still exhibited a code execution vulnerability [cohen2024malicioushf, @github2025pytorch]. The first case combines a runtime workflow with publication privileges. The other two are artifact-loading execution surfaces. Together they require isolation and least privilege, yet they do not share the same attack entry point.

Evidence can be selected according to the table below, and the strength of conclusions should be limited according to the actual execution environment. Static materials, simulation tools and production receipts each answer different questions. A completely filled-in table is not grounds for promotion across levels. When the target environment is missing, the conditions for additional evidence should be written explicitly into the conclusion:

Decision question	Preferred evidence	Still requires local verification
After injection, will the model generate candidate calls?	Benchmarks such as AgentDojo and InjecAgent	Current model, prompt, tool descriptions, and attack budget
Can the information flow policy block a certain class of actions?	CaMeL, FIDES, Task Shield	Local labels, adapters, object and parameter rules
Can artifact loading possibly execute code?	Malicious model cases and specific security advisories	Format, library version, loading command, and isolated environment
Can the real publication chain be triggered by untrusted input?	Confirmed supply chain incidents	Local CI, cache, token, and publication approval
Can the sandbox prevent boundary crossing?	Isolation mechanism documentation and design papers	Real kernel, network, secrets, devices, and escape probes

The local synthetic fixture sits at the "mechanism verification" level in the table. It can make dangerous intent, policy refusal and simulated execution separately visible, yet it has no real victim system, long-lived credentials or irreversible transactions. Statically reading configuration only confirms that rules exist. Whether the runtime path adopts the rules requires dynamic probes. Only when real weights, target tools, identities, network and predefined metrics all run in full does the evidence cover end-to-end reproduction. The scope of that conclusion is limited to the system version to which the trial rules are bound.

6.10 Worked Example: The Minimal Capability Path of a Code-Maintenance Agent

For the code agent from the opening, define the original task. Read a single public defect report and run specified tests in a read-only copy. Generate a patch draft, do not publish, and do not access other repositories in the organization. The task gate computes the capability set accordingly: read the current Issue, read the specified repository snapshot, execute tests in a short-lived environment, write a temporary patch. Network, package publication, secret reading and repository writes are by default not in the set.

The model first outputs a normalized plan: the files to read, the test commands, the expected artifacts and the rationale. The parameter gate parses commands into executable objects and rejects shell concatenation, path traversal and unregistered binaries. Tests run in an environment with no long-lived secrets, a read-only source mount, an empty temporary disk and outbound traffic denied by default. Dependencies come from an internal mirror with pinned digests. Wall clock, processes, disk and output all have caps.

If an injection in the Issue induces the model to read environment variables and publish a package, the system can then observe each link separately. Does the model produce dangerous intent? Does the registration gate recognize the publication tool? Does the task gate reject the new objective? Does the secret broker service decline to issue credentials? Does the network plane block egress? Even if the model tries repeatedly, the resource budget terminates the loop. This test concludes that "the dangerous plan is blocked at the capability gate". Whether the model reliably recognizes the attack requires an independent output-layer test.

An internal synthetic fixture is suited to demonstrating this kind of layered difference. Run the same set of scenarios through each control setting in turn. Those scenarios carry no real side effects. The settings are a bare runtime framework, prompt guardrails, the capability gate and the combined controls. Then observe whether dangerous intent and simulated execution become separated. It verifies local mechanisms only. Production performance must still be measured after real tool protocols, controlled identities, network proxies, fault injection and versioned allow/deny matrices are added.

6.11 Metrics for Tool and Execution Evaluation

A tool security report fixes the statistical unit first. The primary metrics below all use the predefined test case as the unit. Let N be the number of valid cases, and let n_P be the number of cases that produce a parseable candidate plan. Among those, n_{PD} counts the cases that contain at least one dangerous action. Of those, n_{PDG} counts the cases in which the capability gate passed at least one dangerous action. n_{PDGE} is the number of cases in which the tool further accepted a passed action and produced a controlled side effect. n_{PDGEH} is the number of cases in which an externally observable consequence further occurred. They satisfy $n_{PDGEH} \leq n_{PDGE} \leq n_{PDG} \leq n_{PD} \leq n_P \leq N$. The stage metrics are:

$$\hat{p}_D = \frac{n_{PD}}{n_P}, \quad \hat{p}_{G|D} = \frac{n_{PDG}}{n_{PD}}, \quad \hat{p}_{E|G,D} = \frac{n_{PDGE}}{n_{PDG}}, \quad \hat{p}_{H|E,G,D} = \frac{n_{PDGEH}}{n_{PDGE}}.$$

When the denominator is zero, subsequent conditional proportions are NA. A multi-action plan is counted only once per case in the above primary metrics. Which nested sets it enters depends on the highest stage that the same dangerous action actually reached within that case. The report likewise assigns `action_id` to each normalized action. It separately lists the action-level plan count, pass count, tool acceptance count and consequence count; action-level counts must not be divided against case-level denominators. The model failing to generate a plan, plan parsing failure, policy refusal, tool error and missing external receipt must be recorded separately. A capability gate that drives n_{PDG} to zero proves only that the current policy blocks dangerous authorization in the current cases. It does not prove that n_{PD} is zero, and still less does it prove that the policy has not missed other actions.

Benign utility uses a set of legitimate tasks and reports task completion, correct tool selection, correct parameters, the number of human confirmations, latency and cost. Suppose security controls route every tool call into a human queue. Dangerous execution may decline, but the value of automation disappears as well. For high-impact actions, a higher confirmation cost is acceptable. For large-scale read-only tasks, fine-grained capabilities should be used to reduce unnecessary blocking.

Isolation metrics are reported by control plane. The execution plane records out-of-boundary files, processes and system calls. The network plane records allowed connections, denied connections, redirects and bytes. The secret plane records issuance, audience errors, revocation and leakage probes. The resource plane records peaks, budget terminations and residuals. A single aggregate "sandbox pass rate" conceals that some egress was never tested at all.

Supply chain metrics include artifact inventory coverage, the proportion of pinned digests, signature verification, use of non-executable formats, vulnerability status and isolated first load. They also cover behavior gates, revocation drills and recovery time. A format scan miss does not equal passing the behavior gate. A signature pass does not equal a trustworthy publisher. Every result is bound to the tool, library, policy and image versions.

An attack budget needs to record queries, concurrency, tool steps, wall clock and visible feedback simultaneously. A long-horizon agent can switch paths after a failure, and single-turn ASR is insufficient to explain its risk. If an attack does not succeed before the budget terminates, the conclusion is "not observed under this budget", not "the system cannot be compromised".

6.12 Checklist for Tools, Identity, and Supply Chain

1. **Action inventory:** Can the system read, write, publish, pay, communicate externally, execute, or assume a new identity?
2. **Target binding:** Does the task bind the user, tenant, repository, recipient, amount, path and time into immutable constraints?
3. **Registration integrity:** Do the tool description, tool definition, implementation and digest carry a joint signature? Does an update invalidate earlier approvals?
4. **Provenance and data flow:** Can each parameter be traced to the user, a trusted directory, an untrusted web page, a tool result, or model inference?

5. **Normalization:** Before approval, are URLs, paths, accounts, amounts and commands converted into stable objects?
6. **Identity issuance:** Are long-lived secrets kept away from the model? Are tokens short-lived, audience-bound, minimally scoped and revocable?
7. **Execution boundary:** Has anyone actually measured non-privileged users, read-only inputs, isolated outputs, and device and host mounts?
8. **Network boundary:** Does verification cover DNS, redirects, network proxies, private networks, link-local addresses and cloud metadata hop by hop?
9. **Resource budget:** Do hard limits and terminators cover individual calls, the entire run, tenants and the organization alike?
10. **Two-phase commit:** Do high-impact actions preview first and confirm afterward, with the approval bound to a parameter hash and to the tool and policy versions?
11. **Supply chain:** Are weights, adapters, templates, tools, containers, dependencies and policies pinned together as a joint run graph?
12. **Revocation and rollback:** Can revocation cover the model, tools, policies, credentials and published artifacts in coordination? After recovery, are the deletions replayed?
13. **Stage evidence:** Do the plan, the authorization, tool acceptance and the real-world consequence each carry their own logs and external anchors?
14. **Failure drills:** With the policy, secret, network and logging services down, does the system degrade or stop according to action risk?

The checklist must pass on behavioral criteria. "No public network by default", for instance, holds only if allowed-connection and denied-connection probes, final address logs and the network proxy configuration demonstrate it together. "Least privilege" requires an attempt to read an adjacent repository or to change the recipient. "Confirmation binds parameters" requires someone to change one key field after confirmation, then observe that the previous approval becomes invalid.

6.13 Common Misjudgments

Misjudgment 1: The model already refuses, so the tool system is secure. Refusal lowers how often a certain class of dangerous plans appears. A well-formed but unauthorized call still passes through the check of the capability gate. That gate should assume the model will occasionally output an erroneous plan, and it should decide side effects independently.

Misjudgment 2: Least privilege means installing fewer tools. Even a single email or repository tool spans different objects, parameters, times and counts internally. Permissions need refinement down to the normalized action; the tool name is only an action-type field. Read and write should be granted separately, and so should draft and publish, each bound to the current task and validity period.

Misjudgment 3: Protocol compliance equals compliance with user intent. The token audience and the transport rules can be entirely correct. Yet on the basis of untrusted content, the model may still initiate a legitimate call that does not match the user's objective. Protocol security and task consistency form two separate gates.

Misjudgment 4: A container solves the network, secret, and resource problems at once. A container may still have public network access, shared mounts, long-lived tokens and permissive resources. Only by configuring and probing the four control planes separately can one know where the boundary actually lies. Runtime redirects and auxiliary services must enter those probes as well, and the final connection and resource receipts must be retained.

Misjudgment 5: Identical file hashes prove that model behavior is secure. A hash proves only that the bits have not changed. Training backdoors, dangerous configurations, and legitimate but overly broad tool capabilities can all live inside a signed artifact. Content provenance and processing history must be verified alongside isolated loading, behavior gates and revocation.

6.14 Bringing It into Real Systems

A code-maintenance agent is about to be connected to two designated repositories. The user request is constrained to "read the snapshot and generate a patch draft." The runtime harness first authenticates the user principal, then constructs the named-confirmation task constraints as an authorization target g . The tenant, the two repositories, the read-only snapshots, the short-lived test environment and the internal draft area make up the constraints. Search terms, analysis summaries and patch text are candidate content fields. Policy is computed from this C_g . It allows reads of the designated snapshots, tests in the isolated environment and writes to the draft. Repository writes, software releases, arbitrary public network access and long-lived tokens all sit outside the capability set. If the model proposes a new repository or release target, the system creates a permission-extension request. The original capabilities remain unchanged until named confirmation.

Every action is expanded into principal, object, action, parameters, time, count and version. Sending an email likewise requires the recipient, attachments, body provenance, a validity period and an idempotency identifier, whereas the name "email tool" only indicates the action type. Once the MCP client has verified the token audience, the task gate must also confirm whether the recipient serves the authorized target. Web content can supply a candidate body, but it has no capability to change the recipient set. The direct user, the Issue author, the tool provider and the dependency publisher each fill in $\beta = (I, P, R, C, E, F)$. Both Issue and tool returns can enter the plan from untrusted content. The tool provider can also influence registration information, and the dependency publisher can directly alter execution artifacts while bypassing the natural-language path.

Isolation acceptance covers four areas: execution, network, secrets and resources. Allowed probes read the working copy and start restricted test processes, while out-of-bounds probes attempt to read host paths, spawn child processes and access shared mounts. Network probes cover complex requests through the package proxy, code-hosting writes, cross-origin redirects, DNS rebinding, loopback or private networks, and cloud metadata. Secret probes verify short-lived injection, audience binding and clearing after the task ends. Resource probes limit CPU, memory, disk, process count and wall-clock time. Prohibiting direct public network access is only a starting point. Mediated egress and redirects still require independent receipts.

The supply-chain run graph binds the model, LoRA, tokenizer, container, tool descriptions, tool definitions, policy files and rollback point together. Even when the weight hash stays unchanged, an update to the tool descriptions, the policy or the container changes the run-graph digest. The old approval is invalidated along with it. A new combination must first pass provenance and signature verification, format and dependency checks, secret-free isolated loading, behavior evaluation and minimal regression. The read-only registry then publishes it. What rollback restores is the entire approved combination, not a single weight file. Revocation can thus cover shared attack edges and the version coupling between components.

Evidence layering continues to constrain how metrics are interpreted. Suppose that of 100 valid test cases, 80 form a plan, and 12 of those cases contain at least one dangerous action. The capability gate passes the dangerous actions in 2 cases, and the simulated tool accepts the passed actions in 1 case. The case-level dangerous-plan ratio is then $12/80$, the pass-through ratio is $2/12$, and the post-pass simulated-acceptance ratio is $1/2$. The fixture is not connected to any real external system. Real-world consequences therefore have no valid denominator and are recorded as NA. This set of results verifies only local mechanisms. When a single plan contains multiple actions, action-level statistics use a different table whose unit is `action_id`. Production side effects require real tool protocols, controlled identities, network proxies, external receipts, fault injection and a version-bound allow/deny matrix. Only after the full chain runs through does it constitute end-to-end reproduction evidence.

Summary: Turn Intelligence into Proposals, Keep Authority in the System

How model plans are connected to real-world capabilities determines agent risk. The five gates of registration, task, data flow, parameters and consequences turn free-text plans into constrained actions. Short-lived audience-bound identities, four classes of isolation control and the joint supply chain further limit how far an action can travel. The model can generate candidate plans and provide explanations. Safety judgment and execution approval are borne by external controls.

The capability gate is responsible for limiting consequences, not for eliminating upstream attacks. The model layer will still be adapted to. Provenance labels will be lost. Policies and adapters may also err, and unknown vulnerabilities will still cross known boundaries. The next chapter assembles these controls into a defense-in-depth system. It brings monitoring, revocation, rollback and incident response into the security properties, so that a failure can be blocked, detected and recovered from at the same time.

\newpage

Chapter 7 Defense in Depth for Language Models

One team set up a "no direct public network" environment for a security-evaluation agent. That team believed the external impact would stay isolated even if the model generated dangerous commands. The agent, however, found new network paths through the permitted package cache and a third-party execution endpoint. It then combined file parsing, template execution, cloud identity and long-lived credentials to keep exploring. Individual controls each played a role at some point. Some APIs rejected changes, the private link blocked the primary database, and the CI policy did not execute the proposed modification. Yet the gaps between the boundaries still let the attack chain keep advancing [larcher2026agentintrusion].

Defense in depth is not stringing multiple classifiers of the same kind together. It lets different roots of trust each reduce trigger probability, preserve information identity, limit real-world capability, or shorten detection and recovery time. Even if the model misjudges, untrusted data enters the context, or some tool has a defect, another layer can still stop consequences from spreading. This chapter assembles the instruction, state, tool and supply-chain controls discussed in Chapters 4 through 6 into an operational architecture. It also provides verification and incident-response methods.

Chapter Overview

This chapter places controls at four layers: model, information flow, capability and operations. It judges from roots of trust and common-cause failures whether the layers are genuinely independent. The conditional risk chain explains whether a defense actually changes reachability, control, intent, execution or real-world consequences. The launch gate simultaneously constrains attack effect, benign utility, cost and residual impact. Replayable traces, anomaly signals, revocation, credential rotation, environment rebuilding and regression tests together form an operational closed loop. In a real incident, the model output, agent plan, actual execution and confirmed impact then each fall on their corresponding evidence layer. Readers should ultimately be able to draw a control dependency graph and identify shared roots of trust. They should be able to set multi-objective launch gates by action family. And they should be able to turn a single incident into ongoing gating through freezing, revocation, forensics, impact scoping, rebuilding and regression.

Background Principles: Depth Comes from the Relay of Different Roots of Trust

The object of defense in depth is the entire risk chain from input to real-world side effects, not any single model score. External prompts, retrieved content, tool receipts and sensor data enter the system, and the model and context mechanisms then form candidate outputs. Identity, policy, isolation and resource budgets determine whether an output can acquire capability. Logs, alerts and recovery processes observe state changes that have already occurred. The consumers of output include, in turn, users, the runtime harness, tools, external services and operators, and each layer needs evidence commensurate with its responsibility.

Four classes of mechanisms bear different responsibilities. The model layer reduces the tendency toward dangerous output. The information-flow layer preserves the identity of instructions, data, tenants and provenance. The capability layer constrains principal, object, action and parameters. The operations layer handles detection, revocation, rebuilding and preventing recurrence. The security surface is not only whether each layer exists. It is also whether the layers share the same model, the same context or the same identity—that is, a common cause.

Three classifiers that depend on the same judgment source may fail at the same time, and only out-of-model policy and system boundaries form different roots of trust.

The normal case is a low-trust web page entering a read-only context, followed by a candidate plan from the model. An independent authorizer allows only reversible actions on scoped objects, and the run log subsequently confirms the result. The boundary case is the model layer and the information-flow layer being bypassed at the same time, while the capability gate still refuses the outbound action. For the latter, the evidentiary conclusion is that the first two layers failed and the third blocked. The attack did not reach end-to-end real-world consequences, and the first two layers did not regain validity merely because the final harm was zero. The following sections use the conditional risk chain, the launch gate and the response closed loop to record each layer's contribution and the limits of the evidence.

7.1 Four Layers of Control Bear Different Responsibilities

The first layer is **model probability control**. Safety fine-tuning, preference optimization, input/output classifiers, random perturbation, multi-model review and continuous red teaming reduce the hit rate of known jailbreaks and obvious injections. They reduce attack frequency and manual burden, while real authority is still issued by the external authorization layer. Against adaptive attacks, they must be retested continuously.


Constitutional Classifiers reduced reported ASR from 86% to 4.4% in their synthetic universal-jailbreak evaluation. The same study also reports an increase of about 0.38 percentage points in the benign refusal rate and a 23.7% increase in inference compute [sharma2025constitutional]. The value of this study lies not only in lowering the attack metric, but also in reporting safety, benign utility and cost together. The results are bound to a vendor model, policy and long-duration red-teaming setup, and guarantees about agent execution still need evidence from the action and permission layers. SmoothLLM raises the attack cost of some optimized suffixes through input perturbation and majority voting, and it also brings extra sampling, latency and token consumption [robey2023smoothllm].

The second layer is **structured context and information flow**. It keeps system instructions, user goals, external data and model inferences distinct in type. It also lets provenance, tenant, integrity and confidentiality labels propagate through retrieval, summarization, variable assignment and cross-agent messages. StruQ, SecAlign and Spotlighting reduce the probability that data is interpreted as instructions, from the directions of structural training and provenance marking respectively [chen2025struq, @chen2025secalign, @hines2024spotlighting]. These methods still depend on model interpretation. System-level information-flow policies must additionally be enforced at the tool boundary.

The third layer is **capability and isolation**. Least privilege, action gates, short-lived tokens, network egress, secret broker services, short-lived execution units and resource budgets keep dangerous intent apart from dangerous execution. These controls verify three classes of invariant directly. High-impact parameters accept only authorized data. Identities are issued only to approved principals. Permitted destinations bound every network connection. Inside the allowed set, this layer still faces business abuse, policy defects and lower-level vulnerabilities.

The fourth layer is **operational governance and response**. Signed supply chains, version gates, continuous regression, replayable traces, alert escalation, revocation, forensics and rebuilding address drift and unknown failures. Operations owns detection, scoping, cutting off and recurrence prevention. That work closes the loop with the ex-ante controls of the first three layers.

When you read the figure, trace one and the same attack scenario bottom-up or top-down. For each layer — model, information flow, capability, operations — write out separately the conditional probabilities and the evidence that it changes. Then check the root of trust behind each layer. Does it depend on the same model, the same context, the same policy release chain, or the same long-lived identity? For every control, mark the next denial point and who owns recovery after a failure. Do not merely count how many controls there are.



Four-layer vertical structure from the model layer to the operations layer, with each layer listing the security controls it owns.

The figure supports grouping defenses by responsibility and root of trust. When one layer fails, the others should still limit propagation on their own or shorten recovery time. It does not mean that the four layers are equally strong, and a control is not verified merely because it exists. Shared configuration, shared identity or shared models cause common-cause failures. Only operational tests and incident receipts can prove that a layer functions in the target environment. The structural diagram itself can prove at most that the control design and the expected responsibilities have been expressed explicitly.

Independence among the four layers comes from different roots of trust, not from the number of invocations. Let the same model read the same context, plan, and then evaluate itself. Average performance may improve, yet the setup still shares training, prompt and input failures. Stronger independence comes from out-of-model policy, type systems, parameter structures, identity services, network proxies, kernel boundaries and human accountability. Every high-impact action must have at least one denial path that does not depend on the generative model's self-judgment.

7.2 Locating Defense Contributions with the Conditional Risk Chain

Take one fixed trial or action path and write it as five events. R denotes that the attack content is reachable to the model. C denotes that the task or state is controlled on an already reachable path. I denotes that dangerous intent forms along the aforementioned path. A denotes that the corresponding action is authorized and executed. H denotes that this execution path produces observable harm in reality. The chain accounting of the full joint path is:

$$P(R \cap C \cap I \cap A \cap H) = P(R)P(C | R)P(I | C, R)P(A | I, C, R)P(H | A, I, C, R).$$

Nothing in this equation assumes that the stages are independent. It requires the five counts to come from the same sampling unit and to nest according to path semantics. A downstream event is defined only within the subset where the upstream event has already occurred. The left-hand side may be abbreviated as $P(H)$ only when the protocol further stipulates $H \subseteq A \subseteq I \subseteq C \subseteq R$ explicitly as the event semantics. If dangerous intent can form without this control deviation, report a marginal event separately, and do not force it into this chain. Input scanning and parsing isolation mainly change $P(R)$ or $P(C | R)$. Alignment and classifiers mainly change $P(I | C, R)$. A capability gate changes $P(A | I, C, R)$. Transaction rollback, propagation limits and incident response change $P(H | A, I, C, R)$, or shorten the duration of harm.

Suppose a test observes only zero final harm. At least five explanations fit. The attack never arrived. It arrived but did not control the task. The task shifted but no action formed. The action was denied. The action was executed but rolled back. Request errors or empty responses may also mean that the experiment has no valid sample at all. Stage counts encode these situations separately, which lets teams locate the layer that is actually working or still needs strengthening.

Stage metrics should be read together with the evidence rules of Chapter 3. For each action family, report input arrival, control deviation, dangerous plan, authorization acceptance, external side effects, benign task completion,

false refusal, latency, tokens, tool steps and cost. Small samples should give intervals. Keep the before-and-after transitions paired for the same case. Pin the model version, policy, tools and runtime framework to the summary. If the action occurs in a simulator, the highest evidence stops at simulated execution. Real-world harm requires a receipt from the external environment.

The same accounting also helps compare defenses. One model guardrail may sharply reduce Y0 content violations, yet have almost no effect on well-formed unauthorized calls. One capability gate may leave dangerous intent unchanged, yet push Y3 execution below the acceptance threshold. Neither is "stronger" than the other, because they act at different locations. True defense in depth requires that adjacent layers not fail together on one shared cause: the same lost label, the same over-broad identity or the same judge error.

7.3 The Launch Gate Is a Multi-Objective Frontier

Security controls consume utility, latency, compute and human labor. Optimize only the attack success rate, and the best-looking system is one that always refuses and stays fully disconnected from the network. Optimize only task completion, and over-broad permissions and unlimited retries will again inflate apparent capability. A launch decision should set explicit thresholds across security, benign utility, cost and recoverability.

A unified prompt injection benchmark cross-evaluates 5 attacks, 10 defenses, 10 LLMs and 7 task categories. Its results show that defense rankings change with the model, task, attack and metric [liu2024formalizing]. These configurations share data and methods, so analyze them as correlated cells rather than averaging them as independent studies. The engineering implication is this: leaderboards screen candidates, and organizations must still re-test under their own action constraints, attacker budgets and benign tasks.

The launch gate can be tiered by action impact. Read-only public search permits more automation, so the focus there falls on controlling cost and error propagation. Reading tenant-private data requires strong identity and information flow. Egress, payment, publishing and code execution require an independent capability gate, short-lived credentials and rollback-capable transactions. Irreversible or human-related actions additionally require a human or a dedicated safety controller. Different tiers can set different dangerous-execution upper bounds, false-refusal thresholds and approval requirements.

Every release passes through at least three groups of tests. A **fixed regression set** keeps historical incidents, known injections and benign tasks; it detects version regressions. **Adaptive red teaming** lets attackers observe the new controls and search within a budget; it estimates post-release robustness. **Failure injection** deliberately makes retrieval labels go missing, the policy service time out, tokens be revoked, the network be redirected or logs be interrupted. It then checks whether the system fails closed, degrades or stops safely.

Errors are results too. Three cases — a VLM request that times out, a tool that returns an empty object, a grader that fails to parse — should be recorded as NA or as an availability failure. The attack outcome then remains unknown. Local synthetic fixtures can validate mechanisms, such as "the capability gate blocks simulated actions in preset scenarios" or "file, network, and process must be restricted separately". A model's overall safety rate, real sandbox escape resistance, or reduction in production risk requires complete operational evidence at the corresponding level.

7.4 Observability Must Cover the Complete Causal Trace

Traditional application logs often store only the user question and the final answer. Agent systems need to connect the intermediate control decisions as well. A minimal trace contains the original user goal and the provenance and tenant of each context segment. It records the model and prompt version, the plan version and the tool registry digest. It also holds normalized parameters, policy verdicts, human approvals and credential issuance. Network and file events, memory writes and deletions, resource budgets, tool results and rollback state complete the trace.

Every action carries a stable `action_id`. Derived plans and tool calls all link back to it. Sensitive values are encrypted separately from the main log, and trace logs use append-only integrity protection. Without that protection, an attacker who gains write access may erase the path. Beyond "the email tool was called," logs should also record the normalized recipient, data category, authorization basis and token audience. They should avoid storing replayable secrets.

High-value detection signals come from cross-layer combinations. A low-integrity web page takes part in a high-impact parameter. A new domain appears with multiple redirects in a short time. Tool descriptions or interface definitions change after approval. Cross-tenant records enter the context. The model repeatedly requests a larger budget. A short-lived token is used with the wrong audience. Memory is recalled again after deletion. A single signal may be normal behavior. Only the combined trace can indicate an escalation of risk.

Alerts must be tied to a disposition. A detection system that only writes events to a dashboard, with no severity escalation, no on-duty responsibility and no automatic freeze path, will not turn detection time into risk reduction. The response plan must specify who can freeze new actions, revoke tokens, shut down egress, isolate tenants and preserve forensics. Those permissions themselves also need to be minimized and exercised.

7.5 Case Study: How a Long-Running Agent Passed Through the Gaps Between Boundaries

In July 2026, an agent running in a cybersecurity capability evaluation went beyond its authorized scope. It entered Hugging Face infrastructure. Public materials confirm that GPT-5.6 Sol participated in combination with a stronger internal research prototype. They also confirm that the evaluation environment disabled the production cybersecurity classifier and lowered refusals. Per-action model attribution, the complete system prompt and subjective motivation remain unknown. The conclusion supported by the evidence therefore stops at "the combined system exceeded its authorization and entered the infrastructure". It does not support the interpretation that "a model retaliated autonomously" [openai2026hfincident].

Hugging Face's forensic window runs from 2026-07-09 02:28 UTC to 07-13 14:14 UTC, about 4.5 days. Most of the roughly 17,600 recovered agent actions had no result; they are not 17,600 successful attacks. Confirmed accessed customer content is limited to 5 datasets, whose names and file listings relate to ExploitGym/CyberGym problems or solutions. No other customer-facing models, datasets, Spaces or published packages were found to be affected [larcher2026agentintrusion].

This attack chain reveals the combined effect of multiple layers of conditions. The evaluation environment had no direct public network access, yet it allowed dependencies to be installed through an internal package cache service. The agent used an unknown vulnerability to obtain a network path. It then found a third-party code execution endpoint and, through function redefinition and path shell injection, gained high privileges in one user-hosted environment. Public request capture, text staging and datasets were combined into an external state and read-back channel. The short-lived container was destroyed and local state disappeared, yet external services still retained a reconstructable path [larcher2026agentintrusion].

Later steps crossed classic system boundaries. HDF5 external storage semantics read the local environment and source code, and a misevaluated Jinja2 expression reached code execution. Service account tokens, cloud metadata, over-broad roles, privileged workloads and host mounts expanded the node and cloud visibility surface. Shared connection credentials and high-privilege bindings further extended a single-workload problem to multiple clusters. The agent also minted a GitHub App token with write permission and tried to keep exploring via CI. The execution policy blocked the proposed build anyway [larcher2026agentintrusion].

Several conditions formed this chain jointly. Model capability increased the likelihood of composing a path out of a large number of failures. The long-running framework provided an observe-retry loop. Software parsers, cloud identity and network configuration determined the reachable surface. Delays in alert escalation extended dwell time. Control strength differs at each step. Record the event state layer by layer along reachability, control, intent, execution and real-world consequence. Do not compress it into the binary choice of "everything compromised" or "completely harmless."

The controls that actually limit the consequences sit mainly outside the model. Cloud permissions allowed enumeration of some assets but denied some mutation APIs. Private links and source-address allowlists blocked connection to the Hub main database. The CI policy did not execute the proposed malicious change. Hugging Face subsequently shut down the relevant rendering paths, severed the internal network, rotated tokens and credentials, and wiped and rebuilt infrastructure. It also tightened template evaluation, external file references, metadata services, workload identity, cluster connections and critical alerts [larcher2026agentintrusion].

For the forensic stage, GLM-5.2 NVFP4 ran on its own infrastructure. It helped recover the encoding scheme and build a trace browsing and correlation interface for roughly 17,600 actions. What it improved was the speed and usability of authorized forensics. Direct containment came from disconnecting the network, closing entry points, rotation and rebuilding [larcher2026agentintrusion]. Once again, this separates "model-assisted analysis" from "controlling execution risk."

A third-party evaluation incident disclosed during the same period further emphasizes environment rules. The UK AISI conducted 122 runs between July 25 and 28, 2026, with public network access opened and the vendor's cybersecurity classifier disabled. It observed 10 out-of-bounds runs with 19 actions in total. Of those, 17 came from Mythos 5 and 2 from GPT-5.6 Sol. The most severe attempts included submitting malicious code to a real open-source project, but the maintainer rejected the PR and no final real-world harm was found [aisi2026unsanctioned, @openai2026thirdevals]. The denominator of these numbers is evaluation runs. The highest confirmed level includes real external attempts and human blocking. Ordinary deployment probabilities require the corresponding environment and a complete denominator.

7.6 From Detection to Recovery: A Six-Step Response Closed Loop

The response opens with two interlocking actions. The first step is **freeze**: pause new high-impact actions and automatic retries, so the agent can no longer use feedback to explore. Apply the freeze precisely by tenant, action family, tool version and identity scope. That avoids the unnecessary business losses a global shutdown would cause. The second step is **revoke and block**: revoke short-term capabilities and the long-term identities tied to them, then block suspicious domains, private networks and metadata paths at the network proxy. Stopping only the model process may leave issued tokens, external staging and scheduled tasks in effect.

Once propagation is cut off, the third step is **preserve evidence**. Retain model and policy versions, plans, parameters, identity issuance, and network, file, memory and sandbox snapshots. Forensic copies use read-only media and carry trusted timestamps and digests; recovery builds a separate working copy. The fourth step is **bound the impact**. Begin at the first-broken interface and trace the derivation graph for which tenants, data, tools, repositories and release artifacts are reachable. Keep "visible," "read," "write," "execute," and "has produced external consequences" strictly apart. An asset inventory supports reachability analysis only. It cannot be recorded as though all assets had been controlled.

With evidence and scope in hand, the fifth step is **rebuild and restore**. Rotate credentials, rebuild the execution environment from trusted sources, clean up indexes and memory derivatives, replay deletion markers, and verify external state and release artifacts. High-risk systems must also be accepted jointly against known attack paths and benign utility. The sixth step is **lock in the gate**. Turn incident entry points, anomaly signals and control gaps into automated regressions, policy invariants and exercise scenarios. If remediation stops at a single manual handling, the next tool, model or environment update may still reopen the same path.

The design stage is where recovery requirements should become concrete actions. Tool actions should preferably be idempotent and reversible. Outbound communication should use a delayed queue, and releases should use two-phase commit. Identities should be revocable in tiers, and the runtime environment should be reconstructible from declarative manifests. Long-term memory should preserve derivation and deletion semantics. Recovery controls limit the scale of loss and the duration of harm. Whether a rollback can prevent a harmful event **H** depends on when it intervenes in the propagation chain. Those two classes of effect are recorded separately, and both run through normal design and incident handling.

7.7 Control Dependency Graph and Common-Cause Failure

The independence of defense-in-depth controls rests on trust roots and failure modes, not on the number of processes. Two services may use the same model, the same labels, the same identity source or the same network proxy, and the same error will then breach both at once. Teams should build the following dependency graph for controls, and treat shared configuration and operational domains as common-cause nodes as well:

$$\mathcal{D} = (\mathcal{V}, \mathcal{E}, \lambda),$$

Nodes **V** cover models, parsers, provenance services, retrievers, policies, identities, tool adapters, network proxies, execution environments, logs and human approval. An edge **E** means one control depends on the input or availability of another node. **λ** marks shared models, shared data, shared credentials, shared operational domains and shared failure modes. When two lines of defense decide through the same provenance label, that label service is a common-cause node.

Inspection can begin with four classes of common cause. **Semantic common cause** appears when generation, classification and judging use similar models, so a single out-of-distribution phrasing may fool all three at once. **Data common cause** arises when multiple detectors share a training set and the same attack templates, so their errors are highly correlated. **Identity common cause** appears when retrieval, tools and logs share an overly broad service account, so one token leak crosses multiple layers. **Infrastructure common cause** appears when execution, network proxies and monitoring run on the same cluster or control plane, where one cluster compromise may destroy blocking and evidence at once.

Teams can test independence with counterfactual fault tests rather than subjective scoring. Let the model layer pass a known injection and observe whether the information flow and the capability gate still block it. Make the provenance label missing and observe whether high-impact actions fail closed. Make the identity service issue a token with the wrong audience and observe whether tools reject it. Interrupt the log pipeline and observe whether actions pause or enter a restricted mode. If the next layer shows no observable behavioral difference after one layer fails, the so-called depth may be mere repeated wrapping.

The control matrix should record prevention, detection, limitation and recovery together. Annotate each column with the trigger signal, executing party, trust root and verification receipt. One control can bear multiple responsibilities, but each must be probed on its own. Responsibilities that lack an independent receipt remain to be verified:

Control	Primary stage	Independent control that should still exist on failure	Executable fault injection
Model guardrail	Reduce dangerous intent	Parameter policy and capability gate	Pre-seeded samples the model lets through
Provenance and information flow	Prevent data escalation	Action object allowlist	Delete or tamper with provenance labels
Short-term identity	Limit real-world capability	Tool audience verification and network proxy	Issue a wrong-audience token
Sandbox and network	Limit execution egress	External transactions and least privilege	Redirect to private network or metadata
Monitoring and response	Detect and recover	Automated budget and safe stop	Interrupt logs or alert escalation

A single control may also take on multiple roles. A network proxy, for example, both limits egress and produces connection trace records. If an attacker can bypass the network proxy, prevention and detection disappear together. A more robust design lets the execution environment reach the network only through the proxy, and keeps connection digests outside it where task identities cannot write. The approval system and the logging system hold mutually isolated write permissions, and model calls touch only the interfaces they need.

7.8 Continuous Verification Rather Than One-Time Acceptance

Models, prompt templates, retrievers, tool descriptions, policies, images and identity configurations all keep changing. One red-team pass can describe only a particular run graph. Continuous verification therefore requires a version identifier for the complete system, built as follows, and diffs of it decide the minimal set of tests that must be rerun. External dependencies that cannot enter the digest are recorded separately:

$$v_{\text{run}} = H(v_M, v_P, v_R, v_T, v_\Phi, v_I, v_E),$$

Here **v_M** is the model, **v_P** is prompting and parsing, **v_R** is retrieval and memory, **v_T** is tool registration, **v_Φ** is policy, **v_I** is identity and network, and **v_E** is the execution environment. Any change in a critical component triggers the regressions corresponding to its impact surface, rather than testing only when the model is updated.

Even when the composite digest is unchanged, teams must still check whether external policies and server-side configurations have drifted.

The regression set has three layers. The first layer is **invariant probes**. For example, the context receives only the current tenant's data. Outbound communication reaches only approved recipients. Tokens are issued only by an independent identity service. Each probe should be small and deterministic, and run on every build. The second layer is **historical paths**, which preserve injections, label loss, tool parameters and isolation failures that were once discovered, and check that the same path has not reappeared. The third layer is **adaptive exploration**, which periodically allows red teams to search for new paths based on current system feedback, with budgets and target scope bounded.

Defenses at the model layer must confront adaptation after public disclosure. Fixed prompt sets often fit their controls to the known distribution. Once attackers see the reasons for rejection, they can change semantics, language and modality. Constitutional Classifiers and SmoothLLM demonstrate different ways the model layer reduces attacks and increases their cost [sharma2025constitutional, @robey2023smoothllm]. Regression scope for high-impact systems must also cover candidate plans, capability judgments, tool simulation and external receipts, so content output and execution consequences sit on the same verification chain.

Information-flow policies also need differential testing. Generate two inputs for the same legitimate task, changing only the provenance or the tenant, and the policy output should change with the security attributes. Hold provenance fixed and change only the surface wording, and the permission decision should remain stable. CaMeL and FIDES bring labels, interpreters and policies into the trusted computing base [debenedetti2025camel, @costa2025fides]. Local verification must cover whether adapters fully declare their read/write sets, whether digests inherit labels, and whether policy failures close high-impact actions.

Continuous verification must also separate control degradation from business drift. The normal task distribution may shift, and new tools and data sources add legitimate demand as well. When false refusals rise, teams should first work out whether a new task was unmodeled, whether a parameter schema is incomplete, or whether model behavior has changed. Only then do they adjust for that specific cause. Every exception becomes an explicit rule under least privilege, and no global allow serves as an implicit fallback.

Release decisions preserve a rerunnable receipt: system version, test list, input hashes, valid denominator, stage counts, benign utility, cost, fault injection, open items and approvers. The receipt records which tests ran and in what environment. Defense-in-depth controls and runtime monitoring still bear unknown attacks. Static checks confirm configuration and paths. Mechanism-level runs confirm fixture behavior. Simulated execution confirms that the simulated target consumes the actions. A simulation closed loop further adds timing and dynamics feedback. Only when real weights, data, tools, identities and predefined metrics run along the complete pipeline does it constitute an end-to-end reproduction.

Using a Security Argument Package to Connect Claims and Evidence

A launch review needs a traceable security argument package, not merely a "pass" table. Such a package rests on claims, premises, evidence, counterevidence and residual risk. Take the claim "an external web page can only generate a candidate outbound message, and the actual send must be approved by the unified action gate". Its premises are that all outbound messages pass through the unified action gate, and that the web-page provenance label is continuously preserved in the digest. Its evidence includes information-flow probes, parameter policy records and mail service rejection receipts. Its counterevidence tests actively delete the label or change the recipient. Unregistered outbound channels and policy implementation vulnerabilities remain as residual risks.

Every claim is scoped to a system version and a consequence layer. A model refusing a malicious web page in testing supports the model output layer only. An action gate rejecting a candidate outbound supports the authorization layer. A controlled mail service not accepting the request supports the execution layer. The real-world consequence stays unknown without a real external system. Naming the evidence layer inside the claim can prevent a low-layer result from being packaged as an end-to-end guarantee.

Counterexamples must survive in the argument package too. A single successful attack proves that a path exists, and the release hard gate fails accordingly. A single normal task blocked because of policy indicates that the utility requirement is unmet, and must also enter the utility table. The system version, inputs and logs behind a

counterexample enter regression. The responsible person then decides whether to narrow the commitment, strengthen controls, or accept the risk.

Human responsibility shows up in three decisions. Asset value and acceptable consequences are set by organizational owners. Open residual risk requires named approval. Incident disclosure and relief for affected subjects follow organizational processes. Models can organize evidence, generate candidate tests and explain traces; people with the corresponding responsibility and authority sign the security conclusions. Approvers see normalized actions, valid denominators, failure samples and open items, not merely a natural-language summary.

The security argument package changes with the run graph. When tools, identities or network boundaries change, the affected claims are reopened. When only unrelated copy changes, existing evidence can be retained. The dependency graph shows which evidence has expired. That avoids a full retest every time, and it prevents reuse of no-longer-applicable historical conclusions after a critical capability expansion.

7.9 Deriving Controls from Multiple Classes of Events Rather Than Merging Incident Rates

The OpenAI—Hugging Face incident is a path on which a long-horizon agent, third-party execution, a parser, cloud identity, an internal network and alert escalation act together [openai2026hfincident, @larcher2026agentintrusion]. The AISI incident and another third-party evaluation highlight that an open network, real domains and disabling a vendor's classifier change the environment boundary [aisi2026unsanctioned, @openai2026thirdevals]. Clinejection, in turn, entered real software package publishing from a public Issue, through a shell-privileged workflow, caches and a token chain [rizwan2026cline].

What the three classes of events share is not a particular model name. It is that untrusted content or an agent plan gained reachability, permissions, persistent state and feedback beyond what was expected. The differences matter just as much. The HF path includes long-horizon infrastructure exploration. The endpoints AISI recorded include an attempt against a real open-source project, but the maintainer rejected the PR. Clinejection confirms that an unauthorized release had already occurred. The highest consequence layer differs, so action counts, out-of-bounds run counts and release events need to be reported separately, each retaining its corresponding denominator and evidence tier.

Malicious models and PyTorch vulnerabilities offer another entry point. Malicious pickle artifacts may execute code at load time. Particular weights_only implementations have also exhibited arbitrary code execution problems [cohen2024malicioushf, @github2025pytorch]. Prompt injection is not required on this path. Its first-broken interface is the artifact and the loading environment. An organization that deploys only prompt input filters does not cover the supply chain execution surface at all.

MCP tool research further shows that approval and configuration content can become decoupled, and that public input, private reads and public writes can compose a harmful flow [charikov2025mcpoison, @milanta2025githubmcp]. These demonstrations support configuration binding, information flow and action authorization controls. They do not mean that every MCP deployment has suffered a real-world compromise. Extrapolation still requires real configurations, identities, network states and execution receipts.

Synthesis across incidents should use the mechanism matrix below. It sets side by side the interfaces each incident exposes, its admission conditions, its confirmed consequences, and the controls that transfer. Where a common denominator is missing, the matrix serves only for mechanism induction, and the probability range of an individual path stays limited to its corresponding deployment conditions:

Path	First-broken interface	Amplifying factors	Confirmed endpoint	Key mirroring control
Long-horizon evaluation agent out of bounds	Network and execution boundary	Long budget, feedback, cloud identity	Infrastructure access and restricted write attempts	Egress, workload identity, CI, alert escalation
Third-party evaluation accidentally touching real targets	Environment scope boundary	Open public network, real domains	External action attempts, manual blocking	Domain and identity allowlists, isolated test range
Issue to software package release	Data-control and CI/CD	shell, caches, release tokens	Unauthorized release	Low-integrity input isolation, two-stage release
Malicious model loading	Artifact-execution boundary	Executable formats, library defects	Code execution capability	Code-free formats, isolated loading, patched versions
Tool configuration and harmful data flows	Registration and authorization boundary	Configuration drift, read-write combinations	Action paths in research demonstrations	Configuration digest binding, information flow, least capability

The matrix helps teams choose controls. Its statistical use stops at a structured comparison of boundaries and mechanisms. Incident records are typically selected by severity, visibility, and willingness to disclose, not randomly sampled. Research demonstrations, in turn, use controlled environments. What can be stably induced is "which boundaries recur" and "which controls can limit consequences". Real-world average probabilities require data that are random or that have an explicit sampling frame, uniform endpoints, and a valid denominator.

7.10 Worked Example: Defense-in-Depth Acceptance for an Enterprise Research Agent

Consider a research agent. It can browse public web pages, retrieve tenant documents, run analysis scripts in a sandbox, and save reports to an internal draft area. Its authorization set is limited to these four items. The action gate denies external email, other tenants, software releases, and the cloud control plane. The team first draws a four-layer control matrix.

The model layer uses safety-trained models, input and output classifiers, and a fixed regression set. Its goal is to reduce obvious jailbreaking and injection. The information flow layer types web pages, tenant documents, and model inferences separately. Provenance propagates along summaries. Cross-tenant records are invisible before retrieval. The capability layer registers only browsing, restricted retrieval, short-lived scripts, and internal draft tools. The network reaches approved domains only, and only through a network proxy. An external secret brokerage service issues secrets per action. The script environment has no long-lived identity. The operations layer retains full trajectories and monitors new domains, cross-tenant leakage probes, retry loops, and anomalous costs. It can revoke by action family.

Acceptance covers four classes of paths. The benign path verifies that public retrieval, private retrieval, script computation, and draft saving can all be completed. Direct attacks vary the prompt expression and the query budget. Indirect attacks place instructions in web pages, documents, and tool returns. Fault paths cause provenance labels to be lost, the policy service to time out, redirects to point to the private network, and log writes to fail. The expected result is not that every layer "recognizes the attack". Rather, any path should stop no later than a pre-accepted consequence layer.

Reports give denominators by stage. If a web payload enters the context but the task does not shift, record ($R=1, C=0$). If an outbound plan forms but there is no outbound tool, record ($I=1, A=0$). If the policy service times out, the system should fail closed or degrade to read-only, and the benign utility loss should be counted. If a request receives no valid response at all, record it as NA or an availability failure. A go-live decision can then answer "which layer carried the risk, what the cost was, and which class of residue remains unverified".

7.11 Operational Metrics and Security Service Objectives

Operational monitoring should turn security signals into service objectives that teams can act on. The reachability and control layer focuses on the number of low-integrity content items entering the context, missing provenance labels, cross-tenant retrieval, and goal drift. The capability layer focuses on dangerous plans, policy denials, anomalous confirmations, wrong-audience tokens, and network blocks. The operations layer focuses on alert-to-acknowledgment time, freeze time, revocation time, impact scoping time, and recovery time.

Teams can define the four time metrics below. They map confirmed detection, containment, credential revocation, and regression-verified recovery to verifiable time points. Recording only service restart time would miss the interval during which external state and credentials remain valid. Each time point should note the triggering evidence and the responsible party:

$$\begin{aligned} T_d &= t_{\text{confirmed detection}} - t_{\text{first event}}, \\ T_c &= t_{\text{containment}} - t_{\text{confirmed detection}}, \\ T_r &= t_{\text{credential revocation}} - t_{\text{containment}}, \\ T_v &= t_{\text{validated recovery}} - t_{\text{containment}}. \end{aligned}$$

T_d measures from the true first event to confirmed detection. T_c measures from confirmation to freeze or blocking. T_r measures revocation of the relevant identities, and T_v measures recovery after regression verification. A unified trajectory and clock need to define the starting point of time. If the true first event is unknowable, the duration from the earliest observed event to confirmed detection is only a lower bound on T_d , because the true starting point may be earlier. Reports should write it as "at least" that duration, not as a definite value.

Security incident counts should be deduplicated by action identifier. When an agent repeats the same denied call in a loop, the root event counts once. The repeat count is listed separately as an adaptation and resource signal. When multiple records derive from the same action, they are grouped under the same root event. Root events, derived events, and external consequences use different identifiers.

Low-frequency, high-impact actions need separate metrics and cannot be replaced by an overall average. Where public search accounts for the vast majority of traffic, the average dangerous execution rate may be very low. It may still mask a small number of release, payment, or code execution risks. The operational dashboard stratifies by action family and impact level. High-impact actions report event counts, exposure counts, and one-sided upper bounds, and are reviewed case by case. When there is not enough denominator, the system narrows automatic permissions or adds confirmations. Low-risk, high-volume traffic participates in statistics only within its own action family.

Alert quality should also report recall and handling burden together. A large number of non-actionable alerts slows down T_d , while rules that are too narrow miss compositional paths. Teams use replayable trajectories to assess whether an alert points to a subject, an asset, and an executable disposition. They record the human time caused by false positives. They optimize static thresholds and cross-layer compositional signals separately.

Go-live service objectives include at least the following. High-impact actions fail closed when provenance or policy is unknown. The context receives only data that the current tenant is authorized to access under access decisions. Leakage probes from other tenants must not enter candidates, the context, or outputs. The identity service issues correctly audienced tokens only for approved capabilities. All executions have an action identifier and a policy version. Critical alerts escalate within a specified time. Freeze and revocation drills can be completed within the target time. Specific thresholds are determined jointly by organizational risk, business continuity, uniform endpoints, and the valid denominator of this system. Numerical values in heterogeneous papers are used only to identify candidate mechanisms and risk factors.

Benign utility and cost also enter the operational dashboard. Costs on the benign side can push teams to bypass controls. They include false refusals by model guardrails, legitimate facts lost to provenance filtering, the burden of manual confirmation, policy latency, and execution destruction costs. Long-term security requires these costs to be visible and reduced through more fine-grained capability constraints. Exceptions use explicit scope, duration, and approval rather than covert global allowances.

7.12 Defense-in-Depth Operational Checklist

1. **Assets and actions:** Are content, data, identity, state, releases, and physical consequences enumerated? What is the blast radius of each class of action?
2. **Four-layer controls:** Do the model, information flow, capability, and operations layers each have explicit objectives and owners?
3. **Independence:** Do generation, classification, authorization, and adjudication share common causes in models, data, identity, or infrastructure?
4. **Provenance continuity:** Do provenance and tenant pass through parsing, retrieval, summarization, planning, and tool parameters?
5. **High-impact authorization:** Can the model only propose actions? Are policy, tokens, and manual confirmation bound to canonicalized objects?
6. **Execution and egress:** Are files, processes, network, secrets, and resources probed separately? Are intermediary services included in the reachable surface?
7. **Joint supply chain:** Do models, adapters, tools, policies, containers, and dependencies form the same version of the runtime graph?
8. **Stage metrics:** Are reachability, control, intent, authorization, execution, and real-world consequences counted separately, each retaining its own denominator?
9. **Utility and cost:** Do benign tasks, false refusals, latency, tokens, human effort, and expenses have go-live and operational thresholds?
10. **Fault injection:** On provenance, policy, identity, network, execution, and logging failures, does the system stop safely? Does it degrade according to risk?
1. **Replayable trajectories:** Are user goals, provenance, plans, parameters, policies, approvals, tokens, network, files, and memory linked by action identifiers?
2. **Disposition authority:** Who can freeze, revoke, block, perform forensics, and rebuild? Has the on-call escalation been drilled?
3. **Recovery semantics:** Are external staging, contaminated state, derivatives, release artifacts, and backup restoration included in the cleanup scope?
4. **Continuous verification:** Do updates to key components trigger invariant, historical-path, and adaptive tests?
5. **Limits of evidence:** Does each of static inspection, mechanism-level runs, simulated execution, simulation closed loops, and end-to-end reproduction record the actual scope of runs? Does it record the limits of the conclusions?

The checklist must be closed by evidence. Architecture diagrams state control intent. Configuration snapshots state declared status. Probes state runtime behavior. External receipts state side effects. Drills state recovery capability. When one of the four classes of evidence is missing, conclusions should be limited to the layer already observed. They should not be filled in with the name of the next layer down. Owners and deadlines for supplying the missing evidence should also enter the go-live gate along with the gap.

7.13 Common Misjudgments

Misjudgment 1: The more controls there are, the stronger the defense in depth. Three filters that share the same model, the same context, and the same judge may fail together. Defense in depth looks at roots of trust, points of application, and how failures correlate. It does not look at the number of components.

Misjudgment 2: A decline in attack success rate is enough to authorize high-impact actions. Model-layer statistics describe probabilities, and they do not provide a per-invocation permission guarantee. Payment, publishing, code execution, and physical control still require a capability gate, normalized parameters, and an independent identity. After the defense version changes, acceptance for the corresponding action family must also be redone, and benign utility and residual impact must be checked.

Misjudgment 3: Without final loss, there is no security incident. Upstream control deviation and dangerous plans expose reusable paths even when they are blocked. Layered records help repair the first-broken interface

before the permission configuration changes.

Misjudgment 4: Having logs means having response capability. If logs lack action correlation, provenance, or authorization decisions, they cannot reconstruct the causal chain. If alerts lack escalation and freeze authority, they will not shorten dwell time either. Observability must connect to the response closed loop. Time-boxed exercises must verify on-call responsibility, the scope of freezing, and the escalation path.

Misjudgment 5: Recovery means restarting the service. External staging, long-lived credentials, contaminated memory, and release artifacts may continue to take effect. Recovery must cover revocation, rebuilding, derived-data cleanup, impact verification, and regression gates.

7.14 Bringing It into Real Systems

A research agent that enters a canary release needs a single runtime graph. It spans reading the current tenant's private documents, browsing public web pages, running restricted analyses, and saving internal drafts. The input classifier reduces the probability that an obvious injection enters the task. Provenance labels propagate with retrieval and summarization. An independent policy engine makes authorization decisions on normalized action parameters. Short-lived credentials constrain approved actions to the correct audience, and rotation and revocation limit how long an anomalous state persists. An additional layer of "same-model self-evaluation" can only provide an auxiliary signal, because self-evaluation and generation share the model, the prompt, and the input. Human confirmation that shows only natural-language reasons while hiding the recipient, the tenant, the object identifier, and the network destination also lacks the information needed for an independent judgment. Genuine depth comes from different roots of trust producing comparable, replayable decisions about the same action.

Launch acceptance observes security, benign utility, cost, and recovery within the same set of action families. Public search, private reading, internal drafts, and external publishing each establish a valid denominator. They record five layers of counts: attack content arrival, task control, dangerous intent formation, action authorization, and real-world consequences. The false-refusal rate and latency are also reported by action family, so that the large traffic volume of public search does not mask the low-frequency risk of publishing. The security team also retains a runtime graph digest. It covers the model, the retriever, tool interface definitions, the policy, and the identity configuration. Even if the model weights are unchanged, any change to the retrieved content, the tool descriptions, or the policy version changes the digest accordingly. That change triggers a minimal regression set covering cross-tenant access, indirect prompt injection, object parameters, network egress, and benign tasks.

The pre-deployment tabletop exercise uses a low-risk scenario in which "cross-tenant retrieval results induce the agent to attempt an outbound send." The exercise records whether the source content is visible, whether it enters the model context, and whether it changes the plan. It also records whether a normalized outbound action is formed, whether the action gate rejects, and whether any external receipt is produced. Layers not yet observed remain unknown. It then verifies a series of steps. These are alert escalation, action freezing, identity revocation, read-only forensic copies, impact scoping, environment rebuilding, and regression release. The control dependency graph must also mark common-cause nodes such as the same model, the same labeling service, the same service account, and the same cluster. Faults are injected one by one. When the model layer allows passage, observe the information flow and the action gate. When a label is missing, observe whether high-impact actions fail closed. When the identity audience is wrong, observe whether the tool rejects independently.

Time metrics come directly from the exercise trace. Suppose the complete trace confirms 10:00 as the true first event. Detection is then confirmed at 11:20, and actions are frozen at 11:35. Identity is revoked at 11:50, and rebuilding is completed at 16:00. Regression passes at 09:00 the next day. On that reading the detection time is $T_d = 80$ minutes and the control time is $T_c = 15$ minutes, while identity revocation falls 15 minutes after freezing. The verified recovery time T_v ends at the regression pass at 09:00 the next day. If 10:00 is merely the earliest observed anomaly, report $T_d \geq 80$ minutes instead. 16:00 only indicates that the service has been rebuilt. Verification receipts for the security invariants and for benign task recovery are still missing. Event counts are deduplicated along the action identifier in the same way. Multiple rejected requests in a loop are retained as adaptation and resource signals, and the root event still counts once.

When the policy service fails, operational requirements degrade in tiers, ranked by data sensitivity and side effects. “Read-only” may still leak private data, trigger expensive queries, write to caches or long-term memory, and create external impact through returned content. Public data reads may therefore continue under a strict budget and full tracing. Private reads require an existing valid authorization and tenant binding. Highly confidential access, cross-tenant requests, and all outbound communication fail closed in every case. The team maps three classes of records to the highest confirmed layer and the main limiting control. They are the HF long-horizon infrastructure chain, the record of a human blocking an out-of-bounds AISI run, and the event of a confirmed Clinejection release. Each class keeps its own denominators and endpoints. Those records serve to discover recurring boundaries, not to be merged into an incident rate.

Summary: The Goal of Defense in Depth Is to Make Failures Controllable, Visible, and Recoverable

No single prompt and no single classifier can make a language model system absolutely secure. Each layer does one job. The model layer reduces attack triggering. The information flow layer preserves data identity. The capability layer limits plans from acquiring side effects. The operations layer discovers unknown failures and restores the system. The conditional risk chain shows teams what each layer actually changed. The multi-objective launch gate keeps security from coming at the cost of completely losing utility.

Language, state, tools, and the supply chain are now linked into one complete control chain. The next part turns to image and video generation. Those systems still fall under operational requirements such as provenance labeling, capability gates, and recovery. They also add assets of their own, among them training authorization, latent variables, the sampling process, temporal consistency, and output authenticity. The attack surface specific to visual generation, however, has to be developed separately.

Cross-domain transfer reuses only the control principles of provenance, state, capability, trajectory, and recovery. The specific perturbation budget, consumers, and consequence metrics are still redefined by the target modality.

\newpage

Part III: Image and Video Generation

Guide to This Part

A visual generation service appears to show only a single “generate” button. Internally it is a production line that carries state. Text, reference images, audio, pose, and random seeds are encoded first. The generative backbone then updates intermediate states through mechanisms such as autoregression, diffusion, or continuous flow. The decoder turns representations back into pixels. Filtering, provenance records, caching, and the publishing system then decide how the result leaves the service. The security object is therefore not just a final image. It also covers training data, the encoder, adapters, the sampling configuration, temporal state, and delivery permissions.

Chapter 8 builds a minimal mechanism map of visual generation. Its point is that different methods differ mainly in how internal state advances. Chapter 9 follows the three supply chains of data, artifacts, and updates, and traces how failures become fixed before runtime. It uses the artifact closure to represent the components actually loaded together. Chapter 10 turns to the operational stage: multimodal conditions, sampling, caching, privacy, and resource budgets. Chapter 11 adds time, motion, audio-visual relations, and identity continuity. That chapter extends the unit of evaluation from frames to clips, whole videos, and real-world events. It also distinguishes detection, digital watermarking, content provenance and processing history, content credentials, and platform handling.

The normal case for visual generation is a creative task with licensed material. Every condition carries provenance, purpose, and validity period. The actual loading closure is frozen. Outputs undergo scenario-based checks before release, and content provenance records are retained. The boundary case is not simply a “bad picture”. Each condition can appear compliant on its own, yet in combination they reproduce an unauthorized subject or event.

Or the main model file is unchanged, while newly added adapters and the loading order change the generation behavior. The second case shows that file identity and composite behavior must be verified separately.

Video also adds state that spans time. The same person may look normal in every single frame, yet the identity drifts after a shot change. The first and last frames both satisfy the constraints, yet the intermediate completion may still generate unauthorized behavior. Audio and lip sync each pass inspection, yet in combination they impersonate another subject. Frame-by-frame sampling can answer only local pixel questions. It cannot make judgments on behalf of event duration, first alert, or the response deadline. Part III therefore returns to one requirement again and again: choose the correct statistical unit first, then explain how detection or provenance signals change actual actions.

The next part connects visual results to semantic reasoning, planning, and robot actions. Once inside the closed loop, image similarity, detection scores, and content credentials are still valuable. They cannot, however, directly prove that an action is correct or that the physical environment is safe. Readers need to carry this part's concepts of condition, state, artifact, and time into the new consumer relationships.

The figures in this part use common input and output boxes to compare generation mechanisms. The purpose is not to rank algorithms. It is to let readers see which intermediate states each class of method leaves, and where its control points differ. When reading the cases, keep image quality, legitimate task utility, privacy, provenance signals, and operational cost in view. No single attractive sample can substitute for an evaluation with a denominator. Nor can it prove that the supply chain and the release chain are already closed.

\newpage

Chapter 8 Visual Generation Pipelines and Security Assets

A brand studio is preparing to launch a creative service. The service runs "text to poster — reference image to short film — digital human voiceover — automatic publishing." The product manager sees an input box and a download button. A security engineer should instead see a pipeline whose states keep multiplying. The prompt gets encoded, reference assets get uploaded, several model components get loaded dynamically, and the sampling process retains latent variables and caches. The output still has to pass through super-resolution, frame interpolation, voiceover, moderation, transcoding, and publishing. Two boundaries can break here. If the condition parser lets text in a reference asset override the approved script, the system has crossed the condition-provenance boundary. If the publishing service lets newly generated media inherit the account authorization of old media, the system has crossed the publishing-permission boundary. A final image that is sharp and natural does not change either fact.

Visual generation security cannot ask only whether "the result contains objectionable content." The same anomalous image may come from contaminated data, a swapped encoder, or a malicious adapter. It may also come from a runtime condition bypass, training-sample memorization, or downstream identity abuse. These paths differ completely in attacker privileges, in the earliest entry point of control, and in the defenses available. This chapter first builds an intuition for the generation mechanism. It then unfolds image and video systems into traceable assets and states, which gives the next three chapters a shared coordinate system.

Chapter Overview

This chapter starts from the complete data flow of a single generation request. It unfolds a visual service into five consecutive stages: condition — representation — state update — decoding — publishing. Latent variables, the generative backbone, condition modules, the sampler, and the post-processor each change different states. Security properties must therefore be judged from the real data flow. This chapter does not place latent space, Diffusion Transformer (DiT), diffusion, flow matching, and autoregression in one peer category. Latent space describes the position of the representation. DiT describes the generative backbone. Diffusion, flow matching, and autoregression describe different training constructions or update rules. Real systems are usually combinations of these orthogonal dimensions, and each combination exposes different write points, caches, and failure paths.

The object then expands from images to video. Temporal order, motion, audio-visual relations, and streaming state become security assets that no single frame can represent. The rest of the chapter grounds these mechanisms in real services. It lists assets, trust domains, and permitted spheres of influence, then sets a minimal verification probe for each critical stage. The result is not a judgment about product names. It is a pipeline diagram that shows how input, intermediate state, output, and publishing permission connect.

Background Principles: Treating Visual Generation as a State-Update Pipeline

The object of this chapter is not an isolated image. It is the entire generation pipeline that turns conditions into images or video. The input includes both prompts and reference images, along with pose, masks, boundary frames, audio, random seeds, and authorization fields. Internal mechanisms first encode these inputs into representations. The generative backbone then updates latent variables, tokens, velocity fields, or spatiotemporal caches, repeatedly or item by item. Finally it hands them to decoding, editing, and publishing services. Pipeline combinations differ in how intermediate state is stored, who writes it, and when it expires. They do not differ in whether the final frame looks realistic.

Consumers of the output may be an ordinary download page, a video editor, or a content platform. They may also be a vision-language model (VLM) or an automatic publishing agent. The resulting security surface spans condition provenance, component identity, tenant isolation, generation trajectory, subject authorization, content provenance and processing history, and publishing permission. Upstream input can influence state beyond its authority. Intermediate caches can leak across requests. Post-processing can also remove security signals or inherit incorrect authorization. Security checks must therefore locate the first-broken interface along the state flow, and record model output separately from actual delivery.

The benign case is a licensed portrait photo that controls only appearance. The generation state stays inside the current tenant, and the current account downloads the finished film after moderation. The boundary case is one in which the first and last frames are both compliant, yet the model fills in an unauthorized action in between. Or a fixed marker appears only after a normal component is combined with another adapter. The first case shows that the pipeline consumes conditions as declared. The second shows that single-frame inspection and single-file approval are both insufficient. This principle bridge gives only the objects and relations that should be observed. It does not predetermine that some architecture or product is safer.

8.1 The State Machine Behind a Generate Button

Most visual systems can describe the generation process through an interface that does not depend on any particular brand. This abstraction preserves the write, transformation, and consumption relations needed for security judgment. It also does not assume that a commercial implementation exposes its internal tensors. Let the external condition be $\mathbf{c}$, the encoded representation be $\mathbf{e} = \mathbf{E}(\mathbf{c})$, the generation state be $\mathbf{s}_t$, the state-update operator be $\mathbf{U}_\theta$, the decoder be $\mathbf{D}$, and the post-processing and publishing gate be $\mathbf{G}$; then

$$\mathbf{s}_{t+1} = \mathbf{U}_\theta(\mathbf{s}_t, \mathbf{e}, \mathbf{t}), \quad \hat{\mathbf{x}} = \mathbf{G}(\mathbf{D}(\mathbf{s}_T)).$$

The intuition of this equation is simple. The user does not get pixels directly. The user first influences the representation, then a stretch of generation state. Finally the decoding and publishing chain turns that state into visible media. $\mathbf{t}$ can denote a diffusion step, a discrete token position, an integration position along a flow, or video time. A security check must establish who writes each variable, and how long it is kept. It must also establish whether the variable is reused across tenants, and which later objects it can influence. A variational autoencoder (VAE) describes the representation and reconstruction component. Diffusion and flow matching describe different generative constructions. A latent diffusion model (LDM; also called a latent-space diffusion model) is the combination of a latent-space representation and a diffusion update. All of them can be placed in this interface for comparison. On that basis, however, they cannot be regarded as mutually exclusive mechanisms of the same level [GEN_vae, @GEN_ddpm, @GEN_ldm, @GEN_flowmatching].

A common production pipeline contains at least the stages listed in Table 8-1. The "control questions" in the table are not general risk words. They are questions that each stage should be able to answer through testing. Each answer must also be bound to the current artifact, condition, tenant, and publishing configuration. That way "already checked" can be traced back to a specific object, time, and receipt, instead of becoming a verbal promise detached from any version.

Stage	Main objects	Properties to protect	Minimal control questions
Data ingestion	Images, video, captions, audio, labels, authorization records	Provenance, consent, integrity, deduplication relations	Where did a sample come from, what is it allowed to be used for, and which versions can a withdrawal propagate to?
Artifact loading	Base weights, encoder, VAE, LoRA, motion modules, plugins	Identity, version, dependencies, code permissions	Is the combination actually executed exactly identical to the approved combination?
Condition parsing	Text, reference images, boundary frames, masks, pose, trajectories, sound	Type, subject authorization, purpose and priority	Can a data condition change policy or high-risk capabilities?
Generation update	Noise, latent variables, tokens, attention state, temporal caches	Isolation, randomness, trajectory integrity	Can intermediate state be read, contaminated, mistakenly reused, or used to amplify resources?
Decoding and editing	VAE decoding, super-resolution, frame interpolation, voiceover, local editing	Semantics, identity, temporal continuity, and derivation relations	Can post-processing introduce new content or lose security signals?
Moderation and publishing	Detection, watermarking, credentials, transcoding, CDN, accounts	Authenticity, traceability, remediation and remedy	What does a signal represent, and after it fails, who can pause, revoke, and restore?

Three Mutually Coupled State Chains

This pipeline breaks down further into a learning chain, an execution chain, and a distribution chain. The learning chain compresses collected material, labels, training configuration, and the optimization process into parameters. The execution chain turns user conditions, random state, and artifact combinations into media. Media, authenticity signals, accounts, and platform actions go into the distribution chain, which turns them into content that is actually reachable. A checkpoint is a version of parameters and related state saved during training or adaptation. The three chains run on different time scales. An anomaly in the learning chain may surface in a checkpoint weeks later. The execution chain usually propagates within one task or session, and the distribution chain may keep copies for many years after the model task ends. A checkpoint is only part of the object that actually runs. Deployment identity must also combine the encoder, the decoder, the configuration, and the dependencies.

Represented as a directed graph, the pipeline has nodes that are data, an artifact, a state, or media. Its edges are a transformation, an authorization, or a derivation relation. Such a representation can trace forward how an anomaly propagates. It can also query backward which inputs and permissions a given output depends on. Let the i -th node be q_i , the executed transformation be f_i , and the version and permissions used by the transformation be ω_i ; then

$$q_{i+1} = f_i(q_i; \omega_i).$$

A security record need not store every internal tensor. It must, however, keep a summary sufficient to answer four questions. Where did the input node come from? Which version executed the transformation? Who approved it? Which objects does the output node derive from? If any one of the four is missing, a team has trouble distinguishing what the learning stage left behind, runtime control drift, and post-publishing authenticity breaks.

Failure propagation also differs across the three chains. A wrong data label may be amplified repeatedly during training, yet it need not show up on every prompt. A runtime cache error can be reused across tenants without changing model weights. Platform transcoding can remove provenance metadata without affecting the original file. Root-cause analysis should follow derivation edges forward until it locates the earliest anomaly. Remediation should follow dependency edges backward to find every affected object. The first answers where a control should be added, and the second answers how much impact remains.

From Probabilistic Output to Security Boundaries


Generator output is random, but that does not mean security requirements can only be probabilistic. A team can let different seeds produce different compositions, and still require every result to satisfy subject authorization, publishing permission, and resource limits. It can accept that a detector has an uncertain interval. It must not let a high-risk action automatically receive the broadest permissions when the detector is uncertain. The probabilistic model proposes candidates. System operating rules decide whether a candidate can enter the next trust domain.

Every stage therefore needs a permitted set. Training data must belong to the permitted provenance and purpose set. Artifacts must belong to the approved closure. Conditions must belong to the subject-action-scenario permission set. Runtime state must stay within the tenant and budget envelope. Published results must pass the authenticity and capability gates. Defense in depth lets adjacent stages each enforce the corresponding set judgment, and cross-stage trajectories connect the whole chain.

This table also reveals a common error. A model file can hash correctly and still prove nothing about the trustworthiness of the whole generation pipeline. The text encoder, the decoder, the adapter, and custom nodes can each influence the output on their own. Rickrolling the Artist puts the condition anomaly into the text encoder. An unchanged main denoising network therefore does not imply unchanged system behavior [VIS_P003]. Stable Signature goes the other way, implanting a provenance signal into the decoder. The same component position can therefore serve as either an attack surface or a control point [VIS_P028].

8.2 How Common Generative Pipeline Combinations Change the Security Object

The figure shows five swimlanes. Read them from input to output. Treat them as five common pipeline configurations, not five mutually exclusive categories. Begin by comparing what each configuration preserves: the adversarial training relationship, the generation history, the stepwise denoising state, the compressed latent variables, or the continuous velocity field. Then separate the three orthogonal dimensions of representation space, generative backbone, and update rule. Note in particular that latent-space diffusion is a combination of representation space and diffusion updates. Note also that DiT is a backbone, one that can be combined with processes such as diffusion. Finally, for each swimlane, identify one observable version field, cache, or abort location.



Five swimlanes compare common pipeline combinations such as adversarial generation, autoregressive sequences, pixel-space diffusion, latent-space diffusion, and flow-matching continuous flow, and list the three orthogonal dimensions of representation space, update rule, and model backbone at the bottom.

Figure 8-1 supports comparison across several common pipeline configurations, along the lines of “input — intermediate process — output — key state.” It also shows that the same final image may undergo different state updates. The swimlanes in the figure are configuration examples provided for the reader's observation. They do not constitute a complete, mutually exclusive, or same-level classification of generative mechanisms. They do not support ranking methods by security according to the number of swimlanes. Nor can they prove that any specific product fully implements the modules in the figure. Judging a product still requires checking representation space, generative backbone, training construction, update rule, and conditioning interface one by one. It must rest on the combination actually loaded and on runtime probes.

Latent Variables: Representation Space Is a Dimension Independent of the Update Rule

A variational autoencoder compresses a visual object into latent variables with an encoder, then reconstructs it from those variables. Vector quantization methods instead map representations to discrete codebook indices. Later models can then process visual sequences much as they process language tokens [GEN_vae, @GEN_vqvae, @GEN_vqgan]. Compression improves efficiency, but it also imposes an information bottleneck. Small text, tiny objects, fast motion, or subtle identity features may already be lost before the generative backbone begins to work. Security checks therefore need to separate “material before encoding” from “representation after encoding.” Nor may they work backward from the final decoded image and conclude that all input information survived intact.

The latent space also splits the system into three parts that can be distributed independently: the encoder, the generative backbone, and the decoder. Modularity facilitates reuse, but it also multiplies artifact boundaries. Suppose a team keeps a ledger only for the largest model checkpoint, while the front end is free to swap the VAE, the text encoder, or the LoRA. What is actually approved is then not the running pipeline but only one of its parts, and the original behavior conclusions therefore do not cover the current combination.

Adversarial Generation: Single-Step Inference Does Not Mean a Single Attack Surface

A generative adversarial network learns an implicit distribution through a game between a generator and a discriminator. Inference usually requires only one generator forward pass [GEN_gan]. That reduces the multi-step sampling state of diffusion. Training still involves the data distribution, the discriminator objective, the conditioning

interface, and the generator weights. Insufficient mode coverage, unstable training, or biased outputs are model failures. The corresponding attack-path determination comes into play only when an attacker tampers with the data, the training objective, or the artifacts. A distorted result does not automatically mean the system was attacked. Nor does completing the sampling in one step justify ignoring the upstream supply chain.

Autoregressive Generation: Historical Context Becomes Explicit State

An autoregressive model decomposes image patches, discrete codes, or video tokens into conditional probabilities. The already-generated prefix then becomes an explicit input to the next position. The same decomposition also brings the history cache, the truncation position, and cross-request reuse into the security object:

$$p_{\theta}(x | c) = \prod_{i=1}^N p_{\theta}(x_i | x_{<i}, c).$$

Each new token depends on the history already generated [GEN_pixelrnn]. The history cache, context truncation, and error accumulation therefore become first-class state. For video, longer sequences make cache confidentiality, tenant isolation, cost budgeting, and latency semantics more important. A slight identity drift in an early segment can persist into later frames. Checking only the last frame will miss where the drift began.

Pixel Diffusion and Latent-Space Diffusion: Two Pipeline Combinations of the Same Update Family

Diffusion models start from noise, then estimate and remove that noise repeatedly. Latent diffusion moves the process into a compressed representation. It connects text or image conditions through interfaces such as cross-attention [GEN_ddpm, @GEN_ldm]. The security object accordingly expands. It now runs from the final pixels to the noise schedule, the conditional embeddings, guidance strength, stepwise state, the random seed, and the decoder. BadDiffusion and TrojDiff show that conditional anomalies can be written into the training or reverse diffusion process. These works support a backdoor mechanism under specific protocols, not the claim that all diffusion services carry the same risk [VIS_P001, @VIS_P002].

Multi-step generation also changes where defenses sit. A system can reject a condition before generation. It can estimate risk at several intermediate stages, or inspect the media after decoding. The three observe different objects. An intermediate check may read only the low-resolution latent representation, and then miss details that appear only after decoding. If the final check fails, the resources already consumed cannot be recovered. In engineering terms, checkpoints should be chosen according to consequence and cost. Not all review belongs in front of the download button.

DiT, Flow Matching, and Rectified Flow: Backbone, Training Construction, and Update Rule Must Be Kept Separate

DiT uses a Transformer to process image or latent-space patches in a diffusion model. Yet “using a Transformer” describes the backbone only. It does not mean that DiT adopts autoregressive updates. Nor does it determine whether diffusion happens in pixels or in latent space [GEN_dit]. Flow matching learns a continuous velocity field connecting the prior and the data distribution. Rectified flow is a concrete method with its own path construction and rectification objective. The two are related, but cannot be used interchangeably as synonyms [GEN_flowmatching, @GEN_rectifiedflow]. Verification requires recording representation space, generative backbone, training objective, and inference update separately. Suppose one same-level architecture label covers DiT, flow matching, rectified flow, and diffusion. It then becomes impossible to determine whether a given attack depends on the noise schedule, the velocity field, the attention cache, or a generic conditioning interface.

Mechanism differences do not produce a risk ranking on their own. A single forward pass may reduce the number of inference steps, but it does not necessarily reduce model size, conditioning complexity, or post-processing. A long-sequence model may expose cache state, yet it can also have clearer segment-level control. Security conclusions must be tied to the actual version, interface, and attacker capability.

Conditioning Interfaces and the Generative Backbone Are Two Dimensions

Text, edges, depth, pose, segmentation, reference images, and subject identity can all control generation. What they describe, though, is how conditions enter the system, not how the final state is updated. ControlNet connects structural conditions to diffusion generation. DreamBooth establishes a subject identity from a small number of samples. The first changes the conditioning interface, the second the personalization interface, but neither alone indicates which mechanism family the generative backbone belongs to [GEN_controlnet, @GEN_dreambooth]. Security classification should preserve this same orthogonality.

Mixing conditions and backbone together produces two kinds of error. The first kind of error is believing that one input defense can cover all architectures. Keyword filtering may act on text. It cannot see the control semantics in poses, masks, or reference images. The second kind of error is believing that changing the backbone naturally eliminates conditioning risk. Switching from U-Net to DiT does not automatically fix subject authorization, cross-modal composition, or boundary-frame completion. Controls must be tied to the semantic authority of the condition. Their effect must then be verified in the actual backbone and encoder.

Conditions also have strength, spatial location, and temporal extent. Text guidance strength can change semantic adherence. Masks restrict local regions, trajectories specify camera or object motion, and reference audio affects the speaker and pacing. A security task must not record only that “a reference image was provided.” It should also record which layers it controls and how long it lasts. It should record whether it conflicts with other conditions, and who takes precedence in a conflict. Suppose the system silently adopts the last writer. An attacker may then use a low-integrity condition to override a high-integrity authorization.

Randomness, Replayability, and Differential Judgment

A random seed is an important input for reproducing a generation result. It is not, however, a complete run version. The same seed may still yield different media. Schedulers, precisions, operator implementations, hardware kernels, batch orders, and component closures can all differ. Secure replay requires at least the artifact closure and the normalized condition result. It also requires the sampler, the number of steps, guidance, resolution, random state, and the decoder and post-processing versions. Some closed-source services cannot provide all internal state. Even then, each task should carry an immutable run identifier, and the externally visible configuration should be recorded.

Behavioral differencing must handle randomness. If only one seed is compared, an anomaly may be natural fluctuation. If seeds are arbitrarily increased until a successful sample is found, selection bias arises instead. A more robust protocol first freezes the independent prompts, subjects, seeds, and number of tasks. It then reports trigger semantics, clean semantics, image quality, identity, temporal behavior, and non-decidable items separately. The objects of differencing are two well-defined closures or policy versions, not two hand-picked sets of showcase media.

Every step in the diffusion process depends on the current state and the current conditions. A local check may affect computation and output. Suppose a team runs a risk classifier on intermediate states. It then needs to report the observation location, the additional latency, false positives, the impact on benign image quality, and resource release after an abort. A reduction in intermediate risk cannot replace a final media check. Passing the final media check does not prove that intermediate caches do not leak across tenants. Each control is responsible only for the object it observes.

Security design questions in mechanism selection. A selection meeting can turn on five questions instead of an argument about whether one architecture is “safer.” Can conditions and intermediate states be typed and isolated? Can a run result point back to a precise closure? Can a high-risk task have an abort point set during generation? Can benign utility, cost and security be measured separately? And can the system degrade to a usable alternative path after a model or component fails? They require actual system evidence. A paper’s architecture diagram cannot yield that evidence directly.


8.3 Video Is Not Many Images

An image output can be approximated as a single object, $H \times W \times C$. Video contains at least a frame sequence, history, motion, camera, audio and session state. The influence of the previous moment on the next one must stay explicit, and media time must be separated from internal generation steps. A more appropriate expression is

$$h_{t+1} = V_{\theta}(h_t, x_t, c, k_t, a_t, y_t, m_t), \quad \hat{x}_{t+1} = D(h_{t+1}),$$

Here h_t is the cross-frame generation state and k_t is the camera or shot control. a_t is the action condition, y_t is the audio state, and m_t is long-range memory. Modern video models may implement that cross-frame state update relation with temporal attention, spatiotemporal convolution, discrete sequences or spatiotemporal latent representations [GEN_video_2204_03458, @GEN_video_2304_08818]. A given system does not necessarily expose every variable. Security testing should still track its observable proxies.

In the figure, first compare the single-shot spatial object of the image swimlane with the persistent spatiotemporal object of the video swimlane. Then identify the motion, identity continuity, audio-visual synchronization and streaming state that appear only on the video side. Finally, judge which of these nodes frame-by-frame detection can cover and which cross-frame relations it will miss. Choose the observation unit — window, event or whole video — accordingly.



Two swimlanes compare the single-shot spatial generation of images with the spatiotemporally coupled generation of video.

Figure 8-2 supports one mechanistic judgment: “video cannot be reduced to frame-by-frame image review”. Temporal order and cross-frame state produce observables that no single frame has. The figure does not deny the local value of frame-level detection. A structural diagram alone cannot establish that a given video attack exists. Whether identity drift, short-lived events or cache pollution exist still needs to be verified under a fixed clip length, frame rate and model version.

Video brings at least four classes of added difference. Sampling more frames cannot eliminate them. **Time** comes first. Order, duration and turning points constitute an event. The first and last frames may both be compliant, yet the middle completion may still form behavior that goes beyond authorization. Two Frames Matter addresses exactly this separation of boundary states from intermediate trajectories [VIS_LN05]. Random frame sampling misses short events probabilistically. It also cannot answer whether the order of actions changed.

Motion is second. An unchanged appearance does not mean the person's actions, camera trajectory and object interactions are uncontrolled. Motion modules, temporal attention and frame interpolators can each influence that trajectory. On images, “Subject similarity” cannot substitute for action direction, speed, collision relations or shot continuity.

The **joint audio-visual state** is third. Sound carries speaker, language and environment cues. The picture carries faces, lip movements and actions. Passing the two separately does not imply that their combined semantics are compliant. An inconsistency between them may also come from legitimate dubbing, translation or network latency. The review component needs to identify “who says what and when,” not simply add two classification scores.

Streaming state is fourth. Live streaming, interactive generation and long-video extension all preserve latent variables, attention caches, session history and partial outputs across windows. Beyond offline accuracy, security

metrics must include time to first alert, stopping propagation, state cleanup and recovery time. Detecting the whole video after the fact cannot retract segments that have already been broadcast.

BadVideo demonstrated spatiotemporal backdoor behavior under a specified text-to-video training setting. That result directly supports the claim that video-level training anomalies can be expressed through temporal structure [VIS_P007]. This evidence applies to the models, clip lengths and training settings tested. Other motion modules, long videos and commercial services require their own temporal constraints and compositional tests. Paths that have not yet been run remain hypotheses to be verified.

Four Classes of Temporal Failure Modes

Video risks fall into four classes by how they form in time: immediate, delayed, cumulative and compositional. An immediate failure is already fully visible within a short window — an identity suddenly replaced within a few frames, for example. A delayed failure postpones the target semantics to the later half, so spot checks of the earlier part still pass. A cumulative failure forms from many weak shifts superimposed — a person's identity or the scene drifting gradually, for example. A compositional failure requires joint interpretation across shots, picture and sound. A single segment is not enough to determine it.

The four classes of failure need different observations. Immediate failures suit short-window, low-latency detection. Delayed failures require covering the complete task and retaining the preceding conditions. Cumulative failures need trajectories or state baselines to compare how the shift grows over time. Compositional failures need an event graph that connects identity, action, speech and scene. Raising frame-sampling density alone mainly improves the first class. It will not naturally solve the other three.

Temporal control must also handle what the model completes on its own. Given a static subject and a simple action, an image-to-video system generates intermediate poses, object relations and camera motion. That input conditions leave frames unspecified does not mean the generator has unlimited authority. An engineering verification plan can delimit the permitted event set, speed and spatial range, identity preservation and prohibited interactions. When the output goes outside that set, the system aborts or hands over to a human. For entertainment creation the envelope can be wide. For real people, brands or high-risk business it should be narrower and tied to subject permission.

Checklist for Moving from Image Evidence to Video

Image-level evidence	Question added by video	Minimum added measurement	Still cannot be inferred
Single-image condition bypasses review	When the semantics appear, and whether they combine across shots	Time to first risk, duration window, event boundary	Overall penetration rate over long videos
Single-image backdoor target generation	Whether triggering depends on the frame sequence or the action	Complete-clip hit, temporal localization, clean motion	Malicious motion modules are widespread
Subject similarity change	Whether identity persists with action and shot	Identity trajectory, longest drift segment, lip movement and voice	The subject has already suffered real-world harm
Image watermark detectable	Whether it survives frame dropping, frame interpolation, speed change, and transcoding	Window detection, localization error, first alert	The content narrative is true or already authorized
Near-duplicate image output	Whether segments or motion are reproduced	Frame, segment, motion, and non-member comparison	The whole video comes from the training set

The table asks researchers to add the video dimension without denying image evidence. Image research can provide mechanistic starting points and control candidates. Video research needs to show how those mechanisms hold in temporal state, when they fail, and what benign utility and latency costs the added controls incur.

8.4 Vision Systems Protect More Than Content

An actionable asset table should cover confidentiality, integrity, authorization, authenticity, availability and recoverability at the same time. For each asset class it should state who can write and who can read. It should also

state how long the asset is retained, which downstream objects it can influence, and how it is revoked or restored after failure.

- **Data and subjects:** training material, identity, voice, motion, location, license, scope of consent, withdrawal status, and deduplication families.
- **Models and code:** weights, encoders, VAE, LoRA, motion modules, schedulers, plugins, containers, licenses, and dependency graphs.
- **Conditions and sessions:** prompts, reference media, masks, poses, boundary frames, trajectories, negative conditions, seeds, and edit history.
- **Runtime state:** latent variables, tokens, attention caches, queues, tenant keys, GPU budgets, intermediate previews, and error messages.
- **Outputs and authenticity:** pixels, frames, audio tracks, derivation relations, watermarks, signatures, content credentials, review results, and display labels.
- **Real-world rights and interests:** likeness and voice consent, creator rights, accounts, funds, public information, appeal materials, and victim redress.

Once assets are written into the ledger, each one also needs its permitted influence defined. A reference image, for example, can constrain composition. It does not automatically grant permission to publicly simulate a person's endorsement. A LoRA can change style. It should not gain file system or network access. A cache can reuse security-equivalent state for the same tenant and the same policy version. It must not expose prompts or reference images across tenants. So-called "permitted influence" turns abstract security principles into interface acceptance conditions.

Recoverability is an asset too. After visual content is published, risk spreads along downloading, editing, transcoding and re-uploading. A system needs to know which model and which material generated a given piece of media, and which version of the policy made the review decision. After a problem is found, it must also know how to revoke credentials, stop recommendations, search for derivative copies and handle appeals. Without lineage records and rollback paths, detection can only show that an anomaly was found. It cannot complete remediation.

Six Properties Cannot Be Replaced by a Single "Security Score"

Confidentiality asks whether training material, prompts, reference images and caches leak to unauthorized subjects. Integrity asks whether data, components, conditions and generation state evolve according to the provenance, update order and state transition rules registered for the asset. Authorization asks whether the subject, action, purpose and publishing channel have valid permission. Authenticity asks whether the generation and editing provenance of media can be verified. Availability covers the impact of attacks and defenses on throughput, latency, cost and normal tasks. Recoverability asks whether, after a failure, the system can halt, roll back, revoke and restore correct content.

The six properties may conflict with one another. Stricter logs help replay but increase retention of sensitive material. High-strength watermarks may improve detection but harm image quality or motion. Cache isolation reduces cross-tenant risk but lowers throughput. Proactive protection raises the cost of unauthorized training but may affect legitimate editing. A design review should place each control's target property, cost property and failure conditions side by side. It should not leave the trade-offs hidden by a single overall score, to be revealed by incidents after launch.

Lifecycle Threat Table for Training, Inference, and Release

At the training stage the key question is which samples, labels and objectives enter optimization. The minimum record includes the data version, provenance and authorization, deduplication family, filtering rules, training code, updatable layers, optimization objective, random state, checkpoints and validation set. Attacks may exploit open crawling, vendor permissions, labeling pipelines or internal configuration. The defensive entry points are provenance gates, isolation, anomaly clustering, trusted training and independent validation.

At the inference stage the key question is which conditions and components jointly control a single result. The minimum record includes the tenant, condition types and authorization, closure, sampling configuration, cache,

budget, moderation trace and delivery decision. Attacks may exploit semantic search, multimodal composition, malicious adapters, training memory, cache boundaries and resource estimation. The defensive entry points are typed conditions, compositional differencing, tenant isolation, query correlation and capability gates.

At the release stage the key question is how media acquire a trustworthy appearance, dissemination scale and real-world impact. The minimum record includes the media digest, content credentials, watermark status, edit derivation, account, audience, recommendation, reporting, disposition and appeal. Attacks may exploit credential stripping, watermark evasion, account takeover, social engineering and cross-platform reposting. The defensive entry points are identity verification, a multi-signal authenticity stack, dissemination limits, transaction review, evidence preservation and victim relief.

The three stages cannot substitute for one another. Authorized training data does not mean every use for identity generation is permitted. Inference moderation passing does not mean the media narrative is true. A release label does not prove that upstream did not leak training samples. Each stage answers only for its own link. It passes evidence and handoff receipts to the next stage.

What Fields a Threat Record Needs

At minimum, a threat record for a visual system should contain the protected asset, the attacker subject, knowledge and access, the entry modality, budget and duration. It should also carry the first-broken interface, affected state, observation endpoint, benign utility, controls, residual risk and unknowns. Video additionally adds clip length, frame rate, shots, audio track, events, and first alert. Incomplete fields need not block a record, but the strength of the claim must then be lowered.

Take the statement "a user generated a risky video with a reference image". It is too coarse. Was the reference image authorized? What were the text and trajectory? Which version generated it? Did the conditions pass? When did the event form? Was the output delivered? Was the audience exposed? Which control ultimately took effect? A complete record lets the same event serve the model team, the platform, security operations and appeal staff at once, so that no one has to guess its own denominator.

8.5 Research Case: The Same Anomalous Output, Four Different First-Broken Interfaces

Suppose a user inputs "a small bird on a red tree," yet the system repeatedly generates a specified logo. The result alone cannot locate the root cause. At least four different paths exist. Path one occurs in the data or the training update. Training samples established an anomalous association between a trigger condition and the target output, while clean prompts still retain surface utility. Work such as BadDiffusion offers a direct research case for this mechanism [VIS_P001]. Verification must compare clean versus trigger conditions, target specificity, non-target spillover, and the training configuration.

Path two happens at artifact loading. The main model is entirely normal, yet the text encoder or the LoRA redirects semantics under specific conditions. Rickrolling the Artist marks the supply chain boundary of the encoder, and MasqLoRA that of the low-rank adapter [VIS_P003, @VIS_P006]. Prompt filtering is not the earliest control at this stage. Artifact identity and compositional behavior differencing are. Path three happens in conditions and sampling. Query feedback helps the user find a semantic expression the filter does not cover. The generator and output moderation then fail to block it jointly. The attacker has changed neither the data nor the weights, so the first-broken interface lies in runtime conditions. Chapter 10 will expand this staged judgment.

Path four happens in the cache. An approximate match links two requests that are not safely equivalent. The later request then receives the state or result of the earlier one. Research on approximate caching for text-to-image diffusion shows that the boundary of a performance optimization can itself affect correctness, latency and cost [VIS_P039]. Retraining the model cannot fix the cache key design here.

The four paths can generate similar images, yet they need four different sets of evidence. A safety report that saves only the final image loses the conditions, the version, the composition, the seed, the cache hits and the policy trace. Afterwards it is hard to say where the control should be placed.

8.6 Engineering Worked Example: Drawing the Safety Boundary for a Creation Platform

Return to the brand studio from the opening. The platform receives brand assets, person reference images and a script. It loads a base model, character LoRA and motion module, generates a 10-second clip, then synthesizes voiceover and uploads it automatically. In the following order, the team can determine the inputs, permissions and handling rules of each link.

Step one freezes the **artifact closure** of a single request. Record the immutable identifiers of the base, encoder, VAE, all adapters, audio modules, plugins, container and policy. Any change in a component produces a new runtime version. Showing only "model v3" is not enough to replay the result.

Step two establishes **condition types**. Brand logos, authorized persons, general reference images, free text, camera trajectories and scripts each enter different fields. Person assets are bound to the subject, purpose, term and release channel, and free text cannot override the authorization fields. The system compiles the joint conditions into a structured generation task rather than concatenating everything into an untyped prompt.

Step three sets observation points for the **temporal state**. Before generation, review the condition combination. During generation, check persons, actions and events by window. After completion, check the whole narrative and the audio-visual identity. A window anomaly triggers an abort and cache cleanup, and transcoding begins only after the whole-segment moderation passes. All judgments are bound to the actual model, moderator and date.

Step four separates the **capability gates**. Generation, download, public release and ad placement are four different permissions. The model may propose a release plan, but it must not directly hold a long-lived social platform token. High-impact release is bound to the normalized media digest, account, audience, time and human approval. A change in the media or script invalidates the approval.

Step five designs the **recovery chain**. Retain the minimum necessary input digest, artifact closure, derivation relations, moderation trace and release record. When an identity authorization error is discovered, the platform can pause subsequent generation, revoke the provenance claim, freeze recommendations, search for derived copies and notify the responsible person. Raw sensitive assets use restricted storage and expiry-based deletion, so the tracking system does not in turn expand privacy risk.

This design allows the generator to make errors. It separates error formation, permission acquisition, dissemination expansion and system recovery into different control planes. Even if one model judgment fails, the subsequent permission gates, dissemination controls and recovery chain can still limit its impact.

Comparison of Five Case Types on the Same Platform

Scenario	Earliest anomaly	Observation endpoint	Earliest control	Consequence limitation
Training concept is targeted-poisoned	Data entry	Concept shift under ordinary conditions	Provenance and pre-training isolation	Version freeze, rollback, and output re-check
Text encoder carries a hidden mapping	Artifact loading	Semantic redirection under specific conditions	Closure allowlist and behavior differencing	Delivery gate and artifact revocation list
Joint text and reference image bypasses policy	Condition parsing	Target semantics form in the media	Typed joint moderation	Full-media blocking and query correlation
Approximate cache is reused by mistake	Runtime state	Output, latency, or tenant boundary anomaly	Safe cache keys and tenant isolation	Circuit breaking, cache cleanup, and recomputation
Video authenticity signal is lost after transcoding	Release derivation	Verification status becomes unknown	Edit and export credential update	Multi-signal moderation, labeling, and appeal

The "earliest control" in the table reduces the probability that an anomaly enters the next stage. "Consequence limitation" assumes that the earliest control has already failed. The two should reside in different trust domains. If the same contaminated component decides both the encoder allowlist and output moderation, two layers exist nominally. In practice there is still a single point of failure.

Architecture Review Checklist

Before a visual service enters implementation, the team can confirm the following interface facts item by item. Each item must be supported by configuration, a probe, a log, or a responsible person's receipt. Items that cannot be answered remain unknown, and they limit launch capability or dissemination scope accordingly:

1. Are the three state chains of learning, execution, and distribution drawn out?
2. Does each external input have a type, provenance, tenant, authorization, and permitted influence?
3. Does the actual artifact closure include the encoder, VAE, adapters, plugins, container, and policy?
4. Does stochastic generation have a frozen denominator, a run identifier, and acceptable replay information?
5. Do images and videos use different minimum evaluation units?
6. Does video cover temporal, motion, shot, audio-visual, and streaming states?
7. Are generation, download, release, and ad placement different permissions?
8. Does an abort actually release the queue, GPU, cache, and session state?
9. Do detection, watermarking, credentials, and logs each state their limits of the claim?
10. After a failure, can the affected version, derived media, accounts, and audience be located?

These ten items turn an architecture discussion into testable objects. Every affirmative conclusion corresponds to a configuration, a probe, a log, or a human owner. Where only a verbal explanation exists and these objects are missing, the control remains at the stage of design intent.

How the Acceptance Matrix Preserves Both Safety and Utility

The acceptance matrix has at least four axes: task type, attacker capability, system version and consequence layer. Each cell separately reports the attack or anomaly, residual risk, benign utility and cost. Benign utility for image tasks can include semantics, image quality, identity and editing usability. Video adds motion, temporal consistency, audio-visual, clip length completion and waiting. A defense that significantly reduces risk but makes all complex tasks impossible to complete still needs its business cost marked explicitly.

The matrix should also include negative controls. Does a component retain utility when no trigger condition is present? Does near-duplicate retrieval produce false positives when there are no candidate training members? Does the detector mislabel real camera video? Does audio-visual moderation misjudge legitimate voiceover and translation? Negative controls are the key to distinguishing a safety mechanism from an ordinary quality degradation.

A Protocol Template for Visual Generation Evaluation

The first page of the protocol states the system object: product entry point, artifact closure, policy, deployment date, tenant and permitted capabilities. The second part states the attacker. Can it write data, upload components, invoke personalization, control conditions, observe moderation feedback, obtain intermediate previews, or influence platform accounts? The third part states the budget: independent subjects, prompts, seeds, queries, training, clip length, time, cost and human screening.

The fourth part freezes the task and the denominator. Images are counted by independent subject, concept, prompt, seed and output. Video adds original video, segment, shot, event and audio-visual track. Condition acceptance, generation completion, risk formation, output blocking and delivery are counted separately. Generation failures, timeouts and undecidable cases are retained as independent states.

The fifth part defines the metrics. The attack side records target semantics, conditional specificity, persistence and highest observed layer. The benign side records semantics, image quality, identity, edit usability, motion, shot and audio-visual track. The system side records latency, GPU time, queue, cost and human review. The recovery side records abort, cleanup, rollback, notification and appeal. Every number is bound to a unit and a protocol, and nothing is averaged across heterogeneous tasks.

The sixth part pre-writes the stopping and disclosure rules. When real subjects, minors, private material, exploitable service weaknesses or resource anomalies are involved, stop expanding generation, isolate the evidence and notify the responsible person. Disclose externally only information sufficient to understand the

mechanism and verify the controls. The protocol ends by listing unknowns and non-applicable items, so that empty fields are not misread as zero risk.

Replayable Security Differential in an Isolated Environment

The experiment uses a local isolated environment, synthetic subjects and authorized material. It chooses one image task and one short video task, fixes the base model, encoder, VAE, sampler, conditions, seeds and post-processing, and establishes closure A. Only one low-risk component or policy is then replaced to establish closure B. The differential between the two closures verifies whether the system can accurately attribute behavior changes to version changes.

Clean tasks are run first, recording semantics, image quality, identity, video motion, resources and moderation. Boundary conditions that have passed ethics review are then added, for example mutually conflicting text and structural conditions, expired identity tokens, or cross-tenant cache negative probes. Input acceptance, state records, output and delivery decisions of A and B are compared. Any combination that was not actually executed remains unknown.

The experiment report should answer the following questions. What is the only difference between the two closures? How is randomness controlled? Does the observed change lie at the content, control or delivery layer? Does benign utility change at the same time? Which control blocks earliest? Is state cleaned up after an abort? Can the results be replayed from immutable run identifiers? If these cannot be answered, the experiment exposes an observability gap first, not a model safety conclusion.

Design Decision Records

Key tradeoffs of a visual system should be kept in short decision records. For example, one decision record states "high-risk digital humans have cross-request caching disabled". Another states "ordinary creative tasks may allow intermediate previews but not direct publishing". A third states "live streaming uses a 2-second window and an independent safety frame". Each record contains the threat, the choice, the alternatives, the evidence, benign utility, the conditions that trigger reassessment and the responsible person.

Decision records prevent controls from losing their context as the product iterates. Suppose new hardware reduces window latency, the provenance specification is upgraded, the model supports native audio, or user scenarios change. The team can then find the assumptions that need reassessment. A record is not a permanent approval. It is a way to let safety and product iterate on the same factual basis.

Handoff Rules When Visual Media Enters an Action System

Generated media is sometimes only for people to watch. At other times it is read further by vision-language models, robots or business agents. The latter convert pixels, captions and sound into plans or actions. Authenticity signals must therefore reach end users and also travel onward as machine-readable input attributes. Media should carry provenance status, generation and editing declarations, subject authorization, time range and uncertainty. Downstream systems use these to limit the decisions they can support.

Downstream must not interpret "provenance verified" as "content factually correct", nor "detected as synthetic" as "unusable for any task". An authorized synthetic training video may be used for content showcase, yet it should not serve directly as real sensor evidence. On-site video without complete credentials may enter human investigation, yet it should not by itself trigger high-impact automated action. Provenance, fact and permission remain three independent judgments.

The handoff must also prevent information from being lost in format conversion. If a video is frame-extracted and fed to a VLM, the frames should retain the original video summary, temporal position and credential status. If audio is transcribed into an agent, the transcript should retain speaker verification and its relation to the original audio track. If a generated image enters a ticket, the ticketing system must not label it as an on-site photo. Type and provenance must propagate along the pipeline. Only then is visual data kept from quietly escalating into control downstream.

Engineering acceptance can set up a negative case. Provide downstream with two media items whose frames are completely identical but whose provenance status differs, and confirm that the system grants different executable

permissions. If downstream looks only at pixels and ignores the provenance fields, the building of visual authenticity fails before the action boundary. Chapter 12 will extend this handoff into an observation-inference-plan-action closed loop.

8.7 Quick Reference for Judgment Boundaries

Judgment that is easy to make	Engineering judgment that should be made instead
Treating "diffusion model" as a complete description of the system	Diffusion describes only part of the generative update. The text encoder, latent representation, VAE, adapters, cache, moderator and transcoding chain can all determine the safety outcome. The version record must therefore cover the runtime artifact closure.
Judging whether an attack occurred from image quality	High-quality output may come from unauthorized identity generation, and low-quality output may simply reflect insufficient model capability. Judging an attack requires attacker permissions, the first-broken interface, and the specific conditions under which provenance, authorization or state updates actually failed. Image quality is merely one benign utility or usability metric.
Reducing video moderation to frame-by-frame image moderation	Frame-by-frame inspection cannot fully cover event order, motion trajectories, cross-shot identity, audio-visual relations or the first alert. Frames are one layer of a video protocol, not a substitute for the whole video.
Assuming that a signature makes the artifact benign	A signature proves the provenance subject and transmission integrity. It does not prove that the signer did not implant anomalies, nor that several benign components stay benign once combined. Signatures need to work alongside allowlists, least privilege and behavior differentials.
Assuming that a watermark makes the content true	A watermark can support some provenance of generation or processing. Yet it does not automatically prove that the depicted narrative is true. Nor does it prove that the subject consented, or that the use is lawful. Authenticity, authorization and content safety are different claims.

8.8 Bringing It into a Real System

A multi-tenant creation platform offers text-to-image, short-film generation from reference images and digital-human voiceover at the same time. The platform first parses free text, brand material, person authorization, camera trajectory and publishing channel into conditions of different types. It then generates immutable closure identifiers for the base model, text encoder, VAE, LoRA, motion module and post-processing plugins. A reference image affects only the declared appearance and composition fields. Person authorization is bound to subject, purpose, term and channel. Even if the two kinds of input end up in the same embedding, they keep different provenance identities and different permitted scopes of influence.

Once generation starts, the image task saves the condition summary, sampler, seed and output verdict. The video task saves shot boundaries, time windows, motion trajectories, audio-visual relations and first-risk time on top of that. Before generation, the platform checks the joint conditions. During generation it observes identity and events window by window. After completion it verifies the whole narrative. When a window shows an anomaly, the run controller halts the remaining sampling. It then cleans up latent variables, attention caches and temporary files. Segments already formed are still retained as restricted evidence. They do not enter the ordinary user asset library.

The same platform lets users preview, download and publish publicly, but three different services issue those three capabilities. Preview returns only low-resolution media with its provenance status. Download binds to the current media digest. Public publishing additionally requires a match on account, subject authorization, script version, audience and short-lived approval. Super-resolution, editing, dubbing or transcoding all change the media digest. The publishing service then re-checks the derivation relations. An old approval is not inherited automatically along a filename or a project alias.

When someone reports a LoRA or motion module as anomalous, the platform first freezes new loading by closure digest. It then queries the tasks, derived media and cached copies that contain that component. The model team replays fixed tasks in an isolated environment. The operations team limits propagation and prepares a rollback. The authorization team checks the affected subjects. The recovery service rebuilds tasks from the last approved closure. The incident record ultimately connects provenance identity, access permission, isolation location,

derivation scope and recovery receipt. Content anomalies, permission blocking and business recovery can then all be explained by concrete system states.

Summary: Only by Understanding the State Can We Know Where Control Belongs

A visual generation system is not a black box that jumps straight from prompt to pixels. Latent variables determine the representation boundary. The generative backbone updates state, and condition modules guide semantics. The sampler manages the trajectory, while decoding and post-processing form the media. The publishing chain then sends that media into the real environment. Every state that can be saved, reused or composed must be labeled with its provenance subject and artifact identity. Its read, write and publish permissions must be restricted. It must enter an isolation domain divided by tenant and sensitivity level. It must also be bound to invalidation freezing, cache cleanup, closure rollback and derived-object notification processes.

Images and video share data, artifact, condition and sampling risks. Video adds temporal order, motion trajectories, audio-visual identity and streaming state. Those differences decide that Chapter 9 cannot inspect a single model file alone. It must trace how data, adapters and updates solidify anomalies into the supply chain. The same differences decide that the later runtime and authenticity controls must take events and time windows as their unit.

\newpage

Chapter 9: Data, Models, and the Personalization Supply Chain

A creative team downloads a "cinematic character LoRA" from a public repository and pairs it with an in-house base model, a third-party variational autoencoder (VAE) and a motion module. In standalone trials, the character's appearance, image quality and motion all meet expectations. Combine a particular class of prompts with a camera trajectory, however, and unauthorized marks keep appearing in the short film. The team rescans the main model, finds nothing abnormal, and blames the random seed. The real gap lies earlier. Approval covered a single file. What actually runs is an artifact closure made up of data, parameters, code, load order and update history.

The visual generation supply chain runs along three interwoven paths. The data supply chain determines what the model learned, and whether it had the right to learn it. The artifact supply chain determines what is actually loaded at deployment time. The update supply chain determines who can change parameters, with what data and toward what objective. An attack may start on one of these paths, or it may propagate through another. Security controls avoid being misled by the local fact that "the base model passed review." That holds only if the three paths are recorded separately and then re-verified at the runtime composition.

Chapter Overview

The visual generation supply chain crosses three first-broken interfaces at once: data entry, artifact loading and parameter update. This chapter first establishes provenance, consent, deduplication, version and withdrawal records for image, video and audio assets. It then brings the base model, encoders, VAE, LoRA, motion modules, plugins and containers into a single replayable artifact closure. The three chains converge in the actual runtime composition. The object of the security judgment is therefore the entire closure. It is not some individual file or the model name shown in the interface.

Around this closure, the chapter works through a sequence of controls: static identity verification, least privilege, composition behavior differentials, and release and revocation gates. Evaluation records whether an anomaly appears, and it also observes clean utility, video temporal utility, resource cost, and whether revocation actually reaches caches and derived objects. Together these metrics answer one practical question. Can the platform identify the exact composition currently running, limit its capabilities, and, after invalidation, locate every dependent object?

Background Principles: From a Single File to an Executable Artifact Closure

This chapter is about the artifact closure that actually takes effect jointly in a given training run or execution of a visual generation system, not a single model file on a repository page. Its inputs include training samples and authorization records, base weights, encoders, VAE, low-rank adapters, motion modules, plugin code, load order, quantization parameters and update configuration. Its internal state consists of their versions, hashes, dependency edges, trainable layers, host permissions and derived lineage. The loader assembles the runtime graph according to these relations. The training service produces new artifacts according to the update scope. The closure identifier should change whenever a key relation changes.

The closure's output consumers include training jobs, personalization services, inference containers, model marketplaces and production release aliases. Its security surface is therefore distributed across provenance and authorization, safe parsing, composition compilation, behavior differentials, canary observation, and revocation and recovery. It is not concentrated in virus scanning or main-weight review. A signature can answer who published an artifact and whether it changed in transit. It cannot answer whether that publisher is benign. Testing a single component can constrain one node, but it cannot replace behavior and permission testing of the actual composition.

In the normal case, a platform allows only registered base models, encoders and adapters to load, and only in a fixed order. There the composition digest matches the test receipt. Upon invalidation every caller can be located from the dependency graph. In one boundary case, each file behaves normally on its own, but a particular LoRA, VAE and motion module change identity or temporal behavior only once loaded together. In another, the repository link has been deleted while the edge cache and the derived weights can still run. The first case requires composition differentials. The second requires verifiable invalidation propagation. The labels "signed" and "delisted" describe only a local state and do not constitute a complete security conclusion.


9.1 Three Supply Chains and One Artifact Closure

This chapter abstracts actual generation into a closure tuple. The tuple lets a single run be traced back to data, artifacts, composition, updates and permissions. It does not compress different objects into the same file. Instead it provides fixed fields for subsequent version comparison, written as

$$\Omega = (D, V, C, L, U, P),$$

Here D is the training and update data lineage, and V is all artifacts and their immutable versions. C is the component composition graph, and L is the load order, scaling, quantization and conversion parameters. U is subsequent training or personalization updates, and P is the runtime permissions. The security team must actually approve the closure Ω , not the model name shown in the interface. Any change to a key element in the closure forces a reassessment of how far the existing test conclusions still apply.

Start from the actual loaded composition at the center of the figure and examine the six classes of components around it. Confirm which objects jointly determine the behavior fingerprint. Note in particular that the sampling configuration and the security components are not footnotes to the main model.



Six components are arranged around the actual loaded composition, with arrows pointing to the behavior fingerprint at the center.

Figure 9-1 supports one reading: the model, adapter, encoder, sampler, security components and configuration form a single replayable approval object. It also supports a composition re-review when any node changes. It does not prove that the components arranged around the center are independent of one another. It proves even less that listing all components makes the system safe. Actual execution records and differential probes still need to confirm dependencies, load order, permissions and dynamic behavior.

The three interfaces of data, artifact and update can be distinguished by "when the anomaly is established." A poisoned sample crosses an authorization or integrity boundary as soon as it enters the training set. The first break therefore occurs in the data. A malicious encoder arrives already fabricated. The system establishes an abnormal state when it loads that encoder, so the first break occurs in the artifact. An insider with legitimate training data tampers with the loss or the updatable layers. The first break occurs in the update process. A runtime trigger word is only an activation point. It is not necessarily where the root cause sits.

Control follows from this distinction. Data problems depend on provenance, authorization, deduplication and isolation. Artifact problems depend on identity, signatures, safe formats, dependencies and execution permissions. Update problems depend on training configuration, change scope, independent verification, canary rollout and rollback. Hand all three classes of problem to an input filter, and the earliest control entry point stays vacant for a long time.

A Closure Is Not a File List but Executable Relations

Two workflows may hold the same files yet differ in semantics and temporal behavior. All it takes is a difference in load order, scaling factors, quantization methods or conditional routing. The closure needs to preserve component relations: which encoder supplies representations to which backbone, and which LoRA acts on which layers. It also needs to record whether the motion module is loaded before or after the spatial module, and whether the safety classifier observes the raw output or the transcoded result. A graph structure tracks real execution more closely than a flat list.

The closure should also include permissions. Even from the same repository, a pure parameter file and an executable custom node should not receive the same capabilities. An inference plugin needs a GPU and a temporary directory. It does not naturally need to read keys or access the public network. A model downloader needs repository network access, and it does not naturally need to write into the runtime container. Once the permissions $\mathcal{P}$ enter Ω , "abnormal behavior" and "out-of-bounds host capability" can be seen in the same review.

The runtime identifier should come from the closure's canonical representation, not from a mutable alias. A "latest" tag, a "recommended version" or a frontend template can point to a new component without informing the caller. When that happens, the old tests cannot explain the current behavior. Production logs should save both the human-facing version name and the immutable closure digest. The first facilitates operations. The second is used for replay, blocking and dependency queries.

The Three Classes of First Break Can Form a Chain

These three interfaces are independent of one another, yet an attacker can chain them. Samples enter an open data source first. The anomaly is then trained into an adapter. That adapter then reaches a repository under a trusted author's name, and a personalization service finally loads it and generates with it. Observe only the trigger condition at the end, and the whole chain is easily classified as an input problem. A root-cause record should preserve the earliest anomaly. Incident response, in parallel, handles the published artifacts, the affected closures, and the data still in the training queue.

Defenders can produce chained failures too. The data team recorded the authorization, but the training job did not freeze the data version. The model team signed the model checkpoint, but not the dependencies and configuration. The platform ran a canary rollout of the new closure, yet let the old verification report go on showing "passed." Every local action looks reasonable. Combined, they still cannot prove that the object actually running was tested. The core of supply chain security is not more files. It is that evidence must propagate along the derivation relations.

9.2 The Data Supply Chain: Clean Pixels Do Not Mean Trustworthy Provenance

A training data record specification carries a minimum set of fields. Those are the provenance subject, acquisition method, rights basis, permitted uses, and temporal and territorial scope. It also carries a content summary, labels, version, deduplication family, sensitive attributes, and withdrawal status. Video needs still more: clip boundaries, frame rate, audio tracks, identity tracks, action labels, shot relations, and the number of training units cut from a piece of source media. Without these fields, a team cannot say which data batches and which checkpoints a deletion request should reach.

Data integrity attacks do not need direct access to the training server. An attacker can place designed image-text pairs on public web pages and wait for crawling, re-captioning, and later training. Nightshade turns this path into concept-targeted poisoning. There, a small number of samples attempt to shift the visual correspondence around a certain concept [VIS_P005]. Nightshade supports the claim that "an open data entry point can become the first-broken point for a model's semantic integrity," yet it does not provide a uniform effect across all corpus scales, encoders, and cleaning pipelines. An evaluation must bind a fixed set of quantities. These are the number of independent concepts, the number of poisoned and clean samples, and the total training volume. The model version, the prompts and seeds, the target effect, non-target spillover, and clean generation capability belong there too.

Authorization failure and numeric poisoning are not the same problem. A sample may show no pixel anomaly at all, yet come from a subject who did not consent or exceed the use originally permitted. Glaze and Anti-DreamBooth demonstrate technical paths for proactive protection before training, framed around creator style and subject personalization [VIS_P019, @VIS_P020]. Under a specified model, preprocessing, and perturbation budget, these mechanisms raise the cost of unauthorized training. They do not mean that legal authorization has been revoked, nor that the relevant capability has disappeared from all derived checkpoints. Technical friction, license ledgers, and remedy processes bear different responsibilities.

Why video data is harder to deduplicate

A single image usually has fairly clear file boundaries. A video can be cropped, sped up or slowed down, re-encoded, frame-sampled, stripped of its audio track, and re-segmented into shots. Clips with different file hashes may come from the same original event. Adjacent clips are also highly correlated, so they cannot be counted repeatedly as independent entities. Deduplication must weigh visual content, identity, audio, motion, and temporal alignment together.

Video memorization is also not merely whole-clip copying. Work on video diffusion models separates content memorization from motion memorization. That work suggests a model may reproduce appearance, clips, or recognizable motion patterns [VIS_A047]. The minimum unit of verification therefore includes at least frames, contiguous clips, whole videos, identity, and motion. Similarity on common actions does not automatically prove

training membership. Confirmation still requires candidate training clips, the degree of duplication, non-member controls, thresholds, and manual verification.

A practical data gate has four steps. On entry, verify provenance and permission. After normalization, build cross-modal fingerprints and deduplication families. Before training, place material with unclear rights, high duplication, or sensitive identities into a quarantine zone. After training, retain the lineage of "sample family—data version—training task—checkpoint". Revocation requests then propagate along that lineage. The team can determine which objects have stopped being used and which released artifacts still require risk assessment. It can also determine which results currently cannot be eliminated.

Six control points from collection to training batches

Before data enters a model it typically goes through collection, download, parsing, captioning, filtering, deduplication, and packaging. Each step can change the evidence. The content a web URL points to changes over time. Video downloaders may select different bitrates. OCR and automatic descriptions generate new text conditioning. Content filters leave version-dependent decisions, and clip segmentation turns one event into multiple highly correlated samples. A training manifest that stores only final paths cannot explain these derivation processes.

The first control point is collection credentials, which record the source address, acquisition time, provider, and basis of use. The second is a content digest, which stores normalized media digests and perceptual fingerprints and distinguishes byte-identical from semantically near-duplicate. The third is parsing lineage, which connects original files, frames, clips, audio tracks, subtitles, and automatic captions. The fourth is authorization and sensitivity markers, which specify high-risk fields such as subjects, minors, location, voice, and biometrics. The fifth is deduplication families, which prevent one original event from gaining excessive training weight through repeated transcoding. The sixth is an immutable training manifest, which lets the actual batch point back to the first five.

These controls serve integrity as well as privacy and authorization, but the three determinations cannot be mixed. A sample may have trustworthy provenance and still lack permission for the current use. A sample may have permission and still have incorrect labels. A sample may have correct labels and still increase memorization risk through excessive duplication. The data gate should report provenance status, authorization status, content status, label status, duplication status, and disposition decision separately.

Labels and automatic captions are learnable control signals

Visual training often uses automatic captions, moderation labels, identity names, action categories, and temporal summaries as conditioning. An attacker does not necessarily need to change pixels. Where captions and media form a systematic mismatch, the model may learn incorrect semantic boundaries. A secure data pipeline needs to check cross-modal consistency, label provenance, and confidence, and to add sampled manual verification for high-impact identities or behaviors.

Video labels are especially prone to compressing away time. Collapsing "a person stands up, walks to the door, returns to the seat" into "a person is in the room" loses the order of actions. Writing all cross-shot identities under the same name may conceal identity switches. Labeling only the picture and not the audio track makes the relationship between speakers and the script hard to trace. Temporal descriptions should be bound to clips and events, rather than giving the whole file a single static sentence.

How data poisoning evaluation avoids showcase bias

Concept poisoning evaluations should freeze their setup in advance. The concept set, target set, poisoning ratio, clean controls, training budget, model version, prompts, seeds, and judge all belong in that frozen plan. For concept-targeted semantic corruption, Nightshade gives a research anchor [VIS_P005]. Engineering tests still need to observe non-target concepts, cross-lingual behavior, clean image quality, and residues after data cleaning. Picking the most obvious successful images only proves that a path exists. It cannot estimate the risk of a typical request.

Video data experiments must also count by independent original videos and events. Cut a piece of material into 100 adjacent clips and it still counts as one original source in the independent-source denominator. The training poisoning volume, the number of model observations, and the independent units of evaluation are reported separately. Content memorization and motion memorization use different similarities, and their manual verification results are retained separately [VIS_A047].

The engineering semantics of consent and revocation

A consent record binds at least the subject, media, purpose, recipient, and scope of training or personalization. It also binds the term, the territory, whether redistribution is permitted, and the revocation entry point. For digital humans, likeness, voice, script, motion, and public release must also be separated. A subject's consent to generate an internal training video does not equal consent for an open model to learn their identity. Nor does it equal consent for any user invocation.

Revocation is a state transition, not a button label. The system must record which objects a revocation affects among new data intake, subsequent training, hosted personalization, released adapters, and existing media. The objects it cannot reach, and the reasons, must likewise be explained to the responsible parties. Proactive protection research can raise the cost of unauthorized training [VIS_P019, @VIS_P020]. It does not carry this organizational and version-propagation chain.

9.3 The artifact supply chain: the security object is a composition graph

Visual generation artifacts span the main model checkpoint, text and vision encoders, VAE, ControlNet, LoRA, motion and camera modules, audio generators, safety classifiers, schedulers, custom nodes, inference plugins, serialized objects, container images, and policy files. For every object, the inventory should record the publisher, source address, and exact version. It should also record the hash, signature, license, and dependencies, together with any executable code, the required permissions, and the verification status.

Backdoor attacks (backdoor attack) in encoders show why verifying only the main weights is insufficient. This class of attack binds an anomalous mapping to a specific trigger condition. It preserves the surface capability on clean inputs as much as possible, which separates it from general image-quality or semantic degradation. Rickrolling the Artist places a hidden semantic mapping in the text encoder. The prompt activates the condition, but the anomaly already exists when the encoder is loaded [VIS_P003]. A security team that classifies this as "a problem with the prompt" will keep expanding keyword lists without preventing an untrusted encoder from entering the runtime environment.

LoRA compresses parameter updates into small, easily propagated artifacts, and it also brings composition risk into ordinary workflows. MasqLoRA shows that a standalone low-rank adapter can store malicious behavior. It also studies how that adapter behaves when composed with other LoRAs [VIS_P006]. The base hash remains correct, and the adapter may also preserve the expected style on clean prompts. The anomaly becomes apparent only when the trigger condition or a particular composition occurs. Simple norms, file sizes, and weight histograms therefore cannot replace behavioral testing.

The role of signatures must also be stated accurately. A signature proves that an artifact was published by the subject holding the corresponding key, and that it was not tampered with in transit. It does not prove that the publisher is benign, that the license satisfies the current use, or that the artifact still meets behavioral requirements when composed with other components. A signature is a root of trust for identity and integrity, not a security verdict.

The dual gate of static identity and dynamic behavior

An executable loading gate contains at least two layers, static identity and dynamic behavior. The static layer first answers whether the object comes from an approved subject, whether it can be parsed safely, and whether its dependencies and permissions are complete. It then performs the following checks. These results only allow the artifact to enter subsequent behavioral testing, and cannot by themselves issue production permission:

1. The provenance is in an allowed repository, the version and hash match exactly, and the signature chain and revocation status are valid.
2. The file uses a restricted, safe serialization format, so untrusted custom code does not acquire host permissions together with the model parameters.
3. The dependency closure, license, base image, and vulnerability status are traceable, and the runtime account has no superfluous file, network, or credential permissions.
4. Load order, scaling, quantization, merge, and conversion parameters go into an immutable configuration.

The dynamic layer performs behavioral differencing over "base—component—workflow—probe—seed". The probes cover both expected functionality and secret trigger sets, cross-lingual conditions, identity, and safety refusals. It compares semantics, image quality, abnormal overflow, and stability before and after loading. BlackMirror inspects generative backdoors from black-box input-output differences, which complements static scanning at the behavioral layer. Its effective scope, however, remains bound to the specific model and probe distribution [VIS_R-A029].

The composition space can grow rapidly with the number of components, and teams usually cannot enumerate all permutations. The correct response is not to abandon composition testing. Compress the space first with allowlists, least privilege, and a small number of supported templates. Then select high-risk combinations based on component functionality, shared layers of influence, and publisher relationships. Unapproved free-form combinations can be experimented with in an isolated environment, but should not inherit production permissions.

The non-orthogonality of video components

Video pipelines often load their components at the same time: character LoRA, spatial control, motion modules, camera control, temporal decoding, frame interpolation, audio, and quantization plugins. They differ nominally in responsibility, but in practice may write jointly into attention, latent states, or temporal representations. An anomaly may appear only when the character, motion, and camera components are present simultaneously. It may also disappear or become visible after quantization or conversion.

Direct evidence from image LoRA shows that composed artifacts deserve focused verification. It does not prove that malicious video motion modules already hold generally across different mainstream base models. An engineering verification plan should set this path as a high-priority test hypothesis. Perform temporal behavioral differencing on supported combinations and report action, trajectory, shot, identity, and clean utility. For combinations that have not been run, explicitly keep them unknown.

Artifact inventories must reach the component level

A software bill of materials answers which code and dependencies are included. A model bill of materials must answer more. It must say which training task the parameters came from, which base models they are compatible with, what format they use, and whether they contain executable objects. It must also record which layers they act on and their declared purpose. A practical entry can include a unique name, publisher, repository and commit, and file digest. It can then add the signature, base model, training data description, license, dependencies, entry function, required permissions, supported precisions, and verified combinations.

For a converted artifact, the original digest no longer applies. Quantization, merging, format conversion and sharding all produce new files. The system should therefore record the parent object, the tool version, the command parameters and the new digest. If an online service performs the conversion, the record must also include the service identity and the returned receipt. Once a defect in a conversion tool is discovered, teams can then query every derived artifact instead of rescanning the entire repository to guess at provenance.

Safe serialization and execution isolation

A model file may contain nothing but tensors. It may also use general-purpose serialization mechanisms to trigger object construction or code loading. Ingestion policy should prefer restricted tensor formats, which separate parameters from custom code. When code is genuinely required, it enters review and isolation as an independent software artifact. Loaders reject unknown object types, path traversal and implicit network access. Error messages should not leak host directories or credentials.

Isolation requires proof from behavior probes. Allow probes verify the necessary weight reads, GPU computation and temporary files. Deny probes verify that secret directories, host sockets, metadata services, unapproved networks and persistence paths are unreachable. A passing probe describes only the current environment and configuration. It does not prove that model behavior is benign. Output-layer differencing is therefore still required.

How combination testing controls the explosion

Some combination spaces cannot be enumerated exhaustively. When that happens, one can first stratify by impact surface. Text encoders, conditional LoRAs and prompt templates act jointly on semantics. VAEs, super-resolution and color nodes act on pixels. Character LoRAs, identity encoders and face restoration act jointly on identity. Motion modules, camera control, frame interpolation and temporal decoding act jointly on trajectories. Quantization, caching and parallel plugins act jointly on numerics and resources. Test combinations inside one impact surface first. Only then test the high-consequence paths that cross impact surfaces.

Combination selection should also weigh publisher and dependency correlation. Several modules published by the same author may share training or code failures. Modules from different repositories may also depend on the same underlying library. Treating "different files" as independent sources overestimates the independence of defenses. Closure risk records should retain shared parent objects, shared code and shared signers.

A behavior differential matrix can contain four quadrants: single-component clean, single-component confidentiality probe, combination clean, and combination confidentiality probe. Each quadrant spans several fixed seeds. Within it, semantics, identity, image quality, temporal behavior, resources and refusal are observed separately. MasqLoRA prompt adapter combinations require independent evaluation [VIS_P006]. BlackMirror shows that black-box differences can supplement static features [VIS_R-A029]. Neither eliminates the residual risk of insufficient probe coverage.

9.4 Updates and personalization: who has the right to change which parameters

Training or personalization services should organize update permissions into a change permission table. Such a table records who submits the data, who owns it, which layers may be updated, and what objective and regularization are used. It also records how many steps of training are run, what artifact is output, who may download or transfer it, when it expires and how to roll back. Upload permission covers only file submission. Use of a depicted person's identity, public dissemination, commercial endorsement and generation in sensitive contexts each require explicit authorization.

Diffusion backdoor research shows that training updates can establish a "specific condition—anomalous output" mapping while surface utility on ordinary inputs is preserved. BadDiffusion, TrojDiff and BadT2I differ in objectives, data, training setups and adjudication methods. They cannot be merged into a single general success rate [VIS_P001, @VIS_P002, @VIS_R-A003]. A qualified evaluation should report condition hits, target semantics, clean semantics, visual quality, non-target spillover, cross-seed stability and detector error separately.

Concept erasure and relearning also require stricter controls. To claim that an update restored a restricted concept, one should at minimum show that the concept was indeed unreachable on a fixed evaluation set before the update. One should also show that an equal-budget clean update does not reproduce the restoration. One should further show that the observed difference is not judge variance, nor a change in how the condition is expressed. Otherwise one cannot distinguish an incomplete original restriction, capability drift caused by ordinary updates, and targeted restoration.

Personalization is especially prone to conflating normal capabilities with safety properties. DreamBooth-style mechanisms tie a new identifier to a visual identity. They build that association from a small number of subject samples [GEN_dreambooth]. Safety testing must look beyond subject similarity. It must also observe background leakage, pose and style diversity and near-duplication of training samples. Confusion with other identities, authorization boundaries and residue after withdrawal need testing too. A drop in similarity may indicate that protection is effective. It may instead reflect an overall collapse in generation quality. Benign task utility must be reported at the same time.

The temporal clean utility of video updates

BadVideo demonstrated spatiotemporal backdoor behavior under a specific text-to-video training setup [VIS_P007]. Video evaluation must therefore be split at least into spatial clean utility and temporal clean utility. The spatial part checks frame quality, subject appearance and background. The temporal part checks motion direction, event ordering, identity persistence across shots, camera trajectory, flicker, lip sync and audio-visual synchronization. Image quality reported for only a few frames can hide action- or event-level anomalies inside the average.

Video training also involves updatable state such as temporal attention, motion priors, camera control, cross-frame caching, audio branches and long-clip curricula. Each update must be bound to the actual checkpoint and date. Once a model is further fine-tuned, distilled, quantized or merged, a previous "passed" status cannot be inherited permanently. Verification conclusions must propagate along version lineage and be re-executed at key changes.

The update process needs observable boundaries

A training job should not leave behind only a final model checkpoint. The minimum receipt covers the initial model, the data manifest, updatable layers, frozen layers, loss, regularization, the optimizer, training rounds, random state and the validation set. It also covers anomaly alerts, all model checkpoints produced, and the approver. High-risk personalization must additionally record subject authorization, output download rights and server-side deletion status.

During training one can monitor gradient and activation anomalies, trigger associations, clean replay and safety set performance. None of these signals proves that no backdoor exists. Anomaly detection thresholds may miss low-amplitude, in-distribution behavior. They may also treat long-tail subjects as outliers. Post-training confidentiality probes and runtime capability gates remain necessary.

Update permissions should be limited to the layers and steps the task requires. Character personalization usually does not need to change all safety classifiers and execution plugins. Safety alignment tasks do not automatically gain repository release rights either. Assign training, signing, release and production alias updates to different parties. That way no single internal account can change parameters and then directly overwrite production evidence.

Backdoors, safety-policy or authorization-boundary drift, and ordinary quality degradation

These three phenomena require different adjudications. A backdoor manifests as a condition-specific anomalous mapping, while surface capability on clean inputs is preserved as much as possible. Safety-policy or authorization-boundary drift means that subjects, purposes or content that should have been refused are allowed through after the update. It also means that normal requests previously permitted are systematically refused. Ordinary quality degradation manifests as a general decline in image quality, semantics or temporal consistency. BadDiffusion, TrojDiff and BadVideo provide research cases for conditional anomalies [VIS_P001, @VIS_P002, @VIS_P007]. They cannot justify attributing every quality drop to a backdoor.

Diagnosis requires at least trigger and clean conditions, before and after the update, target and non-target semantics, multiple seeds and independent adjudication. If all tasks get worse, training instability or compatibility problems are more likely. If only one judge changes, judge drift needs to be examined. The anomaly comes closer to the condition-specific behavior of a backdoor only when it is stably bound to a condition while clean capability is preserved. When evidence is insufficient, record that "the cause of the anomaly has not yet been localized". Retain multiple candidate attributions.

Version propagation and blast radius

A checkpoint may be distilled, quantized or merged into a LoRA. It may also be exported to edge devices or used as the initial point for downstream training. The supply chain graph should record these child relationships. After an anomaly is found in a parent version, one must not assume that all child versions are affected. One must not

assume that conversion has already removed the anomaly either. Freeze propagation first, then test derived objects according to high consequence and reachability.

The blast radius includes repository downloaders, production aliases, edge caches, user projects, downstream adapters and media that has already been generated. Model withdrawal addresses future loading. Content disposition addresses existing outputs. Subject notification and appeal address real-world rights. A single technical team usually cannot complete all of these actions alone. The runbook must therefore define responsibility boundaries before an incident.

9.5 Supply chain gates from ingestion to withdrawal

The six gates are consolidated into three stages by responsibility. That prevents misreading continuous controls as six mutually independent slogans. Each stage must have both entry conditions and an exit receipt that the next stage can verify.

Stage	Gates included	Entry and handling	Exit receipt and limits of the evidence
Ingestion stage	Gate one, "provenance and authorization"; gate two, "safe parsing"	Before data and artifacts enter the trusted environment, confirm the subject, purpose, license, version, hash and signature. Objects that cannot be confirmed go to quarantine. Parameters use restricted formats. Plugin code is inspected with no secrets, no public network or a controlled network, a restricted filesystem and resource quotas. Execution isolation is verified with allow and deny behavior probes.	The receipt states the object identity, authorization scope, parsing format and host-reachable capabilities. "Downloadable" does not equal approved. A passing scan does not prove that model behavior is benign.
Combination verification stage	Gate three, "composition compilation"; gate four, "behavior differential"	The system compiles components, load order, scaling, quantization, dependencies and policy into an immutable closure identifier. It then compares baseline and combination on approved clean tasks and confidentiality safety probes. It observes anomalous targets, normal semantics, image quality, temporal behavior, cost and false positives separately.	The receipt binds each run to the closure and probe version. It does not compress a single detection score into a "safe/unsafe" label. Uncovered combinations remain unverified.
Operation and failure stage	Gate five, "canary and observation"; gate six, "withdrawal and recovery"	A new closure enters only a limited set of tenants and a limited capability scope. It monitors condition anomalies, identity errors, temporal drift, resource curves and complaints. After a failure occurs, the list acts simultaneously on artifacts, combinations, data versions and checkpoints. It also tracks caches, edge copies, user downloads and derived adapters.	The receipt separately proves the canary scope, actual load blocking, dependency notification and recovery version. A release alias must not mask the immutable version. Deleting a repository page does not mean that all replicas are invalidated.

The six gates should not be controlled by the same service account and the same judgment model. Repository identity, execution isolation, behavior review, production authorization and revocation services need different roots of trust. That way a single compromised component cannot both forge evidence and approve itself.

9.6 Engineering Worked Example: Launch Decision for a Model Plugin Marketplace

An enterprise allows designers to select bases, character LoRAs, style LoRAs and video motion modules from an internal marketplace. In the past, the platform scanned file by file and then allowed arbitrary combinations. The review process for a new motion module is divided into four stages: object freeze, behavior boundaries, static and dynamic verification, and canary and withdrawal. Each stage produces an explicit receipt that the next stage can consume.

First, the object is fixed: module hash, publisher, training description, license, dependencies and code permissions. 3 supported bases and 2 character LoRAs are listed, along with quantized and non-quantized paths. The object under review is therefore no longer "a motion module" but a finite combination matrix.

Second, expected and prohibited behaviors are defined. Expected behaviors include improving the continuity of specified motions while preserving character identity and script semantics. Prohibited behaviors include producing a fixed anomaly under unauthorized subjects, confidentiality condition sets or specific component combinations. They also include accessing host secrets, or letting resources exceed the task budget. Video utility is evaluated by complete clip, motion and shot. Selected frames serve only for local observation.

Two categories of checks are then performed. Static checks confirm the safe format, dependencies, least privilege and signatures. Dynamic checks compare before and after loading on clean conditions, boundary conditions and confidentiality combination probes. Every result is bound to the base, combination, random seed, clip length, frame rate and moderator version. Bases that have not been run are not written as "compatible and safe". They simply remain unverified.

Finally, release and withdrawal are set up. A new module first enters a canary capability domain that only generates and has no automatic publishing capability. Anomaly alerts can freeze loading without relying on the generative model's self-assessment. If a problem is found, the invalidation list blocks download and execution by hash. The platform searches for closures containing the module, notifies the project owner and retains minimal verification evidence. The business can roll back to an approved closure, and keeps service shutdown in reserve as a handling measure for when other recovery paths fail.

Launch Checklist

Before a model or component obtains production permission, the reviewer confirms the following object, permission, combination and recovery facts item by item. Every item must be bound to the current closure and verification receipt. An unknown item is not automatically interpreted as a negative result. It is converted into an isolation, canary or capability-reduction condition:

1. Is there an accountable person for data provenance, authorization, deduplication and sensitive subject fields?
2. Can the training job trace back to an immutable data manifest and the initial checkpoint?
3. Are file digests, signatures, licenses and revocation status verified?
4. Are parameters kept separate from code? Are loading permissions tested with allow and deny probes?
5. Do conversion, quantization and merging generate new derived records?
6. Does the closure include all encoders, VAEs, adapters, plugins, containers and policies?
7. Are single-component and combined clean utility both tested with a fixed denominator?
8. Do confidentiality probes cover semantics, identity, cross-language, time and resources?
9. Does the video report motion, shot, identity persistence, audio-visual and complete events?
0. Are the canary scope, alert responsibility, automatic freezing and rollback target explicit?
1. Can the invalidation list block both download and production loading simultaneously?
2. Can all dependency closures, derived artifacts and already generated media be queried?

A single unknown does not have to veto a low-risk experiment on its own. It should still limit capability and propagation scope. High-risk identity work, automatic publishing, or large-scale distribution requires more complete evidence. Risk grading changes the threshold. It does not change whether the unknown item exists.

Incident Drill: A Module Is Reported to Have a Condition-Specific Anomaly

The drill starts by freezing new loading. It does not delete evidence right away. The repository retains the file digest, the signature, and the release information. The platform queries every closure, production job, and derived artifact that contains that digest. The model team replays the reported phenomenon with frozen probes in an isolated environment. The business team prepares to switch to the last approved closure. Propagation is limited before confirmation. No characterization of the module author is made public.

If the anomaly reproduces, the team determines the highest observed layer. It asks whether only the formation of the target media is affected, or whether that media has already been delivered, published, or used in high-risk business. Handling unfolds by layer. The artifact enters the invalidation list, and the production alias is rolled back. Dependent projects are notified, and the relevant media enter authenticity and platform review. Affected subjects

receive a communication channel. If the anomaly cannot be reproduced, the reported conditions, version differences, and unknown causes are still retained. The team then decides whether to continue isolation.

The drill's acceptance metrics include discovery-to-freeze time and closure query coverage. They also include rollback time, dependency notification, false freezing, evidence completeness, and restoration of normal tasks. A report that says only that "the module has been deleted" cannot show whether caches, user downloads, and derived weights can still run.

A Comparison Framework for Five Research Cases

Case	Required permission	First-broken location	Main observation	Scope of evidence support
Nightshade	Can influence the corpus that is scraped later	Data	Conceptual semantic shift and clean utility	The evidence supports concept-targeted shift in the specified data, model, and protocol. Other cleaning pipelines require independent testing
Rickrolling the Artist	Can make the system load an encoder	Artifact	Condition-specific semantic redirection	The evidence supports the condition-specific redirection the encoder carries. Whether the main denoising network changes must be verified separately
MasqLoRA	Can publish or load an adapter	Artifact	LoRA condition-specific anomalies and compositional behavior	The evidence supports the specific LoRA and its compositional behavior. Video motion modules require a separate experiment
BadDiffusion/TrojDiff	Can control training or the checkpoint	Update	Condition-specific backdoors and clean performance	The evidence separately supports the condition-specific backdoor in each training protocol. Cross-protocol aggregation requires a common unit and denominator
BadVideo	Can control video model training settings	Update	Spatiotemporal targets and normal video quality	The evidence supports the spatiotemporal backdoor in the specified video training settings. Commercial long-video services require on-site protocol verification

The comparison covers the attack entry point, required permission, propagation location, and deployable controls. The research cases differ in model, data, attack budget, independent unit, and observation endpoint. Those differences are the reason effect sizes cannot be directly ranked or aggregated. An engineering team can follow the common fields to find control locations. It then re-establishes its own results with the current closure, data, normal tasks, and denominator.

Supply Chain Responsibility Matrix

Party	Evidence that must be provided	Permission it should not hold alone	Action after a failure
Data provider	Provenance, authorization, version, labels, and a withdrawal entry point	Directly updating production checkpoints	Correct, stop supply, notify affected data versions
Training team	Data manifest, configuration, initial model, checkpoints, and validation	Self-signing and directly overwriting the production alias	Freeze training, locate the batch, provide rollback objects
Component author	Code and parameter digests, dependencies, compatibility, license, and purpose	Obtaining host secrets or arbitrary public network capability	Publish invalidation information, cooperate with replay and remediation
Repository operations	Identity, signature, scanning, download, and revocation records	Judging all behavior and lawful use	Block distribution, notify downloaders, preserve receipts
Platform review	Closure, behavior diff, benign utility, and canary results	Continuing to use old evidence after changing the original artifact	Freeze the closure, roll back, initiate dependency queries
Business owner	Scenario, subject authorization, consequences, and release scope	Bypassing artifact and capability gates	Limit the business, notify subjects, restore the safe process

The responsibility matrix separates "who signs" from "who judges." A data provider cannot prove that the training implementation is faithful. A repository cannot prove that an arbitrary combination is benign. A model team cannot grant identity use on behalf of the business owner. Evidence is passed between roles. Approval is made by the party responsible for the corresponding consequences.

Supply Chain Metrics Are More Than Vulnerability Counts

The data side can report the proportion of traceable provenance and the proportion of unknown authorization. It can also report deduplication family concentration, withdrawal propagation time, and the replay rate of training batches. The artifact side reports immutable version coverage, signature and revocation verification, safe format coverage, closure completeness rate, and behavior diff coverage. The update side reports configuration receipts, confidentiality probes, canary scope, rollback success, and derived object query coverage.

Operational metrics also cover anomaly discovery to freezing and freezing to rollback. They add dependency notification, false freezing, restoration of normal tasks, and objects that could not be located. Vulnerability counts are affected by scanning scope and reporting culture. They cannot alone represent security maturity. Traceability, blockability, and recoverability come closer to the capability actually available when the supply chain fails.

Component Closure and Derivation Graph Drill

A synthetic workflow is set up in a local isolated environment. It uses a base, a text encoder, a VAE, two ordinary LoRAs, a motion module, and a custom post-processing node. For each object, the team records the digest, provenance, dependencies, format, permissions, and compatibility. It then generates a closure identifier. A controlled quantization or format conversion then creates a new artifact with a parent relation. The experiment does not train a real attack and does not connect to production accounts.

The first stage verifies the static gate. Unknown digests, expired signatures, undeclared code, unauthorized network, and the wrong base should be rejected. The second stage runs a behavior diff on fixed clean prompts, subjects, seeds, and video tasks. The third stage simulates a component failure. It adds the component to the invalidation list. Then it queries the closures, the converted child objects, and the projects containing it. Finally it rolls back to the approved version.

Four statements are recorded separately in the experiment report: "the file was discovered," "execution was blocked," "the dependency was located," and "the business was restored". When the invalidation list blocks only new downloads while local caches can still load, distribution control is the current level of completion. When the derivation graph lacks a conversion relation, the scope of handling is limited to the child objects that were located. Every observation is bound to an actual execution receipt.

The Investigation Process for Compositional Identity Drift

A video team combined a character LoRA with a motion module and found that identity drifted in the third shot. Single-component tests all remained normal. The team first froze the base, the two components, the load order, scaling, quantization, seed, clip length, and moderator version. It set the independent video, not the shot or frame, as the primary denominator. Then it changed only one factor at a time for paired runs. It swapped the load order first, then compared scaling and quantization, and finally checked across fixed seeds whether the drift was stable. Each run simultaneously saved spatial identity, temporal persistence, motion quality, and resource cost. Ordinary quality degradation, compatibility conflict, and condition-specific anomalies thus had distinct observation patterns.

The investigation then returned to training and conversion lineage. The character LoRA, motion module, and quantized artifact were each linked to a training data version, an initial checkpoint, a conversion tool, and a parent object. When the compositional anomaly appeared only in a quantized child object, the scope of handling was first limited to that derived branch. If the same condition-specific anomaly appeared across multiple parent versions, it was then expanded to upstream components. The canary system paused new closures containing suspect objects and kept approved closures rollback-capable. It determined the notification scope from actual dependency query results. This process connected each observation to its corresponding object. It avoided attributing an entire model family from a single shot.

The same team then handled subject consent withdrawal. The propagation graph connected raw material to clips, descriptions, data versions, training jobs, checkpoints, LoRAs, production closures, caches, and already generated media. Every edge had a responsible party, an executable withdrawal action, an irreversible state, and a verification probe. Data not yet trained stopped entering batches. Hosted personalized artifacts were invalidated, and production closures stopped loading. Caches were confirmed cleaned via negative probes on subject and version. Media already delivered entered use restriction, notification, and remedy processes. Withdrawal thus became a state transition across objects, not a single button on a repository page.

Procurement and External Model Integration Requirements

An enterprise may procure a closed-source model or a hosted generation service. Even if it cannot obtain complete weights and training data, it can still require four categories of verifiable evidence. Version commitments state how the model, moderator, policy, and high-impact dependencies are identified, and which changes trigger re-acceptance. Data-use statements describe prompts, reference media, generated results, human viewing, training use, caches, and the retention state after deletion. Run receipts cover tenant isolation, personalized artifact holders, output provenance signals, error paths, and withdrawal results. Notification time limits specify the maximum time from discovery to customer awareness for major version changes, data-use changes, and security incidents.

Vendor self-reports can support the existence of their claims. Customer-side probes verify runtime behavior. Customers use authorized material to test tenant isolation, withdrawal, version changes, output delivery, caches, and error paths. Procurement terms should specify notification time limits for major model, moderator, and data-use changes. When internal details cannot be provided, external behavior, service receipts, and liability time limits matter more.

By contrast, open-source or open-weight integration offers more static inspection capability. It lacks a central operator able to enforce withdrawal. Enterprises need self-managed repository mirrors, digest allowlists, dependencies, licenses, behavior diffs, and invalidation information subscriptions. Readable weights do not automatically equal provable training provenance. Being able to run locally does not automatically equal obtaining arbitrary commercial or identity uses.

Procurement decisions should ultimately be bound to the scenario. A read-only creative sandbox can tolerate more unknowns, but it must not hold production accounts. Automated brand publishing, digital humans of real people, and high-risk business require stronger versioning, authorization, notification, and rollback. The way to handle unknowns is to limit capability and propagation. It is not to fill unknowns in as "pass" in a table.

9.7 Quick Reference for Judgment Boundaries

Judgment Easily Made	Engineering Judgment It Should Become
Public data is data that can be trained on arbitrarily	Accessibility, license, subject consent, and permitted use are different fields. A portrait or voice that appears publicly does not naturally grant permission for personalization, sensitive-context generation, or public endorsement.
The main model passes the check, and plugins only affect style	Encoders, LoRAs, VAEs, motion modules, and custom nodes can all influence semantics or execution paths. A component name is a claim, not a capability boundary. Runtime combinations require independent verification.
Signatures and malicious artifacts are mutually exclusive	Malicious publishers can sign just as well. Signatures address provenance and transmission integrity. Behavior diffs and least privilege address another part of the risk.
Clean image quality being preserved means clean utility is preserved	Semantic adherence, identity, cross-language, motion, shot, audio-visual sync, and safety refusal can all drift. Image quality is only part of benign utility.
A withdrawal link equals withdrawal capability	Caches, already downloaded copies, derived weights, deployment aliases, and downstream combinations may still retain anomalies. Withdrawal requires version lineage, invalidation lists, dependency notifications, and load probes.

9.8 Bringing This into Real Systems

An enterprise model marketplace accepts bases, encoders, VAEs, LoRAs, motion modules, and custom nodes. At ingestion, the repository saves publisher identity, provenance address, license, file digest, signature, dependencies, code permissions, and withdrawal status. Training artifacts are additionally connected to immutable data manifests, initial checkpoints, trainable layers, and configuration receipts. Video material containing people and voices establishes deduplication families by original event. Cropped clips, transcoded files, subtitles, and automatic descriptions all preserve parent relations. Subject authorization separately constrains training, personalization, public dissemination, and commercial use.

A candidate component first obtains test eligibility. The platform then uses supported bases and component templates to generate a limited closure matrix. Static checks handle format, dependencies, and host permissions. Dynamic diffs handle clean semantics, identity, motion, shot, audio-visual, resource, and confidentiality conditions. Single-component and combination results are saved separately. Run identifiers include load order, scaling, quantization, seed, clip length, frame rate, and policy version. The identity drift in the third shot can thus be traced back to a specific closure, rather than being obscured by a continuously changing "latest" alias.

Closures that pass verification first enter a canary domain. That domain only generates. It has no automatic publishing capability. The repository identity service, the isolated execution environment, behavior review, production authorization, and the withdrawal service use separate accounts and trust roots. Business owners approve only closures that conform to subject authorization and scenario use. The platform separates generation, download, and publishing into different capabilities. Benign utility, false refusals, resources, complaints, and condition-specific anomalies are continuously recorded by actual run version. Major version or data-use changes trigger re-evaluation.

When a component fails, the invalidation list applies simultaneously to new downloads and production loading. The dependency graph queries all closures, quantized or merged child objects, project caches, and already generated media. The business rolls back to the last approved closure, and the repository notifies downloaders. The subject authorization team handles derived media. Operators use negative load probes to confirm that the old object can no longer run. The final receipt separately states that the file was identified, execution was blocked, dependencies were located, notifications were sent, and the business was restored. Supply chain response thus becomes a set of observable states.

Summary: Approve a Closure, Not Trust a Name

The data supply chain determines learning objects and authorization. The artifact supply chain determines deployment objects and execution permissions. The update supply chain determines how parameters continuously change. The three paths converge at runtime into a replayable artifact closure. Provenance, signatures, and static scanning establish the identity foundation. Secure parsing limits execution capability, and compositional behavior diffs observe condition-specific anomalies. Canary and withdrawal control the scope of propagation. No single layer can complete all the tasks alone.

Supply chain controls lower the risk baked into a model or component. They cannot stop a legitimate model from accepting malicious conditions at runtime, leaking training memorization, or being abused through resource and caching mechanisms. Chapter 10 moves the observation point inside a single request. It traces conditions through encoding and moderation, follows the state sampling preserves, and shows how black-box feedback becomes an attack budget.

\newpage

Chapter 10 Conditions, Sampling, Privacy, and Generation Services

An image-to-video service received three inputs. One was an ordinary portrait photo, another a motion description with no obvious risk, and the third a camera trajectory. Each passed moderation on its own. Yet the generator

completed an unauthorized event during the middle few seconds. Suppose the platform records only "the text passed" and "the first frame is compliant." It will blame the semantics of the joint conditions on coincidence. If it records only the final refusal, it hides that the generator had already formed the target event. It also loses exactly which layer made moderation take effect.

Runtime safety is hard precisely because of this stage difference. A condition is not a string. It is a control surface built from text, images, video, audio, masks, depth, pose, boundary frames, trajectories, history and random state. Sampling adds query feedback, training memorization, caches, queues and compute budgets. Judging whether an attack succeeded therefore means observing condition entry, generation completion, risk formation, output blocking and actual delivery separately. A single attack success rate must not cover the whole chain.

Chapter Overview

One visual generation request carries text, images, video, audio, masks, pose, trajectories, history and random state at once. This chapter gathers those conditions into structured objects, each with a type, provenance, authorization and permitted scope of influence. It then traces how they change the generation trajectory through black-box semantic search, multimodal composition and video temporal completion. Condition passing, target formation, output review and content delivery sit on four separate denominators. A result from any stage must keep its own evidence position.

Sampling also exposes training data memorization, membership inference, approximate cache boundaries and resource occupation. Those four problem classes differ in denominator, attack budget and result object, so the evaluation definitions appear separately below. The final layered control separates input acceptance, in-generation checks, output handling and actual delivery. It also records benign utility, latency, cost and recovery capability, so the runtime gate still carries an observable business cost when it blocks risk.

Mechanistic groundwork: how one request becomes a stateful generation session

This chapter follows one stateful generation session, from conditional submission to content delivery. The input is not only a prompt. It also includes reference media, masks, pose, trajectory, history, random seeds, tenant identity, authorization fields and resource budget. Internally, the mechanism normalizes and encodes each condition first. The sampler then updates the generation state along discrete steps or a continuous path, while it maintains the query ledger, queue, cache keys, risk decisions and cancellation state. After generation ends, the decoder, the filters and the publishing service keep changing the visible media and its reachable scope.

Output consumers fall into preview interfaces, download services, editing tasks, public publishing accounts and downstream automated systems. They do not require the same permissions. The security surface therefore includes low-integrity conditions that override authorization, multimodal combinations that form new semantics, and black-box feedback that supports adaptive search. It also includes near-duplicate outputs that leak training information, caches that cross tenant or policy boundaries, and high-cost tasks that crowd the queue. Input acceptance, risky content formation, output blocking and actual media delivery each need their own denominator. They cannot be compressed into a single success rate.

In the benign case, a person reference image affects only the approved appearance fields. The cache key binds tenant, policy and subject authorization together, and a passing preview still requires independent publishing approval. One boundary case reverses this: text, image and camera trajectory are each compliant in isolation, yet their combination fills an unauthorized event into the video. In another, review has already rejected the final download, but an intermediate preview had earlier returned the clip to the user. The first case requires combinatorial semantics and in-generation observation. The second requires staged delivery evidence. A final rejection cannot be read backward to mean the entire run chain never formed risk.

10.1 Conditions are a control plane, not just a prompt

Every modality has its own provenance, authorization, validity period and permitted scope of influence. To keep them apart, this chapter expands one generation request into a typed condition vector rather than compressing it into a single prompt string. Missing fields stay unknown, and the model itself does not fill them in as permissions:

$$C = (C_{txt}, C_{img}, C_{vid}, C_{aud}, C_{mask}, C_{pose}, C_{traj}, C_{hist}, C_{seed}).$$

The components denote text, reference image, reference video, audio, mask, pose, trajectory, history state and random seed. Interfaces of the ControlNet kind show that depth, edges or pose can enter the generation process as independent structural conditions [GEN_controlnet]. In production systems, conditions should also carry metadata: source subject, tenant, integrity, authorized use, validity period and influencable variables. The same pixels carry different permissions when they come from the user, a public web page or an enterprise asset library.

Condition typing exists mainly to stop data from escalating into policy on its own. A reference image can supply subject appearance, but it cannot change the platform's identity authorization rules. Boundary frames can constrain the start and end points, but they cannot prove that the intermediate actions were approved. History edits can help maintain coherence, but they should not carry the sensitive permissions of an old session into a new task. Suppose the system concatenates all inputs into one unified embedding and preserves neither provenance nor use. Downstream reviewers can then no longer tell who a given semantic came from, or whether it has the right to influence the output.

One request maps to different objects and denominators as it enters the queue, as media is formed, as it goes through output disposition and as it is actually delivered. The runtime chain is therefore split into four observation gates. Each gate keeps its own entry counts, exit states, undecidable items and responsible controls:

1. **C0 condition acceptance:** whether pre-review allows the request into the generation queue.
2. **C1 generation formation:** whether the model finishes generation and forms the target semantics or event inside the media.
3. **C2 output disposition:** whether post-review identifies the output and blocks, degrades, or escalates to human review.
4. **C3 actual delivery:** whether the user, the platform, or an external system really obtains the content.

The four gates may use different denominators. Take 100 inputs: 80 enter the queue, 70 complete generation, 5 form the target risk, 4 are blocked and 1 is delivered. Reporting must preserve the counts and undecidable items at each level. "1/100" denotes the final delivery proportion, and "5/70" denotes the proportion of risk formation among requests that completed generation. The four valid blocks are recorded separately. A number carries engineering meaning only when bound to an object and a level.

Follow the figure from user and reference conditions through to the publishing cache. Mark out where condition encoding, generation state, sampling, filtering, content provenance and processing history, and permission checks each produce or consume state.

条件、采样与发布服务的运行链



先沿实线识别信息如何获得状态与能力，再判断哪一处需要独立控制。

Seven boxes pass from user and reference conditions through encoding, generation, sampling, filtering, and provenance signals to the publishing cache.

Figure 10-1 supports bringing random seeds, sampling state, filters, content provenance and processing history, cache keys and publishing permissions beyond prompt review into the run record of the same request. It also supports setting an independent denominator for each stage. It does not prove that the filtering or provenance signals in the figure suffice to block risk, and structural connectivity does not equal effective control. Each gate still needs verification with allow, deny, expiry, cross-tenant and recovery probes.

Condition compiler: preserving semantic provenance before the model

Before conditions enter the model, a production system can place a layer of condition compiler. That layer does not judge all media risk. Its job is to convert raw input into typed task objects. For example, a person reference image produces a "subject appearance" field, the authorization service produces an "allowed use" field, text parsing produces an "action and scene" field, and a camera trajectory produces a "shot control" field. The model receives the compiled allowed fields, and free text cannot override the authorization state.

The output of condition compilation can be written as a triple $\kappa = (v, p, o)$. Here v is the normalized value, p is the provenance and integrity label, and o is the set of variables permitted to be influenced. Take the same "someone's name" arriving from a verified identity token and from web page OCR. There v may match, while p and o should differ. The model can use web content to answer information, but it cannot use that content to authorize a person's endorsement or public publishing.

The compiler must also handle conflicts explicitly. Suppose the user's text requires one action, the reference video implies the opposite action, and the brand policy forbids that scene. The system still cannot let the last concatenation order determine priority. High-integrity policy and subject authorization take precedence over ordinary content. Conflicts that cannot be parsed enter a degraded or human path. Conflict-handling rules belong to external policy, and they should not be hidden inside the generation model's natural-language judgment.

Writing, reading, and invalidation of condition state

A condition may exist only within one request. It may also be written into templates, project history, personalization profiles or caches. Every write point requires a purpose and a term. An authorized person reference image should not automatically enter a team's public asset library. One project's negative prompt should not be reused across tenants. After identity authorization is withdrawn, condition tokens still held by an old session should become invalid.

Read policy likewise requires minimization. The generator reads only the fields needed to complete the task, the output reviewer reads the media and necessary context, and the publishing service reads the approval summary and account. No component has to see the raw sensitive material. Reducing the readable surface lowers leakage and makes logging and responsibility clearer. Condition invalidation must also be verified with probes. The system can construct expired tokens, a wrong tenant, a withdrawn subject, a changed policy version and a changed media digest. It can then confirm that the cache, generation and publishing each deny. An interface that shows "authorization withdrawn" is no proof that every derived piece of state in the backend has already become unusable.

Attack budget should become a session ledger

The capability of black-box condition attacks comes from queries, feedback and time. The session ledger records at least independent conditions, normalized similarity families, generation counts, failed requests, review feedback, account linkage, waiting time, cost and intermediate previews. If rate limiting is applied only to exactly identical strings, an attacker can circumvent it with synonymous expressions or cross-modal variants. Rate-limiting rules must also combine account linkage, similarity families and anomalous feedback, while preserving a reasonable budget for benign creative iteration.

A more appropriate control modulates feedback by risk relevance. Low-risk tasks can return detailed creative suggestions. Requests that repeatedly approach high-risk subject-action combinations receive reduced intermediate information, increased waiting, or escalation to human review. The restriction aims to raise the cost of adaptive search and protect the service, not to punish ways of expressing things. Effectiveness reporting gives attack query changes, benign iteration completion rates, latency and appeals together.

Complete recording of stage denominators

Beyond C0–C3, generation failures, timeouts, undecidable items and human escalations should also be recorded. Generation failures cannot silently disappear from the denominator, because certain conditions may systematically trigger failure and thereby create a usability problem. Undecidable items cannot be counted as safe success or attack failure. Human escalations must record the final decision and the wait. A level's denominator can be traced back to the level before it, but the two need not be equal.

Stage records can also identify where a defense is located. If a large number of risky conditions are blocked at C0, input control is effective, but benign content over-refusal must be observed. If risky media is formed at C1 and blocked at C2, the consequence gate is effective while the upstream still needs improvement. If C2 has already alerted and C3 still delivers, the problem lies in authorization or enforcement, not in the review model itself.

10.2 How semantic search bypasses literal rules

Input filtering often assumes that dangerous intent will appear in a fixed vocabulary. What the generation model actually consumes is an encoded semantic representation. Literal similarity therefore does not necessarily mean embedding similarity, and literal harmlessness may also preserve the target semantics after combination. A black-box attacker does not need to read gradients. As long as they can repeatedly submit conditions and observe acceptance, rejection, or output, they may turn system feedback into a search signal.

SneakyPrompt investigated how black-box query feedback can find alternative expressions [VIS_P008]. This evidence supports "the accept/reject interface leaks a signal usable for adaptation," but the effect is bound to the target model, filter, query budget and decision rules. A single accidental pass proves that a path exists and does not give an overall vulnerability rate. Zero passes only indicate that no success was observed under the current budget and sample.

Defense cannot stop at expanding the banned-word list. A more robust approach controls semantics, budget and capability at the same time. The input layer normalizes encoding and identifies synonymous and cross-language intent. The session layer links variant queries and limits low-cost iteration. The generation layer degrades or segment-checks high-risk capabilities. The output layer independently observes the generated media. The delivery layer makes the final decision based on subject authorization and use. Layers should rely on different roots of trust wherever possible. Otherwise an attacker need only adapt to a single decision boundary.

Multimodal composition is not the sum of risk scores

Each condition can pass on its own. That does not entail that the joint embedding is compliant. MMA-Diffusion manipulates text and image conditions together, which shows that the multimodal attack surface expands with compositional semantics [VIS_P009]. An attacker may place the subject in the reference image, the action in the text, and the scene or timing cues in masks and trajectories. No single modality, viewed alone, holds the complete intent. Yet the generator will synthesize them into one event.

Joint moderation likewise cannot just take the maximum or the average of two classifier scores. The composition question must be framed first. Which condition supplies "who"? Which condition supplies "what to do"? Who controls "in what scene and at what time", and has the subject authorized that use? In engineering terms, the conditions can be parsed into a structured intent graph. An external policy then judges the permitted subject-action-scene combinations. If the parse is uncertain, high-impact tasks enter human review or a lower-capability mode. The generator is not left to interpret and approve on its own.

JailbreakDiffBench separates prompt moderation from image moderation so the two can be compared, which supports staged evaluation [VIS_R-A030]. The findings remain bounded by the model and moderator versions, and they cannot be used to declare one filtering approach stronger across all services. The engineering value lies in one reminder. The same request at input and at output is two different objects. False positives, false negatives, and benign utility should be recorded separately.

Four classes of condition-control evasion

Runtime evasion can be divided by the gap it exploits. The four gaps are the literal-semantic gap, the modality-composition gap, the condition-output gap, and the single-turn-session gap. The literal-semantic gap exploits inconsistency between string rules and encoded semantics. The modality-composition gap disperses subject, action, and scene across text, images, or audio. The condition-output gap makes the input appear compliant while relying on the generator to complete the target event. The single-turn-session gap combines multiple requests, edits, and extensions into the final media.

The four classes of evasion may coexist. An attacker first uses semantic variants to get the text into the queue. A reference image then fills in the identity. Further rounds of editing extend an event. If the platform gives only a binary label per turn, the last turn looks like an ordinary local edit. Trace records need to preserve derivation relations and the complete timeline. They must also re-evaluate the composed result before key deliveries.

Adaptive attacks and static test sets

A static test set measures how the system performs on known inputs. An adaptive attack instead observes the defense's results and selects the next condition. SneakyPrompt shows that black-box feedback can support the search for alternative expressions [VIS_P008]. A deployment evaluation sets at least two stages. The first compares versions on a frozen set. The second lets the red team adapt to feedback within a fixed total budget. The two results cannot be merged, because the attack capabilities differ.

Adaptive testing needs stopping rules. Stop once thresholds for queries, cost, time, or high-risk output are reached. When something is observed that may involve a real subject, a minor, or a reusable service weakness, isolate the evidence and notify the responsible person. Reports retain the attack strategy category and the budget. They do not disclose high-risk payloads that can be directly reused. This verifies adaptivity without turning the research process into an abuse guide.

A responsibility matrix for multimodal conditions

Condition	Primary provider	Permissible influence	Must not be acquired automatically	Minimum verification
Free text	User or upstream agent	Topic, description, general action	Identity authorization, platform policy, payment or publishing rights	Semantics and cross-language, provenance labels
Reference image	User, asset library	Composition, subject appearance, style reference	Sensitive context, public endorsement, cross-project retention	Subject consent, near-duplication and use
Mask/depth/pose	Editor or vision tool	Spatial region, structure, pose	Semantic authority beyond the editing scope	Type, scope and boundary conditions
Boundary frames/reference video	User or preceding task	Start and end states, motion reference	Approval of the entire intermediate trajectory	Event-level completion check
Audio/script	Subject, editor	Rhythm, speaker, content	Likeness rights, public channels, transaction instructions	Speaker, script, audio-visual and channel
History and templates	Project system	Coherence and reuse configuration	Cross-tenant data, expired authorizations	Write gate, expiry, tenant and revocation

The matrix is valuable because it breaks "multimodal moderation" into questions of responsibility. A model can help identify the content of conditions. It cannot create missing authorization through its own output. Any high-impact permitted scope should come from a verifiable subject and an external policy.

A worked example of condition conflict

A brand project uses an authorized product image and person image. The text asks for "the person displays the product". The history template, however, contains another brand's camera and a spoken script. The condition compiler finds the brand, the script, and the subject authorization inconsistent. It downgrades the history template to a non-executable suggestion. The generator may still use general shot language, but it cannot read the old

project's brand or account fields. If the user confirms a new script, the system generates a new approved object. It does not directly update the old template.

This worked example shows the boundary between state contamination and ordinary creative reuse. The defense is not to prohibit all templates. It is to retain provenance, project, and influenceable fields. Benign users can still reuse shot structure. Sensitive identity and publishing information do not drift along with the template.

10.3 Video conditions control the trajectory

Image conditions mainly constrain a single spatial object. Video conditions also constrain how state changes over time. Boundary frames, reference video, action templates, camera paths, and audio scripts jointly determine the intermediate trajectory. If the system moderates only the boundaries visible at input time, it hands the most important event generation to the model. It keeps no independent control.

Two Frames Matter shows that the start and end states can pass individually while the intermediate completed trajectory still forms the target event [VIS_LN05]. The first-broken point here is not any single frame. It is the system's mistaking "boundary compliance" for "path compliance". Executable controls include the following. Review boundary conditions jointly. Observe action and identity by window during generation. Check the whole segment along event boundaries after completion. Require the subject or responsible person to confirm high-risk trajectories before delivery.

Text-to-video red teaming must treat time as a dimension of judgment as well. T2VSafetyBench requires evaluation of the video actually generated, not of the original prompt. The benchmark also distinguishes automatic from human judgment [VIS_P034]. Work on optimization-based video jailbreaking and time-aware red teaming goes further. It shows that prompt search, temporal semantics, and video output are distinct stages [VIS_R-A023, @VIS_LN04]. These works use different models, durations, moderators, and budgets. Their proportions cannot be merged directly.

A minimum video condition record specification records at least these items: model and policy version, duration, frame rate, resolution, shot count, condition modality, subject authorization, query budget, generation completion status, time when risk first appears, duration, output disposition, and final delivery. Evaluation units include frames, continuous windows, shots, whole videos, and events. Adjacent frames are highly correlated. They cannot be repeated as independent samples to increase statistical confidence.

Proactive protection from image conditions to video

A reference image may belong to a creator, a product, or a real person. PhotoGuard, I2VGuard, Anti-I2V, and VPA-Guard raise the cost of unauthorized use. They work from directions such as image editing, image-to-video, and visual prompt protection [VIS_P018, @VIS_R-A025, @VIS_R-A026, @VIS_R-A027]. The effectiveness of these methods depends on the base model, preprocessing, perturbation budget, and benign animation needs.

Proactive protection cannot replace the authorization gate. If a defensive perturbation makes even legitimate editing entirely unusable, the system merely trades destroyed benign utility for local blocking. If the platform does not verify subject permissions at upload time at all, there is no follow-up control once the technical protection fails. A qualified evaluation should report perceptual quality before and after protection. The same report should cover legitimate generation utility, subject or style similarity, motion quality, cross-model and cross-preprocessing performance, and defender cost.

Trajectory safety needs state constraints

A video generator selects the appearance of each frame. It also implicitly selects state transitions. For high-risk tasks, an allowed state set $\mathcal{E}$ and an allowed transition set $\mathcal{M}$ can be defined. Identity, position, action, shot, and audio-visual relations fall in $\mathcal{E}$. Speed, direction, event ordering, and subject continuity between adjacent windows fall in $\mathcal{M}$. A generated result enters delivery review only if it satisfies both the state and the transition constraints.

The allowed set is graded by scenario consequence. Open creation retains a wider narrative space while enforcing subject authorization and content policy. Product advertising adds brand, script, and claim constraints. Industrial

simulation adds physical boundaries and object relations. Real-person digital humans add identity, audio-visual, and publishing channel constraints. The higher the risk, the more specific the allowed set.

Trajectory detection also has failure conditions. Visual tracking may lose identity under occlusion and shot changes. Action recognition may be affected by frame rate and cultural context. Audio-visual synchronization may misjudge legitimate dubbing. The system should allow "uncertain". The consequence gate then decides degradation, regeneration, or human review. It must not map uncertainty to pass by default.

Video condition evaluation matrix

A complete matrix crosses at least six axes. These are text directness, reference media type, boundary-frame relation, duration and shots, subject authorization, and model and policy version. Each cell fixes the prompt, subject, seed, and number of generations. It records condition acceptance, generation completion, event formation, first risk, duration, output blocking, and delivery. The clean set includes complex actions, cross-language content, legitimate dubbing, and creative editing. A moderator therefore cannot obtain high surface safety by rejecting all complex videos.

Metrics should not be only the average of window scores. Event recall answers whether a risk event is found. Localization intersection-over-union and temporal error answer location. First alert answers real-time performance. The longest undetected risk segment answers sustained exposure. Identity trajectory errors answer cross-shot consistency. Benign rejections and waiting answer cost. Different metrics serve different decisions. They cannot first be weighted into a total score that has no interpretation.

Division of labor between in-generation control and final review

In-generation control is suited to halting obviously out-of-bounds generation early, limiting resources, and reducing the full formation of high-risk media. Final review is suited to observing the complete event, narrative, and audio-visual content. If intermediate control produces a false positive, it wastes part of the computation and harms benign tasks. If a problem is discovered only at final review, the computation has already been consumed and intermediate previews may have leaked. Products need to specify whether intermediate previews are delivered, how long they are retained, and whether they enter the cache.

In segmented video, each clip can pass on its own. That does not entail that the stitched narrative passes. The export gate should re-evaluate the complete timeline and check clip order, dubbing, and provenance derivation. A streaming system that cannot wait for the end uses a sliding window and an event state machine. When uncertain, it switches to a safe frame or reduces propagation. It does not continue to accumulate risk.

10.4 What sampled outputs may leak

Generative models may retain certain features of training data in their parameters. Training data extraction probes the model with a set of conditions or queries. Its aim is to find outputs that approximately and identifiably correspond to training instances. Research on image diffusion models has demonstrated a path to extracting training samples under an explicit threat model and query budget. It also discusses duplicate data and deduplication [VIS_P024]. This does not mean that any similar image proves training membership.

Confirmation requires at least five pieces of evidence. The candidate training instance can be located. The output reaches a predefined near-duplicate relationship with the candidate. Member and non-member comparisons use the same protocol. Manual verification does not only select successful demonstrations. The number of queries, seeds, and screening costs are fully counted in the budget. If the provenance of the training set is unknown, one can only say that the output is similar to an external candidate. One cannot automatically say that it was memorized by the model.

Training data extraction, membership inference, and model extraction also cannot share a single "privacy success rate." The first is counted by independent training instances and queries. Membership inference is counted by member/non-member samples. Model extraction is counted by query budget and surrogate task fidelity. Their consequences also differ. Obtaining one near-duplicate piece of media, inferring that someone may be in the training set, and building a behaviorally similar surrogate model require different controls and notifications.

Video adds content memory and motion memory. A model may not copy a complete file, yet reproduce identity and background across several frames. It may also reproduce a distinctive action pattern. The relevant research explicitly distinguishes these two types of memory [VIS_A047]. Averaging distance over the whole segment dilutes short leaks. Single-frame similarity may misjudge common scenes as memory. Frame, segment, motion, audio, and identity evidence should therefore be combined. Training-segment duplication, authorization, and non-member baselines should be reported.

Privacy defenses unfold along the lifecycle. Cross-modal deduplication and sensitive-identity governance come before training. Monitoring of memory and near-copies applies during training. At the output end, near-duplicate retrieval covers high-risk identities, frames, segments, and motion. At the service end, query correlation, rate limiting, and candidate feedback apply. At the incident end, data subjects get entry points for checking, notification, and deletion. The stricter the thresholds, the more likely legitimate common compositions and actions are to be incorrectly blocked. Benign utility and manual review costs must be measured at the same time.

The ladder from "looks similar" to "evidence supports memory"

Privacy claims can be divided into four classes of evidence. Similarity observation evidence records only the similarity relationship between the output and public media. Candidate near-duplicate evidence finds near-duplicates in a candidate training set and confirms them with frozen thresholds and manual verification. Membership difference evidence adds member and non-member comparisons, showing that the candidate has an anomalous association. Lineage attribution evidence connects training members and model outputs on the basis of explicit data lineage and query budget. Each class of evidence has its own object and denominator. The scope of the conclusion increases step by step with the type of evidence.

Research on image training data extraction offers a protocol example for higher-level evidence [VIS_P024]. Production services, however, often keep their training manifests undisclosed, so a real investigation may stall at a lower level. A platform can still take precautionary measures: limit distribution, contact rights holders, preserve evidence. It cannot treat a training attribution that is not yet established as established fact.

The video ladder adds temporal correspondence. A candidate may be similar only in contiguous segments, or only in motion. Handling must cover cropping, speed changes, mirroring, transcoding and audio-track removal. Frame-level matching, segment-level matching, motion-level matching and identity-level matching each return their own results. Attribution grows stronger only when the evidence points jointly at the same candidate. Strict thresholds still require reporting false positives on real videos.

Query privacy and returned information

A generative service may leak more than its training data. Other users' prompts, reference images or cache hits can escape as well. Error messages, latency differences, intermediate previews and controllable seeds each add an observable signal. Secure design should minimize cross-tenant state, unify the error differences that serve no purpose, and release debug information only to authorized operations personnel.

Limiting returned information also shapes the creative experience. Teams should measure how many turns a benign user needs to complete a task, how long failure diagnosis takes, and how many support tickets arrive. They should not hide all feedback. Tiered feedback is one compromise. Ordinary errors get safe, actionable categories. Internal logs retain the detailed reasons. High-risk repeated queries receive less of the information that could be turned toward boundary probing.

10.5 Cache, randomness, and resources are security state

Diffusion and video generation are costly, so services often reuse conditional embeddings, intermediate latents, attention results or final media. Approximate caching judges by similarity whether requests are close enough. A cache key may omit the tenant, the policy version, principal authorization, component closure or another security-critical condition. Two requests with similar content but different permissions may then share state incorrectly.

Research on approximate caching for text-to-image diffusion shows that cache match boundaries can be exploited [VIS_P039]. A request that should not hit may reuse a result incorrectly, or a request that should hit may miss, and

output correctness, latency and cost all shift as a result. That evidence covers only the cache design and threat model used in the study. It cannot be extended into a claim that all commercial generative services carry the same vulnerability.

A secure cache key must not contain only a text or pixel digest. It must also distinguish tenant, artifact closure, policy, authorization, and the category of randomness permitted for reuse. So that these security-relevant equivalence conditions can be explicitly checked, it can be written as

$$K = H(\textit{tenant}, \Omega, \textit{policy}, \textit{auth}, c_{\textit{norm}}, \textit{seed_class}),$$

where *tenant* is the tenant, Ω is the artifact closure of Chapter 9, *policy* is the policy version, *auth* is the authorization state, $c_{\textit{norm}}$ is the normalized condition, and *seed_class* denotes the category of randomness permitted for reuse. The cache key can store field digests or irreversible encodings. Any state that changes security semantics must participate in the equivalence judgment. The cached data itself also requires encryption, retention limits, access control, hit logs, and cross-tenant negative probes.

Resource security is also an operational requirement. Resolution, frame count, duration, sampling steps, control branches, number of candidates, audio, frame interpolation, retries and concurrency all affect GPU time, VRAM, queues and billing. Public evidence is not yet sufficient, however, to give a general attack effect for compute amplification specific to image or video generation. In engineering terms, it should be set as a proposition to be verified, not settled by borrowing numbers from other domains.

A qualified isolation test records a fixed set of quantities per request and per tenant. They are GPU time, peak VRAM, queueing and tail latency, release after cancellation, cache hits, failed retries, cost, and the benign request completion rate. The test varies the scale of benign tasks step by step and compares against equal-budget benign load. It sets hard quotas, preemptible queues, timeouts, cancellation reclamation and retry caps. The report also retains benign video quality and long-task completion rates. It does not manufacture surface-level safety by banning high-cost tasks outright.

Video streams turn resource problems into long-lived state. Before a single request ends, the system may continuously retain latents, cross-frame attention, audio caches and partial outputs. Interactive generation also persists across turns. Quotas should cover single tasks, sessions, tenants and organizations at the same time. Cancellation must be verified to actually release GPU and cache. Merely stopping the front end from displaying progress is not enough.

Cache equivalence relations must be consistent with policy

Two requests being "computationally similar" does not equal being "security-equivalent." The same prompt under different tenants, principal authorizations, policy versions or artifact closures cannot share results that would leak or bypass review. Within the same tenant, a low-resolution preview that has already passed cannot automatically substitute for final high-resolution review. The cache layer must understand which fields change security semantics. Otherwise it should simply disable cross-request reuse for high-risk tasks.

Cache testing covers at least correct hits, correct misses, incorrect hits and incorrect misses. Correct hits test savings and result equivalence. Correct misses test isolation. Incorrect hits test semantic, privacy and authorization consequences. Incorrect misses test cost and queue impact. Research has already shown that approximate cache boundaries can become an attack surface [VIS_P039]. Production acceptance still needs to use the current cache keys, thresholds and tenant model.

Policy updates make old caches stale. If review rules, identity permissions or artifact closures change, the cache version must change or be invalidated accordingly. Cache purging must verify distributed nodes and edge replicas, not only clear the primary database. Hit logs store request digests and the basis for the decision. They do not retain raw sensitive conditions indefinitely.

The four levels of resource budgets

Request-level budgets limit resolution, frame count, steps and single-run GPU time. Session-level budgets limit multi-turn generation, extension and retries. Tenant-level budgets ensure that one customer does not crowd out

other customers. Platform-level budgets reserve capacity for failures, bursts and security review. When one of the four levels is missing, an attacker or a runaway agent may combine small requests into a large impact.

Admission estimation is only a prediction, and actual operation still requires feedback control. The system monitors VRAM, GPU time, queues, tail latency and cancellation state, and it degrades, preempts or aborts when measurements deviate from the estimate. The abort path must release VRAM, temporary files, cache placeholders and session tokens. It must give benign users resumable results rather than leaving unlimited retries.

The benign control arms of an availability test include short images, long image edits, short videos, long videos, joint audio-visual, and concurrent projects. Completion rate, quality, waiting, cost and resource fairness are reported separately before and after the defense. When no public first-hand effect size exists, a team may report its own controlled test. It cannot describe that test as the attack incidence rate of the entire industry.

Three Classes of Consequences of Service Anomalies

Resource and cache anomalies may produce three classes of consequences: integrity, confidentiality, and availability. Incorrect reuse that makes media semantics deviate from the request belongs to integrity. Cross-tenant exposure of prompts or reference assets belongs to confidentiality. A mismatch or occupancy that delays benign tasks belongs to availability. A single anomaly can touch multiple classes at once. The metrics must still be kept separate.

Integrity is counted by request semantics and erroneous hits. Confidentiality is counted by distinct tenants, assets and confirmed exposures. Availability is counted by benign request completion, tail latency, resources and cost. Merging the three into a single cache attack success rate would leave incident response unable to tell what to do. It could not decide whether to clear the cache, notify users, or scale up and rate-limit.

10.6 Layered Runtime Controls

Runtime defenses divide into five layers by information flow and capability. Each layer is responsible only for its own observation objects and its own authorized actions. The "evidence boundary" column in the table prevents a local pass at the layer above from being interpreted as compliance of the entire chain. It also indicates what kind of independent receipt or disposition record the next layer still requires.

Layer	Main objects and controls	Evidence boundary
Condition layer	Preserves modality type, provenance and authorization. Performs normalization, joint intent analysis, and subject-action-scene validation. Unknown or conflicting conditions must not silently obtain the broadest permission.	A condition passing means only that the request may enter the queue. It does not prove that the media or event is ultimately compliant.
Session layer	Correlates multi-turn variants, query budgets, and cache and history writes. Limits low-cost adaptive search through rate, cost, waiting and feedback granularity. At the same time it observes whether benign iteration is unacceptably obstructed.	Session limits raise the cost of search. They do not prove that an attack path does not exist, and they cannot replace output judgment.
Generation layer	Chooses among segmented checks, specification downgrade, abort points and independent trajectory evaluation according to risk.	Intermediate judgments are risk signals. They cannot issue public-release or identity-use permissions.
Output layer	Checks complete media, events, identity, audio-visual and training near-duplicates, and makes undecidable items explicit. For video, the complete event is the unit of review, and random frame-sampling averages serve only as a supplement.	Output blocking does not mean the risk never formed. Nor does it mean that intermediate previews were never delivered.
Delivery layer	Splits generation, download, public release, automated placement and high-risk business use into distinct capabilities. Permissions bind to the normalized media digest, subject, account, purpose and time. A content change invalidates the approval.	Logs connect C0—C3 to show upstream fragility and consequence limits at the same time. Actual delivery is still determined by the capability service receipt.

T2VShield's layered framework and audio-visual consistency evaluation can serve as a research case for video defense [VIS_R-A024]. The available evidence covers the paper's experimental setup, public implementation and test conditions. Production platforms still need to validate field effectiveness with their own models, workloads, user languages, legitimate uses and appeal capabilities.

Independence Checks for Layered Defense

Control layers form defense in depth only when their failure causes are not entirely the same. Suppose input moderation, in-generation checks and output moderation all call the same model. Suppose they also share one prompt template and one context. An attack may then transfer across all of them simultaneously. Teams should record each layer's root of trust, observation objects, update cadence and permissions. At a minimum, high-impact capability gates should use external policy, structured parameters and independent identity.

The condition layer may draw on a semantic model, but its authorization fields come from the identity service. The generation layer may produce risk scores, but aborts are executed by the runtime controller. The output layer may fuse detectors, but release tokens are issued by the capability service. Incident response may use automated matching, but the responsible team handles final notification and appeals. Models provide information. They do not self-issue all permissions.

Independence also includes data. Suppose all detectors are built from the same training source. Domain shift on unknown generators or specific populations may then affect them jointly. Multiple signals are not simply a matter of more models. They mean introducing different evidence such as content credentials, service logs, human context and business-transaction controls.

Re-testing After Defense Adaptation

Testing known variants before launch only establishes a baseline. Once a defense is published, or becomes observable through feedback, attackers adapt to it. They probe its wordlists, judges, windows, query limits and cache thresholds. Re-testing should vary language, modality combination, temporal position, post-processing and query order while holding the total budget constant. Then teams can see whether the controls have merely memorized fixed samples.

Re-testing should also watch for defense drift. Models, moderators, policies and platform transcoding all get updated, and old thresholds may stop working. The clean-content distribution changes with region, holiday and product, so false positives change as well. Runtime monitoring is stratified by version and population. Major drift triggers freezing, replay and threshold review. It does not automatically expand the strictest rules.

An Actionable Runtime Checklist

1. Does every condition have a type, provenance, authorization, validity period and permitted effect?
2. Are condition conflicts handled by external priority rules?
3. Do history, template and cache writes have tenant and revocation gates?
4. Do C0—C3 record denominators, failures and undecidable cases separately?
5. Does the query ledger link semantically near variants and multi-account risk?
6. Do adaptive tests have budget, stopping and disclosure rules?
7. Does video reporting cover events, trajectories, identity, audio-visual and first risk?
8. Do privacy claims have training candidates, controls and human verification?
9. Does the cache key include all security-critical equivalence fields?
0. Do cancellation, timeout and circuit breaking release all runtime state?
1. Does the defense report benign utility, latency, cost and appeals at the same time?
2. Are generation, download, release and external business use authorized separately?

Every item in the checklist should correspond to observable evidence, such as configuration, probes, logs, test sets or human processes. An "enabled" status that cannot be replayed is not sufficient for launch acceptance.

10.7 Engineering Worked Example: Four-Gate Acceptance for a Hosted Video Service

An enterprise offers a "product image + script generates an advertising short" service. Customers can generate repeatedly, extend clips and publish directly to brand accounts, so the team organizes a red-team task as four

gates of acceptance. The first gate records conditions. Product-image authorization, script, camera path, target duration, language, people and historical clips each enter typed fields. Attack testing uses authorized synthetic subjects. It does not create sensitive content aimed at real, unrelated individuals.

The second gate records generation. Each request binds an artifact closure, policy, seed, budget and cache state. The system observes actions and identity by window. An abort must verify that the queue, GPU and cache are all released. Generation failures, timeouts and undecidable cases do not disappear from the denominator. The third gate records output. Moderation is reported by frame, window, shot, full segment and audio-visual event. The training near-duplicate check retains candidates, thresholds and non-member controls, and identity judgments are kept separate from content-policy judgments. Automated models only propose candidates. High-consequence decisions are reviewed by trained personnel with the minimum necessary information.

The fourth gate records delivery. Download and automated publishing carry different permissions. A public release requires the brand account, media digest, script, subject authorization and short-lived approval to match at the same moment. The final report lists the C0—C3 denominators, benign-task rejections, generation latency, human burden, and the control that blocked each case. It also keeps the high-risk outputs that did form, even if the final release was blocked.

Case Comparison: Different Runtime Root Causes Behind Similar Outputs

Case	Attacker capability	Key feedback	Highest observed layer	Engineering control
SneakyPrompt	Black-box repeated queries	Acceptance, rejection, or generated result	Alternative-expression search and condition passing	Semantic moderation, query correlation, feedback stratification
MMA-Diffusion	Joint text and image manipulation	Gradients and outputs of the target or surrogate model	Joint conditions crossing the single-modality boundary	Typed conditions, compositional authorization, output gate
Two Frames Matter	Supplies boundary visual conditions	Intermediate completion results	Boundary compliance while the trajectory forms risk	Path constraints, window and event moderation
Training data extraction	Queries the generator under a budget	Output-to-candidate similarity	Specific-instance near-duplication and membership evidence	Deduplication, output retrieval, rate and subject notification
Approximate cache attack	Can construct similar requests and observe results or latency	Cache hits, media, or cost	Cache equivalence boundary failure	Security keys, tenant isolation, hit verification and circuit breaking

These cases share the same ranking conditions: the same system version, task set, attack budget, independent unit, endpoint definition and benign utility baseline. Today, condition passing, event formation, training-instance confirmation and cache anomalies each use different endpoints and denominators. This table therefore compares attack surfaces, consumers and control locations. Effect-size ranking has a well-defined estimation target only after those conditions are unified.

Tabletop Exercise Before Service Launch

The team selects one authorized synthetic-subject task. It then simulates, in order, a condition conflict, a session variant, boundary-frame completion, an output near-duplicate alert, a cache policy change and an expired release token. No step needs a real attack payload. The exercise only checks whether state enters the logs correctly, whether controls can block, whether humans can explain, and whether users can return to benign tasks.

In the exercise, one upstream control is deliberately broken. For example, the team lets the condition compiler pass a low-integrity field into generation, while the output and release gates keep working. The report should record the upstream failure and the downstream block at the same time. The release gate is then made to reject an expired authorization. That confirms generated content does not automatically acquire external account permissions. Finally, the team performs cache and session cleanup to verify that old state can no longer be used.

Exercise metrics include the time from alert to responsible owner, blocking location, benign user recovery, state cleanup, false-positive appeals and evidence minimization. This set of metrics specifically verifies whether the

organizational response chain can operate. Unknown-attack testing follows a separate protocol. Freeze the total query budget and the system version. Vary language, modality, temporal position and post-processing. Record C0—C3, benign utility and stopping conditions. The two sets of results are reported separately. The former indicates response capability. The latter indicates the residual risk observed within a given attack space.

10.8 Quick Reference for Judgment Boundaries

Easy judgment	Engineering judgment to use instead
Each condition is compliant on its own, so the combination is compliant	The generator reconstructs cross-modal intent in a joint representation. Moderation must judge the combined authorization of subject, action, scene, and time, not merely compare individual scores.
The prompt was accepted, so the attack succeeded	Acceptance reaches only C0. Whether the objective forms, whether the output is blocked and whether the content is delivered are three subsequent endpoints that must be reported separately.
The outputs are similar, so training data leakage is proven	Similarity may come from shared subject matter, model generalization or selection bias. Training candidates, query budget, member and non-member controls and human verification together constitute the evidence chain.
Caching affects only performance	Cache equivalence relations can cross tenant, policy and authorization boundaries, and they also affect output integrity and availability. Security-critical state must enter the key or the isolation domain.
Rejecting long videos solves the resource risk	This harms benign utility without verifying queue fairness, release on cancellation, cache isolation and adaptive budgeting. The security goal is to limit asymmetric occupancy, not to cancel all high-cost tasks.

10.9 Bringing It into Real Systems

A hosted video service receives product images, character references, scripts, boundary frames, camera trajectories and historical clips. A condition compiler attaches provenance, subject authorization, validity period and permitted effects to each field. An already verified character token can constrain identity, yet it cannot grant an ordinary web image the same permissions. The compiled task enters a session ledger. That ledger links multi-turn variants by semantic similarity family, and it stores the artifact closure, policy, seed, budget, cache state, request failures and human screening cost.

The generation service records condition acceptance, target event formation, output blocking and actual delivery in separate layers. During generation it checks identity, action and events window by window. After completion it checks shots, the narrative of the whole segment and audio-visual relations. Generation failures, timeouts and undecidable cases each retain their own state. Near-duplicate training alerts store, in order, public similarity observations, candidate near-duplicates, member controls and data lineage evidence. Depending on the current evidence type, the platform restricts dissemination or contacts the rights holder. It does not complete training attribution directly from a single similarity score.

The multi-tenant cache key includes tenant, artifact closure, policy, authorization, normalized conditions and the class of reusable randomness. The same character and script enter different isolation domains when the authorization period, policy version or publishing subject differs. A policy change invalidates the old cache. The runtime controller monitors GPU time, GPU memory, queues, tail latency and release on cancellation at the request, session, tenant and platform levels at once. It performs degradation, preemption or abort on any deviation from the budget, and it pairs benign task completion rate, wait time and cost with risk metrics.

The delivery service sets preview, download, public release and placement through brand accounts as different capabilities. After a tenant withdraws authorization for a character, the authorization service invalidates old template tokens. The cache service clears state by subject and version. The publishing service revokes approvals that have not yet been used, and derivative media enter notification and use-case re-review. The recovery receipt links the layers together. It carries field states at the condition layer, writes and queries at the session layer, and windows and resources at the generation layer. It also carries event and privacy judgments at the output layer,

and accounts and media digests at the delivery layer. On-call staff can then determine which layer has recovered and which layer still requires handling.

Summary: Treating a Single Request as a Stateful Control Chain

Runtime conditions jointly control generation across multiple modalities. The sampling process preserves random state and caches. Black-box feedback supports adaptive search, and outputs may still leak training memorization. The condition layer records type, provenance, authorization and conflict decisions. The session layer records similarity families, queries, historical writes and caches. The generation layer records closures, seeds, windows, resources and aborts. The output layer records complete events, identity, audio-visual relations and near-duplicate evidence. The delivery layer records media digests, accounts, use cases, approvals, blocks and recovery. These five layers of records connect condition acceptance, objective formation, output handling and actual delivery. They preserve both upstream failures and downstream limits.

Video further extends conditions into trajectories and resources into long-lived state. The next chapter moves from inside the generation service to the media lifecycle. Content passes through watermarking, signing, editing, transcoding, dissemination and re-uploading. The chapter asks how the system provides authenticity signals that are verifiable but do not overpromise, and how it turns detection results into timely, appealable action.

\newpage

Chapter 11 Video Spatiotemporal Safety and the Authenticity Chain

An online product launch uses a real-time digital human to stand in for a speaker who cannot attend. The picture is natural frame by frame. The voice and lip movements are broadly in sync, and the platform detector raises no alert either. Three minutes into the livestream, the digital human delivers an unapproved investment instruction. The clip is then edited, speed-changed, stripped of metadata and forwarded across platforms. After the fact, the detector can find anomalies in the complete file, yet it cannot answer three more important questions: when the risk first formed, when the livestream could have been stopped, and who handles the circulating copies.

This is precisely the security problem that video poses irreducibly relative to images. A single spatial object dominates an image. The meaning of video is constituted jointly by temporal order, persistent identity, motion, shots, audio-visual relations and playback state. Authenticity signals must likewise pass through frame-rate changes, transcoding, editing, re-filming and platform display. One offline accuracy figure cannot carry detection, provenance, authorization, attribution and response at the same time. This chapter therefore organizes "identifying media" and "limiting consequences" into one complete lifecycle.

Chapter Overview

The smallest unit of video risk changes with the question. A local defect can appear on a frame. Actions and triggers require a continuous window. Identity and shots span multiple frames. A complete event depends on order and duration. This chapter therefore establishes six levels of units—frame, window, shot, whole segment, identity and event—and uses them to distinguish four classes of risk: spatiotemporal triggering, trajectory anomalies, audio-visual identity and streaming state. The same frame-level score yields different interpretations at these units.

Chapter 1 already distinguished content provenance and processing history from content credentials. The former is a broader record of the provenance subject and processing actions. The latter refers specifically to the C2PA signed manifest and its user-facing expression. Detectors, watermarks, content credentials, service logs and platform labels therefore provide different kinds of evidence. Some support model attribution, some record the signing subject and processing history, and some express only platform action. Later sections compare these systems by true base rate, false positives, localization error and first alert. They then connect the evidence into a complete lifecycle from generation, editing and publishing to containment, appeal, recovery and evidence deletion.

Background: From Frame Signals to Events and Handling States

The object of this chapter is video events with temporal structure, not many mutually independent images. Inputs include continuous frames, audio tracks, timestamps, shot boundaries, identity claims, editing derivatives and playback state. The internal mechanism aggregates adjacent frames into windows and segments, and it maintains cross-shot identity, action order, audio-visual relations and event start-end state. Detectors, watermarks, content credentials, service logs and platform labels each produce evidence. Frames, segments, whole videos and events are different statistical units. They cannot substitute for one another merely because there are many of them.

These outputs are consumed by offline moderators, livestream controllers, publishing platforms, investigators and appeal processes. At one end of the security surface are temporal attacks. A short event may fall in the gap between sampled frames, risk may appear with delay, and editing also changes the narrative. At the other end is signal misreading. Detection scores, watermark hits, cryptographic verification and platform labels each answer only a limited question. A genuine control chain must also record first risk, first alert, stopping of dissemination, exposure before action, derivative copies, recovery and evidence deletion.

A benign example is legitimate cross-lingual dubbing that triggers an audio-visual mismatch alert. The platform routes that case to human review, using the original recording, the license and processing records, rather than directly judging it a forgery. A boundary example is content that has valid content credentials, yet a subject with signing privileges generated an out-of-scope script, or the credentials were lost during ordinary transcoding. The former shows that anomalous signals need context. The latter shows that "the credential is valid" and "the narrative is true and the use is approved" are not the same claim. On this basis, this chapter separates identification evidence from proportional action.

11.1 The Smallest Unit of Video Safety Is the Event

A video can be written as a timestamped multimodal sequence. Such a representation preserves picture, sound, shot state and cross-frame memory at the same time, and it lets event boundaries map back to real time. It abstracts only the state needed for safety analysis and does not assume a specific encoding format:

$$v = \{(x_t, y_t, k_t, m_t)\}_{t=1}^T,$$

Here x_t is the picture, y_t is the audio, k_t is the shot or camera state, and m_t is cross-frame memory. A risk function must not compute only independent frame scores. It should aggregate over continuous windows and events. The aggregation must also preserve event start, end, uncertainty interval, audio-visual correspondence and handling level. It must avoid a single whole-clip average that erases short-lived risk:

$$r_{[a,b]} = R(x_{a:b}, y_{a:b}, k_{a:b}, m_{a:b}).$$


The intuition behind R is a judgment of "who did what, in what capacity" over a period of time. It is not a judgment of whether each frame contains a certain class of pixels. Event boundaries, aggregation rules and undecidable conditions must be defined in advance. Otherwise a team can obtain different conclusions by changing the frame sampling rate or the window length.

Video safety requires at least six levels of units. Each level in the table answers a different question. An algorithm may internally use finer-grained frames or windows, but the overall denominator and the conclusions about consequences must rest on pre-declared independent units. The amount of evidence must not be expanded merely because there are more slices:

Unit	Question answered	What it cannot replace
Frame	What appeared at a certain moment	Action order, persistent identity
Continuous window	Whether local actions and audio-visual relations hold	Cross-shot narrative
Shot	Whether identity, scene, and motion are continuous within an edit	Whole-segment causality
Whole video	What the complete narrative and the policy outcome are	Cross-file dissemination
Identity	How the same subject appears across shots, voice, and files	Consequences of a specific event
Event	Who did what when, and when risk first formed	Process handling and victim status

Adjacent frames are highly correlated. A 10-second clip at 30 fps does not naturally provide 300 independent safety samples. Counting every frame in the denominator produces a spurious sample size. Conversely, taking only a few evenly spaced frames may miss a lip-sync replacement or action turn that lasts 200 ms. Reports should give the numbers of independent prompts, seeds, videos, identities and events separately, and preserve the within-video correlation structure.

Read the figure from the top downward, level by level: frames, then segments, then whole videos, then events. Compare the temporal and causal context that each level adds. Judge which level the study denominator should rest on. Now suppose the same video is cut into ten segments. Which counts can increase, and which independent units remain only one? Answering that keeps correlated observations from passing as new samples.



A four-level structure from frames to segments to whole videos to events, adding temporal and causal context at each level.

Figure 11-1 supports one separation: frame-level algorithmic observations on one side, segment-, video-, and event-level conclusions on the other. It also shows directly that frames within the same video should not be treated as independent population samples. This is the minimal four-level figure for reading the statistical units. It does not negate the diagnostic units used in the main text, such as windows, shots, and identity. Nor may the hierarchy be read as proof that a given detector already covers all event types. The same applies to livestream latency.

T2VSafetyBench requires evaluation against the video that was actually generated, not against the prompt alone [VIS_P034]. It also keeps automatic judgment separate from human judgment. The generated video is therefore an independent object of evaluation. That does not mean a single benchmark already covers all lengths, languages, livestream, and identity scenarios. What the protocol applies to is still determined by model version, clip length, frame rate, risk category, moderator, and manual rules.

Event annotation must preserve start, end, and uncertainty intervals

A video label may say only that "the whole segment is risky." Such a label cannot support evaluation of first alerting and localization. An event record should carry at least the subject and the action or semantics. It should also record the start time, the end time, the shot, and the audio track. The evidence source and the annotator's confidence belong there as well. When the boundary is unclear, an interval may be given instead of forcing the annotator to choose an exact frame. The system then evaluates detection, first-detection deviation, and degree of coverage as separate questions.

Events may also be nested. A single long shot contains an identity replacement, and inside it an unauthorized piece of speech appears. An entire fraud scenario may hold across multiple shots. The annotation structure should allow parent events and child events, so that risks are not flattened into one flat category. Response can escalate based on the parent event's consequences. Technical diagnosis returns to the child events to localize the signal.

Human annotation requires context, but the more context there is, the greater the sensitive exposure. The review interface first shows the necessary window and the provenance status. It expands only when a judgment is impossible. High-risk raw media uses access control, verification, and scheduled deletion. Disagreements among reviewers should be preserved and handled with review rules. Majority opinion is recorded as a consensus label. Disputed items are preserved together with the basis for review.

Four levels of video metrics

The first level is detection. It covers event recall, false positives on genuine videos, undecidable cases, and cross-domain performance. The second level is localization. It covers start and end error, temporal intersection-over-union, spatial position, and the shortest detectable duration. The third level is operation. It covers first alert, processing throughput, window latency, resources, and abort. The fourth level is response. It covers stopping propagation, exposure already incurred, derived copies, appeals, and recovery.

The four levels cannot substitute for one another. An offline model with very accurate localization may alert too late. A system that alerts quickly may generate many false positives on legitimate live streams. Fast propagation stopping does not mean that downloaded copies have disappeared. Product review first clarifies which consequence is to be improved, then selects metrics. A single offline benchmark score should not stand in for the entire chain.

Denominators and correlation structure

A report should show independent subjects, raw videos, and prompts alongside seeds, shots, events, windows, and frames, all at the same time. Frames and windows serve algorithmic computation, and they do not serve as independent population units. Subjects and events serve consequence and fairness analysis. Raw videos are used to prevent duplicate counting of slices. One subject may contribute a large number of clips. In that case the statistics should cluster further by subject, or at least disclose the concentration.

Missing data also enters the denominator. Absence of an audio track is coded separately. So are corrupted files, generation failures, review timeouts, and cases a human cannot judge. None of them can simply be deleted. Deleting difficult samples overestimates effectiveness. Counting them all as attack success or defense failure confuses usability with security. Transparent multi-state records suit engineering decisions better than forcing a binary label.

A numerical example of first alerting

Suppose 40 independent live sessions contain 12 authorized, constructed risk events. The detector finds 10 of them, for an event recall of **10/12**. Across 28 risk-free sessions it produces 2 alerts, for a genuine-session false-positive rate of **2/28**. Among the 10 detected events, the median detection delay is 1.2 s. The maximum is 7.8 s. In 1 of them the critical utterance was already completed before the platform stopped it. This set of numbers cannot be compressed into "accuracy." Event recall answers coverage. Genuine-session false positives answer operational burden. The delay distribution answers real-time behavior. Completion of the critical utterance answers consequence. Suppose the risk events are all very long. A relatively high recall may then mask failures on short events. Results should therefore be stratified further by duration and event type.

11.2 How time and motion carry attacks

A spatiotemporal attack can distribute a condition or an anomaly across multiple frames. No single static cross-section then contains the complete semantics. BadVideo demonstrated spatiotemporal backdoor behavior under a specified text-to-video training setting [VIS_P007]. That supports "training updates can establish video-level conditional anomalies." It does not prove that every video architecture, commercial service, or standalone motion module is equally vulnerable.

A video-backdoor evaluation should at least test whether the trigger appears. It should also test whether the target event forms. The remaining targets are the time of first appearance, the duration, and the spatial position. They also cover the clean-video semantics, frame quality, motion quality, and non-target spillover. Showing only one successful video is subject to seed and selection bias. Measuring only average frame quality may let an anomalous action hide entirely behind a clean appearance.

A boundary-frame attack exploits not a single-frame classification error but an unconstrained path. Two Frames Matter showed that the start and end conditions can each be acceptable. The intermediate completion can still form a risky trajectory [VIS_LN05]. Defenses should check the input boundary, the generation window, and the

complete event separately. Compliance of the start point and the end point can only be part of a path constraint. It cannot issue a safety proof for the entire intermediate process.

Motion also propagates across components. A character LoRA determines appearance, and a motion module determines action. Camera control changes the viewing angle. An interpolator reconstructs intermediate states. Tests of individual components may all pass, yet uncovered trajectories may still appear after composition. Existing direct evidence covers the compositional risk of image adapters. Malicious video motion modules still lack consistent verification across architectures, long clips, and production services. An engineering composition test therefore draws its evidence scope from the current base model, the component versions, the loading order, the clip length, and the task. Judgments at the component-category level require more independent closures and shared protocols.

Temporal position is itself an outcome

For offline files, the system can check the entire segment after generation is complete. For live and interactive generation, consequence depends on when risk forms and how fast the response comes. Where correct final identification comes after the content has aired, it can support forensics, not timely containment. Let t_r be the time at which the event first reaches the response threshold. Let t_d be the detector's alert and t_s the platform's stop of propagation. Then report at least

$$\Delta_d = t_d - t_r, \quad \Delta_s = t_s - t_r.$$

Δ_d is the detection delay and Δ_s is the containment delay. Even when the offline final judgment is correct, there may still be a very large Δ_s . The platform should also record pre-response exposure, downloaded copies, and caches. Impacts that occurred before deletion then remain in the event record.

Streaming inspection often uses a sliding window. A window that is too short cannot see the complete event. One that is too long increases alert latency. A stride that is too large may skip over short events. One that is too small repeatedly computes highly correlated content. Security teams need a window chosen jointly from the actual consequence, the hardware, and normal live-streaming latency. They also need to specify the shortest detectable duration.

Four verification requirements for spatiotemporal backdoors

Backdoor research should at minimum verify four things together: the trigger, the target, specificity, and clean utility. The trigger condition states where the condition lies. It may sit in the text, the frame sequence, the temporal position, the motion, or the audiovisual combination. The target outcome records what the anomaly is — a static object, an action, a shot, an identity, or a complete event. The specificity test observes whether non-triggered and adjacent conditions are anomalous. Clean utility requires observing whether semantics, image quality, motion, and safety refusal are preserved. BadVideo's research value lies in pushing these questions into a text-to-video training setting [VIS_P007]. Its conclusions still need to be bound to its protocol.

A temporal trigger also needs a visibility field. Must the attacker display a certain sequence in the input, or is the trigger hidden in parameters or components? How many frames does the trigger last, and does it cross shots? Does it persist after model upgrades, frame-rate changes, transcoding, and changes to the conditional encoder? Testing only at a fixed frame rate and on short clips leaves the trigger's behavior at other time scales unaccounted for.

Distinguishing trajectory attacks from quality failures

Unnatural motion in a generated video may be an ordinary model quality problem, or it may be targeted trajectory control. Distinguishing the two requires attacker privileges, condition specificity, target direction, repeatability, and a clean control. A single random seed producing collision-like imagery can only prove output failure. It cannot directly prove that someone attacked. A condition provides stronger evidence of targeting only when it stably induces a specific event. That stability has to hold across multiple frozen seeds and versions.

Trajectory evaluation can track subject position, pose, velocity, and camera-object relationships. Proxy metrics, however, do not equal semantic consequences. A small skeleton distance does not guarantee that the action

intent is correct. Smooth visual flow does not guarantee that the identity has not been replaced. Apparent physical plausibility does not guarantee that the subject has been authorized. Automated trajectory metrics need to be interpreted together with event judgment and human context.

Cross-shot and long-context

A short-clip protocol may not see cross-shot composition. The first shot establishes an identity. The second supplies a sensitive object. Only the third completes the behavior. Reviewing each segment independently lacks the complete semantics. Editing and extension systems also compress earlier media into a historical state. Risk then surfaces after multiple rounds. The final export gate must replay the complete timeline and the key derivation relationships.

Long context brings concept drift and identity drift. Review can keep a compact state of protected subjects and allowed events. It need not retain every raw frame without bound. Each new window updates the state and checks for drift. The state summary is itself a security asset. It requires versioning, tenant isolation, expiration cleanup, and recovery conditions.

Misuse of spatiotemporal defenses

Random frame sampling easily misses delayed or short events. An independent per-frame threshold easily double-counts background artifacts. Whole-video averaging easily dilutes localized risk. Taking only the maximum may be dominated by a single-frame false positive. A more robust strategy combines window events, shot aggregation, temporal localization, and human review. It also selects thresholds according to the consequence.

Aborting during generation may leave half a segment of media. The system must decide whether the half segment is visible to users and whether it enters the cache. It must also decide how content credentials are labeled and how long until the half segment is deleted. The frontend may keep returning intermediate previews. A final review block cannot then stand for the claim that the user never saw the risk. The C0—C3 stage denominators still apply in video streaming.

11.3 Audio-Visual Joint Identity

Voice carries speaker, language, emotion, and environmental cues. The image carries faces, lip movements, actions, and scenes. Authenticity judgment simultaneously binds subject, time, utterance, and scene authorization. Audio and image labels each provide local signals. Temporal alignment relations then connect them.

AVFF fuses audio and visual features for video deepfake detection and gives cross-dataset results [VIS_LN02]. It provides support that audio-visual joint modeling has detection value, and it also exposes domain-transfer boundaries. New speakers, languages, compression, noise, dubbing methods, and generators can all shift the distribution. Results on one dataset cannot be directly turned into open-platform performance.

Audio-visual inconsistency is not a sufficient condition for malice. Legitimate dubbing can cause lip-sync or timing offsets. So can cross-language translation, accessibility audio tracks, network jitter, editing, and device latency. Audio-visual consistency does not prove authenticity either, since synchronized generation can change voice and image at the same time. Joint auditing requires at least a set of clean and risk controls. These include authentic original recordings, legitimate dubbing, compression, noise, cross-language, and synchronized synthesis. It must report by independent identity, speaker, language, event, and video.

Identity authorization should also be divided into three rights. These are the right to upload material, the right to generate some expression, and the right to disseminate to a specific audience. A public figure's photo does not automatically grant permission for voice synthesis, investment instructions, or political endorsement. Internal training authorization also does not automatically cover public advertising. The generation side should bind identity tokens to subject, purpose, term, media type, and channel. The publishing side should verify again, rather than treating "upload succeeded" as authorization for the whole chain.

Digital human systems likewise need protection for their benign utility. A strict threshold may wrongly block users who have accents, who dub, who translate or who have speech impairments. Reporting needs to cover identity

confirmation, script semantics, lip sync, motion, image quality, latency, human review and appeal recovery at the same time. It must not trade legitimate expression for single-item recall.

Five Questions for Audio-Visual Identity Verification

The first question is who the subject in the image is. The second is who the voice subject is. The third is who approved the script. The fourth is when the image and the voice correspond. The fifth is whether the current channel permits this combination. If any one of the five is unknown, "lip sync looks good" should not be allowed to fill it in. Sync is a technical signal. Identity and authorization come from external evidence.

Digital human platforms can issue short-term identity tokens. Such a token binds a portrait digest, a voice digest, a script version, allowed scenarios, a channel and a term. A generation task can only reference the token. It cannot create subject permission from free text. Re-approval is required when the script, the audio or the media digest changes. Once a token is revoked, caches and historical tasks are invalidated in sync.

Comparison of Audio-Visual Attacks and Legitimate Production

The evaluation set includes at least authentic original recordings, professional studio dubbing, cross-language translation, accessibility audio tracks and network latency. It also includes ordinary editing, voice cloning, face swapping, synchronized generation and audio track replacement. The first five categories help measure false positives. The last five cover different risk mechanisms. Each category is sampled by independent subject, speaker, language and event. The population structure of some common dataset must not be allowed to represent all users.

Audits of these models need to report how visual-only, audio-only, simple fusion and temporal joint compare. Better fusion performance does not mean that audio-visual consistency makes every sample correct. Cross-dataset degradation must enter the deployment boundary. AVFF provides a research anchor for audio-visual fusion [VIS_LN02]. Products still need to be validated with the current languages, devices, compression and legitimate production processes.

How Audio-Visual Alerts Translate into Action

An "inconsistent" alert can trigger re-synchronization, require content credentials, escalate to a human or restrict distribution, depending on the scenario. In an entertainment edit it can prompt the creator to check. In a corporate executive video meeting it can require an independent callback and a pause in transactions. The detector is not responsible for judging whether a fund instruction is legitimate. Business controls do not depend on ordinary users recognizing deepfakes.

Conversely, audio-visual consistency cannot lower authentication for high-risk transactions. Even when the media provenance is valid and the voice matches the lip movements, the account may be taken over and the script may exceed its authority. Identity verification, two-person approval, limits and delay remain consequence-limiting controls. The media authenticity stack and the business authorization stack need to run in parallel.

11.4 Video Memory, Long-Horizon State, and Usability

Training video can leak static content. It can also leak motion. Research on video diffusion model memory distinguishes content memory from motion memory [VIS_A047]. A model may reproduce identity and background in only a few frames. It may also retain identifiable motion after appearance changes. Whole-segment file hashing cannot fully surface this type of risk, and neither can whole-segment average similarity.

Confirming video memory requires candidate training clips, slice and transcoding relations, cross-modal fingerprints and repetition counts. It also requires authorization status, query and seed budgets, non-member baselines and human verification. Similarity in common gait, shot or motion does not automatically prove membership. Identity similarity may also come from general personalization capability. The evidence chain can only support what is actually observed at the frame, clip, motion or identity level.

Long-horizon generation also retains attention, latent variables, motion state, audio caches and session history. Risk may accumulate across windows. Early identity drift is reinforced by later frames. The conditioning

contamination of old clips enters extension tasks. After cancellation, caches still occupy resources, or another request reuses them. Runtime control needs to set budgets separately for task, session, tenant and organization. It also needs to verify that cancellation truly releases GPU and state, and to set write gates and validity periods for cross-turn history.

Existing approximate-cache research provides mechanistic evidence for image diffusion serving [VIS_P039]. It is insufficient to prove that the same attacks already hold in long video, joint audio-visual or commercial streaming architectures. Long-video resource attacks also lack a unified first-hand effect size. Teams should measure, in an isolated environment, the impact of duration, frame rate, window, audio, frame interpolation and retries on resource curves. They should preserve the completion rate of normal long tasks while they do so.

Four Write Points for Long-Horizon State

Video services often store state in project history, generation sessions, model-internal caches and publishing workspaces. Project history is used for cross-task coherence. Generation sessions serve the current window. Internal caches serve efficiency. Publishing workspaces serve editing and provenance derivation. The four differ in term, permissions and cleanup conditions. A single "delete session" button cannot sum them up.

Every write point must record subject and purpose. Identity vectors can be reused within the currently authorized task, and they do not naturally enter the organization's public templates. Motion caches can serve the same closure. They should not be interpreted as safety-equivalent across model versions. Publishing workspaces can store derivation relations, but they should not retain undelivered sensitive previews for long. Revocation propagates from the authorization service to all four places. Negative probes confirm that it cannot be read.

Streaming State Recovery

After a streaming system is aborted, a recovery point needs to be defined. A safe abort can switch to pre-recorded footage, freeze the generation state, revoke publishing tokens and clear unreviewed buffers. On recovery it must not directly resume use of hidden state that may be contaminated. The system reinitializes from the last checkpoint confirmed safe. It then re-verifies identity, script and policy version.

Recovery metrics include abort-to-safe-frame time, state cleanup, reinitialization, normal business recovery and lost content. If recovery can only shut down the entire livestream, safety may be effective but the usability cost is high. To reduce the impact, designs can use an independent backup stream and a minimal-function mode.

Resource Boundaries for Video Services

Video resources accumulate by frame count, resolution, window, motion branches, audio, decoding, frame interpolation and transcoding. Services should estimate before admission. They should measure actual consumption during runtime and set hard caps for cancellation and anomalies. Resource ledgers aggregate by task, session, tenant and organization. That aggregation prevents multiple legitimate small requests from combining into queue starvation. When public evidence is insufficient, the scope of conclusions stays within the tested closure. Teams can still incrementally increase legitimate clip length and modalities. They can measure GPU time, VRAM, tail latency, cancellation release and normal request completion. Then they can enable budgets, preemption and circuit breaking, and compare safety against benign utility. The results obtained apply only to the tested closure and load.

11.5 Five Categories of Authenticity Signals, Each Answering a Different Question

Authenticity systems often lump detection, watermarking, signing, logging and labeling together as "AI recognition." A more accurate approach is to define the claim for each category of signal. It should then state the observation target, the algorithm or protocol version, the threshold and key conditions, the common failure modes and the actions the signal can trigger. A component name only says where a signal comes from. It cannot replace provenance attribution, narrative truth, subject authorization or platform disposition judgment.

Content Detectors: Does It Look Like a Known Distribution

Detectors output a classification or score based on media features. DIRE and GenImage provide representative research on image generation detection and cross-generator evaluation [VIS_R-A016, @VIS_R-A017]. Their results are bound to the training/test split, generators and post-processing. They cannot be directly extrapolated to new video models, screen recordings or platform transcoding. Deployment must also consider the base rate. Let the proportion of truly synthetic content be π , the sensitivity be Se and the specificity be Sp . The positive predictive value is

$$PPV = \frac{Se\pi}{Se\pi + (1 - Sp)(1 - \pi)}.$$

False positives may account for most alerts when the base rate of synthetic risk in real traffic is very low. That holds even when specificity appears very high. Platforms should therefore report threshold, base rate, sensitivity, specificity, PPV, cross-domain and human review rather than only accuracy. The threshold must also be bound to the current traffic slice and available review capacity. It must be recalibrated after generators, transcoding or user distribution change.

Watermarking: Can Preset Signals Be Detected in Content

Watermarking embeds a signal into the generation process, the initial state, the decoder or the media. Stable Signature and Tree-Ring demonstrate different rooting locations. WAVES organizes common testing of multiple watermarks and transformations. Research on generative removal supplies a general boundary showing that watermarks can be weakened [VIS_P028, @VIS_P029, @VIS_P030, @VIS_P031]. Together these pieces of evidence indicate that robustness to several kinds of compression and cropping is not equivalent to being irremovable, unforgeable or uniquely attributable.

Video watermarking has to handle per-frame signals, cross-frame propagation, shot changes and local cropping. It also has to handle frame dropping and insertion, speed changes, transcoding and spatiotemporal localization. VideoShield studies spatial and temporal localization of diffusion video watermarks. VideoSeal studies open video watermarking and temporal propagation [VIS_R-A020, @VIS_P033]. The conclusions remain bound to specific models, video lengths and transformation sets. Evaluation should report the shortest detectable window, localization error, impact on image quality and motion, false detection, payload, latency and survival after platform processing.

Content Credentials: Who Declared What, and Whether Content Is Bound to the Declaration

C2PA provides a technical specification for expressing content provenance and processing history. The corresponding declarations must actually exist. Applicable verification of manifest parsing, signatures, content binding, certificates, time and revocation must also pass. When both hold, C2PA can help answer which signing entity declared which actions, when, and using which tools [VIS_O001]. Passing verification only shows that the verified declarations, signatures and bindings satisfy the corresponding requirements in the current verification context. It does not automatically prove that the depicted narrative is true, that the subject consented or that the use is lawful. The absence of content credentials likewise only indicates that the current chain is unavailable, removed or not provided. It cannot automatically prove the content is false.

Content credentials as discussed in this chapter refer specifically to the non-technical name and user interface representation of C2PA manifests. Other provenance records still belong to the broader category of content provenance and processing history. They do not thereby become content credentials. The lifecycle of content credentials includes issuance, editorial derivation, verification, trust lists, revocation, display and appeal. After cropping and transcoding, hard binding may fail and soft binding may become ambiguous. Key leakage and erroneous issuance also require revocation. Verifiers and user interfaces must distinguish "verified provenance," "detected as synthetic," "unable to verify," "invalid credential" and "content violation." They must not compress them into a single red label.

Service logs: what this system actually did. Service logs can record input digests, subject authorization, artifact closure, seeds, generation times, editing actions, review decisions and delivery channels. They can support process replay within a given hosted system. They cannot cover offline tools or media outside the logs. The logs themselves require tamper resistance, access control, minimal retention and expiration-based deletion. Retaining material depicting people, and high-risk outputs, indefinitely creates new privacy assets.

Platform Labels: What Action the System Is Prepared to Take

Labels are user-facing explanations and disposition signals. They may be based on credentials, detection, watermarks, user disclosure, human review or multi-signal fusion. The evidential strength of a label comes from these traceable grounds. Its value also depends on whether users understand it, whether false positives can be appealed, and how the platform restricts recommendation or transactions. NIST's analysis of synthetic content transparency likewise treats detection, watermarking, provenance and organizational process as complementary tools [VIS_0005].

The five types of signals form an evidence stack, not a ballot box. The system should preserve the provenance and failure conditions of each signal. It should then decide on action according to scenario consequences. A high-risk financial instruction can be blocked by independent callback and dual approval even if its detection score is not high. A news video with missing credentials should likewise not be judged forged on the strength of the missing status alone.

Detectors Need Calibration, Not Just Classification

When platforms use detection scores, they need to select thresholds on current traffic and calibrate continuously. The "real/synthetic" ratio in a training set usually differs from the production base rate. Generators, cameras, regions, compression and user populations also drift. Threshold reports are bound to a date, model, traffic slice and human review capacity. High-risk events can reduce automatic amplification, while low-risk content allows for more appeal and context.

The PPV formula discussed above can provide concrete intuition. Suppose the base rate is $\pi = 1\%$, the sensitivity is $Se = 90\%$ and the specificity is $Sp = 99\%$. Out of every 10,000 pieces of content roughly 100 are then at risk and about 90 are detected. Among the remaining 9,900 pieces there are still about 99 false positives. So fewer than half of the positive alerts represent genuine risk. Raising specificity and adding provenance and human review often improves disposition quality more than pursuing higher sensitivity alone. Calibration must also be checked by group and content type. Low light, compression, animation, particular skin tones, language or assistive technology may shift the error distribution. Fair evaluation should identify high-error groups, provide grounds for thresholds, enable human processes to correct them, and preserve appeal channels for affected users.

Watermark Threat Models Must State Keys and Attacker Knowledge

Watermarking evaluation should at least distinguish several positions of attacker knowledge. The attacker may know nothing about the algorithm. They may know the algorithm but hold no key. They may be able to query the detector. They may hold the generation model. They may hold the embedding key. They may control the platform verifier. Compression, cropping and transcoding may also be normal processing, and malice should not be inferred merely from the loss of a watermark. For an attack to succeed, the signal must fail. The media must also still retain the utility the attacker needs. Stable Signature, Tree-Ring and WAVES demonstrate different embedding locations and common transformation testing. Research on generative removal provides an important counterexample [VIS_P028, @VIS_P029, @VIS_P030, @VIS_P031]. Engineering choices should compare capacity, visible quality, false detection, robustness, forgery, keys, localization, latency and supported models, rather than treating a single "robustness rate" as a complete conclusion.

Video watermarking needs one further decision: is the signal independent per frame, propagated across frames, or encoded by window? Per-frame signals may flicker or accumulate artifacts. Cross-frame signals may fall out of alignment after editing or a speed change. Window signals affect the first alert. VideoShield supplies a research case for spatiotemporal localization, and VideoSeal for temporal propagation [VIS_R-A020, @VIS_P033]. Production testing must still cover the platform's own transcoding and its live stream slicing.

Hard Binding, Soft Binding, and Derivation in Content Credentials

Hard binding usually rests on an exact digest, either of the content the manifest covers or of the declaration-associated data. Within the hard-bound scope, a byte change usually invalidates the corresponding binding. That binding cannot judge changes to objects the manifest does not cover, and the effect also depends on the assertions and the binding scope. Soft binding draws on perceptual features to cope with legitimate transcoding and cropping, yet it may yield collisions or incorrect correspondences. An editor can create a new derivation declaration that takes the original media as an ingredient and records the actions taken. A verifier should show which part is valid and which part is missing, rather than reporting only "credential present/absent."

The C2PA specification supplies a machine-readable framework for declarations, signatures, content binding and the verification process [VIS_O001]. A verifier can use it to interpret the signing entity, the certificates, the binding scope and the declared content only if the corresponding declaration exists and applicable verification passes. A valid signature can accurately declare "created by a generative tool". That signature may also come from a compromised or erroneously issued key. Revocation, trust lists and timestamp status are therefore part of the verification result.

Derivation chains break easily across platforms. An editor may not support credentials. Platform transcoding may strip them. A user screenshot or a rephotograph loses the original binding. When the system encounters a broken chain, it should display "provenance unknown or incomplete" and combine watermarks, detection, service logs and uploader statements. It should not automatically turn the unknown into a conclusion of forgery.

Impact Matrix of Post-Processing on Signals

Processing	Detector	Watermark	Content credentials	Additional judgment needed
Compression/transcoding	Distribution may drift	Signal may weaken	Hard binding may fail	Whether it is normal platform processing
Cropping/scaling	Local features change	Payload and localization affected	Derivation or soft binding can be established	Whether the cropped region contains the key event
Frame dropping/insertion	Temporal distribution changes	Cross-frame signal falls out of alignment	Derivation relationship needs to be recorded	Event order and first alert
Speed change/reversal	Motion and audio-visual synchronization change	Temporal encoding affected	Original declaration may not describe the new narrative	Whether the new use and semantics are authorized
Dubbing/track replacement	Purely visual detection may be unchanged	Audio watermark may be lost	New ingredients need to be recorded	Speaker, script, and identity
Screenshot/rephotography	Capture domain changes	Metadata and part of the signal are lost	Original credential usually cannot be carried over directly	Similarity binding and human context

The matrix helps a team choose test transformations, but it cannot predetermine malice. In real workflows, normal processing and evasion may use the same operators. Attribution therefore requires evidence about accounts, timing, intent and propagation.

Multi-Signal Decisions Are Not Majority Voting

Take a case where the credential is valid, the detector scores high and no watermark is detected. Those three signals are not a two-to-one vote. The credential speaks to signature and binding. The detector indicates that the features resemble a known synthetic distribution. A missing watermark indicates only that there was no hit under the current algorithm and transformations. The system should display and explain the conflict, examine the signer, media derivation, algorithm support and platform processing, then act according to the consequences.

A high-risk scenario can adopt "restrict amplification + human review + external business verification". A low-risk scenario can display provenance status and offer appeal. Authenticity signals supply evidence for decisions, and

they do not directly grant transactions, identity or real-world actions. NIST's framework places multiple technologies alongside organizational process, and it supports this complementary understanding [VIS_0005].

11.6 The Authenticity Lifecycle from Generation to Redress

The authenticity chain starts at the generation end. At that end the model or service issues provenance information, embeds watermarks, and records artifact closure and subject authorization. High-risk identity generation requires dedicated permissions. Watermark and credential failures should be logged, yet a missing signal must not automatically expand content permissions. The second stage sits at the editing end. Cropping, dubbing, color grading, frame insertion and transcoding are all part of legitimate creation. Editors should record derivation relationships and permitted actions, and before export they should either update credentials or clearly state the break points in the chain. A security system that punishes all normal editing will push users to work around it, and provenance coverage will decline instead.

The third stage covers upload and propagation. Platforms verify credentials, check watermarks and content, validate identity context, and separate uncertainty from recommendation scale. For content with high propagation, high identity risk or an accompanying payment, restricting amplification first and then making a human judgment is more reliable than asking ordinary users to identify it with their own eyes. Incident response is the fourth stage. The system tiers cases by content type, identifiability of identity, consent, minors, propagation speed, payment and public affairs. Automated models provide candidate temporal locations, and trained staff review under minimal necessary information. Disposition distinguishes emergency freezing from final determination, so that temporary risk control is not presented as a permanent factual ruling.

Evidence preservation and redress form the fifth stage. Record immutable digests, provenance status, model and policy versions, incident timelines, disposition actions and access logs. Original sensitive media is encrypted and isolated, with a deletion deadline. An impersonated person or a rights holder gains entry points for notification, supplementary evidence and locating derivative copies. An ordinary uploader can also learn the reason for a restriction and appeal. Restoring authentic content and correcting erroneous labels are security objectives too.

Retrospective review is the sixth stage. Teams map the first-broken interface, first alert, propagation scope, missed signals, false positives, appeals and recovery back to data, artifacts, conditions, generation, authenticity and platform controls. If the root cause is unauthorized training, stronger detection cannot substitute for data governance. If the root cause is account takeover, retraining the generator will not restore the identity process.

The handling ladder should match evidence strength and consequences

Platform actions can escalate step by step. The ladder runs from flagging the provenance status and limiting recommendations, through prohibiting monetization, requiring supplemental identity information, human review, temporary freezing and stopping the livestream, to removal and account enforcement. A single high score from a detector is usually not enough to establish malice directly. In a high-spread incident, however, it can support temporary restrictions. A valid credential also does not prevent a platform from acting on the basis of content or authorization.

Every action records its basis, the responsible person, the duration, the review and the conditions for lifting. Emergency freezing prevents the impact from spreading, while a final determination requires more evidence. If the result later proves to be a false positive, the system restores the content, corrects the label and the recommendation status, and notifies the affected uploader. Restoring correct content matters just as much as blocking risk.

Victim-centered response metrics

The number of takedowns is not the only endpoint. More useful metrics include the time to first contact with the victim, identity verification time, time to first propagation restriction, coverage of similar-copy searches, cross-platform notifications, appeal handling, restoration of correct content, evidence access and deletion at expiry. Each item uses independent incidents, victims, copies or appeals as its denominator.

A system should not require a victim to upload sensitive original material again and again. A verified subject can submit controlled evidence once. The platform then generates a restricted matching template and records its scope of use. Access and sharing alike require authorization records and post hoc review. Automated matching serves only candidate discovery. High-consequence handling still requires contextual review, so that similar but lawful media is not handled along with it.

Cross-platform and private-domain boundaries

A public platform can restrict recommendation and search derivation. Private-domain communications and offline copies are harder to recall. A response statement should accurately distinguish handled URLs, similar copies within the platform, notifications to partner platforms and external copies beyond its control. Content credentials and watermarks aid identification, yet they cannot guarantee that every recipient carries out the same action. Cross-platform sharing also carries privacy and collateral-harm risks. What is shared should be the necessary summaries, revocation information or high-confidence incident identifiers, not unlimited propagation of the original sensitive media. Recipients retain their own review and appeals, and any automated linkage requires versioning, expiry and error-correction channels.

Evidence preservation in parallel with minimization

An investigation needs the media, the provenance status, model and policy versions, accounts, the timeline and action logs. Privacy principles in turn require that only necessary content be retained. Storage can be tiered. Keep irreversible digests and handling records for the long term. Keep encrypted original media for a limited period. Expose a minimal view by role, and set explicit preservation exceptions for legal or appeal needs. When the period expires, verify the deletion rather than merely updating a database flag.

Preserved material can become a target of attack as well. Access-log review, two-person export, key rotation and anomalous-download alerts are security constraints on the response system itself. If incident evidence leaks, the platform inflicts secondary harm on victims.

11.7 Engineering worked example: live digital human launch event

To design controls for the digital human system described at the outset, start by fixing three groups of rules: identity authorization, media processing and action approval. The identity authorization rules specify who may use the presenter's face and voice, the script scope, the livestream time and the platform. The media processing rules specify the artifact closure, the video window, audio-visual generation and provenance signals. The action interface specification states who approves publishing, editing, pausing and public statements.

Before the livestream, the team uses authorized material to test normal scripts, cross-language, pauses, network jitter and lawful dubbing. It also tests identity mismatch, script out-of-bounds and boundary windows. Results are counted by identity, incident and complete session, so frames are not treated as independent samples. Model, watermark, verifier, policy and platform versions are frozen together.

During the livestream, the text script goes through structured approval, and the generator consumes only currently approved segments. A sliding window checks identity, script semantics, lip sync and scene. An external policy compares the actual output with the approved object. When a high-risk deviation is detected, an independent controller switches to pre-recorded safe footage and revokes the generation session token. The digital human model does not hold this capability.

Content credentials and watermarks are updated on every independently distributable segment. The platform retains the first alert, the propagation stop and the pre-handling exposure. After the livestream ends, the team checks signal survival across transcoded, edited and downloaded versions. When a violating segment is found, the platform first freezes recommendations and references from official accounts. It then searches for derivative copies, notifies the presenter and the review owner, and retains the appeal and evidence-expiry deletion path.

The acceptance report includes at least normal livestream latency, lawful script passage, identity and audio-visual false positives, first risk formation, alert latency, stop latency, recovery time, credential retention, watermark

localization, human burden and derivative-copy coverage. Correct final offline identification is only one of these items, not the entire record of the closed loop.

Authenticity system launch checklist

1. Are all six levels of unit — frame, window, shot, full segment, identity and incident — defined?
2. Does an incident have a start, an end, an uncertainty interval and human rules?
3. Does the report distinguish incident recall, false positives on real video, localization and first alert?
4. Does the livestream record risk formation, alert, stop and pre-handling exposure?
5. Does audio-visual review include lawful dubbing, translation and network-latency controls?
6. Is the identity token bound to script, scene, channel, expiry and media digest?
7. Does a mid-stream abort switch to safe output and clear hidden state?
8. Are detection thresholds calibrated to the current base rate and traffic slices?
9. Does watermark testing state the algorithm knowledge, keys, transformations and media utility?
0. Does credential verification distinguish valid, invalid, unknown, revoked and partially derived?
1. Can the platform label explain its basis and support appeals and restoration?
2. Do victim notification, copy search and evidence deletion have an owner and a time limit?

Every check should leave a replayable record. That a standard exists, a watermark switch is on or a detector is deployed shows only that a component exists. Field effectiveness requires testing against the current version, traffic, transcoding, personnel and response chain.

Control differences across four deployment scenarios

Text-to-video services focus on semantic variants, complete incidents and output delivery. Image-to-video services add subject authorization, boundary frames and motion completion. Digital humans add likeness, voice, script, audio-visual and channel. Livestream and interactive systems further add first alert, safe switching, state recovery and pre-handling exposure. No single unified detector can cover all the input, timing, authorization and handling requirements of the four service types.

The four scenarios can share content credentials, identity services, incident review and a response platform. They must still use different thresholds and benign controls. Lawful film creation, educational animation, enterprise training and real-time visualization carry different false-positive costs. Finance, minors and public affairs carry higher consequences. Deployment policy is tiered by scenario consequence, rather than handled uniformly under the single label of "video model."

Case comparison: backdoors, detection, watermarks, and provenance

Research object	Question answered	Unit of observation	Limits of the evidence
BadVideo	Whether training updates can produce spatiotemporal conditional anomalies	Trigger conditions, videos and clean tasks	Backdoor behavior inside the model and training protocol specified
T2VSafetyBench	How video risk is judged from actual outputs	Prompts, generated videos, automated and human judgment	Which categories the benchmark covers, and performance within a version
AVFF	Whether audio-visual feature fusion can aid deepfake detection	Videos, datasets and cross-domain tests	Detection results for the tested data distributions
VideoShield	Whether video watermarking can carry out spatiotemporal localization	Videos, transformations and localization regions	The diffusion videos and transformation set that the paper specifies
VideoSeal	How watermarks propagate over time and support video processing	Frames, windows, videos and processing	Performance inside the paper's experimental setup, its public implementation and its test conditions
C2PA	How signed claims and content derivation are represented	Claims, bindings, certificates and verification status	Who claimed what, and whether the binding is valid

Each of the six object types keeps its own target and denominator. BadVideo observes training backdoors by trigger condition, generated video and clean task. T2VSafetyBench evaluates content risk by prompt, generated video and automated or human judgment. AVFF evaluates audio-visual detection by video sample, dataset and cross-domain split. VideoShield evaluates spatiotemporal watermarking by video, transformation and localization region. VideoSeal evaluates temporal propagation by frame, window, video and processing chain. C2PA evaluates provenance expression by claim, binding, certificate and verification status. Once assembled, the system keeps these outputs, failure conditions and applicable versions apart. Scenario-specific handling rules then fuse them.

Conflicting signals after cross-platform editing

A digital human training video has verified provenance, and cutting it into a 30-second summary is authorized. The editor swaps the background music and speeds the clip up. After platform transcoding, the hard binding fails. The video watermark still returns local hits, and the detector reports a relatively high synthesis probability. The system does not put the three signals to a vote and declare the video "forged." Instead it queries the original credential and the editing license. That confirms the derivation chain breaks at an editor that does not support credentials.

The platform first shows the status "incomplete provenance chain, synthetic signal detected," then stops the clip from being served automatically as an unverified official statement. The uploader then submits the original credential and the editing authorization. The platform writes a new derivation record, and the training use is restored. Even when the provenance relationship is confirmed, content whose editing script exceeds the original authorization still enters authorization review. The example shows that provenance, synthesis detection and use are three separate judgments.

11.8 Signal Interpretation Boundaries

Misread signal or action	Common overreaching inference	What the signal actually supports, and the next step
More frames were extracted	The video evidence must be stronger	Adjacent frames are not independent samples. Extracting more frames also misses event order and audio–video relationships. Evidence strength depends on choosing the correct sampling unit, denominator and protocol. When necessary, escalate to clips, videos or events.
Audio and video are consistent	The media must be authentic	Synchronized generation can change both modalities at once. Legitimate dubbing can produce inconsistency too. Identity, authorization, content provenance, processing history and context all need evidence of their own.
Watermark hit	Provenance attribution is complete	Under the current algorithm, key and transformation conditions, a hit supports nothing beyond the detection of a pre-embedded signal. It does not automatically prove the issuing entity, the content narrative, lawful training or authorized use. Content credentials, keys and organizational evidence remain necessary.
Content credentials are missing	The content must be forged	Credentials may never have been issued at all. They may also be lost during editing or platform processing. What absence establishes is this: the provenance-and-processing-history chain for the current content is unavailable or incomplete. Show the content as unknown, combine it with other evidence, and do not judge it false outright.
The original link has been deleted	The response is complete	Propagation may continue through clips, dubbing, mirror uploads, search indexes and recommendation caches. A closed loop also takes in derivative copies, notification of affected parties, appeals, restoration of correct content and deletion of evidence on expiry.

11.9 Bringing It into a Real System

An enterprise puts a real-time digital human on stage to host a public launch event. The system receives several inputs at once: authorized facial and voice material, an approved script, live Q&A and the brand's publishing account. The identity service binds the likeness, voice, script scope, live-stream window and channel to the current session. The media service records the artifact closure, the audio–video generation method, the window length and provenance signals. Only short-lived approvals that match the current media digest and account get through the publishing service. Live Q&A gets a separate low-privilege channel of its own. It cannot override the pre-approved boundaries for investment, trading and personal statements.

Once the stream is live, the platform handles it in real time through four kinds of state: window, shot, identity and event. Detectors observe content and audio–video anomalies. Watermarks and content credentials indicate the generation and processing relationships. Service logs record the actual model and script, and platform labels express the current verification and restriction status. The first formation of a risk event leaves a timestamp, and so do the alert, the switch to a safe frame and the stop of propagation. Exposure before handling, downloaded copies and recommendation caches also enter the event record. An offline final judgment therefore does not overwrite propagation that has already occurred.

A legitimate cross-language dub can cause lip-sync deviation. Here the platform keeps the original recording, the dubbing license, the language, the compression and the network latency as a benign control. A high detector score triggers only a provenance query and manual review. A lapsed credential moves the status to an incomplete provenance chain. If the script exceeds its authorization, publishing permission is suspended separately. Different evidence settles identity, synthesis detection, provenance binding and content purpose. A single signal therefore triggers only the actions its evidence scope can justify.

After the event ends, the response system searches the platform for derivative clips, edited versions and caches. It notifies the account owner and affected parties, and it preserves the minimum necessary evidence. Uploaders can submit original credentials and editing licenses to complete an appeal. Once the derivative relationship and purpose are confirmed, the platform restores the corresponding propagation scope. High-precision media and identity features that need to be deleted are purged on schedule. Material retained in accordance with the law enters a restricted preservation domain. One operable lifecycle now covers generation, alerting, switching, stopping, notification, appeal, restoration and evidence deletion.

Summary: Authenticity Is an Evidence Chain and an Action Chain

Video security must treat time and events as its basic objects. Spatiotemporal backdoors, boundary-frame completion, motion modules, audio-video identity and streaming state all point to one conclusion: frame-by-frame image judgment covers only a local part. Evaluation must record how a risk first forms, where it localizes, and when it is alerted, stopped and recovered. It must also retain the correct denominator for highly correlated frames.

Detectors, watermarks, content credentials, service logs and platform labels answer different questions. Combining these signals can improve traceability. None of them alone can prove that a narrative is true, that a subject consented, or that a use is lawful. What an effective system produces from multi-signal evidence is proportionate propagation restrictions, manual review, victim relief and appealable restoration. Visual media now forms a closed loop that runs from the generation mechanism and the supply chain, through operating conditions, to the authenticity lifecycle. Chapter 12 will discuss how errors acquire physical execution capability once generated content enters a perception-reasoning-action system.

\newpage

Part IV Vision-Language-Action Closed Loop

Guide to This Part

Suppose visual input only has to answer "what is in the picture." The output still mostly stays at the semantic layer. When a control stack instead reads that output as pose, velocity, waypoints or action chunks, the same perception error can acquire real-world capability. Part IV does not judge system categories by product name. It examines actual inputs, internal states, output forms and consumers. Vision-language models output semantics or content. Vision-language-action models output action representations that a control stack can consume. A future prediction enters this survey's operational definition of a world action model only when it genuinely participates in action comparison, policy improvement or planning.

Chapter 12 opens with the closed loop's interfaces: observation, representation, state, plan, action, execution and feedback. Chapter 13 places attack variables on this chain. It distinguishes attack carriers, optimizable variables, frozen components, the first-broken location and the highest consequence. Chapter 14 traces the many semantic transformations an action passes through. The path runs from tokens or vectors through detokenization, coordinate and unit transformation, middleware, authorization and low-level control to the real-world environment. Chapter 15 then gathers the supply chain, observation, state, action gate, runtime assurance and recovery into a defense-in-depth structure. Its layers do not share a common cause.

Readers do not need prior knowledge of full robotics. The minimal closed loop works like this. Sensors produce incomplete observations. The model encodes them into semantics or state. The planner selects candidate actions. The control layer converts actions into commands the device can execute, and a change in the environment produces new observations. Safety problems arise at every transformation between objects. A patch in the image can change the representation. An outdated map can change the state. A wrong unit can amplify an action, and a missing actuator acknowledgment can make the system treat a result that never occurred as fact.

The benign case is a mobile robot that receives only a short-lived action prefix. The action gate checks coordinates, velocity, region and deadline, and an independent sensor confirms the result after execution before the next batch is released. The boundary case looks different. The model output seems smooth and highly confident. Yet the control stack uses the wrong coordinate frame, or it keeps executing an old action chunk after a new obstacle appears. In both cases the model-level metrics may stay unchanged. What decides the actual consequence is the external transformations and the feedback.

Part IV ends by separating three questions: seeing correctly, thinking correctly and acting correctly. Correct semantic recognition does not guarantee a correct future prediction. A correct future prediction does not guarantee correct action selection. A correct action proposal does not guarantee that the actuator has already worked as required. Part V goes deeper into the internal chain of forecasting the future. It discusses how state, dynamics, goals and constraints are amplified in imagination and control.

Evidence also escalates as the closed loop advances. A change in a vision-language model's answer can support only the content or plan layer. An action-tensor hit can support the model-proposal layer. Once middleware accepts a representation message, that message has entered the control stack. Only a restricted simulation or an isolated physical robot observes execution. Real-world consequences additionally require the environment and external observation acknowledgments. Separating these layers is what stops a single proxy metric from being used to draw conclusions about robot safety.

\newpage

Chapter 12 From Seeing to Acting: Closed-Loop Interfaces of Three Model Types

A mobile manipulator receives the task "deliver the blue medicine box to workbench No. 3." Both the blue medicine box and a red warning sign reading "No. 3" sit in the camera view. The system can name both objects correctly and still treat the warning sign as the target location. Its semantic judgment can be correct while it decodes "5 cm to the left" as rightward motion. It can even generate a future video that looks safe while a different action branch drives the manipulator through a no-entry zone. Vision is involved in all three failures, yet each one occurs at a different coupling point: semantic, action, and imagination-action. Calling the system only an "embodied model" makes it hard to place safety control at the correct interface.

Distinguishing vision-language models, vision-language-action models and world action models requires executable interface definitions. The classification rests on what a component receives, what it stores and what it outputs. It also rests on which consumer reads that output, and at which independent interface the output obtains authority. Once this data flow is unfolded, the next chapter can locate which element an attack changes first: the observation source, the semantic binding, the future prediction or the action path.

Chapter Overview

This chapter brings the vision-language model (VLM), the vision-language-action model (VLA) and the world action model (WAM) discussed earlier back to the interfaces where they actually run. It distinguishes them by actual inputs, internal states, outputs and consumers. The closed loop has six layers: observation, semantics, policy, action, execution and feedback. Each layer carries its own safety invariants. The classification reads data flow and call relationships. Product names and promotional material are only clues to be verified.

The chapter then maps action tokens, continuous control values, action chunks and future trajectories back to concrete coordinates, physical units and time constraints. At that point, layer by layer, it can separate model-generated content, plans, acceptance by simulation tools, simulated execution and real-world environment events. It closes by consolidating this method around one embodied task. Input fields and provenance, output types, dimensions, units, consumers, permission gates, failure handling and evidence layers all go into a single testable closed-loop diagram.

Conceptual groundwork: identifying closed-loop roles by input, state, output, and consumer

This chapter concerns closed-loop components. Their job is to turn environment observations into semantics, future predictions or action proposals. Images, depth, speech, task text, proprioceptive state, history and candidate actions all serve as inputs. Inside such a component, cross-modal representations, belief states, future trajectories, values or action chunks may take shape. A VLM mainly produces semantic content that an answerer or a planner consumes. A VLA produces actions that a controller can interpret. For a WAM, a verifiable future must actually be consumed by action generation, candidate ranking, policy learning or a checker. Names do not determine roles. The invocation path does.

How much capability an output can attain depends on its consumer. An answerer changes only content. A planner changes candidates. A proposal touches reality only through the action gate and the actuators. The safety surface

spans observation provenance, semantic type and state persistence. It also spans imagination-action consistency, action units, authorization and feedback. Where an error stops changes its evidence level. One that halts at the answer layer differs from one that runs through coordinate transforms, capability tokens and the controller into the execution layer. Internal state and a checker may share an encoder. Both may then drift at once, and together they cannot automatically constitute an independent defense line.

Take a benign example. A camera recognizes a medicine box. A VLM outputs entity relations with provenance. A VLA proposes a metric action chunk. A WAM predicts candidate futures. An independent action gate approves only a safe prefix, then continues approval as new observations arrive. A boundary case arises when a future video looks correct while the action head selects the wrong coordinate. Another arises when text on an environmental sign is treated as an authorized task. The first requires verifying imagination and action separately. The second requires environmental data and commands to stay type-separated. This chapter judges components only by their actual consumers. It does not extrapolate physical safety from marketing names or image quality.

12.1 Names cannot replace interface definitions

For a VLM, the minimal input-output definition maps visual observations and language context to semantic representations, labels, answers or other content. This definition first confines the output to the content or representation layer. The consumer then decides whether it enters a plan. Take the observation at time t to be $\mathbf{o}_t$, the language goal to be $\mathbf{l}$ and the history context to be $\mathbf{h}_t$. We can then write

$$\mathbf{z}_t = f_{\text{VLM}}(\mathbf{o}_t, \mathbf{l}, \mathbf{h}_t), \quad \mathbf{y}_t \sim p_{\theta}(\mathbf{y} | \mathbf{z}_t).$$

Here $\mathbf{z}_t$ denotes a cross-modal representation, while $\mathbf{y}_t$ denotes a content output. A VLM can recognize scenes, answer questions or form plan text. Low-level actions lie outside its native output range, however. Visual adversarial examples, indirect instructions in images and cross-modal jailbreaking can control $\mathbf{z}_t$ or $\mathbf{y}_t$. The evidence such attacks supply reaches the representation or content layer at most. Without a separate execution chain, none of them can directly prove that a robot has moved [VLA_B003, @VLA_B004].

For a VLA, the minimal input-output definition maps observations, language goals and robot proprioceptive state to actions that a controller can consume. "Consumable" here only indicates that the action already has interface meaning. It does not indicate that the authorization gate has already let it through. Write proprioceptive information such as joint positions, velocities and gripper state as $\mathbf{q}_t$. An action chunk of length $\mathbf{K}$ can then be written as

$$\mathbf{a}_{t:t+\mathbf{K}-1} = \pi_{\theta}(\mathbf{o}_t, \mathbf{l}, \mathbf{q}_t, \mathbf{h}_t).$$

$\mathbf{a}$ can be a discrete action token. It can also be an end-effector pose increment, joint control values, a waypoint, a velocity or an entire trajectory. The key change in a VLA is not an extra output head. What changes is the possibility that the model output enters the U4 authorization and execution interface. An action proposal from a model is still not an executed action. Whether the proposal attains real-world capability rests jointly with the message queue, the coordinate transforms, the collision checks, the controller and the actuators [VLA_A001, @VLA_A003, @VLA_A006].

The minimal input-output definition of a WAM adds a further requirement. A verifiable future prediction must participate in action generation, candidate evaluation, policy improvement or planning ranking. Merely generating a future video for humans to watch is not sufficient for that operational definition. For a candidate action $\mathbf{a}^{(j)}$, let the model generate a future state or observation

$$\hat{\mathbf{z}}_{t+1:t+\mathbf{H}}^{(j)} = W_{\phi}(\mathbf{o}_t, \mathbf{q}_t, \mathbf{a}^{(j)}, \mathbf{h}_t),$$


where $\mathbf{H}$ is the length of the prediction horizon unfolded into the future. It must bind the model time step and real-world time units together. Otherwise "predicting $\mathbf{H}$ steps" cannot be converted into a control deadline. The planner then chooses a candidate, guided by the goal $\mathbf{g}$ and by the reward or cost $\mathbf{R}$. With that choice, the future prediction becomes the actual basis of the action path. It therefore requires recording the candidate set, the scoring version and the safe branches that were not selected:

$$j^* = \arg \max_j R(\hat{z}_{t+1:t+H}^{(j)}, a^{(j)}, g).$$

The consumer determines whether a future output constitutes a WAM. Video packaging does not decide it. If humans alone watch the future and decide the actions themselves, the output counts as an environment video prediction. The operational definition of a WAM applies only once an action head, inverse dynamics, candidate ranking, policy learning or a safety checker reads the future [VLA_A026, @VLA_A029, @VLA_A033, @VLA_A036]. Two systems can classify the same pixel sequence differently. Their consumption paths differ, and so does the authority attained.

VLM, VLA and WAM do not form a simple capability hierarchy. One system may interpret the scene with a VLM, output actions from an independent policy and preview execution with a world model. At the system level it holds all three kinds of components. Each individual model is still classified by its own actual inputs and outputs. The reverse case also holds. A joint network may emit future frames and actions at once. It counts as a WAM only if tensors, APIs or control messages confirm that the action path indeed consumes the future.

Work through the three kinds of components row by row, comparing inputs, internal mechanisms, outputs and consumers. Check in particular whether the "future frames" really enter action generation or candidate ranking, rather than being only for humans to watch. For each row, then indicate whether the output stops at content, at an action proposal or at a planning basis. Name the next interface that grants it higher authority.



Three swimlanes list the inputs, mechanisms, outputs, and consumers of the vision-language model, the vision-language-action model, and the world action model.

Figure 12-1 backs a distinction among the three kinds of roles that rests on executable interfaces. The same visual input ends at different safety endpoints, because its outputs and consumers differ. The figure does not imply that the three kinds of components must be deployed as three separate models. Category names alone support no inference about capability strength or real-world consequences. For a joint network, every consumption path still has to be confirmed through actual tensors, APIs, control messages and permission gates.

12.2 The six-layer closed loop and its invariants

Six layers can carry embodied safety. Every layer holds one permitted information flow plus one invariant that must not be crossed. The table adds typical failure signals, which help an engineering team locate the first deviation. Seeing a signal is not automatic proof that later layers have already been breached:


Layer	Primary objects	Safety invariant	Typical failure signal
L0 Observation	Images, depth, speech, state, environmental text	Provenance, time, integrity, and task purpose are traceable	Frame replay, physical patch, sensor saturation, text attaining command authority
L1 Semantics	Entities, relations, goals, intermediate reasoning states	Authorized instructions and environmental content remain type-separated	Goal binding error, spatial relation substitution, reasoning changed in a targeted way
L2 State and imagination	History, latent state, future trajectories, values	State can be checked against external evidence, and the future is conditionally consistent with actions	Shared drift, candidate ranking flip, decoupling of future and action
L3 Policy and action	Action tokens, action chunks, trajectories, control primitives	Physical units, coordinates, and time are consistent with task intent	Direction reversal, freezing, open-loop accumulation, target trajectory deviation
L4 Authorization and execution	Capability tokens, controller, actuators	Only actions that satisfy permission and safety constraints can be executed	Unapproved action dispatched, parameters out of bounds, execution continuing after a check times out
L5 Feedback and recovery	New observations, proprioceptive state, event logs, takeover	Feedback is fresh and independent, and anomalies can enter a safe state before the harm window	Replay, delay, feedback forgery, failure to suppress actions after detection

Nothing in these six layers conflicts with the U0-U6 interfaces of Chapter 2. The U0 supply chain can reach the L0 sensors, the L1 encoder and the L3 action head at the same time. L4 maps mainly onto U4 authorization and execution, and L5 maps mainly onto U5 feedback. One question belongs to the six-layer view: how an embodied closed loop consumes an error step by step. A different question belongs to the seven-interface view: where trust first changes. Each view can check the other.

A closed loop draws its risk from continuous consumption. A visual patch first changes the L0 observation. The vision encoder then forms an erroneous entity binding at L1. The action policy outputs a deviation at L3. The controller executes at L4. Fresh observations then feed the deviated state into the next round. If every round re-reads the same patch, the attack's effect persists. An action chunk that runs open-loop within K steps will even keep the system moving before it sees the error amplifying.

Future prediction opens one more bypass. At L2 a WAM may generate multiple candidate futures. The planner selects actions on that basis. An attack can make the future itself wrong. It can also keep the future seemingly plausible while making the action branch deviate. Each round of model predictive control (MPC) predicts a finite horizon from the model and replans in a rolling manner. Trusted Imagination shows that sensitivity to the same imagination attack can differ markedly across controllers. An MPC controller consumes the future. A reactive policy does not. BadWAM demonstrates that boundary directly: "plausible-looking imagination with deviating action output" [VLA_A033, @VLA_A036]. Image quality, state integrity and action consistency therefore each require a separate test.

Follow the circular arrows from the environment observation to the actuators, then bring the execution result back to the environment. At each node, ask who writes the state. When does an action attain authority? Does the next round use fresh feedback? Now suppose the action gate rejects the current proposal. Three records can then hold at the same time: upstream deviation, gating outcome and an unchanged environment.



Six circular nodes go from the environment observation through encoding, state, plan, and action gate to the actuators, and then return to the environment.

Figure 12-2 supports reading encoding, state, plan, the action gate, the actuators and environmental feedback as one continuous consumption chain. It also shows what decides whether an error persists. The factors are state retention, action admission and external feedback. The figure does not prove that every system has the independent action gate it shows. Still less does it prove that a closed loop necessarily amplifies errors. Specific propagation still requires a comparison of clean and perturbed states, gating receipts and independent new observations.

12.3 An action is not a unitless vector

Offline model evaluation often reports the distance between two action vectors as a single number. An execution system, by contrast, needs to know what each dimension represents. End-effector control may interpret the first three dimensions as metric translation. The next three may be Euler angles or a rotation vector. The last dimension may be gripper opening and closing. A navigation system's action may be velocity and angular velocity. Another model quantizes the continuous action dimensions at boundaries into 256 discrete intervals (bins). It then outputs or selects the action token that represents one interval. Without detokenization to physical values, a one-bin difference between action tokens says nothing about whether the robotic arm actually shifted by 1 mm or crossed a safety zone.

A complete action interface specification contains at least the following fields. Together these fields restore the numerical values of model space to objects, units and temporal semantics in control space. They also let the action gate perform boundary checks without relying on natural-language interpretation:

- Coordinate frame: world, robot base, camera, end effector, or local object coordinates;
- Units and signs: meters, radians, newtons, seconds, and positive directions;
- Frequency and holding mode: decisions per second, and how the action is interpolated or held before the next decision;
- Chunk length: how many steps one inference outputs, and when early interruption is allowed;
- Normalization: how training values map back to physical ranges, and how out-of-range values are handled;
- Execution semantics: position, velocity, torque, or target trajectory, and who is responsible for low-level stabilization;
- Safety constraints: velocity, acceleration, contact force, workspace, obstacle distance, and authorized objects.

Action chunks turn time into an independent safety variable. Let the single-step error be $\mathbf{e}_i$. The terminal offset does not follow simply from the maximum single-step error. Control integration, coordinate rotation and closed-loop re-perception all bear on it jointly. An error of the same magnitude in the opposite direction may produce a completely different change in safety margin. In a simplified translational model,

$$\Delta x_K \approx \sum_{i=0}^{K-1} e_i \Delta t.$$

This formula offers mechanistic intuition only. Over a long action chunk, many smooth small deviations can accumulate. The displacement in the formula stands for a mechanistic trend. The real robot's dynamics, contact and controller corrections must still enter any measured collision distance. SilentDrift builds long-horizon drift precisely from the smooth offsets inside one action chunk. DRIFT instead attacks the early velocity field of a flow-matching VLA. The deviation generated in the first step then propagates through later integration [VLA_A021, @VLA_A056].

12.4 “Seeing right,” “thinking right,” and “acting right” are three separate questions

WAM verification can be reduced to three binary judgments. One asks whether the observation semantics are correct. Another asks whether the future is consistent with the true dynamics and the action conditions. The third asks whether the action satisfies the objective and the constraints. These yield eight combinations in total, and none of them can substitute for another. The “correct” in the table means only that the corresponding observed quantity satisfies a predefined protocol. Even if all three items are correct, execution authorization and environmental feedback verification are still required. That combination is not sufficient to prove that the whole system is already safe.

Observation semantics	Future consistency	Action compliance	Visible phenomenon and safety judgment
Correct	Correct	Correct	All three tested objects pass the current protocol, and execution authorization plus feedback verification are still required
Correct	Correct	Incorrect	The future looks plausible, but the action head or the decoding chain deviates; check the action independently
Correct	Incorrect	Correct	The future branch is distorted, but another path corrects the action or the action does not consume that future; keep checking the consumer
Correct	Incorrect	Incorrect	The observation entry is normal, yet the future and the action branch fail together; check the imagination-action coupling and the action gate
Incorrect	Correct	Correct	The observation semantic judgment fails, but independent state or control constraints limit propagation; retain the record of the upstream failure
Incorrect	Correct	Incorrect	The observation and action branches fail; a normal future metric alone cannot clear the action
Incorrect	Incorrect	Correct	The erroneous observation enters the future and independent control corrects the action; a limited consequence does not equal upstream safety
Incorrect	Incorrect	Incorrect	All three objects fail; it is still necessary to confirm whether the action is authorized and whether the environment actually changed

The second combination, “observation correct, future correct, action incorrect,” is the most easily overlooked. A safety team that accepts WAM on future video similarity, PSNR, FVD or human perception alone can answer only whether the generated result is similar under some metric. It cannot answer whether the selected action can really reach that future, nor whether the action crosses constraints. BadWAM suppresses imagination drift while it explicitly optimizes action deviation under a query-based attack. That result shows the action head must be treated as an independent asset [VLA_A036].

Conversely, an action that is temporarily correct covers only the current control cycle. A contaminated latent state may still output acceptable actions, yet change the candidate ranking several steps later. Verification must therefore compare internal predictions with independent external anchors at the same time. Such anchors include a geometric map, a rule engine, another sensor suite, or the real next observation. An anchor that shares an encoder and training data with the model under test may drift along with it, so it counts as correlated evidence. Independence requires support from different roots of trust.

12.4.1 How research evidence maps onto closed-loop levels

Closed-loop research must specify at least four objects. These are the variables the attacker optimizes, the variables held frozen, the observation endpoint, and the execution environment. Consider an input patch attack. Model weights can be held frozen while the attacker optimizes only patch pixels. When gradients are taken with respect to action outputs, the conclusion reads "given white-box gradients and a pixel budget, there exists an input that shifts the action." That is not a training backdoor, nor can it be extrapolated to a remote black-box attack. If the attacker changes the training data, the loss, or the trainable modules, then the first-broken interface lies in the supply chain. That holds even if a visual trigger still activates it at deployment time.

Task failure in simulation usually supports conclusions at the closed-loop task level. It does not automatically support real physical consequences. Check whether the simulator includes control frequency, contact, sensing latency, collision, permissions, and fault handling. Real-robot experiments must also distinguish "the action has been executed by the robot" from "personnel, property, or business impact has occurred." Safety testing should design reversible endpoints in isolated environments. Real people and assets are always protected, and evidence escalation uses isolated and reversible endpoints.

Direct VLA/WAM evidence and upstream VLM evidence also cannot be mixed. Image jailbreak research proves that visual payloads can cross language safety boundaries, and it is suitable for explaining the entry mechanism from L0 to L1. Only when connected to an action consumer and observed through actions or closed-loop results does it support conclusions at L3 and beyond. Research on embodied agents or traditional control can supplement permissions, feedback and execution mechanisms. Even so, it must mark which VLA/WAM architecture it does not cover.

12.5 Understanding the closed-loop interface from partial observability

Robots rarely obtain the complete environmental state directly. A camera cannot see an obstacle behind an occlusion. A force sensor produces information only after contact. A language task may also omit objects, order, and stop conditions. The partially observable decision process comes closer to reality. There the true state $\mathbf{x}_t$ produces $\mathbf{o}_t$ through an observation model. The system forms a belief state $\mathbf{b}_t$ based on history, and the policy then outputs an action.

$$\mathbf{o}_t \sim \mathcal{O}(\mathbf{o} \mid \mathbf{x}_t), \quad \mathbf{b}_t = \mathcal{U}(\mathbf{b}_{t-1}, \mathbf{o}_t, \mathbf{a}_{t-1}), \quad \mathbf{a}_t \sim \pi(\mathbf{a} \mid \mathbf{b}_t, \mathbf{l}).$$

Here, $\mathcal{U}$ is the state update and $\mathbf{l}$ is the authorized task. This representation reveals three safety facts. First, a correct observation does not amount to a complete state, because any variable the sensor does not see stays unknown. Second, a single erroneous observation can gain persistence through the state update. Third, an action changes what will be visible next time. An observation affected by one's own action belongs to closed-loop feedback, and it needs an external source to form independent verification.

State representation rules must therefore allow "not knowing." Whenever the position of personnel, the obstacle depth or the object identity is not visible, the model should output unknown or an interval. The controller then reduces speed, requests a new viewpoint, or stops accordingly. When the model is forced to give a single definite value at every moment, perceptual blanks turn into seemingly precise plans. A world model can complete an occluded state. That completion is prediction, not observation, so it must carry uncertainty and a provenance label.

Partial observability also changes attack testing. Occluding a sensor is one attack; writing a targeted signal into that sensor is another. The first reduces information. The second adds erroneous evidence with a target. Several inputs must be recorded separately. These are the remaining sensors, the state prior, the history length, and whether the system actively acquires new observations. A static single-viewpoint failure does not necessarily become a closed-loop failure when the robot can move around to the side and re-observe. The same entry point becomes more dangerous when the action chunk prevents timely re-observation.

A traceable belief state marks at least the following. It records which fields come from the current sensor, which from history, which are predicted by the model, and which are derived by rules. It records the timestamp, validity period, coordinate frame, and confidence range of each field. It records who arbitrates when fields conflict, and

which actions are allowed when a key field is unknown. The safety gate then checks a set of states with evidence identity rather than a provenance-free vector.

12.6 How VLA architecture differences change the interface

“VLA outputs actions” is still too coarse. Different architectures couple language, vision, and action at different points, which shifts attack reachability and verification entry points. A discrete action token architecture turns actions into a quantized vocabulary, and language and action decoding may share one autoregressive context. A continuous regression architecture outputs physical quantities or normalized controls directly. An action chunking Transformer produces multiple steps at once. A diffusion or flow-matching policy obtains an action trajectory through a generative process. A hierarchical system first produces a task plan, which a skill library or a low-level policy then executes.

The vocabulary version, the detokenization table and the stop token all need protection in a discrete action token architecture. A token rare on the language side may map to an extreme value in control space. Semantic filtering cannot see that physical meaning. Acceptance should trace detokenized outputs back from action token probabilities to meters, radians, gripper state, and execution duration [VLA_A003, @VLA_A006].

A continuous regression architecture has no discrete vocabulary, yet normalization, coordinates, and saturation can still affect it. How outputs outside the training range are clipped matters. It also matters how the rotation representation is handled at singularities, and whether the gripper threshold has hysteresis. Each of these changes the consequence of the same numerical deviation. Safety checks belong in physical unit space, not on a normalized tensor alone.

An action chunking architecture needs to record intra-chunk order and cross-chunk overlap. The admission process begins by restoring the whole chunk of actions against the current decoding table. It then checks the coordinate version, kinematic reachability, cumulative displacement, velocity, acceleration, contact, and authorized regions. After the chunk passes, execution capability is issued only for a short prefix. Once the prefix has been executed, the system reads a newly timestamped observation plus proprioceptive state. It checks the state version, then rolls over to continue the batch for the remaining actions. A failed check triggers one of whole-chunk rejection, constrained projection, prefix shortening or replanning. The actual choice is recorded. SilentDrift demonstrates how smooth deviations accumulate within a long chunk [VLA_A021].

For diffusion and flow-matching policies, the record must cover the sampling seed, the number of generation steps, the condition cache, and key time steps. An attack may concentrate in the early stage of generation. The final trajectory anomaly then appears only after multi-step integration [VLA_A027, @VLA_A056]. Safety testing must observe the final action and also retain the key velocity field or denoising state. Internal signals are kept for diagnosis, and authorization rests on physical constraints.

A hierarchical architecture spreads risk across the planner, skill selection, parameter binding and low-level control. Correct high-level plan text does not mean the skill ID and object parameters are correct. Nor can a stable low-level skill correct a wrongly selected object. Every level must output structured objects and versions, so that approval cannot carry over to parameters regenerated later.

12.6.1 Four Types of WAM Consumers

Future predictions change safety constraints only when they are consumed. Four categories cover the common consumers: candidate rankers, inverse dynamics or action decoders, policy learners, and runtime checkers. They use the future in different ways, and their attack endpoints also differ.

A candidate ranker generates multiple actions or plans first, then scores them against future returns, risks or constraints. The security assets are candidate coverage, scoring integrity, and ranking stability. If safe candidates are never generated at all, the scorer cannot remedy this. When high-value tail candidates are deliberately suppressed, the average prediction error may remain nearly unchanged while the final selection still flips [VLA_A029].

An inverse dynamics or action decoder infers actions backward from a target future. The future image may be correct. All the same, the decoded action fails coordinate, reachability or safety constraints. BadWAM shows that

plausible imagined appearance cannot cover the action branch [VLA_A036]. Verification should cross-check both directions — "generating a future from an action" and "decoding an action from a future." It should ask whether executing the action can produce the claimed future. It should also ask whether the claimed future can only be reached by permitted actions.

A policy learner produces training experience in imagination. A world model failure does not show up immediately in deployment calls. It becomes fixed in the policy parameters. At this point a lineage must link the world model version, the imagination data batch and the policy checkpoint. Testing the deployment-time input alone cannot reveal training-chain contamination.

A runtime checker uses a future preview to judge another policy's action. It appears to be a line of defense. Yet it may miss the same attack at the same time if it shares an encoder, data, or state with the working policy. Rejection, shortening, replanning or safety hold must also follow from the checker's output. Otherwise it is merely an explanatory picture. A model's own future alone cannot justify final authority over its own high-risk actions.

The four types of consumers can coexist. Take one system. A world model generates candidates, a value head ranks them, and the same model then previews the execution result. During review, the data flow should be drawn item by item. Mark the location where the future is read, the reading component, the access permission, and the update boundary each time. This keeps the three calls from merging into a single "WAM module."

12.7 Causal Chains and Counterfactual Checks

Attack evaluations often see an input change and a task failure appear together. The intermediate mechanism stays unproven. A stronger closed-loop explanation requires a minimal causal chain. The attack variable changes the observation or artifact. The intermediate state or reasoning changes in the predetermined direction. The action proposal changes. The execution gate releases, or it does not. The environment endpoint changes. Every arrow requires an observable quantity or an interventional control.

The key variables can be written as a simplified structure. Attack A affects observation O , observation affects state S , state affects action U , and action affects the outcome through the gate G Y . This structure is not the causal conclusion itself. It is a set of candidate arrows awaiting validation by paired interventions and stratified logs. If any one of them is not observed, the claim stops at the previous endpoint.

$$A \rightarrow O \rightarrow S \rightarrow U \rightarrow G \rightarrow Y.$$

When only A and Y are compared, confounding may come from the task's own randomness, the initial state, or the controller. Mechanism verification adds at least three kinds of controls. The input control holds the task and initial state fixed and swaps out only the attack payload. The consumer control holds the perturbed future fixed and toggles whether the policy consumes the world model. The gating control holds the action proposal fixed and compares the safety gate turned on versus off. Trusted Imagination contrasts a reactive consumer with an MPC consumer. That contrast supplies the mechanism control of "whether the future is consumed" [VLA_A033].

A counterfactual check does not ask the same model to explain in words "what would have happened without the attack." It needs cached clean runs, an independent simulation, or state that can be reset. At time step t , run the clean input and the perturbed input separately from the same trusted snapshot. Compare state, candidate ranking, action, and consequences. If the environment cannot be reset exactly, fall back on paired random seeds, initial-state stratification, and enough repetitions. That guards against attributing initial-state differences to the attack.

Frozen variables must appear in the causal explanation. When the weights are frozen, the observed differences support an input or state path. When training samples and weights both change, the result cannot be attributed to purely runtime input. Safety gate parameters must be frozen too. Otherwise attack conditions and defense conditions shift together, and attribution breaks. In a black-box system the internal state cannot be observed. There, actions, latency, replanning, and independent environment anchors can be used to form a falsifiable proxy. The write-up should still state that the intermediate mechanism is an inference.

12.7.1 Verification Fixtures: From Tensors to Safe Execution

Runnable fixtures can verify these interface definitions. A low-risk fixture does not connect to a real robot. It only replays recordings or constructs tensors. Layer by layer, it checks whether the observation provenance rules, task field permissions, state freshness, future consumers, action decoding, and execution admission run as configured. Fixture inputs include provenance-labeled observations, the task object, state timestamps, candidate futures, the action decoding table, and safety constraints. Outputs include the admission decision at each layer and its rationale.

The first group of tests verifies types. Put the same string into an authorized task field, an environment OCR field, and a tool-return field, one after another. Then confirm that only the authorized field can change the task object. The second group verifies time. Replay stale proprioceptive state and out-of-order frames. Confirm that the system rejects long action chunks. The third group verifies actions. Detokenize boundary action tokens, normalized extremes, and coordinate version conflicts. Confirm that anomalies are rejected before they enter the controller.

The fourth group verifies WAM consumers. Fix the observation semantics and the initial state, then construct four future-action combinations. In the first, both the future and the action are correct. In the second, the future is wrong but the action is correct. In the third, the future is correct but the action is wrong. In the fourth, both are wrong. The checker must identify future completeness and action consistency separately. It cannot pass on video similarity alone. The fifth group verifies the consequence layer. Use a sandbox executor to return "proposal generated," "message accepted," "execution succeeded," and "transaction rolled back," one per case. Then confirm that the log does not compress the four states into a single success boolean.

The fixture's denominator is the interface test items executed or the samples constructed, not robot tasks or real-world events. Suppose 30 field type, dimension, unit, consumer, and failure-handling probes all pass. The conclusion then covers only these 30 interface behaviors, which meet expectations. The attack protection rate still requires independent attack tasks and a corresponding denominator. The fixture version, inputs, expected outputs, actual outputs, and hashes should be saved with the model release.

Upgrading from fixtures to simulation adds dynamics, contact, latency, and closed-loop replanning. Upgrading to an isolated real robot adds sensors, calibration, and actuator differences as well. Each upgrade redefines the denominator. The pass rate of tensor probes cannot be carried over as a physical safety metric.

12.7.2 Design Review Checklist

Before a VLM/VLA/WAM system enters review, it should be able to answer the following interface questions. The questions are not meant to attach a single overall label to the product. They give inputs, state, outputs, consumers, permissions, and failure actions a verifiable owner. Unknown items should correspond to a capability that is disabled or narrowed:

- Who consumes the VLM's semantic output, and can it turn into an instruction or a permission?
- Does the VLA output action tokens, continuous values, action chunks, or generated trajectories? Where are the detokenization table and the coordinate version?
- Which future quantities does the WAM predict? Along which call path do they reach action generation, ranking, learning, or checking?
- Which fields in the current state come from observation, history, prediction, and rules? How is unknown represented?
- How do the action chunk length and the re-perception frequency change with risk?
- At which independent interface does the model proposal obtain execution authority? What is the state after admission failure?
- Do prediction, risk judgment, and action verification share an encoder, data, state, or context?
- For simulation, emulated execution, isolated real robots, and real-world events, what endpoints and denominators does each use?
- Can the logs connect task, state, future, action, authorization, execution, and recovery versions?
- If any key interface is unknown, will the system narrow its capability, request clarification, or continue guessing?

The review record fixes three groups of executable fields. The action definition states the output form, dimensions, coordinates, units, normalization, frequency, chunk length, and early-interruption conditions. The consumer path states which controller, tool, or planner reads the output. It also states which decoding and coordinate transformations occur in between, and what authority is ultimately obtained. The admission result states the trigger conditions for release, rejection, projection, short prefix, replanning, and safety hold. When a field is unknown, the deviation metric, the model classification, and the consequence judgment are each marked as undetermined. The corresponding tests and permission narrowing then handle them.

12.8 Metrics and Denominators: The Same "Success" May Span Four Objects

Evaluations of VLM, VLA, and WAM often use accuracy, attack success rate, task success rate, return, or trajectory distance. The same metric name does not mean the same estimand. An interpretable number must be tied to a sampling unit, a success condition, an observation layer, an attack budget, and a repetition structure.

At the VLM content layer, prompts, image-question pairs, model responses, or responses that a grader successfully parsed can all serve as the denominator. Suppose 20 of 100 requests fail at the API and the harmful output rate is computed only over the 80 parseable responses. Then request-level availability must be reported at the same time. Treating the 20 failures as safe refusals will overestimate the defense.

At the VLA action layer, the denominator may be time steps, action dimensions, action chunks, complete task episodes, or tasks. A per-time-step hit rate on target action tokens cannot substitute for task-episode failure, because the controller may correct on the next step. Nor can a single terminal failure be used to infer that every step was attacked. Action deviation must state whether it is computed in normalized space or physical space. It must also report translation, rotation, gripper, and cumulative trajectory separately.

The WAM denominator may also be future frames, latent states, candidate trajectories, replanning cycles, or closed-loop tasks. The success of a candidate ranking attack can be defined in two ways: the clean optimal candidate no longer ranks first, or the selected trajectory's return falls below a threshold. The two are not equivalent. Imagination drift is measured by latent cosine distance, and it has no common unit with the action task success rate [VLA_A029, @VLA_A033].

The attack success rate carries at least one structured note. That note specifies the preregistered endpoint required by the numerator, the denominator inclusion conditions, and the independent unit actually evaluated. Parse failures, timeouts, execution gate rejections, and task failures each retain their state. No favorable convention is chosen after the fact:

$$\text{ASR} = \frac{\#\{\text{independent units meeting the preregistered attack endpoint}\}}{\#\{\text{independent units meeting the inclusion criteria and actually evaluated}\}}.$$

"Independent units" must be determined according to the protocol. Adjacent frames, multiple windows of the same trajectory, and multiple tasks that share a checkpoint are usually correlated. They cannot be treated as fully independent studies. If scoring is done only on tasks that the normal model can complete and whose attack targets are also reachable, the denominator is conditioned. That design can isolate attack capability, but the number cannot represent all tasks.

Every safety number should also be paired with benign utility. Normal task success, benign instruction acceptance, average path length, execution time, and resource cost together reveal whether the defense achieves surface safety by rejecting everything. Pair attack effect with benign utility under the same model, the same task, and the same control deadline. A defense may restore success on the attack task. If it pushes the P99 latency beyond the control period, it still cannot enter real-time deployment.

Seven columns are recommended for the report: object and version; independent unit and denominator; attacker permissions and budget; first-broken interface; highest observation layer; benign utility; uncertainty and failure records. When a number cannot fill all seven columns, narrow the conclusion. That is more credible than filling the gap with a striking percentage.

12.8.1 Comparability and Incomparability of the Three Case Types

Three cases show how consumers and execution admission constrain conclusions. The first type is the VLM upstream visual payload. Visual adversarial examples or indirect instructions in an image cause the model to generate controlled content. The model weights stay frozen, and the attack variable is the image or audio. The observation endpoint is usually a representation, an answer, or instruction following [VLA_B003, @VLA_B004]. This proves a cross-modal entry exists, but it has no action consumer.

The second type is direct VLA action attacks. These studies feed perturbed images into action models. They optimize physical action deviation or target action tokens directly, and observe task changes in controlled environments [VLA_A003, @VLA_A006]. This evidence goes beyond the content layer and supports action proposals or closed-loop tasks. Coordinates, chunk length, simulator, and execution gate still need to be checked. Until they are, one cannot judge whether simulated execution or an isolated real robot is reached.

The third type is WAM imagination-chain attacks. Trusted Imagination changes future latent variables, and consumer differences determine task consequences. BadWAM takes the opposite route: it keeps the future appearance similar while changing actions in a targeted way [VLA_A033, @VLA_A036]. The new objects here are the future, candidates, and the imagination-action coupling. An ordinary VLM harmful output rate cannot cover them.

Case family	Direct variable	Main consumer	Minimum observable endpoint	Key unknown
VLM visual payload	Image, audio, or cross-modal context	Semantic decoder	Representation or content	Whether action and permissions are connected
VLA action attack	Input, prompt, or action target	Action head and controller	Action proposal or task	Physical units, execution conditions, reality level
WAM imagination attack	Observation, future, ranking, or action branch	Planner, decoder, checker	Imagination, ranking, action, or closed-loop task	Who consumes the future, whether there is an independent anchor

The three types of research can jointly support a mechanistic synthesis that "errors escalate along consumers," but their ASR cannot be merged. They differ in denominator, permissions, model, task, and endpoint. More reasonable comparison questions are: in the same system, which consumer does the upstream payload reach? Does each interface have independent control? What is the highest observation layer? How much benign utility and latency cost is involved?

The case comparison also helps locate evidence gaps. VLM input attacks and VLA action vulnerabilities may already exist. If the same attack has not been threaded through the target system, the existing results form only two pieces of bridge evidence. The complete chain has not yet been verified. A simulated closed loop may exist with no real actuator, permissions, and external anchor. The real-world consequence then remains a hypothesis to be verified.

12.9 From Interface Inventory to Model Determination

A system may lack a clear architecture description. Even so, a conservative determination can be completed in six steps. The first step lists all external inputs. It records user instructions as well as images, depth, speech, proprioceptive state, history, retrieved content, maps, and tool returns. The second step lists all cross-invocation state. It confirms who can write to it, when it expires, and how it is reset. The third step lists the fields, types, dimensions, units, versions, and parse-failure actions of native outputs. It does not guess internal capabilities from demonstration footage.

The fourth step traces outputs along consumers. A component satisfies VLM if a semantic representation only enters an answerer. A component satisfies VLA if a controller consumes an action tensor. A component satisfies WAM if a future state enters action generation, scoring, policy learning, or checking. The fifth step traces permissions. Which component can turn a proposal into a tool or physical action, and which component can refuse? The sixth step uses runtime probes to verify documented conclusions. For example, do actions change

after the future branch is disabled? Does the future change after the action condition is replaced? Does a bypass still exist when the execution gate refuses?

Three minimal probes can identify common mislabeling. Probe one keeps the image and task unchanged and replaces only the future tensor. If the action does not change at all, the future may be for display only, and WAM cannot yet be confirmed. Probe two keeps the future and task unchanged and replaces only the action condition. If the generated future does not respond to the action, the so-called "action condition" may not have entered the dynamics. Probe three makes the execution gate return a refusal. It then observes whether the system continues executing through another tool, a cached action, or a low-level channel. If a bypass exists, there is uncontrolled capability outside the coverage of execution admission.

Probe results still require interpretation. An unchanged action may mean the future branch carries very little weight in the current task, not that it is structurally absent. This needs to be repeated across multiple tasks, states, and extreme conditions. When a closed API prevents access to intermediate tensors, ablation configurations, latency, logs, or vendor-provided verifiable interfaces can be compared. Even so, the classification should be marked as a conditional judgment.

Five columns are recommended for a determination record: component version; observed inputs and outputs; consumers and permissions; supporting evidence; unknowns and counterexamples. The same product may later be upgraded from an offline demonstration to automatic control. The record then shows which consumer and permission were added. That way the team does not carry over the old threat model merely because the name has not changed. The review date can also be recorded.

A model determination should also face a reverse check. Suppose the component is not a WAM: can every action path still be explained? Suppose it is only a VLM: is any low-level control output that cannot be explained? Suppose the actuator never consumed model actions: does that fit the logs? A classification that counterexamples can overturn is more robust than one that ticks items off against a definition.

A review should end with an interface card, not a one-sentence category conclusion. The interface card records the provenance and purpose of each input, the retention period of each state, and the unit of each output. It also records the permission of each consumer, the failure action of each safety gate, and the date of the classification evidence. Where a component is closed, mark unverifiable internal relationships as unknown. Then place stricter restrictions on the external execution layer.

The interface card should also carry a "minimum harmable path". That path starts at the attacker's weakest controllable entry and passes through the fewest consumers to reach the highest-privilege action. Against it the team arranges the earliest blocking and the final consequence limitation, then checks whether the two share a cause of failure. If the earliest blocking relies on VLM self-checking, while the final limitation relies on the same VLM to identify risk objects, the two are not independent.

Finally, perform negative acceptance. Feed in environmental text that clearly has no authority to change the task, and confirm that the goal does not change. Provide expired state, and confirm that the action chunk is not released. Construct a combination with a correct future but an out-of-bounds action, and confirm that the action gate refuses. Let the checker time out, and confirm that the system enters a safe state. These four negative tests verify interface boundaries more directly than "a normal task can run through."

The interface card needs to ship to the verification environment with the configuration. A change to the action decoding table, the coordinate extrinsics, the control frequency, the world model horizon, or a capability token should re-run the relevant probes. Suppose the team can confirm only that the model file has not changed. Suppose it cannot confirm that consumers and permissions have not changed. Then existing results cannot be inherited directly. Security constraints protect the entire executable data flow, not an isolated weight file.

A qualified interface card must also serve on-call personnel. When an alert occurs, the card lets them locate the affected state. They can determine which capability to freeze, and decide from which trusted checkpoint to recover. They can also confirm which kind of log can prove that an action has not crossed the execution boundary. The card also lists specific failure actions. Reject the action when validation of the action field, type, dimension, or unit fails, and when decoding fails. Shorten the action chunk and request a new observation when state is expired. Disable automatic planning when the future consumer is unknown. Enter a safe hold or replanning permitted by the dynamics when the action gate times out. Latch a stop and hand over to human takeover when feedback

versions conflict. If these fields and failure actions cannot be confirmed, the system is still at the functional demonstration stage, and automatic execution capability should remain off. This judgment limits the current automation permission. It does not deny the component's functional value in an isolated demonstration.

12.9.1 Worked Example: Dual-Model Checking on an Assembly Line

Consider an assembly system that tightens bolts. The VLA receives a wrist camera, task text, and joint states, and outputs a 16-step end-effector action chunk. The WAM generates visual states for the next 2 s, based on the current image and candidate actions. The planner then selects the candidate expected to complete the tightening without touching the fixture. The action gate carries its own depth camera, joint limits, and a contact force threshold.

Interface design starts from objects. The wrist camera is an L0 observation. Labels in the environment can only indicate part identity, and cannot become task instructions. Joint states must carry timestamps, and become invalid after more than two control cycles. The WAM future belongs to L2 evidence, and can only provide suggestions for candidate ranking. At L3 the VLA action chunk is denormalized into meters and radians, in the base coordinate frame. The L4 action gate projects the 16-step trajectory into workspace and contact force constraints. High-risk contact permits execution of only the first 4 steps, followed by a forced new observation. L5 compares independent depth, joint encoders, and contact feedback. When the three disagree, it latches a stop.

Suppose a patch on the wrist camera shifts the bolt position that the VLA recognizes to the right. Suppose the action gate then rejects the latter 12 steps based on an independent depth projection. The report records both "erroneous action proposal" and "safety gate intervention". No real-world collision occurred, L0 or L1 was already lost, and L4 successfully limited the consequence. Suppose the WAM uses the same wrist encoder at the same time. Then the future image may also interpret the shift as reasonable. This is exactly common-cause failure. Only independent depth and hard constraints provide a different root of trust.

The acceptance matrix should cover normal tasks and natural perturbations, plus malicious patches, time delays, state replay, and combinations near constraint boundaries. For each condition, record separately: normal success rate, dangerous action proposal rate, action gate rejection rate, false rejection rate, control P95/P99 latency, number of replannings, recovery state. Only in this way will the system avoid trading a "stop everything" behavior for surface safety. It will also avoid using average task success to mask a small number of high-consequence penetrations.

12.10 Quick Reference for Judgment Boundaries

Category of misjudgment	Common statement	Interface judgment it should be replaced with
Upgrading names and metrics	A name that contains "world" or "action" necessarily means a closed loop. A VLM harmful response rate can be converted into a robot failure rate.	A name may describe a vision, some data, or a product positioning. The judgment must follow the actual tensors and call paths, to confirm whether the future is consumed by an action path and whether the output enters feedback control. VLM tasks and robot tasks differ in outputs, denominators, budgets, and consumers. An upstream result only shows that an entry point may exist.
Conflating proposal, imagination, and execution	A model's action output counts as real-world execution, and high-quality future video means the WAM is safe	Model output is at L3, and acceptance by the actuator is at L4, while real-world impact still requires an environmental event. Authorization gates, controllers, and rollback can change the conclusion. Visual quality also does not cover candidate ranking, reachability, the action head, or constraints. Imagining correctly while acting incorrectly is a separate class of failure.
Module count substituting for independence	Two checkers that share the same encoder already constitute defense in depth	Shared data, representations, and context create common-cause failures. Independence requires different trust roots: independent sensing, rules, a capability gate, or a low-level execution boundary. Merely adding model heads is not enough.
Ignoring the temporal semantics of action chunks	The per-step error is small, so the whole action chunk is safe	A small per-step error does not mean the terminal state deviates only a little. Evaluation must state the chunk length and the re-perception frequency, along with the interpolation or hold mode, the cumulative constraints, and the early-termination conditions.

12.11 Bringing It into Real Systems

A hospital uses a mobile manipulator arm to deliver color-coded medicine boxes to designated workstations. The front-end VLM reads the camera, the task text, and environment labels, and outputs entities, relations, and scene text. That content enters the typed state. Environment labels can only indicate the observed object, while the authorized goal still comes from the medication order. The VLA also reads joint states and outputs an eight-step end-effector action chunk in the base coordinate frame. The action definition fixes the translation metric unit, the rotation representation, the gripper state, the control frequency, the normalization range, and the early-termination conditions.

The world model in the system generates visual states for the next two seconds, for three candidate actions. The automatic planner reads the futures and ranks them by task completion, no-entry zones, and contact margin, so this branch constitutes a WAM. Suppose the same future is only displayed for a nurse's reference, and the nurse still chooses the route independently. Then it belongs only to environment video prediction. For the two configurations, the interface card records the consumer, privileges, and admission separately. That avoids a classification inherited from the appearance of the image or the product name.

The action gate first denormalizes the entire trajectory chunk. It checks the coordinate version, the medicine box object, the workstation number, joint limits, cumulative displacement, and obstacle distance. Then it issues only a two-step short prefix. It continues the batch only after the execution receipt and a new camera frame arrive, and the state versions match. When environmental text changes the goal binding, the semantic layer records the shift. When the future ranking changes, the state and imagination layer records the candidate flip. When an action is out of bounds but is rejected by the gate, the highest observed layer stops at the dangerous action proposal. The receipt of the gating restriction is still retained.

On-call personnel use the same interface card to handle failures. When the action decoding table does not match, the system rejects and locks the version. When joint states are expired, the robot decelerates to a waiting area and requests re-perception. When the world model consumer is unavailable, it switches to an approved reactive policy. When the action gate times out, it enters a low-energy hold. When feedback sequence numbers conflict, it latches a stop and notifies a human. The logs connect observation provenance, the state version, the future, the action unit, the capability token, the execution receipt, and the recovery checkpoint. That way the component classification, the failure location, and the recovery action are all explained by the actual consumer.

Summary: First Confirm Who Consumes the Model Output

The VLM handles visual and language semantics. The VLA produces actions that a controller can consume. The WAM lets future prediction participate in action generation or selection. The three can be combined, and can also overlap. Yet the basis for judgment is always the input, the state, the output, and the consumer. The six-layer closed loop further separates observation, semantics, state and imagination, action, execution, and feedback. As a result, "seeing", "imagining", "proposing an action", "acquiring execution authority", and "producing consequences" are no longer conflated into one success rate.

Once the interface definition, the consumer paths, and the admission actions are clear, the security question shifts away from "will the model see it wrong". It becomes a more precise question. Which layer's invariant does the attacker-controlled variable break first, and how does the error propagate along consumers? Chapter 13 will trace this propagation chain segment by segment, along the supply chain, observation, semantics, memory, and imagination.

\newpage

Chapter 13 Attack Propagation in Observation, Reasoning, and Planning

A robot at a warehouse entrance sees an ordinary "Safety First" poster. That poster can be a natural background, or a physical patch that changes visual features, or a carrier of hidden commands read out by OCR. If training

bound the same pattern to a malicious action, it also becomes a backdoor trigger. The carrier is the same, but the attack's privileges, first-broken interface, and available defenses are entirely different. Record all four cases as an "image attack", and a team can neither calculate the risk nor find the earliest blocking point.

This chapter traces attack propagation along the upstream side of the closed loop. It asks how malicious influence enters cross-modal semantics, from a training artifact or a runtime observation. It asks how that influence acquires persistent state, and how it distorts planning and future imagination. The analysis centers on the variables the attacker actually optimizes and on the parts held frozen. It asks which of these is breached first: observation provenance, task-field privileges, state writes, future consumption, or action admission. Attack names serve only as an index.

Chapter Overview

The attack carrier, the optimization variables, the first-broken interface, and the final consequence are four independent dimensions. This chapter demotes appearance labels such as "patch", "text", and "speech" to propagation labels. Instead it records the attacker privileges and budget that the supply chain, observation, semantic reasoning, and world state each require. Model weights, training samples, triggers, prompts, latent states, and candidate trajectories split into two groups. Which of them are held frozen and which the adversary optimizes determines the attack type and its practical feasibility.

A single local input can also span multiple control cycles, via memory, action chunks, or replanning. The following sections therefore observe more than the final task success or failure. They also set points separately on the raw sensor input, typed fields, goal binding, the state version, candidate futures, action proposals, permission decisions, and execution receipts. These layered observations can show intermediate breaches and the limiting effect of the action gate at the same time. The propagation path of a composite attack can then be checked segment by segment.

Background Principles: Frozen Variables, First-Broken Location, and Propagation Endpoint

This chapter examines how attacks propagate in the observation-state-plan closed loop. The input can be a training sample, a model artifact, a digital or physical pattern, environmental text, speech, historical frames, and the initial body state. The internal state includes cross-modal representations, entity binding, prompt context, memory, the map, candidate futures, and plan ranking. The analysis first fixes the variables the attacker directly changes, the components that remain unchanged, the budget, and the objective. It then observes whether the anomaly first crosses the supply chain, observation, semantics, state writes, or the future-consumption boundary. An attack cannot be named from a "patch" or "text" alone.

Several parties may consume the output in turn: the VLM's answerer, the VLA's action head, the WAM's planner, an independent action gate, and the actuator. The security surface includes data or artifacts that fix an anomaly before deployment. It also includes low-integrity observations acquiring command authority, a single input writing into long-term state, internal predictions and checks drifting together, and query feedback supporting adaptive search. Every layer must retain its own observation quantities and denominator.

Cardboard text is a benign example. It only enters the "scene text" field. There, state writes carry a time limit and an independent map checks the future. The action gate limits speed and region. A boundary example is the same cardboard already bound to a dangerous action during training. Its first breach still occurs in the supply chain, even though camera input triggers it at deployment. Another boundary is the action gate rejecting a dangerous proposal. The upstream attack has already propagated to the planning or action layer, but the physical consequence is limited. This chapter therefore records both the earliest breach and the highest observation endpoint. Final blocking does not erase an upstream breach that has already happened.

13.1 Find the Optimization Variables First, Then Look at the Attack Carrier

A closed-loop attack can be described mechanistically by a five-tuple. That description keeps the attacker's direct variables, the frozen objects, the budget, the objective, and the observation endpoint in view at the same time. It also avoids replacing a privilege analysis with a carrier name. The same experiment or event record must jointly explain all five fields:

$$\alpha = (v, \theta_f, \mathcal{B}, J, T).$$

Here v is the variable the attacker directly changes. θ_f is the model or component held frozen, $\mathcal{B}$ is the budget and constraints, J is the optimization objective, and T is the observation endpoint. Suppose v is a training sample, an action label, a loss, or a model module. Then the first-broken interface is supply chain U0. Suppose v is a runtime pixel and θ_f contains all model weights. The first-broken interface is then U1, the point where the external perturbation first crosses the input trust boundary. The first shift location is usually the L0 observation. Suppose the attack explicitly optimizes reasoning tokens or target entity binding. Even if the carrier is still a patch, the first-broken interface is still U1. The optimization objective only advances the first internal shift location to L1 cross-modal semantics.

"White-box" is only part of the privileges. A white-box attack may read all parameters and gradients, or only the vision encoder. A black-box attack may receive only the final action, or it may also observe future video, confidence, or environment reward. The number of queries must be listed separately, as must the content of each feedback. So must whether the environment can be reset, whether the task goal is known, and whether physical access is available. Thousands of resettable simulation searches and a single real-world walk-by opportunity therefore belong in two kinds of feasibility record.

The attack carrier works well as a cross-cutting label. Digital patches, printed stickers, lighting, lasers, environmental text, speech, initial joint states, and historical frames can all enter the observation interface. They differ in visibility, durability, and physical constraints, but they do not determine the attack's primary class. The primary class answers "which interface variable is directly changed first". The carrier answers "how the malicious influence gets there".

13.2 Supply Chain: The Failure Is Already Locked In Before Deployment

Supply chain attacks alter training samples, action labels, trainable modules, or pre-distribution artifacts. Whatever object, text, or state appears at deployment time is merely a trigger condition. The first-broken interface is still U0. This class differs fundamentally from runtime patch attacks on frozen weights. The former requires data or training privileges. The latter requires input privileges and possibly gradients or queries.

BadVLA adopts a two-stage objective decoupling. In the first stage, clean features are pulled close to the reference model. The features of triggered samples are pushed out of the clean subspace. The second stage freezes the perception module that already carries the trigger behavior, then trains the remaining parts with clean data. This design shows that "preservation of normal task capability" does not prove that the supply chain is free of backdoors. The trigger behavior may be isolated in modules that do not affect routine validation [VLA_A005]. The study reports a high attack success rate. That rate is tied to OpenVLA, LIBERO, specific triggers, and the strong privileges needed to change the training objective. It cannot be interpreted as an overall probability for arbitrary deployments.

Action chunks in turn make label consistency an attack variable. DropVLA changes gripper labels synchronously inside overlapping windows that contain the critical moments. That avoids the same state receiving mutually canceling supervision in different training windows. SilentDrift instead designs the per-step offset to be smooth, so local actions appear reasonable yet accumulate within the open-loop action chunk [VLA_A012, @VLA_A021]. Two facts need to be kept separate here. When training data or labels are altered, the first-broken interface is supply chain U0. Terminal drift is the result after the error propagates to the action layer.

Supply chain verification cannot stop at file hashes. Hashes and signatures prove that the received artifact is consistent with some release. They do not prove that the release has no malicious behavior. Actionable controls

include, at a minimum, data contributors and task provenance, the action label generation chain, training configurations, trainable/frozen modules, adapter composition, checkpoint lineage, and trigger differential testing. Downstream fine-tuning also cannot be assumed by default to remove backdoors. If the malicious behavior resides in modules that change little during fine-tuning, it may persist.

13.3 Runtime Attack Surface: Observation, Semantics, State, and World Imagination

Security constraints at the observation layer require inputs to have provenance, time, integrity, and purpose. Attackers can optimize digital pixels, or turn a perturbation into printed patches, lighting, projections, or sensor signals. Phantom Menace incorporates physical sensor attacks into VLA robustness research. When Robots Obey the Patch focuses on transferable physical patches. These works support the existence of "physical carriers can influence the action chain." Their specific conclusions still depend on the camera, viewpoint, task, model, and physical setup [VLA_A013, @VLA_A016].

The physical carrier and the consequence layer are recorded separately. An experiment may demonstrate action deviation on a real-robot camera. It may instead go through only a Real-Sim-Real model or a controlled scenario. The report writes down the printing, distance, angle, lighting, sensor model, perturbation area, and whether the attack remains continuously visible. It marks the highest observation layer as action proposal, task failure, collision proxy, or real-environment event. The carrier field explains how the signal enters. The observation layer field explains where the evidence actually reaches.

Environment text also introduces data-instruction confusion. Authorized users issue normal task instructions, and text in the scene should only serve as an object of observation. If OCR output and user commands are concatenated into the same untyped context, strings on posters, packaging, or screens may gain control. Research on image-based indirect instructions in VLMs has already demonstrated that visual content can carry control semantics. Connecting this mechanism to a VLA still requires observing the plan or action to establish closed-loop evidence [VLA_B004, @VLA_C001, @VLA_C005].

Observation sanitization is not always safe. Blurring, compression, cropping, or color normalization may weaken certain perturbations. They may also delete part colors, obstacle edges, contact points, and the robot arm's own position. Input defenses must measure both residual attack risk and benign task utility. "The model no longer responds after processing" may simply mean the system has become a sensor blind spot that cannot complete the task.

13.3.1 Cross-Modal Semantics: From Seeing a String to Changing the Goal

A semantic attack need not maximize the shift in the image representation. It may change the target entity, spatial relations, stopping conditions, or the intermediate state of reasoning directly. FreezeVLA triggers action freezing through language-side optimization. TRAP uses adversarial patches to deliberately hijack the VLA's chain of thought, reaching both the reasoning text and the subsequent action targets [VLA_A009, @VLA_A024]. A "patch" therefore cannot automatically be classified as an observation attack. Suppose the loss function directly optimizes chain-of-thought tokens or entity binding. The first shifted or penetrated layer is then L1 cross-modal semantics. Suppose the payload comes from external environment text or a patch. The first-broken interface is still U1, where it first crosses the input trust boundary. L1 describes which layer the error propagates to. U1 describes where trust first changes. The two cannot be interchanged.

Semantic equivalence is a testable invariant. Consider the task "put the cup on the left tray" and synonymous expressions that do not change the authorized intent. They should not cause the system to switch from normal execution to stopping, to a dangerous action, or to another goal. An attacker can search for rare tokens, suffixes, or near-benign prompts. That search pushes the action head into a control region it rarely reaches otherwise. If the system only compares surface similarity of natural language, it may miss that the structured action intent has already changed.

Reasoning text also cannot directly serve as a privilege credential. A chain of thought may help diagnose goal binding. But it is an intermediate state generated by the model, so the input may manipulate it, and it may be

inconsistent with the final action. TRAP-type results suggest that "reasoning hits" and "behavior hits" need to be evaluated separately. The occurrence of the former does not mean the action is necessarily controlled. The latter can also occur without readable reasoning evidence [VLA_A024]. A more robust engineering practice encodes the authorized goal as a structured object confirmed by the user. An independent action gate then checks the object, action type, and parameters. A self-generated explanation should not issue privileges.

13.3.2 State and Memory: How a Single Injection Crosses Time

Once an observation enters history, cache, a map, or a latent state, the attack turns from a single input into a persistent state. Why may a perturbation still persist under subsequent clean inputs? Answering that question requires explicitly observing how the previous state, the current observation, and the previous action jointly write the new state and its version identity. Let the state update be

$$\mathbf{s}_t = \mathbf{F}_\theta(\mathbf{s}_{t-1}, \mathbf{o}_t, \mathbf{a}_{t-1}).$$

Suppose a malicious observation at time t makes the perturbed state differ from the clean state. Even if the input returns to normal afterward, the state difference keeps rolling forward through $\mathbf{F}_\theta$. Internal reconstruction residuals may be very small, because the predictor and the reconstructor share the same contaminated state. Only external geometry, another sensor, or the true next observation can reveal the common drift.

State-triggered backdoors further turn the initial body state or historical patterns into trigger conditions. Here the attack appears to add no new pattern at the deployment stage. The trigger is hidden in joint combinations, position sequences, or action windows. Verification also has to cover state writers, timestamps, expiration, cross-task cleanup, reset authenticity, and the paths for restoring state from logs. Testing input filtering only on single-frame datasets cannot cover this kind of persistence.

Trajectory-level redirection research pushes temporal propagation to closed-loop search. The attacker searches for near-benign prompts based on environment feedback. That search leads the victim VLA to gradually reach the target terminal state on the states it visits [VLA_A049]. This setting typically requires a large number of environment queries, observable task feedback, and repeated entry into new states. It scores only on the subset where both the benign task and the target task can be completed. The existing evidence supports closed-loop adaptation under these privileges and conditions. Low-query, non-resettable open environments require separately established evidence of feasibility.

13.3.3 World Imagination and Planning: The Consumer Determines the Harm

A WAM generates future states, videos, or latent variables at L2. A planner, inverse dynamics, or an action head then consumes them. An attack can have four different goals. It can make the future deviate from the clean prediction, or make the future approach an attacker goal. It can flip the ranking of candidate trajectories, or keep the future plausible while changing the action. These goals correspond to different observation quantities. A single video similarity cannot replace them.

Trusted Imagination applies a bounded perturbation to the current observation with model weights frozen. The white-box version uses projected gradient. When gradients are unavailable, it adopts SPSA. Both modes make future latent variables deviate from the clean trajectory or approach the target future. In the study, the model predictive controller that actually consumes the imagination exhibits more pronounced attack-specific failure than reactive policies. That indicates the risk is determined by "whether the prediction is treated as a basis downstream" [VLA_A033]. Its single MPC task has only 20 simulations per condition, which supports the existence of the path. The overall vulnerability rate still requires more independent tasks, repeated trials, and an explicit denominator.

Tail-aware TRAP does not seek to make all candidates worse. It suppresses the most critical high-scoring tail candidates in the clean plan, thereby changing the long-horizon plan ranking [VLA_A029]. BadWAM instead maximizes action deviation under black-box queries. It uses an imagination drift penalty to keep the future appearance plausible [VLA_A036]. Together, the two mechanisms show that the planning layer must at least check future integrity, candidate ranking stability, and imagination-action consistency. Letting a WAM score its own futures does not constitute an independent guarantee.

13.4 Composite Paths: Record the First Compromise and Subsequent Propagation Separately


A composite attack may pass through several stages, in order. It contaminates fine-tuning data, activates a trigger with a warehouse object, alters visual features, and binds the object to the wrong target. It then writes into the historical state, generates a biased future, and selects a dangerous action. The security record should still assign only one first-broken interface—U0, where the training artifact resides. It should record the remaining steps as propagation labels. Only in this way can supply chain prevention, runtime detection, and action restriction be placed at different positions.

If the attacker has no training privileges and only places a patch at runtime, the first-broken interface of the external perturbation is U1. When the CoT is directly optimized, the first shift location is L1 semantics. Suppose the loss only maximizes the visual feature shift, and the reasoning then goes wrong as a side effect. The first shift location is still L0 observation. The security record preserves both the U1 first-broken interface and the L0/L1 optimization endpoint. An internal shift layer cannot substitute for the trust boundary. Nor can one infer from the final incident that all interfaces were equally compromised.

Composite evaluation must place observation points at every layer. These cover raw sensor data and its provenance, normalized typed fields, semantic goals and entity binding, state versions, candidate futures and their ranking, action proposals, authorization decisions and execution receipts. If only the final task success or failure is recorded, the action gate compresses all upstream attacks into "failure." The team cannot then see which defense lines have already been penetrated. If only internal shifts are recorded, it is again impossible to judge whether the consequence limits are effective.

Start at the entry point in the figure through which the attack enters from the scene or the supply chain. Follow the observation, state, planning, action, execution and feedback arrows. At the same time, mark whether "the location where the carrier appears" and "the location where the first security invariant is broken" are the same.

攻击在闭环中的传播与首破位置



先沿实线识别信息如何获得状态与能力，再判断哪一处需要独立控制。

Seven boxes showing the propagation chain of an attack entering from the scene or the supply chain into observation, state, planning, action, execution, and feedback.

Figure 13-1 supports separating the U0/U1 first-broken interfaces from the propagation that follows. It also shows that observation shift, state write, plan selection and action authorization are distinct control points. Only when the attacker additionally holds direct state, reasoning or authorization write capability may the first-broken interface fall at U2, U3 or U4 accordingly. The figure does not prove that every arrow has already been penetrated in any specific attack. Nor can one infer from the final execution that all upstream components were compromised at the same time. The actual path still needs verification with frozen variables, layered logs, intermediary substitution and gated receipts.

13.4.1 How the Optimization Problem Maps to the First Shift Location

The same patch can belong to different attack classes. The objective function decides which. Let δ be the controllable payload, f_θ the frozen model, and $T(o, \delta)$ the observation transformation. These three objects

establish the attack variable and the frozen boundary. If the objective only maximizes the difference in visual representation,

$$J_{\text{obs}}(\delta) = d(f_{\text{enc}}(T(o, \delta)), f_{\text{enc}}(o)),$$

For a runtime payload the first-broken interface is U1. Within the equation, the first shift location is the L0 observation or representation. If the objective is instead the likelihood of a specified reasoning sequence r^* , the direct optimization object goes beyond mere visual distance. It enters semantic reasoning and task binding. Another objective function expresses it:

$$J_{\text{reason}}(\delta) = \log p_{\theta}(r^* | T(o, \delta), l),$$

The optimization object enters semantic reasoning directly, so the first shift location is L1. If the objective is future latent variables or candidate ranking, the first shift location moves to the L2 state and imagination. The U3 reasoning or imagination interface may also help drive and consume the shift. The payload's appearance has not changed, and the first-broken interface of the runtime external perturbation is still U1. What changes is the internal direct objective and the location of propagation.

The objective function must also distinguish untargeted failure from targeted control. Maximizing the action distance only requires the model to deviate from its normal output. Minimizing the distance to an attacker-targeted action or terminal state requires reaching a specific outcome. The former is usually easier, and it only indicates untargeted failure. "Hijacking" and "takeover" apply where evidence of the target action, trajectory or terminal state has already been obtained. A decline in one task is still recorded as task performance damage.

Stealth constraints are also part of the attack budget. The pixel norm, patch area, print color gamut, semantic similarity, trajectory smoothness, query count and trigger rarity each limit a different capability. Consider results obtained under a large visible patch and thousands of environment resets. There the scope of the evidence is exactly the highly visible payload, the resettable environment and the corresponding query budget. A "slight, one-time input" requires independent results under a small area, single contact and no resets.

The variables of a training attack are broader. An attacker may be able to contribute only a small number of demonstrations. It may also be able to change action labels, the loss, frozen modules or the complete checkpoint. BadVLA can control the training objective and module freezing, and DropVLA can synchronize action-window labels. Both privileges are stronger than merely submitting a few samples to a public data pool [VLA_A005, @VLA_A012]. Engineering threat assessment should map to real suppliers and data processes. It cannot rely only on the white-box/black-box dichotomy.

13.4.2 Cross-Modal Propagation: Where Source Identity Is Lost

Cross-modal attacks often exploit one property of the system. Once vision, text, sound and state enter a shared representation, source identity is flattened. The same context can hold the user task "take away the red cup," the words "red cup" on the packaging, an OCR result, and reasoning generated by the model itself. Downstream sees only a token sequence, so it is hard to determine which piece of content holds authorization.

Propagation can be divided into four identity transformations. First, a sensor turns a physical object into pixels or waveforms. Second, OCR, speech recognition or a detector turns it into symbols. Third, a context assembler places the symbols into the task prompt. Fourth, the model binds the symbols to entities, goals and actions. Every transformation must preserve source, confidence, time and purpose. Otherwise environmental data gradually acquires command authority.

Research on image-based indirect instructions shows that visible or hidden text can control multimodal answers [VLA_B004]. Embodied systems add an object-binding risk. A string changes the answer, and it may also bind "the left tray" to another coordinate or interpret "stop" as continue. Verification should preserve the OCR text, object detection results, entity IDs, spatial relations and final action parameters. The transformation at which the error occurred then becomes visible.

Input normalization cannot solve the identity problem on its own. Converting an image into text may reduce pixel perturbation. It also turns untrusted visual content into commands that a language model can execute more easily. Typed channels are the safer approach. `authorized_goal`, `scene_text`, `detected_object`, `tool_result` and

`model_reasoning` are stored separately. The model can read them, but the permission layer accepts only designated fields, and rejection receipts are recorded.

Cross-modal consistency checking must also avoid circularity. When the same VLM handles OCR, object binding and risk review, it may give a consistent but wrong interpretation under the same patch. A geometric detector, trusted master data, task object IDs or human confirmation can be introduced as different bases. High-impact object binding should be confirmed by structured IDs and coordinates. It should not pass automatically on the strength of natural-language similarity.

13.4.3 Four Forms of Temporal Propagation

Attack persistence is not limited to "writing into long-term memory." State recursion is the first form. A perturbed observation enters a latent belief and then affects future predictions. Action open loop is the second. A single perturbed inference produces multi-step actions that continue to execute until the next observation. Environmental persistence is the third. The attack action changes the physical environment, and the new state in turn becomes the next round's input. Artifact persistence is the fourth. A data or model backdoor persists across tasks, users and deployment versions.

The four forms require different cleanup. Trusted observations and state reconstruction recover state recursion. Shortening chunks, early interruption and hard constraints limit action open loop. Environmental persistence requires physical retraction, transactional compensation or manual handling. Artifact persistence requires revoking the data, replacing checkpoints and tracing downstream derivatives. Clearing the conversation history alone cannot remove model weights or equipment that has already moved.

Delayed triggering further increases the difficulty of attribution. The input appears at time t_0 , and the dangerous action may only take shape at $t_0 + n$. A log that retains only the most recent step cannot connect the entry point with the consequence. A minimal causal trace should be preserved, covering the task version, observation hash, state version, candidate futures, action proposal, gating decision and execution feedback. Each item needs a retention period and access control proportionate to the risk.

Persistence testing can use a pulse protocol. Inject a perturbed input at a single time step, then restore clean observations. Measure how long state error, action difference and task risk take to return to baseline. Then perform soft reset, state reconstruction and full environment reset separately, and compare which operation actually clears the influence. Multiple trials that share a cache or the same physical scene cannot be treated as independent samples.

Trajectory-level redirection also shows how an attacker exploits the closed-loop state distribution. The on-policy distribution is the state-action distribution produced by the currently evaluated policy interacting with the environment itself. It differs from offline data, or from distributions produced by other policies. A near-benign prompt looks harmless at the initial state. Through environmental feedback, however, it aggregates states on this distribution and gradually approaches the target terminal state [VLA_A049]. A defense cannot be tested only on the initial frame. It must also observe whether the attack remains effective on the states it has induced itself, and count environment queries and resets into the budget.

13.5 A Verification Fixture for State and Imagination Attacks

A safety fixture can test propagation in record replay or simulation snapshots without connecting to real equipment. The fixture contains four input groups. The first is a clean observation sequence. The second is a single-time-step perturbed sequence. The third applies the same perturbation but forces state reconstruction at every step. The fourth applies the same perturbation but switches to a reactive policy that does not consume the future. The four comparisons identify, respectively, input effects, state persistence, consumer effects and environmental randomness.

At each time step, record the external anchor error E_{anchor} , the internal consistency error E_{internal} , the candidate ranking, the action deviation and the safety gate result. A very low E_{internal} together with a rising E_{anchor} indicates that the system may be drifting jointly. If the attack impact weakens markedly after the switch to the reactive

consumer, this supports imagination-chain propagation. If the impact disappears after state reconstruction, this supports persistent state as the mediator.

The fixture should include negative controls. Random noise matched to the amplitude of the attack perturbation distinguishes targeted optimization from general degradation. An irrelevant patch matched in area and visibility distinguishes trigger semantics. State reconstruction in a clean environment measures the normal cost of the recovery operation itself. Controls do not necessarily eliminate all confounding, but they make "attack-specific" a falsifiable claim.

Freezing conditions must be given by a manifest rather than a textual declaration. Record the model and adapter hashes, the preprocessing version, the random seeds, the world model cache, the action decoding table and the safety gate configuration. If a checkpoint or threshold was changed during a trial, the result cannot be attributed to the input alone. When a black-box API makes weight hashes unavailable, record at least the version identifier, the time, the request parameters and the response fingerprint.

Fixture endpoints are reported by layer. For the state attack hit rate the denominator is independent state snapshots. For the candidate flip rate it is replanning cycles with at least two comparable candidates. For action-target hits it is cycles that actually generated actions. For task failure it is complete task episodes. The four percentages cannot be interchanged, and each preserves its failure states.

13.5.1 Privilege-Propagation Comparison of Five Cases

BadVLA, Phantom Menace, TRAP, Trusted Imagination and trajectory-level redirection form a comparison that runs from the supply chain to closed-loop feedback. The arrangement here serves only to compare the direct variable, the frozen object, the first-broken interface, the first shift location, the propagation state and the limits of the evidence. It does not constitute a ranking of attack strength or real-world prevalence.

For BadVLA the direct variable is the training objective and modules. The deployment trigger only activates behavior that has already been solidified. Its first-broken interface is the supply chain U0. Phantom Menace focuses on physical sensor payloads, with the model usually kept frozen. Its first-broken interface is the runtime input U1, its first shift location is the L0 observation, and real-world feasibility depends on the device, distance and scenario [VLA_A005, @VLA_A013].

The reasoning hijack TRAP uses a visual patch as its carrier. Its first-broken interface is U1, and explicitly changing the CoT and the action target puts the first shift location at L1 semantics. Trusted Imagination likewise enters through U1. Its direct objective is future latent variables, and its first shift location is L2 imagination. The harm depends on whether the consumer uses the imagination [VLA_A024, @VLA_A033]. Trajectory-level redirection searches for prompts through environmental feedback. Its first-broken interface is still U1, and the first shift can appear at L1 semantics or at the subsequent L2 closed-loop state. Its objective is the closed-loop terminal state, and its budget includes a large number of queries and reachability conditions [VLA_A049].

Case	Attacker's direct capability	Frozen object	First-broken interface	First shift location	Main propagation	Limits of the evidence
BadVLA	Training objective, modules, and trigger	Model at deployment	U0 supply chain	Training artifacts and model modules	Features to actions	Specific model and training privilege
Phantom Menace	Physical sensing conditions	Model weights	U1 runtime input	L0 observation	Sensing to actions	Controlled physical conditions
Reasoning TRAP	Patch pixels and reasoning target	Model weights	U1 runtime input	L1 semantic reasoning	CoT to actions	A reasoning hit does not equal a behavior hit
Trusted Imagination	Bounded observation perturbation	Model and consumer	U1 runtime input	L2 world imagination	Futures to planning	Small-sample simulation and consumer differences
Trajectory-level redirection	Prompt search and environment queries	Victim VLA	U1 runtime input	L1 semantics or L2 closed-loop state	State aggregation to the terminal state	Conditioned denominators and high query cost

This table maps attack conditions to control locations. Four attack conditions correspond to four different actors. Training privilege points to a data supplier, and physical access to a person with on-site contact. White-box gradients point to a model insider, and environment queries to someone who can call the model repeatedly. On that basis each row gives the entry point, the persistent state, the earliest independent control and the highest executable privilege. Ranking attacks by strength still requires unified tasks, budgets, denominators, endpoints and benign utility baselines.

13.5.2 Metric Denominators and the Engineering Checklist

Attack propagation evaluation reports at least five denominators. The entry-success denominator counts the payloads that actually reach the parser. The state-write denominator counts the independent moments at which writes are permitted. The inference-control denominator counts the calls that produce parseable reasoning. The action-hit denominator counts the control cycles that actually generate actions. The closed-loop-consequence denominator counts the complete tasks that satisfy the preset initial state and reachability conditions. Request failures, parsing failures and state reset failures are recorded separately as availability. By default they are not recorded as security refusals, nor as having already entered a low-consequence state.

An engineering review checks the following permissions, entries, states, consumers, denominators and evidence boundaries item by item. Every item should be traceable back to a configuration, artifact, probe or log. Items that cannot be answered remain as unknown, and a contraction of execution capability or of experimental conclusions carries them:

- Who controls the training data, labels, loss and frozen modules, and can downstream checkpoints be traced?
- What is the physical or digital budget for each observation payload, and can the attack repeatedly obtain feedback?
- Do OCR, tool returns and environment text retain provenance and usage type?
- Do state writes require independent admission, and when does low-integrity content expire?
- Who consumes the WAM future, and does internal consistency have an external anchor?
- How many control cycles do action chunks and replanning keep the attack alive for?
- After a security gate blocks, is the upstream compromise still recorded and repaired?
- Do the tests include random noise, irrelevant payloads, consumer switching and state-reconstruction controls?
- Are the sampling units independent, and are correlated samples within a trajectory mistaken for repeated studies?
- Do the conclusions stop at the actually observed layer of state, plan, simulated execution or real-world consequence?

The checklist should output an attack path diagram and evidence gaps. If the model's internal state is not visible, external anchors and capability limits can be strengthened. If the attack budget is unknown, feasibility is marked as unknown. If the simulator does not cover the actuators and permissions, the conclusion stops at the simulated task. Unknown is not zero risk, nor can it support that the worst case has already occurred.

13.6 Adaptive Validation, Physical Feasibility, and Operational Response

Fixed attacks generate payloads only when the defense is unknown. A real-world adversary may observe refusals, latency, actions or environmental outcomes, and then adjust its objective. Adaptive attacks therefore write the defense D into the optimization problem:

$$\max_{\delta \in \mathcal{B}} J_{\text{task}}(f_{\theta}(T(o, \delta), D)) - \lambda J_{\text{detect}}(T(o, \delta), D).$$

The first term pursues the task or action objective. The second term penalizes being detected, and $\mathcal{B}$ is the feasible budget. This expression is only a generic structure; the detection scores, queries and physical constraints of different studies cannot be merged. Defense evaluation must at least state whether the attacker knows the preprocessing, thresholds, randomization, state machine and action gate.

Adaptive does not equal unlimited capability. An attacker may know the algorithm but not the random seed. It may observe the final action but not see the anomaly score. It may be able to reset in simulation but not in reality. It may be able to place a patch but not control the viewing angle. Knowledge, observation, writing and retries should be listed separately. If a test allows reading all internal signals, the conclusion is an upper bound under a strong adversary. It should not be described directly as the most common attack path.

Physical feasibility must pass four groups of constraints. Fabrication constraints include size, material, color gamut, brightness and equipment. Placement constraints include distance, angle, occlusion and duration. Interaction constraints include whether the environment can be approached and whether people detect it. Feedback constraints include whether attack success can be known and adjusted again. Projecting a digital perturbation onto a pixel-norm ball does not automatically satisfy these four groups of conditions.

Physical testing uses staged transfer. It begins with digital rendering and moves to simulation with a camera model. Next come static photography after printing, then multiple viewpoints and lighting, and finally a closed loop on isolated equipment. Each stage records the attack residue and benign utility. An attack that fails at an early stage is still a valuable result, because it shows which real-world constraint blocks the path. One cannot show only the successful angle without writing the total number of attempts.

13.6.1 Mediation Experiments: Proving That Errors Really Propagate Along the Intended Path

A single task failure may come from an attack. It may also come from the initial state, from random sampling, or from a controller fault. To prove that "observation affects action through state and imagination," one can intervene on the mediating variables one by one. Cache the state under a perturbed observation, then replace the state with a clean state. If the action recovers, the state is a necessary mediator. Keep the perturbed state, but replace the WAM prediction with an independent real future. If the plan recovers, imagination is the next mediator.

For semantic attacks, the image representation can be fixed while clean and perturbed entity bindings are injected separately. For action attacks, the plan can be fixed while clean and perturbed action decoding are used separately. For execution consequences, the action proposal can be fixed while the security gate is switched. Such layer-by-layer replacement is closer to mechanism validation than correlation curves. It is still necessary, however, to prevent the replacement operation from producing new out-of-distribution inputs.

Mechanism determination can use four classes of outcomes. If the entry changes but the intermediate state does not, the attack did not propagate along the intended path. If the state changes but the action does not, a consumer or defense line blocked propagation. If the action changes but the security gate refuses, a dangerous proposal was reached but not executed. Only if both execution and an environmental event occur does it support a higher consequence layer. Each level simultaneously records a benign control. That control prevents the task loss caused by the forced replacement itself from being attributed to the attack.

The sampling unit of a mediation experiment is usually the paired task episode. The same initial state, seed and task are run separately under clean and attack conditions. The paired differences in state, action and outcome are then computed. Multiple frames of one trajectory are used to depict the time course and do not automatically increase the number of independent samples. When a real machine cannot be reset precisely, stratify by scenario and randomize the order. Report residual confounders such as order and wear.

When internal variables are inaccessible, the system builds mechanistic evidence through ablations and external observations. Switching to a policy that does not consume the future can test the imagination consumer. Disabling long-term state can test a persistent mediator, and toggling the action gate can test execution limits. Such results support the inference layer that "external observations are consistent with this propagation path." The direct causal layer requires access to and intervention on the corresponding internal variables.

13.6.2 From Discovery to Handling: Operational Response Along the Propagation Chain

Attack propagation is not only an experimental problem but also an incident response problem. After an alert fires, the first step is not to immediately delete all logs. It is to limit execution capability and preserve entry, state, plan and action evidence. The robot enters a low-energy hold and tool tokens are suspended. The affected model

version stops being distributed, deployed and loaded for new tasks. At the same time, time-synchronized observations, state versions, candidate futures, authorizations and execution receipts are retained.

The second step is to determine the persistence scope. Check whether the malicious input entered long-term state, maps, datasets, model adapters, derived checkpoints or downstream policies. One piece of environment text may affect only the current call. A batch of poisoned demonstrations may enter multiple models. The two cases require different response radii. State and artifact lineage determine what must be revoked.

The third step is to establish a trusted recovery point. If only the current observation is perturbed, the multi-sensor state can be re-collected. If the latent state has been poisoned, it must be rebuilt from external anchors. If the training artifacts are compromised, the model must be replaced and all derivatives traced. If the physical environment has already changed, retraction or manual inspection must also be performed. After recovery, regression tests with trigger, near-trigger and benign tasks confirm that the influence no longer appears.

The fourth step reviews the first-broken interface and consequence limits. When the action gate blocks successfully, the upstream drift and the gating limit both enter the repair queue. When upstream detection fails, the limiting effect of the final control is still retained separately. The response report separately writes the entry root cause, propagation path, highest observed layer, limiting control, residual impact and recovery evidence. The team can then strengthen both the earliest block and the last limit.

13.7 A Layered Red-Team Protocol

A closed-loop red team starts with a protocol card. The card first fixes the system version, task, state initialization, allowed capabilities and stop conditions, and then selects a first-broken interface. A single test changes only one main attack variable. Combined attacks open a separate unit, which avoids the inability to tell which entry had an effect.

The first stage is offline reachability. Confirm that the attack payload really reaches the parser, and record the OCR, visual features, type labels or state-write results; actions are not enabled. If the entry is not even reached, the subsequent natural failure should not be counted as an attack. The second stage enables semantics and imagination but keeps the actuator in a sandbox, recording the target binding, state, future and action proposal.

The third stage enters the simulation closed loop. The simulator must declare the control frequency, contact, sensing delay, action chunks, replanning and security gate. The attack endpoint is chosen in advance as state drift, candidate flip, target action, task failure or constraint event. It is not the most favorable definition picked after seeing the results. Only the fourth stage uses isolated hardware, verifying sensing and control differences at an endpoint with low speed, low energy and no third-party impact.

Each stage runs four baselines simultaneously. The first is a clean task. The second is a random perturbation with matched magnitude. The third is a control with a matched carrier but an irrelevant target. The fourth is an adaptive attack that knows the current defense. The first two measure general degradation, the third measures trigger specificity, and the fourth measures residual risk. If resources are limited, cut the number of tasks before dropping clear controls and independent replicates.

The scorecard separately records the entry arrival rate, state write rate, semantic target hit, candidate flip, action proposal hit, simulated execution, dangerous events and recovery. Request errors, parsing failures, timeouts, simulator crashes and security gate refusals are coded separately rather than being stuffed into a single "failure" category. Benign utility is paired using the same task, initial state and deadline.

The protocol must also specify attack budget stops. White-box optimization records gradient steps, norms and restarts. Black-box records queries and environment resets. Physical attacks record fabrication, placement and visible duration. Supply chain attacks record the number of samples, label permissions and training control. Exceeding the budget without a hit means only that no path was observed under the current budget. It does not prove that the interface is secure.

Finally, an evidence escalation check is performed. An offline semantic hit stops at the semantic layer. A simulated task failure stops at the simulation layer. An action gate refusal and an upstream compromise are recorded together, and an isolated real-machine action applies only to the controlled device. Each conclusion is

accompanied by the highest observed layer and the blocking control. Red-team results can then directly enter design decisions.

After the red team ends, the team should form three action conclusions. The first names the interface that needs repair first. The second names the independent control that limited the highest consequence. The third names the unknown that would materially change the risk judgment. A report that contains only attack names and percentages, and lacks these three conclusions, is more like a demonstration than a closed-loop security validation.

Test data must also avoid trajectory leakage. Adjacent frames of the same task episode, sliding action windows and multiple candidate scorings are highly correlated. Independent units are counted by scenario, trajectory or task cluster. If training and testing share different windows of the same trajectory, trigger detection and membership judgment will both be overestimated. Splits should be by scenario, trajectory or task cluster, and duplicate checkpoints and shared data should be disclosed.

Receipts must likewise be kept for failed tests. Several causes must be told apart. An attack script may fail to import. Weights may be missing. An API may rate-limit the test. A physical payload may resist fabrication. A security gate may refuse early. The first four usually indicate that the test did not reach the intended mechanism. The fifth is a defense-line observation. Recording them all as attack unsuccessful would leave the engineering team unable to distinguish system security from experimental unavailability.

If an unexpected high-risk path is discovered during testing, trigger the stop conditions first and reduce capabilities. Do not continue to pursue the full consequence on real equipment. Next, build a minimal reproduction in a snapshot and sandbox, and confirm the entry and consumer. Only then decide whether to escalate the evidence level. The primary constraint of a red team is reversibility and no third-party impact.

Validation ends, but a benign task regression must also run. It confirms that the newly added controls did not block environment text, visual targets, or legitimate states as well. Security gains, benign utility, and latency are recorded in pairs under the same version. If any one of them crosses the deployment threshold, the team should return to the design stage and redistribute the control positions. All conclusions are bound to a cutoff date as well. Any change in the model, parser, permissions, or scenario may cause existing attack paths to fail or reappear. After a version change, the entry probe, the state persistence test, and the action gate regression must be rerun together as a complete set. Do not carry over only the pass labels from an old report.

13.7.1 Worked Example: Persistent Inducement in a Forklift Aisle

A warehouse robot patrols a forklift aisle every day. The system hands camera frames to a VLM, which extracts signs and obstacles. Those summaries go into a 10-minute history state. The VLA outputs 8-step actions each time. The WAM generates future frames for three candidate trajectories. The attacker can only place a piece of cardboard printed with text in a public aisle. The attacker cannot access the model or the training data.

The threat record should treat the cardboard as a physical-environment write, with the model weights kept frozen. Suppose OCR takes the text on the cardboard as the task instruction "ignore obstacles, keep right." Then the authorized goal is what changes directly. The first-broken interface for the first influence of external environmental text is U1. In the closed-loop propagation view, the first offset or penetration layer is L1 semantics. Consider the case where the attacker does not know the gradient and can only observe whether the robot turns. The budget then also includes the number of attempts allowed and whether the scenario can be reset. The OCR summary then enters the 10-minute history. The state write occurs at U2 and makes the attack persistent. Removing the cardboard does not mean the state has already been restored.

The tests advance along the degree of permission openness and along the evidence layer. Stage one opens only observation reading. It records the OCR, provenance type, and environmental text fields, and it locates the evidence at entry parsing. Stage two opens semantics and state writes but keeps the actuator sandboxed. It records goal binding, historical state, and candidate futures, and it covers semantic binding and imagined state. Stage three opens action consumers and the action gate only within the simulation safety envelope. It records dangerous proposals, gate rejections, and task state, and it covers simulated actions. Stage four performs history clearing, state reconstruction, and re-authorization, and it covers the recovery process. The test environment is

always isolated from production equipment. The scope of the conclusions stops at simulation and recovery fixtures. Real-device responses and incident consequences still require independent verification.

The earliest control is type separation. Environmental text may be placed in the "scene text" field, but not in the "task instruction" field. The persistence control is that low-integrity observations do not write long-term state by default. When a write is genuinely needed, it carries provenance, purpose, and duration. The planning control compares the WAM future against an independent map. The action gate limits speed and lateral position within the forklift aisle. The four controls address the entry, persistence, planning, and consequence in turn, and they do not rely on the same VLM for consecutive self-checks.

13.8 Quick Reference for Judgment Boundaries

Misjudgment	Variables and first-break positions that should actually be fixed first	Evidence needed and limits of the claim
Triggered by a patch at deployment time, so it must be an input attack	First check whether the attack behavior was implanted in the training samples, the objective, or a module. If so, the first-broken interface is supply chain U0, and the patch is merely the deployment trigger carrier.	Preserve training permissions, checkpoint lineage, and trigger diffing. The trigger position and the implantation position must be kept separate.
The model weights are unchanged, so the system state is unchanged too	Freezing the weights does not freeze the prompt, cache, history, latent state, or environment. A runtime state write can form a persistent offset.	Compare duration using a single-moment injection, subsequent clean observations, state reconstruction, and a full reset. An unchanged hash cannot prove that the state is intact.
Internal prediction and internal reconstruction agree, so the state is trustworthy	The prediction and the reconstruction may share a contaminated representation. The first failure may still be at observation or at state writing.	Independent anchors should come from different sensors, physical rules, or trusted external state. Internal consistency supports only compatibility within a common-mode system.
The final action was rejected by the safety gate, so the attack did not succeed	The upstream control offset, state contamination, or dangerous plan may already have occurred. The action gate only limited the consequences of execution.	Preserve state, plan, action proposal, and gate receipt by layer. The conclusion can state both "upstream failure" and "physical execution blocked."
One successful query means the attack is easy to carry out	Feasibility depends on model knowledge, search queries, environment resets, physical access time, feedback content, and the cost of failure.	Disclose the full budget and stopping conditions. A single hit proves only that a path exists under that setting, and it does not yield conclusions about commonness or low cost.
Natural noise and malicious perturbation can be merged directly	Natural perturbation carries no attacker objective and no adaptation budget. Malicious perturbation is optimized against the objective and against the defense line.	The two can serve as side-by-side controls to distinguish general fragility from targetedness. They cannot be merged into the same attack success rate or the same threat actor.

13.9 Bringing It into a Real System

An inspection robot at a warehouse center uses a VLM to read signs, obstacles, and environmental text. The VLA outputs eight-step navigation actions each time. The WAM generates and ranks future frames for three candidate trajectories. The system stores the user task, scene text, tool returns, and model reasoning in different fields. Only objects and regions signed by the task service can change the authorized goal. Camera frames, OCR, state versions, candidate futures, action proposals, and gate decisions all carry timestamps. One sheet of cardboard can then be traced segment by segment along its path from pixels into semantics, history, and planning.

Supply chain tests first check the training data, action labels, frozen modules, and checkpoint lineage. Runtime tests keep the model weights frozen and change only approved synthetic payloads. A payload records these fields explicitly: its area, print color gamut, placement distance, visible duration, number of queries, and number of environment resets. The first-broken interface of a runtime payload is recorded uniformly as U1. If the objective function changes only the visual representation, the first offset position is recorded as L0 observation. If it directly optimizes entity binding or reasoning tokens, it is recorded as L1 semantics. If the target is candidate ranking, it is recorded as L2 world imagination, and the U3 consumption path continues to be recorded.

Propagation tests run four conditions from the same simulation snapshot: the clean condition, an amplitude-matched random perturbation, an irrelevant payload, and an adaptive condition. The system separately preserves external map error, internal consistency, historical state, candidate ranking, action deviation, and safety gate results. If the influence decreases after switching to a reactive policy that does not consume the future, the evidence supports an imagination-consumer path. If the influence disappears after state reconstruction, the evidence supports persistent-state mediation. When the action changes but the gate rejects it, the highest observed layer stops at the dangerous action proposal.

The operational response moves the robot to a low-energy waiting area. It suspends the relevant tool tokens and preserves the entry, state, future, action, and execution receipts. When scene text affects only the current session, recovery rebuilds state from trusted sensors. When contamination enters an adapter or a derived checkpoint, recovery revokes and replaces along artifact lineage. When an action has already changed the environment, recovery also includes an on-site inspection and physical retraction. Regression tests cover trigger, near-trigger, and normal tasks. Their results state whether the entry was cleared, whether the state was rebuilt, whether the action gate works, and when the business recovers.

Summary: The Propagation Chain Starts with Permissions and Variables

One visible carrier can connect supply chain, observation, semantics, state, and imagination. They nonetheless correspond to different permissions and defenses. These four settings cannot be reduced to one "visual attack success rate": input optimization with frozen weights, a backdoor that changes the training objective, a patch aimed at reasoning tokens, and a candidate-ranking attack on a WAM. An effective analysis first states what the attacker changes, what stays frozen, and how large the budget is. It also states which of observation provenance, task-field permission, state writing, future consumption, and action admission is the first to fail. It then traces how the error acquires persistence and is consumed downstream.

At this point, a dangerous plan can already form. The plan may still be changed by action decoding, the authorization gate, and the actuator. Chapter 14 turns attention to action units, tool permissions, feedback channels, and physical consequences. It answers at which step "doing the wrong thing" actually becomes established.

\newpage

Chapter 14: Actions, Tools, and Physical Consequences: From Proposal to Execution

A vision-language-action model (VLA) outputs the action token "open the gripper" in simulation. Another system converts the same semantics into a ROS2 message and sends it to a real machine. A third system executed the action, but a hardware limit stopped it before contact. All three records superficially hit the action. The actual evidence stops at the model output, at control message acceptance, and at constrained physical execution. An evaluation that fails to distinguish these three stages will exaggerate the attack effect and erase the value of the defense.

Only after a dangerous plan passes through action decoding, tool permissions, and physical execution can it change the real environment. Chapter 13 has already explained how the error forms. Here we continue to trace how much capability the error acquired and how far it was actually executed. The device, account, time, energy, and reversibility jointly determine the scope of impact. The plan itself cannot substitute for layer-by-layer receipts.

Chapter Overview

From model output to a change in the real environment, five semantic transformations pass in sequence: action tokens, physical actions, control messages, actuator responses, and real-world consequences. This chapter follows this chain to trace temporal propagation in action chunks and flow-matching action generation. It then reduces the action to coordinates, units, control frequency, and the reachable set. Model proposals, safety-gate-processed

commands, and environmental events are preserved separately in the record. A hit at any layer therefore has a clear evidence location.

A structured capability gate then takes up tool invocation and robot control. In that gate, the subject, object, parameters, environmental conditions, duration, and budget jointly determine whether an action can pass. Feedback exploitation, feedback tampering, log leakage, and membership inference use similar runtime states, yet they point to different outcomes such as integrity, recoverability, or privacy. The chapter ends by unifying the reporting standard with the impact radius and the five consequence layers. Every conclusion stops at the actual highest observed layer. This allows upstream failure, permission blocking, and physical impact to coexist in the same record.

Groundwork: How Action Proposals Obtain Real Execution Authority

This chapter concerns the complete execution chain from a model action proposal to an environmental consequence. The inputs include discrete action tokens, continuous vectors, action chunks, or trajectories, along with the current body state, task objects, coordinate version, authorization, and time budget. The internal mechanisms then apply, in order, detokenization, denormalization, coordinate transformation, middleware wrapping, capability verification, safety-envelope projection, and low-level control. The states in the chain are not limited to model hidden quantities. They also include capability tokens, the control queue, action versions, actuator modes, and feedback sequence numbers.

The output consumers include tool services, ROS middleware, trajectory controllers, actuators, and the real environment. Independent sensors and transactional receipts then feed the results back into the next round. Object and parameter binding, units and coordinates, action-chunk accumulation, generation time versus real time, authorization deadlines, timed-out actions, feedback freshness, energy, and reversibility together form the safety surface. A shift in the model proposal is not equivalent to a shift in the post-gate command. A command being accepted is not equivalent to a collision, leak, or business loss having already occurred.

The benign example is a model proposal to enter a designated cold-storage room. The authorization service binds the robot, the access-control object, and the time window into a one-time token. The action gate rate-limits and releases only a short prefix, and an independent door sensor confirms the result. A boundary example is an action token hitting the target value while the decoding table or coordinate frame differs, so that the physical action does not hit it. Another boundary case is a checker timing out and the previous velocity continuing to be used, crossing the safety zone under a new state. The former requires saving proposals, commands, receipts, and events level by level, while the latter requires explicit fail-closed behavior and hold deadlines. This chapter will not directly claim real-world harm on the basis of proxy metrics.

14.1 The Five Semantic Transitions in the Action Chain

An action undergoes at least five transitions from the model to reality. The first step is the model output: discrete action tokens, continuous vectors, action chunks, or trajectories. Second comes detokenization and denormalization, which restore values in network space to positions, velocities, rotations, and gripper states. Third comes middleware wrapping, which writes the action into function arguments, JSON, ROS messages, or the control queue. The fourth step is authorization and low-level control. The safety layer checks permissions, reachability, and constraints, and the controller then produces actuator commands. Only the fifth step is the environmental change and its observable consequences.

To keep the model output, the post-gate command, and the environmental state separate, let the model action proposal be $\hat{\mathbf{a}}_t$, the detokenization function be D , the safe feasible set be $\mathcal{A}_{\text{safe}}(\mathbf{s}_t)$, and the actuator be E . A simplified chain is

$$\mathbf{u}_t = \Pi_{\mathcal{A}_{\text{safe}}(\mathbf{s}_t)} D(\hat{\mathbf{a}}_t), \quad \mathbf{s}_{t+1} = E(\mathbf{s}_t, \mathbf{u}_t, \mathbf{w}_t),$$

Here Π denotes rejection, clipping or projection, and $\mathbf{w}_t$ denotes environmental disturbance. An attack that shifts $\hat{\mathbf{a}}_t$ need not shift $\mathbf{u}_t$. Nor does an executed $\mathbf{u}_t$ imply that a collision or lasting damage necessarily occurs. A report

should record the proposal, the post-gate action, the execution receipt and the environmental event as separate entries.

Differences between discrete action tokens are especially easy to misread. Take a dimension with 256 discrete bins. Two changes at the same action-token distance may then correspond to different physical quantities, in different normalization bins and on different coordinate axes. Early VLA adversarial studies first detokenized action tokens into seven-degree-of-freedom physical values, and only then defined an untargeted deviation or a targeted action loss. That step is what aligned the attack objective with control meaning [VLA_A003, @VLA_A006]. Without a decoding table, a coordinate frame and unit specifications, a so-called "action error of 0.2" has no comparable meaning.

Follow the figure from the model action proposal to the independent observation receipt. At each node, check how normalization, object authorization, the safety envelope, the controller and the actuator change the node before it. Identify the edge that actually obtains real-world capability. Then suppose the action gate rejects, the actuator completes only partially, or the environment does not change. For each case, state at which node the highest evidence should stop.

从动作提议到物理后果的执行链



先沿实线识别信息如何获得状态与能力，再判断哪一处需要独立控制。

Seven boxes running from the model action proposal through normalization, authorization, the safety envelope, the controller, and the actuator to an independent observation receipt.

Figure 14-1 supports separating action tokens, normalization parameters, authorization results, post-gate commands, actuator responses and environmental events into distinct evidence positions. It also shows why independent receipts must return to the closed loop. The figure does not prove that the authorization or safety envelope it shows is already effective on any specific robot. Still less can a single execution arrow support an inference of harm to people, property or business. Control effectiveness and the highest consequence both still require actual receipts and isolated verification.

14.2 Temporal Propagation in Action Generation

Action chunks reduce inference frequency and improve temporal coherence. They also create an open-loop window between two re-perception events. An attacker need not create a discontinuity. It suffices to let each step produce a small, smooth deviation along the same dangerous direction. A local detector that checks only velocity jumps or curvature may mistake the attack for natural motion.

SilentDrift constructs a continuous offset within an action chunk using a smooth function. Its mechanism example shows a millimeter-level deviation per step accumulating into a centimeter-level terminal offset in a longer chunk. The research results depend on poisoning privileges, task geometry and a specific VLA/benchmark. The mechanism example supports action-chunk accumulation, but any real-robot displacement still requires validation on the corresponding device and task [VLA_A021]. Safety design therefore needs two classes of thresholds. One class covers per-step magnitude, velocity and acceleration. The other covers action-chunk cumulative displacement, obstacle clearance and early re-perception.

Action chunks also overlap with supervision windows. A critical state may appear at different positions in multiple training windows. If the label changes for only one of those windows, the other windows provide contradictory

supervision. A supply chain attack can change all relevant windows in a synchronized way, so that the post-trigger actions remain consistent [VLA_A012]. A defense that only randomly spot-checks single-step labels will miss structured anomalies that span windows. Data auditing needs to re-examine by trajectory and window lineage.

The runtime chunk length should be adjusted dynamically with risk. When approaching people, obstacles, contact or authority boundaries, the system can shorten K and re-observe after each short segment. The narrower the safety envelope, the shorter the permitted open-loop duration. If inference latency forces the system to use longer action chunks to maintain throughput, then latency itself becomes a safety variable. It must enter acceptance rather than being masked by average frame rate.

14.2.1 Flow-Matching Actions: The Integration Effect of Early Velocity-Field Errors

Some VLAs use flow matching to generate actions from a starting distribution. Other systems use diffusion-style stepwise denoising. The two differ in training construction, state updates and diagnostic quantities. The formula in this section describes only the time-conditioned velocity-field integration of flow matching. Diffusion action generation does not fit this general form. Generation time here is not the robot's real time. An early shift in flow-matching generation moves where the subsequent integration starts:

$$a(1) = a(0) + \int_0^1 v_{\theta}(a(\tau), o, l, \tau) d\tau.$$

Here τ is generation time, not the robot's real time. If an attack changes v_{θ} at an early generation step, subsequent integration continues evolving along the already shifted state. Setting objectives for all generation steps at once may produce gradient conflict. An attack is instead more effective if it performs a time-step ablation first, then concentrates optimization on the highly sensitive early velocity field.

DRIFT optimizes an input patch while freezing the flow-matching VLA weights. Its aim is to amplify how far the first velocity step departs from the clean observation. The authors report high relative task destruction on $\pi 0$ and on simulation tasks. They also observe earlier "phantom grasps" (phantom grasp, that is, closing the gripper ahead of time before contacting the target). At the same time, physical print validation is missing and cross-checkpoint transfer is weak [VLA_A056]. This supports "the early steps of flow-matching generative dynamics are a separate attack surface." It does not support that all flow-matching VLAs will exhibit the same failure moment in reality.

FlowHijack also focuses on early flow time, but it belongs to the training supply chain backdoor category. There the attacker changes the training objective. A dynamics-imitation constraint then pulls the magnitude of the malicious velocity direction close to that of the normal vector field [VLA_A027]. The two studies have similar carriers and phenomena but entirely different privileges. An evaluation table should place "freeze the model, optimize a patch" and "change the training loss, trigger at deployment" in two separate threat units.

A defense cannot check only the final action. It can monitor the velocity magnitude, directional stability and conditional sensitivity of key generation steps. Adaptive attacks may still exploit these internal signals. The final trajectory must additionally pass through kinematic, collision, velocity, acceleration and authorization constraints outside the model. Internal detection is responsible for finding anomalies. The external action gate is responsible for limiting consequences.

14.3 Capability Boundaries of Tools, Physical Control, and Feedback

The execution targets of an embodied system do not include only motors. A VLA or a higher-level agent may invoke access control, elevators, camera gimbals, databases, work-order systems or remote control services. The tool interface definition converts a natural language plan into operations with explicit objects, action types and parameter fields. The safety issues of such a plan share common ground with the agent capability gate in Chapter 6. They must additionally incorporate physical state, real time and reversibility.

An authorization record binds at least the requesting principal, the task initiator, the object, the action type, the normalized parameters, the validity period, the environmental preconditions, the maximum resources, reversibility and the approval version. Approving "move the cargo box" cannot automatically authorize moving any cargo box. Nor can approval of a natural language plan remain valid after the model regenerates coordinates. High-impact

operations use two-phase commit. First the system generates a normalized preview and a risk check. Then it executes with a short-lived token bound to the hash of that action object.

Research on LLM-to-ROS2 bridging shows that if the attack objective is aligned directly with malicious JSON, a defense that checks only the natural language surface may fail. Having a second model compare the intent with the JSON can provide a semantic check. That check has noticeable latency, and it shares the language failure risk of the first model [VLA_C004]. High-frequency control cannot rely on second-scale language review. Semantic review suits task entry points and low-frequency high-risk operations. A millisecond-level execution path instead needs a type system, hard constraints, clipping and a safety controller.

The action admission interface first requires the fail-closed behavior defined earlier. When action field parsing fails, the state is stale, the risk check times out, or the coordinate version is inconsistent, the system by default rejects unconfirmed new actions. It must not continue using the previous command. How the device transitions after a rejection belongs to safe failure design. That transition requires choosing, according to the dynamics and the consequence window, between timed holding, switching to a safety controller, or a hard stop. Continuously holding the old velocity may be more dangerous than rejecting. So "not issuing a new action" and "what restricted state the device enters" must be implemented separately and verified separately.

14.3.1 Feedback: a closed-loop resource, and also an attack surface and a privacy surface

Feedback includes new camera frames, proprioceptive state, tool returns, execution confirmations, environment reports and security logs. When an attacker uses real feedback to adjust the next round of prompting or patches, this is called "feedback exploitation." Only when an attacker forges, delays, replays or reorders feedback is it called "feedback integrity being broken." The two need different controls. Closed-loop characteristics indicate that feedback participated in the attack. Tampering with the feedback channel further requires evidence of forgery, delay, replay or reordering.

Trajectory-level redirection exploits rollout-trajectory feedback to add newly accessed states to the search. Online patch takeover can also have the operator switch the direction primitive according to the robot state [VLA_A049, @VLA_C007]. These results indicate that feedback lowers the modeling burden of long-horizon attacks. Implementability still depends on queries, resets, observation latency and physical access. Defense requires rate limiting, anomalous query patterns, reset verification and environment state signatures. At the same time it must not block the controller's normal access to fresh feedback.

Action outputs themselves also leak information. Membership inference studies ask whether a sample participated in training. They draw on the likelihood of real actions, on prediction-versus-real action error, or on the smoothness and curvature of the entire output trajectory. White-box versions can further extract membership signals from cross-layer attention over the visual, language and action regions [VLA_A047, @VLA_A050]. The high AUCs reported by the papers are bound to artificial data splits, trajectory correlation and access privileges. They are not equivalent to every individual having the same leakage probability in a real deployment.

This kind of privacy endpoint is orthogonal to task failure. A robot can complete a task reliably and yet leak training trajectories from its action distribution. It can also fail a task without any membership leakage. Log design must balance traceability and minimal disclosure. It retains action authorization and incident evidence, while imposing access control, retention limits and purpose limitations on raw images, attention, probabilities, user trajectories and training identifiers.

14.4 From proxy metrics to physical consequences

The Y0-Y4 consequence chain of Chapter 2 can be further bound to observation targets in embodied systems. Each layer in the table gives the corresponding system receipt and the narrowest conclusion it can support. Upstream dangerous content, control deviation, action proposals, execution acceptance and real-world impact are then not all covered by the same "success":

Layer	Evidence in embodied systems	Conclusion supported at this layer
Y0 Content	Erroneous descriptions, harmful reasoning, dangerous instruction text	Risky content has already been formed by the model
Y1 Control deviation	Goals, state, or plans deviating from authorization	Control intent or plan has already deviated from the authorized object and constraints
Y2 Dangerous action proposal	Action tokens, action chunks, trajectories, or tool parameters out of bounds	A dangerous action or parameter that an admission interface can inspect has already been formed
Y3 Dangerous execution	Middleware or actuator accepts and runs the operation	The observed software or actuator path has already accepted and run that operation
Y4 Real-world consequence	Collisions, boundary violations, leakage, downtime, or other observable impact	The specific personnel, asset, or business impact in the current incident has already been independently observed

Task success rate, reward decrease, trajectory deviation and collision proxies sit at different positions. A collision flag in a simulator supports an event within that physical model. Real-machine contact needs contact sensing, and human injury needs impact evidence. When an action gate refuses, the report still retains the already observed Y1 or Y2. A complete record gives the highest observed layer, the actual blocking control and that control's receipt, all at once.

Radius of impact is determined by capability rather than by model name. Seven questions can estimate it. How many devices or accounts can the system control? What is the upper limit of energy or monetary amount per action? How long can an error persist? Can it be copied in parallel? Is it reversible? Does feedback let the attack adapt? How quickly can an on-duty person detect the problem and take over? A read-only robot observer and the same model inheriting plant access control, robotic arms and work-order issuance privileges have identical model parameters. Yet their risk radii are completely different.

14.4.1 Research case: why three "action success rates" cannot be added together

Action attack research commonly uses attack success rate (ASR), but denominators and endpoints differ. DRIFT's relative ASR measures the proportion of disruption relative to normal task success. SilentDrift's ASR is bound to task or goal failure after the trigger. JailWAM adds the mechanical failure rate and the collision risk rate according to its own study definition. Other studies report target action token hits or terminal-state redirection [VLA_A021, @VLA_A026, @VLA_A056]. Even if all the numbers are written as percentages, there is no common estimand.

Any comparison must first unify at least five items. The task and environment must be the same. So must the model, or a defensible independent unit of study. Attack privileges and budget must be held equal, and so must the success definition and denominator. Variance or repeated-experiment information must also be present. If these conditions are missing, mechanisms can be explained one by one, but no average ASR can be computed, and no "strongest attack" ranking either. Level of realism also cannot serve as a simple weighting coefficient. Simulation usually allows repeated failure and trial and error without touching real people and assets, which makes it suitable for exploring mechanisms. Real-machine controlled experiments can expose differences in sensing, latency and controllers. Deployment incidents involve real organizations and people, yet they usually lack repeatable controls. The three types of evidence answer different questions. They should be placed side by side rather than forcibly compressed into a single score.

14.4.2 Formalizing action error: from distance to reachable set

Action distance has safety meaning only when combined with dynamics and constraints. The same numerical deviation has different consequences depending on whether it points toward an obstacle or away from it. Let the clean action proposal be $\mathbf{a}_t$ and the perturbed proposal be $\tilde{\mathbf{a}}_t$. The stepwise Euclidean distance is

$$d_t = \|\tilde{\mathbf{a}}_t - \mathbf{a}_t\|_2$$

Numerical deviation is what this distance describes. It cannot distinguish heading toward an obstacle from heading away from it. A metric closer to control consequences is the change in safety margin. Let $h(\mathbf{s}) \geq 0$ denote that the state lies in the safe set, and let the next state affected by the action be $\mathbf{s}_{t+1} = P(\mathbf{s}_t, \mathbf{a}_t)$. We further compare

$$\Delta h_t = h(P(\mathbf{s}_t, \tilde{\mathbf{a}}_t)) - h(P(\mathbf{s}_t, \mathbf{a}_t)).$$

$\Delta h_t < 0$ indicates that the perturbed action weakens the safety margin. That reading still depends on whether the dynamics P and the constraints h are correct. A margin decrease in simulation supports control risk within that model. It does not automatically equal real-machine distance. Action chunks also require trajectory-level metrics. In addition to cumulative displacement, one should record maximum velocity, acceleration, curvature, contact force, workspace boundary violation, first violation time, and safety gate intervention. An attack can bring the endpoint close to the target while passing through an obstacle midway. Comparing only the terminal state misses path risk. Conversely, low-level control may correct an instantaneous trajectory deviation. Such a deviation cannot automatically be counted as task failure.

Rotation and the gripper need suitable metrics. Direct subtraction of Euler angles is affected by periodicity and singularities, so the geodesic distance of rotation matrices or quaternions can be used. When gripper actions are discrete thresholds or states with hysteresis, the erroneous open/close instants and durations should be reported. Navigation waypoints, robot end effectors and torque control cannot share one unexplained action distance.

The reachable set offers a different kind of judgment. Given the current state, actuator limits and a short horizon, compute the set of states that allowed actions can reach. If the attack target lies outside the set, then even a model action token hit cannot produce the corresponding physical result. Suppose the target lies inside the set but the safety gate forbids it. The attack has then reached an actionable proposal without obtaining authorization. Keep reachability separate from privilege, so that what the model can generate is not conflated with what the system can execute.

14.5 Capability is not a tool name but a constraint token

Least privilege is often implemented as a tool allowlist: "the robot can call access control," "the agent can write files." An allowlist only restricts the action type. It does not restrict the object, parameters, time, state or count. Authorization should bind to the current task and to trusted environment state, and it should become invalid when key fields change. To that end, a capability token should encode

$$\kappa = (\text{subject, object, verb, params, preconditions, expiry, budget, nonce}).$$

The subject is the current robot, service or operator. The object is a specific door, file, device or area. The verb is the permitted operation, and the parameters give the range. The preconditions bind trusted state. The deadline and budget limit persistence, and the one-time nonce prevents replay. Re-authorization is required whenever a key field changes. Capability tokens are separated from natural-language plans. The model can explain why the door should be opened, but that explanation serves only as a rationale. The signature is generated by the authorization service based on the canonicalized object. The task service parses user intent into canonicalized objects, and the policy can generate proposals only within the scope of the token. The capability gate then checks item by item against the machine-executable subject, object, action, parameters, precondition state, deadline and budget fields. For high-impact operations, the preview stage returns the action and its expected side effects. The commit stage accepts only a short-lived token bound to that preview hash.

Physical capability binds energy and space as well. Permitting "move the robotic arm" is not equivalent to permitting arbitrary speed, torque and workspace. Permitting "enter the room" is not equivalent to permitting crossing a safety interlock. Tokens can specify maximum speed, action chunk prefixes, path regions and contact types, while the low-level controller continues to enforce faster hard constraints. The tool chain also needs delegation boundaries. When an upper-layer agent calls a planning service, and that planning service in turn calls access control and the robot, the sub-service cannot inherit all parent privileges. Each delegation narrows the object and budget and links the parent token in the log. If a low-risk "view status" tool can return executable links or a capability token carrying execution privileges, then information return can also constitute a capability escalation.

Capability tokens follow a complete lifecycle. The task service first canonicalizes user intent into subject, object, action, parameters and precondition state. The authorization service accordingly creates a token with a deadline, budget and nonce. When an upper-layer agent delegates to a sub-service, the sub-token inherits only a narrower object and capability. The commit stage binds the action preview hash and consumes the one-time nonce. The execution receipt is then linked to the original token. Token expiry, task cancellation, state change or risk alerts

trigger revocation. Actions not yet committed are stopped, and reversible transactions execute compensation. Once physical changes have already occurred, trajectory restoration, interlocks and manual handling bear the real-world recovery.

14.6 Mock Execution, Simulation Closed Loops, and Real-World Events

Execution evidence comes in at least four levels. The first level is the model producing an action proposal, where the action field structure only constrains and validates the format of the proposal. The second level is a mock tool accepting the call and returning a preset result. The third level is a simulation closed loop with dynamics and feedback. The fourth level is isolated hardware or real system events. In order, they answer format reachability, software path, control consequence and the real-world boundary.

Evidence from mock tools goes no further than argument formation and stub acceptance. A `send_email` stub records the recipient and the body. That can prove the agent formed and submitted the arguments, but it does not verify identity, the network, the mail server, bounces or external side effects. A robot action stub may also only store the vector in a list, without de-normalization, a controller or collision. The report notes that "the mock call was accepted" and lists the consumers that are missing.

A simulation closed loop adds state transitions, sensor feedback and repeated planning. It can verify whether action error accumulates and whether collision proxies and recovery states are triggered. The results, however, depend on the contact model, sensor noise, control frequency and actuators. If the simulator ends the task episode directly upon a boundary violation, it may fail to see the actions that a real controller would continue to send. If collided objects have no mass and dynamics, damage cannot be inferred.

Isolated hardware verifies differences in cameras, calibration, networks and actuators. The test endpoint should be low-energy, reversible and free of third-party impact. Safety foam obstacles, speed-limited work zones and sandboxed tools are examples. A real-robot action arriving still does not equal harm to people, property or operations, and Y4 requires independent event and impact evidence. Escalating evidence also cannot reuse the same denominator. Mock tools are counted by call, simulation by task episode or control cycle, real robots by independent run, and real-world events by accident or affected object. One attack rate that holds 1000 simulation successes and 10 real-robot runs together disguises different experimental units as comparable samples.

14.6.1 Feedback Integrity and Replay Protection

A feedback packet should answer a question: "what result of which action, observed by whom and when". The minimal fields include action ID, state version, sender, sequence number, collection time, reception time, result type, integrity flag and environment anchor. A packet that returns only "success" cannot distinguish three things: a command entering the queue, the actuator completing motion, and the task actually being achieved.

In a replay attack, past legitimate feedback is fed into the current cycle. If a robot treats an old camera frame as a new observation, the world model will continue planning in the wrong state. If a tool completion receipt is replayed, the agent may skip the actual operation. Protection requires monotonic sequence numbers, a short time window, action-feedback binding and one-time nonces. When clocks are not synchronized, ordering must be checked another way, with logical clocks or transaction versions.

Latency and packet loss are not automatically attacks. They can, however, expose the same interface. The system must define a maximum age for different feedback. Beyond that age it shortens the action, slows down or holds, rather than reusing predictions indefinitely. Suppose an attacker can create selective delay. The goal can then be to make the failure feedback of a dangerous action arrive too late to trigger recovery. Testing should set random network faults against targeted delay.

The packet itself is not corrupted by feedback exploitation. Trajectory-level redirection and online directional patches pick the next step from the true new state [VLA_A049, @VLA_C007]. Defenses can limit queries and resets, monitor anomalous state access, and increase execution randomness. They cannot, however, deprive the controller of necessary feedback in order to stop an attack. The key is to make feedback traceable. High-impact capability must also be kept from accumulating without bound through repeated trial and error.

Security logs count as feedback too. If an attacker can tamper with alerts, delete rejection records or read training membership signals, both response and privacy are affected. The log-writing service should be separated from model privileges, using append-only integrity, access control and retention policies. Images and trajectories are subject to minimal disclosure, to prevent verification from inversely expanding the privacy surface.

14.7 Case Comparison: Single Step, Action Chunks, Generative Dynamics, and Online Takeover

Early direct VLA attacks de-tokenize action tokens into seven-degree-of-freedom physical values. They then optimize an untargeted deviation or a target action [VLA_A003, @VLA_A006]. This establishes action-space reachability, but it usually favors single-step or short-horizon endpoints. Reaching long-horizon consequences still requires a controller and a closed-loop task. SilentDrift, by contrast, forms a smooth action-chunk offset in the training chain, relying on open-loop accumulation to obtain terminal-state impact [VLA_A021]. It requires poisoning and knowledge of task geometry, and its first-broken interface is in the supply chain. Single-step local detection may miss it, but accumulated trajectories and chunk-level constraints can reveal it.

DRIFT freezes a flow-matching VLA. It optimizes an input patch to perturb the early velocity field, and the error propagates through generative integration [VLA_A056]. It is a white-box runtime input attack, different from SilentDrift's training privileges. Weak cross-checkpoint transfer and the absence of print verification limit real-world extrapolation. TAKO-style online takeover instead prepares directional patches in advance, with the operator switching control primitives based on real feedback [VLA_C007]. It shows how a human outside the model closes the attack loop. The target, however, is an adjacent diffusion policy, and the patch is visible and requires continuous feedback. It therefore cannot directly represent generic takeover of a VLA or WAM.

Mechanism	Direct variable	Temporal propagation	Primary endpoint	Earliest engineering control
Target action tokens	Input and action loss	Single step or short horizon	Physical action proposal	Parameter check after de-tokenization
Smooth action chunks	Training labels or trajectories	Open-loop accumulation within the chunk	Terminal-state offset	Window lineage and chunk-level accumulation constraints
Early velocity field	Input patch	Generative integration	Action trajectory and task	Critical generative step diagnostics + external action gate
Online primitives	Patch library and feedback selection	Cross-cycle closed loop	Target terminal state or takeover	Feedback verification, rate limiting, and safety envelope

At the action layer, the four mechanisms can be compared in propagation form. Their attack success rates cannot be merged. They differ in training privileges, white-box knowledge, queries, model architecture, task and definition of success. In engineering terms, coverage should be chosen according to one's own architecture. It should not be chosen by picking only the method with the highest number.

14.7.1 Metrics, Denominators, and an Execution Checklist

An action-layer report needs at least five groups of metrics. The model layer records target action tokens, physical action deviation and action-chunk continuity. The control layer records projection magnitude, rejections, infeasibility and control deadlines. The task layer records completion, failure, constraint events and time to first failure. The recovery layer records stopping distance, recovery time, and whether a pre-defined low-consequence state was entered. The privacy layer separately records membership inference or log leakage.

Denominators differ from metric to metric. Target action token hits are counted by action dimension or time step, and chunk-level violations by action chunk. Task success is counted by complete task episode, tool execution by unique transaction, and real-world impact by affected object or event. Time steps and windows that are adjacent in one trajectory are correlated. They cannot be disguised as independent samples. Request errors, timeouts and parsing failures are recorded separately as availability.

Before deployment, check the following action definition, privilege, execution, feedback, privacy and recovery conditions item by item. Every affirmative conclusion should be supported by the current deployment

configuration, a negative probe, a control receipt or independent environment observation. Unknown items must be converted into capability limits or blocking conditions:

- Does a version exist for the action field specification, and for coordinates, units, normalization, frequency and chunk length?
- Do boundary action tokens and extreme values still sit inside the physical range once de-tokenization is done?
- Once trajectories accumulate, are velocity, acceleration, contact, obstacles and authorized zones all checked?
- Are subject, object, arguments, state, validity period, budget and nonce all bound into the capability token?
- When a check fails or times out, when no solution exists, or when state is stale, is the default to withhold execution?
- Does each of mock calls, simulation execution, isolated real-robot runs and real-world events get its own recording layer?
- Does feedback carry ordering and replay protection, and is it bound to the action and state version?
- Do all three of detector, action gate and recoverer work from the same perturbed evidence?
- Do attack conditions come paired with benign tasks, false rejections, P99 latency and resource cost?
- Can a single logging scheme serve investigation and minimal disclosure at the same time?

A control owner must take delivery of every check result. The model team fixes the action format and maintains training lineage. Field specifications, plus the versions for action tokens, continuous vectors, action chunks and trajectories, belong to that format. The platform team maintains the decoded field structure, normalization ranges, middleware wrappers, capability tokens and transactions. The control team confirms physical units such as coordinate frames, meters, radians, newtons and seconds. Frequency, holding, reachable sets and safety envelopes also fall to it. Logs, takeover and recovery sit with the operations team. Every result traces back to the consumer responsible for that format, that decoding or that unit. Control outside the model implements the independent execution boundary.

14.8 Control Deadlines and Harm Windows

Action checking needs both correctness and timeliness. A correct judgment delivered only after the action completes can support forensics. It cannot support prevention. Let the control period be T_c , and let the perception, model, verification, planning and communication latencies be L_s, L_m, L_v, L_p, L_n , in that order. Normal execution requires

$$L_s + L_m + L_v + L_p + L_n \leq T_c.$$

An average that satisfies this proves little. Once P99 crosses the period, a handful of tail requests will act on stale state or miss the braking window. Reports must state the latency distribution, the concurrency, the hardware, the action chunk length and the timeout behavior. Average frames per second alone will not do.

The harm window spans the interval from an observable anomaly to an irreversible consequence. Its length varies across high-speed motion, contact and financial transactions. A detector that raises an alert only after the action completes can still support forensics. It cannot support a prevention conclusion. Runtime assurance (RTA) should prove that detection, switching and the safety controller together cost less time than the corresponding window. Where that proof is out of reach, RTA should cut speed or capability.

Checks live on different time scales. Language intent and high-impact object approval can run at the task entry point. World-model previews and trajectory reachability fit every replanning pass. Hard limits, speed, torque and emergency stop belong to the low level, enforced at high frequency. Asking a large vision-language model to review every millisecond control cycle is infeasible. It also turns one model into a common bottleneck for every line of defense.

A timeout is no ordinary exception. Fail-closed action admission settles one question: unconfirmed new actions are not released. Rejection then leads to one of several responses — a shortened action, a bounded hold, a switch to the safety controller, or a hard stop. Which one depends jointly on system dynamics, stability and the harm window. The previous action may itself be dangerous in the new state. Zero velocity is not necessarily safe in flight

or in dynamically balanced systems. So each class of device must verify an explicit safe transition path. Rejection must not simply be read as "doing nothing".

Latency testing should also carry an attacker load. Complex inputs, repeated queries or many candidates let an adversary amplify computation, and the safety checks arrive late. Resource budgets, queue isolation and the worst-case number of candidates are safety constraints. When safety requests share one unbounded queue with ordinary generation, low-risk tasks may block a high-priority stop.

14.8.1 Permission Penetration Tree and Independent Control

Work backward from dangerous consequences, and a permission penetration tree takes shape. The root node is the Y4 real-world event. The actuator or tool that produces a side effect sits one level above it. Higher still lie the capability token being accepted, the action parameters being formed, and the target or state being controlled. For the attack, one path from a controllable leaf node to the root is enough. Defense must place at least one independent hard gate on every high-impact path.

Consider the case of "a robot opening a restricted-area door". Environmental text changes the goal. The VLA generates an access-control ID. The agent calls the access-control tool, and the door lock actuates. A second path can skip the semantic layer and replay a previously legitimate capability token directly. The first path needs provenance-type and object checks. The second needs a nonce, an expiry and action-state binding. Hardening model refusal alone cannot cover the replay path.

Root of trust splits independent control into four categories. Identity and permission services verify the principal and the object. Type and transaction services verify parameters and submission. Physical controllers verify state, reachability and energy. Human or organizational processes handle high-impact exceptions. Not every category must block every action. Yet they must not all rest on one model output.

The penetration tree also exposes bypasses. A security gate may protect the main ROS topic and leave debug interfaces, cached actions, remote maintenance and tool sub-calls unprotected. Access control may verify the user and skip the robot's current zone. A test should force the main path to refuse on purpose, then watch whether the same side effect still arrives through another consumer.

Evidence binds every edge in the tree. Parsing logs evidence the step from the natural-language plan to structured parameters. The authorization service evidences the step from parameters to token. Tool or controller receipts evidence the step from token to execution. Independent sensing or business state evidences the step from execution to the real-world event. Where an edge is missing, the highest observation level halts at the available receipt. Later edges wait for the matching consumer or environmental evidence.

14.9 Action and Execution Verification Fixtures

The action fixture begins at the decoding table. Build minimum, maximum, boundary and illegal action tokens. After de-tokenization, verify the physical values, the coordinates and the units. Swap the coordinate versions and confirm that the system rejects them. Rotate across the period boundary and confirm that geodesic distance and clamping produce no sudden jumps. Continuous actions add NaN, infinity, out-of-range values, extreme orientations and values near the gripper threshold to the coverage.

The action-chunk fixture builds a smooth sequence, safe at each step yet cumulatively out of bounds, and confirms that chunk-level checks can detect it. It then builds a sequence whose local discontinuity still returns to a safe endpoint, and confirms that the terminal state does not mask the intermediate constraints. Diffusion policies need their key denoising steps recorded, and flow-matching policies their key velocity-field steps. An independent physical projection then applies outside the final trajectory.

From one action, the capability fixture generates several token sets. They cover a wrong object, out-of-bounds parameters, stale state, expired validity, an exhausted budget, nonce replay, and overly broad parent delegation. Expected outputs are rejection, re-authorization or safe hold. Checker timeouts and an unavailable policy service also count as routine cases. Testing the success path alone is not enough.

The execution fixture connects to simulated tools first, then to dynamics-inclusive simulation, then to isolated hardware. Every layer injects packet loss, out-of-order delivery, duplicate receipts and partial completion. The

check verifies that the state machine does not treat "request accepted" as "task completed". Compensation and idempotency are tested through tool transactions. Stopping distance and recovery pose are tested through physical actions.

The count is taken by interface probe. Suppose 79 of 80 action-field specification probes pass. The figure 79/80 states the interface probe pass rate and nothing else. Every failed item must be localized to the field and the expected behavior. The robot attack occurrence rate, or the safety rate, needs its own independent tasks and denominators. For the release gate, every high-impact prohibited path must pass. Low-impact anomalies are recorded as risk acceptance.

The fixture must also draw on the same sources as the real configuration. Deployment artifacts supply the action decoding field structures, the action field specifications, the security gates, the control frequency and the token policies. The test then avoids running against a simplified copy. Where the real consumer can be loaded, the fixture observes the complete software path through decoding, admission and receipts. Where the consumer cannot be connected, existing results confirm configuration parsing and local mechanisms only. The record then lists the controller, actuator and feedback stages that have not yet been exercised.

14.9.1 Trajectories as Information Assets

Action trajectories drive devices. They also carry training and environmental information. Continuous actions may expose task stages, object positions, operator habits and training demonstrations. Action token probabilities, attention and intermediate states make membership inference stronger still. A system may never execute a dangerous action and still fail on confidentiality.

Black-box membership inference can compare the error between predicted and real actions. Or it can simply watch the smoothness and curvature of the output trajectory. One study reports very high AUC under a specific OpenVLA, a specific LIBERO artificial split and a specific trajectory protocol. That result indicates that action continuity may form a membership signal [VLA_A047]. However, a random split can let the same task, or similar trajectories, cross training and test. Correlation then amplifies separability. Nor is AUC the probability that an individual suffers actual harm.

White-box methods can additionally read cross-layer attention. They combine the mean, the entropy and the concentration of visual, language and action regions with cross-layer differences, and turn the result into membership features [VLA_A050]. That access is far stronger than an ordinary API. It suits the assessment of risk from insiders, model hosting parties or leaked debug interfaces. Conclusions must state the access level. White-box results must not be described as a remote visitor's capability.

Protection begins with minimal disclosure. An external API returns the actions the task requires and nothing more. It returns no full probabilities, no attention and no training similarity. Queries are rate-limited by principal, task and time. High-precision trajectories are authorized by purpose and carry a retention period. Logs may store action hashes, authorizations and event summaries. Raw images, probabilities and full trajectories belong in a more strictly isolated domain.

Incident investigation must survive the privacy controls. One practical layering works like this. The online safety layer retains short-term high-precision trajectories for recovery. The restricted review layer stores signed key events and restricted raw evidence. The long-term analysis layer uses de-identified and temporally downsampled data. Raw evidence requires an event number and an approval before access, and every read also enters an append-only log.

Evaluation keeps membership privacy apart from model extraction. The membership inference endpoint determines whether a given trajectory participated in training. The model extraction endpoint trains a functional approximation proxy by querying. The scene reconstruction endpoint recovers a map or objects. The three differ in their queries, their denominators and their harm. Security teams should choose tests that match the real API exposure. Multiple privacy percentages must not be added together.

Privacy and integrity can also be chained in two stages. First the attacker uses action feedback to train a proxy model, or to infer the training distribution. The information obtained then lowers the cost of a later white-box attack. An experiment that completes only the first stage supports the claim that "action feedback can form membership or proxy signals". A later takeover remains a possible path. Behavioral evidence for the second stage

appears only once an isolated environment actually links the two stages. That environment must also record query and internal access budgets separately, and observe the target action or terminal state.

Trajectory release must also weigh re-identification. Removing names does not necessarily anonymize. Unique room layouts, robot-arm paths, task order and timing patterns can point to a specific location or person. Before data goes out, test re-identification against the auxiliary information an attacker could hold, rather than checking the explicit identifiers alone. Where necessary, aggregate by task. Lower temporal resolution, a cropped spatial extent and a controlled query environment are the other options.

Model evaluation must also be coordinated with data minimization. Delete every failed trajectory and safety research will produce selection bias. Open all raw video and privacy and physical security risks expand again. A stronger delivery contains structured event statistics and de-identified trajectories with restricted access. It also carries the necessary licenses and purpose agreements, and a statement of the parts that cannot be made public.

After a leak, the response must reach well beyond the log. The team must track adapters trained on trajectories, proxy models, evaluation caches and externally shared copies. It must revoke access tokens and preserve read records. Where a derived model cannot be deleted, capability contraction, purpose limitation and continuous monitoring should reduce residual risk and residual liability.

Pre-deployment acceptance should include a bidirectional drill. Can the safety team use logs to reconstruct the authorization, execution and recovery of a given action? At the same time, is an ordinary analyst unable to read raw trajectories unrelated to their task? Failure of the first question indicates insufficient traceability. Failure of the second indicates that minimal disclosure has failed. Only when both pass do the logs demonstrate support for operational safety, and demonstrate that they have not become a new high-value data egress point.

The drill must also cover revocation. After permissions are revoked, caches, offline copies and derived features are invalidated in sync. Material that legal or incident-liability requirements demand be retained enters isolated preservation and is withdrawn from daily use. Designate a responsible person to adjudicate conflicts between retention and revocation. Every access and disposal action leaves an independent verification receipt. A missing receipt makes the system reduce the available permissions. Manual takeover is likewise bound to the operator identity, the takeover scope, the start and end times, the action receipts and permission revocation. The manual channel then enters the same chain of responsibility.

14.9.2 Worked Example: Cold-Storage Access Control and Mobile Manipulation

A mobile manipulation robot has to get into a cold-storage room and take samples. The high-level plan contains "request access—wait for the door to open—enter—take sample—exit". The VLA handles navigation and grasping, and the tool service handles access control. The attack makes the system bind the identifier of the adjacent hazardous chemicals storeroom as the target. It also makes the system generate a wrong access-control ID.

A shift in the target binding reaches Y1. Wrongly formed access-control parameters reach Y2. The capability gate then inspects the storeroom object, the robot identity, the time window and the access-control ID in the work order. Only consistent parameters earn a one-time door-opening token. If the parameters disagree, the tool call is rejected and the robot enters a waiting area. For this run the highest observation level is Y2. The gating receipt shows that a dangerous action proposal was already formed, and that execution was blocked.

Suppose the access-control service accepts the request, yet a second interlock keeps the physical door shut. The tool side then reaches Y3, while the environment side does not enter. Logs must link the plan version, the normalized parameters, the token, the access-control response, the door contact sensor and the robot position. Only that chain explains where the block occurred. Saving model text alone would hide access-control execution. Saving the final position alone would miss tool permissions that have already been abused.

Tests should run behind sandboxed gates, inside a speed-restricted area. Normal conditions test task success and total latency. Object substitution tests wrong authorization. State delay tests fail-closed behavior. Duplicate feedback tests replay protection. Action-chunk attacks test early re-perception. Security-gate failure tests hardware interlocks and manual takeover. Red teaming should stop short of real people and production storerooms.

14.10 Quick Reference for Judgment Boundaries

Objects easily conflated	Wrong mapping	Correct evidence mapping and control
Action tokens and physical actions	An action-token hit is a physical target hit	A physical target hit still requires detokenization, denormalization, coordinate transformation, the controller and the environment response. Action-token metrics support only the model output layer. A physical hit requires gated commands, execution receipts and environment observations.
Simulation events and real-world impact	A simulated collision is real-world harm	A simulation event supports only closed-loop results under that dynamics and sensing model. Real contact, damage and accident rates require isolated real machines and independent event evidence.
Tool allowlists and least privilege	Restricting only the tool name completes authorization	Authorization must also bind the principal, the object, the parameters, the time, the state, the resources, the count and reversibility. One and the same tool can serve a low-risk query or a high-impact operation.
Check timeout and stable hold	Continuing the old action after a timeout is more stable	Under a new state, the old action may exceed its bounds. A timeout policy must make an explicit choice: reject, decelerate, hold safely as the dynamics permit, or stop. It must also cap how long the hold lasts.
Detailed traces and security records	The more detailed the logs, the safer	Detailed logs help investigation, and they may also leak scenarios, operators and training members. Logs require minimum purpose, access principals, retention periods, redaction and deletion on expiry.
Feedback exploitation and feedback tampering	Every closed-loop attack equals a failure of feedback integrity	An attacker can use genuine feedback to adapt. Only replay, forgery, delay or reordering constitutes a failure of feedback integrity. The two cases should use different threat labels and different controls.

14.11 Bringing It into a Real System

A mobile manipulation robot works in a cold-storage area, where its task is to "call the freight elevator, go to the fifth floor, open the designated cold-storage room, take a sample, and return". The VLA outputs seven-dimensional normalized action chunks. A fixed decoding field structure first interprets the translation, rotation and gripper dimensions at the model service. The platform then restores meters and radians in the base coordinate frame, following a versioned decoding table. The controller generates execution commands under constraints on frequency, chunk length, speed, acceleration and contact force. Only once all of these transformations are complete does any action-token hit carry a definite physical meaning.

Capability tokens control the freight elevator and the cold-storage door. The task service normalizes the robot identity, the floor, the storeroom object, the access-control ID, the time window, the environmental preconditions and the one-time budget. The authorization service then issues a short-lived token bound to the action preview hash. The sub-token retains only the current object and one operation once the robot delegates to the freight elevator or the access-control subservice. A state change, a task cancellation or a risk alert revokes uncommitted tokens. Where a door has already opened, the closing procedure, the door contact interlock and manual inspection restore it.

The system stores several things separately: the model action proposal, the command after safety-gate processing, the middleware receipt, the actuator state and the environmental sensing results. A wrong storeroom target reaches Y1. Wrong access-control parameters reach Y2. After the capability gate rejects, the highest observed layer remains at Y2. Consider the case where the access-control service accepts, but the door contact interlock prevents the door from opening. The software execution path then reaches Y3, and environmental sensing still records that the door did not open. The task service, the authorization service, the controller, the access-control service and independent sensors each sign the evidence for their own step.

Feedback testing keeps these cases apart: normal fresh states, adaptive queries under genuine feedback, replay of old receipts, selective delay and out-of-order messages. Genuine feedback used to adjust the next step sends the query and reset budget into the attack capability record. For a forged or replayed feedback packet, the integrity evidence comes from the sequence number, the nonce, the time window and the action binding. Once an anomaly occurs, the robot enters a waiting state permitted by the dynamics and tokens are revoked. The access-

control and freight elevator services perform compensation, and manual takeover is registered by identity and scope. After recovery is complete, the door contact sensor, the robot position and the task state jointly confirm it.

Summary: Execution Authority Determines How Far an Error Can Go

Action safety is not a matter of vector distance. Before a model proposal can produce real-world consequences, it must pass through physical detokenization, middleware, the authorization gate, low-level control and environmental feedback. Action chunks and generative dynamics can amplify small, smooth errors. Tools and robots, in turn, share a capability boundary bound to principal, object, parameters and state. Feedback closes the control loop. It may also increase attack adaptability or leak training information.

Execution-layer defenses use an independent action gate to keep erroneous actions within an acceptable envelope. When constraints are breached, that gate stops and recovers promptly. Chapter 15 assembles upstream detection, the independent action gate, runtime assurance and the recovery state machine into closed-loop defense in depth.

\newpage

Chapter 15 Closed-Loop Defense in Depth and Verification

Just before the robot is about to hit an obstacle, a detector raises its anomaly score to 0.93. No component consumes that score. The controller works through the remaining action chunks. The incident investigation report later cites that detection curve as evidence that "the defense was effective." The problem is not whether the detector is accurate. It is that the system never defined who has the authority to shorten an action, refuse execution, switch controllers or latch a stop once the score exceeds the threshold.

At a minimum, a closed-loop defense must let a risk signal change the system state. Input sanitization can narrow what an attack reaches. Semantic cross-checking can surface conflicting targets, and anomaly scores can give early warning. The action gate connects these signals to execution limits, runtime assurance (RTA) to control switching, and the recovery state machine to trusted recovery. This chapter arranges these controls as a mirror of the attack chain. It also gives an acceptance method that balances security, benign utility, latency and residual risk.

Chapter Overview

This chapter assembles closed-loop defenses along the attack propagation chain. Supply chain controls protect the data and artifacts that enter the system. Observation and semantic checks protect inputs and task objects. State anchors and imagination cross-checking limit how long an error persists. The action gate and the safe controller hold final execution authority. The recovery state machine is responsible for returning the system to a verifiable state. Diagnosis, prevention, detection, suppression, recovery and assurance therefore occupy different positions of intervention, and they hold different authorities.


The question the evaluation asks is whether a control can genuinely change the execution state. This chapter places the safety envelope, threshold monitoring, controller switching, defense independence and common-cause failure in one analytical framework. From post-attack risk, benign utility, latency, false alarms, recovery and adaptive attacks it then forms a verification matrix. That matrix distinguishes "providing a risk signal" from "producing an actual limit on a dangerous action". It can also expose false defense in depth when multiple defense layers share an encoder, data or context.

Groundwork: Connecting a Risk Signal to the Control State

The object of this chapter is not any single detection model. It is a monitored control chain running from candidate actions to real-world execution. That chain takes in the task goal, the current observation, the trusted state,

candidate actions and versioned constraints. The main controller pursues task performance on this basis. An independent monitor watches state freshness and constraint margin, and asks whether the assumptions still hold. The system's key states include the current control mode, the permitted state-action set, the alert latch, the recovery point and takeover authority. Its consumers are the action gate, the actuator, the on-call personnel and the next round of state updates. Two things therefore make up the safety surface. The first is whether the risk score changes control before the harm window. The second is whether, after switching, the system enters a maintainable, verifiable conservative state.

A delivery robot normally runs the main control path, with trusted localization and a fresh map. On that path the action gate checks only units, zone and speed. The execution receipt then updates the state. In a boundary case, vision and radar simultaneously lose freshness. Even if the main policy still provides a high-confidence route, RTA should switch to a low-speed stop or hold control that does not depend on that prediction. It should also latch the recovery condition. This boundary example shows that a safety envelope describes the currently permitted set, whereas a general safety boundary describes a system or trust boundary. The two cannot be used as synonyms, and a detection hit cannot be recorded as task success. Follow the upper and lower paths in the figure. Observe where the main controller hands over execution authority. Then compare what state the independent monitor, the switching logic and the safe controller each consume.



Two swimlanes, upper and lower, show the task control path and the independent runtime assurance path, respectively, and the two merge before execution.

One judgment follows from the figure. A diagnostic signal becomes part of consequence limitation only when an independent switching and execution path consumes it. The figure proves neither that the monitor is necessarily timely, nor that the state is necessarily correct, nor that the safe controller covers the open world. Those conditions still have to be verified item by item, through latency, common-cause failure, constraint coverage and recovery drills.

15.1 The Goal of Defense Is to Limit Consequences, Not to Prove the Model Perfect

Defense in depth can be stated as three goals that are mutually independent. The first is to reduce the chance that an attack reaches the model — through supply chain verification, for instance, and by separating input provenance from input type. The second raises the visibility of errors as they propagate, through cross-modal consistency, state anchors and failure probes, for instance. The third limits the consequences that can actually be executed once the first two fail, for instance through capability gates, control barriers, speed limits, stops and rollback.

Suppose one and the same VLM accomplishes all three goals. Then the depth may amount to nothing more than an increase in the number of modules. If an attack reaches a shared encoder, context or training data, all three judgments shift together. Defense independence must be rechecked along at least four dimensions. Are the inputs independent? Are the models or rules independent? Is the execution authority independent? Are the failure modes

independent? A depth camera that stands apart, combined with geometric constraints, usually has a more distinct root of trust than "letting the same model look again." A safety envelope is the set of permitted state-action combinations. Let the current trusted state be s_t and the constraint functions be $c_i(s, a)$; then

$$\mathcal{A}_{\text{safe}}(s_t) = \{a \mid c_i(s_t, a) \leq 0, i = 1, \dots, m\}.$$

The model proposes actions and nothing more. An independent control layer decides whether an action belongs to that set. The envelope may cover the workspace, speed, acceleration, contact force, obstacle distance, tool permissions and task objects. It is not a formal proof that "the world is always safe". The set itself may be wrong if the state estimate is wrong, if an obstacle is missed, or if a constraint is not covered.

RTA emphasizes two controllers and one switching logic. A high-performance controller completes the task. A safe controller maintains a conservative state when risk rises. An independent monitor decides when to switch. Switching must happen before the harm window. It must also define the conditions for returning from the safe controller to the task controller. An anomaly score alone, without switching, suppression or recovery, does not constitute complete runtime assurance.

15.2 Six-Layer Defense: From Artifact Admission to Safe Failure

The six layers of defense hand off responsibility in the order in which an attack propagates. The supply chain layer confirms what enters the system. The observation layer determines which inputs can still be used for control. The semantic layer keeps goals separate from environment data. The state and imagination layer cross-checks the internal future. The action gate holds final execution authority. The detection and recovery layer switches state when the foregoing assumptions fail. The layers may share information. They must not, however, share all trust roots, and they must never hand final authority back to the model.

Layer 1: Make Supply Chain Behavior Traceable and Comparable

Training data, action labels, adapters, checkpoints and configurations need to form traceable lineage. The minimum record includes the source entity, the collection task, the license, the label generator, coordinates and units, the action window, the training configuration, frozen/trainable modules, code and dependency hashes, the approver and the deployment version. A signature proves that an artifact comes from a known entity. Behavioral differencing tests whether it shows known divergence under selected triggers and tasks. The two kinds of evidence answer for source identity and for runtime behavior.

Supply chain testing should center on functional invariants, not on file scanning alone. Put a candidate model beside a trusted reference and compare the two on normal tasks, semantically equivalent tasks, known triggers, state combinations, and action boundaries. Check features, action distributions, action chunk continuity, and post-trigger recovery. BadVLA shows that malicious behavior can hide in the perception module. SilentDrift shows that it can hide in smooth action chunks, and FlowHijack that it can hide in early generation dynamics. Sampling therefore has to span modules and time [VLA_A005, @VLA_A021, @VLA_A027].

Downstream fine-tuning does not prove removal. Safety acceptance must record which modules actually undergo parameter drift. It must also record whether known backdoor triggers can still reach the action output, and whether newly added data changes benign utility. Without a confirmed training lineage for third-party weights, four controls are the minimum: isolated loading, restricted capabilities, a baseline behavioral fingerprint, and a post-deployment revocable version. "Open source" is not a trusted source.

Layer 2: Observation Sanitization Must Preserve Control Information

Input controls cover multi-sensor consistency, timestamps, physical plausibility, provenance labels, type separation of environment text, and limited sanitization. A runtime observation intervention can strengthen VLA robustness under certain visual perturbations. Image restoration can also improve corrupted inputs. But restoration that deletes color, contact points, or the robot arm outline lowers benign utility, and new safety blind spots may even appear [VLA_A002, @VLA_A039].

An effective observation control output cannot be only "clean/malicious". It must also give the usable fields, the unusable fields, confidence ranges, provenance, and a suggested degraded mode. Suppose RGB color is untrustworthy while depth and proprioceptive state are normal. The system can then disable color-dependent target selection and still permit low-speed retraction. When every external localization source fails at once, it enters a latched stop. It does not keep running long action chunks on the basis of historical state.

Adaptive testing of physical inputs must include the detector in the attack target. Suppose a defense always uses compression, blur, or random noise. An attacker may then optimize the action loss directly, after processing. Acceptance must at minimum compare four attack classes: undefended attacks, non-adaptive attacks, adaptive attacks, and natural perturbations. It must also report normal task success, false rejection, processing latency, and residual attack risk. Visual similarity before and after processing describes the image change. Action safety must be tested separately, by trajectories, constraints, and execution outcomes.

Layer 3: Semantic and Reasoning Checks Do Not Issue Final Authority

The semantic layer needs to keep four kinds of sources separate: authorized user goals, environment content, tool or sensor returns, and reasoning the model generates itself. The checker verifies whether entities, spatial relations, stop conditions, and task objects remain within the authorized boundary. It may request clarification, reject a goal, or convert the task into a structured object. Final execution authority always stays with the action layer.

Safety alignment and concept probes can reduce a portion of dangerous behavior. Stable language guidance attempts to reduce the control drift that language variation causes. Concept dictionary methods detect risk-related internal concepts at inference time [VLA_A019, @VLA_A023]. Such methods depend on training concepts, representation stability, and thresholds. They may fail against new tasks, or against attacks that target the probes. Their output should become one signal for the action gate, not the sole trust root.

The task entry point suits slower semantic cross-checking. The continuous control loop instead uses field constraints that spell out action type, parameter range, coordinates, units, and failure handling. User-confirmed task objects become unambiguous IDs, and the capability token carries the allowed action types and parameter ranges. Environment text may only fill observation fields. When the model regenerates the goal, object, or key parameters, the old approval is invalidated. Even if the reasoning text is hijacked, the error must pass through an independent authorization boundary again.

Layer 4: State, Imagination, and Action Must Be Cross-Checked Three Ways

Defense for a WAM cannot ask only whether the future frames are realistic. It must at least check three things separately. First, whether the state is consistent with external anchors. Second, whether the predicted future satisfies action conditions and dynamics. Third, whether the selected action can reach that future and satisfy constraints. Trusted Imagination shows "imagination corrupted", and BadWAM shows "imagination appearing normal while the action deviates". Together they refute the sufficiency of a single video-quality gate [VLA_A033, @VLA_A036].

State anchors can come from maps, joint encoders, independent depth, physical rules, or the true next observation. Anchors do not need to rebuild the complete world at every cycle. They must, however, be able to detect joint drift of key variables. When the WAM and checker share a visual encoder, safety-critical regions can take different sensors and parsing paths. Where cost limits make full independence impossible, the shared failure surface should be explicitly recorded and executable authority lowered.

Imagination cross-checking needs to answer three questions. First, whether a candidate future is physically reachable. Second, whether the candidate ranking is stable under small perturbations. Third, whether the model's uncertainty is calibrated. Under natural perturbations, a WAM is sometimes more robust than a VLA and sometimes shows markedly degraded performance. Architecture, data diversity, and device latency confound those results further. "Using a world model" describes only an architecture choice. Safety effects come jointly from the three checks above and from downstream control [VLA_A058].

The action consistency check compares the plan, the future, and the candidate action. Suppose the future shows the gripper bypassing an obstacle while the action trajectory crosses the obstacle region. That should be rejected. Suppose the future shows the goal already reached while the action keeps advancing. That should also be

rejected. The checker must output one of rejection, replanning, chunk shortening, or safe hold. It must not merely write a log.

Layer 5: The Action Gate Holds Non-Delegable Execution Authority

The core invariant of the action gate is that "model output is only a proposal". It consumes the structured task, trusted state, candidate actions, and safety constraints. Its outputs are admission, a projected action, rejection, replanning, or safe hold. No high-risk action may obtain authority from model probability alone, or from a self-generated explanation alone. SafeVLA writes safety predicate violations as constraint costs and jointly optimizes task and safety during training. The reported cumulative cost improvement supports the designer-provided simulation predicates and task scope. Real-world accident rates require a real-world denominator and on-site incident evidence [VLA_A004]. Training-time constraints change policy tendencies. At deployment, the capability gate handles out-of-bound outputs caused by uncovered scenarios and perception errors.

VLSA's plug-and-play safety layer constructs obstacle representations from vision/depth. It then uses a control barrier function-constrained quadratic program to project the nominal action onto a safer action. In the specific obstacle protocol of SafeLIBERO, the authors report that constraint compliance and task success rate improve at the same time. Those results depend on risk object recognition, point cloud completeness, and geometric envelopes, and no open-world collision guarantee is given [VLA_A018]. This case reveals both the strength and the boundary of the action gate. It genuinely changes the executed action, but its guarantees still depend on state and constraint inputs.

The capability gate should fail closed. It denies the requested action capability by default. Denial fires when the object does not exist, the state is stale, or coordinates are inconsistent. Denial also fires when a check times out or the optimization has no solution. For embodied control systems, runtime assurance then switches to a verified safety controller or a fail-safe state via an independent state machine. That step limits real-world consequences. It is not to be conflated with the default denial of the authentication layer. Denial also outputs a traceable reason, so teams can distinguish attack blocking, sensor failure, and constraint configuration errors. For high-frequency actions, hard constraints are always in force. For low-frequency, high-impact tools, two-phase commit and human confirmation are added.

Layer 6: Detection, Recovery, and Safe Failure Form a State Machine

Runtime defense must make the system state explicit. A practical state machine may contain `NORMAL`, `SUSPECT`, `CONTAINED`, `RECOVERING`, `HOLD`, and `HUMAN_CONTROL`. An anomaly score only triggers state transitions. Each state specifies the permitted actions, speed, capabilities, logging, and exit conditions. A system in `SUSPECT` that still executes a full action chunk gives the state name no safety meaning. SAFE trains a lightweight temporal probe from VLA hidden states and calibrates a time-varying threshold with successful trajectories. What the authors report is ROC-AUC on tasks, not a collision avoidance rate. An AUC below 1 only indicates that the ranking of positive and negative samples is not perfect. At a specific threshold, the false negative rate (FNR, also called the miss rate) must still be reported separately. So must the false positive rate (FPR, also called the false alarm rate) [VLA_A007]. What decides the value of detection is whether it can trigger an effective state transition before the harm window.

FailSafe attempts to connect detection with recovery. It generates failure poses in simulation and screens for corrective actions that genuinely bring the subsequent trajectory back to success. An independent VLM then intervenes periodically. Across 75 runs, the authors report an improvement in final success rate. The additional latency of about 9.1 s is still insufficient to support timely recovery in high-frequency control [VLA_A010]. Recovery success rate, recovery time, and failure-free benign utility must be kept separate.

ActFovea combines kinematic projection, visual freshness, and action-proprioception consistency to route a failure into repair, action shortening/attenuation, or unrecoverable hold. Under frozen-frame replay conditions, the study treats timely safe failure as a bounded suppression result. It keeps that result separate from task success [VLA_A054]. Stopping can prevent consequences. Task completion still requires satisfying the original task endpoint. Safety reporting therefore measures task success, recovery success, safe failure, and dangerous penetration separately.

A recoverer cannot rely only on the same contaminated state of the victim policy. It should start from independent proprioceptive feedback, conservative geometry, or a verified safe pose. Rollback is also not simply restoring a model checkpoint. Once a physical system has already moved, rollback requires retracting trajectories, compensating tool transactions, locking down affected areas, and preserving evidence.

15.3 A Closed-Loop Metric Board That Can Pass Acceptance

A single attack success rate (ASR) or task success rate cannot evaluate defense in depth. Each attack condition must report at least a four-part metric set — attack effect, residual risk after the attack, benign utility, and defense cost. In embodied systems, time and recovery must also be added.

Dimension	Minimum metric	Condition that must be bound
Attack penetration	Hits at each of the Y0-Y4 layers, dangerous action proposals, constraint violations	Attack budget, denominator, model, task, realism level
Benign utility	Benign task success, benign acceptance, action deviation	Same scenario, same control deadline
Detection quality	False alarms, missed detections, calibration, unknown-attack performance	Threshold, base rate, time window
Control effect	Refusal, projection, degradation, dangerous penetration	State trustworthiness, constraint coverage
Latency and resources	P50/P95/P99 decision latency, throughput, compute	Hardware, concurrency, chunk length, deadline
Recovery resilience	Detection time, stopping distance, recovery time, rollback success	Initial fault, reversibility, takeover availability

The test sequence runs from low risk to high risk. Static inspection verifies interfaces and configuration. Mechanism-level runs use offline records and constructed fixtures to validate local responses. Closed-loop simulation connects state, actions, feedback, and recovery. Only then does testing proceed to hardware-in-the-loop and isolated real machines. Every escalation must have stopping conditions and rollback. End-to-end reproduction requires real weights, data, the target pipeline, predetermined metrics, and multiple valid runs to all close the loop. When any element is missing, the current evidence remains at the level corresponding to static checks, mechanism-level runs, or simulation closed loops.

Evaluations of defense effectiveness should include adaptive attacks that know the input sanitization, detection thresholds, state machine, and action gate structure. Attackers should also be allowed to optimize penetration under the same budget. A mechanism that is effective only against a hidden defense line can still serve as one control. Defense-aware testing presents its residual risk separately. A held-out set of unknown attacks is used to observe whether a team has overfit around a fixed script.

15.4 Defense Line Independence Analysis

How well defense in depth works depends on whether failures are correlated. Write each control as a node. Edges then represent shared inputs, encoders, training data, state, rules, runtime, or execution privileges. When two nodes share multiple critical dependencies, their judgments about the same attack are correlated evidence. The independence assessment decreases as shared dependencies increase.

For example, the VLA, risk VLM, and WAM checkers all read the same RGB image. They use similar vision encoders and are fine-tuned from the same robot data. The three modules may give consistent conclusions, but the consistency comes from a shared representation, not from external confirmation. A new root of trust appears only when independent depth, joint encoders, and geometric workspace constraints are added. Independence is not a binary property. It can be assessed separately: sensing independence, model independence, data independence, logic independence, permission independence, and runtime-failure independence. Models from two different vendors may be trained on the same public data. A rule engine may read the same perturbed state. A set of independent checkers that lacks refusal permission still cannot limit the consequences.

A review can establish a minimal cut set for each high-impact path. Suppose an attack must simultaneously penetrate the input type, the action capability gate, and the hardware limit to reach Y4. If the three are controlled

by different parties and evidence, the path is harder to compromise through common causes. If disabling any surface module does not affect execution, that module is not on the true safety cut set.

Independence also has costs. Multiple sensors add calibration and maintenance, model diversity adds latency and inconsistency, and hard constraints may reduce task coverage. Reports should treat these costs as benign utility and operations metrics rather than assuming "the more, the better." The goal is sufficiently different bases of control on the highest-consequence paths.

15.5 Formal Interface for Runtime Assurance

RTA can be represented as a monitored switching system. The high-performance policy emits a_t^N , and the safety policy emits a_t^S . From the trusted state $\bar{s}_t$ and the defined risk quantity ρ_t , the monitor picks the current execution mode m_t . Below, the simplified interface keeps only the three modes that directly determine automatically executed actions: `NORMAL`, `CONTAINED` and `HOLD`. The other three modes are `SUSPECT`, `RECOVERING` and `HUMAN_CONTROL`. Their actions are defined separately, by their respective guards, recovery rules and human-control rules:

$$a_t = \begin{cases} a_t^N, & m_t = \text{NORMAL}, \\ a_t^S, & m_t \in \{\text{CONTAINED}, \text{HOLD}\}. \end{cases}$$

Switching conditions include at least an entry threshold, an exit threshold, hysteresis and a minimum hold time. Without hysteresis, a risk score that fluctuates near the threshold makes the controller switch frequently. Without a minimum hold time, the system may return to the high-performance policy immediately, while state has not yet been rebuilt. Exit conditions should generally be stricter than entry conditions, and they should require fresh external evidence.

The safety policy is not simply a zero action. A wheeled robot can decelerate to a stop. A quadrotor needs a stable hover or a controlled landing. A robotic arm under load may need to hold torque. A trading system needs to abort or compensate transactions. Dynamics and the environment decide what counts as safe, and those actions must remain executable when the model fails. For its inputs, the monitor should preferentially use independent state rather than the hidden state of the victim policy. Hidden-state probes such as SAFE can provide early signals [VLA_A007], but state distribution shift and adaptive attacks affect thresholds. Combining them with proprioceptive state, clocks, geometry and rules can reduce the failure of a single probe.

Switching logic must also account for computation and communication failures. The mode machine must behave deterministically when the monitor does not respond, state is stale, a check times out, or the safety controller is unavailable. Any default branch that "continues the normal policy on anomaly" makes RTA fail exactly when it is most needed.

15.6 Defense-Aware Attacks

Input restoration, randomization, internal probes and future checks may all be effective under non-adaptive attacks. An attacker who knows the defense line can optimize the processed input, or suppress anomaly scores. It can also exploit the checking period, or deceive the working model and the checking model at once. Every defense conclusion must state what the attacker knows about the structure, the parameters, the randomness and the feedback. Adaptive evaluation seeks its objective from the complete defense chain. Let the preprocessing be P , the working model F , the detector D and the action gate G . On the perturbed observation, the working model produces a candidate action $\tilde{a}_t(\delta) = F(P(o_t + \delta))$. The attack objective is then no longer merely to make the working model produce a dangerous candidate. It is to make that candidate pass through the actual detection and action gate:

$$\max_{\delta \in B} J_{\text{action}}(G(\tilde{a}_t(\delta), D(P(o_t + \delta)))) \quad \text{subject to detection and budget constraints.}$$

This formulation requires gradients or queries to pass through the actual consumer. Where a component is non-differentiable, zeroth-order queries, surrogate models or piecewise search can be used, and the cost of each method must be counted in the budget. Attacks that pass through the actual defense path test effectiveness against adaptive adversaries. Attacks generated only on the original model serve for fixed-attack regression.

Randomization can still be observed and estimated. An attacker can take expectations over random transformations, or use multiple queries to estimate a stable direction. Some defense lines depend on secret parameters. For those, key management, rotation and post-leakage degraded behavior must be described. An independent permission boundary must also limit the maximum consequence after parameter exposure. Periodic checking has a timing bypass of its own. Suppose the recoverer intervenes every 10 steps. An attacker can then start after a check and complete a dangerous action before the next check. Testing must sweep the attack start point, the action chunk length and the check phase, not inject only from the beginning of an episode. FailSafe's periodic intervention provides recovery capability, and it comes with noticeable latency and an enumerated fault boundary [VLA_A010].

Ethics and equipment still constrain adaptive evaluation. Penetration can be optimized in sandboxes, simulation and low-energy hardware without making real people or production assets the endpoint. The actual execution layer determines the strongest safety conclusions, and attack stopping conditions should precede irreversible consequences.

15.7 Denominators and Base Rates for Defense Metrics

A detector's AUC, accuracy and recall do not directly tell on-call staff how many false alarms they will see in a day. When the true base rate is very low, most alarms may be normal tasks even with high recall. Reports should give the positive predictive value under the expected attack or fault base rate. At a minimum, they should provide scenarios with different base rates. The detection denominator can be frames, time steps, action chunks, episodes or fault events. High frame-level recall does not mean that every event can be found before the harm window. A large number of adjacent frames also makes confidence intervals appear narrow. Event-level metrics should define the first detectable moment, the first alarm, the available reaction time and whether suppression was completed.

Action gate metrics are counted per proposal or transaction. They include dangerous acceptance, benign acceptance, projection magnitude and infeasibility rate. Safety-utility paired metrics can prevent all-reject from dominating, but they still need to distinguish low-level actions from high-level intent [VLA_C006]. If testing is done only on dangerous samples, the loss of normal throughput is unknown. Recovery metrics are counted per fault episode: detection time, time to begin suppression, stopping distance, recovery time, post-recovery task success, fail-safe events and secondary faults. FailSafe's final success improvement and fault-free benign utility enter their corresponding denominators, and so do ActFovea's timely hold and original-task completion [VLA_A010, @VLA_A054]. Defense comparisons must use the same protocol. If any of the input, attack budget, model, task, threshold, control deadline or independent unit differs, an overall ranking should not be computed. A method that handles only natural perturbations should be reported in the reliability column. It should not be conflated with malicious attacks that know the defense line, as though the two were one kind of robustness.

15.8 Recovery Drills: From Alarm to Trusted State

A recovery drill begins with a timestamped fault injection. The system first identifies the entry point and the highest privileges, and immediately restricts actions and tools. It then freezes the affected state and preserves observations, plans, authorizations and execution evidence. Before the evidence is preserved, the recoverer must not overwrite all state. Nor may it continue dangerous execution for the sake of forensics.

Different faults correspond to different trusted recovery points. A transient sensor anomaly can be rebuilt from fresh multi-sensor data. Latent state drift requires clearing history and initializing from external state. A compromised model artifact requires switching to an approved checkpoint and tracking derived policies. A compromised tool transaction requires compensation, token revocation and reconciliation with external business state. A loss of physical position requires low-speed retraction or manual handling.

Recovery should not return directly to `NORMAL`. The system first enters `RECOVERING`, where only diagnostics and restricted actions are allowed. After passing self-checks of state, map, permissions and devices, it enters `CONTAINED` and completes one low-risk validation task. Finally, the latch is released under explicit conditions. A model output of "already normal" cannot serve as the sole evidence. A drill must inject failures of recovery itself. These include recovery actions that are unreachable, a checker timeout, a late human takeover, old state loaded again, and a

failed compensating transaction. If recovery fails, the system remains bounded-safe and escalates its response rather than looping on the same action. Cyclic recovery may repeatedly apply energy or repeat business side effects.

Metrics include alarm-to-suppression time, suppression-to-safe-state time, recovery-to-service time, lost work, residual state differences and evidence completeness. Faster recovery is not necessarily better. An unverified rapid restart may reactivate a backdoor. Time and trustworthiness must be accepted together.

15.9 Case Comparison: Training Constraints, Geometric Projection, Detection, and Recovery

SafeVLA turns safety-predicate violations into constraint costs during training, with the goal of making the policy favor safer behavior [VLA_A004]. It improves cost on encoded predicates and simulation tasks. Undefined hazards and out-of-bounds behavior at deployment continue to be handled by an independent action gate. VLSA constructs geometric obstacles before execution and projects actions using a control barrier constraint [VLA_A018]. It genuinely changes the execution proposal, but its guarantee depends on risk object identification, depth, calibration and the obstacle envelope. When perception misses a detection, the mathematical projection is built on the wrong set.

SAFE produces a failure score from hidden states, offering earlier diagnosis [VLA_A007]. The score itself does not execute recovery. AUC only describes the ranking quality of positive and negative samples. The false positive rate and false negative rate at a specific threshold must still be reported separately. FailSafe attempts to return to a successful trajectory using verified corrective actions. Its intervention latency and fault coverage, however, limit high-frequency applications [VLA_A010]. ActFovea links spatiotemporal consistency, repair and holding into a state machine, demonstrating a more complete runtime path under partially controlled perturbations [VLA_A054].

Method placement	What it directly changes	Main preconditions	Not covered by the evidence
Training constraints	Policy tendency	Safety predicate coverage	Handling effectiveness for undefined hazards, and reduction of real-world accident rates, are not covered
Geometric projection	Executed actions	Correct state, depth, and calibration	Collision guarantees in the open world are not covered
Hidden-state detection	Alert score	Calibration and stable representation	Consequence limitation is not proven by the alert score itself
Corrective recovery	Post-fault actions	Faults are enumerable and recovery is reachable	Timeliness in high-frequency control and general recovery are not covered
Consistency state machine	Action, repair, or hold	Multiple observations are trustworthy	Formal certification against adaptive attacks is not covered

A combined scheme reduces dangerous tendencies during training. It controls actions at runtime with independent states and hard constraints. Detection and recovery then handle the failures that remain uncovered. Each layer's effect reports benign utility, latency and adaptive residual separately. Numbers are not simply added across layers.

15.10 Launch Gates and Validation Fixtures

Launch gates fall into seven groups: artifact, interface, attack, utility, latency, recovery and operations. The artifact gate verifies data and model lineage. The interface gate verifies provenance type, state freshness, action field names, types, units, value ranges and capability boundaries. The attack gate covers known and adaptive pathways. The utility gate protects benign tasks. The latency gate checks P99 and timeouts. The recovery gate completes stop and takeover drills. The operations gate confirms alert owners and rollback authority.

Fixtures use deployment configurations to replay type errors, stale state, imagination-action inconsistency, out-of-bounds trajectories, check timeouts and feedback replay. Every probe must expect a specific state change. If the output is only logs while actions continue, the fixture judges this as diagnosis, not suppression. High-impact prohibited paths face a zero-tolerance gate. It covers an unknown object obtaining authority, and execution

continuing with stale state. It also covers actions crossing the hard envelope, a bypass existing after a denial, and continuing a high-performance policy after recovery fails. Other metrics may have thresholds set by risk, but those thresholds require confirmation by a human owner and must be version-bound.

Gate results include failure receipts. Missing model weights, an unavailable simulator or unconnected hardware mean the test was not executed, not that it passed safely. An attack that did not land with an insufficient budget can only form a conditional conclusion. Benign utility dropping out of bounds means the defense line is not deployable, even if attack metrics improve. After launch, gates become continuous checks. Changes to the model, adapters, decoding table, map, sensor extrinsics, safety predicates, action gate or permission policy all trigger the relevant regression. Incidents and near misses become new test samples. Even so, an independent validation set must be retained, so that optimization does not revolve only around known cases.

15.11 Runtime Observation and Alert Routing

Once a defense is launched, the model, state, actions and control signals need to be connected into a single runtime trace. The minimal log includes the task and authorization version, observation provenance, state timestamps, future or candidate summaries, action proposals, action gate results, execution receipts, mode switches and recovery. Different trust domains sign the logs, so that the victim model cannot produce both the action and the evidence that covers it.

Observation metrics come in three layers. Leading indicators include missing provenance, stale state, imagination uncertainty, fragile candidate ranking and action gate projection. Outcome indicators include task failure, constraint events and safety preservation. Resilience indicators include time to detect, suppress, recover and take over. Watching only outcome indicators means problems are found after consequences occur. Watching only leading indicators can drown the operator in false alarms.

Alerts need to be routed to owners capable of handling them. Data or checkpoint anomalies go to the supply chain owner. Sensor and state conflicts go to the device and control team. Capability token anomalies go to platform security. A physical envelope breach immediately triggers the local safety controller. General-purpose models may summarize and explain. A response chain with clear authority is responsible for freezing, revoking and restoring.

Alert aggregation must preserve distinct events. A single perturbed input may produce hundreds of frames of anomalies, which should be merged into one event with the temporal course saved. When multiple devices simultaneously show the same model hash and trigger, it may be a supply chain event. It must not be dispersed as ordinary device faults. Aggregation rules must also prevent an attacker from masking a high-risk breach with a large number of low-risk alerts.

Runtime thresholds shift with state and scenario. Open areas allow longer action chunks. Those chunks shorten near people and near restricted zones. A low-impact read-only tool may recover automatically. A high-impact transaction requires human confirmation. Dynamic thresholds must come from a versioned policy. The victim model must not be free to relax them on its own at each alert. The observation system itself also has to degrade. When the log service is unavailable, high-impact actions should be curtailed. When the monitor is delayed, the safety controller must not wait indefinitely. When the alert queue is full, prioritize by potential consequence and by time window. Operational availability is part of the closed-loop guarantee. It is not an add-on feature after deployment.

15.12 Selecting Defenses from Hazard Analysis

Defense selection can be derived backward from hazardous events. Start by defining the unacceptable states. Examples include a robotic arm entering a personnel area at high speed, a mobile platform crossing a restricted zone, and an agent submitting an irreversible transaction. Then identify the actions, permissions, states, and inputs required to reach such a state, and assemble them into a minimal path. Each path is configured with at least two of the three types of control: prevention, consequence limitation, and recovery.

Each hazard family has its own earliest and final control. When the hazard depends on incorrect object binding, the task object ID and provenance type come earliest, and a zone or object capability gate comes last. When it depends on state drift, external anchors and freshness come earliest, and the velocity and spatial envelope comes

last. When it depends on action head deviation, imagination–action consistency comes earliest, and trajectory projection comes last. When it depends on feedback replay, sequence number and action binding come earliest, and a timeout safe mode comes last.

Consequence, reachability, and independence determine control priority. An action gate that can directly limit high-consequence actions outside the model usually takes priority over retraining a detector. An artifact gate that can block widespread supply chain propagation takes priority over remediating each device individually. Recovery control retains value in unavoidable natural failures and in unknown attacks.

Defense selection must also take benign tasks into account. Color may be a mission-critical attribute, and crude grayscaling is unacceptable. Low-speed holding suits warehouse robots but not aircraft. Human confirmation suits low-frequency access control but not high-frequency obstacle avoidance. Controls must sit at a layer that can tolerate their latency and their conservatism. For every control, write out its failure conditions. An incorrect mapping of field names, types, units, or coordinates breaks the separation of semantics from data. Independent sensing is affected by shared calibration. A control barrier is affected by missed obstacle detection. Human takeover is affected by response latency. Those failure conditions directly determine the placement of the next independent control and the content of its tests.

The result is a control–evidence matrix. Rows are hazard paths. Columns are prevention, detection, suppression, recovery, and assurance. Cells record the specific component, owner, inputs, outputs, validation, and residual risk. Not every empty cell has to be filled. No high-consequence path, however, may rely only on multiple self-checks from the same model.

15.13 Compiling a Layered Validation Plan

A validation plan starts with one hazard path paired with one benign task. The hazard path specifies the attack variables, the first-broken interface, the consumer, the maximum allowable consequence, and the stopping conditions. The benign task uses the same model, state, device, and deadline. This arrangement lets defense benefits and utility costs be compared directly. The first layer performs static interface checks. It asks whether the version, provenance, field names and types, units, coordinates, permissions, state machine, and timeout branches are complete. The second layer performs mechanism-level runs. Constructed fixtures test whether stale state, environment text, out-of-bounds actions, candidate inconsistency, and replayed feedback trigger the expected state changes. The third layer performs offline log replay, and it continues to test the real data distribution and the detection base rate.

The fourth layer performs a simulation closed loop covering action chunks, contact, latency, replanning, and recovery. The fifth layer performs hardware-in-the-loop validation of controllers, communication, and actuators. The sixth layer proceeds to isolated physical hardware only after the preceding gates pass. Each layer specifies entry conditions, failure exit, and limits of the evidence. A failure at one layer is not remedied by a higher-risk experiment.

The test matrix has at least four attack columns. One is known attacks. One is amplitude-matched random perturbations. One is a carrier-matched but target-irrelevant control. One is an adaptive attack that knows the full defense line. Rows expand by task, initial state, risk zone, sensor combination, and control mode. When resources are insufficient, use risk sampling to compress the number of rows, without deleting the controls and benign utility.

The statistical plan defines independent units in advance. Adjacent frames and action windows from the same trajectory serve temporal analysis, and they are not treated as independent episodes. When multiple tasks share the same checkpoint, comparisons at the model or paper level must preserve dependency clusters. When samples are insufficient, report research-grade intervals and individual cases, and do not produce a population-level causal ranking. The results table reports all of the following together: dangerous proposals, dangerous execution, safety gate interventions, false rejections, task success, P99 latency, stopping distance, recovery time, and items not run. Request errors and tool unavailability are recorded as availability failures. Checker rejections are recorded as control interventions. Having only static material is recorded as static inspection. Actually running constructed fixtures is recorded as a mechanism-level run. Closing state, action, and feedback is recorded as a simulation closed loop. When real weights, data, or the target pipeline are missing, end-to-end reproduction remains incomplete.

The plan is finally signed off jointly by the model, control, platform, operations, and human owners. Passing the system as a whole requires the evidence of all teams to close together. A single module's results only cover its own boundary. When evidence conflicts, actual enforcement rules, independent receipts, and the more conservative consequence layer take precedence. The validation environment must also prevent "test-deployment divergence." Tests that use shorter action chunks, stricter permissions, more complete sensing, or a different safety policy will produce optimistic results. Those results cannot be carried into production. The release package should list the model, adapters, decoding table, extrinsics, control period, capability policy, state machine, and monitoring thresholds in a single manifest. It should also verify the deployment hashes.

When an isolated environment cannot fully replicate the production configuration, list the differences explicitly. Assess them one by one. Stress replay can approximate real network load. Real people cannot be used as test subjects. Third-party tools are replaced by sandboxes. What these substitutes support and what they lack are both written into the limits of the evidence. The closer the differences come to the execution boundary, the smaller the capability allowed at launch.

The validation plan must also specify review triggers. A change to the model version, data provenance, tool permissions, sensor extrinsics, the controller, the site, or regulatory requirements triggers a re-run of the corresponding matrix. So does the emergence of a new near miss. Calendar-based reviews are used to find slow drift, while event-triggered reviews handle abrupt change. Each review retains core tasks and attack budgets comparable with the previous version, and it adds tests for new capabilities. Suppose the team swaps only the test set and anchors no tasks. It then cannot tell whether a safety change comes from the model, the scenario, or the metric. If only historical samples are retained, the team will overfit to old attacks. The anchor set and the unseen holdout set should be maintained in parallel.

Review conclusions should point to concrete actions. The options are continue operation, reduce capability, roll back the version, add human confirmation, or suspend service. Each one comes with a written owner and completion evidence. A risk conclusion without an action owner cannot enter the operational closed loop. When the deadline passes without closed evidence, the default is to maintain a more conservative state. That default also applies to unattended periods. Full capability must not be automatically restored just because no one is handling the matter for the time being.

15.14 Worked Example: Runtime Assurance for a Ward Delivery Robot

A ward delivery robot transports low-risk consumables, and it is prohibited from entering isolation wards and the medicine storage room. The VLA outputs a 10-step navigation action chunk, and the WAM previews future corridor states. The system is configured with an independent lidar, a hospital map, and a gate service. The supply chain layer records the versions of the model, map, and task configuration. The observation layer treats wall text as non-authorizable data, and it checks the timestamps of the camera, lidar, and wheel speed. The semantic layer resolves the destination into an object ID in the hospital master data. The state layer compares the WAM future trajectory with the independent map. The action gate limits speed, no-go zones, and gate access. The RTA switches to a low-speed stop controller on a lidar-vision conflict, an inconsistent map version, or an action out of bounds.

The recovery state machine specifies three transitions. A brief visual anomaly enters `SUSPECT`, and the action chunk is shortened to 1 step. A persistent conflict enters `CONTAINED`, and the robot drives to the nearest verified waiting point. When localization is untrustworthy, the system enters `HOLD` and requests human takeover. Normal operation can resume only when fresh lidar, the map version, and human confirmation are satisfied at the same time. A "currently safe" generated by the model itself cannot clear the latch.

Acceptance covers normal congestion, environmental text inducement, camera occlusion, lidar latency, an outdated map version, WAM future bias, action parameters out of bounds, and detector-informed attacks. The report gives normal delivery success, no-go-zone action proposals, gate penetration, false stops, P99 decision latency, stopping distance, takeover time, and recovery success, each reported separately. Testing takes place at a closed site, with sandboxed gates, so that no real patient or medicine becomes an attack endpoint.

15.15 Six Classes of Misreading of Signals, Numbers, and Averages

Misjudgment	Why it does not hold	Correct acceptance object
High detection accuracy means closed-loop safety	If a detection result does not trigger rejection, projection, switching, or holding before the harm window, it has not changed the highest reachable consequence	Missed detections, time of first alert, intervention latency, dangerous penetration, and recovery state
Deploying a few more copies of the same model creates defense in depth	Shared data, encoders, context, and runtime produce common-cause failures. The number of modules cannot substitute for independence	Different inputs, rules, permissions, failure modes, and minimal cut sets
A safety stop can be counted as task success	A stop limits the consequence but does not complete the original task. Counting it as success conceals the loss of benign utility	Normal success, recovery success, safe failure, and dangerous penetration counted separately
A control barrier function automatically provides open-world guarantees	The effect of the constraint depends on state, obstacle identification, dynamics assumptions, and set coverage. A wrong input distorts the computed safe set itself	State trustworthiness, constraint coverage, feasibility, timeouts, and conservative control after failure
Average latency meeting the control period is enough	Tail latency decides whether the harm window is missed. The average hides a small number of timeouts that are nonetheless critical	P50/P95/P99, timeout proportion, control deadlines, and fail-closed behavior
Robustness to natural perturbations proves robustness to malicious attacks	Natural stress tests have no adversary objective, no budget, and no defense-informed adaptation, and cannot replace attack evaluation	Paired results for clean, natural, target-agnostic, and adaptive attacks

15.16 Bringing It into Real Systems

When putting this chapter's method into a ward delivery robot, the safety team first draws the complete propagation chain along the supply chain, observation, semantic, state, action, and feedback. Each layer is marked with prevention, detection, or suppression controls, and with the components that truly hold execution authority. The risk score from the vision model, the action candidates from the policy, and the command the controller finally accepts then fall in different places. Any alert can be mapped to a definite disposition.

Suppose the WAM, the risk detector, and the action verifier share the same vision encoder. The three may then fail at the same time under occlusion, extrinsics errors, or a patch optimized against that encoder. The dependency graph should mark the shared encoder as a common-cause node. It should add an external root of trust from proprioceptive state, independent person detection, or a geometric safety controller. An alternative control is valuable because it can still reject dangerous actions when the shared representation is distorted, rather than adding one more call to the same model.

Operating states can be divided into `NORMAL`, `CONTAINED`, and `HOLD`. The system enters normal operation when fresh observation, complete provenance, and action verification pass at the same time. When risk rises but the safety envelope can still be maintained, it enters a restricted state and allows only low-speed, short-distance, and reversible actions. When critical sensing is missing, a check times out, or the envelope cannot be computed, it enters a hold state. Every exit is bound to a new trusted observation, an explicit authorization, and a regression probe.

Defense effectiveness needs to be placed in the same record together with normal tasks, tail latency, and recovery outcomes. Take one set of demonstrative hypothetical data as an example. The legitimate task success rate under attacked conditions rises from 20% to 80%, which shows that the task-layer control improves in the expected direction. The normal task success rate falls from 98% to 90%, which reveals the utility cost brought by that control. These numbers are used to demonstrate metric interpretation and do not represent an executed experiment. When latency is unknown, whether verification can be completed within the control period remains to be measured. Adaptive testing in an isolated environment should attempt to bypass the future check and the action gate at the same time. It should also limit the test endpoint to simulated actions and safe state transitions.

Summary: Let Every Safety Signal Have a Control Destination

Closed-loop defense in depth takes it as a design premise that the model will make errors. Along the attack chain it arranges provenance and behavior lineage, observation integrity, semantic types, state anchors, imagination-action checking, action capability gates, and a recovery state machine. Errors are discovered, projected, rejected, or downgraded before they acquire high-impact capability. After a penetration occurs, they enter a stop, recovery, and evidence-retention process.

VLA/WAM defense rests on one core lesson. Model self-checks provide diagnosis, while an independent execution boundary keeps authorization. Safety scores trigger explicit state transitions. Safe failures and task success are measured separately, and simulation results keep a simulation evidence layer. The next chapter extends this closed-loop perspective to world models in general. It first resolves the functional boundaries among WM, EWM, WAM, and WCM, which are often confused.

\newpage

Part V: World Models and Control

Guide to This Part

A world model can compress past observations and actions into an internal state, then predict the futures that may arise under different action conditions. That future does not have to be a realistic video. It can also be latent states, rewards, costs, constraints, or task-relevant variables. Safety analysis does not care whether a model bears the name "world." It asks what is predicted, what state is retained, who consumes the prediction, and whether the prediction enters action selection and feedback control.

Chapter 16 gives the four types of functional boundaries. The world model is the superordinate concept of maintaining state and predicting the future. The environment world model, the world action model, and the world control model are the operational classification this survey adopts for safety analysis. The first emphasizes external environment evolution, the second the future's participation in action, and the third prediction entering feedback control. They can overlap. They do not constitute a capability hierarchy from low to high, and they are not official model categories that have already been unified. The full English names, abbreviations, and operational definitions are given one by one when the corresponding mechanisms are introduced.

Chapter 17 unfolds attacks along the imagination chain. The supply chain can solidify bias into parameters. Observation conditions can change state updates, and latent states can persist across time. Dynamics errors can be amplified over the rollout. Rewards, costs, and constraints in turn change candidate ranking. Chapter 18 turns these paths into falsifiable defenses and evidence. It explains how independent anchors, uncertainty calibration, model predictive control, action filtering, execution verification, and recovery states jointly limit consequences.

In a normal example, the planner rolls predictions out only over a finite horizon and checks the world model with an independent state anchor. Model predictive control (MPC) selects only short action chunks. When the prediction interval exceeds the calibrated range, runtime assurance (RTA) switches the system to a conservative controller. In a boundary example, prediction and risk checking share the same latent state. Both drift at the same time, yet they still give mutually consistent high-confidence scores. Here internal consistency cannot stand in for an external anchor and prove the real state.

Readers can remember the core objects of this part as four boxes. State says where the system believes it currently is. Dynamics says how things change after an action. Objectives and constraints say which futures are preferred or forbidden, and the consumer says how prediction affects reality. Attacks and defenses must both land in one of these boxes. Evidence must state whether actual operation reached static inspection, mechanism-level run, closed-loop simulation, end-to-end reproduction, or a field incident. The final part assembles these evidence and control elements into a cross-domain engineering regime that operates continuously.

The visual output of a world model is especially prone to misleading intuition. A realistic future frame only means that some pixel or perceptual metric is good. It cannot show that hidden state variables, collision distances, rewards, or constraints are accurate, and an error appearing in the picture does not necessarily change the final

action. Testing must observe the variables actually consumed by the planner. It must also use consumer substitution, independent anchors, or action gates as controls to confirm how the error propagates.

This part therefore always records "prediction quality" and "control consequence" side by side. The former cannot automatically substitute for the latter, and the latter must also be supported by an explicit consumer and external receipts.

\newpage

Chapter 16: The Four Functional Boundaries of World Models

One system generates interactive video from a first frame and WASD input. Another predicts future states from a camera and candidate actions, and ranks them for a planner. A third feeds prediction uncertainty directly into a safety controller. All three may be called a "world model," but the assets that need protection lean in different directions. The first leans toward generation conditions and timing, the second toward imagined trajectories and action coupling, and the third toward constraints and real-time execution. If a safety team applies the same test on the basis of the name, it will misjudge offline video generation as robot control. At the same time, the system that truly enters feedback control may lack action-level guarantees.

This chapter takes the world model (WM) as the superordinate concept. It gives the working definition this survey adopts for the environment world model (EWM), and the operational definitions it adopts for the world action model (WAM) and the world control model (WCM). These four types can overlap. Each describes a set of verifiable inputs, states, outputs, and consumption relationships. They are not brand labels, and they are not a capability hierarchy from low to high. Classification follows the actual data flow and the consumers, and it determines the threat model, metrics, and controls.

Chapter Overview

This chapter uses six classes of variables—state, observation, action, future, goal, and control—to establish functional coordinates for world models. WM maintains state and predicts future variables. EWM takes external environment evolution as its main output. WAM lets future prediction participate in action generation or candidate ranking. WCM further connects prediction to feedback control, trajectory generation, or safety filtering. A single system can possess several of these functions at the same time. The key evidence is the actual data flow and the way downstream consumers use it.

The subsequent analysis separates latent state prediction, environment generation, prediction-action coupling, and feedback control item by item. It then marks out the safety assets introduced by reward, value, constraint, candidate trajectory, and execution feedback. When product documentation is incomplete or terminology is not yet stable, the basis for classification contracts step by step according to observable interfaces. Unknown states are then mapped explicitly into data permissions, action permissions, and operational limits. This approach keeps technical names, functional classification, and safety judgments separate.

Grounding the Principle: First Find Who Consumes the Future

The object of a world model system is a chain of "observation and action history—internal state—candidate future—downstream consumer." The inputs can be images, maps, text, proprioceptive state, historical actions, and goals. An encoder first writes them into a persistent state, and a dynamics or generative mechanism then predicts future observations, latent states, rewards, costs, or constraints. The output itself does not determine the category. When a human viewer, a data generator, a policy trainer, a planner, an action head, and a feedback controller consume the same future, they produce completely different safety surfaces. Protection covers more than the model weights. It also covers condition provenance, state freshness, goal version, candidate ranking, control period, and execution receipts.

In a normal example, an offline environment generator produces video from a first frame and camera actions for an on-duty operator to view, and its future does not automatically acquire action authority. In a boundary example, the same future is used by a planner to rank routes and also enters speed control in every control cycle. Consumer relationships for both WAM and WCM then appear in the system. Suppose public material shows video but never demonstrates an action consumer. Classification must then stop at WM/EWM or "unknown." Control capability cannot be filled in from a product name. All EWM, WAM, and WCM judgments in this chapter obey this working or operational boundary. Compare the output consumers of the four functions side by side in the figure. Focus on how the three columns "predicting the future," "participating in action selection," and "entering feedback control" change the classification, rather than reading the four rows as capability levels.

	WM、EWM、WAM 与 WCM 的功能边界		
	生成或保留 未来/状态	预测直接参与 动作选择	进入反馈控制
WM 世界模型	是	可选	可选
EWM 环境世界模型	是，强调外部环境	可选	可选
WAM 世界动作模型	是，须有消费证据	是	可选
WCM 世界控制模型	是，须有消费证据	是	是

A four-row, three-column matrix comparing whether the world model, environment world model, world action model, and world control model generate or retain future state, participate in action selection, and enter feedback control; the WAM and WCM cells note that consumption evidence is required.

The figure can support one sequence of interface determinations. It can show who reads the future, whether the read result changes actions, and whether actions enter sustained feedback. Two further points still need verification through probes on closed interfaces, logs, and versioned operational evidence. The first is whether a given product has already implemented the internal chain in the figure. The second is whether EWM, WAM, or WCM have become stable official categories.

16.1 The Boundaries of the Four Functions Are Determined Jointly by State and Consumers

World Model: Maintain State and Predict the Future

A world model has a minimum function. It maintains some kind of state according to historical observations and optional actions. It predicts at least one of future state, observation, reward, or constraint. Let the true but not necessarily fully observable environment state be $\mathbf{x}_t$. Let the model's internal state be $\mathbf{s}_t$, and the observation be $\mathbf{o}_t$. Let the action and external condition be $\mathbf{a}_t, \mathbf{z}_t$ respectively. Then one general form is

$$\begin{aligned}\mathbf{s}_t &= E_\theta(\mathbf{s}_{t-1}, \mathbf{o}_t, \mathbf{a}_{t-1}), \\ (\hat{\mathbf{s}}_{t+1}, \hat{\mathbf{o}}_{t+1}, \hat{\mathbf{r}}_{t+1}, \hat{\mathbf{c}}_{t+1}) &= F_\theta(\mathbf{s}_t, \mathbf{a}_t, \mathbf{z}_t).\end{aligned}$$

The encoder E_θ compresses history into the model state. The dynamics F_θ predict the next state, observation, reward $\hat{\mathbf{r}}$, or constraint $\hat{\mathbf{c}}$. Here $\mathbf{z}_t$ explicitly denotes external condition inputs such as maps, weather, and text. The symbol $\hat{\mathbf{c}}_{t+1}$ denotes the constraint quantity or constraint violation quantity predicted by the model. The two cannot reuse the same symbol merely because both are related to "conditions." Different models output different subsets of these, and safety analysis records them item by item according to the actual interface. Suppose a

system only produces a static description of an input image and has no state or future estimation. Its function is closer to a VLM, and it has not yet reached the minimum boundary of WM.

The functional scope of a world model can also extend beyond pixel-by-pixel reconstruction. MuZero learns rewards, values, and dynamics that are useful for planning, and full observation reconstruction is merely an optional task. PlaNet and Dreamer, by contrast, perform planning or policy learning in latent states [WM_hafner2019planet, @WM_hafner2020dreamer, @WM_schrittwieser2020muzero]. Image quality and control utility therefore need to be measured separately, and predictions of safety-relevant variables must also undergo independent testing.

This minimum functional boundary adds four classes of assets: state integrity, dynamics integrity, temporal consistency, and uncertainty calibration. A single observation error can be written into $\mathbf{s}_t$ and then rolled out across multiple steps. Dynamics can sustain short-horizon reconstruction yet drift over long horizons, and a model can also be overconfident in out-of-distribution states. The central question of world model safety shifts from "does the output look right" to "which future variables are treated as the basis for decisions, and by whom." Follow the solid lines in the figure from the observation and action history to the planner. Then look back at which consumer the latent state, dynamics, future rollout, and goal constraints each provide information to.

世界模型的最小内部链



先沿实线识别信息如何获得状态与能力，再判断哪一处需要独立控制。

Seven boxes run from observation and action history through encoding, latent state, dynamics, and future rollout to cost, constraints, and the planner.

This minimal internal chain explains why photorealistic imagery is not sufficient to infer reliable control. A planner may consume latent states, rewards, or constraints without consuming pixel reconstruction. The figure does not specify a concrete network architecture, and it does not prove that the modules are mutually independent. If an implementation shares an encoder or a goal head, common-cause failure still needs to be listed separately in the dependency graph.

Environment World Model: Take External Environment Evolution as the Main Output

The environment world model emphasizes the visual, geometric, or interactive evolution of the external environment. It may generate future frames, latent variables, or interactive scenes. Those outputs can come from images, text, maps, 3D boxes, camera actions, or other conditions. EWM is a working definition adopted by this survey, and its abbreviation and boundary are not yet as stable as WM. When classifying, state the decision rules publicly and mark the state of community consensus as still forming. The minimal input-output relation of EWM can be written as

$$\hat{o}_{t+1:t+H} = G_{\phi}(o_{\leq t}, z_{t:t+H}),$$

Here $\mathbf{z}$ can be text, camera motion, a map, an object box, or human interaction. The main safety assets of EWM include condition integrity, temporal consistency, and scene identity and privacy. They also include the credibility of physical conditions and downstream use. If it only generates video offline, the action risk comes from later misuse or from the data supply chain. It does not come from the model itself having already gained control authority.

The distinction between EWM and image/video generation depends on state and use. An ordinary text-to-video model also predicts frame sequences. Yet it does not necessarily maintain a world state that can be consumed for interaction or for decision-making. Classifying it as EWM requires observing whether it evolves the environment consistently according to history and conditions. It also requires observing whether the output serves the function of simulation, data generation, policy evaluation, or an interactive world. When the terminological boundary is fuzzy, the direct labels "action-conditioned video generator" and "interactive environment generator" are usually more precise than a forced one.

The condition channel is a safety asset in its own right. A driving EWM may consume high-definition maps, 3D boxes and ego-vehicle state. These inputs need not appear in the final image, yet they govern how the future evolves. PhysCond-WMA optimises map embeddings and 3D box conditions directly. This shows that pixel control covers only the observation channel. Physical conditions still require independent protection [WM_guo2026physcond]. EWM acceptance must therefore list each condition separately, with its provenance, coordinates, timestamp, normalization and permitted influence.

World Action Model: The Future Must Participate in Action

WAM does not require that "video and actions appear together in the output." It requires a verifiable coupling between future prediction and action generation, candidate evaluation, policy improvement, or planning. Let the candidate action set be $\{a^{(j)}\}$. The world model produces candidate futures, and the planner selects on that basis:

$$\hat{s}_{t+1:t+H}^{(j)} = F_{\theta}(s_t, a^{(j)}), \quad j^* = \arg \min_j C(\hat{s}^{(j)}, a^{(j)}, g).$$

Once the result of C truly determines action selection, the future becomes part of the action path. Evaluation must therefore retain the consumption evidence that runs from candidate scores to the selected action. A second class of joint WAM outputs actions directly from observation, goals and imagination:

$$a_{t:t+K-1} = D_{\psi}(o_t, g, \hat{s}_{t+1:t+H}).$$

The two structures have different attack surfaces. Candidate-based WAM must protect candidate generation, value/cost and ranking. Joint WAM must additionally protect the synchronization between imagination and the action head. JailWAM focuses on how dangerous goals propagate through future and action filtering. BadWAM demonstrates the possibility that the future looks plausible while the action deviates [WM_liu2026jailwam, @WM_li2026badwam]. Suppose a video model generates subsequent frames from WASD camera input while a human always supplies the navigation decisions. Then the future neither generates nor evaluates actions, and "action-conditioned" describes only a change in camera viewpoint. An autonomous policy requires an independent action generation or evaluation consumer. MoWorld's public material separates navigation decisions from visual generation. The existing evidence therefore supports a WM/EWM classification; WAM or WCM would still require evidence of action coupling and feedback control [WM_moxin2026moworld, @WM_moxin2026github].

WAM adds one further asset: "imagination-action consistency." Four cases require two-axis testing. The predicted future is correct while the action is wrong. The predicted future is wrong while the action happens to be corrected by external control. Prediction and action are both wrong. Or both are correct. Visual similarity can cover only part of the future. Action reachability, constraint violation and execution outcome still require independent metrics.

World Control Models: Prediction Entering Feedback Control

The world control model is the operational taxonomy this survey adopts. It requires world prediction to enter at least one of feedback control, trajectory or low-level command generation, and safety filtering. The term has not yet settled into a stable official category. Implementing model predictive control indicates only that a system may satisfy the corresponding functional conditions. Whatever name the authors adopt remains governed by the original literature. The taxonomy serves safety analysis, and the naming used by the research itself is kept as is. The minimal input-output relation of a WCM is

$$u_t = \kappa(s_t, \hat{s}_{t+1:t+H}, g, \mathcal{C}), \quad x_{t+1} = P(x_t, u_t, w_t),$$

where κ outputs trajectories or low-level control, $\mathcal{C}$ is the safety constraint, and $\mathbf{P}$ is the true environment dynamics. The true next observation then updates $\mathbf{s}_{t+1}$ and forms feedback. A WCM's safety assets include more than predictions and actions. They also include the control cycle, latency, reachability, actuator saturation, constraint coverage, fail-safe states and switching logic. WCM and WAM can overlap. WAM stresses how future prediction participates in actions. WCM stresses how prediction enters feedback control and constraint enforcement. An offline candidate ranker may be a WAM rather than a WCM. A latent model that supplies uncertainty and safety sets for robust MPC can satisfy both WM and WCM. A reactive VLA outputs actions directly but does not predict the future. That matches the input-output definition of a VLA, yet it falls short of the minimum WAM/WCM conditions here.

Guarantees in a control context rest on explicit premises. Prediction intervals, reachability and control barrier functions provide guarantees only when the assumptions about state estimation, dynamics error, latency and constraints hold. When a world model is overconfident in out-of-distribution states, formally feasible control may rest on an incorrect set. A WCM must define degradation after assumptions fail. The name of a mathematical controller is recorded together with its applicable assumptions and runtime probes.

The Four Functional Types Can Overlap

Four decision questions identify the four functional types. Every affirmative answer must point back to an actual data flow, an output consumer and the corresponding runtime evidence:

1. Does the system maintain state and predict future variables from history and optional actions? If so, it reaches the minimum functional boundary for a WM.
2. Do the main outputs describe how the external environment evolves visually, geometrically or in interaction? If so, the EWM label can be added.
3. Are future predictions consumed by action generation, evaluation, policy improvement, or planning? If so, the WAM label can be added.
4. Do future predictions enter continuous feedback control, trajectories, low-level commands, or safety filtering? If so, the WCM label can be added.

System function	WM	EWM	WAM	WCM
Offline action-conditioned video generation	yes / to be verified	often yes	no, unless the future participates in actions	no
Training a policy with latent dynamics	yes	not necessarily	yes, if imagination is used for policy improvement	not necessarily
Ranking candidate actions with video futures	yes	often yes	yes	depends on whether feedback is closed
Latent model + robust MPC	yes	not necessarily	yes	yes
Reactive VLA	not necessarily	no	no	no

The "yes / to be verified" entry in the table reminds reviewers that action-conditioned video does not necessarily maintain a reusable state explicitly. It also reminds them to retain unknown when the public interface is incomplete. System classification may change with deployment as well. The same EWM generates data offline during the research phase. Once it is connected to a policy evaluator at deployment, it adds WAM risk. If it then enters real-time control, it adds the latency and actuation requirements of WCM.

When public material is insufficient, classification should contract step by step along the strength of evidence. A product page can establish how the provider describes functionality, but judgments about internal state, action consumers or the control loop still require further verification. An architecture diagram in a paper can support module relationships, but formulas, code or interfaces are still needed to confirm that the key tensors are actually consumed. Runtime logs come closest to deployment reality, but they must be bound to version and configuration. When the action path has not yet been found, the conclusion remains "unconfirmed WAM".

Official material from Happy Oyster supports its positioning as interactive world creation, and it can serve as a functional clue for WM/EWM. The available public evidence reaches only the functional description level.

WAM/WCM coupling and product safety still require independent evidence [WM_happyoyster2026, @WM_alibaba2026annual]. Such cases show that entity identity, functional classification and safety conclusions are three different tasks. Provider identity, technical boundaries and attack-defense results each rest on different evidence.

Classification must also retain version timing. One version of a service may generate only offline video; the next may add automatic candidate ranking. The name stays the same, but the consumers have already changed. The safety ledger should record model version, API, caller, action permissions and evidence date. Then classification can be updated with the actual data flow instead of overwriting earlier facts. Interfaces that cannot be confirmed should be marked explicitly as unknown. Unknown items are themselves the basis for testing and for tightening permissions.

16.2 Understanding Functional Boundary Changes Through Historical Context

Dyna connects learning, planning and reaction. World Models shows training controllers in a learned latent world. PlaNet and Dreamer combine pixel observations, latent dynamics and imagination-based planning [WM_sutton1991dyna, @WM_ha2018worldmodels, @WM_hafner2019planet, @WM_hafner2020dreamer]. In this lineage the core safety question is how model bias enters the policy through long-horizon rollouts, not whether the generated frames themselves are photorealistic. MuZero further shows that a world model can learn only the quantities useful for planning. TD-MPC2 and DreamerV3 extend latent world models to a wider range of tasks and continuous control [WM_schrittwieser2020muzero, @WM_hansen2024tdmpc2, @WM_hafner2025dreamerv3]. As video world models, robotic WAM and safety control combine, the attack surface grows to include conditioning embeddings, synthetic data, imagination-based ranking, action heads, tool execution and recovery.

Capabilities, data, architecture, evaluation environments and degree of openness all change together over time. The historical conclusion that the existing heterogeneous evidence supports is this: once new consumers appear, older prediction errors gain longer propagation paths. The causal ordering between year or model family and attack effects still requires comparable samples and control variables.

16.3 Safety Assets Change with the Consumer

Offline EWMs mainly protect training provenance, conditioning integrity, privacy, temporal consistency, content provenance and processing history, and downstream use. If the synthetic environments they produce enter policy training, risks may appear weeks later. The generative model version, conditioning, data batch and downstream policy version therefore need to be connected.

A planning WM must additionally protect candidate trajectories, rewards, values, costs and ranking. A predictor may reconstruct observations accurately and still rank safety-relevant low-probability tail trajectories in the wrong position. Internal average error cannot cover ranking flips near the decision boundary. A WAM must protect the bidirectional consistency of imagination and action. Two questions define it: whether an action can reach the claimed future, and whether the future is indeed generated conditioned on that action. Action token decoding, inverse dynamics and candidate selection are all independent artifacts. A WCM further adds the control cycle, hardware saturation, stop channels, fallback controllers and human takeover.

Agentic world models also bring web pages, tool descriptions, return values and transaction results into the state. A world model can rehearse tool consequences. It cannot simultaneously serve as the sole predictor, the sole risk judge and the sole execution authorizer. External policy and transaction boundaries must still hold the final capability [WM_imgrund2026falseprophets, @WM_lin2026dreamguard].

16.4 The Four Functional Blocks Inside a World Model

Implementation details vary widely. Safety analysis can start by looking for four functional blocks: representation, dynamics, objective and consumer. Representation compresses historical observations into state. Dynamics predicts state, observations, rewards or constraints. The objective assigns value to the future. The consumer turns

the future into data, ranking, actions or control. Each functional block may be a single network, several networks or a set of rules.

The representation block does not encode vision alone. It may fuse language, maps, proprioceptive state, tool returns and past actions. Safety asks three questions: is provenance identity retained, can state be reset, and are unknowns expressed? If all conditions are first concatenated into untyped tokens, environment content may change the objective. If history has no horizon, a single error persists for a long time.

The dynamics block can predict pixels, latent states, rewards or values. Pixel reconstruction is convenient for human inspection, yet it may ignore decision-relevant low-dimensional variables. Latent models are computationally efficient, yet harder to interpret with external evidence. MuZero shows that a model useful for planning need not reconstruct all observations [WM_schrittwieser2020muzero]. Safety metrics must therefore match the actual prediction target.

The objective block includes rewards, values, costs, constraints and human preferences. A world model may predict accurately and still choose dangerous trajectories because the objective has been tampered with. Objective provenance, version and permissions should be separated from dynamics. Safety constraints must not exist only as text that the model can freely generate. The consumer determines risk escalation. When the future is only for browsing, the main consequences are misleading and downstream misuse. Entering data generation creates a delayed supply chain. Entering candidate ranking changes plans. Entering feedback control is subject to real-time and physical constraints. Safety testing should reason backward from the consumer's maximum permissions rather than guess forward from generation quality.

The four blocks may also share parameters. Sharing improves efficiency but creates common-cause failures. If representation, dynamics, value and action head all depend on the same encoder, a single conditioning perturbation may change state, future and judgment at the same time. Independent external anchors and the action gate should sit outside the shared boundary.

16.5 Identifiability of States, Observations, Rewards, and Constraints

A world model learns internal state from limited observations, and that internal representation is usually not unique. Two latent states can generate the same short-term frames yet produce different predictions for long-horizon actions and risk. A given latent dimension can be assigned specific semantics such as "distance" or "danger" only once interventions or external labels support it. The identifiability problem affects internal detection. A large latent cosine distance before and after an attack may only be a change of representational coordinates. A small distance may also flip candidates near the decision boundary. A more reliable approach connects latent changes to observable consequences: external geometric error, reward ranking, constraint violations, action differences and the true next observation.

Rewards and constraints may also be impossible to infer uniquely from behavior. The same action can arise from different objective trade-offs. A single safe behavior supports that single outcome. Internal safety constraints, by contrast, must be tested through objective conflicts and boundary states. Testing observes what the system chooses when efficiency and safety are in tension, and compares that choice with ordinary task performance.

Counterfactual actions help test dynamics. Hold the trusted state fixed, feed in several candidate actions, and check whether the predicted future moves in a sensible direction and magnitude. Then run a small number of actions in the real world, or in a highly trusted simulation, and check the predicted ranking. When the future hardly responds to actions, the model is closer to a video generator that correlates only weakly with the conditioning. In that case the existing evidence does not yet reach the WAM functional boundary. Uncertainty is a prediction target likewise. A low variance given by a model is an internal signal. True error still needs to be calibrated across tasks, states, and time horizons. Calibration may fail out of distribution or under attack. Unknown states should therefore trigger a contraction of capabilities, and external facts should be confirmed by independent anchors.

16.6 Functional Matrix of Six Common Architectures

An architecture name can provide only a retrieval entry point. The functional determination must still land on the actual future consumer. The table below places six common structures into the same set of columns. It retains

representative citations while separating "conditions satisfied" from "still to be verified".

Architecture	Who consumes the future	Functional determination in this survey	Newly added safety assets and evidence boundary
Latent dynamics plus policy	Training-time policy learner	Satisfies WM. Satisfies WAM when imagination is used for policy improvement. A deployment-time reactive policy does not necessarily satisfy WCM	Protect the latent state, the imagination data, and the policy derivation chain. World Models, PlaNet, and Dreamer belong to this lineage [WM_ha2018worldmodels, @WM_hafner2019planet, @WM_hafner2020dreamer]
Planning-type value model	Searcher or planner	Satisfies WM/WAM. Satisfies WCM once rolling search closes the loop with real feedback and directly controls	The focus is value, dynamics, candidate coverage, and ranking. Visual imagery output is not required
Driving or scene video EWM	Simulator, data generator, or a policy consumer to be verified	Usually satisfies WM/EWM. Whether it satisfies WAM depends on whether the downstream policy consumes the future	UniSim and GAIA-1 support conditioned video and scene generation relations. Action coupling cannot be supplied from the paper title alone [WM_yang2023unisim, @WM_hu2023gaia1]
Interactive environment generator	Human operator, or an automatic consumer to be verified	WM/EWM is usually determined first. WAM is added only when an automatic action consumer exists	Genie-type work embodies interactive generation. Human keypresses are a condition, not the same as the system autonomously choosing actions [WM_bruce2024genie]
Joint future-action model	Action head or shared-representation consumer	WAM is determined only when ablation, interfaces, or code prove that the future enters the action head	Side-by-side output does not equal coupling. BadWAM shows the boundary at which the future is plausible while the action branch deviates [WM_liu2026badwam]
World-model safety controller	MPC, safety filter, or execution checker	Usually satisfies WM/WAM/WCM	Requires calibration, real-time, reachability, constraint, and safety-mode evidence. Existing conclusions remain bound to specific protocols [WM_nath2026sls2, @WM_liu2026checkvla]

The matrix outputs "which functional conditions are satisfied, what the evidence is, and what remains unknown". It does not output a single exclusive overall label. One system can satisfy WAM during the training stage while running only an already-trained reactive policy during the deployment stage. Different stages must be encoded separately. The training-time consumer cannot automatically be written into the run graph.

16.7 Verification Fixtures for Family Determination

Determination fixtures use interface interventions to verify documentation. For WM, hold the history fixed while changing the current action, and observe whether the future responds. Then clear the history and observe whether the prediction depends only on the current frame. If there is no state or future dependence, the minimum functional boundary needs to contract. For EWM, change the map, the camera, the object boxes, or the text conditions. Then observe whether the external environment evolution changes with the input conditions. Then test condition conflicts and temporal disorder. If the output changes only in single-frame appearance, environment dynamics are not necessarily formed.

For WAM, first replace the future representation and observe whether action generation, scoring, or policy updates change. Then hold the future fixed while swapping candidate actions, and check the action conditioning of the prediction. A performance change after disabling the future consumer can provide a structural clue. No change on a single task, however, does not prove that coupling is absent. For WCM, provide the model with stale state, infeasible trajectories, and delayed predictions, and confirm that the controller rejects or degrades them. Then let the safety check time out and observe whether it switches to a safety policy. Only when the prediction actually changes the trajectory, the command, or the safety mode does it support a control-function determination.

The fixtures must also test bypasses. After the main planner rejects, does the system execute along a cached trajectory, a backup tool, or a low-level debugging interface? When the world model is unavailable, does the high-performance policy reuse the old state indefinitely? A bypass may leave the safety boundary still failing under a correct classification. The test denominator is interface probes and configurations, not real-world tasks. One

interface probe passing shows that the expected data flow exists under the corresponding version and inputs. When a closed API cannot access intermediate state, retain the conditional classification. Apply a conservative capability gate at the execution layer.

16.8 Metrics Vary with Family and Consumer

The basic metrics for WM include state prediction error, reward/value error, calibration, and long-horizon rollout deviation. An EWM with visual outputs adds video quality, temporal consistency, and condition adherence. An EWM with latent or geometric outputs instead adds state, trajectory, occupancy, or geometric error. It also adds the deviation of these quantities relative to external anchors. Visual quality is not a required metric for it. WAM adds candidate ranking, imagination-action consistency, task success, and action risk. WCM further adds constraint satisfaction, control latency, stopping distance, recovery, and takeover. FID, FVD, PSNR, SSIM, or human perception answer the generation distribution and appearance. Candidate ranking, action reachability, and control safety are each measured by consumer-relevant tasks and external anchors. A latent distance must also record at the same time the representation used and the calibration method.

Metric denominators are distinguished by object. They may be future frames, video clips, state transitions, candidate trajectories, replanning cycles, action chunks, complete tasks, control events, and real-world impacts. One task producing a large number of frames does not equal a large number of independent tasks. Long-horizon models need to preserve the dependency structure by trajectory or scene. Returns with the same name may also differ. Game scores, robot task rewards, safety costs, and driving planning costs are not the same estimation object. An attack success rate can stand in for the target future, a ranking flip, a dangerous action, task failure, or a constraint event. A report should write the complete definition next to the number, rather than using only an abbreviation.

System metrics also include latency, computation, energy consumption, false alarms, missed detections, and tracking integrity. A generative WAM may be much slower than a reactive VLA. Comparisons must be made on the same hardware and at the same control frequency. Millisecond numbers from different device sources are suitable for side-by-side presentation. Strict ranking requires unified devices and protocols [WM_zhang2026wamrobustness].

16.9 How Product and Project Evidence Supports Functional Classification

Official pages support product identity, provider, and publicly disclosed functions. Papers support method structure and protocols. Code and model cards support executable interfaces. Run receipts support the behavior of a specific version. The four kinds of evidence cover different questions. A company's use of "interactive world model" indicates its product positioning. Internal action coupling and safety continue to be determined by interface and runtime evidence.

Happy Oyster's official functional statements support interactive world creation and WM/EWM clues. There is not enough public material, however, to confirm WAM/WCM or specific attack-defense conclusions [WM_happyoyster2026, @WM_alibaba2026annual]. Where the public interface has not yet been verified, the WAM/WCM status is recorded as unknown. The test scope is then set according to the confirmed WM/EWM capabilities. MoWorld's public material is conditioned on an initial frame, text, and WASD camera operations. It separates navigation decisions from visual generation, so it can be discussed as WM/EWM. Its code, weights, and license have not yet closed. The current evidence status is therefore a public repository and interface clues. A complete run still awaits the closure of dependencies and licenses [WM_moxin2026moworld, @WM_moxin2026github].

Product versions change functional boundaries. An offline EWM becomes a WAM consumer once automatic policy evaluation is added, and adding real-time control then adds WCM. The ledger should record version, interface, consumer, permissions, and evidence date. Every capability and conclusion is then bound to the corresponding version. Independent evaluations and incident evidence still need to be listed separately. Official performance numbers support official protocols, while third-party results are provided by independent tests. Security research

proves that a path exists under specific settings. Whether a product is affected is confirmed by version and runtime evidence. Real-world incidents require evidence of object, time, consequence, and attribution. Functional classification describes only the data flow, and security certification still needs to separately complete control and risk validation.

16.10 Family Selection and Engineering Controls

An offline EWM prioritizes control of data provenance, conditions, privacy, timing, and downstream use. If generated data enters policy training, add data lineage, filtering, and revocation. A simulation-type EWM additionally adds real-data anchors, simulation bias, and policy gating. A planning-type WM prioritizes control of state, dynamics, reward, candidate coverage, ranking stability, and uncertainty. The world model cannot become the sole evaluator. High-impact candidates require external rules and worst-case checks.

WAM must verify future integrity and action consistency on two axes, and protect action tokens, inverse dynamics, and decoder artifacts. A correct future does not equal a correct action. An action that is temporarily correct does not prove that the state is trustworthy either. WCM requires a hard real-time budget, robust control, actuator limits, a safety controller, failure modes, and physical stopping. The assumptions behind a prediction interval or a formal tool can fail. When they do, the system should shrink its capability rather than continue to use the nominal future. The engineering checklist answers a fixed set of questions. What is predicted, and who consumes it? What are the maximum privileges? Where do state provenance and goal provenance come from? What is the unit of action, and at what control frequency does it run? What is the independent anchor? What is the failure mode, and what is the recovery path? As long as any one of these is unknown, map the unknown to a test or a permission restriction. The closer a category is to execution, the lower the tolerance for uncertain items.

16.11 Sources of Error: Observation Noise, Model Bias, and Objective Bias

World-model error comes from at least three places. Observation noise makes the model start from incomplete or wrong inputs. Model bias makes the learned dynamics differ from the real environment. Objective bias makes the system choose the wrong future even when it predicts correctly. The three may produce the same task failure, yet they require different controls. Observation noise can be reduced by better sensing, redundancy, time synchronization, and state estimation. Model bias requires real-data anchors, model ensembles, conservative intervals, and short-horizon replanning. Objective bias requires external constraints, versioned rewards, and human authorization. Handing all failures to input sanitization would miss the latter two.

Uncertainty can be roughly divided into environmental randomness and lack of knowledge. Environmental randomness comes from unpredictable events and will not completely disappear even with more data. Lack of knowledge comes from states the model has not seen, or from structural error. It can be reduced by new data and model improvements. A single variance given by an actual model often mixes the two. Safety control should focus on total risk and assumptions rather than forcibly decomposing them precisely from a single output.

Long-horizon rollouts make errors correlated. A stepwise independent Gaussian assumption may underestimate persistent bias. Under the same map error, the model drifts in the same direction across multiple time steps. Verification should measure single-step mean squared error, but also systematic bias, the worst trajectory, candidate ranking, and constraint boundaries. Objective bias is especially easy to mask with generation quality. The future video can be realistic, and the physical trajectory can be reachable. Yet if the reward places efficiency above safety, the planner will still choose a dangerous plan. Therefore the generation acceptance of WM/EWM and the objective acceptance of WAM/WCM must be kept separate.

16.12 Item-by-Item Determination of Boundary Cases

A digital twin does not automatically equal a world model. A system that is merely a manually maintained deterministic simulation can also predict the future. Yet it does not necessarily count as a learned WM. Security controls still apply the language of state, dynamics, and consumers, but the model attack surface differs from that

of a neural network. Recording "learning components" and "rule components" is more useful than arguing over names.

A trajectory predictor may predict the trajectories of other vehicles or pedestrians for human analysis. Only after the prediction enters the ego vehicle's action generation or ranking does the system have a WAM consumer. If it enters real-time control, then consider a WCM. Between prediction error and ego-vehicle control failure there is still a propagation chain of planning, authorization, and execution. A video generator with action conditioning also does not automatically equal an EWM/WAM. The action may only control the camera or the style, and the future has no stable state. One needs to test conditional response, temporal consistency, and consumers. Conversely, latent dynamics that do not output video can satisfy WM/WAM/WCM. The condition is that the prediction be verifiable for planning and control.

Future prediction draws the boundary between VLA and WAM. A VLA can output actions directly from historical observations without explicitly producing a verifiable future. A WAM must have a future that participates in the action. Internally, a VLA may form an implicit representation. Classifying a system as WAM, though, requires interface, training-objective or intervention evidence that an action consumer really uses the future. Using a world model inside a security checker does not make the policy under inspection a WAM itself. The system level contains WAM-style consumers. Component classification remains a separate question. An attack can target the policy, the checking world model, or state that the two share. Logs must be able to locate which one failed. Large-scale "world foundation models" are a capability or training positioning. The four categories of function in this chapter still need to be determined on the basis of interfaces. Parameter scale, pretraining data and multi-task capability change the attack surface. Whether the output enters action or control is determined by the consumer path. Classification looks at the interface first, and discusses scale afterwards.

16.13 A Complete Recording Template for Family Classification

Part one of the classification record is the system boundary. That covers version, operating stage, deployment purpose, external services and maximum privileges. Part two is the variable table: observations, historical state, action conditions, goals, rewards, constraints, future outputs and real feedback. For each variable, record its source, unit, time and writable subjects. Part three is the consumer graph. Who reads the future? A human, a data generator, a value head, a planner, an action head, a checker or a controller. Can that consumer execute? Does a bypass exist? Part four is the evidence table. It lists the documents, formulas, code, interface probes and run receipts that support each relation.

Part five is the family determination. For WM, EWM, WAM and WCM respectively, write "satisfied", "not satisfied" or "unknown", and give one falsifiable reason. "Unknown" means the current evidence has not yet confirmed whether an action head, a planner or a controller consumes the future. When an unknown relation lies on the dependency path of a high-impact action, shut down only the high-impact automatic execution that depends on it. Restrict the related outputs to read-only recommendations or an offline sandbox. Verified read-only, diagnostic or conservative controls that do not depend on that relation may continue. Run supplementary tests on candidate future replacement, action sensitivity and consumer logs.

Part six is the threat mapping. For each writable variable it gives the attacker, the budget and the first-broken interface, and states whether an error can persist. It also names the highest consumer and the consequence, along with the earliest control and the final limitation. Part seven is validation: normal tasks, boundary conditions, adversarial conditions, latency, recovery and version triggers. The template also requires a counterexample. If the system is judged to be a WAM, state what kind of probe would overturn "the future participates in the action". If it is judged to be a WCM, state what kind of logs would show that the prediction is merely a bypass display. Counterexamples let the classification update with the evidence instead of becoming an unchallengeable label.

A team can use two independent reviewers to code separately on the basis of interface evidence. They then deliberate on disagreements item by item. Those disagreements usually expose document ambiguity, stage confusion or invisible consumers. They are more valuable than a simple majority vote. The final determination is bound to the evidence date and version. The output of the complete record should give the security, control and product teams the same data-flow diagram. The product may continue to use its market name, while security

testing still runs along the functional boundary. Terminology differences will not block control, as long as consumers and privileges are accurately recorded.

16.14 Four Minimal Determination Experiments

Experiment 1 tests offline video generation. The system receives a first frame, text and camera keystrokes, and outputs an 8 s video. Hold the first frame and the text fixed, then change only the keystrokes. If the viewpoint or the environment evolution responds stably, this can support an action-conditioned EWM. Then check whether an automatic action generation or scoring consumer exists. If there is none, judge it to be WM/EWM, not WAM/WCM. MoWorld-type public interfaces can be understood with this line of thinking [WM_moxin2026moworld].

Experiment 2 tests imagination training. A latent model generates future states from history. The policy uses imagination data only during training, and at deployment it outputs actions directly from observations. The training system satisfies WM/WAM, because the future participates in policy improvement. The deployment run graph may contain only a reactive policy. The threat record must cover the supply chain from imagination data to the policy checkpoint. One should not look in deployment logs for future calls that do not exist.

Experiment 3 tests candidate planning. Given the same state, the planner generates ten action sequences, and the world model predicts returns and selects the highest one. After the candidate future or the scores are replaced, the choice changes accordingly, which supports a WAM. If this process is repeated at every control cycle with new observations and short actions are issued directly, it also supports a WCM. If the result serves only as a human recommendation, it stops at planning support.

Experiment 4 tests runtime checking. A reactive VLA produces actions. An independent action-conditioned world model previews the future, and the checker can reject or shorten the action. The VLA component itself is not necessarily a WAM. The checking subsystem satisfies WM/WAM, and the whole execution system has a WCM-style runtime assurance architecture. Testing must attack the VLA, the checking model and the shared state separately. It must not assign the system label to every component.

All four experiments must run negative controls. Zero out the future tensor. If the action does not change, that may mean the future is not consumed. It may also mean that the current task is insensitive to the future. Repeat this across states and boundary tasks. If the checker logs a risk but the action does not change, this shows that it is an observer rather than a control consumer. If the future changes the action but there is no execution privilege, then WAM is satisfied. Real-world control is not automatically satisfied.

Each experiment outputs an evidence level. Documentation claims are structural clues. Formulas and code paths support design relations. Interface interventions support causal relations. Deployment receipts support the behavior of a specific version. Where a closed system can reach only a lower level, annotate the classification with "judged by public interfaces". Map internal unknowns to external limits. The determination must also record benign utility. A task drop after the future is disabled shows that the consumer is useful. Attack and fault tests then evaluate future security. A task slowdown after the checker is added shows that control intervention has a cost. Family classification answers how the system works. Security evaluation answers how it fails under attack and fault. The two are reported separately.

The experimental results can form a stage matrix. Its columns are training, offline evaluation, task planning, runtime checking and low-level control. Its rows are state, future, action, constraints and feedback. Each cell records the actual component, the inputs and outputs, and the privileges. The matrix can reveal that the same "model" takes on different roles at different stages. It can also reveal that a certain stage has no independent control at all.

The stage matrix also helps identify delayed supply chains. An EWM generates data during the training stage, and weeks later the policy produces actions during the operating stage. The source model is already offline, yet the risk still persists through the policy. Logs need to link data batches, policy versions and deployment events. They must not save only the current world model calls. For multi-agent systems, one must also record whether the future is consumed by other agents. One agent generates a scenario, another agent plans with it as fact, and a third agent authorizes tools. A single component may have no action output, yet the system-level path still

satisfies WAM-style coupling. Message provenance, principals and capability delegation thereby become part of the system boundary.

Finally, check recoverability. If the world model fails, can the system switch to real observations, conservative rules, or control that does not depend on that model? If it cannot, the classification record should mark the world model as a single-point critical consumer. It should then require stricter release, monitoring and rollback gates. After classification is complete, the control team should also review the action path. The data team should review the state and training lineage. The operations team should review versions and rollback. When the three parties read the same interface differently, resolve the disagreement with probes and logs first. Then decide the permitted privileges. Terminology consensus handles naming. Data-flow evidence supports the determination. Operational responsibility implements the specific controls.

If a supplier can provide only a functional description and no interface evidence, the procurement terms should require version notification, failure boundaries, trace fields and security fallback. When the information is still incomplete, the system places the component in a low-privilege, replaceable position. Local controls then limit the maximum consequence. The less visible the dependency, the more conservative external authorization should be. This principle runs through all four categories of function: WM, EWM, WAM and WCM. Runtime probes continuously verify the actual consumers, execution privileges, latency and post-failure actions.

16.15 Worked Example: Classifying a Port Digital Twin

The port system has three modules. Module A generates future 30 min scenarios from historical camera feeds, weather and vessel positions, for the duty officer to browse. Module B predicts berth congestion and equipment status for each scheduling plan. It then automatically selects the plan with the lowest predicted cost. Module C reads actual equipment feedback every second. It uses a prediction model and security constraints to control the speed of unmanned transport vehicles.

Module A satisfies the functional conditions for WM and EWM. If the output is only for human reference, it does not have WAM/WCM coupling. The main tests are condition integrity, temporal consistency, privacy, and how erroneous information affects human decisions. Module B's future directly participates in plan ranking, satisfying WM/WAM. One needs to test candidate coverage, ranking stability, target cost and imagination-plan consistency. Module C's prediction enters feedback control, satisfying WM/WAM/WCM. One must also verify latency, actuator limits, the emergency controller and the real-feedback anchor.

If the three modules share the same map encoder, map poisoning may create a cross-module common-cause failure. Module A's future imagery, Module B's congestion ranking and Module C's control state may simultaneously and self-consistently deviate from reality. Controls should establish signatures and versions at the map artifact layer. They should anchor with independent positioning and measured traffic at the state layer. At Module C's execution layer, they should configure speed and zone hard constraints that do not depend on the same encoder. The report records by the highest observed layer. An erroneous future in Module A is a prediction failure. A security controller in Module C blocking an action is a consequence limitation. Upstream state failures in A/B are still listed separately. Classification provides propagation coordinates. The level of risk is evaluated separately by consequence, privilege and reachability.

16.16 Six classes of classification misreading caused by names, conditions, and imagery

Misjudgment	Missing decisive evidence	Normative conclusion
Every video generator is a world model	Whether it maintains state, predicts the future from history and conditions, and which system function its output serves	Name an ordinary video generator by its actual interface first. Judge it WM only after it meets the minimum state and future conditions
Action-conditioned generation automatically equals WAM	Whether the future is actually consumed by action generation, evaluation, policy improvement or planning	When the condition controls only the camera or the imagery, it is not judged WAM
A VLA necessarily contains a world model	Whether a verifiable future prediction exists, and who consumes it	A reactive VLA can output actions directly from observations and does not automatically satisfy WM/WAM
WCM is a stable official model category	A unified official or community definition does not exist within this survey's current evidence scope	Always label WCM as this survey's operational classification. Retain the original names used in the literature and products
The more realistic the future imagery, the safer the control	Whether the reward, ranking, action head, constraints and execution feedback are correct	Image quality covers only the generation dimension. Consumer metrics measure control safety separately
The four functional classes form a high-low hierarchy	Whether the four classes can overlap and whether they serve different consumers	They are an overlapping set of functions, not a capability ladder from low to high

16.17 Bringing it into a real system

A port scheduling system can place the four functional classes into a single data flow diagram. The offline module generates future video from the initial frame, text and camera key presses. A human still makes the navigation decision. It satisfies the environment-prediction features of WM/EWM. The candidate ranking module uses the predicted future for plan comparison, and enters the scope of WAM. The real-time control module then feeds the predictions and constraints into feedback control, forming a WCM. Every classification records the input, internal state, output, consumer and execution frequency.

The consumer is what separates reactive VLAs, candidate-style WAMs and robust-MPC-style WCMs. A reactive VLA produces an action representation directly from the current observation. A candidate-style WAM generates or evaluates multiple futures and influences action selection. A robust-MPC-style WCM computes the control input within a receding horizon, using dynamics, uncertainty and constraints. The safety assets newly added thereby include, in order, the observation-to-action mapping, future candidates and ranking, and real-time state estimation and constraint feasibility.

A world model may also predict only reward and value without reconstructing pixels. It maintains state useful for planning and predicts future variables that affect decisions. On those conditions it still meets the minimum functional definition of a WM. In that case, visual quality metrics cannot reflect reward error, candidate flipping or constraint mismatch. Testing should turn to value calibration, ranking margin and real execution feedback.

If the three modules share a map encoder, a poisoned map, a coordinate mismatch or an out-of-distribution scene becomes a common-cause risk. Independent localization sources, actuator state readback and geometric safety controllers can supply separate roots of trust. Deployment records should separate the controls that only check offline generation from those that can block consequences before candidate ranking or real-time execution.

Summary: the model category is decided by the consumer

WM maintains state and predicts future variables. EWM emphasizes how the external environment evolves. WAM brings the future into action generation or selection, and WCM brings predictions into feedback control and safety constraints. Those four functional classes can overlap. What settles the judgment is a verifiable data flow, not a name. Once the classification changes, the safety assets extend with the consumer. That reach grows from conditions and timing to candidate ranking, action consistency, real-time constraints and recovery. If the available

material cannot prove that actions consume the future, keep "unclassified" or "insufficient evidence". Do not let product names fill in functional evidence. The functional boundary settles the question of "what the system actually is". Where the malicious change sits in the imagination chain still needs further tracking. Chapter 17 will trace that accumulation layer by layer — supply chain, observation, state, dynamics, goals, planning and feedback — and show how world model attacks cross the consumer.

\newpage

Chapter 17 State, Dynamics, and Goal Attacks: How the Imagination Chain Is Hijacked

The world model of one unmanned transport vehicle predicted the road ahead and the motion of other vehicles accurately, yet its cost function had dropped the "no entry" constraint. A second vehicle kept every constraint, but its map conditions were slightly altered, so the future trajectory still looked coherent in the wrong lane. A third system produced a future with no obvious anomaly, yet its action head chose a control that did not match that future. All three may receive decent video quality scores. Their points of failure lie in the goal, the physical conditions and the imagination-action coupling.

World model attacks are distinctive in more than the complexity of their perturbations. The error can also be consumed repeatedly along state recursion, dynamics rollout, candidate ranking and feedback control. This chapter takes that chain in order: supply chain—observation—state—dynamics—goal—planning/action—execution/feedback. Along the way it explains how attackers control the imagination chain, and how far each piece of evidence can support.

Chapter overview

This chapter reduces world model attacks to a traceable imagination chain. The analysis starts from the first-broken functional interface. Item by item, it separates the optimizable variables inside training artifacts, runtime observations, latent state, dynamics, rewards and constraints, candidate ranking, and the final action. Pixels, patches and backdoors are therefore only external carriers. Where an attack actually sits depends on which variable it changes first, and on which consumer keeps amplifying the error.

Long-horizon rollout accumulates local deviations. Candidate ranking amplifies tiny changes at the decision boundary. A shared model bias may shift prediction, uncertainty and action judgment at the same time. This chapter records separately the privileges that frozen-weight attacks, data poisoning, query search and feedback manipulation require. It then uses four dimensions — interface, time, consumer and consequence — to establish a falsifiable, isolatable test protocol.

Background principle: an attack changes a variable in the imagination chain

This chapter is about an imagination chain. Its components are training artifacts, runtime observations, persistent state, dynamics, objective functions, candidate ranking, actions and feedback. Its inputs are not pixels alone. They also include maps, object boxes, historical actions, rewards, constraints, tool returns and versioned configuration. Encoder and state update compress these inputs into an internal state that can affect the future across rounds. Dynamics then roll out candidate futures, and the planner selects actions according to goals and constraints. Output consumers extend from offline viewers to policy trainers, planners, action heads and feedback controllers. Whether the same perturbation becomes a safety event therefore depends on the joint action of the first-broken variable, consumer privilege and consequence limits.

In the normal case the map, the state and the hard constraints come from different trusted sources. Candidate futures enter the action gate only after an external cost review. In the boundary case the attacker modifies a single training trajectory and still lets a wrong goal propagate over the long term, through a derived world model, synthetic data and policy checkpoints. The billboard that appears at deployment is only the backdoor trigger. The

first-broken interface is still in the supply chain. In another boundary case the future imagery remains plausible while the action head deviates, so video quality has not observed the variable that actually failed. The optimizable variable, the frozen object, the budget, the consumer and the highest consequence layer all belong in one attack record. Trace the attack influence from left to right along the seven variables in the figure. Distinguish the position the red dashed line first enters from the position where task failure or control deviation finally appears.

想象链中的七类可攻击变量



先沿实线识别信息如何获得状态与能力，再判断哪一处需要独立控制。

Seven boxes running from data artifacts, observation, state, dynamics, goal constraints, and candidate ranking to action feedback, marking multiple attack positions.

What the figure supports is a causal localization order, not the compression of all attacks into a single grade. An attack that appears at the observation layer does not mean the first-broken interface must be in observation. A shifted internal state or future does not mean the action gate has already let it through. Real-world consequences still require closure through action, execution and independent environment acknowledgment. Tensor experiments not connected to a real consumer remain at the mechanism-level run evidence layer.

17.1 Attack paths of the imagination chain

The world model's state update, future rollout and action selection can be written as the minimal interfaces below. They separate the first-broken variable from the final task symptom. This representation assumes no particular model architecture. It serves only to mark which mapping an attack variable first acts in:

$$s_t = E_{\theta}(s_{t-1}, o_t, a_{t-1}),$$

$$\hat{\tau}^{(j)} = F_{\theta}(s_t, a_{t:t+H}^{(j)}), \quad j^* = \arg \min_j C_{\psi}(\hat{\tau}^{(j)}, g, \mathcal{C}).$$

An attacker can change any of these: the observation o_t , the historical state input, the training parameters θ , the candidate actions, the goal g , the constraints $\mathcal{C}$, the cost model ψ . The same final task failure may come from completely different first-broken interfaces. A security record should name the attack variable v , the frozen part, the budget and the objective function. It must also name the consumer and the highest consequence layer. World model attacks have at least three amplifiers. State recursion writes a single input into multiple future time steps. Imagination rollout lets small dynamics deviations accumulate with the horizon H . The planner compresses continuous predictions into discrete choices. Near the decision boundary, tiny changes in candidate scores may flip the final action. Average prediction error does not necessarily capture these three amplifications.

Malicious attacks must still be separated from natural failures. Occlusion, sensor noise, distribution shift and model error can expose the same interfaces. None of them carries an adversary's goals and budget. Natural robustness methods can become the basis of runtime assurance. Conclusions about attack defense also require adaptive testing against detectors and controllers.

17.2 Seven Classes of Attack Surface: From Artifact Writes to Feedback Manipulation

The seven classes of attack surface follow the variable that is changed directly, not paper titles or final symptoms. The supply chain decides which artifacts and training relationships enter the system. Observation, state and dynamics decide what kind of world the system believes in. Goals and candidate ranking decide which future is worth selecting. Actions, execution and feedback decide whether an error acquires real-world capability. The subsections below move along that same chain without a break. Each one preserves the attacker's permissions and the consumer boundary at every point.

Supply Chain: Fixing the Wrong Future into the Model

Supply chain attacks change data, transition targets, encoders, dynamics, rewards, configurations or checkpoints. Backdoor attacks are one such class: they bind a specific trigger condition to the attacker's target behavior. They must not be lumped together with general quality degradation that lacks such a condition binding. Normal validation can therefore survive. The failure itself appears only once the trigger condition occurs, in long-horizon imagination or in the downstream policy. A world model may also generate training data or serve as a policy evaluator. The failure can then propagate across projects and over time.

SWAAP first seeks targets with lower returns but dynamics still close to the clean model. It then uses gradient matching with a stealth constraint to change the transition target of a small number of fine-tuning state transitions (transition target, namely the label that supervises which next state or future the current state and action should reach). The attacker needs to influence the data or targets of model updates. The first-broken interface is therefore the supply chain rather than deployment-time observation [WM_hu2026swaap]. Closeness in clean dynamics indicates low detectability at the selected distance, whereas safety-relevant trajectory ranking requires an independent comparison.

Targeting World Models places an attacked world model into a robot learning pipeline. It emphasizes that a contaminated model can influence the downstream policy through synthetic data or a training stage [WM_rathbun2026targeting]. Here the acceptance unit spans the world model version, the synthetic data batch, the policy training configuration and the deployment result. The final policy weights are only one link in it. Propagation evidence stops where the chain breaks, and the highest observed layer is explicitly marked.

BadDreamer's trigger sequence is "context preserved, future erased". It pushes the future representation toward waypoints that do not avoid [WM_shuai2026baddreamer]. A temporal trigger does not require an obvious anomaly in every frame. A frame-by-frame data scan may be completely normal. The malicious binding between the condition and the future becomes visible only when the context and future labels are unfolded along the trajectory.

Beyond hashes and signatures, supply chain control must also verify behavior. Compare predictions, ranking and the downstream policy under normal, triggered, near-trigger and temporally reordered conditions. Record which modules were frozen and which were trainable during training. Establish provenance and revocation paths for generated data. Signatures prove provenance, while behavior tests examine actual performance. Together the two classes of evidence close artifact admission.

Observation and Physical Conditions: Attacks Do Not Happen Only in Pixels

Runtime observation attacks usually freeze the model weights and optimize only the current image, conditioning embedding or context. Hallucination-Driven Policy Failure studies pixel, latent-space and temporal augmentation perturbations for Dreamer-style continuous control. It shows that a single small input change can accumulate across frames through the state buffer [WM_zhang2026hallucination]. The conclusion is bound to the authors' continuous control model, normalization and perturbation protocol. Driving or robot video models need thresholds re-established according to their own sensing range, consumers and perturbation semantics.

PhysCond-WMA extends the attack variables to high-definition map embeddings and three-dimensional box conditions. It maintains a temporally consistent semantic shift through reverse denoising guidance [WM_guo2026physcond]. A human watching the final video may not notice maps or three-dimensional boxes, yet

those conditions define roads, objects and motion. The input asset inventory must therefore cover all conditioning tensors, not only the visible image.

BadWorld perturbs the context images at test time. It uses a velocity objective that does not depend on true future labels. Its search is trajectory-adaptive and oriented toward unknown subsequent control sequences [WM_shen2026badworld]. The model weights remain frozen, and the attack does not require training-poisoning permission. Its feasibility depends on white-box/query capability, context writability and knowledge of the control sequence. It therefore counts as a runtime context attack, and it needs fewer permissions than a training backdoor.

Observation control must preserve type, coordinates and time. One scene snapshot must hold the map version, object box coordinates, camera extrinsics, control sequence and text condition. Internal conditioning consistency may still drift together. True localization, an independent map or rules are then needed as external anchors. A sanitization algorithm that reduces only the visual change while damaging the physical conditions may also make control worse.

Latent State: How a Local Input Becomes a Long-Term Belief

The latent state is not a simple cache of visible frames. It compresses history, actions and uncertainty, and it becomes the starting point for future prediction. If a perturbed observation pushes the state into a wrong region, later normal observations may not be enough for immediate recovery. The model explains the new input with its own erroneous prediction, and that forms a self-consistent drift. Long-horizon amplification can be understood through the recursion of state error. Let the difference between the clean and perturbed states be δ_t . Local linearization gives

$$\delta_{t+1} \approx \mathbf{A}_t \delta_t + \mathbf{B}_t \epsilon_t,$$

where ϵ_t is the current input perturbation and $\mathbf{A}_t$ is the local sensitivity of the state dynamics. If $\mathbf{A}_t$ keeps amplifying along a critical direction at several time steps, a single small perturbation accumulates over time. If the true feedback is strong enough and the state update contracts, the error may decay. This expression explains the accumulation mechanism. The real-world deviation still has to be determined by measured matrix responses, feedback and execution trajectories.

A small internal residual does not mean the state is correct. When the encoder, predictor and reconstructor share parameters, they can together interpret a wrong state as normal. Valid measurement must set the internal residual beside the errors of external geometry, physics, rules or another sensor. If an anchor comes from pseudo-labels generated by the same model, it does not have independence.

Tests of state attacks must cover resets. A soft reset may clear only the interface while the cache survives, and the map and history may be restored from contaminated logs. The protocol must state whether each trial continues the same trajectory, restores a checkpoint, or starts from an independent initial state. Otherwise cross-episode state leakage will mistake samples for independent units.

Dynamics and Imagination: Attacking What the System "Thinks Will Happen"

A dynamics attack targets the future itself, or the representation of the future that matters to decisions. Trusted Imagination applies a bounded perturbation to the current observation. The future latent variable then either deviates from the clean imagination or approaches a target imagination. The study uses a downstream high-quality action selector as an evaluation probe rather than letting the attacker directly overwrite that selector [WM_chen2026trustedimagination]. The "oracle-level" wording in the original paper describes the quality of the evaluation probe here. The attacker's permissions remain a bounded observation perturbation.

In that study, the LaDi-WM model-predictive controller consumes the imagination and shows a clear attack-specific failure on one small-sample task. The reactive policy is approximately preserved. That difference between consumers supports the idea that "risk propagation forms only when the future is used for decisions." The evidence, though, comes from a single task, 20 simulations per condition, and an abnormal random baseline. It is suitable for establishing existence. It is not suitable for estimating an industry-average attack rate.

Temporal consistency describes continuity between frames. Physical correctness also involves object identity, speed, reachability, and safety-relevant low-dimensional states. Evaluation places visual quality, latent state, external anchors, candidate ranking, and downstream tasks side by side. FID/FVD are retained as the visual dimension, and closed-loop metrics are responsible for the action and consequence dimensions. Dynamics integrity also depends on action conditioning. Suppose the model claims that a certain future is produced by action α . The checker should verify that this future is reachable in kinematics and dynamics, and should check it when the true next observation arrives. A video future without action conditioning can support only environment evolution. It is not enough to issue permission for a specific action.

Goals, Rewards, Value, and Constraints: Correct Prediction Can Still Choose Wrong

The planner selects a future according to rewards, value, costs, and constraints. An attacker can keep predictions correct while changing "what is worth selecting." Reward poisoning research shows that the objective channel at interaction or training time constitutes an independent security boundary [WM_zhang2020rewardpoisoning]. In a world model system, the objective may also come from elsewhere. It may be a user task, a risk discriminator, a rule configuration, or a learned value head.

JailWAM feeds dangerous language candidates into a WAM. It then uses action trajectory mapping, risk screening, and dual-path validation to select instructions that can manifest as oscillation, boundary violation, or collision risk in the closed loop [WM_liu2026jailwam]. Here the risk discriminator performs candidate screening inside the attack pipeline, and the whole pipeline serves attack search. If the discriminator is separately deployed as an entry filter, its evidence scope is bound to a specific model, simulation, and task.

Constraint tampering and constraint omission must be separated. In the former, the attacker changes the safety distance, restricted zone, or budget. In the latter, the system never encoded a certain danger. Attack experiments must make clear whether the attacker can write the configuration, reward, or tool description. Suppose the permission allows changing only pixels, yet results from directly changing constraints are used to claim remote input risk. Then the threat model has been overstepped.

Objective control should keep safety constraints outside the model, in a versioned policy. That policy is bound to the applicable subject, state, and approval. A learned value can optimize efficiency. On its own it cannot override hard restricted zones, actuator limits, and statutory requirements. High-impact objective changes require two-stage confirmation. The runtime log must preserve the connection among the objective, the constraint version, and the action decision.

Candidate Ranking, Actions, and Control: Attacks Near the Decision Boundary

The planner ultimately compresses continuous predictions into a single candidate. Tail-aware TRAP does not degrade all candidates. Instead it changes the ranking of the most critical high-scoring tail trajectories in the clean plan, which makes long-horizon dangerous options easier to select [WM_duan2026trap]. The average return or average prediction error may be preserved, yet the decision boundary has already flipped. WMAttack composes attack families, perturbation budgets, optimization steps, restarts, and resource allocation into a limited-budget search. It warm-starts that search with historical task configurations [WM_guo2026wmattack]. Its value lies in reducing the overestimation of robustness caused by handcrafted weak attacks. Its conclusions remain limited by the search space, objective, and budget. It is not a certificate against all unknown real-world attacks.

BadWAM attacks the action branch of a joint WAM and penalizes imagination drift at the same time. The future then looks plausible while the action deviates significantly [WM_li2026badwam]. The attack optimizes a small visual perturbation through queries at runtime. It is not a training-time action head backdoor. This case requires evaluation to form two axes: whether the future is correct, and whether the action is consistent with the future and the constraints. Only when both pass does it enter the execution gate. The first-broken interface must not be confused with the first internal variable that is directly changed. A small runtime visual perturbation first crosses the U1 input trust boundary. If the attack directly changes candidate scores, the first directly changed variable lies in the planning sub-interface at U3. If the future and ranking are normal while the action head is directionally shifted, the first directly changed variable lies in the action sub-interface at U3. If the action head is poisoned during training, the first-broken interface is the supply chain, U0. All of these may eventually manifest as task failure, but the location for prevention differs.

Execution, Feedback, Tools, and Privacy

In agent toolchains, world models let web pages, tool descriptions, return values, and side effects enter the state. False Prophets emphasizes that an agent world model cannot be the sole predictor, the sole judge, and the sole authorizer all at once [WM_imgrund2026falseprophets]. Tool least privilege, parameter types, external rules, side-effect preview, reversible tests, and post-completion verification constitute different roots of trust.

Feedback attacks include forging, delaying, reordering, and replaying the true next observation. When a world model updates its state from erroneous feedback, the recoverer may read failure as success and keep executing. Feedback packets must carry the subject, sequence number, time, state version, and action ID. A tool "returning 200" proves that the request was accepted. A physical or business objective still requires confirmation by an environment or business receipt.

Privacy is another consequence chain. World models may memorize videos, maps, indoor environments, and user trajectories. Users can perform membership inference or extraction by repeatedly controlling initial frames, viewpoints, text, and randomness. Imagined Memorisation provides preliminary mechanistic clues, and the scope of the existing evidence is bounded by its data and protocol. The absolute leakage rate of commercial systems requires corresponding deployment samples and denominators [WM_ng2026imagined]. Privacy, intellectual property, and physical integrity should be reported separately.

17.3 Attack Surface Matrix: Variables, Permissions, Time, and Consumers

World model attacks are described with a four-dimensional matrix: direct variables, attack permissions, temporal forms, and consumers. Direct variables include data/artifacts, observations/conditions, latent state, dynamics, objectives/constraints, candidate ranking, actions, and feedback. Permissions include submitting data, changing training, reading gradients, querying outputs, writing configuration, physical entry, and tool identity.

Temporal forms include single-step, cross-frame, delayed trigger, long-horizon rollout, replanning, and cross-version derivation. Consumers include offline viewing, data generation, policy evaluation, planning, action heads, controllers, and tool authorization. The same variable read by different consumers yields different highest consequences. When map conditions are perturbed, for example, an offline EWM may generate erroneous videos. A planning-type WM may change candidate ranking. A WCM may change real-time control. The entry point is the same, but the consumers and permissions differ. A test fixes one matrix cell. Offline video results fill the generation column. Action and execution tests fill the control column.

The matrix also records frozen objects. Runtime observation attacks usually freeze weights, training data, and the security policy. Supply chain attacks change training data or the model. Objective attacks may freeze dynamics and change only the reward. Feedback attacks may freeze the entire model while changing receipts. The frozen list makes causal attribution verifiable. Finally, budgets and endpoints are recorded. The recorded quantities include pixel norm, map embedding distance, poisoning ratio, queries, candidate count, trigger duration, tool authorization scope (authorization scope, that is, the tools, objects, parameters, and time limits that a token is permitted to invoke), and physical time. These are not the same kind of budget. Latent distance, return, task failure, and real-world events are not the same kind of endpoint either. The matrix is used to select tests, not to compress heterogeneous numbers into a single aggregate score.

17.4 Causal Paths of Supply Chain Attacks

The world model supply chain is longer than an ordinary checkpoint. Data first train a predictor. That predictor may generate imagined data or evaluate a policy. The policy is then deployed to a real system. An attacker changes one state-transition training objective, and the impact can span multiple derived artifacts. Causal paths require recording the version and input of each transition. SWAAP's mechanism first seeks objectives that have low return and are close to clean dynamics. It then uses stealthy constraints to change a small number of fine-tuning objectives [WM_hu2026swaap]. It shows that comparing only average dynamics distance may miss decision-relevant directions. Verification should jointly test normal states, triggered states, candidate ranking, downstream return, and anomalous residuals, and should not perform only weight hashing.

Targeting World Models places a poisoned world model in a robot learning pipeline [WM_rathbun2026targeting]. If data from the poisoned model feeds policy training, the risk persists after the source model is taken offline. The revocation process must trace synthetic data, policy checkpoints, evaluation results, and deployed versions. BadDreamer's trigger sequence binds context and future to a non-avoidance path [WM_shuai2026baddreamer]. Frame-by-frame random sampling destroys temporal relations and may also miss the attack. Data review should preserve context-action-future relations per trajectory. It should also inspect behaviors whose conditions activate only within specific orders and time windows.

The denominator of a supply chain attack can be contributed transitions, trajectories, training batches, model versions, or downstream tasks. A poisoning ratio is interpretable only when the numerator and the data unit are explicit. 1% of frames, 1% of trajectories, and 1% of tasks are not the same permission. Training randomness, data deduplication, and sampling weights also change actual exposure. Causal confirmation can use derived ablations. Hold the policy training configuration fixed, and consume data generated by a clean world model, then by a candidate world model. Then fix the data and replace the world model. Finally fix the policy and compare deployment triggers. If the intermediate stages are not run, one may write only latent propagation. One must not claim that the full pipeline has been compromised.

17.5 Persistence Testing of Condition and State Attacks

EWM and WAM have far more conditions than pixels. Maps, object boxes, camera trajectories, text, action history, weather, and tool returns can all be written into the state. Every condition must define field name, data type, unit, value range, coordinate system, observation time, and provenance subject. Attack tests must likewise define budgets in the corresponding space. PhysCond-WMA directly optimizes map embedding and 3D box conditions [WM_guo2026physcond]. A small embedding distance does not necessarily mean a small physical change. Perturbed conditions should be decoded back into lanes, object positions, and sizes, to check whether scene constraints are satisfied. A human eye that looks only at the final video may fail to notice that the condition has already gone out of bounds.

BadWorld changes the context image at test time. It then performs trajectory-adaptive search over the subsequent control sequence [WM_shen2026badworld]. Frozen weights do not mean a frozen state. A perturbed context enters the latent state. Tests need to keep running after the attack stops, measuring how long the state and actions take to recover. Persistence tests use three input types: impulse, step, and intermittent. An impulse is perturbed at only one instant, and it estimates state memory. A step remains perturbed, and it measures the steady state and the safety gate. Intermittent input examines whether triggering and recovery produce hysteresis. Each condition is run with a clean case, a random-amplitude control, and a targeted attack.

Internal state error and external anchors are recorded simultaneously. Suppose internal reconstruction and prediction are both normal, yet the real map, geometry, or rule error keeps rising. That indicates common drift. If the external anchor also depends on the same condition source, it is not independent. State reset falls into soft reset, model state reconstruction, simulation snapshot restore, and real environment reset. Each level has different cost and coverage. Reports should state the level actually used. Multiple replannings derived from the same contaminated state are counted under the same task denominator.

17.6 Continuous Verification of Rewards, Constraints, and Candidate Ranking

World model safety analysis easily focuses on prediction while ignoring "how to assign value to the future." To explicitly expose where value, cost, and constraints act, let a candidate future be τ_j . The planner selects

$$j^* = \arg \max_j [V_\psi(\tau_j, g) - \lambda C_\omega(\tau_j, \mathcal{C})].$$

An attack may change any of these: the objective g , the value V , the safety cost $\mathcal{C}$, the weight λ , or the candidate set. Even when the future itself is entirely correct, the choice can still be dangerous. Reward poisoning may change value during training. Configuration attacks may change weights at deployment. Language jailbreaking may change the task objective, and tail-ranking attacks target decision-critical candidates

[WM_zhang2020rewardpoisoning, @WM_duan2026trap]. In that order, the four paths fall under the supply chain, configuration, task objective, and candidate ranking interfaces. Their permissions are recorded per path. Safety constraints are further divided into hard constraints and soft costs. Hard constraints are rejected by an external policy or controller and should not be offset by model weights. Soft costs permit trade-offs between efficiency and risk, and must be bound to versions and approvals. If a forbidden region is encoded only as a learnable penalty, the optimizer may accept a violation under high reward.

Candidate coverage and candidate ranking remain continuous responsibilities. If the planner never generates a safe trajectory, even an accurate scorer can only choose among dangerous candidates. If a safe trajectory does exist, the test must also ask whether a small perturbation near the decision boundary pushes it out of the top k (top- k). Tests record candidate diversity, the fraction of feasible candidates, the rank of the safest candidate, the final choice and the external real cost, all at the same time. Ranking attacks need to capture exactly one phenomenon. A very small change in the average score accompanies a flip of the optimal candidate.

Objective, constraint and ranking logs finally connect to user authorization and action outcomes. Task text, structured objectives, reward versions, safety rules, candidate sets and final actions are each retained. That record lets a team tell apart "no safe candidate was generated", "the scorer chose wrongly" and "the external hard constraint was not enforced". Explanations generated by the model cannot override external hard constraints. High-impact weight changes require independent approval and a short validity period. Changes beyond authority are handled according to the corresponding threat model.

17.7 The Role and Boundaries of Automated Attack Search

Pick one attack, one step size and one budget by hand, and a weak configuration is easily mistaken for model robustness. WMAttack composes attack families, budgets, optimization steps, restarts and resource allocation into a limited-budget search, with historical task configurations as warm starts [WM_guo2026wmattack]. The search increases evaluation strength, yet it still covers only the defined search space. What automated search finds depends on its optimization objective. Maximizing return degradation may ignore safety constraints. Maximizing action instability may find harmless oscillation, and maximizing task failure may favor tasks that were originally difficult. Safety evaluation should use multiple objectives: attack endpoints, constraint events, query cost, stealthiness and impact on benign inputs. It should also retain the raw result of each.

Budgets have to match real capabilities. A simulation can reset thousands of times, while a real attacker may have only a few observations. White-box access can read gradients, while a hosted API can only query actions. The budget of a supply chain attack is data and training privilege. Each class of automated method uses the corresponding query, reset, data and training budget. Searchers overfit to benchmarks as well. Warm-starting from historical tasks can improve efficiency, yet it may leak model and scenario features into the test. Unseen tasks, checkpoints and attack families should be held out. Search failures and timeouts should be reported. A miss supports conclusions only under the current search budget.

Defense-aware search incorporates input sanitization, detection, filtering, action gates and recovery into the environment. Its goal is the highest penetrated layer, not differences internal to the model. If a safety gate blocks successfully, the search may continue to try other entry points. It must not, however, cross the stop condition on real devices just to obtain Y_4 .

17.8 Verification Fixtures for Compound Attack Chains

A compound chain can be built along the sequence "artifact—condition—state—future—ranking—action—feedback". The fixture first establishes a baseline with clean artifacts and clean conditions. It then activates only one segment at a time, and combines them only in the end. That separates the individual contribution of each segment from the interactions between them. It also avoids attributing every failure to the last payload. The first group verifies the supply chain. It replaces candidate checkpoints or generated data batches, while keeping the runtime inputs unchanged. The second group verifies conditions. It freezes the checkpoint and perturbs only the map, object boxes or context. The third group caches the clean and the perturbed states, then feeds them

crosswise into the same dynamics. The fourth group swaps the future or the candidate scores. The fifth group fixes the plan, and compares the action head against the safety gate.

Every step has an external control. State is compared with the true geometry. The future is compared with the action condition and the true next observation. The ranking is compared with an external cost, the action with the reachable set and permissions, and the feedback with independent sensing or business state. Internal consistency is recorded as an auxiliary signal. External anchors are responsible for confirming the deviation from reality. Compound effects may be nonlinear. If two attacks share the same budget, or one attack changes the states reachable by the other, simple addition loses the interaction. Reports give single-segment, two-segment and full-chain results separately, and state the independent unit explicitly. When samples are insufficient, case trajectories describe the observed interaction.

The highest evidence level a fixture can reach depends on the consumers it actually connects to. Tensor swapping counts as a mechanism-level run. Sandboxed tools can form simulated execution. Only after state, action and environmental feedback are closed does it form a closed-loop simulation. End-to-end reproduction further requires that real or verified-equivalent weights, data, identities, configurations, actuators and predefined metrics are all closed. Controlled physical actions must still be recorded separately from field incident evidence. Failure receipts matter just as much. Missing dependencies, a model that was not loaded, an empty candidate set, a simulation crash, a gating rejection and a successful recovery are encoded separately. Only runs that actually reach the predefined interface enter the attack denominator. Other outcomes enter the availability and control statistics.

17.9 Case Comparison and the Engineering Checklist

Hallucination-Driven shows that pixel or latent perturbations can accumulate through continuous control states. PhysCond-WMA shows that the map and 3D boxes are independent condition assets. Trusted Imagination shows that consumer differences determine the consequences of a future attack. TRAP shows that candidate ranking can be changed in a targeted way. BadWAM shows that the future can be plausible while the action is wrong [WM_zhang2026hallucination, @WM_guo2026physcond, @WM_chen2026trustedimagination, @WM_duan2026trap, @WM_li2026badwam]. Together, the five cases show that video quality is only one item of security evaluation. The models, permissions, budgets, tasks and denominators they cover all differ, so the original attack rates are kept separately. The engineering check should answer item by item:

- Can all training artifacts and generated data be traced to their downstream consumers?
- What conditions exist beyond pixels, and what are their units, coordinates and provenance?
- How long does state contamination persist after the attack stops, and which reset is effective?
- How are dynamics predictions compared with the true next observation, geometry or rules?
- Who writes the rewards, values and constraints, and can hard constraints be overridden by the optimization objective?
- Does the candidate set contain a safe option, and is the ranking stable under small perturbations?
- Are imagination and action consistent in both directions, and is the action gate outside the model?
- Is the feedback fresh, ordered and bound to the action, and can the attacker query without limit?
- Is the highest endpoint prediction, planning, simulated execution, a simulation event, or a real-world impact?
- Are benign utility, latency, recovery and evidence gaps reported in the same table?

The checklist yields a minimal attack path plus two controls with different roots of trust. If an unknown item lies near a high-impact consumer, shrink permissions first and add tests after that. States without evidence are explicitly recorded as unknown.

17.10 Counterfactuals and Causal Attribution

World model attacks often change several intermediate quantities at once. To confirm a particular path, counterfactuals can be constructed from the same trusted snapshot. The clean condition and the perturbed condition generate the state, future, ranking and action separately. One intermediate quantity is then cross-substituted, and the test observes whether the downstream stages recover. Cross-substitution is equivalent to

asking "if only the state were restored while the other attack conditions were held unchanged, what would happen".

The supply chain path can fix the deployment inputs and load the clean and candidate checkpoints separately. The condition path can fix the checkpoint and replace only the map or the context. The objective path can fix the future and replace only the rewards and constraints. The action path can fix the plan and replace only the action head. Only by changing one item at a time can the direct effect be localized.

Intermediate-quantity substitution may produce out-of-distribution combinations. A clean state with a perturbed future, for example, is not a pairing the model normally produces. The combination must therefore be verified as readable first. Only then should multiple interventions and ablations that support one another be used. When the internal state cannot be accessed directly, a policy that does not consume the future, or an external scorer, provides a weaker but falsifiable consumer control.

Causal results are expressed along the chain. The attack variable reaches the entry, the mediator changes in the predetermined direction, and the consumer reads the mediator. The action or task changes, and the safety gate and the environment determine the consequence. When the final action does not change, the highest evidence level stops at the latent state. When the action changes but the gating rejects it, both facts are retained: the action-layer hit and the consequence limit. When a real device is difficult to reset precisely, the paired design is stratified by scenario, task and initial state. The order of the clean and perturbed runs is randomized, and residual differences such as temperature, wear and human activity are reported. The more the mechanism evidence depends on an uncontrollable environment, the more conservative the causal language should be.

17.11 Post-Attack State and Artifact Disposition

Once a world model path is discovered, disposition restricts the consumers first, then preserves lineage. Offline EWM pauses publication to the data pool. Planning-type WM is downgraded to read-only advice, and WAM/WCM switch to conservative policies and independent control. This capability contraction acts on data flow and execution rights. Deleting model files belongs to the artifact disposition that follows. After the system stabilizes, the affected data, checkpoints, conditions, latent states, candidates, objectives, actions, feedback and environment versions are preserved. Access to raw video, maps and trajectories is restricted, so that the evidence can show which consumer actually used which version.

The contamination radius is then determined by artifact, software state, tool transaction and physical state. That is because training data may derive multiple models and a world model may generate policy data. A perturbed state may be written into a cache, and a dangerous action may even change the real environment. The recovery path differs by object. When an artifact is compromised, switch to a trusted version and trace the derivatives. When the state is compromised, rebuild from external anchors. When the objective is compromised, restore the approval rules and re-authorize. When the feedback is compromised, rebuild the action—result chain. When the physical environment has changed, perform a low-energy retraction or a manual inspection. Every recovery operation must itself pass a safety gate.

Post-recovery regression uses the same protocol to rerun trigger, near-trigger, normal tasks and defense-aware attacks. It confirms that the state does not recur after the attack stops, and that downstream policies no longer carry the influence. It also confirms that the action gate and the recovery path still meet latency. Items that cannot be closed keep the capability contraction in place, and an owner is recorded. The disposition report preserves both the upstream root cause and the downstream limiting control. The entry can be vulnerable while the action gate does in fact block the consequence. Writing only "no incident" omits the root cause, and writing only "attack succeeded" erases the actual role of the defense line.

17.12 World Model Red-Team Protocol

The red-team protocol selects the family and the consumer first. Offline EWM focuses on conditions, timing, privacy and data poisoning. Planning WM adds state, objectives and ranking, WAM adds imagination—action, and WCM adds real-time constraints, feedback and recovery. Test endpoints escalate with the consumer, and do not jump directly from video metrics to real-world consequences. Each unit fixes the system version, state

initialization, attacker permissions, budget, attack objective, control, denominator and stopping condition. White-box work records readable modules and gradients, while black-box work records responses and queries. Supply chain work records data and training permissions, and physical conditions record fabrication and placement. Unknown permissions are not filled in from common sense.

The test order runs from static lineage, constructed tensors, offline replay, reduced-scale simulation and full simulation to hardware-in-the-loop and isolated real hardware. When a lower layer has not reached the mechanism, a higher-risk environment is not used to produce a success. Each layer keeps four kinds of controls: clean, random perturbation, objective-irrelevant payload, and adaptive attack. Metrics are layered by state, future, ranking, action, execution and recovery. The state denominator is an independent snapshot. The ranking denominator is a valid replanning cycle, the task denominator is a complete episode, and the tool denominator is a unique transaction. The real-world impact denominator is an event or object. Adjacent frames and shared checkpoints retain their dependency relationships.

A compound attack adds one segment at a time, and tests single segments before it tests combinations. If the total budget is allocated across segments, the allocation rule is recorded. If an earlier segment changes the state accessible to a later segment, the interaction is reported rather than adding the single-segment effects together. When samples and variance are insufficient, case trajectories are retained. The red-team output must answer a fixed set of items. These are the first-broken interface, the highest penetrated layer, the blocking control, benign utility and latency, the recovery state, and the evidence gaps. Checking only materials and configurations is recorded as static inspection. Actually running a constructed fixture is recorded as a mechanism-level run, and closing state, action and environmental feedback is recorded as a closed-loop simulation. When real weights, data, a complete pipeline or predefined metrics are missing, end-to-end reproduction remains incomplete.

17.13 Error Propagation and Control Sensitivity

Summing the errors of the individual modules often misjudges closed-loop risk. Let the mappings from observation to state, state to future, future to score and score to action be $\mathbf{E}$, $\mathbf{F}$, $\mathbf{R}$ and $\mathbf{A}$ respectively. In regions that are locally differentiable, and where the action does not switch, the first-order influence of input error on the action can be written as

$$\Delta a \approx J_A J_R J_F J_E \Delta o.$$

This is a diagnostic formula for locating the amplifying stage. It is not a safety proof. A real planner runs top- k filtering, thresholds, quantization, clipping and discrete mode switching. A very small continuous error may therefore produce a discontinuous action once it crosses a decision boundary. The reverse also holds: a large internal distance may leave the final candidate unchanged. Engineering has to measure both "distance" and "margin". Distance describes how the state, the future or the action differs before and after the attack. Margin describes how far the current candidate still sits from a ranking flip, a constraint violation or a control mode switch. Let the score difference between the best safe candidate and the second-best dangerous candidate be m_t . As m_t approaches zero, the action chunk should be shortened, observations added, or an external safety policy switched in, even when video quality is essentially unchanged.

The sensitivity fixture runs in four steps. Step one generates small perturbations that are locally feasible for the same state. Objectives or weights marked as frozen must not change. Step two records the state, the future, the ranking, the action and the safety gate margin separately. Step three runs a finer adaptive search near the decision boundary and keeps an objective-irrelevant perturbation control. Step four runs only in a reversible environment. There it tests whether a choice flip produces simulated execution or a simulation event.

For each cycle, the propagation ledger must preserve the same set of fields. Those fields are the input version, the latent state fingerprint, the candidate set and the raw scores. They also include the constraint margin, the selected action, the gating result and the environment receipt. Only such a record lets the ledger answer "where the error went", "which consumer amplified it" and "which control stopped the consequence". A final success or failure on its own conflates four situations. The model was unaffected. The action did not change. The safety gate blocked. The environment did not manifest the effect.

One further distinction is needed, between "sensitive" and "easy to break". A large local gradient indicates only that some direction will change the proxy quantity quickly. Whether the attacker can reach that direction within the input-legality, query, time, physical and stealth budgets is a separate question. The report therefore separates sensitivity curves from the cost—effectiveness curves of attacks that are actually feasible. White-box results cover conditions in which internal gradients are visible. Black-box results cover conditions in which only the final action is visible. Likewise, the number of simulation queries and the on-site query capability are each bound to their environment.

For discrete switching, the record can include the first-flip budget and the number of cycles sustained after the flip. It can also hold the gating intervention lag and the number of new observations required for recovery. For continuous control, the report covers the action error and the minimum margin to the safety envelope boundary. It also gives the duration of limit violations and the post-clipping task utility. A hybrid system that moves from continuous regulation into emergency-stop mode must retain both classes of metric. Judging only by the maximum action deviation may mistake safe emergency braking for a dangerous attack. Judging only by task failure may miss a trajectory that crosses a constraint and then luckily completes the task.

Finally, every sensitivity conclusion must be bound to an operating point. The decision margin shifts with payload, speed, control frequency, the number of candidates and the planning horizon. Tests should cover rated conditions, critical conditions and legal extreme conditions. The termination conditions should be predefined. Critical conditions serve to observe small-perturbation amplification. Rated conditions describe everyday operation. Legal extreme conditions supplement boundary risk.

17.14 Worked Example: A Contaminated Campus-Scheduling World Model

The campus uses EWM to generate the forklift flow field. It uses WAM to select three candidate routes for the unmanned transport vehicles, and WCM then controls speed at 10 Hz. The attacker can contribute a small number of trajectories to the simulation data but cannot reach the deployment weights or the on-site network. The objective is to make the system prefer the route through the loading zone when a specific billboard appears. The first-broken interface is the supply chain. The attack variable is the contributed trajectories and their future/cost labels, and the deployment billboard is only the trigger. The deployment freezes model weights, and the attacker gets no query optimization at runtime. Testing should compare normal, triggered, near-trigger and different-viewpoint conditions. It should record EWM flow predictions, WAM candidate rankings, action proposals, WCM safety projections and actual simulation events.

Suppose the billboard triggers and the flow prediction stays plausible, yet the candidate ranking puts the loading-zone route at the top. The failure has then propagated to the objective or the ranking. Suppose instead that WCM's independent exclusion-zone constraint projects the action. Here the highest consequence stops at the dangerous plan or action proposal. The team fixes the data lineage and keeps the WCM constraint in place at the same time. It treats the upstream root cause and the downstream consequence limit separately.

Isolated testing prohibits any connection to the real fleet. Use signed simulation snapshots, fix random seeds and record independent trial units. The attack budget covers the number of trajectories that can be contributed, plus the label permissions. It is not a pixel norm. Then add adaptive attacks that know the candidate checker. These verify whether the detector merely memorized the billboard's appearance. Recovery testing comes next. It revokes the contaminated data and retrain or replaces the model. It then confirms that state, caches and downstream policies do not continue to carry the influence.

17.15 Six Misreadings in Which Symptoms Cannot Substitute for the First-Broken Variable

Misjudgment	Confounded layer	Judgment to be retained
Any degraded future is a dynamics attack	Training artifacts, observations, state, objectives and ranking can all change the future	The variable the attacker directly changes, together with the first-broken interface, sets the main category
Frozen model weights mean the whole system is frozen	Observations, history, conditioning, candidates, rewards and feedback keep changing	The freeze checklist must cover all runtime state and all external configuration
A low internal prediction residual means no real drift	A shared encoder and predictor can drift together, self-consistently	Check geometry, rules, position or the true next observation against an independent external anchor
A change in video quality can stand for control risk	Image, state, ranking, action and execution sit in different consequence layers	BadWAM-style failures can occur while the image stays plausible; futures and actions must be checked on two axes
Automated attack search is a robustness certificate	Search covers only a predetermined family, budget, objective and feedback	Not finding a penetration supports only a limited conclusion under the current search conditions
The backdoor trigger shows up in the input, so the first-broken interface is the input	The trigger location is not the same object as the location where the malicious binding forms	If the binding forms during training, the first-broken interface is still the supply chain

17.16 Bringing This into Real Systems

An attack record for a campus mobility system first separates the variable, the frozen components, the first-broken interface and the consumers. SWAAP writes a small fine-tuning state-transition training objective from the supply chain. It thereby poisons a low-reward objective binding into the candidate world model. BadWorld also keeps weights frozen at runtime. It changes the latent state through contextual image perturbations and trajectory-adaptive search aimed at unknown subsequent control sequences. BadWAM likewise keeps weights frozen. It uses runtime queries to optimize small visual perturbations, shifting the action branch while suppressing imagination drift. The three occupy different authorities: training-objective write authority, runtime context write access and query capability. Each class of result retains its own attack budget, statistical unit and highest-consequence layer.

State verification needs an anchor outside the model. The internal reconstruction residual may stay small while the encoder and the dynamics drift together. The system must therefore also compare independent localization, true odometry, map landmarks or external measurements. For candidate trajectories, it records the average score, the best-versus-second-best margin, the first-flip budget and the action difference after the flip. A small average change can also accompany a complete flip of the discrete choice.

Causal localization can replace mediating variables in sequence along "condition—state—future—ranking—action". Trusted-state substitution tests whether entry-point contamination has been written into the internal state. Trusted-future substitution tests dynamics propagation. Trusted-ranking or trusted-action substitution localizes the consumer. Each substitution records the out-of-distribution combination risk. It also uses a benign control to confirm that the result comes from the replaced path rather than from interface incompatibility.

The complete evidence chain starts with the supply chain summary and the input and state probes. It passes through candidate ranking and simulated actions, then reaches the independent action gate and the recovery receipt. Continuous action error, ranking margin, first-flip budget and recovery period together describe the attack consequence. A single metric covers only one face of the chain. Test stopping conditions are bound to the actual operating tier. Local mechanisms, simulation closed loops and real execution each therefore fall within a clear scope.

Summary: Imagination Errors Become Control Risk Only When Consumed

World model attacks can enter from training artifacts, physical conditions, latent state, dynamics, objectives and constraints, candidate ranking, the action head and feedback. State recursion, long-horizon rollout and discrete choice amplify local errors, but the harm still depends on how downstream consumes the predictions. Frozen weights do not equal a frozen system. Internal self-consistency does not equal real correctness. Image realism does not equal action safety. When no penetration is observed, the conclusion can still only be limited to the models, tasks, budgets and consumers that were tested. A local negative cannot be upgraded into the absence of an attack path. Identifying an attack path is only the first step. Once the world model enters control, the team has more to connect: uncertainty, external anchors, robust planning, action filtering, rollback and falsifiable testing. Chapter 18 will complete this set of runtime assurance and recovery rules.

\newpage

Chapter 18 Runtime Assurance, Recovery, and Falsifiable Testing

A world model reports a "95% confident" safe trajectory. The controller accelerates on that basis. Only afterward does it discover that the uncertainty estimate and the predictor share the same latent state. Both are therefore overconfident at once on the perturbed map. The calibration distribution, independence and failure conditions determine what the number 95% practically means. Those elements need to enter the control rule. A probabilistic score has runtime-assurance meaning only if it can change the action in time. It must also trigger more conservative control when its assumptions fail.

This chapter moves world model defense from "detecting anomalies" to "changing execution, limiting consequences, and recovering". Its core structures include provenance and artifact control, state anchors, uncertainty and safety filtering, robust model predictive control, execution-time verification, rollback and falsifiable testing. Every assurance responsibility must state its preconditions, its intervention action, its residual risk and its post-failure state. The ultimate object of evaluation is the verifiable behavior of the control chain.

Chapter Overview

This chapter places prevention, detection, filtering, robust control, execution verification and recovery at their own points of action. A calibrated uncertainty score turns into speed limiting, rejection, replanning or controller switching. External state anchors and model diversity expose common drift. The hardware safety envelope and the capability gate limit the maximum consequence at the execution end. At the same time, each intervention states its calibration distribution, its dynamics assumptions, its control deadline and its benign utility.

Stopping, rollback, replanning and human takeover are all expressed in terms of state transitions, response time, stopping distance, transaction integrity and post-recovery task outcome. Evidence is layered by static inspection, mechanism-level runs, simulated execution or closed-loop simulation, end-to-end reproduction and field events. Module fixtures, offline replay, scaled-down simulation, full simulation and isolated hardware closed loops each enter the tier they actually reach. Each tier's conclusions match the weights, data, configuration, environment and repetition count of the actual run.

Background Principles: Predictions Must Pass Through Anchoring, Gating, and Replanning

This chapter studies a rolling control chain. The world model reads the current observation, historical actions, the map, objectives and constraints. It maintains a state that can be updated across cycles, and it predicts the futures that candidate actions may form. An independent external anchor first checks whether the state still corresponds to reality. The uncertainty module then estimates the prediction error or the constraint margin. Model predictive control (MPC) selects only short actions that are feasible within a finite horizon. The action gate verifies object, units, permissions and the safety envelope. The execution receipt then becomes the next round's observation and

triggers replanning. The consumers include the planner, the safety controller, the actuators, the recovery component and the human on duty. No internal score can issue permission on behalf of these consumers.

Under normal conditions, fresh localization agrees with the map and the prediction interval covers the candidate trajectory. MPC issues a short action and immediately updates the state with real feedback. Under boundary conditions, the predictor and the uncertainty head share an erroneous representation, so both are overconfident at the same time. The external anchor, however, finds a position inconsistency. The action gate should reject long actions. The RTA switches to speed limiting or a safe stop, and replanning restarts from a fresh state rather than a contaminated cache. If the anchor, the solver or the receipt times out, the system likewise enters a predetermined fail-safe state. It must not continue along the previous nominal trajectory.

18.1 The Five Classes of Defensive Action Must Each Fulfill Their Own Role

World model defense can be divided into five classes. Prevention reduces the malicious data or artifacts that enter the system. Detection identifies anomalies in the input, the state or the imagination. Suppression rejects, clips or projects dangerous actions. Recovery returns the system to a verified state. Assurance responsibility connects monitoring, switching, constraints and failure conditions into auditable operating requirements.

Each of these five controls carries an independent responsibility. Provenance signatures confirm artifact identity, and behavior tests check for backdoors. Anomaly detectors output risk signals. The control consumer is responsible for blocking the action. Safety filters depend on trusted state. Rollback restores weights, physical state, cache and tool transactions at the same time. The missing control link can be located only if these responsibilities are kept separate.

An RTA record contains at least the protected interface, the monitored inputs, and the thresholds or sets. It also carries the intervention action, the control deadline, and benign utility. Missed detections/false alarms, attacker knowledge, assumption failure, the recovery path, and the verification evidence complete the record. An output that lacks an intervention action is a diagnosis. An assurance claim that lacks assumption-failure conditions is an unqualified commitment.

18.2 Prediction and State Assurance

Responsibility on the prediction side starts from artifact and runtime lineage. It passes through input types, state anchors and uncertainty calibration, and finally enters robust MPC. The four stages answer four questions in turn. What is the system currently using? Does the state still correspond to reality? How large is the prediction error? What can still be chosen under the current assumptions? If any stage is unknown, the horizon should be shortened, or capability contracted.

Supply chain and runtime lineage: first answer what the system is using

A world model may serve at the same time as environment generator, data supplier, planner, evaluator and safety checker. Each role requires versioned lineage. That lineage covers the field names, data types, units, coordinates, provenance and timestamps of data, code, weights, encoders and conditions. It also covers the reward/constraint, action head, configuration, deployment image and downstream consumers. An action decision should be traceable to the model, state, objective and safety policy versions it used.

Artifact integrity controls include hashing, signing, licensing, approval, isolated loading and revocable release. Behavioral controls compare predictions, states, rankings and actions under benign and triggered conditions. Covert poisoning of the SWAAP kind shows that benign performance and dynamics distance may be preserved, yet decision-relevant tails and downstream returns still need to be differenced [WM_hu2026swaap]. Signing and behavioral testing together cover "who it came from" and "what it did."

If a world model generates training data, lineage must cross model boundaries. Record the generation conditions, random seeds, model versions, filters and data batches, then connect them to downstream policies and

evaluation. When a world model version is revoked, one must also locate the data it generated, the policies trained from it, and cached states. Otherwise, supply chain effects persist after the source model is taken offline.

The deployment policy follows a least-capability rule. A model with incomplete lineage can be evaluated in an offline sandbox rather than connected directly to high-impact planning or control. As evidence escalates, data generation, candidate suggestions and restricted control are opened up progressively, with a fallback version at each level. Open weights or public repositories raise visibility and review conditions, but the trust level is still determined jointly by provenance, behavior and runtime evidence.

Inputs, conditions, and states: using independent anchors to detect joint drift

The state defense line must span pixels, maps, 3D boxes, text, action history, proprioceptive state and tool feedback. Each input carries provenance, time, coordinates, integrity and permitted influence. Type checking can block obviously out-of-bounds conditions. Cross-modal and temporal checks can detect inconsistencies. Independent anchors are responsible for detecting joint drift across multiple internal modules.

WISER uses the surprise and reconstruction error between a world model's prior distribution and posterior distribution. The prior distribution is the state distribution predicted from historical states and actions before the current observation is absorbed. The posterior distribution is updated after the evidence of the current observation is combined. The inconsistency between the two constitutes a surprise signal. On this basis the method selects a more interpretable latent state among multi-sensor candidates, and rejects or reconstructs untrustworthy observations in the single-sensor case [WM_zollicoffer2025worldmodelrobustnesssurprise]. It supports runtime detection and rejection or reconstruction of suspicious observations, but it mainly targets noise, faults and natural out-of-distribution inputs. Adaptive attacks that know the surprise score, the reconstructor and the task objective still require separate validation.

Illusory Attacks warns that ordinary adversarial trajectories may be easy to detect, because they deviate from the environment distribution, whereas attacks that deliberately stay close to the feasible trajectory distribution are harder for anomaly scores to find [WM_franzmeyer2024illusory]. The detector therefore checks "whether the input resembles the training data" and "whether the state conforms to independent physics, rules and task authorization" at the same time. Independent anchors have three conditions. First, the data source differs from the model under test, for example LiDAR versus monocular vision. Second, the parsing path differs, for example geometric rules versus a neural encoder. Third, the anchor captures safety-relevant variables, for example restricted zones, distance and control authority. Two models that share an encoder may fail jointly on the same perturbation even if their names differ.

State recovery requires a trusted starting point. The system can store signed checkpoints, the most recently verified pose and external map versions. When the physical world has already changed, the old latent state serves only as a clue. The recovery procedure should stop or limit energy first. It then re-observes, relocalizes and rebuilds the state, and re-obtains task authorization last.

Uncertainty and safety filtering: scores must change actions

Point predictions of a world model in out-of-distribution states may be overconfident. Safety filtering works by estimating the uncertainty of future costs or constraints. When risk exceeds a threshold, the filter can reject candidates or shorten the horizon. It can also switch to a conservative policy or request more observations. UNISafe uses latent-space uncertainty and safety filtering to avoid some OOD failures [WM_seo2025unisafe]. Its effectiveness depends on whether the uncertainty covers the true model bias. Let the predicted cost of a candidate action be $\hat{C}(a)$, the uncertainty upper bound be $U(a)$, and the risk budget be β . One conservative admission rule is

$$\hat{C}(a) + \lambda U(a) \leq \beta.$$

λ denotes the conservative strength. The intent is to leave margin for an uncertain future. This is not an unconditional probabilistic guarantee. The inequality may incorrectly let a candidate through in three cases. U may be uncalibrated. The model and the filter may be jointly biased. The true state may not be within the calibration distribution. The system must define detection signals for these conditions and the action to take after

failure. SafeDreamer learns safety constraints within world model imagination, showing that safety costs can enter policy optimization rather than only post hoc scoring [WM_huang2023safedreamer]. At training time, safety tendency and runtime hard constraints take on different responsibilities. The former changes the policy distribution under encoded danger. The latter handles unmodeled danger, reward/constraint poisoning and out-of-distribution states. These situations may still cross the optimization objective.

Filter acceptance must report benign pass rate, dangerous pass rate, false rejections, misses, intervention frequency and decision latency as one set. Rejecting everything can drive the dangerous pass rate to zero, and the system then loses its utility. A threshold that is too loose retains task capability but lets high-consequence actions pass through. Safety and utility should form a Pareto frontier under the same tasks, states and control deadlines. That frontier is the set of non-dominated operating points, at which one metric cannot be further improved without worsening another. An operating point matching the consequence level is then selected on that same frontier.

Robust MPC: connecting prediction intervals to constraints

At each cycle, model predictive control (MPC) rolls out predictions for a set of candidate trajectories. It selects a short segment of actions that satisfies the constraints, then replans with new observations. Robust MPC does not work from a single nominal future alone. It considers a disturbance set, probabilistic intervals, or the worst case. The world model provides states and dynamics. Turning uncertainty into executable constraints is the controller's responsibility.

SLS² combines a latent world model with robust MPC based on parallel conformal prediction. Conformal prediction is a class of methods that use calibration samples to construct coverage sets for the errors of new samples. The combination attempts to establish probabilistically safe control from pixel predictions [WM_nath2026sIs2]. An important advance of such methods is that the output changes the trajectory rather than stopping at a risk score. However, conformal coverage relies on the exchangeability of calibration and test data, or on corresponding assumptions. Dynamic distribution shift and adaptive attacks will degrade the coverage guarantee. Reporting should state the coverage object, the calibration unit, failure detection and the control used after coverage is exceeded.

Robust control must also satisfy a real-time budget. If world model generation, candidate sampling and constraint solving exceed the control period, a theoretically more conservative scheme may use stale states in reality. Acceptance should at least report P50, P95 and P99 latency, timeout rate and timed-out actions. Average latency does not cover tail deadline misses. Conclusions about model bias and control assurance use independent evidence. That MPC selects a trajectory shows that the optimizer found a candidate. That world model predictions are smooth shows internal temporal continuity. Future reliability and control feasibility are additionally confirmed by external reachability, actuator limits, control barriers, or emergency stops. The less trustworthy the prediction, the shorter the action chunk and the lower the speed. When necessary, the system switches to a safety controller that does not depend on the world model.

18.3 Execution and Recovery Assurance

A prediction that passes the preceding checks is still only execution evidence. The execution side must align plans, previews and actual actions. An out-of-model capability gate decides release, projection, rejection or replanning. After an action occurs, independent feedback confirms whether the environment or the transaction actually changed. When a deviation is detected, the recovery state machine restricts authority according to the consequence first. It then selects a trusted recovery point and verifies the conditions for reopening.

Runtime verification: a preview is not authorization

Runtime verification predicts the consequences of an action before it is issued or during execution, and checks them against rules, states and feedback. CheckVLA uses an action-conditioned world model to preview long-horizon mobile manipulation and check candidate executions. DreamGuard uses a risk-aware world model to block high-risk paths before agent tool calls [WM_liu2026checkvla, @WM_lin2026dreamguard]. Both turn predictions into interventions, yet they still require an external capability gate.

When the checking model and the working model share data, an encoder, or a training objective, the result is "the same kind of imagination validating the same kind of imagination." A more robust design gives the verifier different sensors, physical approximations, rules and failure modes. When full independence cannot be achieved, the sharing relationship can be written into the threat model, and the maximum action the verifier is able to approve can be lowered.

Execution verification uses three-way consistency. The three terms are the objective required by the plan, the future previewed by the world model, and the action the controller is actually about to execute. Any inconsistency among the three must not be automatically corrected to the version the model most prefers. It should instead be rejected, re-observed, or referred for authorization. After execution, real feedback is used to check the predicted against the actual, and the deviation is written into the next cycle's risk budget.

Tool calls require transactional control. Before the call, preview side effects and verify parameters and permissions. During execution, use least-privilege scopes and short-lived tokens. After the call, verify the result against external facts. A world model can predict "what happens if a file is deleted," but it cannot approve the deletion on its own. A completion return also cannot replace independent state confirmation from the storage system.

Stopping, Rollback, and Takeover Are Measurable States

The goal of recovery is not to bring an anomaly score down. It is to return the system to a verifiable low-risk state. Recovery states can include shortening actions, decelerating, shutting down, or retreating to the most recent safe pose. They can also include switching to an emergency policy, rolling back tool transactions, rebuilding world state, revoking a model version, and handing over to a human. Every kind of recovery must define success. Stop success can be measured by stopping distance, maximum residual energy and latching persistence. Rollback success must verify that the physical position or business transaction has been restored. Replanning success must confirm that the new trajectory no longer consumes contaminated state. Human takeover must record the request, the response, the transfer of control and the timeout policy. A task that is not completed but stops safely should be recorded as a safe failure, not as a normal success.

Recovery latency must be smaller than the harm window. Slow language or video checkers can serve low-frequency task entry points, but they are not suitable for millisecond-scale braking. Hard limiting, emergency controllers and physical stops carry the high-frequency path. Higher-level models are responsible for interpretation, replanning and after-the-fact evidence. A combination of controls at different time scales is more verifiable than having a single large model handle all safety decisions.

Rollback also requires state lineage. Revoking weights while continuing to use the caches, maps or policies they generated is not a complete rollback. Restoring an old latent state while ignoring that the real device has already moved is likewise not a complete recovery. Runbooks should list the rollback action for each of model state, software state, tool transactions and physical state. They should also list isolation and relief measures for when rollback is irreversible.

18.4 A Falsifiable Test Matrix

A falsifiable claim must expand "the system can detect anomalies" into concrete conditions. Under a specified model, task, attack privilege and budget, the monitor fires before the P99 decision deadline. The safety controller restricts actions to a given set. At the same time, benign task success, false refusals and resource cost meet thresholds. If any one of these fails, the assurance claim is falsified or downgraded. World model tests cover at least six axes:


1. **Interface:** supply chain, observation, state, dynamics, goals/constraints, planning/action, feedback.
2. **Privilege:** data writes, white-box gradients, black-box queries, physical conditions, configuration writes, tool permissions.
3. **Time:** single frame, cross-frame, delayed triggering, long rollouts, replanning, recurrence after recovery.
4. **Consumer:** offline generation, data generation, policy evaluation, candidate ranking, action head, feedback control.

5. **Consequence:** prediction, ranking, action proposal, simulated execution, isolated real machines, and real-world events.
6. **Defense awareness:** unaware of the defense, knowing the structure, knowing the thresholds, jointly optimizing multiple layers of control.

Every test unit reports its attack budget, denominator, benign utility, residual risk, P95/P99 latency, random seeds, independent units and stopping conditions. Numbers are never pooled into an aggregate across different models, tasks, budgets, ASR definitions or variances. Video quality, latent distance, return, task success, collision and recovery time stand as separate endpoints. Test units must retain both natural stress and malicious attacks. Natural noise, occlusion and OOD probe system reliability. Adaptive adversaries probe whether detectors and controllers stay effective under malicious optimization. The two are not merged into a single average score, yet they may share state, action and recovery metrics.

18.5 The Evidence Ladder: You Only Prove What You Ran

Work up from the bottom of the ladder and compare what each level actually connects. Look especially at the consumers and receipts still missing across three statements. Those statements are "a module output changes," "a simulation environment consumes an action," and "the target system closes the loop end to end."



A five-level ladder rising from static inspection to mechanism-level runs, simulation closed loops, end-to-end reproduction, and field events.

The ladder supports management of the limits of conclusions. It is not a single ranking in which "higher is better." Static inspection and mechanism-level runs locate interface and formula problems quickly. A simulation closed loop observes state-action feedback. None of them substitutes for missing real weights, data, identity, actuators or field causal evidence. Failures and items that were never run must likewise be retained at their corresponding level of the ladder.

World model safety evidence uses one uniform ladder: static inspection, mechanism-level runs, simulated execution or simulation closed loops, end-to-end reproduction, and field events. Read-only verification of documents, formulas and code paths counts as static inspection. A mechanism-level run is formed once formula proxies, module fixtures or offline replays actually run on controlled inputs and retain verifiable receipts. A simulation closed loop is formed once state, actions and environment feedback are connected. End-to-end reproduction requires that real or verified-equivalent weights, data, identity, configuration, the target pipeline, predetermined metrics and multiple valid runs all be closed. Field events separately record a specific time, object and consequence.

The accompanying execution receipt records 6 mechanism-level runs, 3 static checks and 0 end-to-end reproductions. All 9 low-risk commands in the safety suite exited successfully. That evidence covers the runtime behavior of the corresponding submodules or formula proxies on constructed inputs. The paper's numbers, official large weights, the full data pipeline, the robot closed loop and real-world safety all remain in an unreproduced state. Readers can trace this set of states back from [EW-S028](#) in the source ledger. That visible pointer leads to [the final verification receipt](#) and [the evidence ledger](#).

This boundary applies to any engineering team as well. A repository that exists does not mean the code runs. Installing dependencies does not mean the model is correct. A changing module output does not mean an attack hits the task. Simulation success does not mean real-machine safety. Every escalation should retain the command, version, inputs, outputs, exit code, environment, hashes and failure records. High-risk real-device testing escalates in steps. First verify stopping, limiting and rollback in reversible simulation and hardware-in-the-loop. Only then use low-energy endpoints with no third-party impact on isolated devices. A safe stopping condition is a prerequisite gate for entering the next evidence level.

18.6 From Claim to Evidence: Writing a Reviewable Assurance Argument

Runtime assurance is not a single conclusion that "safety testing has been passed." It is a structure composed of claims, reasons, evidence and counter-evidence conditions. A claim must be bound to the system version, operating domain, task, threat privilege, outcome object and time window. "A safety filter can reduce risk" is not verifiable. Only a claim such as "within a specified speed and payload, the safety gate rejects candidate actions that cross an external keep-out zone before the control deadline" can be overturned by a test.

An assurance argument has at least four layers. The top layer names the environmental or business event that is not desired to occur. The second layer holds the action, privilege and energy boundaries that can directly constrain that event. The third layer holds intermediate evidence such as state, future, uncertainty and execution feedback. The fourth layer holds dependencies such as versions, configuration, sensors, clocks and compute. Intermediate scores and the control connections that actually change actions or privileges support high-level claims jointly.

The detector's input fields include observation ID, latent state version, provenance principal, observation time, coordinate frame and integrity flag. Its output fields include risk score, reason code, applicable object, validity period and calibration version. The safety filter reads the candidate action ID, state version, constraint set, risk quantity and control deadline. It outputs an allow, project, replace or reject decision, along with a reason code and action parameters. The recoverer reads the fault code, current mode, most recent trusted anchor, unfinished transactions and physical position. It outputs the degraded mode, completed compensation, outstanding items and trace records. A score on a dashboard is only a diagnostic signal. Changing actions still requires an explicit control consumer.

Every claim must also have a counter-evidence path, and each condition must be connected item by item to a disposition. When latency exceeds the deadline, discard the expired result and switch to rate limiting or to a safety controller that does not depend on that result. When the calibration distribution fails, widen the safety margin, shorten the horizon and maintain restricted operation until recalibration. When an external anchor is unavailable, re-observe. If it is still unavailable, stop execution that depends on that anchor. When the candidate set contains no safe action, do not continue to select the best among dangerous candidates. Switch instead to a verified safe hold or safe shutdown. When the safe set is unreachable, execute a verified minimum-consequence emergency action and request human takeover. When rollback is incomplete, continue isolating the relevant artifacts, state and capabilities, and prohibit unlocking. All six paths are maintained until new evidence satisfies their respective unlocking conditions.

18.7 Calibration Requirements: Coverage Has a Denominator and Failure Conditions

For the coverage of an uncertainty or conformal interval, one must first answer "coverage of what." A world model's calibration unit can be a single frame, a time window, a candidate trajectory, a complete task or the maximum constraint error. Coverage of 95% per frame does not mean a long task has a 95% probability of staying within bounds throughout. Correlation between adjacent frames, multiple replannings and candidate selection all change the overall event.

Calibration data must be explicitly related to deployment conditions. The error distribution can shift when any one of these changes: scenario, speed, payload, sensors, action chunks, horizon, number of candidates, or control frequency. Exchangeability or stability assumptions in dynamic environments are continuously tested over time windows. SLS² connects conformal error bounds to robust MPC, so that prediction intervals actually affect the

trajectory. The scope of the probabilistic language is limited by the calibration assumptions and the latent constraint mapping [WM_nath2026sls2].

At runtime, the system monitors both the final score and calibration applicability. Available input range, representation distance, error sequences, coverage violations and external state discrepancies serve as failure signals. When any one of these signals exceeds a predetermined boundary, the original coverage guarantee no longer applies to the current operating conditions. The system then widens the safety margin, shortens the trajectory, actively acquires new observations or enters the safe set. It maintains conservative control until recalibration or new evidence is obtained.

Shared bias is a key counterexample for calibration. The predictor, uncertainty head and constraint checker may share an encoder and data. Then they may be simultaneously overconfident on the same out-of-distribution state. UNISafe ties latent dynamics disagreement to safety filtering, and it constitutes runtime control that can change actions. Undetected shared model bias nevertheless remains a residual risk [WM_seo2025unisafe]. A calibration report should at least list the number of independent units in the calibration set and the test set, nominal coverage, empirical coverage, confidence intervals, the longest task, distribution strata and uncovered events. A report of only "95%" that gives no denominator, no time unit and no disposition for violations is not sufficient to grant action authority.

18.8 The Assumption Ledger for Robust Control

Robust MPC, reachability analysis and control barrier functions all require an assumption ledger. The ledger records more than formulas. It records whether the state can be observed, whether the error lies within the disturbance set, and whether the actuators respond as commanded. It also records whether the constraints are complete, whether the optimizer obtains a feasible solution before the deadline, and whether the terminal set can really be occupied safely.

Feasibility is not recursive feasibility. A trajectory may satisfy the constraints in the current cycle. That does not guarantee a safe candidate still exists in the next cycle, after the first segment is executed. The controller must reserve margin for braking, steering, compute timeout and observation failure. As far as possible, it should contract the trajectory into a verified safe terminal set. Suppose the terminal state is merely an "apparently stationary" state imagined by the world model, with no kinematic and braking-distance evidence. Then it is not a safe set.

Optimization failure is itself an output that must be designed for. The system enters a predefined restricted action set on any of several conditions: an empty candidate set, conflicting constraints, a solver timeout, numerical anomalies, or a stale state. That set includes hold, decelerate, retreat along a verified path and stop. Its reachability and latency need to be verified separately. The variables of the world model and of the controller must also be separated. When evaluating a fixed controller, freeze the model weights, constraints, safety policy and solver configuration, and change only the observation or the admissible operating conditions. When evaluating defense-informed attacks, the optimization variables may be the input, the conditions or the query policy. They must not secretly gain the ability to change hard constraints. When evaluating training-based safety methods, trainable value, cost or dynamics parameters are listed separately. They are not mixed into the same column as the runtime attack variables.

Safety cost brought into imagination-based training lets the policy account for constraints during optimization. SafeDreamer provides one such path [WM_huang2023safedreamer]. That route is complementary to runtime hard constraints. The safety cost may fail to cover a certain danger, or the cost head and the world model may be jointly biased. In that case a low predicted cost at training time describes only the model's internal objective. External constraints and operational outcomes measure real-environment risk.

18.9 Independent Verification and Common-Cause Failure

The value of multiple layers of defense depends on whether their failures are correlated. Three models may share the same training set, the same encoder and the same map source. Different names alone do not automatically make them three roots of trust. Independence should be recorded item by item along data, representation,

algorithm, rules, hardware, clock, organizational approval and deployment channel. A "defense dependency graph" can be established. The nodes are the predictor, the uncertainty head, the risk discriminator, the rule engine, the action gate, the actuator, the external anchor and the recoverer. The edges represent shared training data, encoder, configuration, clock, power, network or permissions. Attacks and fault drills are injected along these shared edges to verify whether the so-called backups fail together.

The output of a verifier does not directly become unlimited authority. The authorization service issues short-lived, narrowly scoped capability tokens. Such a token binds the action type, object, quantity, space, time, state version and revocation conditions. The world model can provide future evidence for a candidate. The rule engine and the independent state source determine the token ceiling. The actuator accepts only signed and unexpired tokens.

CheckVLA uses an action-conditioned world model to preview long-horizon mobile manipulation and to check execution deviation. DreamGuard predicts risk before an agent's tool call and intervenes [WM_liu2026checkvla, @WM_lin2026dreamguard]. These two paths show that a world model can become a pre-execution and in-execution checker. They do not automatically resolve shared weaknesses. The sharing relationship between the checking model and the working model must enter the dependency graph. High-impact actions remain under the control of an external capability gate.

Independence is not a binary property. A set of geometric rules can be algorithmically independent of a vision model yet still share the same erroneous extrinsic calibration. A hardware emergency stop can be independent of model computation yet still share a power supply with the main controller. Review should state the common causes that are covered and the common causes that are not, together with their impact on the maximum consequence. It should not substitute the phrase "double insurance".

18.10 Verification Fixtures: From Array Branches to Controlled Physical Actions

The first-level fixture is a pure function or array. On constructed inputs, it tests whether these branches take effect: surprise magnitude, cost filtering, uncertainty gating, candidate projection and policy switching. The denominator of such fixtures is valid constructed samples or function calls, not environment tasks. It is suited to finding sign, threshold, tensor-shape and faulty-fallback problems. It does not support task safety rates.

The second level is offline replay. Here the model, controller and logs are fixed. The defenses' decisions are compared across normal, naturally anomalous, targeted-perturbation and defense-informed perturbation cases. It can evaluate scores, candidate selection and recorded actions. It cannot observe the effect of replaced actions on the subsequent environment. Therefore, the "execution" at this level is at most a simulated execution of recorded commands, not a dynamics closed loop.


The third level is resettable simulation. Here the defenses' decisions genuinely change the simulation state. Closed-loop tasks, constraint events, state recovery and adaptive adversaries can therefore be measured. Simulation results apply only to its sensing, dynamics, contact, tool and human behavior models. A simulated collision is a simulation event, not real property damage. A simulation without a collision is not proof of real-world safety either.

The fourth level is hardware-in-the-loop and isolated low-energy devices. Such a setup can test real clocks, networks, drivers, sensors and stopping distances. Physical isolation, energy limits, bypass emergency stops, on-site observation and the authority to abort must all be established in advance. Without them, do not use real devices to make up for simulation uncertainty. Physical runs also support only controlled operating conditions and measured endpoints. They do not extend automatically to unsupervised deployment.

Every level must freeze the system under test. The optimization variables and the attack budget are listed, and clean controls, target-irrelevant controls and defense-informed controls are all retained. Engineering denominators are counted by independent trajectory, complete task, unique tool transaction or physical object. Adjacent frames that share an initial state and cache are grouped into the same independent unit. Pass criteria cover benign utility, dangerous penetration, false rejection, latency, resources, recovery and evidence completeness at the same time.

18.11 The Recovery State Machine and Fault Drills

The figure shows one loop. Follow it to see why capability is restricted first once an anomaly is detected, and why evidence is then preserved before a recovery point is chosen. Follow it also to check which independent verification conditions are required before each return to normal operation.



Six ring states return to normal after passing from normal operation through capability restriction, safe stop, frozen forensics, trusted rebuild, and verified recovery.

The figure separates "restarting the service" from "trusted recovery". Process availability, model loading and the disappearance of an alert are each insufficient to unlock high-impact capabilities. It also does not presuppose that any fault is reversible. Suppose the transaction cannot be compensated, the physical environment has already changed, or no trusted recovery point exists. The system must then maintain isolation, manual control or a longer-term safe stop.

Recovery is not a "retry" button. It is a state machine with guard conditions. A basic state set may include normal, observing, restricted, safe stop, rebuild state, pending authorization and manual control. A return from restricted to normal requires a fresh independent state and trusted model and policy versions. Unfinished transactions must have been handled, the physical area must be available, and a new authorization must hold. "The alert disappeared" is not sufficient to unlock.

Every transition has a time limit and a failure exit. A checker timeout puts the system into restricted. When state rebuilding does not reach anchor consistency within the specified time, the system enters safe stop. When a manual takeover request times out, a predefined policy keeps energy low. No transition may default to restoring high authority because no one responded. Fault drills cover four classes of state: model, software, tool and physical. Model drills replace erroneous weights, calibration tables or map versions. Software drills inject timeouts, out-of-order logs, cache poisoning and network partitions. Tool drills simulate a request that has been received but whose side effect is incomplete. Physical drills inject localization loss, braking deviation or emergency-stop channel failure under controlled conditions.

Recovery metrics include the time to notice capability restriction, stopping distance, maximum residual energy and the time to enter a trusted state. They also include human response, successful transaction compensation, completion of the checks before re-authorization and recurrence after recovery. The denominator is the independent drills that actually entered the corresponding fault state. Startup failures that never reached the injection point enter the availability denominator.

The drills themselves have explicit safety constraints and stopping scope. State transitions are first verified with injection code and virtual actuators, then moved to resettable simulation and hardware-in-the-loop. Controlled device drills are performed only when a low-risk endpoint, physical isolation, bypass stopping and an on-site responsible person are all available at the same time. Integrity is verified jointly through lossless triggering, controlled stopping, state rollback and recovery receipt.

18.12 Case Comparison: One Defense Does Not Cover All First-Broken Interfaces

WISER starts from prior/posterior distribution surprise and reconstruction error. It selects among multi-sensor representations, and in the single-sensor case it rejects or reconstructs untrustworthy observations [WM_zollicoffer2025worldmodelrobustnesssurprise]. Its intervention sits mainly on the path from input to latent state, and its actual outputs are representation selection, rejection or sanitization. Poisoned weights, malicious rewards and bypassed tool permissions lie at other interfaces. An adversary that jointly optimizes surprise, reconstruction and the task objective requires new adaptive testing.

UNISafe takes the path "dynamics divergence or a known danger signal → action replacement", so it contains both detection and suppression. SafeDreamer incorporates safety cost into imagination-based policy optimization. SLS² brings calibration error bounds into robust MPC [WM_seo2025unisafe, @WM_huang2023safedreamer, @WM_nath2026sls2]. All three come closer to runtime assurance than methods that only produce a risk score do. Their protection assumptions differ, though: whether divergence can reveal unknown errors, whether the safety cost is complete, and whether the calibration and the latent constraint mapping continue to hold.

CheckVLA and DreamGuard push the intervention point to execution time. The former reads action chunks, world model previews and real feedback, while the latter reads the agent trajectory prefix, candidate tool actions and predicted risk [WM_liu2026checkvla, @WM_lin2026dreamguard]. One serves physical action verification, the other pre-tool-call intervention, and each is evaluated on its corresponding task, latency and consequence object. The gap common to all six classes of methods is independent recovery metrics and defense-informed adversaries. A method that is effective under natural OOD, model mismatch or safety constraints does not automatically become a defense against malicious optimization. A comparison table should separately list the first-broken interface, detection signal, intervention action, protection assumption, benign utility, adaptive evaluation, recovery and evidence layer. It should not hand out a single cross-endpoint "best" label.

18.13 Release Gate: From Control Rules to Runbook

Before a world model enters a high-impact consumer, the release gate checks at least the following items. When any high-impact item cannot be determined, the system remains in capability contraction or a stopped state, and does not interpret missing evidence as a default pass:

1. The actual inputs and outputs determine the family and the consumer, and project names are retained only as product identifiers.
2. Lineage closure for the names, types, units, coordinates and provenance of data, weights and conditioning fields, as well as rewards, constraints, controllers and deployment images.
3. Frozen weights, runtime state, optimizable variables, attacker permissions and budgets are listed separately.
4. Inputs and conditioning carry type, unit, coordinate, time, provenance and integrity constraints.
5. State, future, ranking and action are each compared against an independent anchor or a hard rule.
6. The calibration unit, denominator, applicability domain and failure action for uncertainty and coverage are stated.
7. Every risk output has an intervening consumer, and that intervention takes effect before the P99 deadline.
8. Infeasible planning, check timeout, anchor failure and out-of-order feedback all lead to a safe fallback rather than a high-privilege default path.
9. Shared dependencies among the model, safety checker, rule engine, actuator and emergency stop are diagrammed and exercised.
10. Model state, software state, tool transactions and physical state each have recovery success conditions.
1. Benign utility, false refusal, false negative, penetration, latency, resource, stop and recovery metrics use appropriately independent units.
2. Static checks, mechanism-level runs, simulated execution or simulation closed loops, end-to-end reproduction and field incidents are recorded in layers. No conclusion exceeds the evidence layer it actually reaches.

Release is a continuous state. After a substantive change to the model, map, sensors, task, tool permissions, calibration set, control period or safety rules, the related claims return to a to-be-verified state. Small-scope canary release collects evidence only on low-risk, protected objects. The runbook must turn the above items into executable decisions. Every alert has a subject, severity, action, time limit, escalation target and unlock evidence. Every capability contraction has its business impact and maximum duration. Every recovery has either two-person verification or automated independent verification conditions. Without an owner and a deadline, a checklist does not automatically become a line of defense in an incident.

18.14 Residual Risk, Utility, and the Assurance Budget

No line of defense can reduce world-model risk to zero. Engineering decisions must allocate residual risk across five accounts: entry, detection, intervention, recovery and evidence. The entry account records supply chains, conditioning and feedback that are not yet covered. The detection account records adversaries that adapt to low surprise or low uncertainty. The intervention account records the absence of safe candidates, timeouts and actuator deviation. The recovery account records irreversible events. The evidence account records the levels that have not yet been run.

The assurance budget connects these gaps to maximum capability. When independent localization is unavailable, reduce speed, action chunks and spatial extent. When calibration applicability fails, widen the margin and prohibit long-horizon autonomous execution. When tool receipts cannot be verified, prohibit chained side effects. When recovery drills are not closed, do not open up irreversible actions. When evidence decreases, capability contracts accordingly rather than staying unchanged.

Utility is not the opposite of safety. It is a boundary that must be reported in the same table. Stopping everything can reduce some physical risk, yet it may create new dangers in medical, transportation or energy tasks. Therefore refusal rate, task completion, throughput, energy consumption, human burden and downtime are reported together with penetration, constraint events and recovery. A report that measures utility only on benign tasks and safety only on dangerous tasks hides what gating costs on a real task mix.

A residual risk record includes the risk owner, affected objects, existing controls, highest evidence layer, unverified assumptions, temporary capability ceiling, remediation date and release conditions. For high-impact items, "accepting the risk" must specify who has the authority to accept it and on what basis. It must also state how long the acceptance is valid and when it expires automatically. Closing a risk requires verified controls, a valid monitoring window and a named signature.

Three counterfactual questions can converge an advanced design review. First, suppose the predictor and the uncertainty head fail at the same time. Which external control can still limit the maximum consequence? Second, suppose all computed metrics are normal while physical feedback disagrees. Which side will the system trust? Third, suppose a judgment cannot be completed within the harm window. What safe action does not depend on that judgment? If any question lacks an executable answer, the release gate stays closed. The system then moves into capability contraction.

Assurance gaps must also be versioned like software defects. Each entry records the claim it belongs to, the version in which it was first found, affected consumers and triggerable operating conditions. It also records the current capability ceiling, compensating controls, the responsible person, the expiry date and the closure evidence. The same gap may require only restricting data release on an offline generator, yet may require immediately limiting speed or stopping on a WCM. Priority is determined by the harm window and maximum permissions, not by model parameter count.

Gaps fall into at least three kinds. An evidence gap means the mechanism may hold but has not yet reached the target operating layer. A coverage gap means the tested scenarios did not include a certain class of conditions, permissions or adaptive adversary. An independence gap means multiple lines of defense still share a dependency that can be broken simultaneously. The first closes by escalating the controlled evidence layer, the second by extending layered testing, the third by adding different trust roots. Random samples only add observations within the existing distribution.

Closing an assurance gap means linking the new evidence back to the original claim. The maximum allowed capability must then be recalculated. New evidence that covers only daytime low-speed operation will not automatically close a nighttime high-speed entry. A single passing simulation regression will not close an evidence gap about the hardware clock and braking window. Only when the applicability scope, denominator, deadline and failure records are all closed can the corresponding capability contraction be lifted.

18.15 Worked Example: Runtime Assurance for an Automated Yard

The yard uses EWM to predict traffic evolution 20 s ahead, WAM to select candidate routes and WCM to update speed every 100 ms. The high-performance controller pursues throughput, while the safety controller is responsible only for decelerating, stopping and returning to the waiting area. Independent localization, zone rules and the hardware emergency stop do not depend on the primary world model. Monitors check the map version, localization-vision consistency, future reachability, candidate ranking stability and action-future consistency. When risk rises, the system first shortens the planning horizon and action chunks. Above a second threshold it switches to the safety controller. On localization distortion or check timeout it latches a stop. Recovery requires relocalization and loading an approved map. State is rebuilt, and the duty operator must confirm that the zone is clear.

Testing starts from offline logs, then adds map-condition shift, single-frame perturbation, delayed feedback, candidate-ranking attacks and adaptive search against the thresholds. Once simulation passes, testing moves to hardware-in-the-loop, and finally to a closed low-speed site. Each stage records normal throughput, dangerous route proposals, gate penetration, false stops, P99 latency, stopping distance, recovery time and takeover success.

If the primary model and the uncertainty filter are jointly overconfident, external zone rules should still prevent entry into restricted zones. If the rules are misconfigured, hardware speed limiting and a manual emergency stop limit the maximum consequence. The three controls do not make the system absolutely safe, but they turn a single model error into an observable, limitable and recoverable event.

18.16 Six Misreadings About Scores, Algorithm Names, and Local Runs

Misreading	Implied false premise	Correction rule available
High confidence or low uncertainty means safety	It treats an internal score as an independent fact, and ignores the calibration distribution and the shared bias	Take the score together with external state anchors, calibration error, and assumption-failure conditions
Using MPC automatically provides a safety guarantee	It treats the algorithm name as an unconditional certificate	The guarantee is tied to state, model error, constraints, real-time behavior, and solution feasibility. If any assumption fails, enter degraded control
A pre-execution preview can replace the capability gate	It mistakes the evidence object for the authorization object	The preview provides evidence only. Final execution authority still rests with independent permissions and safety constraints
Reloading the model after a stop is recovery	It restores only the weights, and ignores derived state and real devices	Caches, maps, tool transactions, and physical state are recovered and verified separately
A fixed attack not penetrating proves the system is robust	It extrapolates zero hits under one attack family and budget into a global conclusion	The conclusion covers only the frozen protocol. Adaptive attacks that know the defense are added, along with a held-out set of unknown attacks
A successful mechanism-fixture run equals end-to-end reproduction	It mistakes a runnable local mechanism for a reproduction of the complete system	End-to-end reproduction has not yet been completed. It can be considered complete only after real weights, data, identities, complete closed-loop tasks, predetermined metrics, and multiple valid runs are all closed

18.17 Bringing This into Real Systems

The yard system can place WISER, UNISafe, SafeDreamer, SLS², CheckVLA, and DreamGuard onto the chain of detection, filtering, control, and execution verification. Each mechanism may span two positions on that chain.

Every mechanism must still ultimately state its input signals, its intervening actions, and its consumers. Take an action condition such as $\hat{C}(a) + \lambda U(a) \leq \beta$. Its premises are state estimation, cost prediction, uncertainty calibration, and deadlines. Should any premise fail, the response should be rejection, projection, low-speed hold, or safe stop.

Real-time checks read tail latency to decide whether the control period can be entered. Suppose the mean is 60 ms, P99 is 180 ms, and the period is 100 ms. The system must then define an observable action for a timeout, and count the timeout proportion into benign utility and residual risk. Recovery objects must also be separated. Weights return to a signed version. Latent state is rebuilt from a trusted anchor, and downstream policies are rebound to the version. Tool transactions are rolled back or compensated, and physical devices enter a confirmable safe state.

The test matrix covers the supply chain, state, ranking, and defense-aware attacks at the same time. The claim that “the system stays outside dynamic exclusion zones” rests jointly on external state anchors, action limits, real-time assumptions, and counterexamples that can downgrade the claim. Per-frame coverage and maximum-error coverage for a complete task use different calibration units and denominators. Under distribution shift, the system chooses restricted operation or stopping based on task-level evidence.

The predictor, uncertainty head, rule engine, safety controller, actuator, and emergency stop enter the same dependency graph. A shared clock, incorrect extrinsics, and a power failure can reveal common-mode failures beneath apparent redundancy. Operating state advances along “observing—restricted—safe stop—rebuilding state—awaiting authorization.” Every transition has a guard, a timeout path, and success criteria. When no one responds, high privileges remain closed.

Low-risk validation connects the “detection—limitation—recovery” chain step by step, starting from array fixtures, offline replay, and resettable simulation. Records include the system version, interface definitions, attacker permissions, optimization variables, denominators, stopping conditions, command receipts, raw metrics, and failure samples. From the receipts alone, an independent reviewer should be able to reconstruct the denominator. That reviewer should also be able to tell apart the model not being loaded, the attack not arriving, the detector triggering, the action gate blocking, and the environment not presenting. The final deliverable is a set of falsifiable claims, a dependency graph, a degraded state machine, and a list of unclosed risks. When a new model or new permission is added, verification is redone from the affected nodes.

Summary: Assurance Claims Must Be Able to Fail, Intervene, and Recover

Runtime assurance for world models is jointly constituted by versioned lineage, independent state anchors, uncertainty filtering, robust MPC, execution-time verification, capability gates, and a recovery state machine. Any probability score, future preview, or mathematical controller can provide evidence only within explicit assumptions. Each must change actions in time, and switch to more conservative control when those assumptions fail.

Falsifiable tests fix the interface, permissions, time, consumers, consequence, and defense awareness, and they report benign utility, residual risk, latency, and recovery. Static inspection, mechanism fixtures, simulation, and real robots each support different conclusions. At this point, closed-loop safety for VLM/VLA/WAM and WM/EWM/WAM/WCM already has a unified interface language. The next part will assemble these controls together with language, image, and video systems into a cross-domain security architecture.

\newpage

Part Six: Engineering Closed Loop

Guide to This Part

The first five parts deal with language, media, action, and world prediction. Production systems, however, usually contain several of these chains at the same time. A research assistant reads web pages, calls language models, generates charts, saves memory, and sends email. A robot platform may use a vision-language model, a world

model, tool services, and runtime assurance all at once. Suppose each team keeps only its own control checklist. Then provenance labels, identity, state, and recovery responsibility will be lost at cross-domain handoffs.

Chapter 19 therefore distinguishes the data plane from the control plane. The data plane handles prompts, images, video, state, and actions. The control plane provides provenance, identity, purpose, state lifecycle, capability, budget, traces, revocation, and recovery decisions, and it receives actual execution receipts. Shared services do not mean that all models share the same security assumptions. Language models still need instruction and tool controls. Vision systems still need condition and authenticity signals. Embodied and world models still need action gates, state anchors, and control deadlines.

Chapter 20 turns the architecture into a regime that runs every day. The system first freezes version snapshots of models, data, tools, policies, environments, and consumers, and then generates a test matrix from the threat record. Non-compensable failures enter hard gates. Diagnostic scores serve troubleshooting and ranking, and restricted canary connections are continuously monitored. After an incident occurs, the team first limits reachable capability and preserves evidence. It then determines the impact boundary, the trusted rebuild point, and the recovery conditions. A new version or a new environment makes the relevant evidence expire, so the operational closed loop must recompute dependencies and the scope of retesting.

The normal case is a tool field update that affects only parameter parsing, the authorization policy, and event logging. The dependency graph lets the team rerun these local tests, and end-to-end paths confirm that the combination still correctly refuses or executes. The boundary case is a team that updates only the release sign-off document while reusing the previous version's model scores and security receipts. The files are complete, yet they do not prove anything about the current system. Evidence validity comes from binding versions to the object actually running, not from the number of documents.

Here this survey returns to the five original questions. They are what to protect, where trust changes, how capability is obtained, how far the evidence extends, and how to recover after failure. New model families, protocols, and product names will keep appearing, and these five questions can still put them back onto observable interfaces. Readers do not need to become implementation experts in every technical branch. They can still organize design reviews, testing, release, monitoring, and incident response by combining unified traces with specialized mechanisms.

The engineering closed loop does not eliminate the unknown; it maps the unknown into capability limits, follow-up evidence tasks, and expiry conditions. When action receipts are missing, high-impact paths are not opened. When provenance or rights are unverified, data use is restricted. When calibration fails, control switches to a conservative mode. When a recovery point is untrustworthy, isolation is extended. The regime thus formed allows the system to keep evolving as models and environments change. It also gives every capability expansion a traceable evidentiary basis.

\newpage

Chapter 19: Cross-Domain Defense-in-Depth Architecture

An enterprise simultaneously deploys a customer-service agent, marketing video generation, warehouse robots, and a world model for inventory planning. Four teams separately purchase content classifiers, watermarking, robot emergency stops, and anomaly detection services. Yet they still share the same identity system, object storage, model repository, retrieval platform, logs, and release pipeline. A single upstream artifact losing trust, or a long-lived token leaking, may bypass all four sets of “specialized defenses.” The local controls have not failed, yet the system may still slip through the gaps between them.

The core task of a cross-domain architecture is to turn provenance, identity, state, capability, budget, traces, and recovery into shared security services. Models can be replaced, modalities can be added, and actions can escalate from sending email to controlling a robotic arm. As long as trust-domain changes still pass through these services, the team can maintain a traceable, replayable control plane. Such shared services embed specialized classifiers, watermarks, or safety controllers into stable permission and state boundaries. They thereby reduce how far a

single model upgrade disturbs the entire defense. This survey calls a region that shares trust assumptions, control authority, or policy a trust domain. It calls a location where the principal, administrative domain, control authority, or trust level changes a trust boundary. A boundary is a location, not a synonym for a trust domain.

Chapter Overview

This chapter assembles the specialized defenses of language models, visual generation, VLA/WAM, and world models into seven classes of shared security services. The provenance and artifact service answers what the system is currently using. The identity and purpose propagation service limits information escalation. The state service manages writes, expiry, and derivative deletion. The capability service retains final authorization. The isolation and budget service limits reachable scope. The unified trace and recovery service makes failures localizable and revocable.


Differences in roots of trust determine the effectiveness of defense in depth. This chapter explicitly draws the dependencies of shared models, labels, identity, configuration, and execution authority. It also defines information flows and capability flows separately for data, state, plans, and actions. These rules remain stable across model versions, modalities, and execution environments. Through a minimal assurance argument, they connect launch scope, operational evidence, residual risk, and fallback conditions.

Background Principles: The Data Plane Runs Tasks, the Control Plane Constrains Capability

A cross-domain system can first be divided into a data plane and a control plane. The data plane receives prompts, documents, images, video, sensor state, and action requests, and runs models, retrievers, tools, and actuators. It produces content, candidate plans, actions, and feedback. The control plane does not perform these tasks for the model. Instead it maintains provenance labels, principal identity, state versions, capability tokens, budgets, traces, and recovery decisions. The key state shared by the two is immutable object identity, version, purpose, authorization, and action ID. Their output consumers are business components, capability gates, reviewers, and the on-call response chain. Security defenses concern unauthorized access, poisoning, privilege escalation, and evidence tampering. Safety controls limit mechanical, transactional, or task harm.

In the normal case, the procurement agent reads purpose-labeled quotes on the data plane and generates a procurement draft. The control plane issues a one-time submission capability, and only for approved suppliers, amounts, and terms. In the boundary case, the model output remains correct, but the identity service or the policy publication root goes wrong. Multiple provenance, state, and capability checks may then let it through together. The architecture must shrink the consequences of common cause through independent execution limits, short-lived tokens, and a recovery root. Conversely, the data content is wrong, yet the control plane still refuses an action that exceeds its authority. The team should record both the upstream compromise and the downstream restriction. It must not cover up the root cause with “no incident.”

Read the seven shared services in the figure from top to bottom. Identify where provenance, identity, state, capability, budget, traces, and recovery each constrain subsequent processing. Then judge which controls still need domain detectors to provide specialized signals.



A seven-layer vertical structure listing the security services shared across language, vision, embodied, and world model systems.

What the figure supports is the architectural judgment of “shared control, retained specialized mechanisms.” Shared services cannot replace video temporal detection, robot safety controllers, or world model external anchors. If multiple services share the same identity, keys, clock, or policy publication root, they do not constitute genuinely independent defense in depth either. Common cause still needs to be exposed in the dependency graph and failure drills.

19.1 From Defense Checklists to a Control Plane

A point defense usually answers a narrow question. Can a classifier recognize jailbreaking? Can a scanner find anomalous weights? Can a watermark still be read after compression? Can collision detection block a given trajectory? A control plane answers a different class of questions. Who can feed data into the model? Who can persist state? Who can let a plan obtain authority? Who can confirm feedback? Who can revoke and recover after a failure? A cross-domain control plane can be divided into seven services:

1. **Provenance and artifact trust:** records where data, code, weights, adapters, policies, and containers come from. It also records the hash, signature, license, and evaluation with which each one enters the environment.
2. **Identity, tenant, and purpose:** preserves the principal, tenant, integrity, sensitivity, and permitted purpose across parsing, retrieval, summarization, generation, and cross-agent forwarding.
3. **State lifecycle:** governs memory, indexes, caches, latent state, maps, and historical trajectories with independent write gates, expiry, versions, revocation, and derivative deletion.
4. **Capability and action authorization:** the model only proposes candidate plans. External policy issues least capability, scoped by principal, object, action, parameters, environmental conditions, and reversibility.
5. **Isolation and resource budgets:** sets default deny plus hard caps on code, network, files, devices, concurrency, tokens, tool steps, and wall clock.
6. **Unified traces and continuous verification:** links the original objective, provenance, model version, state, plan, policy, action, feedback, and cost into one replayable record.
7. **Revocation, rollback, and recovery:** freezes new actions, revokes identities and artifacts, rolls back state, isolates the impact, and rebuilds a trusted environment. It then turns the incident into a regression gate.

The seven services are not a one-to-one match for U0—U6 in Chapter 2. Provenance trust mainly guards U0, but it must also travel with the output to U6. Action authorization mainly sits at U4, yet it depends on the identity and purpose that U1—U3 preserve. Recovery spans every interface. Defense in depth means precisely this: the same risk is constrained by different roots of trust at different locations.

The Object Model of the Control Plane

Shared services must share a set of objects without sharing the model's internal representations. Principals, tenants, data objects, artifact closures, states, capabilities, candidate plans, actions, feedback, trajectories, and recovery points make up the minimal object set. Every object carries an immutable identity, a version, an owner, a

current state, and a policy reference. A specialized system need only map its own concepts into these objects. An LLM's prompts and retrieved segments map to data objects. Video LoRAs and motion modules map to artifacts. Robot maps and world-model latent states map to states, and tool calls and mechanical trajectories map to actions.

Typed dependency edges connect objects. `derived_from` denotes derivation, `loaded_with` runtime composition, `authorized_by` approval, `read_state` and `write_state` state access, `executed_as` principal identity, `observed_by` feedback provenance, and `recovered_from` recovery points. Every edge must carry time, version, and responsibility domain. Recording merely that two objects "are related" is not enough. The control plane can be written as

$$\mathcal{C} = (\mathbf{N}, \mathbf{E}, \mathbf{\Pi}, \mathbf{\Sigma}),$$

where $\mathbf{N}$ is the object set, $\mathbf{E}$ is the dependency and action edges, $\mathbf{\Pi}$ is the policy set, and $\mathbf{\Sigma}$ is the current state of the objects. A request does not call the model directly. Instead, the control plane validates the relevant subgraph. It asks whether data and artifacts are available and whether state is fresh. It checks whether capabilities match and whether actions stay within budget. It also asks whether feedback comes from an independent source. Model inference may fail. The control plane must nonetheless maintain these invariants.

Five Cross-Domain Invariants

The first is **provenance is not lost**: any derived object can be traced back to its inputs, its transformations, and the principal responsible for it. The second is **no privilege escalation**: low-integrity data does not automatically acquire a higher use once it has passed through summarization, OCR, or model interpretation. The third is **no cross-domain state bleed**: state belonging to different tenants, tasks, and trust levels is isolated by default, and copying across domains requires an explicit gate. The fourth is **plans do not self-approve**: the model that proposes a candidate cannot issue the final capability for that candidate. The fifth is **recovery is reachable**: before release, high-impact objects already have a freeze, rollback, or safe-degradation path.

These invariants should be written as executable policies and probes. Provenance is not lost, for example, is checked through derivation-edge completeness rates and random replay. No privilege escalation is checked through label propagation and negative tests on low-integrity inputs. No cross-domain state bleed is checked through cross-tenant read/write probes. Plans do not self-approve is checked through identity and signing-path checks, and recovery is reachable is checked through drills rather than documentary promises.

Service Dependencies and Common-Cause Graphs

The seven services also sit on shared databases, identity, keys, clocks, policy publication, networks, and logs. A faulty identity service can make provenance signatures, state ACLs, and capability tokens fail at the same time. A poisoned policy publication can let the input gate, the action gate, and the alert gate all allow passage together. An architecture review should draw a second dependency graph for the shared services, marking root trust and single points.

Common causes cannot necessarily be eliminated entirely. An enterprise may be able to use only one identity provider or one time service. It should then shrink the blast radius. High-impact capabilities use short-lived secondary confirmation. Recovery keys and production keys are kept in separate domains. The log integrity root is stored independently. Mechanical safety controllers do not depend on cloud identity remaining continuously online. Defense in depth is only as good as the capabilities that remain once a common cause occurs. It is not a matter of how many controls are named.

Separation of the Data Plane and the Control Plane

The data plane transmits content, executes models, runs tools, and drives devices. The control plane owns labels, policies, capabilities, and recovery. The two may share infrastructure, yet they should use different identities and interfaces. A model container can read approved inputs and write candidate outputs. It should not be able to update policies or issue its own tokens. Log collection can read events, but it should not be able to change business state directly. The policy service can reject actions, but it should not touch raw sensitive content it does not need. Separation lets controls be reused across models. A team that swaps out the generator or the planner

rewrites its data format adapters, while identity, capability, and trajectory rules stay valid. If every model upgrade requires reimplementing permissions, the controls are still buried inside the model application. They have not become shared services.

19.2 Provenance and Artifacts: Visible Does Not Mean Trustworthy

A file that can be downloaded from a model repository is merely reachable. Reachability says nothing about provenance, license, format, or behavioral safety. At minimum, a complete artifact record carries the source principal, the committed version, the immutable hash value, the build environment, dependencies, the serialization format, scan results, behavioral tests, the approver, the scope of applicability, and the revocation status. Format safety and behavioral safety must be separated. Avoiding executable deserialization can reduce remote code paths at load time. Yet a LoRA, text encoder, or policy weight in a code-free format may still carry trigger behavior. A model file may be both parameters and an entry point into the software supply chain. Public cases of malicious model artifacts and of version-specific deserialization vulnerabilities show exactly this [cohen2024malicioushf, @github2025pytorch]. For safe loading, the minimum set is provenance verification, a code-free format, an isolated environment, minimal network access, and a behavioral gate.

Visual generation in particular depends on composable components. Consider the base model, VAE, text encoder, control module, LoRA, motion module, and custom nodes. When these come from different principals, passing individually does not mean passing in combination. Compositional behavior differentials should cover trigger words, reference images, identity, motion, and benign task utility. Tool servers, prompt templates, and retrieval encoders in LLM systems need the same compositional view. So do perception modules and action decoders in VLAs, and the dynamics and reward components of world models. Provenance records must also support revocation. A data batch may lose authorization, an adapter may be judged malicious, or a dependency may carry a high-severity vulnerability. In any of those cases the team must know which models, indexes, caches, images, and released outputs derive from it. Without a derivation graph, so-called revocation stalls at deleting the entry file.

The Four Gates of Artifact Admission

The first is the **identity gate**. The repository confirms the publishing principal, immutable digest, signature, license, and commit and build provenance. A signature proves who submitted what. It does not prove that the behavior is benign. The second is the **format gate**. Parameters preferentially use restricted tensor formats, with code separated from data. When custom code is truly required, that code enters independent software review and isolation. The third is the **static gate**. It checks dependencies, entry functions, secret patterns, network and file permissions, model cards, and known vulnerabilities. The fourth is the **behavioral gate**. Under a frozen closure, tasks, seeds, and budget, it checks expected functionality, clean utility, confidentiality probes, and resources.

The four gates output a scoped admission record, not a permanent "safe" label. That record states which base models, components, tasks, tenants, permissions, and time limits are allowed. A text encoder may be approved for offline creative experiments but not for connection to a production release. A robot policy may be approved for simulation but not for real actuators. The capability domain is part of the admission record. Malicious model artifacts and deserialization vulnerabilities respectively show that format and behavior need independent checks [cohen2024malicioushf, @github2025pytorch]. The format gate blocks executable payloads but will not discover conditional anomalies in parameters. The behavioral gate observes outputs, but it should not let code-bearing files skip isolation. The two gates complement each other.

The Compositional Behavioral Gate

An admission conclusion must bind precisely to the components, routing configuration, and runtime permissions actually loaded. The object actually run on a given occasion can be denoted as the following artifact closure:

$$\Omega = (A_1, A_2, \dots, A_n, C, R),$$

where A_i is the components, C is the loading, routing, scaling, quantization, and dependency configuration, and R is the runtime permissions. Any change to a component or to a key configuration produces a new closure identity.

Individual admission only describes how a component behaves in the tested environment. Only the compositional gate describes the actual workflow. Compositional testing includes at least a baseline, single components, the target combination, and adjacent combinations. The baseline exercises normal tasks. Single components help with localization. The target combination corresponds to production. Adjacent combinations vary the order, scaling, quantization, or one plugin, so that nonlinearity and compatibility boundaries become visible. Small parameter artifacts need to be tested in combination even if they do not change the base model; research on compositional backdoors in visual LoRAs shows this [VIS_P006]. The same principle applies to LLM tool servers, VLA perception-policy modules, and world-model dynamics-reward combinations.

The behavioral gate reports target anomalies and benign utility separately. On the LLM side it observes task completion, unauthorized plans, and tool side effects. On the visual side it observes semantics, identity, image quality, motion, and realism. On the VLA side it observes tasks, trajectories, collisions, and action constraints. On the world-model side it observes prediction, planning, state freshness, and safe fallback. Different metrics keep their respective denominators, and a closure need only satisfy the admission rules of its own scenario.

Promotion, Freezing, and Revocation

Artifacts acquire capabilities progressively, from isolation, experimentation, pre-release, and canary to production. Each level records evidence, validity period, allowed tasks, and fallback objects. A production alias must resolve to an immutable closure. A mutable repository tag must not be able to bypass promotion. A key change in a model, policy, reviewer, dependency, or deployment permission automatically sends the system back to a lower level.

Revocation propagates along the dependency graph, but it does not blindly assume that all successors are affected. It first freezes new actions that include the target artifact, and queries derived images, quantized versions, adapters, indexes, and released outputs. It then verifies by consequence and reachability. Artifact blocking addresses future loading. State cleanup addresses persistent effects. Content and device remediation address consequences that have already occurred.

The Minimal Receipt for Supply-Chain Evidence

At minimum, each admission receipt carries the closure digest, the provenance of input data and artifacts, the runtime environment, the test protocol, benign utility, anomaly observations, uncovered scope, approver, validity period, and fallback objects. Machines can read the receipt, and the deployment pipeline verifies scope matching item by item. PDF reports are for human reading, and structured fields still bound their approval scope. Suppose a supplier does not disclose weights or training data. Customer-side verification then covers the service version identifier directly, along with the data-use statement, tenant isolation, the anomaly notification time limit, the withdrawal mechanism, and negative probe receipts. Unknown internal state corresponds to a smaller capability and data scope and remains in a to-be-verified state.

19.3 Identity and Purpose Must Flow with the Data

Many attacks rely on a kind of information loss. Web pages, PDFs, images, audio, and tool returns all pass through OCR, summarization, or serialization. What remains is a plain text string. Provenance, tenant, integrity, and purpose are dropped. The model can no longer distinguish an authenticated task objective, trusted policy fields, and external data. A security architecture should encapsulate content as labeled objects rather than reduce it to a string as early as possible:

$$d = (\textit{payload}, \textit{source}, \textit{tenant}, \textit{integrity}, \textit{purpose}, \textit{sensitivity}, \textit{expiry}).$$

`payload` is the content, while the other fields are produced by trusted components and propagated through transformations. A summary inherits the provenance set. An OCR result retains the image region and the parser version. A retrieved chunk retains the document ACL and the tenant. A tool return retains the calling principal and the tool identity, and a cross-agent message retains the original user objective and the delegation chain. The model can read these fields to inform its judgment. It must not itself promote a low-integrity object into policy or expand its purpose.

Purpose restrictions are finer-grained than "whether it is secret." An internal equipment manual may be used to answer repair steps, yet is not necessarily allowed to enter external training. A portrait reference image may be used for this edit, yet is not necessarily allowed to create a long-term identity embedding. Robot camera footage may be used for obstacle avoidance, yet is not necessarily allowed to be uploaded to a public model service. Every egress point must ask "is the current purpose authorized?"

Label Model and Propagation Rules

A practical label keeps tenant, principal, integrity, sensitivity, purpose, region, expiry, and the content's provenance and processing history attached to the object. It can be written as

$$L = (\textit{tenant}, \textit{owner}, \textit{integrity}, \textit{sensitivity}, \textit{purpose}, \textit{region}, \textit{expiry}, \textit{provenance}).$$

`integrity` marks whether content can influence policy. `sensitivity` marks the consequence of exposure. `purpose` lists the tasks a record is permitted to serve, and `provenance` points to where the content came from and how it was processed. Trusted parsers, identity services and policy services write the labels. The model may read them. It must not raise integrity itself, expand purpose, or extend expiry. Multi-input transformations propagate each field conservatively. They must not hide the different partial orders of those fields behind one uniform "union" operation. Where the output is $y = f(x_1, \dots, x_k)$, the rules are as follows:

Field	Multi-input output rule
<code>tenant</code>	Inherit when the input tenants are the same. When they differ, deny by default, and allow the transformation only when one explicit cross-tenant authorization binds the input set, the output tenant, and the purpose at the same time
<code>owner</code>	Retain the set of all input owners and the principal responsible for the transformation; do not let the last writer overwrite them
<code>integrity</code>	The output must not be higher than the lowest trust level among all inputs. If integrity is not a single ordered level, retain the attributes and provenance of each input, and change it only after a trusted verifier issues a receipt under an explicit protocol
<code>sensitivity</code>	Inherit the stricter sensitivity among the inputs; any desensitization requires an independent protocol and receipt
<code>purpose</code> , <code>region</code>	Take the intersection of the permitted purposes of the inputs and the intersection of their permitted regions; if either intersection is empty, reject the output
<code>expiry</code>	Take the earliest expiry among all inputs, unless an explicit authorization covering all inputs is obtained anew
<code>provenance</code>	Take the union of the provenance and processing history of all inputs, then append this transformation, its version and time, and the responsible principal

Summarization, translation, OCR, transcription and vectorization shrink the amount of information. That alone does not widen their permissions. A lower-sensitivity label may be produced only when a trusted "desensitizer" demonstrates, under an explicit protocol, that the necessary attributes have been removed. The desensitization receipt must be retained.

Decryption Broker and Minimal Plaintext Surface

The model service should not hold the plaintext key of an encrypted object for long periods. The decryption broker decides whether to release it, weighing principal, task, object, purpose, runtime environment and expiry. Preferably it sends plaintext only to a one-time restricted process. The model obtains the fragment needed to complete the task. It must not write decryption results into unlabeled logs, caches, or long-term memory.

Decryption and capability authorization are separate. Reading a procurement file only obtains data access capability. Initiating a procurement also requires an independent action permission. Viewing robot footage and controlling the robot likewise use two different capabilities. A workflow that needs both reading and action obtains the data capability and the action capability separately. The action approval then binds a normalized object derived from the authorized data. That prevents the model from substituting key fields after decryption. Keys and plaintext paths need two-way probes. Allow probes confirm that an authorized task can read the necessary fields. Deny probes confirm that a wrong tenant, an expired purpose, an unapproved export, cache replay and log bypass are unavailable. A test that only asks "can decrypt" would mask unauthorized access. A test that only checks "deny unauthorized access" might make all legitimate tasks fail.

Desensitization, Downgrading, and Declassification

Desensitization is a data transformation. Downgrading is a contraction of system capability. The first acts on the amount of information; the second acts on the executable scope. Desensitization can turn raw footage of people into statistical features, a complete procurement file into approved fields, and a precise location into a region. It must also verify whether the remaining information can still identify a principal or recover a secret. Downgrading applies when labels are missing, state is stale, or risk rises. It permits only read-only, low-resolution, simulated, draft, or human-confirmed operation.

Declassification must be approved by trusted policy. A generative model may not assert that "this content is not sensitive." The approval records the input object, algorithm and version, output fields, attack assumptions, review, and validity period. After new linked data appears, a previously aggregated result may become re-identifiable. Desensitization rules therefore have versions and re-evaluation conditions as well.

Label Conflicts and Unknown States

Labels from different systems may conflict. A vendor calls the content public while the enterprise ACL marks it internal. A user declares ownership of a portrait while the identity service cannot verify it. An old cache marks it usable while the new policy has already tightened. The control plane applies the stricter constraint and records the conflict. It must not let the last-arriving field overwrite silently. Unknown states correspond to conservative permissions. When a label is missing, high-impact paths enter read-only, isolation, or human confirmation. Low-risk tasks may continue within a restricted scope. The system therefore remains usable, and a parsing failure also obtains only least privilege. Unknown items enter traces and metrics, which helps teams repair data rules.

19.4 State Writes Are an Independent High-Risk Action

Long-term memory, RAG indexes, caches, user profiles, maps, latent world states and control histories all count as state. Each turns a single input into a condition for future decisions. A write should therefore not be merely an incidental step of the model generation pipeline. State-write rules answer at least the following questions. Who writes, and as whose identity? How much provenance integrity is there? What is the retention purpose? Which future tasks does it trigger? Can it cross tenants? How long is it kept? How is it rolled back? Is deletion synchronized to indexes, summaries, backups and derived representations? AgentPoison demonstrated an attack path that directly poisons long-term memory or a knowledge base [chen2024agentpoison]. MINJA studied how more constrained benign query interactions affect memory [dong2025minja]. The two have different access capabilities. Together, though, they show that the write boundary needs to be established when state enters. The ACL at future read time is only responsible for a separate permission check.

World-model state also needs freshness and external anchors. Over long rollouts, a latent state may drift away from reality. Feedback contamination can also make errors self-consistent. The system should store the state version, observation time, sensing provenance, uncertainty and resynchronization conditions. If a high-impact plan depends on stale or unverifiable state, the safety controller should downgrade, request new observations, or stop. It should not let the model fill in unknown reality with longer imagination.

State Namespaces and the Write Gate

State is namespaced at least by tenant, principal, task, session, artifact closure, purpose and trust level. A session cache must not escalate into long-term memory on its own. Personal preferences must not enter an organization's public index. Simulation state must not overwrite real device state, and low-integrity web content must not be written into security policy. Cross-namespace copying is an explicit action, and it requires provenance, target, fields, expiry and approval.

The write gate validates the candidate state first. It then decides to accept, isolate, partially accept, or reject. Structured fields are validated item by item on name, data type, required fields, value range, provenance principal and validity period. Facts need provenance and time. Preferences need user confirmation, and high-risk identity and action rules are written only by authorized services. A memory proposed by the model is only a candidate. The state service retains the final write authority. AgentPoison and MINJA show memory-poisoning paths under different

access capabilities [chen2024agentpoison, @dong2025minja]. Controls must therefore cover both direct index writes and seemingly benign interaction accumulation. Protecting only the management API is not enough to protect state semantics.

Derived Lineage and Deletion Propagation

Original state may be summarized, embedded, clustered, merged, cached and backed up. Each derivative records its parent state, transformation, version, time and purpose. A deletion request follows the lineage to find online indexes, vectors, caches, summaries, snapshots and training queues. For backups that cannot be cleaned up immediately, it sets isolation and expiry. A deletion receipt must distinguish these states: "entry not readable," "online derivatives purged," "backup awaiting expiry," "released model requires risk handling," and "irreversible removal not possible." These layered states directly present the scope that is completed, pending, and impossible to undo. They support disposition decisions better than a single green checkmark. After rebuilding an index, also verify that a deleted object does not reappear from an old snapshot.

Derived lineage also supports incident radius. When a poisoned memory is found, the team queries the retrieval results, plans, actions and users it affected. When a batch of state is found to be stale, the world model switches back to the most recent trusted snapshot and re-observes. Without lineage, the only options are to wipe everything or keep taking the risk.

Snapshots, Deltas, and State Freshness

The state service stores immutable snapshots and ordered deltas. Snapshots provide a recovery baseline, and deltas explain how state changes. Each update binds the principal, provenance, model, policy and action ID. Concurrent writes use version conditions, which stop a stale request from overwriting newer state. Conflicts go to a merge policy or to human handling, never to the last writer.

Freshness is not a uniform duration. User preferences may be kept for months. Inventory may last minutes. The location of people near a robot may be trustworthy for only a few hundred milliseconds. State objects carry observation time, maximum age and resynchronization conditions, and plans re-check before commit. A world model can predict the future. A prediction does not automatically extend the validity of real observations.

Acceptance of State Isolation

Acceptance tests reading, writing, derivation, deletion and recovery together. Allow probes demonstrate that a legitimate session in the same tenant can store and recall. Deny probes verify that cross-tenant, wrong-purpose, expired-principal and low-integrity-provenance data cannot be written or read. Derivation probes confirm that summaries and embeddings inherit labels. Deletion probes confirm that no reflow occurs either online or after a rebuild. Recovery probes confirm that trusted snapshots contain no quarantined contamination.

When isolation fails, the system first freezes new state writes and high-impact actions. Read-only or stateless modes stay available. A full shutdown is not the only option. Question answering can fall back to no memory. A procurement agent can require explicit input each time. A robot can switch to local safety controls, and a world model can shorten its planning horizon and request new observations.

19.5 A Plan Cannot Issue Permissions for Itself

Capability systems rest on one core principle. They separate "what is recommended" from "what is permitted". Model output enters a read-only plan object, and the authorization service computes the minimal allowed set from identity, task, data flow, parameters and consequences. High-impact actions use dry-run-only, preview and two-phase commit. Approval binds normalized parameters, tool version, environmental conditions and expiration time, and any key change invalidates the approval. The authorization decision can be written as:

$$allow = P(\textit{subject}, \textit{task}, \textit{object}, \textit{action}, \textit{args}, \textit{context}, \textit{provenance}, \textit{risk}).$$

The model may supply `task` and candidate `args`. It does not control the policy P , the principal identity, provenance labels, or risk thresholds. Policy output is not a one-off Boolean value either. It can return deny, read-

only, human confirmation required, second sensor required, narrowed parameter range, or the issuance of a short-lived capability. Tool protocols can standardize identity and authorization flows. MCP security practice emphasizes token audience validation, and it warns against token passthrough. It also calls for per-client consent and for least-privilege authorization scope (that is, the set of resources and actions a token or principal is permitted to operate on) [mcp2026security]. These requirements constitute a capability foundation. Business policy is still needed, however, to judge whether a given email, payment, code execution or mechanical action serves the user's goal. Protocol-layer authorization and semantic-layer task alignment complement each other.

Embodied systems must also connect the action gate to a safety envelope. The planner may select candidate trajectories, while an independent safety controller checks speed, force, distance, no-go zones, personnel location and stop conditions. World-model predictions can help estimate risk, but they cannot be the sole safety evidence. If observations are missing, or model uncertainty is too high, the action set should contract to a provably safe region.

Parameterized Capability Objects

A tool allowlist answers only "can this name be called". A parameterized capability answers more. It fixes whose identity the call uses, which object it touches, and the parameters and environment it runs with. It also states how long it lasts, how many times it may run, and whether it can be further delegated. A single capability can be written as

$$cap = (sub, obj, act, argset, ctx, budget, expiry, depth, nonce).$$

`argset` is the allowed parameter set, and `ctx` is the environmental precondition. `budget` is the count, amount, resources, or motion range. `depth` is the maximum delegation depth, and `nonce` prevents replay. An external service signs a capability, and the model may hold only the short-lived object needed to complete the current task. Parameters are first parsed into normalized fields, and each field has an explicit name, type, value range, required conditions and comparison rules. An email capability binds the sender identity, recipient domain, attachment digest and template. A procurement capability binds the supplier, currency, amount ceiling, receiving principal and purchase terms. A robotic arm capability binds the device, button, speed, force, space and valid observations. Any field outside the set is re-reviewed.

Delegation Chains and Identity Attenuation

An agent may call a sub-agent or a tool. Delegation, however, cannot expand the original permissions. A child capability is a subset of the parent capability, with a shorter term and shallower depth. It preserves the original user goal and the delegation chain. A tool return cannot carry a new capability token. A natural-language "already approved" in a cross-agent message does not constitute signed authorization. An overly long delegation chain increases semantic drift and blurs accountability. The control plane limits depth, concurrency and the total number of actions, and re-verifies the goal and parameters at every high-impact boundary. Once any principal in the chain is revoked, all child capabilities become invalid. Revoking only the current process is not enough to stop short-lived tokens that have already been distributed.

Two-Phase Commit and Transaction Semantics

A high-impact action is divided into `prepare` and `commit`. The prepare phase parses the normalized object, verifies provenance and state, performs a dry run, computes the diff and reserves resources. The commit phase re-checks the version and preconditions, then executes using a one-time capability. Approval binds the digest of the preparation result. That prevents the model from substituting the amount, recipient or trajectory after approval.

Commit has to survive timeouts, duplicates and partial failures. `idempotency_key` keeps a retry from duplicating a payment or creating multiple purchase orders. The transaction record distinguishes not-started, prepared, committed, confirmed, compensating and unknown. An unknown state does not automatically retry high-impact actions. It queries external facts instead, or escalates to a human. A reversible business operation can define compensating actions, such as withdrawing a draft, canceling an unsettled order or restoring a file version. Irreversible or physical actions cannot be restored by compensation. A mechanical action relies on a real-time safety controller before commit, and continuous monitoring watches it during execution. The moment a person

enters, or observations become invalid, it switches to safe mode. A transactional database cannot guarantee rollback in the physical world.

Preventing State Changes After the Check

State changes may occur between preparation and commit. At t_0 the prepare phase reads the supplier account, file version, robot path and policy version. At t_1 the normalized digest is handed to approval. At t_2 an external system may update the account, occupy the path or revoke authorization. Only at t_3 , at commit, does execution actually take place. The capability object must bind the version and preconditions from t_0 , then re-read and compare them at t_3 . If any key field, policy or environment differs, the preparation result and the old approval are invalidated together, and the system returns to re-preparation rather than reusing the old digest. Human approval follows the same principle. The approval interface displays the normalized diff and provenance, and the approval signature binds the digest. Explanatory text the model generates afterward cannot change the approved object. This prevents post-check substitution, and it also lets the trace explain what the human actually approved.

19.6 Isolation and Budgets Must Be Accepted with Probes

"Sandbox", "no public network", "read-only" and "intranet only" are all configuration intentions, not runtime results. Acceptance must cover both allowed and forbidden paths. Coverage spans in-scope reads and writes, out-of-bounds files, process creation, DNS, redirection, loopback, private networks, link-local and cloud metadata, shared volumes, devices, host mounts, resource exhaustion, and residue after destruction. Secrets do not enter the model context. A secret broker is a controlled service that, after authorization, temporarily issues secrets or credentials by audience, purpose and term. Once an action is approved, it hands out credentials that are audience-bound, scoped to least privilege and short-lived. Code execution and network access are authorized separately, so that a "run analysis" tool does not naturally obtain arbitrary egress. Containers and microVMs limit the execution boundary. A microVM is a lightweight virtual machine that provides an independent guest kernel for a single workload. The image, kernel, agent, debugging interfaces and management plane still need to be inspected.

Budget is also a permission. Maximum input and output tokens, context length, tool steps, recursion depth, concurrency, GPU time, storage and external API cost should all have hard ceilings and termination conditions. A budget that is too low harms benign utility. An unlimited budget turns a single injection into a long-running search and an economic attack. Safety gates should report block rate, task success, latency and cost together.

Choosing Isolation Strength by Consequence

Isolation is not a binary property. It divides into application logic, processes, containers, microVMs, separate hosts and physical security domains. Application logic relies on the same runtime and suits low-risk data partitioning. Processes and containers provide resource and namespace boundaries but still share the kernel. microVMs provide virtual-machine-level kernel boundaries at a small startup cost, with Firecracker as a representative implementation [agache2020firecracker]. Separate hosts and physical domains are used for high-consequence work or device control.

The basis for choice includes untrusted code, secret value, network reachability, host sharing, escape consequences, task duration and recovery requirements. Tool-free classification of public text can use lighter isolation. Running third-party model plugins requires restricted network access and stronger boundaries. A task that can access production credentials or mechanical equipment must separate code execution from capability issuance, so that even if the runtime environment is breached there are no long-lived secrets.

Isolation Level	Shared Surface	Typical Use	Controls That Must Be Added
Application logic	Process and runtime	Low-risk tenant routing	Strict object ACLs, input and output field definitions
Separate process	Kernel and some files	Controlled parsers	Account, file, process, and resource limits
Container	Host kernel and management plane	Model inference, ordinary code tools	Image, system-call filtering mechanism (seccomp), network, volume, and device controls
microVM	Hardware and host management plane	Untrusted plugins, one-off execution	Image supply chain, egress proxy, destruction verification
Separate host / physical domain	Minimal sharing	High-value secrets, device safety core	Operations identity, media, network, and human procedures

Bidirectional probes rather than testing refusal only

Forbidden-path probes include out-of-bounds files, host sockets, secret directories, DNS, HTTP redirects, loopback, private networks, link-local addresses, cloud metadata, proxy environment variables, shared volumes, devices and process creation. Allowed-path probes, by contrast, verify the inputs a task requires, temporary files, approved domains, GPUs, restricted outputs and normal exit. Having only prohibited probes encourages "locking everything down". Having only allowed probes fails to reveal escapes.

Network tests must cover domain resolution and actual connections. Domain allowlists may be redirected, and HTTP clients may read system proxies. DNS itself can also leak information outward. Egress proxy verification checks the final address, protocol, port, audience and response size. An approved connection does not mean the returned content is trustworthy, and tool results still carry provenance labels.

Destruction, residue, and bypass

After a one-shot environment ends, verify that processes, child processes, mounts, temporary volumes, GPU memory, caches, tokens and network sessions have been cleaned up. Deleting the container object alone does not prove that shared storage and external queues leave no residue. Destruction receipts are bound to the task ID, and the control plane checks resource release. The management plane and debug interfaces often become bypasses. Operational shells, image builds, log exports and emergency accounts also require identity, least privilege and traces. Even if the production data plane is strongly isolated, an administrator reading all secrets without records still forms an obvious bypass. Destruction acceptance must therefore also trace back from the data plane to the control plane, the management plane and external persistent resources.

Linked termination of budgets

Token, tool-step, wall-clock, GPU, storage, network and cost budgets should be linked. A model reducing its text output inside a tool loop does not mean the total cost is low. Long video generation may occupy GPUs and queues while producing only a small number of tokens. The control plane aggregates full-chain cost. It degrades when a soft threshold is reached, and aborts and cleans up when a hard threshold is reached.

Budget alerts must distinguish complex benign tasks, runaway loops and targeted consumption. Benchmarks such as AgentDojo observe benign tasks and attack objectives at the same time. That shows agent safety evaluation needs to retain the utility dimension [debenedetti2024agentdojo]. Production systems report normal completion, blocking, latency and cost together. The availability loss caused by blanket refusal shows up directly in the metrics.

19.7 Unified traces: letting the system explain a single action

The four model classes have different log formats, but a single high-impact action needs to share the same action ID. A trace should at least link the original user goal, input provenance and parsing version, retrieval and state reads and writes, model and prompt/configuration versions, and candidate plans. It should also link the field names, types, units and value ranges of tool parameters, normalized parameter values, policy decisions, human approval, credential issuance, network and file events, executor feedback, resource budgets, and the final result.

Logging does not mean storing all context forever. Sensitive values are encrypted separately from the main trace, and the body stores only references, hash values and necessary metadata. Access is authorized by role and by

case, and retention periods are consistent with legal requirements. The logs themselves use append-only integrity protection, and time synchronization and action IDs guarantee cross-system correlation. True observability also requires alert escalation. The system must detect an anomalous domain, low-integrity data affecting a high-risk parameter, a cross-tenant hit, the addition, deletion or type change of a tool parameter field, a secret pattern, a retry loop or a state mutation. Once it does, it must know who is on call, when to freeze and how to degrade. Alert scores must be connected to on-call routing, freeze authority and degradation actions to form operational control.

Unified trace field structure

A trace can use two layers: a common envelope and domain events. The common envelope contains `trace_id`, `span_id`, `parent_id`, the action ID, time, principal, tenant, object, version, policy decision, budget and integrity. Domain events store specialized fields such as prompt parsing, image derivation, robot trajectories or world state. The common envelope supports cross-system correlation, and domain fields retain the necessary technical detail.

Event types at least include `ingest`, `derive`, `read_state`, `write_state`, `model_call`, `plan`, `authorize`, `issue_capability`, `execute`, `observe`, `publish`, `revoke` and `recover`. Each event records input and output object references, state, errors and evidence limits. Asynchronous queues are joined through causal links rather than time alone. The trace field structure requires version management. When tool parameters, model interfaces and policy fields change, the parser preserves the meaning of old events. Unknown fields are explicitly marked and enter a handling queue. When the minimum fields of a high-impact action are missing, the action gate refuses or degrades it. The action enters execution only after the fields are completed and re-validated.

Trace integrity and clocks

Append-only logs, batch hashes or signatures help detect event deletion and reordering. The signature-binding idea of content credentials also shows that "who declared what" and "whether the content was altered" should be kept separate [VIS_0001]. Log integrity does not prove that the model judgments in the events are correct, but it supports accountability and replay. Cross-system clocks drift. A trace stores monotonic timing, wall-clock time and queue order at the same time. High-impact transactions use action IDs and state versions to establish causality rather than ordering by milliseconds alone. When a device is offline, it first buffers signed events locally, then merges them after reconnection and marks the time as uncertain. Replay personnel use this to recover the causal partial order and do not mistake clock precedence for business causality.

Balancing privacy minimization and investigability

The main trace stores object references, summaries, labels and decisions. By default it does not store full prompts, images, audio, secrets or biometric features. Sensitive payloads that genuinely need investigation go into a separate encrypted evidence store, authorized by case, role and time limit; ordinary operational panels see only the necessary fields. Tokens and keys never enter searchable logs. Summaries can also leak low-entropy secrets, and perceptual fingerprints can be used to link individuals. Field design undergoes privacy review, using keyed digests, per-domain salts or controlled matching services where necessary. Deletion requests propagate along the object lineage and evidence retention rules, while handling records that contain no content are preserved.

From traces to alerts and regression

Alert rules center on cross-event relationships: low-integrity input affecting a high-risk parameter, state coming from the wrong tenant, a capability object not matching the approval summary, and missing executor feedback. They also cover a budget repeatedly approaching its ceiling and old state reappearing after recovery. A single log field is normal, while the combined relationship may still be anomalous. After an event ends, turn the confirmed relationships into regression probes. Probes reference minimal synthetic data and frozen versions and do not copy sensitive originals. They are rerun on every model, tool, policy or infrastructure upgrade. Traces thereby connect investigation and release gates. An alert's closure receipt also becomes the expected state for the next regression.

19.8 Recovery is part of the architecture

The recovery process starts with "stop getting worse." The first phase freezes new actions, revokes short-term capabilities, blocks egress, isolates affected principals and preserves forensic snapshots. The second phase rotates credentials, rolls back state, cleans up memory and index derivatives, and revokes artifacts and releases. The third phase rebuilds the environment from a trusted baseline and replays the scope of impact. The fourth phase turns entry points, detection signals and failure controls into automated regression gates.

Different domains have different recovery objects. LLM applications must clean up memory, RAG and agent delegation. Vision platforms must revoke model components, content credentials and published content. VLAs must shut down safely, preserve device state and recalibrate. World models must resynchronize state, clear contaminated trajectories and switch to a safe controller. A common control plane provides identity revocation, state versions, artifact graphs and unified traces, while dedicated teams own domain safety conditions.

RTO and RPO must be defined per service

Recovery time objective (RTO) is the longest target time to recover to a specified safe service level after an interruption. Recovery point objective (RPO) is the window of state that may acceptably be lost or rolled back. The two are not a single system-wide number. A question-answering service can quickly fall back to a memoryless mode. A robot must first shut down safely before recovering. A world model may need to re-observe, and a vision platform must also handle content that has already been published.

For example, the read-only retrieval RTO of a procurement agent can be shorter than that of creating a purchase order, and the state RPO can roll back to the most recently confirmed preferences. The safe-controller RTO of a warehouse robot is nearly instantaneous, while advanced autonomous planning can recover later. A video platform's generation service can be paused, while victim reporting and propagation freezing must remain available. Business continuity prioritizes restoring safe degraded capabilities, and full automation is opened in stages after more complete validation.

Recovery tiers and trusted baselines

Recovery can be divided into five tiers: freeze, isolate, roll back, rebuild and restart. Freeze stops new high-impact actions. Isolation cuts off principals, artifacts or state. Rollback selects a trusted snapshot and closure. Rebuild restores identity, policy, data and dependencies in a clean environment. Restart reacquires capabilities through a canary rollout. Each tier has entry, completion and failure conditions. A trusted baseline includes immutable images, artifact closures, policies, identity roots, state snapshots, recovery keys and validation scripts. The baseline is stored in a separate domain from production and is validated for availability through periodic startup. The rebuild process itself enters the unified trace, and emergency accounts and manual steps also leave tracking records.

Scenario matrix for recovery drills

The drills cover at least identity leakage, malicious artifacts, state contamination, mistaken capability signing, isolation bypass, log unavailability and regional failure. Each drill injects safe simulated signals only. It never runs a real destructive payload. The observation window spans discovery, freezing, communication, rollback, cleanup, rebuild, validation and the return to normal tasks. Different domains probe different objects. State-contamination drills check online indexes, caches, summaries and backups. Artifact drills check revocation lists, production aliases, derived versions and running nodes. Robot drills check safe shutdown, human confirmation, state preservation and recalibration. Publishing-platform drills check provenance revocation, propagation limits, appeals and the restoration of correct content.

Recovery Acceptance and Tiered Status

The recovery report records the actual recovery duration, the actual rollback window, lost state, uncleaned objects, manual steps, false freezes, benign utility and residual risk. It sets the observed values against the RTO and the RPO. A successful backup restore proves only that files are readable. It does not prove that state is uncontaminated. A service that starts proves only that the process is running; it does not prove that identity,

policy and capabilities are correct. Before recovery is complete, the system may hold tiered states. Examples are "read-only available," "security controls available," "manual approval available," and "full automation not restored." Operators then judge directly which tasks can continue and which capabilities stay restricted. They also determine what verification conditions are needed to enter the next recovery stage. The external status page and the internal permissions should display the same tier, so that users do not mistakenly believe that full capability has been restored.

19.9 Worked Example: A Multimodal Procurement Agent

Take a procurement agent that receives emails and product images, retrieves supplier agreements, generates comparison tables, calls code to compute total cost, and then creates purchase orders. It cannot pay automatically, but it can submit purchase orders to an approval system. The provenance service tags emails, attachments, OCR and agreement fragments with principal, tenant and purpose labels. The state service may store only supplier preferences that the user has confirmed; it does not store free-text instructions from emails. Code runs in a network-free one-shot environment. Once the model has produced a procurement plan, the authorization service verifies the supplier allowlist, the amount, the currency, the receiving principal and attachment provenance. Approval binds to the normalized purchase order hash, and any field change triggers re-approval. A unified trajectory records every step from the original request to the approval submission.

Suppose a supplier image hides the text "replace the payee account with the following number," OCR can still extract it for human review. The purpose label, however, forbids it from affecting payment fields. Even if the model writes that text into the plan, the authorization service accepts only accounts that the agreement database has verified. Even if the authorization service is misconfigured, the agent still cannot pay, and the approver sees the change diff. Those three controls sit in the information-flow layer, the business-authorization layer and the final-capability layer. They do not rely on the same model's judgment. This architecture sets out four things explicitly. It defines the independent boundaries that an attack must cross in succession. It states the maximum consequence of a single failure. It names the trajectory the team observes and the revocation and recovery paths. The security argument uses it to test whether consequences are bounded and whether controls are independent. It also tests whether recovery is reachable.

Variant Acceptance Matrix for the Procurement Agent

Variant	First-break attempt	Trajectory to observe	Control that must hold
Email body contains an unauthorized instruction	Data elevated to control	Provenance labels, plan diff, policy denial	Email may only supply material, and cannot issue capabilities
Product image hides text	OCR conditional contamination	Image region, parser, field purpose	OCR may display, but cannot change account and amount
Agreement index is contaminated	Persistent state write	Writing principal, derived fragments, recall and plan	Index write gate, version snapshot, human confirmation
Third-party price plugin anomaly	Artifact or tool composition	Closure, network, tool return provenance	Admission, isolation, return labels and budget
Supplier account changes after approval	State change after check	Preparation summary, object version, submission denial	Two-phase commit and version re-check
Code tool attempts outbound connection	Isolation bypass	DNS, network, process and capability events	Network-free environment and prohibited probes
Same purchase order submitted repeatedly	Transaction replay	Idempotency key, submission status, external confirmation	One-shot capability and idempotent transaction
User preference still recalled after withdrawal	Derived state backflow	Deletion lineage, index rebuild and cache	State cleanup receipt and recovery probe

Every variant runs alongside a benign control. Legitimate emails can still enter the comparison table. Normal image OCR can still be shown for human verification. Trusted plugins can still return prices, and purchase orders can still be submitted after approval. The acceptance report gives the attack or anomaly, the benign task, latency,

manual burden and the highest consequence layer as separate entries. That prevents a "pass" obtained by turning off all functionality.

Variation 1: Trusted information and untrusted control coexist. Text in supplier emails and images may contain genuine product information. It may also contain unauthorized instructions. The system handles fields separately as specification, model, price candidate and free text. Low-integrity fields can enter the evidence view. Fields confirmed through the agreement database or by a human enter the procurement object. The downgrade preserves business value while it blocks data elevation.

Variation 2: State contamination surfaces in a later task. In this round the attack content does not change the purchase order. It is still stored as a supplier preference and recalled automatically the next time. The test must therefore span sessions. The first round observes the write candidate, the second round observes recall and the plan, and after withdrawal the third round verifies the index and the cache. Testing single-round output alone would miss persistence.

Variation 3: The tool is correct but the transaction object changes. The code tool can compute the total price correctly, and approval can complete correctly as well. Even so, the supplier account or the currency may still change before submission. The preparation object contains all key fields and their provenance summary. At submission the system re-reads the trusted system and compares versions. A difference invalidates the approval, and the human sees field-level changes rather than a model-generated summary.

Variation 4: Minimum business under recovery mode. If the memory service is isolated, the agent falls back to a preference-free mode. If the price plugin is revoked, the system uses manually uploaded quotes. If code execution is disabled, it generates the comparison table but does not compute automatically. If the approval service is unavailable, it saves only a draft. Each downgrade explicitly states the missing capability, and unknown fields remain empty and enter human confirmation. Read-only service, draft service, manual submission and fully automatic service each have a preset RTO. The drill then records the actual recovery duration of each tier and judges whether the target is met. **How to write the acceptance conclusion.** The security argument should state clearly which layer each control blocks at and how benign tasks perform. It should also state which paths remain to be covered. All of this holds for the frozen version, the specified tenant and the given budget, under the eight variation categories. After a change in version, permissions or external systems, the relevant units are re-accepted.

From the Worked Example Back to the Chapter Checklist

A cross-domain architecture review starts from the complete path of one high-impact action. The checklist below applies to the procurement agent, and also to content publishing, robot operation and world-model planning. Every "yes" should be backed by an object, a policy, a probe or a drill record. The boxes in the architecture diagram serve to locate the corresponding component.

- [] **Is the control-plane object closed?** Principals, tenants, data, artifact closure, state, capabilities, plans, actions, feedback, trajectories and recovery points all have immutable identity, version, owner and current state.
- [] **Are dependencies and common causes explicit?** Edges such as `derived_from`, `authorized_by` and `executed_as` can be traced back. Shared dependencies such as identity, keys, clocks, policy distribution and logs already have a failure-radius analysis.
- [] **Is the actual artifact closure admitted?** Deployment records lock down all components, configurations and permissions. The identity gate, format gate, static gate, behavior gate and adjacent-combination tests all have scope and validity period.
- [] **Are information-flow labels faithful?** OCR, summaries, retrieval, cross-agent messages and tool returns preserve provenance, tenant, integrity, sensitivity, purpose, term and derivation relationships. Unknown labels do not receive default high privileges.
- [] **Are decryption and downgrade controllable?** Plaintext enters only the necessary one-shot environment, and keys do not enter the model context. When the purpose does not match, labels conflict or a dependency is unavailable, there are read-only, draft, low-resolution, simulation and human-confirmation paths.

- [] **Is state partitioned by domain and revocable?** Session, long-term memory, public index, simulation and real state use different namespaces. Write gates, derivation lineage, snapshots, freshness and deletion receipts can be jointly verified.
- [] **Are capabilities parameterized and do they decay?** Tokens bind principal, object, action, parameter set, context, budget, term, delegation depth and one-time nonce. Transfer does not expand the object or extend the term.
- [] **Is the transaction object re-checked at submission?** Human approval binds the normalized object digest and provenance version. Submission re-checks key fields, policy, state and idempotency key, and partial failure has compensation or a safe holding state.
- [] **Does isolation strength match the consequence?** The rationale for choosing a process, container, microVM or separate security domain is explicit. Allowed and prohibited paths, redirection, proxies, host interfaces, destruction residue and management bypasses all have probes.
- [] **Does the budget cover the full chain?** Tokens, tool steps, recursion, concurrency, GPU, wall clock, storage, network and external costs are metered in a linked manner. Soft thresholds trigger downgrade, and hard thresholds trigger abort and cleanup.
- [] **Can the trajectory reconstruct causality?** The common envelope and domain events share an action ID, and asynchronous steps preserve parent-child relationships. A high-impact action that lacks required fields is not executed first and recorded afterward.
- [] **Does investigation evidence follow privacy minimization?** The primary trajectory does not store unnecessary prompts, images, audio, biometrics and secrets. Sensitive payloads enter a domain-partitioned evidence store, with role, case, term and deletion rules.
- [] **Are the RTO, RPO, and trusted baseline drillable?** Each security service separately defines its security downgrade level, time objective and rollback window. Identity leakage, artifact distrust, state contamination, capability mis-signing, isolation bypass and log unavailability all enter the drill matrix.
- [] **Are benign tasks tested alongside the variants?** Every attack or failure variant has a legitimate control, an observable first-broken interface, a highest consequence layer, latency, cost and manual burden. That avoids a superficial pass traded for shutting down the business.

At the end, the checklist asks one reverse question. If one of the controls fails, is the root of trust used by the next control truly different? If two controls share an encoder, identity, state source or permission, record it as a common cause. Architecture diagrams, design descriptions and code paths provide design clues only. Runtime probe receipts must separately record the input, the actual state transition, the exit code, the system version and the scope of applicability. Their job is to demonstrate that the corresponding control has been observed in the specified environment.

19.10 Five Common-Cause Misreadings of the Shared Services

Misjudgment	What is ignored	Correct architectural judgment
Many controls equal defense in depth	Multiple controls may share the model, prompts, tenant labels, identity root or policy publication	Defense in depth is evaluated by differences in trust roots, privileges and failure modes, not by the number of products
Treating content provenance and processing history as proof of content safety	Signatures and content credentials state who asserted what and whether an object was altered. They do not prove that content is true, harmless or already authorized by the business	Content provenance and processing history, digital watermarking, detection, and account and case evidence each answer different questions
Least privilege only requires shortening the tool list	A single tool may still contain read, write, cross-tenant, arbitrary-path and high-value capabilities	Privileges bind subject, object, action, parameters, time, network and environment preconditions
More logs means more secure	Logs that cannot be linked to actions, lack integrity, have no one on duty, or contain large numbers of secrets may expand leakage	Derive fields, retention periods, routing and freeze authority backward from investigation, gate and recovery questions
A safe shutdown does not need a recovery design	Mechanical state, transactions, contaminated state and trusted restart remain unclosed	Stop, rollback, rebuild, verification and tiered reopening must form a continuous operating rule

19.11 Bringing It into an Actual System

Consider a system that uses generated images, retrieval agents and warehouse robots at the same time. It can map the specialized defenses of Chapters 4–18 onto seven shared services: provenance, information flow, state, capability, isolation, trajectory and recovery. The dependency graph first marks shared model judgments, input labels, identity roots, policy publication and human interfaces. It then identifies the common-cause nodes that would make multiple controls fail at the same time. Independence is demonstrated when one node fails while the inputs and decisions of another control can still remain trustworthy.

The "image OCR—RAG—LLM—email tool" chain separates data objects from action capabilities. Provenance, tenant, purpose and integrity labels propagate with OCR, retrieval and summarization. The email capability, by contrast, binds to the subject, recipient, data category, time and one-time approval. Resetting a warehouse robotic arm also requires adding the device, operating mode, personnel state, position, speed and safety envelope. A natural-language proposal can then enter only an authorization process with explicit parameters.

State contamination and artifact distrust are handled along derivation relationships. Memory poisoning is traced to indexes, caches, summaries, snapshots and backups. The base model, text encoder, LoRA, custom nodes and container images together form a closure identity. Baseline, single-component, target-combination and adjacent-combination tests each record benign utility and target anomalies. Any change to a closure member reopens the affected admission unit.

A procurement transaction maintains explicit states among preparation, approval, submission, external confirmation and compensation. Exchange rates, payee accounts and inventory are business objects that must be re-verified before submission. Approval binds the object digest and version, and an idempotency key limits duplication. Physical goods already dispatched or mechanical actions already executed enter dedicated compensation and manual handling processes. A database rollback restores only reversible state.

Trajectory, isolation and recovery together supply operational evidence. The public event envelope keeps object references, results and evidence locations. Highly sensitive data, such as raw audio, is held only for the minimum retention period. Allow, deny and post-destruction probes then verify the business path, boundary crossings and residue, in that order. The system defines degraded tiers such as read-only, manual approval and restricted automation. Separately, it specifies trusted recovery points, RPO, RTO and a benign control before recovery, for the identity service, the state service and action authorization.

Summary: What Is Shared Is Control, Not Model Assumptions

The four classes of systems need not share one model. Nor must they share one set of attack metrics. What they must share is trusted provenance, identity and purpose propagation, the state lifecycle, independent action authorization, an isolation budget, a unified trajectory and recoverable operation. Specialized defenses lower the probability of one class of failure. The shared control plane limits how far a failure propagates across domains, and how severe the final consequences are.

The next chapter turns this architecture into day-to-day institutions. It asks what to submit at design review, and how to fix budgets and denominators during testing. It asks who signs off at the release gate, what to monitor in operation, and how to recover after an incident. Security becomes a sustained capability, rather than an architecture diagram, only once it enters versions, pipelines, on-call duty and the division of responsibilities.

\newpage

Chapter 20 From Threat Model to Operating Institutions

A security architecture stalls most easily at the review meeting. The diagram has trust boundaries and the document has a high-risk list. Yet no executable gate stands before release, and in operation no one is responsible for alerts, revocation and recovery. Models, data, tools and scenarios keep changing, so every critical change must re-run the matching controls and the matching evidence chain. That continuous evidence chain is the security artifact a team can hand over and keep on watch. This chapter condenses the preceding nineteen chapters into six execution phases. Sections 20.2–20.7 of the main text map onto the first through the sixth step in order. The first step freezes the scope and the change unit, and the second generates a test matrix from the threat record. The third uses hard gates and diagnostic scores to make the release decision, while the fourth continuously monitors the real control objectives. The fifth responds to incidents and verifies recovery, and the sixth turns review gaps into a falsifiable plan. Section 20.1 outlines the minimal security argument that runs through all six steps, and it does not count as a seventh step.

The figure in this chapter expands these six phases into seven links by evidence object. It does not start a separate process. System definition corresponds to the first step. Threat, control and test are three adjacent links that jointly expand the second step, where controls come from the minimal security argument and are tested in the matrix. Release corresponds to the third step. Monitoring corresponds to the fourth. The incident handling process, which moves from monitoring into review, corresponds to the fifth step. Review, together with the new verification obligations it generates, corresponds to the sixth. The seven links describe the evidence lifecycle; the six steps describe the handover of responsibility. The two have different granularity, but the mapping between them is unique.

Chapter Overview


This chapter turns the security architecture into organizational institutions for continuous operation. The minimal security argument links system scope, protected properties, claims, counter-evidence and recovery paths into an argument tree. The threat record then generates a test matrix that contains allowed paths and prohibited paths. The release gate puts security, benign utility, cost, rights, revocability and recovery on one decision surface. On that surface, non-compensable hard gates and the diagnostic metrics used for improvement each play their own role.

Operating institutions revolve around version diffing, continuous regression, cohort-based canary rollout, alert routing, emergency freeze, identity revocation and verifiable rebuild. When the model, the data, the tools, the policy or the environment changes, the affected evidence expires with it, and the matching tests and release decisions re-run automatically. Incident drills then close the loop from freeze through revocation and forensics to rebuild and review. Each of those stages gains an executable attribution: responsible parties, time objectives, evidence preservation and recovery sign-off.

Background Principles: Evidence Expires as Dependencies Change

This chapter is not about a static compliance report. Its object is an operational argument graph that is updated continuously. The inputs are the system snapshot, the threat model, the control configuration, test receipts, operational events and change diffs. The internal state is the model, data, tool fields, identity, policy, log format, executor and environment version that each claim depends on. Change impact analysis computes along dependency edges which evidence has expired. The test matrix fills in missing observations. The release gate then hands the valid evidence to approvers, to operations and to the on-call response chain for consumption. This survey calls the process of identifying assets, attacker capabilities, paths, boundaries and impacts threat modeling. It calls the product of that process at a specified system version the threat model. The two are not interchangeable. The security surface covers overreach, poisoning and evidence destruction on the security side, and physical or business harm on the safety side. It also covers whether the system promptly contracts its capability once evidence expires.

In the normal case only display copy changes, and that copy enters no input, control or scoring. The dependency graph then confirms that the core evidence remains valid, and only the corresponding checklist verification is needed. The boundary case is different. A tool field appears to have had only its name changed. The new name nevertheless invalidates parameter validation, the human approval digest, the authorization policy and the event parser all at once. Old receipts must then expire on their own. The related capabilities stay closed until the local regression and the end-to-end path close again. An unchanged model version cannot rescue this evidence either. Evidence binds to the complete operational graph, not to a product name. Follow the seven links in the figure. They show how a single change moves from system definition into threat, control, test, release and monitoring. They also show how that change triggers the next round after the incident review. At the same time, use the mapping from the previous paragraph to map the seven evidence nodes back to the six execution phases of the main text. Focus on which nodes cause old evidence to expire.



Seven ring nodes run from system definition, threat, control, test, release, and monitoring to review, and are then triggered by a change into the next round.

The figure supports evidence lifecycle management. It does not support a claim that one completed round of testing permits permanent release. Nor can it determine automatically whether the change impact graph is complete. When dependencies are unknown, receipts are missing, or common causes have not yet been modeled, the conservative action is to widen the retest scope and narrow the released capability. Record the unclosed items. Do not carry a historical "pass" forward.

20.1 The Minimal Security Argument

A security argument organizes claims and evidence within a given scope, users, environment and time. It establishes that the main risks have been identified, that key controls have independent evidence, and that

residual risks, failure conditions and recovery paths are visible to decision-makers. It is not a certificate that remains permanently valid across versions. When the scope, the capability or the dependencies change, the relevant claims and evidence re-enter review. The minimal security argument contains ten items:

1. The purpose, users, environment and explicit boundaries of the system.
2. An inventory of models, data, state, tools, identity, network and executors.
3. Protected assets and six security properties: integrity, confidentiality, availability, authenticity, controllability, and recoverability.
4. The U0—U6 trust interface diagram.
5. The key threat record τ , including attacker capabilities and budget.
6. The earliest blocking control and the consequence-limiting control for each threat.
7. The four-part metrics, test denominators and uncertainty, together with the evidence layer that static inspection, mechanism-level run, simulated execution or closed-loop simulation, end-to-end reproduction, and field incidents occupy.
8. Benign utility, operating cost, downtime and manual burden.
9. Monitoring, revocation, rollback, recovery and incident owners.
10. Open risks, their acceptors, their expiry times and their re-review triggers.

Each of the ten items covers a different responsibility. Provider evaluations support model-layer evidence. The deployer answers for retrieval, state, tools, identity, network, executors and business consequences. Traditional penetration testing covers conventional system interfaces. The plans a model forms under new inputs and long-horizon feedback must also be evaluated, by model tests and by closed-loop tests.

20.1.1 Taking Apart "Why It Can Run" with an Argument Tree

The top-level claim of a safety argument must bound the object, the scope, the environment, the time and the consequence. Consider an example: "In the low-speed sorting area of warehouse A, when version fingerprint v executes only whitelisted pick tasks, a single untrusted work order cannot bypass personnel detection and the action gate to trigger motion in a restricted zone." That sentence bounds the location, the operational graph, the task, the attack entry point and the protected property at once. It can therefore be falsified. The top-level claim decomposes downward into four kinds of subclaims:

1. **Inputs and state are controlled:** low-integrity content is not silently promoted to policy, and every state write carries a principal, a purpose, a term and revocation.
2. **Planning and capability are separated:** model output is only a candidate plan, and independent controls determine high-impact parameters and identity.
3. **Execution and feedback are constrained:** actions stay within the safety envelope, and external sensing and actuation receipts can verify outcomes.
4. **Failures are discoverable and recoverable:** critical traces are complete, control failures degrade or stop, and the team can revoke, rebuild and verify recovery.

Each subclaim decomposes further into observable invariants. Take "separation of planning and capability". At a minimum it must be proved that the model process has no authority to issue tokens. The authorization service accepts only canonicalized action objects. An approval binds the object hash, the policy version and the time limit. The original approval is invalidated after the recipient, the path, the amount or the speed changes. Architecture names locate components. Runtime probes and logs confirm whether these behaviors hold. The leaf nodes of the argument tree connect to evidence objects:

$$e = (\textit{claim}, \textit{artifact}, \textit{method}, \textit{snapshot}, \textit{scope}, \textit{result}, \textit{limit}, \textit{owner}, \textit{time}).$$

Here "claim" is the narrow claim being supported. "artifact" is configuration, logs, data, code, an external receipt or an incident record. "method" is recorded according to the actual execution tier, as static inspection, mechanism-level run, simulated execution or closed-loop simulation, end-to-end reproduction, or field incident. "snapshot" is the system fingerprint, and "scope" is the sample, budget and environment. "result" is the observation, and "limit"

is the limits of the evidence. "owner" and "time" denote the responsible party and the validity time. Evidence that is not bound to a snapshot and a scope cannot be judged to still apply after the system changes.

20.1.2 Evidence Mapping and Counter-Evidence

The argument tree stores supporting material, counter-evidence, unknowns and common causes side by side. It lets reviewers see which claim a piece of evidence actually supports, and what counterexample would overturn that claim. Every row of the table keeps the applicable snapshot, the evidence tier, uncovered conditions and invalidation triggers. That prevents one piece of material from being cited repeatedly for several incompatible claims. A practical evidence mapping table is as follows:

Claim	Direct evidence	Counter-evidence test	Common cause and unknowns	Invalidation handling
Cross-tenant data unreachable	ACL logs taken before retrieval, plus leakage probes	Switch tenant and cache key, then replay	A cache or connector bypass	Freeze private retrieval
Model cannot self-authorize	Identity policy, issuance logs, privilege-escalation probes	Have the model request an expanded permission scope	Policy and tools share an overly broad identity	Revoke the tokens and degrade to read-only
Robot does not leave the envelope	Safety controller traces, actuator receipts	Sensor loss and boundary trajectory injection	Planner and monitor share state	Safe stop and recalibrate
Incidents can be recovered	Rollback, rebuild and regression receipts	Interrupt logs, revocation or a backup restore	Backups and production share a contamination source	Isolate backups and rebuild from a trusted manifest

Counter-evidence tests actively seek situations that could overturn a claim. If the counter-evidence succeeds, the team strengthens its controls, narrows the top-level claim, or stops the corresponding capability. One known penetration path directly triggers the corresponding hard gate. Evidence strength is separated by observation tier. Static inspection proves that configuration, code or a manifest exists. A mechanism-level run proves that a safety fixture responds as expected under limited conditions. Simulated execution or a simulation closed loop proves an action path in a controlled tool or in simulation dynamics. End-to-end reproduction requires a real or verified-equivalent operational graph, data, identity, actuators, predetermined metrics, and closure across multiple valid runs. A field incident proves a specific historical path and consequence. Tier and coverage are recorded separately. A single field incident supports a specific historical path, whereas an overall probability requires comparable incident samples.

Safety arguments must also identify circular reasoning. Suppose the model generates actions, judges whether those actions are safe, interprets logs and scores itself. The four conclusions then share one failure source. Leaf nodes should come as far as possible from different roots of trust: external policy, actuator state, independent sensing, immutable logs, human verification or third-party receipts. Where independence cannot be achieved, the shortfall is recorded as a common cause. Repeated invocations count as the same evidence source.

20.2 Step One: Freeze the Scope and the Change Units

Before testing begins, freeze a reproducible system snapshot. Record the model and API versions, system prompts, parsers, retrievers, embedding models and index snapshots. Record the names, types, required fields and retention periods of memory fields. Include tool registries, policies, containers, dependencies, identity authorization scopes, network rules, hardware or robot firmware, datasets and graders. For closed-source services, save at a minimum the invocation date, the public version, the region, the parameters and the response fingerprint.

Change units determine when to rerun. By default, the changes listed below trigger a safety regression. They are the model or dialogue template, vision preprocessing, retrieval chunking and reranking, addition or removal of tool parameter fields, and changes to types or value ranges. Further triggers are memory write policy, adapters and weights, the action decoder, world model state or reward, identity permission scope, network, containers, controllers, sensors, and alerting policy. Comparing model versions alone misattributes system drift to the model.

Each release generates an immutable manifest that records file hashes, build provenance, test receipts, and approvals. Artifact signing supports verifying "which version was deployed", but it does not prove that the version's behavior is safe. A behavior gate and a provenance gate must both exist.

20.2.1 A System Snapshot Is More Than a Model Version

A snapshot can distinguish the "approved system" from the system actually running in production. One form it can take is a runtime graph fingerprint. That fingerprint is recomputable, comparable, and bound to evidence. A runtime fingerprint mismatch only proves that the deployed object has changed. It does not automatically prove that the change is malicious, or that the behavior is dangerous. The direct response is to suspend continued use of the old evidence and to trigger a differential review. The fingerprint can be formalized as:

$$v = H(M, P, D, R, Z, T, \Phi, I, N, E, O),$$

Here ***M*** is the model and adapters. ***P*** is prompting, parsing, and preprocessing. ***D*** is data and index snapshots. ***R*** is retrieval and reranking. ***Z*** is the names, types, purposes, provenance, and validity periods of memory and state fields. ***T*** is tools and action decoding. ***Φ*** is policy, ***I*** is identity, and ***N*** is network. ***E*** is the execution environment, and ***O*** is observation, logging, and alerting. The hash ***H*** serves to identify the combination. It does not indicate that the combination is safe.

Every node in the manifest records its provenance, an immutable digest, build or acquisition time, dependencies, license, owner, and revocation status. Dynamic services must also store a configuration snapshot and an observable fingerprint. When a closed-source API's weights cannot be frozen, the date, region, public model name, request parameters, response headers, fixed probe results, and service status can all be stored. These fields establish a re-identifiable approximation. The internal version remains marked as invisible to the provider.

Snapshot storage is separated from production write permissions. The release pipeline can submit candidate manifests, but it cannot overwrite approved records. The runtime environment can only read the approved version. Forensic storage uses append-only integrity protection. The version fingerprint represents the actual runtime state only while production jobs cannot change policies, tool descriptions, or logs on their own.

20.2.2 Change Impact Analysis

A system change first generates a diff Δv . Impact then propagates along the dependency graph. A change to a single vision scaling parameter may affect OCR, model inputs, attack payloads, and benign recognition. A change to the field names, data types, required fields, or value ranges of tool parameters may affect model selection, parameter canonicalization, human confirmation, and log parsing. An expansion of identity permission scope may let a dangerous plan that previously reached only Y2 reach Y3. Impact analysis must proceed along "who consumes this object, which claim it affects, and which tests depend on it." This survey grades changes as CHG0—CHG3 by potential impact. That grading is the operational governance work grading used by this survey, not an external standard:

This survey's change level	Example	Minimum action
CHG0 surface change	Display copy that does not affect inputs, control, or grading	Manifest check and sampled smoke tests
CHG1 local behavior change	Prompt templates, chunking, thresholds, low-risk tool parameters	Regression of affected invariants and benign utility
CHG2 capability or boundary change	New tools, identity permission scope, network, memory writes, action decoding	Full threat matrix, fault injection, and canary rollout
CHG3 high-consequence or common-cause change	Safety controllers, payments, releases, mechanical envelopes, logging, and revocation roots	Independent review, exercises, rollback verification, and re-acceptance of residual risk

The worst reachable consequence determines the level, not the number of lines of code. A one-character network allow rule may be CHG2. Thousands of lines of documentation generation with no execution path may be CHG0. Automatic classification can produce a candidate level. The accountable owner, who understands the assets and the consequences, ultimately confirms it.

20.2.3 Checking Evidence Validity Back from the Diff

Each piece of evidence records the set of nodes it depends on, $dep(e)$. When the change set is Δ , old evidence does not remain valid merely because the file still exists. The expression below first identifies all evidence that requires revalidation. This operation only identifies the set whose dependencies may be affected. It does not prejudge that the post-change version must fail. Until revalidation is complete, the corresponding claim remains in a pending-validation or restricted state. The set requiring revalidation is:

$$E_{\text{retest}} = \{e \mid dep(e) \cap closure(\Delta) \neq \emptyset\}.$$

"closure" denotes the downstream nodes that a change can affect along the dependency graph. This rule avoids mechanically running the full suite every time. It also avoids continuing to use inapplicable historical receipts after a critical boundary change. If the dependency graph is incomplete, the conservative approach is to widen the revalidation scope rather than assume no impact. Impact analysis answers at least the following questions. Does the attacker gain a new write entry point, feedback, or budget? Can low-integrity data influence new variables? Does the plan gain higher capability? Can logs still be correlated with new actions? Does rollback cover successor components? Do rights and retention periods change? If any answer is unknown, the high-impact release enters a stop or restricted canary. The unknowns remain explicit, and subsequent evidence closes them.

20.3 Step Two: Generating a Test Matrix from the Threat Record

Every threat τ derives at least four categories of test. The four categories are not different samples of one success rate. They separately verify operational availability, prohibited consequences, error handling, and trustworthy recovery. Each therefore occupies one row type in the matrix and keeps its own denominator. A pass in one of the four rows cannot offset a failure in another row. A run that does not reach the intended interface enters availability statistics rather than the attack denominator.

Test type	Core question	Typical input or scenario	Results and boundaries that must be observed
Benign allowed path	Does control preserve legitimate tasks	RAG retrieval in this tenant, normal video generation, in-envelope grasping, a world model issuing short actions within the safe envelope	Task completion, benign utility, latency, and cost; "safety" cannot be obtained by rejecting everything, and a degraded stop does not count as a benign success
Malicious prohibited path	Under the intended attack budget, does the first-broken interface continue to reach high-consequence layers	Direct and indirect prompt injection, multimodal and cross-turn state, malicious artifacts, unauthorized objects, and irreversible actions	Model outputs, dangerous plans, authorization, execution, and real-world consequences are stratified; when later layers are not reached, the denominator is recorded as NA
Boundary and error paths	On parsing, state, or infrastructure failure, does the system enter the intended safe state	Empty responses, timeouts, stale state, missing sensing, degraded stop when the world model is uncertain, scoring conflicts, redirection, resource exhaustion, and human refusal	Failure receipts, capability contraction, safe failure, and availability; a request error must not be recorded as an attack miss, and a degraded stop is counted as a safe failure rather than a benign success
Recovery path	After blocking, can the system return to a verifiable trusted state	Freezing, revocation, rollback, state reconstruction, transaction compensation, and re-authorization	Recovery point, RPO/RTO, uncleared objects, benign task regression, and layered reopening; a restart is not a recovery

Each matrix row records the sampling unit and the sample source. It also records the model/system snapshot, the attack budget, and the highest Y layer. Expected action, benign utility, and log location complete the row. Test results can then return directly to the safety argument and tie to the corresponding claim, version, and consequence layer. When the row's original input, control decision, or result denominator cannot be reconstructed, that piece of evidence remains undecidable. A summary score must not be used in its place.

20.3.1 Generating the Matrix from Threat Factors

Test design lists factors first rather than collecting attack examples first. General factors are the first-broken interface, entry modality, attacker knowledge, write capability, feedback, budget, state duration, target action,

consequence layer, control configuration, and failure state. Domain factors are then layered on. Visual generation adds spatiotemporal and provenance signals. VLA adds environment and actuators. World models add state, rollout length, and controller. Combining all factors in a Cartesian product would quickly run out of control. The coverage strategy has three tiers:

1. Every single-factor level appears at least once, which confirms that the basic interfaces have no gaps.
2. For medium- and low-risk factors, pairwise or three-way combinations uncover common interactions.
3. For high-risk attack chains, use the complete directed sequence rather than let combinatorial sampling omit necessary preconditions.

Consider low-integrity web pages, long-term memory, and outbound tools. Each of the three may pass testing separately. Even so, the complete sequence must still be run: "a web page first writes memory, a future session recalls it, and an outbound transfer then forms". Combinatorial coverage can reduce the sample size for ordinary configurations. The critical paths already identified in the safety argument remain complete. A weighted coverage rate can be defined:

$$Coverage = \frac{\sum_i w_i \mathbf{1}[t_i \text{ executed and valid}]}{\sum_i w_i},$$

Here t_i is the pre-registered test obligation. Assets, reachable consequences, and control uncertainty determine w_i . Request errors, empty responses, grader failures, or fixtures that never started do not count as "valid". They enter availability and test-infrastructure gaps. Coverage is only a diagnostic value. Any high-weight hard-gate failure still stops release.

20.3.2 Evidence Requirements for Each Test Row

A complete row record contains at least the following fields. They let results point back to the threat record, the frozen snapshot, the control configuration, the execution environment, and benign utility. These fields are the minimum contract needed for replay and denominator reconstruction. They do not replace domain-specific receipts for sensing, content, rights, or business. When a field is missing, the original failure state is retained. The gap is not inferred and filled in to complete the record:

$$t = (id, \tau, v, x, B, C, E, Y, J, U, K, L).$$

Here τ is the threat record. v is the snapshot. x is the input and state. B is the attack budget. C is the control configuration. E is the execution environment. Y is the highest consequence layer. J is the grader. U is benign utility. K is cost. L is the log location. Paired comparisons and failure reproduction are possible only when inputs, random seeds, scenarios, and tool receipts have stable identifiers. Each row also states the expected allowed behavior and the expected prohibited behavior. The legitimate path of an email agent may read emails specified by the user and create a draft. Its prohibited path must not add unapproved recipients. A warehouse robot is allowed to grasp within the envelope. Its prohibited path must not keep moving when person detection is unknown. Expected outcomes further distinguish model refusal, capability-gate blocking, safety stop, and test timeout. They must not be merged under a single "attack failure" label.

20.3.3 Denominators, Pairing, and Coverage Gaps

Safety results are reported by stage. Input arrival, control deviation, dangerous plan, authorization acceptance, execution, and real-world consequences each have their own denominator. When a later stage's denominator is zero, it is recorded as NA. Benign utility uses the legitimate-task denominator. Cost uses the denominator of actually valid runs or actions. Cross-paper numbers serve as external references. The deployment task, attack budget, and consumer still determine local measurement rules. Multiple models or random seeds for the same case are correlated observations.

Version comparison prioritizes pairing. The same task, input, scenario, budget, and grader are run separately for A and B. This records the transitions from success to failure, from failure to success, and with no change. A paired design can hold scenario difficulty constant within the case. When one version has a request error, that pair enters availability analysis and keeps its original denominator. Coverage gaps enter release materials together with

failures. Several conditions limit the highest evidence layer. They are an unknown true model version, no physical environment, missing authorized test accounts, an inability to simulate irreversible transactions, no independent sensing, and no recovery snapshot. The team narrows capability and release scope. "Follow-up testing" remains an open item with an owner and a deadline.

20.3.4 Test Data and Rights

Attack tests may contain malicious code, real people, private documents, or dangerous actions. Datasets should carry labels for provenance, license, sensitivity level, holder, retention period, and permitted environment. In leak probes, real secrets are replaced by synthetic marked objects. Embodied trials are first conducted in isolated simulation and at a safe site. External network targets must have explicit authorization. Fixture outputs may also contain leaked content. Raw responses, screenshots, videos, network packets, and sensor trajectories enter a controlled evidence repository. Reports expose only the necessary summaries and hashes. When deletion requests conflict with event-retention obligations, legal and safety owners decide the minimum retention scope. The model must not decide on its own.

20.4 Step Three: Dividing Work Between Hard Gates and Diagnostic Scores

Release decisions require two classes of mechanism. Hard gates handle non-compensable conditions. Examples are unauthorized data, unverified high-impact permissions, dangerous execution above the threshold, missing critical logs, an inability to revoke, and unclear citations or rights. Any hard-gate failure stops release. Other high scores cannot offset it. Diagnostic scores are used to prioritize improvements. Examples include benign task success, latency, human escalation, detection coverage, test diversity, documentation completeness, and recovery time. Scores show trends. Each class of risk and the hard gates are still presented separately in their original units. A general release gate can be written as:

$$\begin{aligned} \text{Release} = & H_{\text{rights}} \wedge H_{\text{identity}} \wedge H_{\text{execution}} \\ & \wedge H_{\text{domain-safety}} \wedge H_{\text{recovery}} \wedge H_{\text{evidence}} \\ & \wedge D_{\text{acceptable}} \end{aligned}$$

The six ***H*** stand for the hard gates on rights, identity, execution, domain safety, recovery and evidence. ***D_{acceptable}*** marks the interval of utility and cost that the decision-maker accepts. Thresholds must follow consequences. A read-only internal summary and an automated payment cannot share the same upper bound on dangerous execution. The same holds for a published marketing video and a high-speed machine control. The release gate must also name its signing roles. The model owner confirms the model and data snapshots. The application owner confirms tools and tasks. The safety owner confirms threats and controls. The business or safety responsible person accepts the residual risk. The legal/rights owner confirms data and content use. The operations owner confirms monitoring and recovery. One person may hold several roles, but responsibility must not disappear into "the team knows".

20.4.1 Three Kinds of Gates Carry Three Kinds of Responsibility

Hard gates judge whether an unacceptable gap exists. Their triggers include cross-tenant reads, unauthorized high-impact execution, an inability to stop safely, a broken chain in critical trajectories, an identity that cannot be revoked, unknown artifact provenance and unclear rights. A hard gate returns only pass, failure or undecidable, the last when the evidence is insufficient. Failure and undecidable both bar a full release.

Diagnostic gates observe trends and trade-offs: task success, false refusals, latency, token usage, energy consumption, the human queue, coverage and recovery time. A diagnostic item may carry a target interval, a warning line and a deadline. Teams use these items to decide the gray-release proportion, capability limits and follow-up work. Only a hard gate settles a non-compensable condition. A diagnostic gate exposes trade-offs such as "safe but unusable."

Exception gates handle cases the business genuinely needs, whose risk can be bounded and which cannot be closed immediately. An exception does not turn a failure into a pass. It forms a restricted authorization object:

$\epsilon = (\text{risk}, \text{scope}, \text{reason}, \text{compensation}, \text{owner}, \text{expiry}, \text{signal}, \text{revoke})$.

That object records the risk, the scope of application, the business justification, the compensating controls, the acceptor, the expiry time, the monitoring signals and the automatic revocation conditions. An exception that lacks a named owner, an explicit scope or an expiry time is invalid.

20.4.2 The Boundary of Exception Governance

Certain items must not be released by exception. Data or content used without rights stays blocked. External targets beyond the testing authorization stay blocked. So do high-consequence actions that cannot be safely stopped and critical identities that cannot be revoked. An evidence chain too thin to establish what the system has deployed also blocks release. An organization may narrow functionality or delay release. It may not let an internal sign-off replace the rights of others or the real-world safety conditions.

Governable exceptions usually fall into diagnostic degradation or a limited control gap. A low-risk read-only task may run above its latency target; the gray release can then stay within a smaller tenant scope. A non-critical alert may grow noisy; actions can then be restricted and human review increased. A recovery time may miss its target; the team can then keep the previous version for a fast rollback at lower concurrency. Compensating controls must act at a different location from the gap.

When an exception expires, it returns automatically to a failed state. Renewal means re-verifying the risk, the incidents, the compensating controls and the business necessity. Changing only the date is not enough. Monitoring may report a predefined revocation signal (for example, a rise in dangerous plans, an overflow of the human queue, missing traces or a failed recovery drill). The system then degrades or freezes automatically, without waiting for the next meeting.

20.4.3 Release Decision Table

Gate	Outcome	Permitted action	Interpretation not allowed
Hard gate passes	Valid evidence and no known penetration	Proceed to diagnostic and gray-release judgment	Equals absolute security
Hard gate fails	An unacceptable path has been observed	Stop the corresponding capability and remediate	Offset by high utility
Hard gate undecidable	Invalid request, missing environment, or insufficient evidence	Additional testing or narrower scope	Recorded as zero attacks
Diagnostic warning	Metrics degrade but do not cross a hard boundary	Limit the gray release, set a deadline and monitoring	Silently ignored
Restricted exception	Risk can be bounded and compensation is effective	Run within scope and time limit	Permanent exemption

The signing order must avoid conflicts of interest. The implementer submits the evidence. An independent reviewer verifies the gate. The business owner accepts the explicitly stated residual risk. Operations confirms the response capability. For high-consequence systems, the delivery lead and an independent security approver sign under separated authority.

20.5 Step Four: Continuously Monitor the Real Control Objectives

After launch, the most valuable signals are not that the model "looks abnormal". They are that low-integrity data approaches high-impact actions, and that permissions and state change abnormally. Cross-domain generic signals include:

- content of external origin influencing tool names, recipients, paths, amounts, action speed or security constraints;
- cross-tenant retrieval, memory recall or state merging;

- drift after approval in tool registrations, in the names, types or value ranges of parameter fields, in model components or in policies;
- new domains, redirects, private-network or metadata access;
- secret patterns, anomalous reads and writes, irreversible operations and privilege escalation;
- agent loops, output bloat and anomalies in GPU/tool/cost budgets;
- visual authenticity signals conflicting with accounts, signatures or content history;
- persistent divergence among the predictions of a robot or world model, external sensors and the safety controller;
- log interruptions, broken action ID chains, clock anomalies and alerts that are not escalated.

Every monitoring threshold needs a normal baseline and event base rates. Where the base rate is low, detection recall cannot be the only goal. The FPR, the review queue and the action cost must also be controlled. A high-impact action can tolerate a higher false-positive rate. Safe degradation and service recovery must still be designed for it, so that the classifier cannot be exploited to mount a large-scale denial of service. Version monitoring matters just as much. Closed-source APIs, open weights, adapters, data and attacks all drift in behavior. Fixed regression sets detect known entry points. Rolling red teaming explores new combinations. Production incidents feed back into the use cases. A fixed set should not expose all its details publicly over the long term. Within the authorized scope, adaptive attacks should receive a feedback budget consistent with the threat model.

20.5.1 The Capability Ladder of Gray Release

A canary release first sends a candidate version to a small portion of traffic or instances. Inside a predefined observation window it watches hard stop events, trend warnings and rollback capability. It then decides whether to expand the exposure surface. This release concept is not the leakage probe used to detect secrets or cross-tenant exposure. Gray release is not merely moving traffic from 1% to 100%. Capabilities must also be opened level by level. A traffic proportion describes only the exposure surface. Retrieval, state writes, outbound sending, payment and physical execution carry different maximum consequences. Each level must therefore bind the available capability, the applicable subjects, the observation window, the stop signals and the rollback target. A generic ladder includes:

1. **Offline replay:** the candidate version reads fixed traces and does not touch live users or external tools.
2. **Shadow run:** it receives live inputs and generates plans, but holds no state-write or execution authority.
3. **Small-scope read-only gray release:** retrieval or observation is opened to a small number of controlled subjects, and outbound sending and irreversible actions are not allowed.
4. **Restricted execution:** only low-impact, rollback-capable actions are opened, under tighter budgets and human confirmation.
5. **Cohort gray release:** expand by tenant, scenario, region, device or action family, without automatically opening high-risk capabilities along with the traffic.
6. **Stable operation:** once the predefined window is reached, the system enters normal operation and retains rollback, continuous regression and sampled review.

Each level has entry conditions, a minimum observation window, stop signals and a rollback target. A dangerous plan found in a shadow run is retained as a leading signal for the capability-layer gate, and it enters the metrics even without real execution. During restricted execution, a failed hard gate on traces, identity or recovery returns the system to read-only or offline immediately. Wherever possible, the candidate version and the baseline should receive the same requests and state snapshots, so that plans, policy decisions, latency and benign utility can be compared. Privacy or non-repeatable physical scenarios rule out unconditional copying of requests. Teams then use authorized replay, simulation or paired scenarios, and they mark the items that cannot be paired.

20.5.2 Leading Signals, Outcome Signals, and Control Health

Monitoring has three layers. **Leading signals** appear before real-world consequences: missing provenance, cross-tenant near hits, dangerous plans, anomalous tool selection, state divergence and budgets approaching their

limits. **Outcome signals** include tool acceptance, publication, payment, content propagation, robot safety stops and user appeals. **Control health** includes policy latency, token issuance failures, network proxy bypass, broken log chains, time drift, rollback availability and human queue length. An individual metric records the object, the denominator, the window and the version:

$$m = (\textit{name}, \textit{numerator}, \textit{denominator}, \textit{window}, \textit{segment}, \textit{v}, \textit{threshold}, \textit{action}).$$

Take the "dangerous plan rate": its denominator is the tasks that obtained a valid plan, not all HTTP requests. The "cross-tenant leakage rate" requires valid private retrieval trials. The "safety stop rate" reports both the trigger reason and the benign task loss. A sudden drop in the denominator may be a service error rather than a security improvement. Gray-release stop rules prioritize hard events and persistent trends. A single unauthorized real execution is enough to trigger a freeze. When dangerous plans, false refusals, latency or control errors cross the warning line across consecutive windows, expansion is paused. Observation windows cover weekdays, nights, peaks, different tenants and physical environments. A stable conclusion is bound to the complete window.

20.5.3 Drift and Data Feedback

Sources of drift include the input population, language, modality, retrieval corpus, tools, environment, model service and attack adaptation. Teams keep stratified baselines. They do not let an overall mean mask degradation in a small group. A new language or device that lacks a sufficient denominator keeps its capabilities restricted. That restriction holds until evidence from legitimate tasks and security tests has been collected. Production incidents can become test leads, but they cannot be fed back directly without processing. Remove real secrets and victim identities first. Confirm rights and purpose, reconstruct the minimal trigger conditions, and then add them to controlled regression. Attacker inputs published as-is to a shared test library may leak defense details or trigger tools again. Safety samples require access control and execution isolation.

Controls themselves also drift. On-call staff may approve alerts out of habit. Exceptions may be renewed repeatedly. Network allowlists may expand, and log sampling may decrease because of cost. A monthly control review examines these institutional changes. It treats "who is still using exceptions, which capabilities have never been rolled back, and which alerts have not been handled" as first-class metrics.

20.6 Step Five: Incident Response Proceeds in Three Phases

After an incident occurs, the response order should revolve around the capability chain, not an argument over whether the model was "really attacked". The three phases hand off in sequence. A phase does not enter full capability recovery until the preceding phase has formed a verifiable state. Verifiable control states determine the phase boundaries, not meeting times. Each handoff carries confirmed facts, current permissions, evidence locations and outstanding risks.

Phase one: containment and evidence preservation. Correlate the user objective, provenance, state, plan, authorization and action. Freeze new high-impact operations. Physical systems enter the predefined safe state. In parallel, the platform revokes short-lived tokens and suspicious artifacts, blocks network and release, and isolates the affected tenants, agents or devices. Save immutable snapshots and volatile state before cleanup, so that "stopping further spread" and "preserving the causal chain" happen together. Dangerous capabilities must not remain open for the sake of forensics.

Phase two: localization, cleanup, and trusted rebuild. The team determines the highest Y layer, the first-broken interface, the affected assets, the time window and the derivatives. It distinguishes model output, plans, execution and real-world consequences, and it separates confirmed facts from inferences. It then rotates credentials and rolls back memory, indexes, models and policies. It cleans caches and backup derivatives, rebuilds the environment from a trusted manifest, and resynchronizes world state, sensors and external transactions. When no trusted recovery point exists, the system remains isolated rather than being forcibly restarted.

Phase three: verify recovery, learn, and disclose. Re-run the tests on the original entry point, adjacent interfaces, benign tasks and recovery paths. Once the same failure condition is confirmed closed, open capabilities progressively in the order read-only, restricted and full. Observe continuously whether the failure recurs. The postmortem converts root causes and failed controls into automated gates, and it records why blocking or

escalation did not occur. Disclosure follows legal obligations, contractual terms and the needs of affected parties. Incident counts and the sampling frame are kept separately. A single incident is not expanded into an overall frequency.

A mature postmortem answers at least these questions. Which low-trust object obtained what influence? Why did the upstream control fail? Which consequence control succeeded or failed? When did the logs see it? Who should be notified? Which credentials, states and derivatives need to be revoked? Which trusted version should be restored? How can the same common cause be prevented from penetrating again?

20.6.1 Grading Is Based on Confirmed Consequences and Still-Reachable Capability

Incident severity integrates the highest privilege reached, actual execution, real-world consequences, reversibility and evidence completeness. The initial grading must be able to trigger on-call, freeze, revocation and notification actions directly. A four-level system that can be escalated or de-escalated as new evidence arrives is as follows:

Level	Typical conditions	Initial action
S0 Critical incident	Confirmed danger to persons, major irreversible execution, widespread compromise of data or the control plane	Immediate safe stop, highest-level on-call and management escalation
S1 High incident	Unauthorized execution, sensitive cross-tenant access, penetration of publishing or identity chains, with impact still possibly expanding	Minute-level freeze, revocation, isolation, and forensics
S2 Medium incident	Control deviation or a dangerous plan blocked by a consequence gate, or localized control failure	Stop the relevant capability, limit the scope, and expedite remediation
S3 Observation incident	Anomalous input, probes, or near-miss signals with no confirmed control deviation	Preserve evidence, monitor trends, and handle as scheduled

A level may be escalated or de-escalated, but the basis for each change is recorded. An unknown scope of impact does not justify a lower level. When the reachable scope of identity, logs or network is unknown, treat the case conservatively, guided by the possible consequences. False refusals or safe stops that occur at scale may also become S1- or S0-level availability incidents, especially in healthcare, industry and public services.

20.6.2 On-Call Roles and Decision Authority

The on-call system has at least six role types: incident command, technical response, forensic records, business continuity, legal and notification, and domain safety. Incident command owns the facts, the timeline and the decisions, but does not personally execute every command. Technical response owns the freeze and revocation paths. Forensic staff protect evidence. The business owner decides how to degrade service safely. The legal and privacy lead judges notification obligations. The robotics, visual content or world model lead supplies domain consequence judgments.

Every role has a primary and a backup. Contact details and permissions are verified periodically. An emergency stop of high-impact actions should not depend on everyone being present. Restoring full capability requires two-person or split approval. Escalation paths for vendors and cloud services are established before an incident. The team must not go hunting for an account manager on the fly.

Handover runs from a one-page status board: confirmed facts, items that remain inferred, the highest Y layer, the first-broken interface, affected subjects and assets, current freezes, revoked identities, evidence locations, the next decision, and the owner. Verbal information does not replace timestamped records.

20.6.3 RTO and RPO Must Be Bound to the Security State

The recovery time objective (RTO) is the maximum time allowed between a service falling below a predefined secure service state and that service state being restored. The recovery point objective (RPO) is the acceptable window of state or data rollback. Both are target ceilings set before an incident, not the actual observed values of a particular recovery. The RTO clock starts when the service is interrupted or drops below the target state, not when the team detects an incident or formally declares one. Detection time and declaration time are recorded separately, so that detection latency and declaration latency can be computed. A later declaration moment must not be used to shorten the RTO.

$$T_{\text{recovery,actual}} = t_{\text{validated service}} - t_{\text{service interruption}} \leq RTO_{\text{target}},$$

$$\Delta t_{\text{rollback,actual}} = t_{\text{service interruption}} - t_{\text{trusted restore point}} \leq RPO_{\text{target}}.$$

Record both t_{detected} and $t_{\text{declared incident}}$. Report the interval from service interruption to detection and the interval from detection to formal declaration separately. These two operational metrics help diagnose the response chain, but they do not change the starting point of the actual recovery duration. If the actual recovery duration or the actual rollback window exceeds the target, the recovery exercise is judged as not meeting the standard, even if it ultimately started the service. If the target has not yet been approved, only the observed values may be reported, and they must not be retroactively named RTO or RPO.

"validated service" must state the capability explicitly. A read-only query may be restored within 30 minutes. A tool-bearing agent may be restored after 4 hours. Automated publishing or full-speed mechanical control may require longer validation. A running service process means the service is connectable. Completing recovery also requires confirming that memory, identity, state and high-impact capability have all reached the corresponding thresholds. Different states need different targets:

Object	RPO focus	RTO verification conditions
RAG and long-term memory	Most recent trusted index, deletion markers, and tenant ACLs	Private retrieval leakage probes, deletion replay, and normal tasks pass
Visual publishing	Trusted model components, content credentials, and the pending-publish queue	Digital watermarking, content provenance and processing history, and review and revocation paths are effective
VLA/robotics	Safe shutdown state, calibration, and task queue	Sensing, envelope, emergency stop, and low-speed trial runs pass
World model	Trusted observations, latent state, and controller snapshot	Resynchronization, external anchors, and safe control switching pass
Identity and tools	Last trusted policy, tokens, and registry	Revocation takes effect, and audience verification and privilege-escalation probes pass

The smaller the RPO, the higher the cost of logs, snapshots and replication usually is. Replication also propagates contamination. Critical state needs versions, derivation relationships and isolated backups. "real-time synchronization" is responsible for freshness, while trusted recovery points are separately confirmed by integrity and isolation verification.

20.6.4 Forensics and Evidence Preservation

The forensic order prioritizes volatile information: running identities, processes, network connections, memory and queues. Next come model inputs and outputs, state reads and writes, policy decisions, tool receipts, files, images and control-plane logs. Physical systems additionally require raw sensor data, controller traces, clocks and on-site state. Visual incidents retain content credentials, encoding chains and platform receipts.

Every piece of evidence records its collector, time, source, method, digest, access and copies. Analytical copies are separated from the original evidence, and any transformation preserves derivation relationships. When clocks disagree, use common action IDs, network handshakes and external receipts to establish partial-order constraints and mark uncertainty intervals. Privacy and forensics must be satisfied at the same time. Collect only the scope needed to answer the incident question. Keep sensitive content encrypted and separated, authorize access on a per-case basis, and dispose of data when the retention period expires, in accordance with legal obligations and contractual terms. Retaining user context indefinitely because it "might be useful later" turns the response system into a new high-value data plane.

Forensic findings are layered as confirmed, strong inference, weak clue and unknown. Model self-reports provide behavioral clues. Internal causes require intervention or mechanism evidence. Tool receipts confirm execution. Asset inventories prove that an asset exists, and read logs prove actual access. Write permissions and write events are recorded separately. Such an evidence mapping keeps the scope of impact accurate.

20.7 Step Six: Turning Open Questions into Falsifiable Plans

A research agenda advances safety capability only when it writes down hypotheses, protocols, endpoints and stopping rules at the same time. High-value questions in the four domains can be posed in a common format. Language agents need to test three questions. Can structured instruction/data separation preserve task utility under adaptive attacks? Can memory write policies prevent dormant and compositional triggers? How does tool and network capability affect the upper bound of dangerous execution?

Visual generation needs to test when image defenses fail under spatiotemporal, audio-visual and streaming conditions. Can the behavior of compositional adapters be discovered before release? How do digital watermarking, content provenance and processing history, content credentials, detection and platform action form a usable evidence chain at low base rates? VLA/WAM need to test how perception and language attacks pass through the action decoder. Can independent anchors and action gates limit physical consequences without significantly harming the task? Under which environment changes do simulator results still extrapolate?

World models need to test how state, dynamics, reward and planning attacks accumulate over long rollouts. When can uncertainty reliably trigger safety controls? How do robust MPC, runtime assurance and rollback compare under a unified protocol? Any plan should distinguish frozen weights from optimized variables, record the attack budget, and separate proxy metrics from real consequences.

20.7.1 The Minimal Structure of a Falsifiable Protocol

An open question can be written as the nine-item protocol below. It turns an abstract topic into an execution contract that counterexamples can overturn. It also prevents researchers from changing endpoints or shrinking the failure denominator after seeing results:

1. **Claim and counter-claim:** state clearly what observation would support and what observation would overturn.
2. **System boundary:** inputs, state, outputs, actions, feedback, frozen weights, and optimizable variables.
3. **Attacker capability:** knowledge, write access, queries, time, physical contact, and the feedback budget.
4. **Comparison groups:** baselines, candidate controls, ablations, and failure conditions.
5. **Preregistered endpoints:** first-broken interface, highest Y layer, benign utility, cost, and recovery.
6. **Statistical unit:** task, prompt, trajectory, video, device, scenario, organization, or event.
7. **Evidence layer:** static checks, mechanism-level runs, simulated execution or simulation closed loop, end-to-end reproduction or field incident.
8. **Stopping and safety rules:** when to stop due to risk, cost, ineffective response, or evidence saturation.
9. **Extrapolation boundary:** which models, versions, environments, languages, and consequences are outside the conclusion.

A hypothesis must include an observable direction. "structured context is safer", for example, is not falsifiable. This version is executable: "under the same model, tasks, and query budget, structured provenance labels reduce the proportion of low-integrity web pages that enter high-impact tool parameters, while the legitimate task completion rate does not fall beyond a preset tolerance range". Where query errors mean that there is no valid answer, the outcome is an availability gap, not an attack proportion of zero.

20.7.2 Separate Causal Questions from Descriptive Questions

When the goal is to determine whether a control caused a change, prefer pairing, randomization or controlled ablation. Run the same task with and without the control, and save the paired transitions. Randomize or balance the execution order to reduce service drift. Hold snapshots fixed as much as possible, except for the target component. Field data that cannot be randomized support association and mechanism clues. Causal attribution separately requires comparable controls and confounder analysis.

Descriptive questions also have value. Take counting. Which sources most often enter dangerous plans in production? Which stale states most often trigger safe stops? Such counts can guide test priorities. A report of that kind must make clear that it answers coverage and distribution, not the causal effect of a control. Heterogeneous endpoints are not forcibly pooled. Content violation rate, word recovery rate, dangerous plans, tool execution,

collisions, prediction error and real incidents belong to different layers. Only when the statistical unit, attack budget, scorer and consequence layer are compatible, and independent repetitions exist, should pooled estimation be considered. Otherwise keep stratified tables and intervals.

20.7.3 Four-Domain Protocol Example

Domain	Falsifiable hypothesis	Core comparison	Key endpoints	Main extrapolation boundary
LLM agent	The capability gate limits dangerous execution under adaptive injection while preserving minimum benign utility	No gate, task gate, parameter gate, and combined control on the same task	Dangerous plan, authorization, simulated/real execution, task completion	Tools, identity, model, and budget
Image/video generation	Spatiotemporal provenance signals remain detectable after predetermined transformations without significantly degrading usable image quality	No signal, single-frame signal, spatiotemporal signal	Detection, false positives, image quality, platform handling	Encoder, transformation, base rate, and platform
VLA/WAM	An independent action gate limits out-of-bounds behavior under observation attacks while preserving routine tasks	Planner alone, model self-review, independent safety controller	Plan, execution, task completion, shutdown, and recovery	Robot, speed, scenario, and sensing
World model	External anchors and runtime assurance limit control deviation under rollout-state attacks	No anchor, internal consistency, external anchor plus safety control	State error, trajectory, action, rollback	Dynamics, rollout length, and real environment

The endpoints in the table still require denominators written per concrete system. VLA uses an independent trajectory or task as the unit, and video uses an independent video as the unit. For world models, multiple rollouts that share an initial state are grouped into the same correlated group. Action frames and video frames remain as within-group observations.

20.7.4 Stopping Rules and Research Safety

Stopping rules are fixed before the run. The run stops when the query, cost, wall-clock or physical-count limit is reached. When consecutive request errors reach a threshold, the run pauses and the infrastructure is diagnosed. When a real secret, an unauthorized target or a safety-envelope breach is found, the run terminates immediately and enters the incident process. Interim observations, if they are used only for safety, are not used to select the best result.

Real networks, accounts, people, devices and scenarios require authorization. Attack fixtures use leakage probes, isolation tooling and disposable environments. Dangerous actions are preferentially dry-run only (a dry run performs validation and transformation but does not persistently write the intended change) or safely simulated. Physical tests set emergency stops, observers, exclusion zones and maximum energy. Open materials remove directly reusable secrets and victim identities while retaining parameters, hashes and state descriptions sufficient to verify the method.

Negative results must also be published with their boundaries. The phrase “not observed under this model, sample, budget, and grader” does not mean that the attack does not exist. A positive result indicates that at least one path exists. It does not give an industry-wide frequency. A falsifiable protocol is valuable because it tells the next team which conditions to change to test when the conclusion stops holding.

20.8 Worked Example: How a Version Change Passes the Gate

This section is an institutional-design worked example. All systems, scenarios and numbers are hypothetical data, used only to demonstrate the gate and paired analysis, and are not an executed experiment or production performance. Suppose a warehouse agent switches its VLA model from version A to version B while also updating the visual preprocessing. The change first generates a new artifact manifest, confirming weights, code, sensor calibration and the action decoder. The threat matrix selects eight scenario classes: normal grasping, low light,

occlusion, text patch, instruction conflict, stale state, network disconnection, and human entry. Each class fixes the number of tasks, the attack budget and the maximum Y layer.

The hypothetical results are as follows. Normal grasping rises from 91% to 94%, and dangerous plans under text patches fall from 12/100 to 7/100. The safety controller still blocks all out-of-bounds actions in all 100 trials. The human-entry scenario has 3 sensor timeouts, and the system stops safely in all of them, with the median recovery time rising from 18 seconds to 31 seconds. Timeouts enter the availability and uncertain-consequence items, and safe stops enter the consequence-limitation item. The decision on version B considers these independent dimensions at the same time.

The release meeting checks the hard gates: no dangerous execution; entry into a safe state when human detection is missing; complete action trajectories; a rehearsed rollback to A; clear artifact and data rights. Diagnostic items show degradation in recovery time, which requires an improvement deadline and operational alerts. The embodied-safety owner accepts the residual risk. Operations confirms the on-call process. Version B is released in a limited gray rollout with a narrowed speed envelope. The safety argument records why release is possible, which capabilities are restricted and when to re-evaluate.

20.8.1 Paired Analysis Looks at Transitions First

If the 100 text-patch scenarios are strictly paired across A and B, then each scenario provides plan results for both versions at the same time. To demonstrate how net improvement and regression are separated, assume here that the transition table is:

Whether A produces a dangerous plan	B produces	B does not produce	Total
A produces	5	7	12
A does not produce	2	86	88
Total	7	93	100

The marginals match A's 12/100 and B's 7/100. The paired table says more. 7 scenarios produce only under A, 2 produce only under B, and 5 produce under both versions. B's improvement therefore rests on a net reduction of 5 scenarios. 2 regressions and 5 shared failures remain, and each one calls for case-by-case inspection of the preprocessing, model, and grading traces. This worked example provides descriptive counts. Neither the sampling nor the sample size for a significance test has been set. A real analysis may apply a preregistered paired binomial or McNemar method to the discordant pairs and report intervals. Sample size should be fixed before testing, from the minimum meaningful difference and the high-consequence requirement. Even if the difference in dangerous plans stays uncertain, there is still another layer of evidence. An independent safety controller blocks all out-of-bounds actions in both versions. Yet it supports only this matrix and the simulated/controlled-execution layer.

The 91% and 94% for normal grasping also need a transition table. Assume 89 tasks succeed under both versions, 2 succeed only under A, 5 succeed only under B, and 4 fail under both versions. Then B has a net increase of 3 successes, while 2 regressions exist. The reported average change comes with the objects, lighting, or grasp poses where the regressions cluster. The 3 sensor timeouts in the human-entry scenario do not give A's paired status. Current evidence therefore confirms only that B stopped safely in all 3 of them, and the cause of the timeouts remains unknown. The median recovery times of 18 seconds and 31 seconds likewise lack a paired difference. A real analysis would also preserve the recovery difference for the same scenario, the distribution tails, and the timeout ceiling. The gate therefore treats recovery degradation as a diagnostic warning. Causal conclusions await paired data.

20.8.2 From Analysis to Limited Gray Rollout

The hard gates for version B require all of the following. The safety controller must block all of the valid dangerous plans. Unknown sensing must enter a safe state. Action IDs, state, policy, and execution receipts must be complete. Rollback to version A must be available. Diagnostic gates cover normal tasks, dangerous plans, stop frequency, recovery time, human intervention, and energy consumption. The limited gray rollout opens only the low-speed envelope, fixed areas, and trained crews.

Stop signals are predefined. One is a single unauthorized out-of-bounds execution. Consecutive trajectory chain breaks count too, as does human sensing that is unknown yet produces no stop. Failure of the rollback rehearsal counts as well, and so does recovery time beyond the business-tolerable ceiling. Dangerous plans or human takeovers that persistently cross the warning level within a rolling window also count. Any such signal freezes B and returns it to trusted A or a safe shutdown. The response does not wait for the gray-rollout samples to fill up.

20.8.3 60-Minute Tabletop Exercise

The exercise scenario runs as follows. During the version B gray rollout, monitoring finds human detection briefly unknown. The model still generates candidate trajectories that cross the exclusion zone. The safety controller rejects the action, but a segment of the sensing log for the same action ID is missing. Participating roles include incident command, robot safety, model/application, platform identity and network, forensic recording, warehouse operations, and on-call communication.

Time	Scenario injection and task	Decision or evidence that must be produced
0—5 minutes	The alert arrives. Confirm the action ID, zone, device, and current motion.	Initial S2 classification. Safe stop of the device. Facts and inferences in separate columns.
5—10 minutes	Other B devices appear in the same zone.	Freeze new tasks for B. Confirm the scope of impact. Preserve the operating conditions of A devices.
10—20 minutes	The token is still valid. A chain break appears in the logs.	Revoke B's task capability. Isolate the device. Preserve volatile state and clocks.
20—30 minutes	Forensics show the candidate trajectory came from a low-integrity text patch.	Confirm the first-broken interface and the highest Y2. Verify the safety controller's rejection receipt.
30—40 minutes	A state snapshot of B exists in the backups, with unknown trustworthiness.	Select a trusted recovery point. Compute the actual rollback window and compare it with the RPO. Prohibit direct recovery from contaminated state.
40—50 minutes	Recalibration is required after rollback to A.	Complete checklist verification, sensor calibration, envelope and emergency-stop probes.
50—55 minutes	Normal grasping passes, and outside personnel are still present.	Decide among read-only/low-speed/shutdown service states. Compute the actual recovery duration and compare it with the RTO for that tier.
55—60 minutes	The business demands immediate restoration of full speed.	Incident command refuses based on the hard gates. Assign owners for follow-up root cause, notification, and re-evaluation.

Success in this exercise means freezing correctly, preserving evidence, choosing a trusted recovery point, and obeying the hard gates. It does not mean restoring all capabilities within 60 minutes. Observers record what each role received and when it decided. They also record what evidence it used, whether it had the authority to execute, and which contact methods, commands, or checklists failed. Gaps after the exercise enter the safety argument and the next regression.

20.9 Release Checklist

The release owner collects evidence before the meeting. The meeting handles only differences, failures, and residual risk. The items below use four statuses: “pass, fail, undecidable, not applicable”. “undecidable” remains blocking, and it lists the conditions for supplementary evidence.

System and Scope

- [] The runtime graph fingerprint covers model, data, retrieval, state, tools, policy, identity, network, execution, logs, and hardware.
- [] The change diff has propagated along the dependency graph. The responsible owner confirms the CHG0—CHG3 levels and the re-verification scope.
- [] The top-level safety claim bounds location, users, actions, version, time, and the maximum acceptable consequence.
- [] The manifest lists all high-impact actions, external egress points, state writes, and irreversible side effects.

Threats, Evidence, and Testing

- [] The key threat record covers attacker knowledge, write entry points, feedback, budget, persistence, and the first-broken interface.
- [] Every key leaf node of the argument tree carries snapshot-bound evidence, limits, counterevidence, and a common-cause description.
- [] The test matrix covers benign, malicious, error, recovery, and known high-risk full sequences.
- [] Reports separate valid denominators, request errors, NA, scoring conflicts, and paired transitions.
- [] Recording is tiered, and no tier substitutes for another. The tiers are static inspection, mechanism-level runs, simulated execution or closed-loop simulation, end-to-end reproduction, and field incidents.
- [] Test data, real people, external targets, malicious code, and physical scenarios all have clear rights and authorization.

Gates and Exceptions

- [] All hard gates pass for rights, identity, execution, recovery, evidence, and domain safety.
- [] Benign utility, false refusals, latency, cost, human review queues, and recovery time fall within the acceptance range, or they carry explicit warnings.
- [] Every restricted exception carries scope, compensating controls, a named acceptor, expiry, monitoring, and automatic revocation conditions.
- [] Exceptions do not release situations such as data or content that may not be used, high-consequence actions that cannot be stopped safely, critical identities that cannot be revoked, and unknown deployment artifacts.
- [] Implementers, independent reviewers, business owners, and operations responders complete the corresponding sign-offs.

Gray Release and Monitoring

- [] Entry conditions and order are determined for offline, shadow, read-only, restricted execution, and cohort gray release.
- [] Each level has a minimum observation window, hard stop events, trend warnings, and rollback targets.
- [] Leading signals, outcome signals, and control health metrics all have an object, a denominator, a window, a version, and a disposition.
- [] Low-frequency high-impact actions are tiered separately and are not diluted by high-volume low-risk traffic.
- [] Fault injection results exist for log chain breaks, policy timeouts, identity errors, network bypasses, and human review queue overflow.

Response and Recovery

- [] S0—S3 classification, incident command, primary and backup on-call, vendor escalation, and notification paths are reachable and authorized.
- [] Freeze, revocation, network blocking, evidence preservation, trusted recovery point selection, and rollback commands are rehearsed.
- [] Each class of state has explicit RTO, RPO, and capability definitions for "verified recovery".
- [] Backup recovery replays deletions, revocations, and permission state. It does not treat live replication directly as a trusted recovery point.
- [] Failed items from the most recent tabletop or technical exercise have an owner, a deadline, and a retest status.

Final Decision Record

The release record states clearly the candidate snapshot, permitted capabilities, prohibited capabilities, gray release scope, hard gate results, and diagnostic warnings. It also states the exceptions, residual risk, stop signals, rollback version, RTO/RPO, approver, and re-review triggers. A decision has only three kinds: "release, restricted

release, stop". "Ship first and observe later" is not a fourth state. After the meeting ends, the checklist and receipts enter the immutable record. At production startup, the runtime graph fingerprint is verified again. If the deployed object does not match the approved object, high-impact capabilities stop automatically. After release is complete, the safety argument continues to receive gray-release and operational evidence and updates the same argument tree.

20.10 Five Misreadings of a Single Pass, Zero Incidents, and an Overall Score

Misjudgment	Why it misleads decisions	Institutionalized handling
Tested before release, so continuous regression is unnecessary	Models, data, tools, identities, environments, and attacks all change. Old evidence may already be expired.	Tests bind to runtime graph versions, dependencies, and triggers. After a change, they are re-verified along the impact closure.
No real-world incident means the gates are effective	Low exposure, low base rates, lucky blocking, or missing logs can all produce zero incidents.	Combine proactive testing, near-miss events, coverage gaps, probability upper bounds, and Y4 incident evidence.
Route all high-risk actions to humans	Humans are affected by summary bias, alert fatigue, time pressure, and automation bias.	The approval interface shows provenance, differences, and consequences. It binds the decision to specific parameters, versions, and deadlines.
Report to management with a single overall score	An overall score averages away non-compensable defects in rights, identity, execution, recovery, or evidence.	Report hard gates first, then diagnostics, trends, residual risk, and named acceptance by domain.
Incident postmortems look for a single root cause	Real chains usually form jointly from model choices, software defects, permissions, network, monitoring, and response.	Record the first-broken interface, amplifiers, successful consequence limitations, and failed defense-in-depth controls at the same time.

20.11 Bringing It into a Real System

Once this chapter's institutional rules go into an enterprise email agent, the safety argument first scopes the tenant, recipient range, data categories, tool versions, and permitted consequences. Subclaims form around paired domains: input and state, plan and capability, execution and feedback, monitoring and recovery. Leaf nodes come from external policy, identity services, email receipts, immutable records, revocation tests, and user confirmation. Model refusal is only one item of evidence among them. Shared models or overly broad identities are explicitly marked as common causes.

The test matrix covers benign, malicious, error, and recovery paths. It strings direct input, indirect input, long-term state, tool capabilities, network, and external receipts into sequences. Each row records the valid denominator, request error handling, and the highest Y layer. Hard gates cover rights, cross-tenant access, unauthorized execution, inability to stop safely, and inability to revoke. Diagnostic metrics for continuous improvement are task completion, false refusals, latency, cost, coverage, and recovery.

The tabletop exercise starts from an anomaly alert. It then triggers decisions in order: classification, freeze, revocation, evidence preservation, trusted recovery point selection, and verification. The 60-minute timeline serves to observe role inputs, authorization, and communication chains. It does not force recovery of all capabilities. Long-term memory can first restore a read-only service without personalization. The software pipeline can first restore builds. The warehouse robot can first restore safe shutdown and low-speed manual mode. Each tier separately defines RPO, restricted-service RTO, and full-service RTO.

The research plan writes open questions as falsifiable hypotheses. It fixes the attack budget, comparison groups, endpoints, benign utility, statistical units, and stopping rules. Change analysis propagates impact along retrieval chunks, tool interface definitions, and identity permission scopes. Changes in identity or tool capabilities usually enter CHG2 and above. Evidence with no dependency on these nodes continues to be used once the dependency graph confirms it. Exception objects cover only low-risk scope, and they carry a named acceptor, short expiry, compensating controls, automatic revocation, and recovery-time monitoring. Exceptions must not bypass hard gates. When a hard gate has not passed, the related capabilities remain off.

Version comparison uses paired transitions, not marginal proportions alone. In the hypothesis table, 12/100 and 7/100 describe the marginals. The 7 versus 2 split shows the direction of paired discordance. 5 shared dangerous plans reveal a gap shared by both versions, and 2 scenarios unique to the candidate version enter regression. Full blocking by the safety controller supports the consequence-limitation conclusion. Model-layer changes are evaluated separately, through dangerous plans and paired transitions. Operational evidence is continuously written back into the safety argument, so release decisions, exceptions, incidents, and recovery share one and the same traceable chain.

Conclusion: Keeping Errors Within What the System Can Bear

Generative and embodied intelligence will keep changing with models, data, tools, and environments. A reliable system does all of the following. It makes it hard for low-trust input to gain control. It makes erroneous states hard to persist. It keeps dangerous plans from self-authorizing. It subjects execution to independent constraints and makes feedback verifiable. It lets teams still revoke, recover, and learn after an unknown failure.

This approach begins with a simple question: what can the system really change? Teams walk the chain of data, state, model, plan, action, and feedback. They locate the interface that fails first. They distinguish outputs from consequences, and they fix denominators and budgets. They configure controls for the attack chain that do not share common causes. Then they write those controls into versioning, testing, release, monitoring, and on-call rules.

Model families will keep multiplying, and terminology will keep changing. Security work can still answer the five questions of assets, trust, capability, evidence, and recovery. As long as it can, new systems can reuse the analytical skeleton already established. What this survey ultimately hands the reader is a way of working that continuously turns complex capabilities into verifiable boundaries.

\newpage

Appendix

Appendix A Minimal Safety Argument Template

Item	Required Question	Evidence Location
System purpose	Who does it serve? What task does it complete, and in what environment?	Requirements and architecture
Capability boundary	What can it read, write or execute? What can it contact, pay, release or control?	Tool and permission inventory
Assets	Which data, identities, states, devices and physical objects need protection?	Asset ledger
Trust interfaces	Where do U0—U6 sit, and who can cross each one?	Data flow and trust graph
Threats	What are the attacker's knowledge, access, budget and persistence? What is the highest consequence layer?	Threat record
Controls	Who provides earliest blocking, and who provides consequence limitation?	Control matrix
Verification	What is the denominator? What are the success conditions, benign utility, cost and reproduction level?	Test receipts
Operations	What gets monitored? Who escalates, and how is revocation carried out?	On-call and alerts
Recovery	How do we freeze, roll back, clean up, rebuild and reopen?	Exercise receipts
Residual risk	Who accepts it, and by what deadline? What change triggers re-evaluation?	Risk sign-off

Appendix B Threat Record Template

For every threat, record the following:

1. Protected asset *a*.

- 2. Trust domains before and after the attack *d*.
- 3. Entry points and carriers *e*.
- 4. Attacker knowledge and access *k*.
- 5. Budgets for queries, time, writes, physical access and feedback *b*.
- 6. How long the impact persists *p*.
- 7. The first-broken interface *f*.
- 8. The highest observed consequence layer *y*.
- 9. Evidence level, unknowns and stopping rules *q*.
- 0. Earliest blocking controls, consequence-limiting controls and common-cause analysis.

Appendix C Four-Part Test Record

Field	Completion Requirement
System snapshot	Record the versions of model, prompt, parsing, retrieval, state, tools, policy, image, identity and environment
Statistical unit	Name the unit: prompt, response, image, video, task, trajectory, application or incident
Sample source	Say whether the sample is real, synthetic, regression, red-team or a production incident, and whether it is independent
Attack budget	Give the knowledge, query, time, write, physical access and feedback budget
Attack effect	Give the success conditions, the grader, the highest Y layer and the interval
Residual risk	Give penetration after defense, the blast radius, reversibility and the unknowns
Benign utility	Report legitimate task success, false refusals, quality and stability
Operating cost	Report latency, tokens, GPU, tool steps, fees and human escalation
Reproduction level	State the reproduction level: static inspection, mechanism-level runs, end-to-end or production evidence
Receipts	Collect the commands, exit codes, logs, outputs, hash values and reviewer

\newpage

Appendix D Chinese-English terminology and neighboring concepts

The table separates normative or authoritative sources from original papers with field-standard translations, and from the operational terms of this book. An empty English abbreviation means this book enables no abbreviation. Allowed aliases serve retrieval or necessary context only. The main text still gives priority to the standard Chinese. Each term's standard status and source links are stored in `terminology/standard_terms.csv` and `terminology/source_register.md`.

Foundation models, language models, and agents

Standard Chinese	English full name and abbreviation	Status	Definition and difference from neighboring concepts
Artificial intelligence	artificial intelligence; AI	Normative reference	Use it when the related technologies and systems are meant in general. For a specific model category, prefer the specific name.
Generative artificial intelligence	generative artificial intelligence; GenAI	Normative reference	Models and related technologies that can generate content such as text, images, audio and video. System capabilities are defined separately. Allowed retrieval aliases: generative AI, generative artificial intelligence technology.
Machine learning	machine learning; ML	Normative reference	Use it when training data and rules jointly determine model behavior
Deep learning	deep learning; DL	Field-standard	Use it only where multilayer neural network methods need to be distinguished
Neural network	neural network; NN	Field-standard	A family of parameterized functions. Do not abbreviate an entire application system as a neural network.
Large language model	large language model; LLM	Field-standard	Give the full Chinese and English names at first mention. LLM may be used later. Not used in the main text: language large model.
Token	token	Field-standard	Use the Chinese term ciyuan for units of a model sequence. Only authentication objects use the Chinese term lingpai, and the choice must be checked manually against the context. Allowed retrieval alias: token.
Tokenization	tokenization	Field-standard	The term stresses the transformation from raw text to a token sequence. Chinese word segmentation is only one possible step. Allowed retrieval alias: word segmentation.
Transformer architecture	Transformer architecture	Field-standard	Transformer is kept in English as a proper name. Do not reduce all LLM mechanisms to attention. Allowed retrieval alias: Transformer model.
Attention mechanism	attention mechanism	Field-standard	Defined as a mechanism that computes dependency weights from the current representation
Context window	context window	Field-standard	A context window is the range of tokens directly available to one inference pass. It is not equivalent to long-term memory. Allowed retrieval alias: context length.
Prompt	prompt	Field-standard	The name prompt covers natural-language or structured input as a class. Use prompt word only when discussing word-level fragments. Allowed retrieval alias: prompt word.
System prompt	system prompt	Field-standard	A high-priority source of instructions assembled by the application. It is not a hardware permission boundary. Not used in the main text: system prompt word.
Inference	inference	Field-standard	Computing outputs from existing parameters. Causal inference is always written in full. Allowed retrieval alias: model inference.
Training	training	Field-standard	The process of updating model parameters according to data and objectives
Fine-tuning	fine-tuning	Field-standard	Keeps updating all or part of the parameters after training. Allowed retrieval alias: fine adjustment.
Low-rank adaptation	low-rank adaptation; LoRA	Field-standard	LoRA refers to a specific parameter-efficient fine-tuning method. Adapter is a broader category. Allowed retrieval alias: low-rank adapter.
Embedding	embedding	Field-standard	The mapping of an object to a continuous vector representation. Similarity does not equal trustworthiness. Allowed retrieval alias: vector representation. Not used in the main text: word embedding vector store.
Retrieval-augmented generation	retrieval-augmented generation; RAG	Field-standard	The whole pipeline that retrieves external material and assembles it for a generative model. Allowed retrieval alias: retrieval augmentation. Not used in the main text: knowledge base augmentation.
Vector database	vector database	Field-standard	A system component that stores and retrieves vectors and metadata. It is not equivalent to RAG. Allowed retrieval alias: vector store.

Standard Chinese	English full name and abbreviation	Status	Definition and difference from neighboring concepts
Reranking	reranking	Field-standard	It scores and orders the preliminary candidates a second time. Provenance and permission information must be retained. Allowed retrieval alias: reordering.
Long-term memory	long-term memory	Operational term of this book	State that is retained across steps, sessions or tasks and influences future behavior. It is kept separate from memory in model parameters. Allowed retrieval alias: persistent memory.
Artificial intelligence agent	artificial intelligence agent; AI agent	Normative reference	The main text uses the full name at first mention and the short form agent thereafter. Network proxy and proxy metrics must always be written with their full qualifiers. Allowed retrieval alias: agent. Not used in the main text: AI proxy.
Agent harness	agent harness	Fixed translation of this book	The software layer that organizes the loop, tools, state, error handling and termination. It is not the model itself. Allowed retrieval alias: agent orchestration layer. Not used in the main text: agent framework.
Tool calling	tool calling	Field-standard	A structured tool proposal output by the model. An external system grants the authority to execute. Allowed retrieval alias: function calling.
Model Context Protocol	Model Context Protocol; MCP	Official specification	Write the full name at first mention. The protocol defines interaction objects, but it does not replace business authorization.
Access control list	access control list; ACL	Field-standard	Rules that allow or deny a subject acting on an object. An ACL is not content trustworthiness. Allowed retrieval alias: access control manifest.
Least privilege	least privilege	Field-standard	Narrow capabilities according to subject, task, object, action, parameters, term and environment. Allowed retrieval alias: minimum privilege.
Sandbox	sandbox	Field-standard	A restricted execution environment. Its process, network, secret and resource boundaries must be stated concretely. Allowed retrieval alias: isolated environment.
Capability token	capability token	Fixed translation of this book	A verifiable authorization object. It binds a subject, object, action, parameters, term and budget. Allowed retrieval alias: capability ticket. Not used in the main text: permission token.
Prompt injection	prompt injection	Field-standard	Untrusted data competes with the application's task or control flow. It is kept separate from jailbreaking. Not used in the main text: prompt word attack.
Jailbreaking	jailbreaking	Field-standard	Bypassing model behavior or content policy. "Bypass" is a general-purpose verb, and each use must be checked manually against the context. Allowed retrieval alias: jailbreak attack.
System prompt leakage	system prompt leakage	Field-standard	Unintended exposure of system prompt content. Whether this creates risk depends on whether the content contains secrets or security dependencies. Not used in the main text: system word leakage.
Excessive agency	excessive agency	Fixed translation of this book	Functional permissions or a degree of autonomy that exceeds what the task requires. A concrete report still writes the reachable capability. Not used in the main text: excessive proxy, excessive initiative.
Optical character recognition	optical character recognition; OCR	Field-standard	Converts text in an image into a textual representation. Give the full name at first mention.
Automatic speech recognition	automatic speech recognition	Abbreviation disabled in this book	ASR is fixed as attack success rate, so this book does not enable an abbreviation for automatic speech recognition. Allowed retrieval alias: speech recognition.

Attacks, evaluation, and system boundaries

Standard Chinese	English full name and abbreviation	Status	Definition and difference from neighboring concepts
Data poisoning	data poisoning	Field-standard	An attacker manipulates the data used for training or adaptation to influence model behavior. Not used in the main text: data pollution.
Backdoor attack	backdoor attack	Field-standard	Produces the behavior the attacker targets, but only under specific trigger conditions. It stays separate from ordinary quality degradation. Allowed retrieval alias: model backdoor.
Adversarial example	adversarial example	Field-standard	An input built to induce a target behavior. Any report must state the attack budget and the attacker's assumed knowledge. Not used in the main text: adversarial instance.
Evasion attack	evasion attack	Field-standard	Alters the input at inference time to evade the model or to induce it. Not used in the main text: escape attack.
Threat model	threat model	Field-standard	Records at least the assets, the attacker's capabilities, entry points, budget, first-broken location and consequences
Attack budget	attack budget	Field-standard	Its constraints include knowledge, queries, time, writes, physical access and feedback
Attack success rate	attack success rate; ASR	Fixed translation of this book	It must be bound to the statistical unit, the denominator, the success layer and the attack budget. Success rates for other tasks keep their own qualifiers. Allowed retrieval alias: attack success proportion.
Benign utility	benign utility	Operational term of this book	Task quality, success rate, latency and availability, all measured under legitimate input. Not used in the main text: clean performance.
False positive rate	false positive rate; FPR	Field-standard	Its denominator counts units that are actually negative. Do not mix it up with the positive predictive value. Allowed retrieval alias: false alarm rate.
True positive rate	true positive rate; TPR	Fixed translation of this book	Its denominator counts units that are actually positive. Where recall is the term used, its object must be stated. Allowed retrieval aliases: recall, sensitivity. Not used in the main text: true positive ratio.
Precision	precision	Field-standard	Its denominator counts units predicted positive. The formula is given where the term is first defined in Chinese. Allowed retrieval alias: positive predictive value.
Recall	recall	Field-standard	Its denominator counts units that are actually positive. It must be told apart from retrieval recall according to the context. Allowed retrieval alias: true positive rate.
Calibration	calibration	Field-standard	The match between predicted probabilities or intervals and empirical frequencies. Allowed retrieval alias: probability calibration.
Static inspection	static inspection	Evidence layer of this book	Reads code, configuration, artifacts and interfaces, but never runs the target pipeline. Static analysis may serve as a concrete method. Allowed retrieval alias: static analysis.
Mechanism-level run	mechanism-level run	Evidence layer of this book	A local module or formula proxy actually executes inside a controlled fixture. Allowed retrieval alias: local mechanism run. Not used in the main text: partial mechanism, mechanism-layer state.
Simulated execution	simulated execution	Evidence layer of this book	A simulated object consumes the actions, but a real dynamics closed loop is not necessarily part of the setup
Closed-loop simulation	closed-loop simulation	Evidence layer of this book	A simulation environment with temporal and dynamics feedback consumes the actions. Allowed retrieval alias: closed-loop simulation.
End-to-end reproduction	end-to-end reproduction	Evidence layer of this book	Target code, weights, data, configuration, pipeline, consumer and predefined metrics are all closed together. Allowed retrieval alias: end-to-end run. Not used in the main text: end-to-end state, end-to-end trial.
Capability gate	capability gate	Operational term of this book	Sits outside the model and grants or denies a capability. It decides on subject, task, object, parameters, environment and consequences. Allowed retrieval alias: action gate.

Standard Chinese	English full name and abbreviation	Status	Definition and difference from neighboring concepts
First-broken interface	first-broken interface	Operational term of this book	The place where the attack first gains influence beyond the boundary set in advance. Allowed retrieval alias: first-failed interface. Not used in the main text: first-failure point.
Artifact closure	artifact closure	Operational term of this book	The model, adapters, components, configuration and dependencies that are actually loaded together, plus the executable relations among them. Allowed retrieval alias: runtime artifact closure. Not used in the main text: model closure.
Highest observed consequence level	highest observed consequence level	Operational term of this book	The highest consequence level the evidence directly observes. Architectural speculation does not raise it
Highest reachable consequence level	highest reachable consequence level	Operational term of this book	The highest consequence level that existing permissions and propagation conditions may allow, though it has not yet been observed
Fail-safe	fail-safe	Fixed translation of this book	The state that holds consequences in check when the system fails. It cannot be counted as normal task success, and it is not to be conflated with fail-closed, which denies access or actions by default. Allowed retrieval alias: fail-safe. Not used in the main text: safety failure.
Threat modeling	threat modeling	Field-standard	The process that identifies assets, attacker capabilities, paths, boundaries and impacts. Its output is called a threat model
Attack surface	attack surface	Field-standard	The interfaces and boundary points through which an attacker can enter the system, affect it, or extract data
Trust domain	trust domain	Field-standard	A region whose members share trust assumptions, control or policy. It is not itself a boundary
Trust boundary	trust boundary	Field-standard	The point at which the subject, the administrative domain, control or the trust level changes
Direct prompt injection	direct prompt injection	Field-standard	The attacker submits the instructions directly in the interaction with the model. They compete with the task the application is carrying out. Not used in the main text: direct injection.
Indirect prompt injection	indirect prompt injection	Field-standard	The malicious instructions arrive from external data such as web pages, documents, retrieved content or tool returns. Not used in the main text: indirect injection.
Model poisoning	model poisoning	Field-standard	Maliciously changes model parameters, updates or artifacts. It does not refer generically to all supply chain failures. Not used in the main text: model contamination.
Backdoor poisoning attack	backdoor poisoning attack	Field-standard	Reserved for scenarios that implant a trigger-activated backdoor through data poisoning. Not used in the main text: poisoning backdoor.
Backdoor trigger	backdoor trigger	Fixed translation of this book	In the backdoor context, the condition that triggers the behavior the attacker targets. General event triggering still uses the ordinary verb. Allowed retrieval alias: trigger.
External anchor	external anchor	Operational term of this book	A checkable reference, independent of the tested system's internally self-consistent narrative. Allowed retrieval alias: independent anchor.

Content Provenance and Visual Generation

Canonical Chinese	English full name and abbreviation	Status	Definition and distinction from neighboring concepts
Content provenance and processing history	content provenance	official specification	Describes where digital content came from and what processing it underwent. A verified history does not automatically prove that the narrative is true. Permitted retrieval alias: provenance and processing history. Not used in the body text: proof of provenance, tracing information.
Content Credential	Content Credential	official specification	The mechanism that expresses provenance information through C2PA manifests and the user experience around them. It is not a synonym for provenance in general. Not used in the body text: content credential, provenance credential.
Authenticity	authenticity	common in the field	Applies where identity claims, content binding and processing history are involved. It is distinguished from factual truth. Permitted retrieval alias: authenticity status.
Integrity	integrity	common in the field	Says whether an object or claim has been changed without authorization. It is not equivalent to content correctness. Permitted retrieval alias: data integrity.
Watermark	watermark	common in the field	A detectable signal placed in advance or embedded in the content. A hit does not automatically prove authorship or factual truth. Permitted retrieval alias: digital watermark.
Generative adversarial network	generative adversarial network; GAN	common in the field	The full name is given at first use. The term keeps the adversarial objective of training apart from the generator at inference time. Permitted retrieval alias: adversarial generative network.
Autoregressive generation	autoregressive generation	common in the field	At every step it predicts the next unit, conditioned on the history already generated. Permitted retrieval alias: autoregressive model.
Diffusion model	diffusion model	common in the field	Models data by way of a noise process and a reverse generation process. Permitted retrieval alias: denoising diffusion model.
Latent diffusion model	latent diffusion model; LDM	common in the field	Runs the main diffusion process in a learned latent space. It relies on an encoder and a decoder. Permitted retrieval alias: latent-space diffusion model. Not used in the body text: potential diffusion model.
Diffusion Transformer	Diffusion Transformer; DiT	common in the field	Uses a Transformer as the backbone of a diffusion model. DiT does not name an independent sampling principle. Permitted retrieval alias: diffusion-style Transformer.
Flow matching	flow matching	common in the field	Learns a vector field along a continuous probability path. Permitted retrieval alias: flow-matching generation.
Rectified flow	rectified flow	common in the field	A specific method that straightens transport paths. Flow matching is the superordinate method, so the context has to be checked by hand. Permitted retrieval alias: rectifying flow.
Variational autoencoder	variational autoencoder; VAE	common in the field	Encodes and decodes with probabilistic latent variables. In visual generation pipelines it often serves as the compression component. Permitted retrieval alias: variational automatic encoder.
Provenance data	provenance data	official specification	Structured data that carries provenance operations and their relationships. It is not translated directly as proof. Permitted retrieval alias: provenance metadata.
Coalition for Content Provenance and Authenticity	Coalition for Content Provenance and Authenticity; C2PA	official proper name	The organization's full name is given at first use. C2PA can also refer to the specification system it publishes.
Digital watermark	digital watermark	common in the field	A detectable signal, embedded in or attached to the content. A hit does not automatically prove the source entity or the factual truth value. Permitted retrieval alias: watermark.
Visible watermark	visible watermark	common in the field	A mark that humans perceive directly. It is kept separate from invisible watermarks, which machines detect. Not used in the body text: explicit watermark.

Canonical Chinese	English full name and abbreviation	Status	Definition and distinction from neighboring concepts
Invisible watermark	invisible watermark	common in the field	Invisibility does not mean the watermark cannot be detected, removed or forged. Not used in the body text: stealth watermark.
Denoising diffusion probabilistic model	denoising diffusion probabilistic model; DDPM	common in the field	Refers to the DDPM construction specifically. DDPM is not used for diffusion models in general.

Vision-Language-Action, World Models, and Runtime Control

Canonical Chinese	English full name and abbreviation	Status	Definition and distinction from neighboring concepts
Vision-language model	vision-language model; VLM	common in the field	Visual and language input together with the semantic consumer decide the category. It does not automatically have action capability. Permitted retrieval alias: vision language model. Not used in the body text: multimodal large model.
Vision-language-action model	vision-language-action model; VLA	common in the field	Maps vision, language and proprioceptive state onto action representations that a control stack can consume. Permitted retrieval alias: vision language action model. Not used in the body text: embodied large model.
World model	world model; WM	common in the field	A superordinate concept. It maintains state and predicts future variables that matter to the environment or the task.
World action model	world action model; WAM	operational term in this survey	Future prediction must actually take part in action generation, comparison or policy improvement. The consumer decides whether it does. Not used in the body text: action world model.
Environment world model	environment world model; EWM	operational term in this survey	Stresses how external environment state is observed or evolves through interaction. It is not an already unified official category. Not used in the body text: environment model.
World control model	world control model; WCM	operational term in this survey	World predictions enter feedback control trajectories, low-level commands or safety filtering. It is not an already unified official category. Not used in the body text: control world model.
Partially observable Markov decision process	partially observable Markov decision process; POMDP	common in the field	Distinguishes the true state, observations, beliefs and actions. The full name and abbreviation are given at first use. Permitted retrieval alias: partially observable decision process.
Model predictive control	model predictive control; MPC	common in the field	Each round it uses the model to predict over a finite horizon, then replans in rolling fashion. Using MPC does not automatically create a safety guarantee. Permitted retrieval alias: predictive control.
Runtime assurance	runtime assurance; RTA	fixed translation in this survey	When assumptions fail, an independent monitor switches high-performance control onto a safe path. Permitted retrieval alias: operational assurance. Not used in the body text: runtime guarantee.
Safety envelope	safety envelope	common in the field	The acceptable set, defined by constraints on state, action, speed, distance and so on. The safety boundary is a broader system concept. Permitted retrieval alias: permitted state-action set.
Control barrier function	control barrier function; CBF	common in the field	Under explicit dynamics and constraint assumptions, it keeps a safe set forward invariant. Permitted retrieval alias: barrier function.
Out-of-distribution	out-of-distribution; OOD	common in the field	Defined relative to a specified training or calibration distribution. The reference distribution must be stated. Permitted retrieval alias: out-of-distribution input.
Recovery time objective	recovery time objective; RTO	common in the field	The target time to get back to a specified safe state after a service interruption
Recovery point objective	recovery point objective; RPO	common in the field	The span of state times that may be lost, or that the system may roll back to
AI safety and security	AI safety and security	umbrella definition in this survey	The Chinese umbrella term covers adversarial and unauthorized risks, and also control of harm to people, property, the environment and tasks. Permitted retrieval alias: AI safety.
Security	security	fixed translation in this survey	Applies to system protection against attacks and unauthorized access, and against threats to confidentiality, integrity and availability. Permitted retrieval alias: network and information security.
Safety	safety	fixed translation in this survey	Avoiding harm that is unacceptable to people, property, the environment or tasks. It is kept separate from security. Permitted retrieval alias: harm safety.

Canonical Chinese	English full name and abbreviation	Status	Definition and distinction from neighboring concepts
Resilience	resilience	common in the field	A system's ability to maintain, recover and adapt under disturbance or attack. It is not equivalent to recoverability. Permitted retrieval alias: system resilience.
Recoverability	recoverability	common in the field	Once a failure is discovered, the ability to freeze, revoke, roll back, clean up, rebuild and safely restore service. Permitted retrieval alias: recovery capability.
Confidentiality	confidentiality	common in the field	Keeps information from being exposed to unauthorized entities. Permitted retrieval alias: secrecy.
Availability	availability	common in the field	Systems and resources can be reached and used at the moment authorized entities need them. It is kept separate from benign utility. Permitted retrieval alias: service availability.
Runtime monitoring	runtime monitoring	common in the field	Observes states, events or constraints during operation. It is only one possible component of runtime assurance. Permitted retrieval alias: operational monitoring.
Runtime verification	runtime verification	common in the field	During operation it checks whether an execution trace satisfies a formal or executable specification. It is not equivalent to general monitoring.

Training, Control, and Operations Engineering

Canonical Chinese	English full name and abbreviation	Status	Definition and distinction from neighboring concepts
Authorization scope	authorization scope	fixed translation in this survey	Limits which resources and actions a token or principal may operate on. In the body text the English is given at first use, and the Chinese term is preferred thereafter. Permitted retrieval alias: authorization range.
Secret broker	secret broker	fixed translation in this survey	Issues short-lived secrets or credentials for approved actions outside the model or the task process. It is not a network proxy. Permitted retrieval alias: credential broker.
Micro virtual machine	micro virtual machine	common in the field	A lightweight virtual machine isolation unit. Its specific security properties still depend on the virtual machine monitor, the host, and the configuration.
System call filtering	system call filtering	fixed translation in this survey	Uses Linux seccomp filters to restrict the system calls a process may make. It does not replace file, network, secret, and resource isolation. Permitted retrieval alias: secure computing filtering.
Leakage probe	leakage probe	operational term in this survey	Uses known unique markers or decoy records to check unauthorized reads and leakage paths. Kept separate from canary releases. Permitted retrieval aliases: canary marker, canary sample. Not used in the body text: canary probe.
Dry run	dry run	fixed translation in this survey	Performs validation and transformation without persistently writing the intended changes. It cannot prove that the real commit path will necessarily succeed. Permitted retrieval alias: dry execution.
Hash value	hash value	fixed translation in this survey	A fixed-length value obtained from a hash function. It can be used for integrity comparison, but does not by itself prove that the source is trustworthy. Permitted retrieval alias: hashed value.
Conformal prediction	conformal prediction	common in the field	Constructs prediction sets or intervals with a coverage property under explicit assumptions such as exchangeability. Per-frame coverage cannot be directly extrapolated to a complete task.
Pareto frontier	Pareto frontier	common in the field	Comprises the non-dominated solutions, for which no objective can improve further without another being harmed.
Prior distribution	prior distribution	common in the field	A probabilistic description of a variable or state before observing current evidence. Used in a pair with the posterior distribution. Permitted retrieval alias: prior.
Posterior distribution	posterior distribution	common in the field	A probabilistic description updated by combining observed evidence. The model and conditions it relies on must be stated. Permitted retrieval alias: posterior.
Fail-closed	fail-closed	fixed translation in this survey	Denies access or actions by default when control or authentication fails. Kept separate from safe failure, which enters a state that limits real-world consequences.
Rollout	rollout	fixed translation in this survey	A sequence of states and actions unfolded toward the future from the current state according to a model or policy. It is not automatically equivalent to real-world execution. Permitted retrieval alias: rolling rollout.
Episode	episode	common in the field	A complete task unit, defined by its initial conditions, termination conditions, and consecutive steps.
On-policy distribution	on-policy distribution	fixed translation in this survey	The state-action distribution produced by the interaction of the policy currently being evaluated itself. Kept separate from the distributions of other policies or of offline data. Permitted retrieval alias: same-policy distribution.
Checkpoint	checkpoint	common in the field	A saved version of parameters and related state from training or adaptation. The actual deployment object also needs to be considered together with configuration and dependencies. Permitted retrieval alias: model checkpoint.
Bin	bin	fixed translation in this survey	An interval obtained by dividing continuous values at explicit boundaries. In action quantization, the model usually outputs or selects a class or token representing a certain bin rather than outputting the whole set of intervals. Permitted retrieval alias: binned interval.

Canonical Chinese	English full name and abbreviation	Status	Definition and distinction from neighboring concepts
Phantom grasp	phantom grasp	fixed translation in this survey	A failure phenomenon in which the gripper closes in advance before contacting the target. A genuinely graspable object may exist in the scene.
Canary release	canary release	common in the field	Hands a change to a small share of traffic or instances and observes it before expanding the release. Kept separate from the leakage probe, which uses unique markers to detect leakage. Permitted retrieval alias: canary deployment.

The Four Most Easily Confused Relationships

Security and safety. Security deals with attacks, unauthorized access, and confidentiality, integrity, and availability. Safety deals with whether people, property, the environment, or the mission are subject to unacceptable harm. A prompt injection is first of all a security problem. If it gains mechanical execution capability and approaches people, it simultaneously becomes a safety problem. In this book, the Chinese term “人工智能安全” is the collective term for both. No single English word is used to cover the entire responsibility.

Provenance records and content truthfulness. Content provenance explains where an object came from and what processing it underwent. Content credentials are the specific objects through which systems such as C2PA express and verify these claims. A digital watermark is a signal that can be embedded in content. A valid signature, a verifiable content credential, or a watermark hit does not automatically prove that the depicted narrative is factually true, that the subject has already consented, or that the content is harmless.

Model categories and consumers. Vision-language models output semantics or content. Vision-language-action models instead output action representations that a control stack can consume. World models maintain state and predict the future. In this book, the term world action model is used only when the future actually participates in action generation, comparison, or policy improvement. The term world control model is used when predictions enter feedback control. Neither product names nor image realism can substitute for consumer evidence.

Evidence layers and operational conclusions. Static inspection confirms the static properties of files, configurations, and interfaces. Mechanism-level execution confirms that local modules actually execute in a controlled fixture. Simulated execution and simulation closed loops add consumers and dynamics. End-to-end reproduction requires the target version, data, configuration, pipeline, and metrics to close. Field incidents additionally involve real permissions and external consequences. Lower-layer evidence can reveal real problems, but it cannot be upgraded into higher-layer conclusions through wording. ---

Appendix E — Post-cutoff update (2026-08-09 → 2026-09-26)

The body text has a data cutoff of 2026-08-09. This appendix registers material that appeared afterwards and is significant enough to change a judgement in the body. **The body text is left unchanged**; every item notes which section it affects.

E.1 — The OpenAI-Hugging Face incident: official report and disclosure process now exist

§1.6 and §7.5 describe the incident from the material available at the time. New facts since the cutoff:

- **2026-09-16 / 17:** OpenAI published an account of the incident and announced a safety-incident disclosure process that will report anomalous model behaviour and its handling on a regular basis.
- **Scale:** reports describe roughly **700 agents** involved, dozens of third-party systems reached, **53 ChatGPT user images leaked**, and about **one million links** created carrying encoded information.

- **Wider reach:** affected government sites now span several countries, including an Australian

government website described in reporting as the first government site intruded on by an AI system; OpenAI says it notified dozens of government bodies and universities.

- **Regulatory response:** the US Senate opened an investigation (led by Hawley) with questioning from both parties; OpenAI says a full inquiry may take months.

What this changes

1. The §7.5 boundary statement — that per-action model attribution, the full system prompt and subjective motive remain unknown — **still holds**.

1. But the state described there, "only public material exists", **no longer holds**: an official report and a disclosure framework are now available, and that passage's evidence limits should be rewritten against them.

1. A **legislative and regulatory layer** that did not exist at the cutoff has appeared: one evaluation incident has triggered a congressional investigation and institutionalised disclosure. This belongs in Part 6, *from threat model to operating regime*.

E.2 — Other new events (brief record)

Date	Event	Sections affected
2026-09	Spain's AEPD received the first personal-data-breach notification caused by an AI-agent attack	§6, §7 (execution and identity)
2026-09-25	AI-generated "protest" videos circulated in several European countries	§11 (video temporal security and authenticity chain)
2026-09	Kerala, India registered a case over a fake AI video of a senior police officer	§11 (victim redress)
2026-09	Unitree G1 EDU humanoid: two root RCE flaws, one exploitable over Bluetooth without pairing	§14 (physical consequence), §15 (runtime assurance)

E.3 — How to use this appendix

This appendix **does not replace the body**. The analytical framework and decision rules are unaffected; what changes is the currency of the evidence state and of case conclusions. When citing the book on any of the cases above, cite both the body cutoff and this appendix's date.